You make **possible**

Let's Get Started with ACI Service Insertion

Minako Higuchi, Technical Marketing Engineer @DCBG

BRKACI-2486

Agenda

- ACI Contract security
- ACI L4-L7 service integration
  - Firewall Design Options
  - Load Balancer Design Options
  - Multi-Pod/Multi-Site Design Options
- Q&A

ACI Contract security

You make networking **possible**

Cisco ACI - Logical Network Provisioning

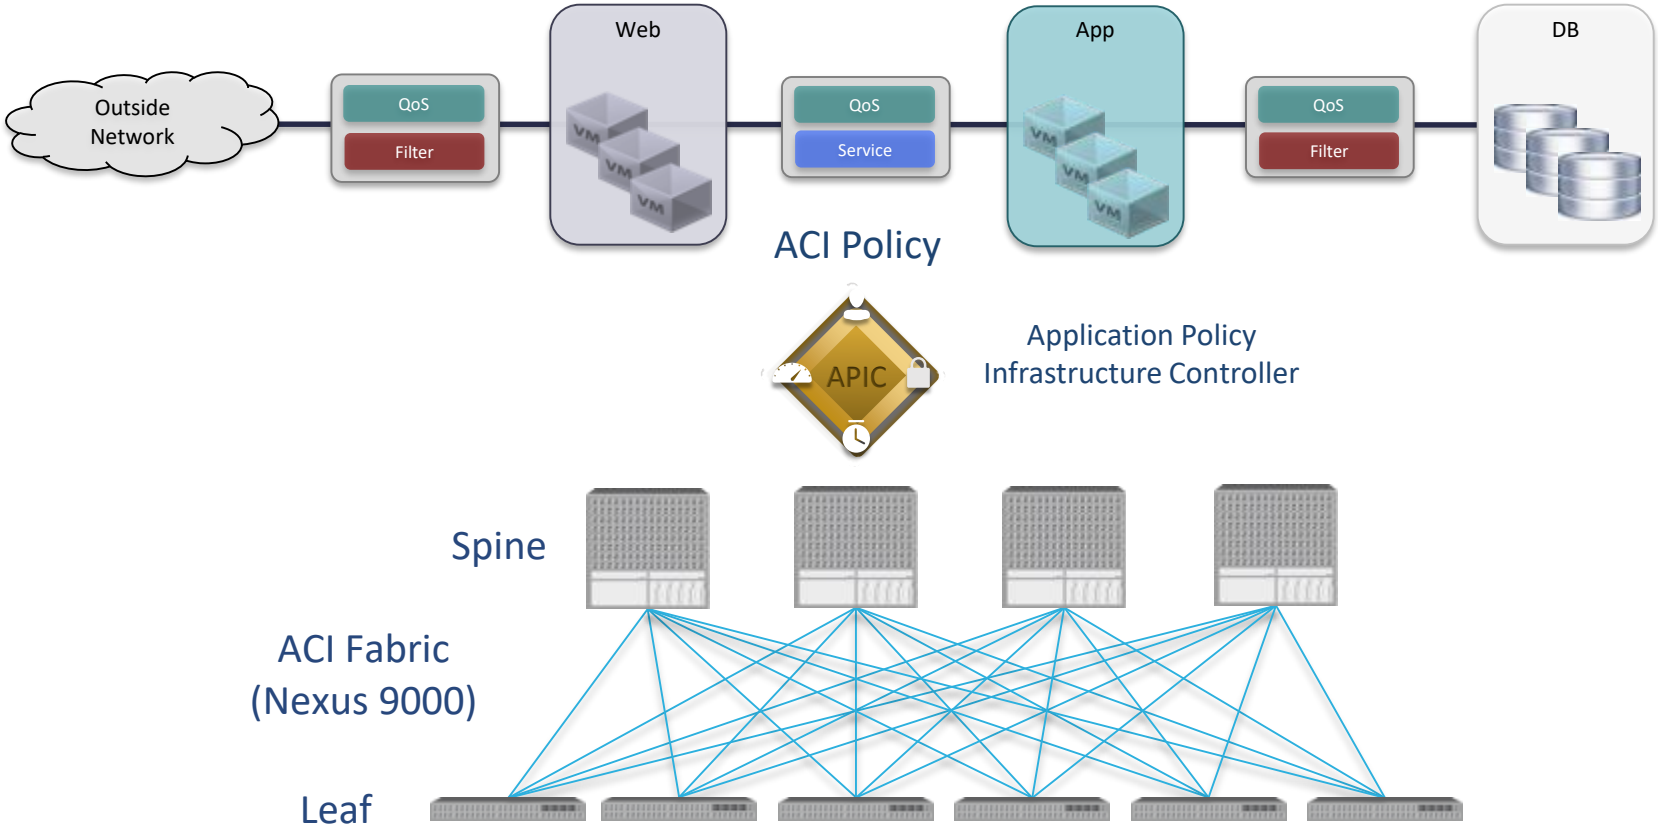

Cisco ACI Policy Constructs
EPG (End Point Group) and Contract

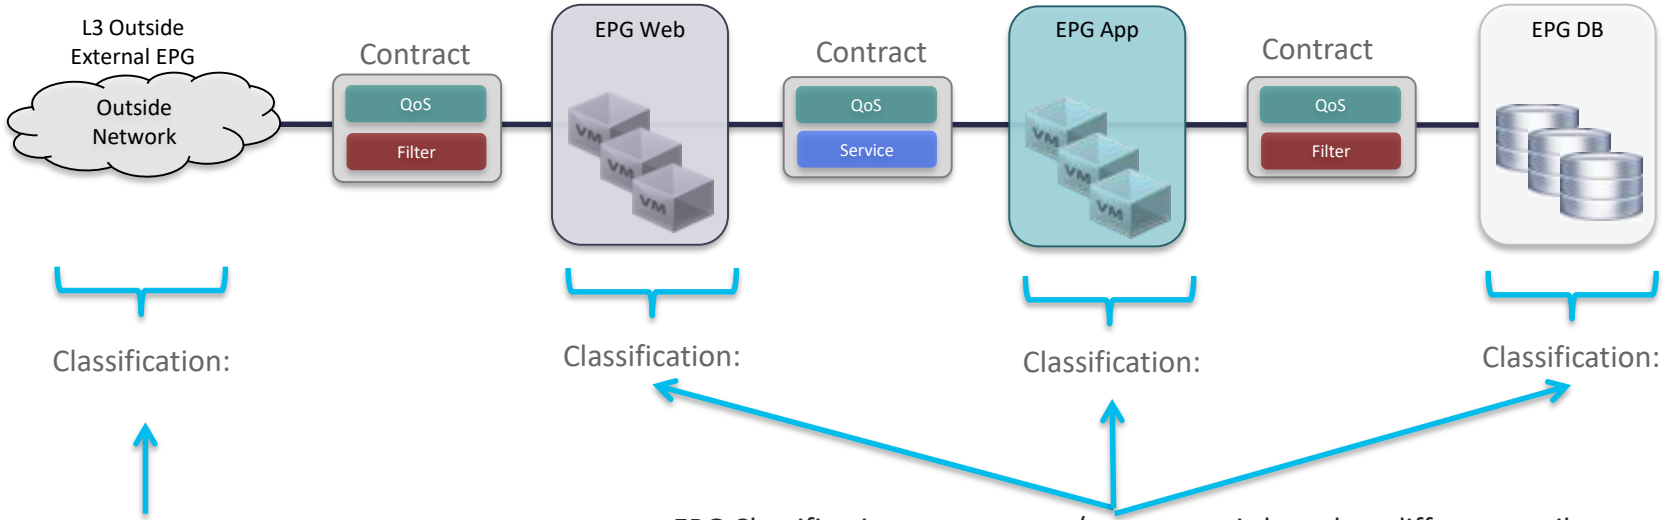

EPG Classification on L3 Outside is based on IP address Network/Mask

EPG Classification on an access/server port is based on different attributes

- Port + VLAN, Port + VXLAN, Network/Mask
- IP, MAC Address
- IP, MAC, VM Attributes for VM's

Physical, Virtual, container endpoints can co-exist in same EPG

Cisco ACI Policy Constructs

Tenants, Application Profiles, Bridge Domains, VRFs

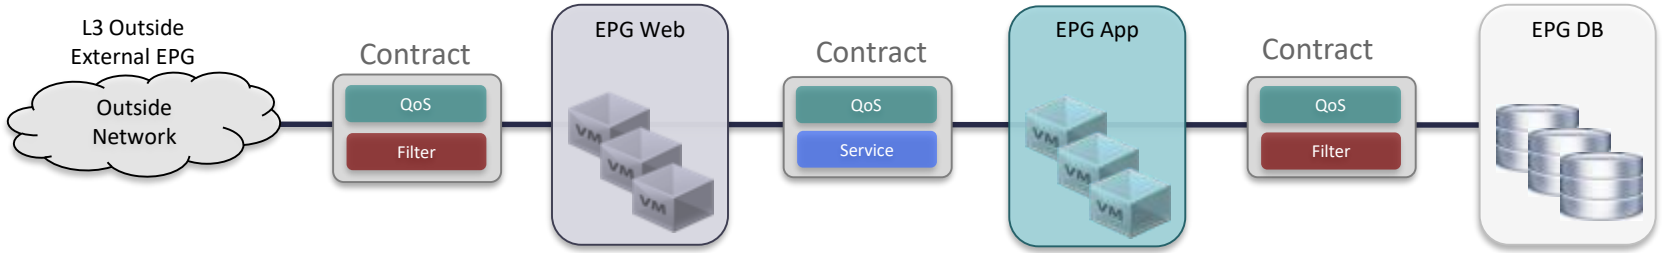

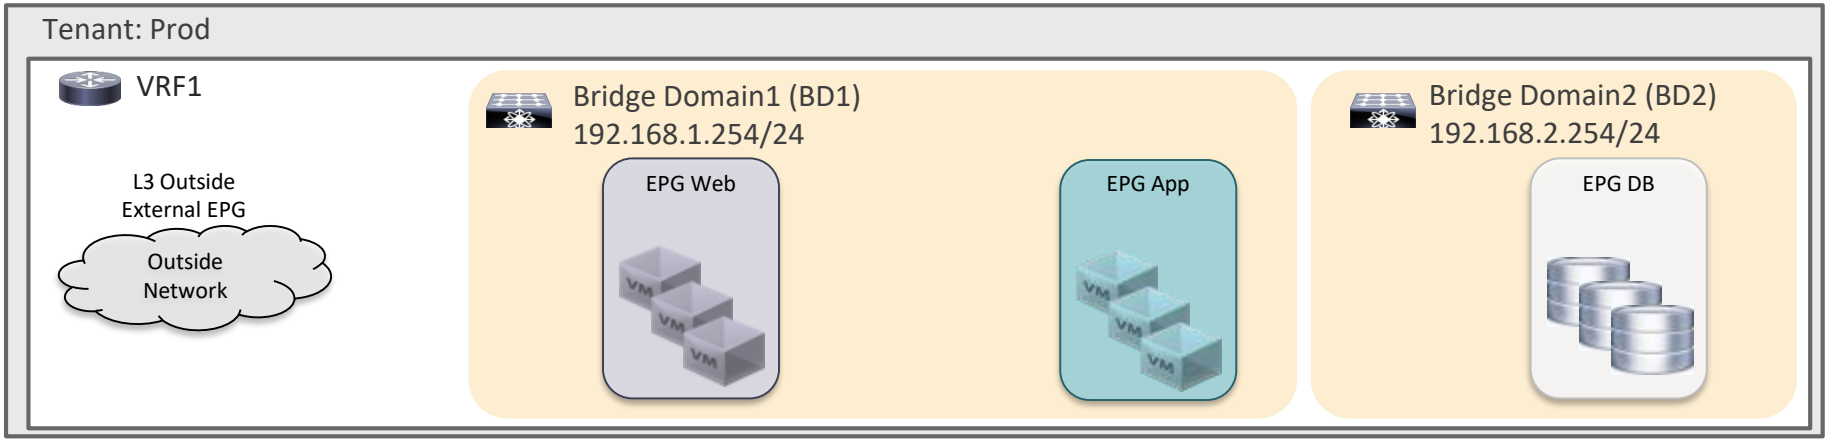

Micro-Segmentation with ACI

- Micro EPG (uSeg EPG)
  - EPG classification based on IP, MAC, VM attributes
- Intra-EPG isolation
  - Deny traffic between endpoints in same EPG
- Intra-EPG contract
  - Contract enforcement on traffic between endpoints in same EPG

Micro EPG (uSeg EPG)

- EPG classification based on IP, MAC, VM attributes
- Endpoints assigned to the uEPG regardless of the encapsulation/port

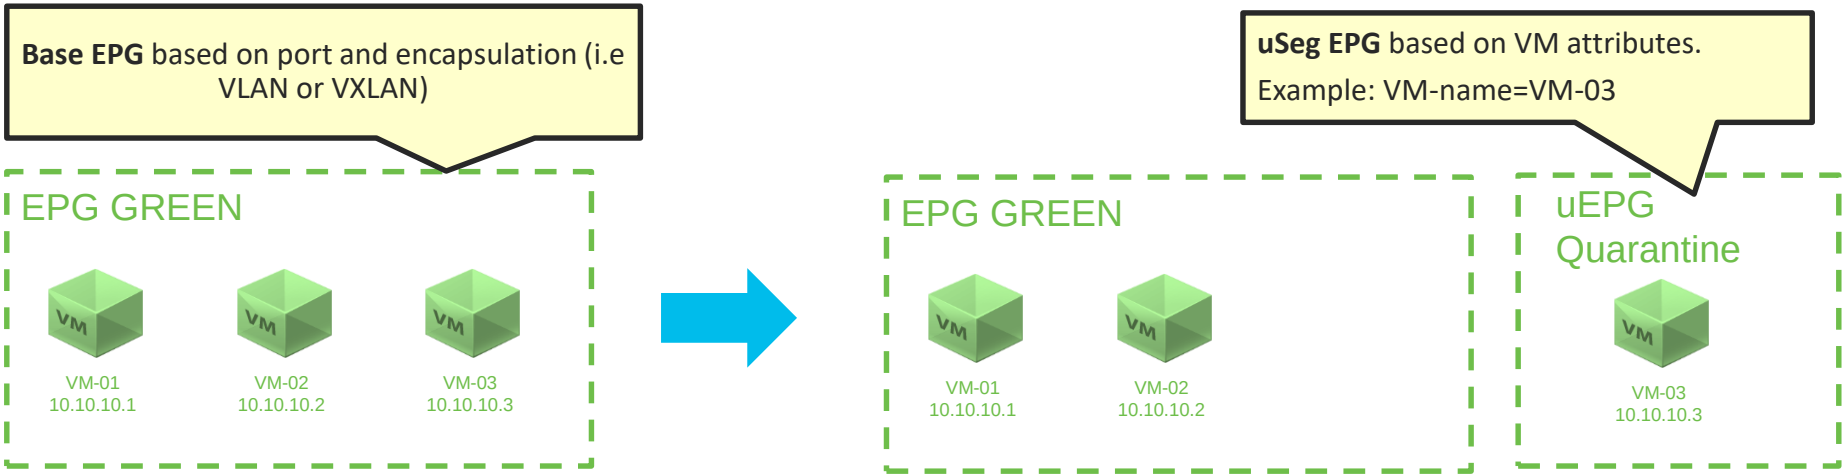

Intra-EPG Isolation and Intra-EPG Contract

- By default, endpoints in same EPG can talk without contract (permit-all)
- Intra EPG isolation is an option to deny traffic within an EPG (deny-all)
- Intra EPG contract is an option to filter traffic within an EPG (filter)

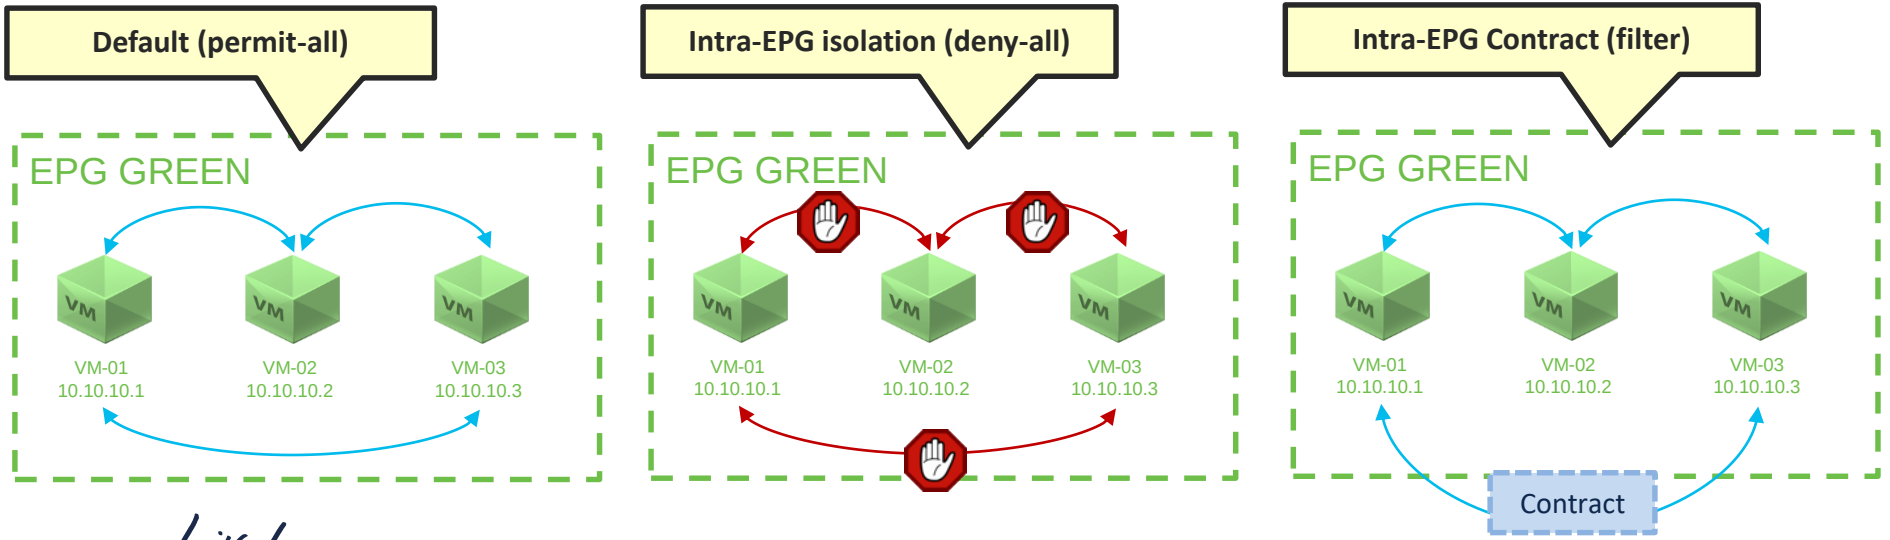

ACI L4-L7 Service integration

You make networking **possible**

L4-L7 Design Tips

- Understand desired traffic flow
  - North-South FW?
  - East-West FW?
  - Service Chain order?
  - Is there IP and/or port translation?

- Are there devices located in multiple DCs?

L4-L7 Design Options
Understand Requirements

- Firewall/IPS
  - Firewall: Layer 1(inline), Layer 2(Transparent) or Layer 3(Routed)?
  - Gateway: ACI or Firewall?
  - Insertion: VLAN/VRF stitching or PBR?
  - HA option: Active/Standby, Active/Active Cluster or Independent Active nodes
- Load Balancer
  - Load Balancer: Layer 3
  - How to handle return traffic: LB as Gateway, SNAT, PBR or DSR?
  - HA option: Active/Standby
  - VIP: Is VIP in self IP subnet range?

Firewall Design Options

You make networking **possible**

Firewall Design Options

L2 FW
VLAN stitching

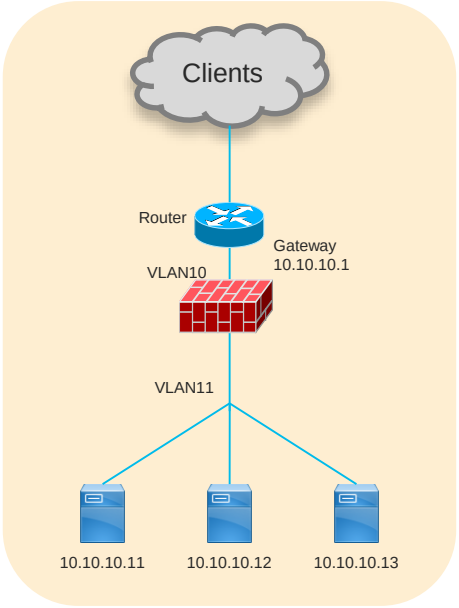

L3 FW
FW as gateway

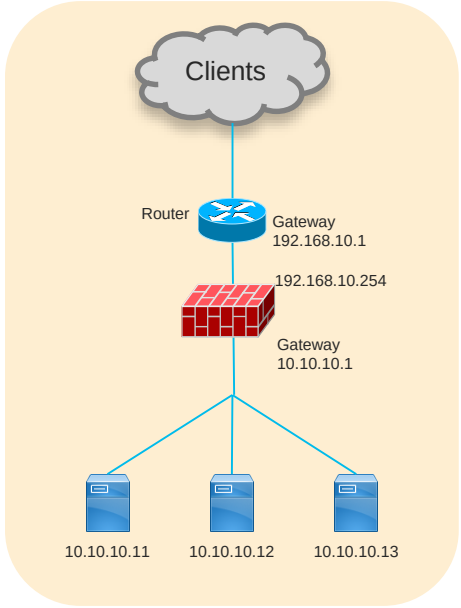

L3 FW
Fabric as gateway
VRF sandwich

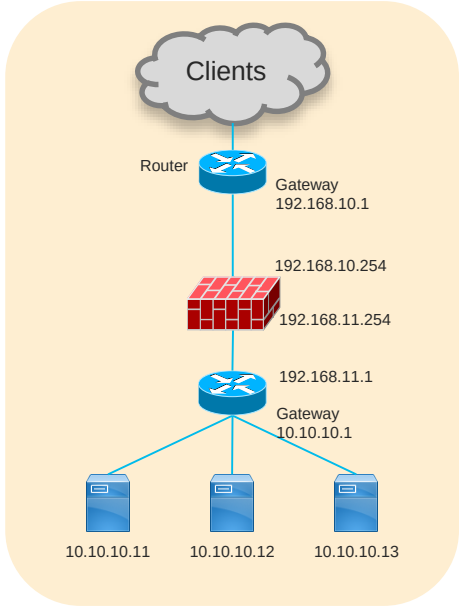

L3 FW
Fabric as gateway
PBR

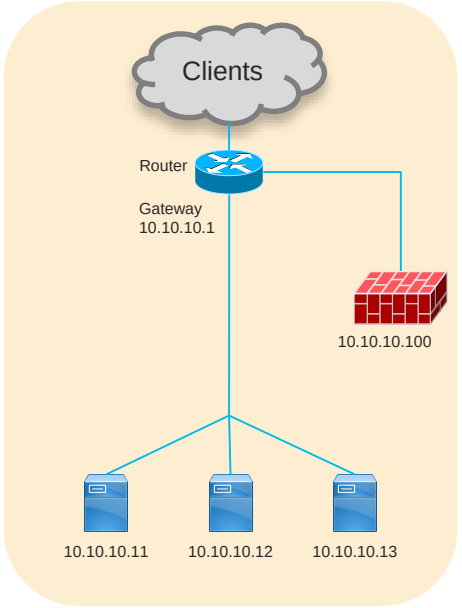

Option 1: L2 Firewall with VLAN Stitching

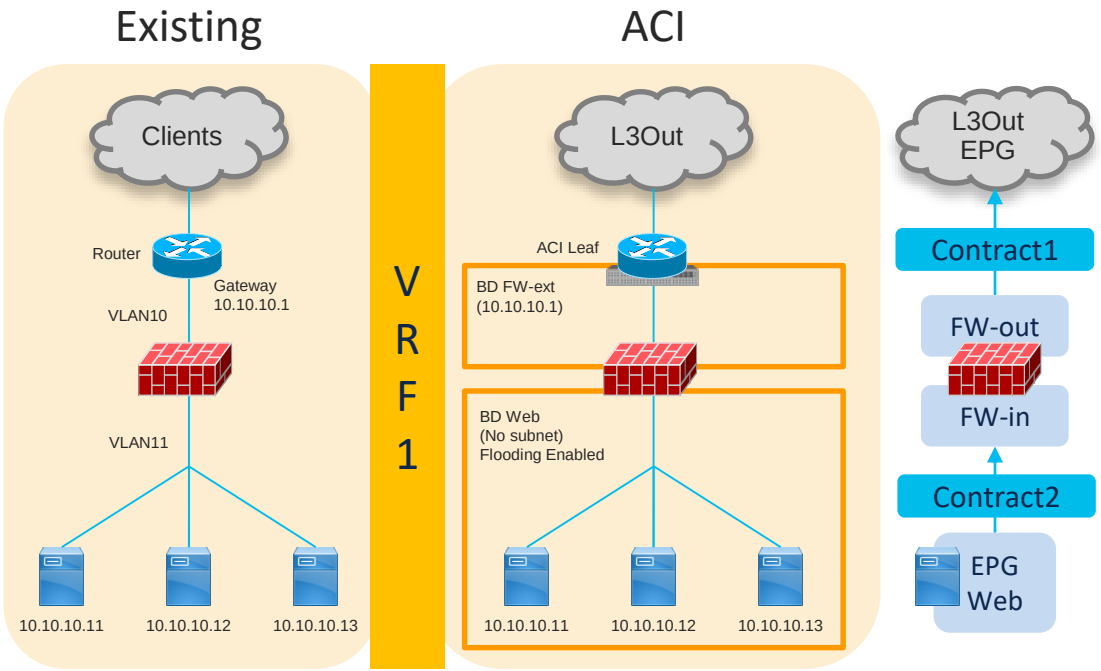

- Traditional VLAN stitching
- FW and EPG are in same BD
- ACI as L3
- All inter-BD traffic goes through FW
- Simple
- Service Graph is not mandatory

- L1/L2 PBR available in 4.0 that requires ACI as gateway and dedicated service BDs.

Option 2: L3 Firewall with the Firewall as the Default Gateway

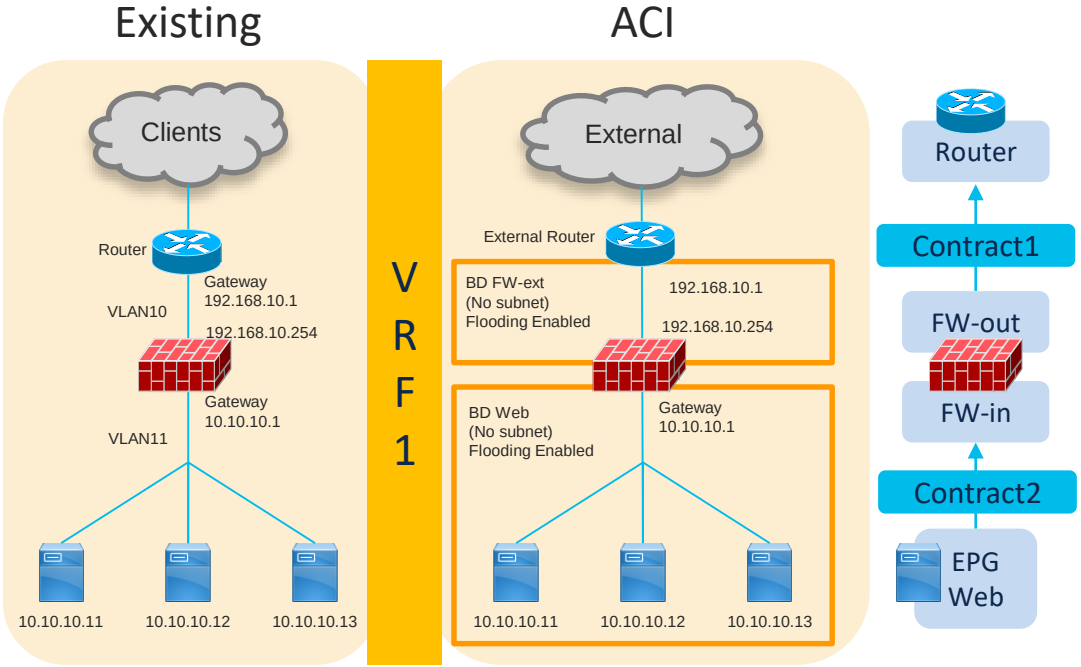

- FW as gateway
- FW and EPG are in same BD
- ACI as L2
- All inter-subnet traffic goes through FW
- Simple
- Service Graph is not mandatory

Option 3: L3 Firewall with the Fabric as the Default Gateway – "VRF sandwich"

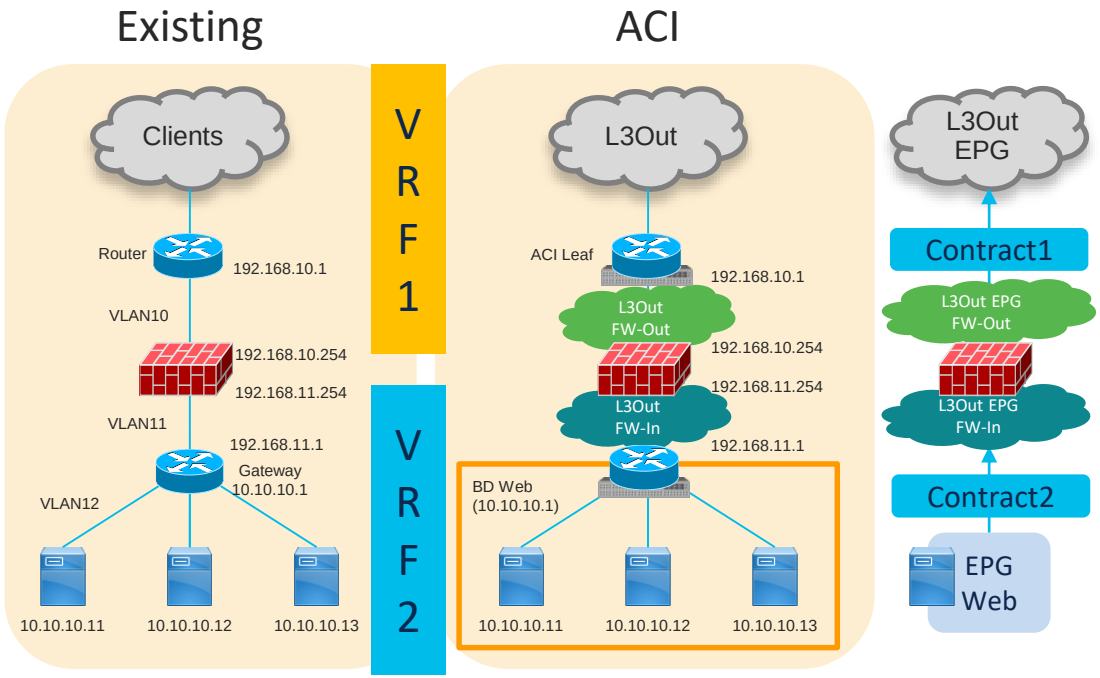

- Traditional VRF sandwich.
- FW is in L3out
- ACI as L3
- All inter-VRF traffic goes through FW
- Require multiple VRFs and L3outs
- Service Graph is not mandatory
- Good for North-South FW

Option 4: L3 Firewall with the Fabric as the Default Gateway, Redirect with PBR

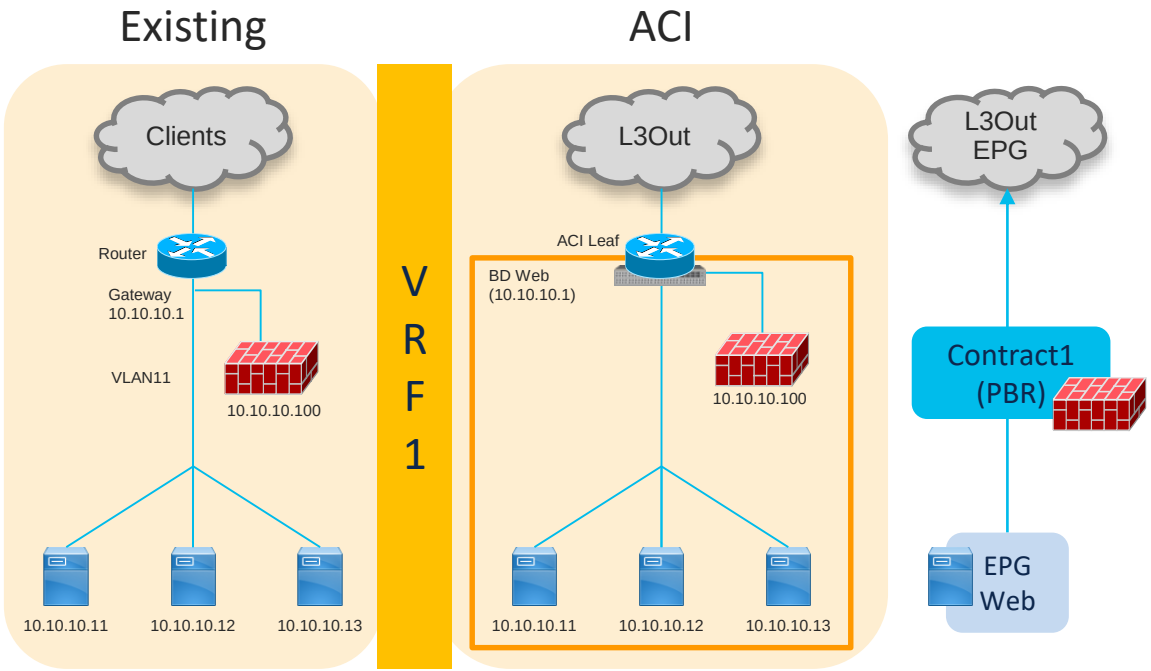

- PBR (Policy Based Redirect).
- ACI as L3
- FW is in BD
- Specific traffic goes through FW
- FW can be two or one arm mode

- Good for East-West
- Requires the use of Service-Graph
- Service device can be in same or different BD with servers

ACI PBR Use Cases

- Inspect specific traffic

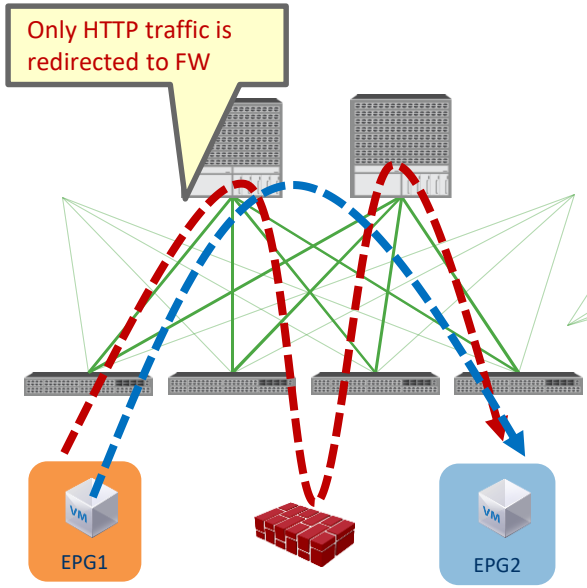

- Use different Firewall

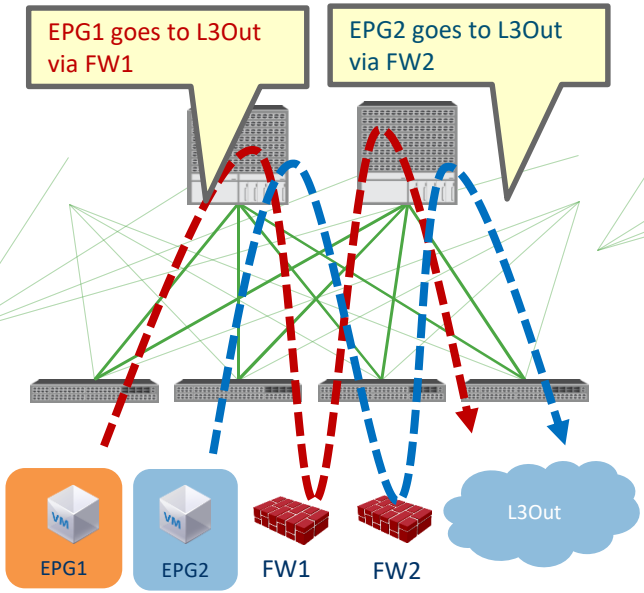

- LB without SNAT

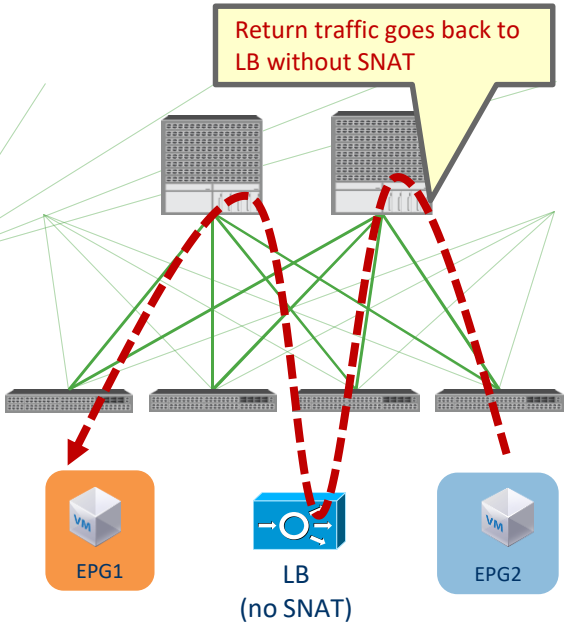

Symmetric PBR: Scale Firewall Easily

- Ensure incoming and return traffic goes to same firewall

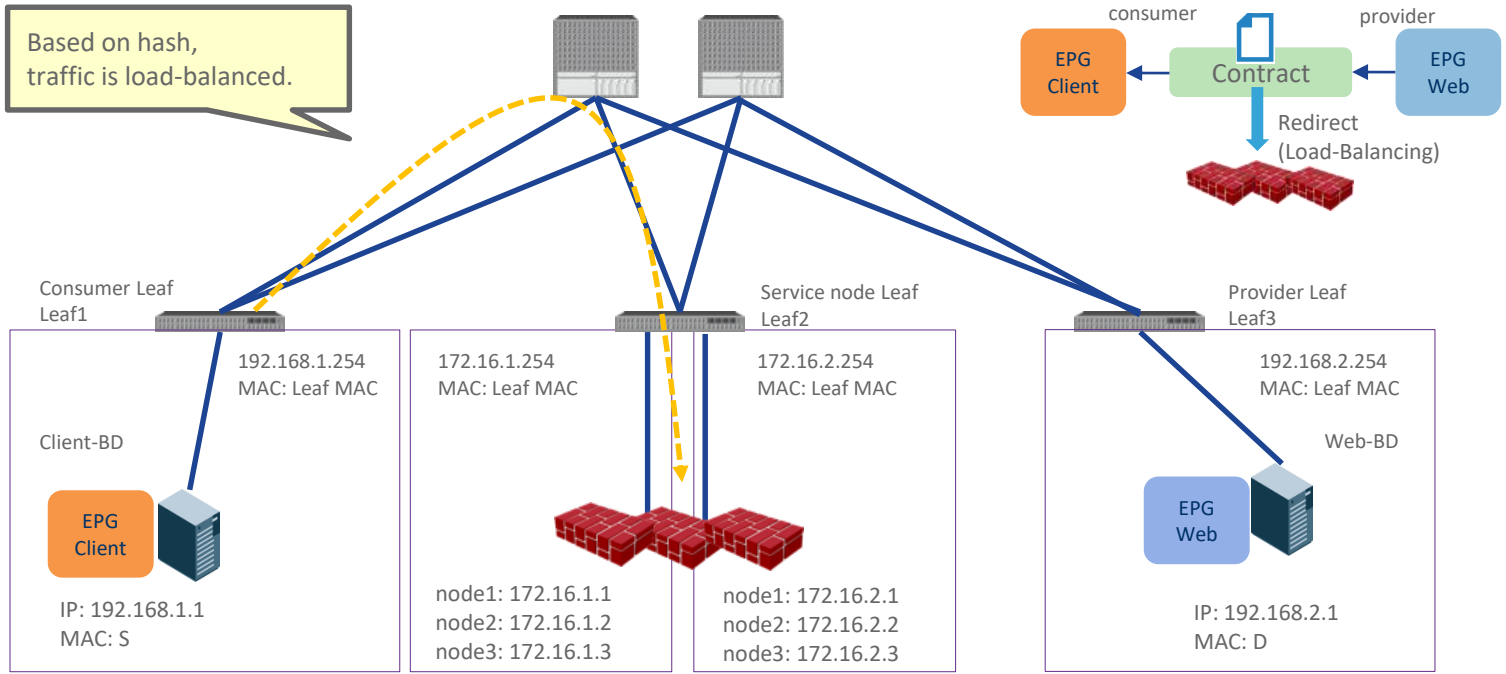

HA Options

Active/Standby Cluster

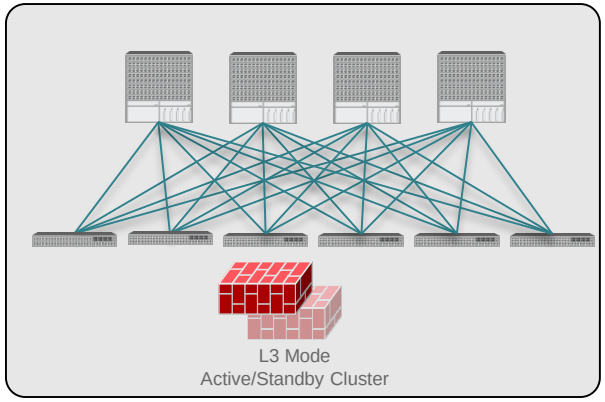

- PBR is not mandatory
- The Active/Standby pair represents a single MAC/IP entry.

Active/Active Cluster ('Scale-Up' Model)

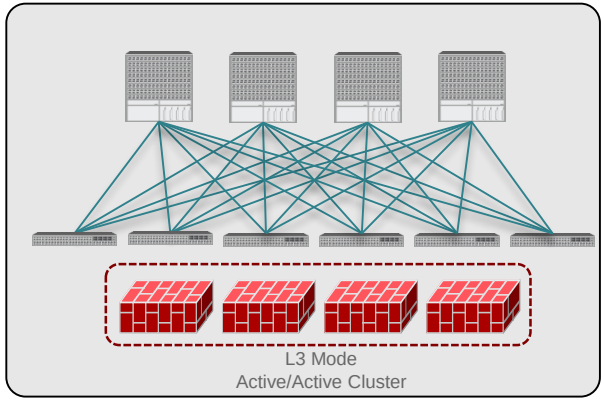

- PBR is required if the cluster is stretched across pods.
- The Active/Active cluster represents a single MAC/IP entry.
- Spanned Ether-Channel Mode supported with Cisco ASA/FTD platforms

Independent Active Nodes ('Scale-Out' Model)

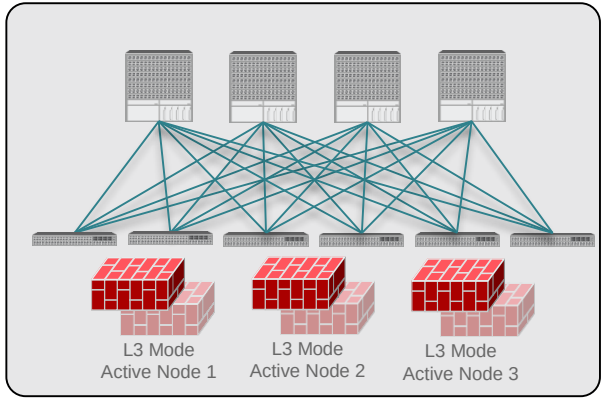

- PBR is required.
- Each Active node represent a unique MAC/IP entry.
- Use of Symmetric PBR to ensure each flow is handled by the same Active node in both directions

Copy Service

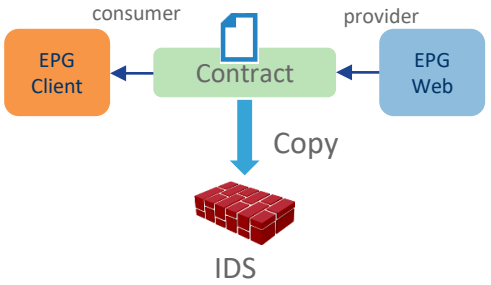

- APIC 2.0
- Service Graph is mandatory and EX/FX hardware is required
- Copy specific traffic

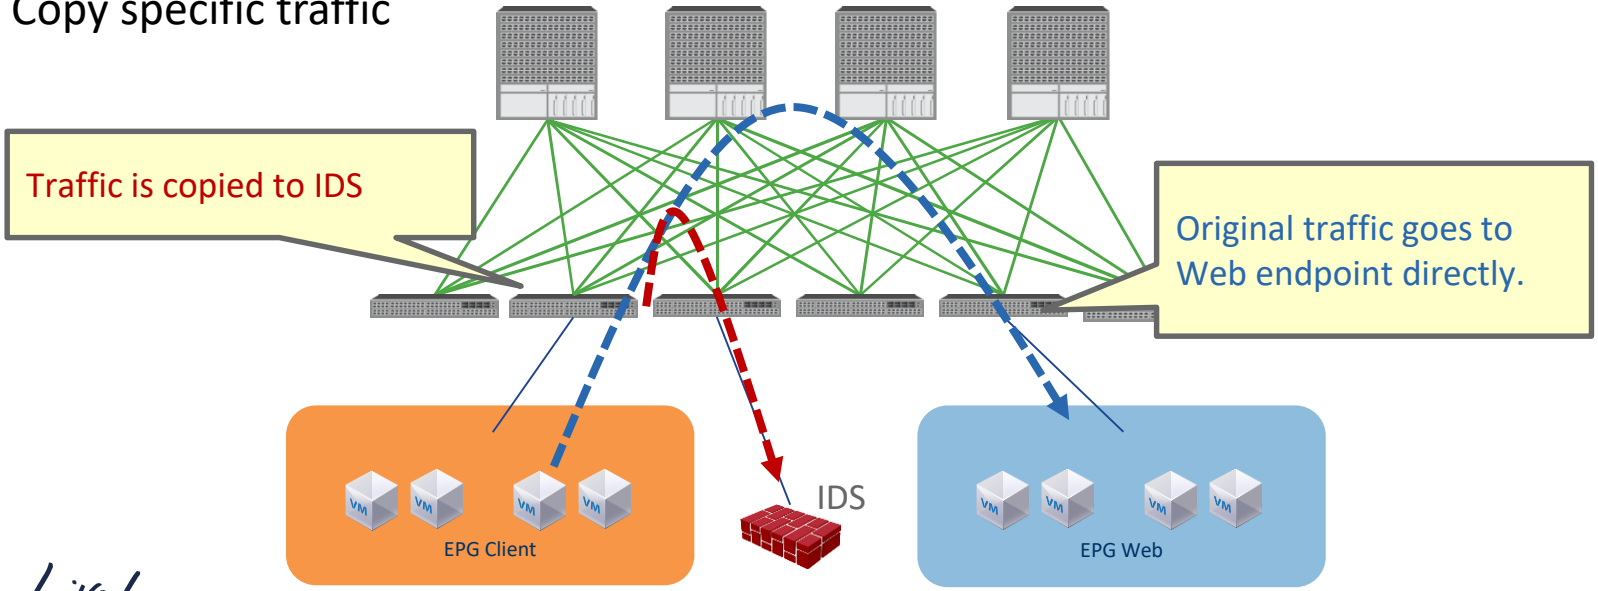

PBR Design FAQ

You make networking **possible**

Can We Use PBR for L3out EPG?

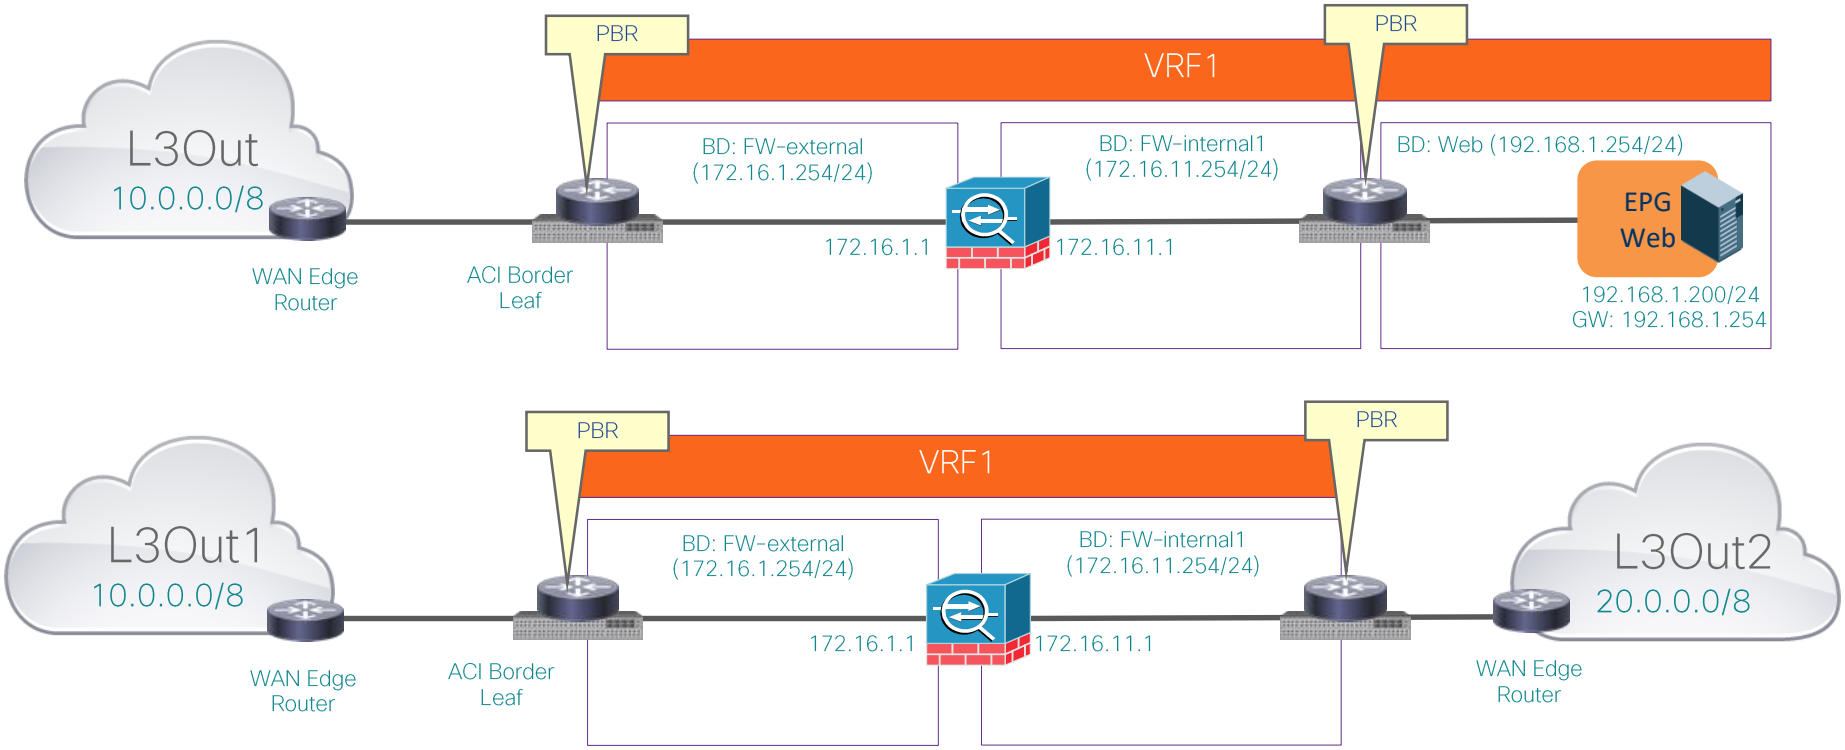

Can We Use PBR for EPGs in Same Subnet?

- Inspection between endpoints in same subnet

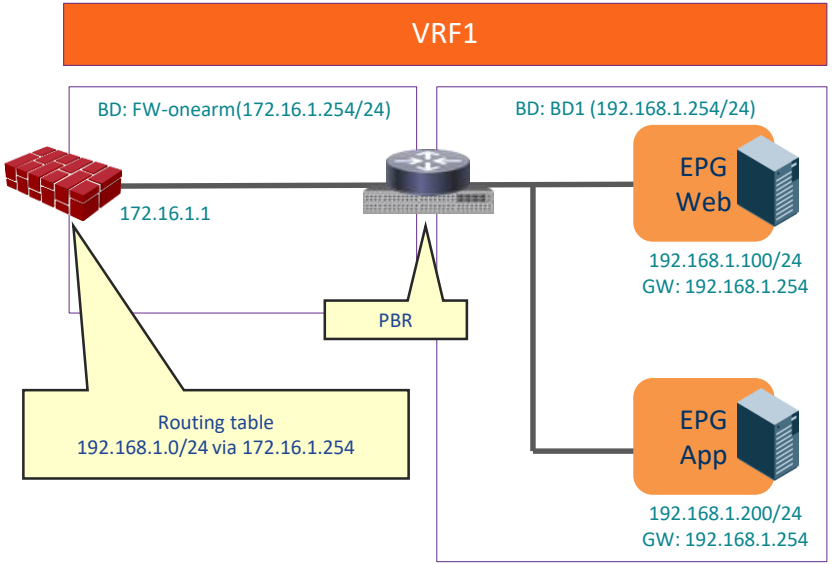

- Inspection between endpoints even in same EPG

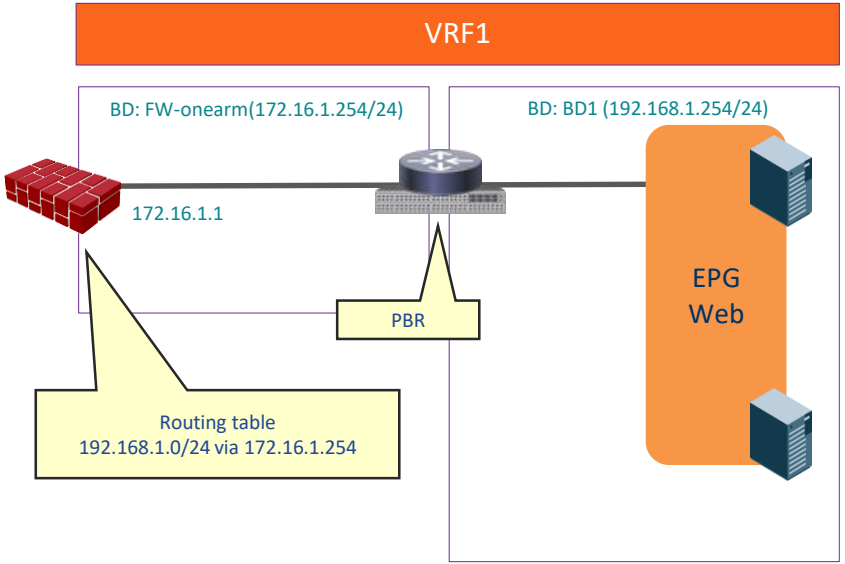

One-Arm vs Two-Arm?

- One-Arm
  - Simple routing design on service node
  - Some firewall doesn't allow intra-interface traffic by default
- Two-Arm
  - Need to manage routing design on service node
  - Different security level on each interface

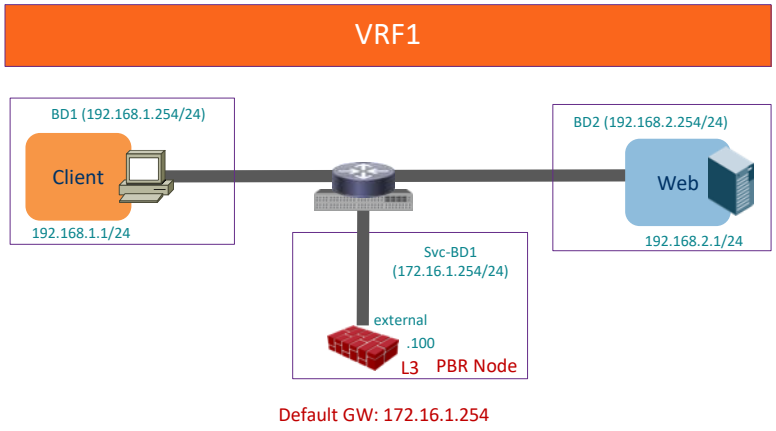

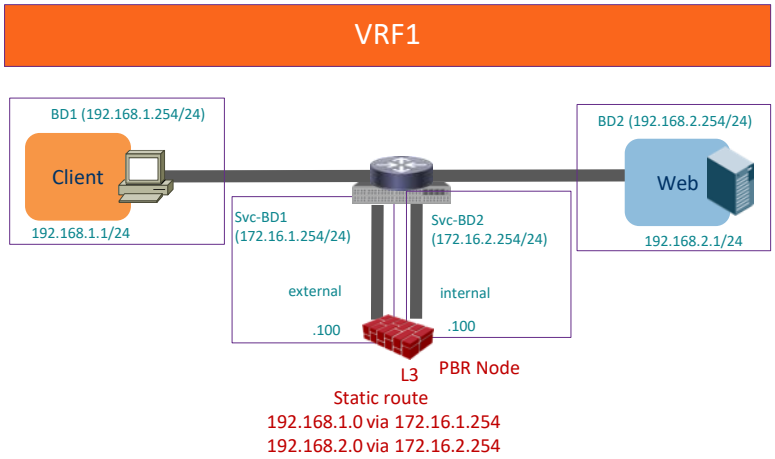

Can We Reuse Same PBR Node Multiple Times?

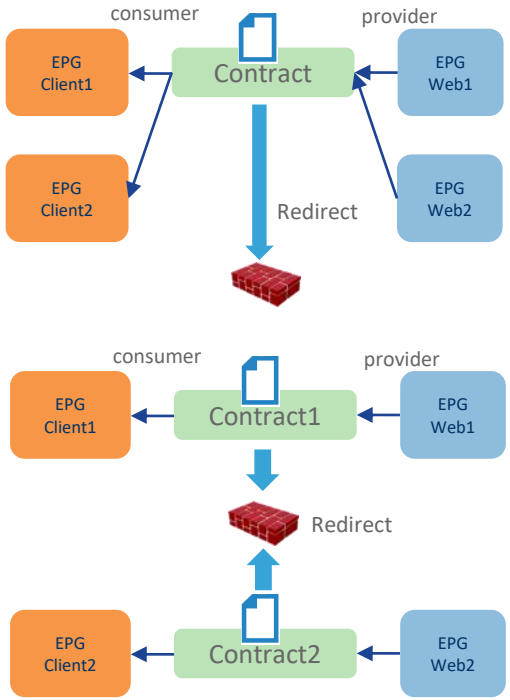

- Multiple consumer/provider EPGs
- Multiple contracts using same PBR destination and Service Graph.

- Note
  - Depending on routing design, one-arm mode deployment may be required.

Can We Insert Firewall to Any-To-Shared-Service?

- vzAny is useful if we have a security requirement that is applied to all EPGs in same VRF and also it helps to reduce policy TCAM consumption.
- Prior to 3.2, PBR with vzAny (consumer) is supported.
- In ACI 3.2, PBR with vzAny (provider) is also supported.
- Use case: Insert Firewall everywhere.

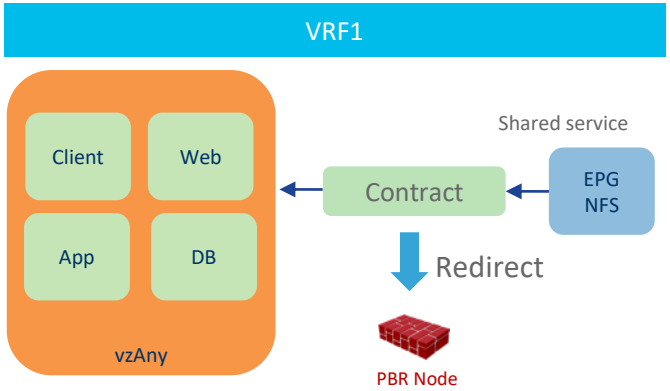

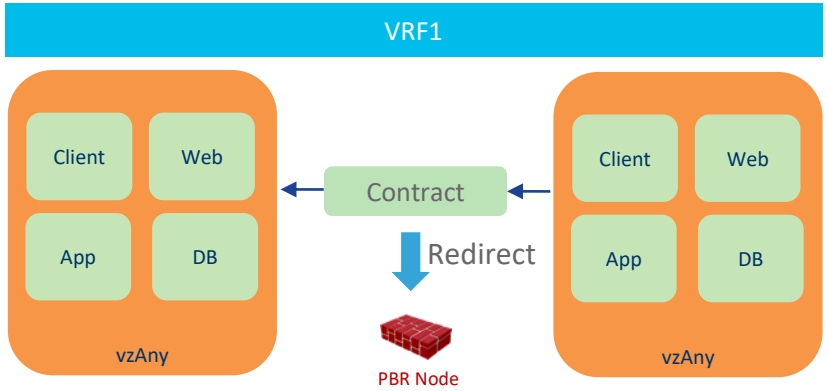

Can PBR Node be in Consumer/Provider Subnet?

- Prior to APIC version 3.1, PBR node must be different than the consumer/provider BDs.
- Starting from APIC version 3.1, this requirement no longer mandatory. (Need EX/FX Leaf)

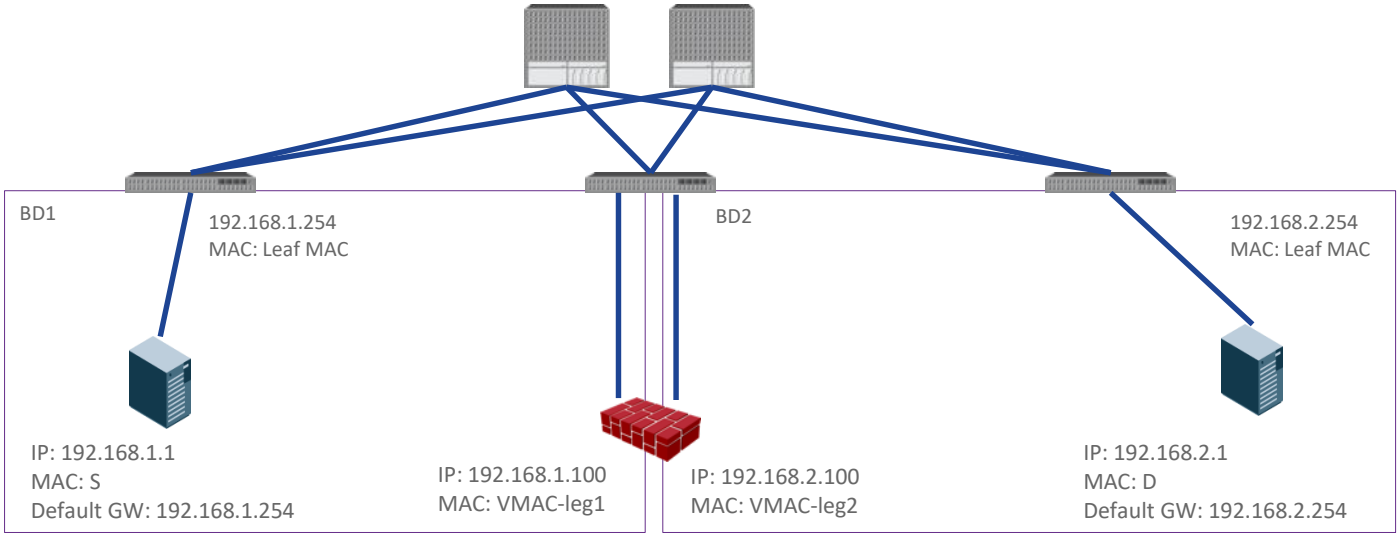

Can We Concatenate Services? Multi-Node PBR

- Prior to ACI 3.2: Concatenating PBR nodes is not supported.
  - For example, both 1st and 2nd node can't be PBR nodes. Either one of them can be.

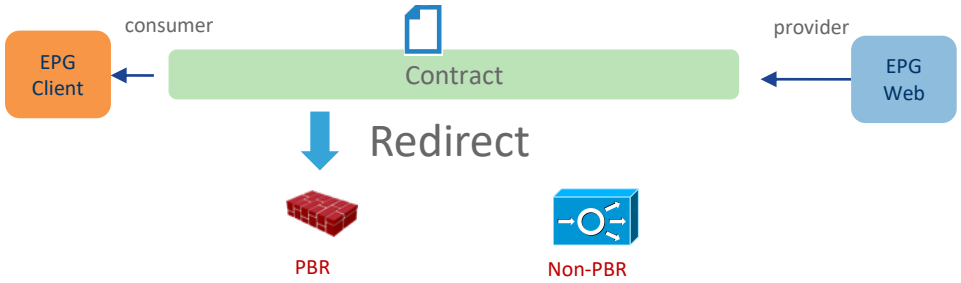

- ACI 3.2: Support more than 1 node PBR in a Service Graph.

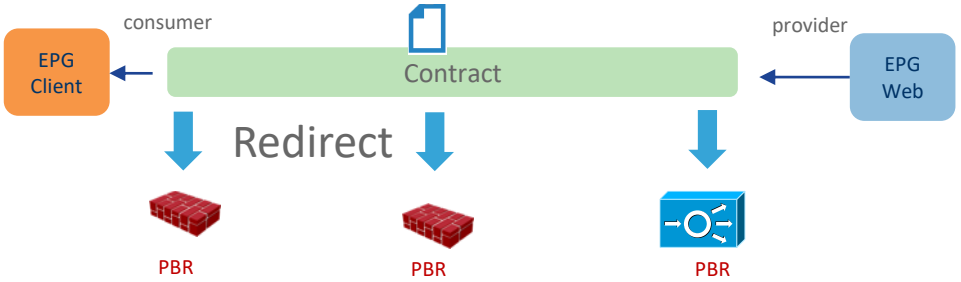

Load Balancer Design Options

You make networking **possible**

Load Balancer Design Options

Two-arm (inline)
LB as Gateway
No SNAT/PBR

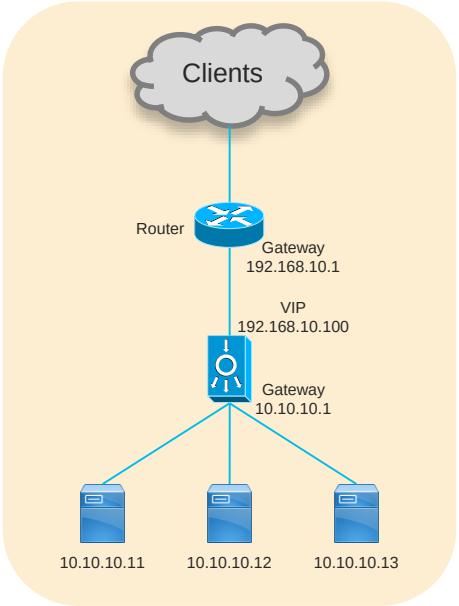

Two-arm (inline)
Fabric as Gateway
VRF sandwich

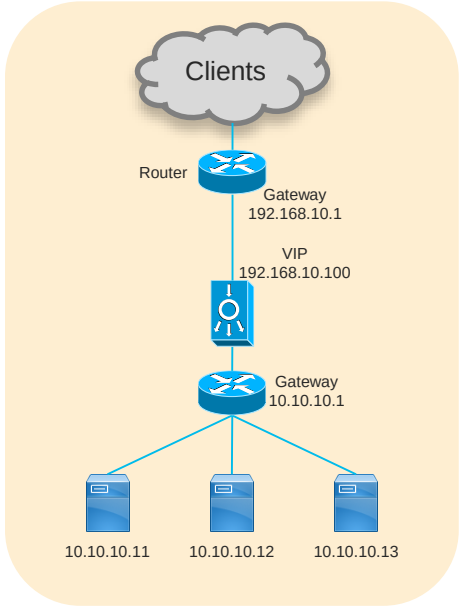

Two-arm
Fabric as Gateway
SNAT/PBR

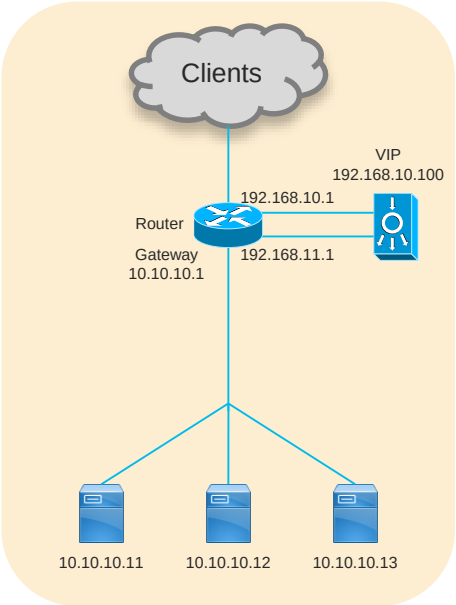

One-arm
Fabric as Gateway
DSR/SNAT/PBR

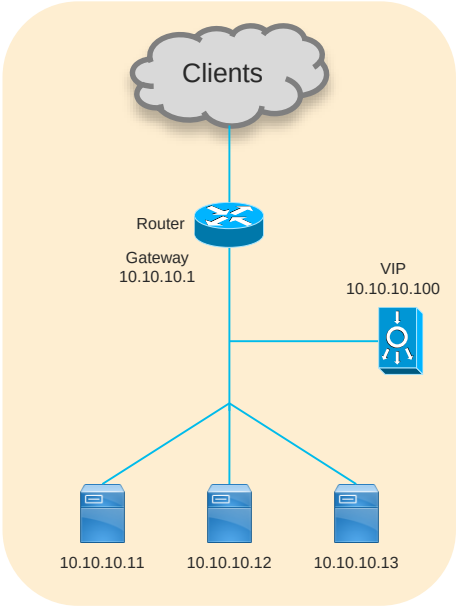

Option 1: Two-Arm (Inline) with the SLB as the Default Gateway

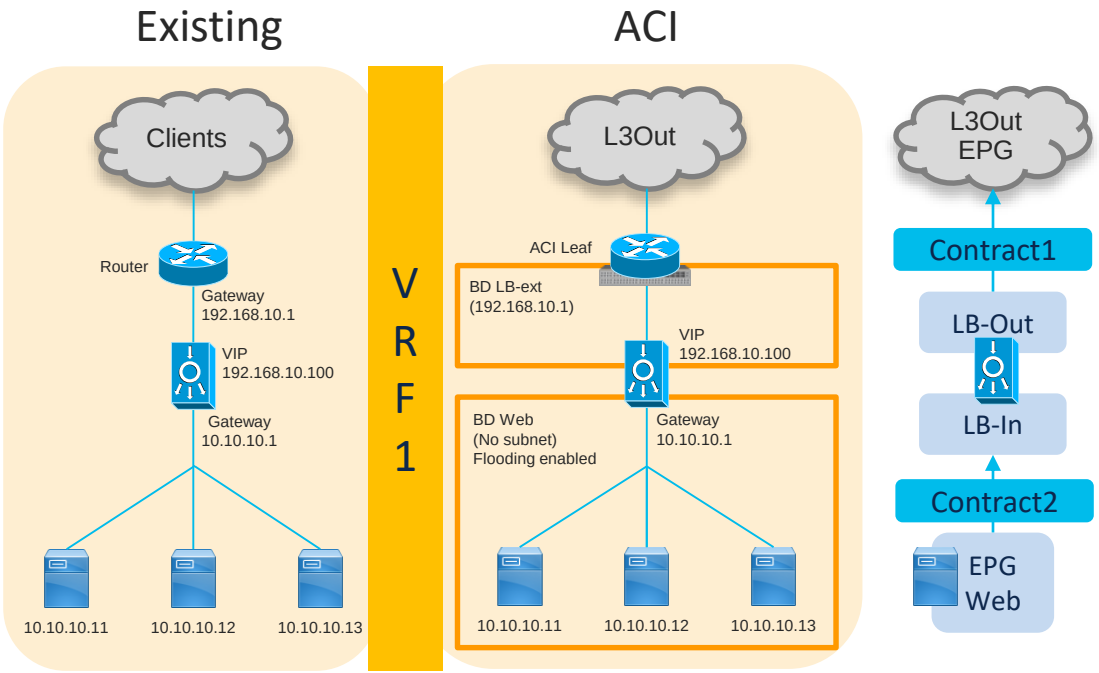

- LB and EPG are in same BD
- ACI as L2
- All inter-BD traffic goes through LB
- Simple
- ACI can be L3 for external side of LB
- Service Graph is not mandatory
- SNAT/PBR is not required

Option 2: Two-Arm (inline) with the Fabric as the Default Gateway

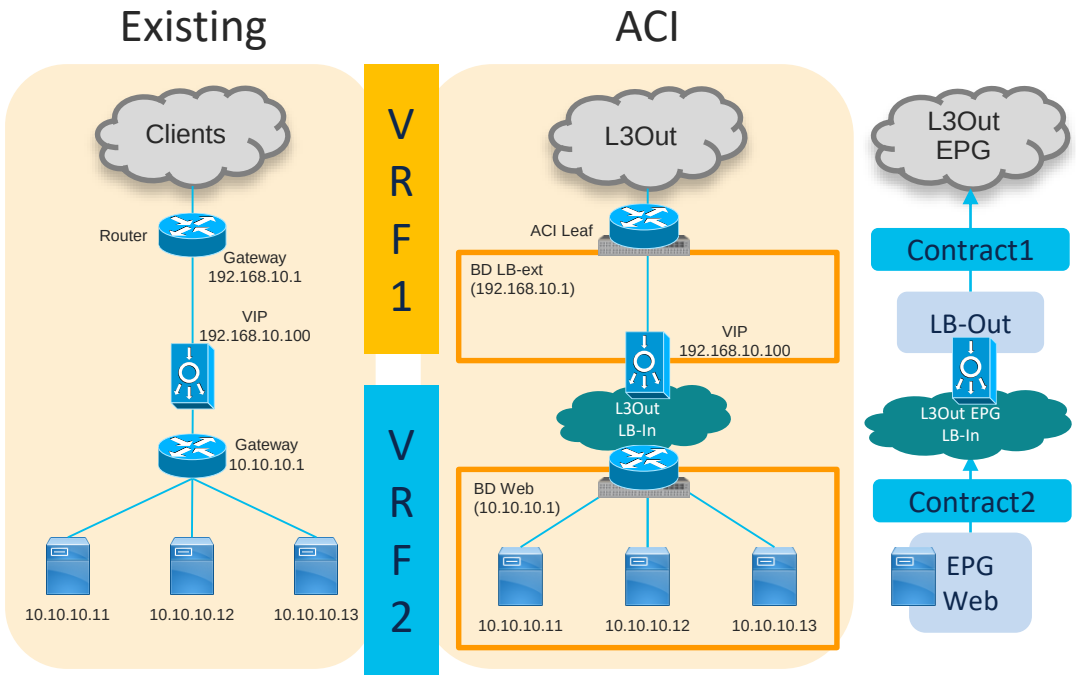

- Traditional VRF sandwich
- ACI as L3
- All inter-VRF traffic goes through LB
- Service Graph is not mandatory
- SNAT/PBR is not required

- If SNAT is enabled on LB using LB internal interface as NAT IP, LB-in can be in a BD. VRF2 and L3Out LB-in are not required.

Option 3: Two-Arm with the Fabric as the Default Gateway – SNAT/PBR

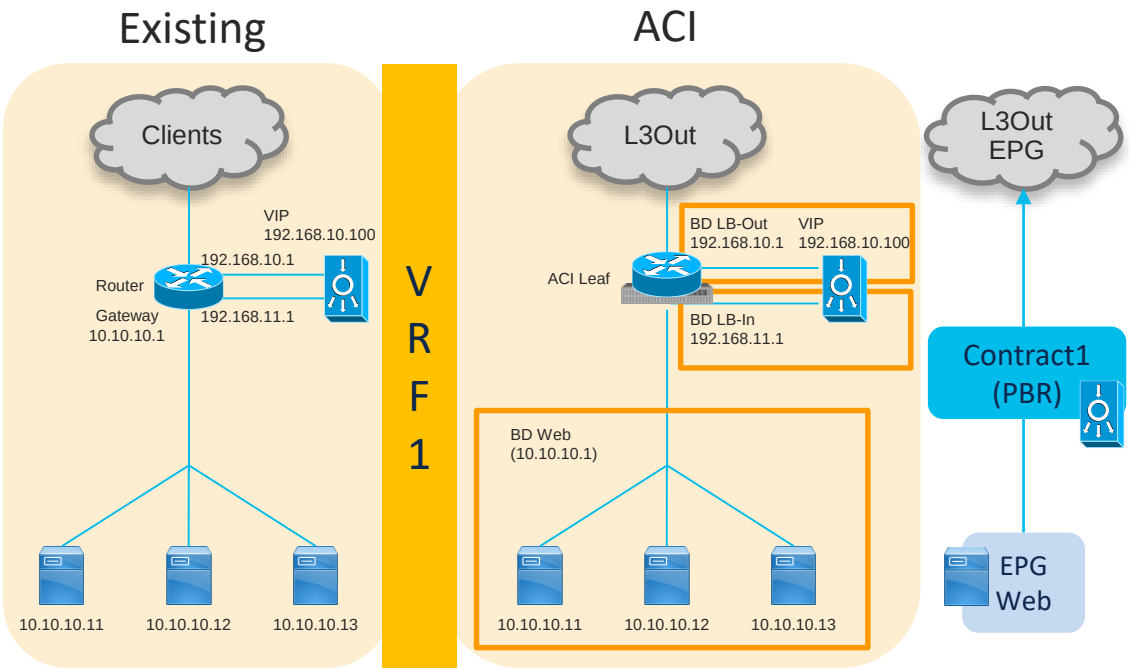

- PBR or SNAT is required
- ACI as L3
- Service device can be in same or different BD with servers
- If it's PBR:
  - Service Graph is required
  - Specific traffic goes through LB

Option 4: One-Arm with the Fabric as the Default Gateway - L2DSR/SNAT/PBR

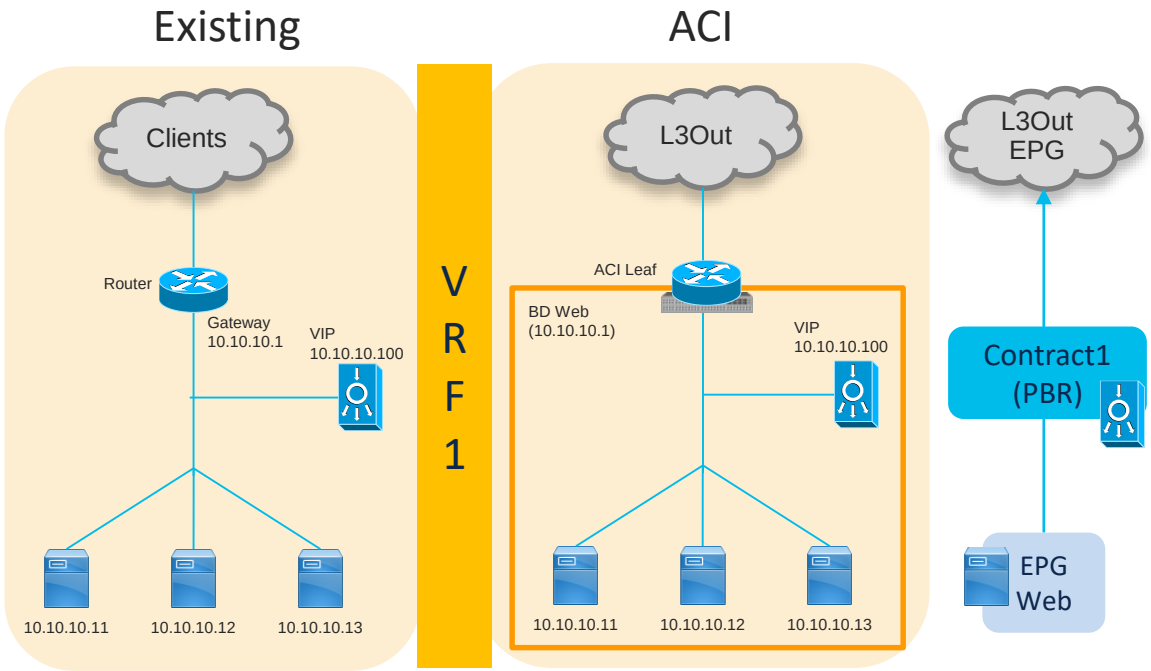

- L2DSR, PBR or SNAT is required
- ACI as L3
- Service device can be in same or different BD with servers
- If it's PBR:
  - Service Graph is required
  - Specific traffic goes through LB

What if the VIP is not in LB Interface IP Subnet Range?

Use L3Out (or /32 static route on BD)

Two-arm (inline)
LB as Gateway
No SNAT/PBR

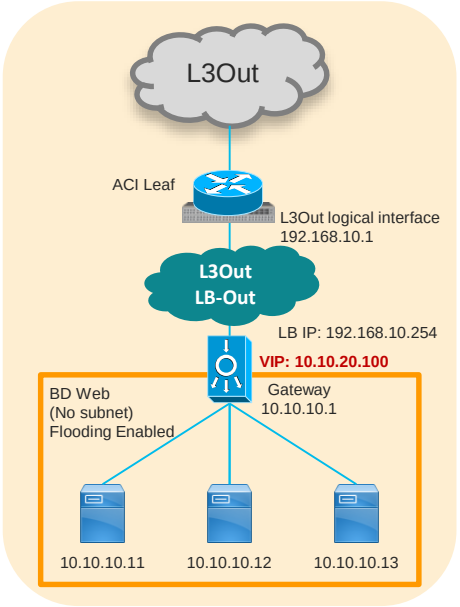

Two-arm (inline)
Fabric as Gateway
VRF sandwich

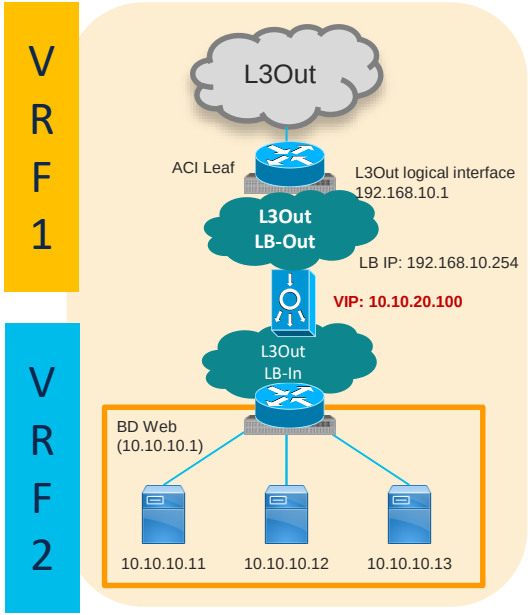

Two-arm
Fabric as Gateway
SNAT or PBR(After 5.0)

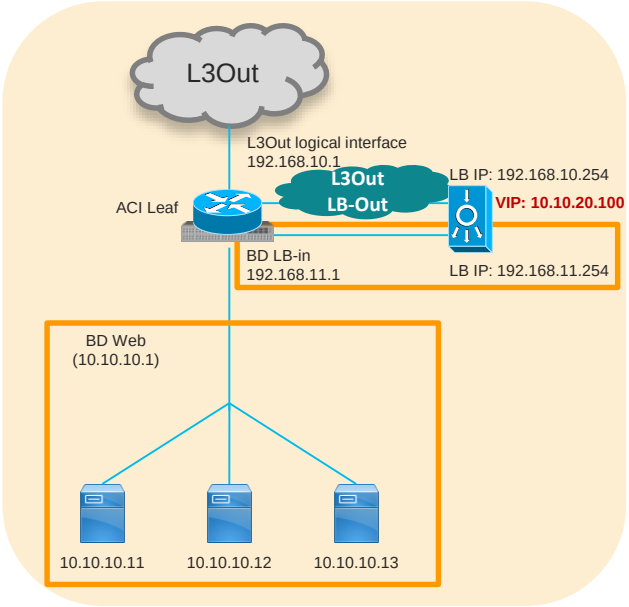

Multi-location Data Centres

You make the power of data **possible**

ACI Anywhere

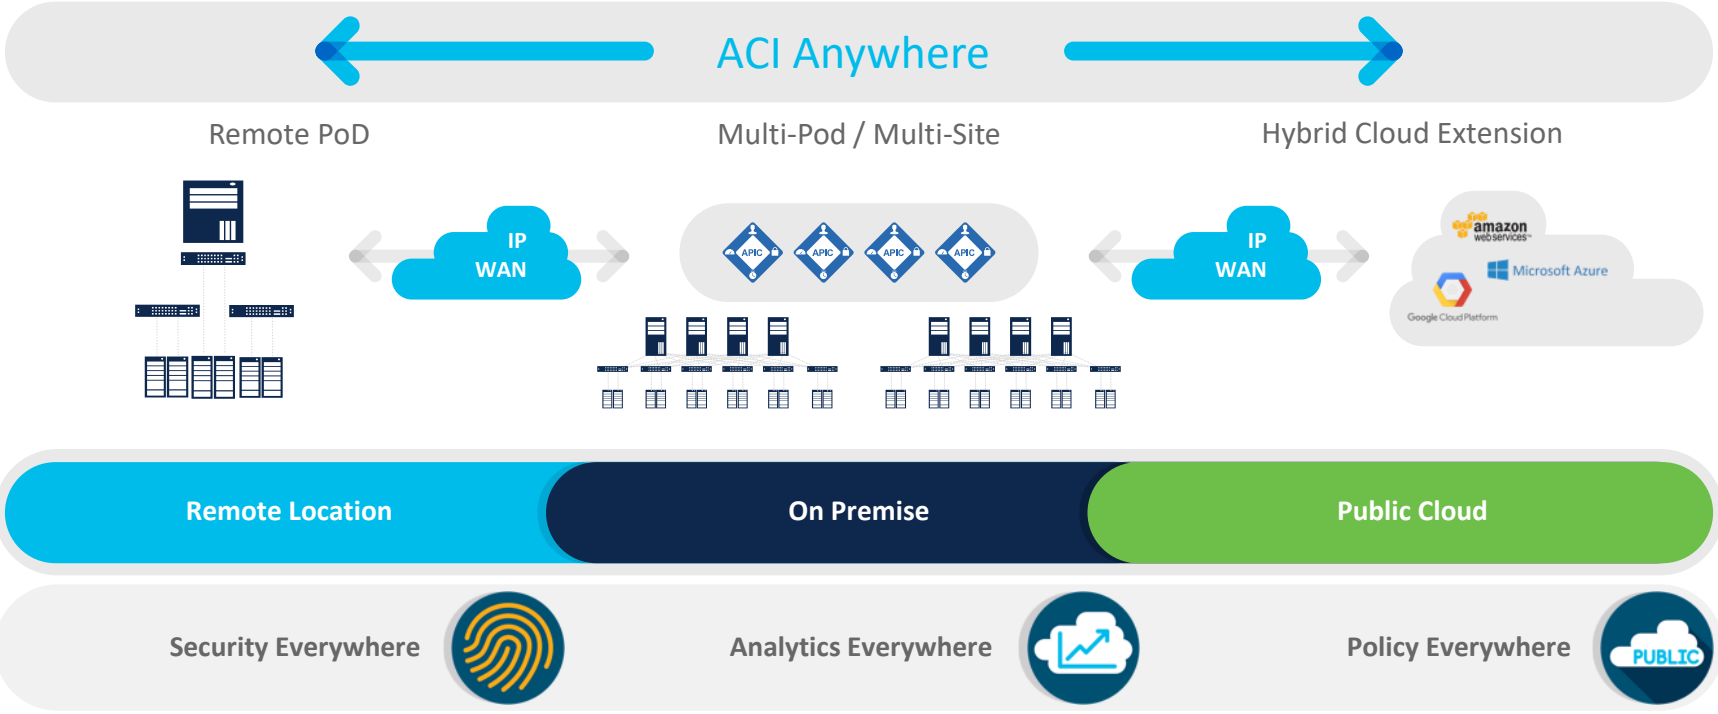

Service Insertion in Multiple DC Locations
What is the Challenge of Service Insertion in Multiple DC Locations?

- Traffic Symmetricity is important

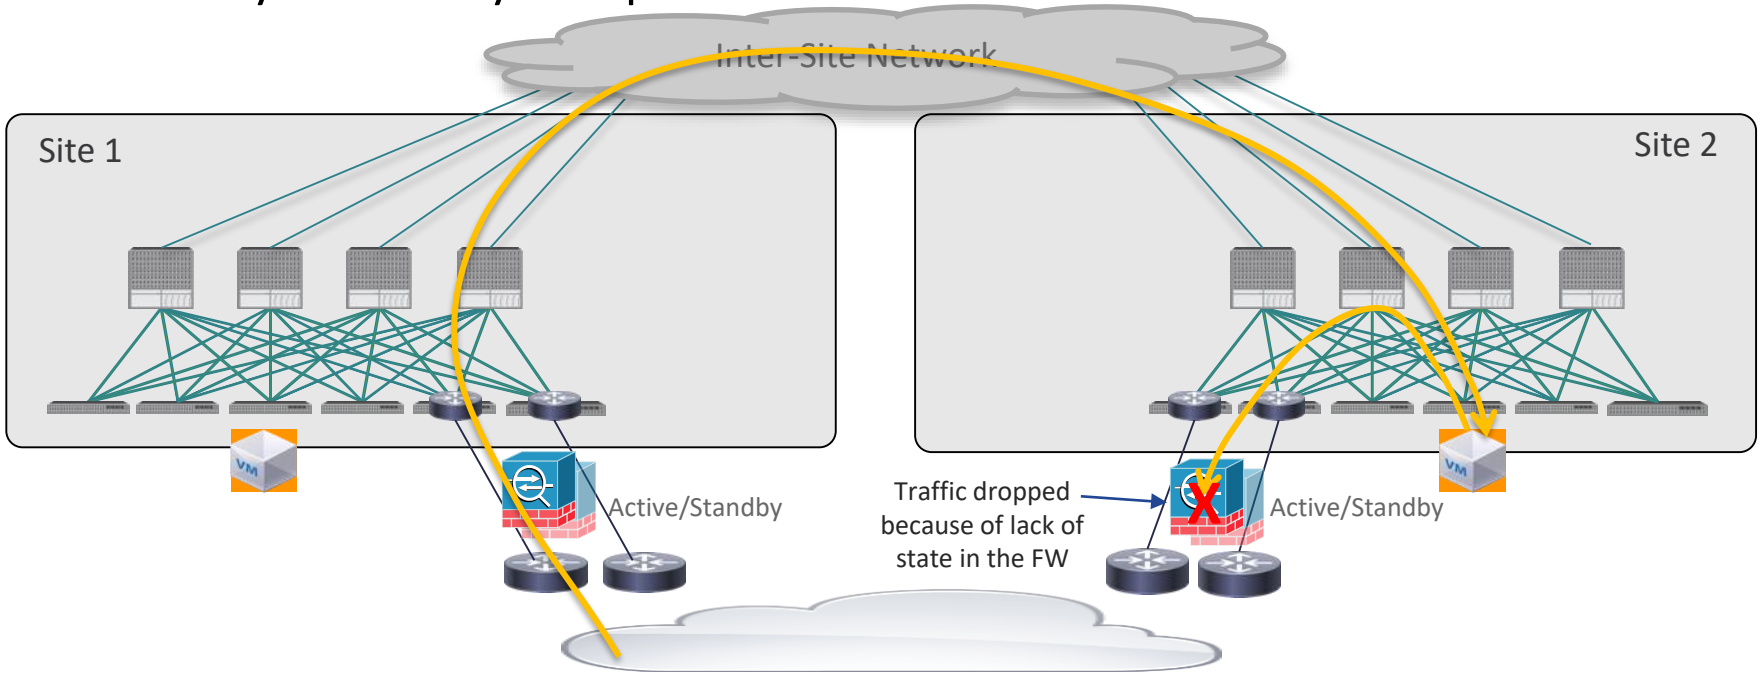

Multi-Pod and Network Services
Integration Models

For Your Reference

Typical options for an Active/Active DC use case

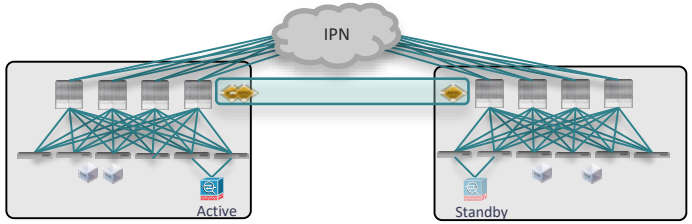

- Active and Standby pair deployed across Pods
- No issues with asymmetric flows

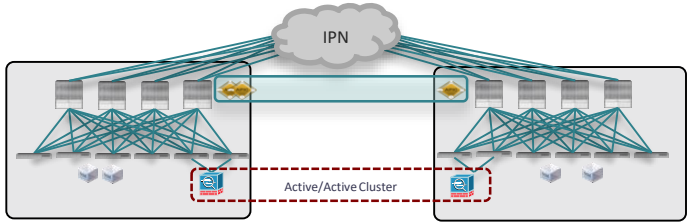

- Active/Active FW cluster nodes stretched across Sites (single logical FW)
- Requires the ability of discovering the same MAC/IP info in separate sites at the same time
- Supported from ACI release 3.2(4d) with the use of Service-Graph with PBR

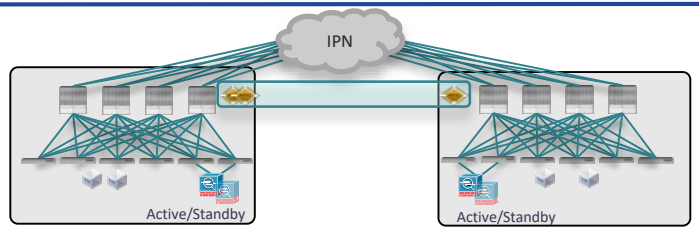

- Independent Active/Standby pairs deployed in separate Pods
- Use of Symmetric PBR to avoid the creation of asymmetric paths crossing different active FW nodes

ACI 3.2

Active/Active Cluster Across Pods
Anycast IP/MAC with PBR

- All the active FW nodes have the same IP/MAC identity, so one of them will be picked

  By default one of the nodes local to a Pod is selected (based on IS-IS metric toward the IP address)

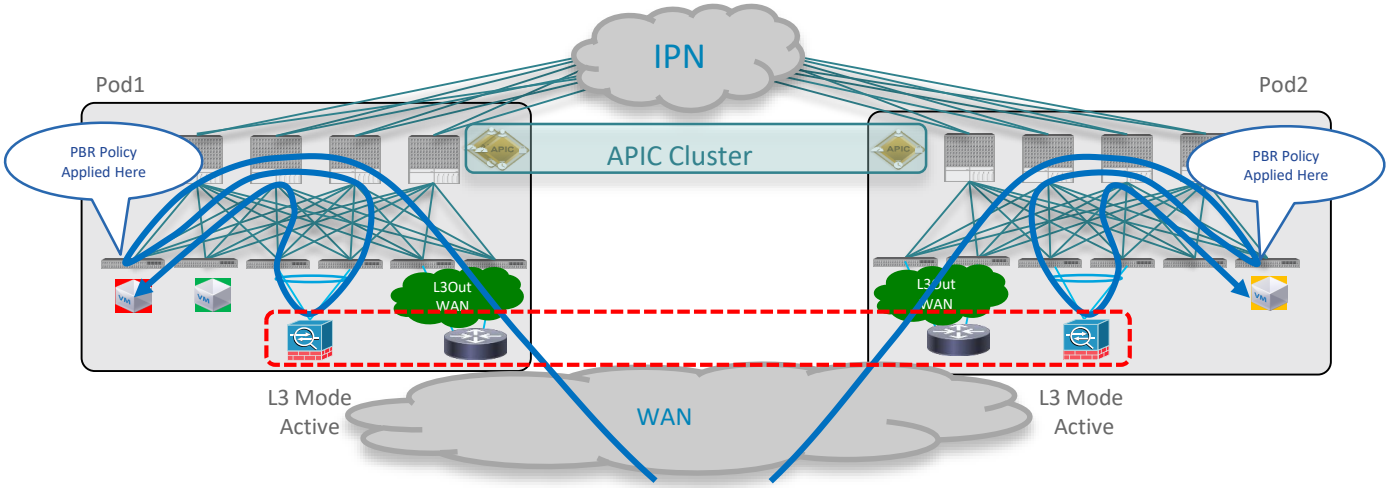

Without Anycast IP/MAC Feature

X

THIS IS NOT WORKING WITHOUT ANYCAST SERVICE

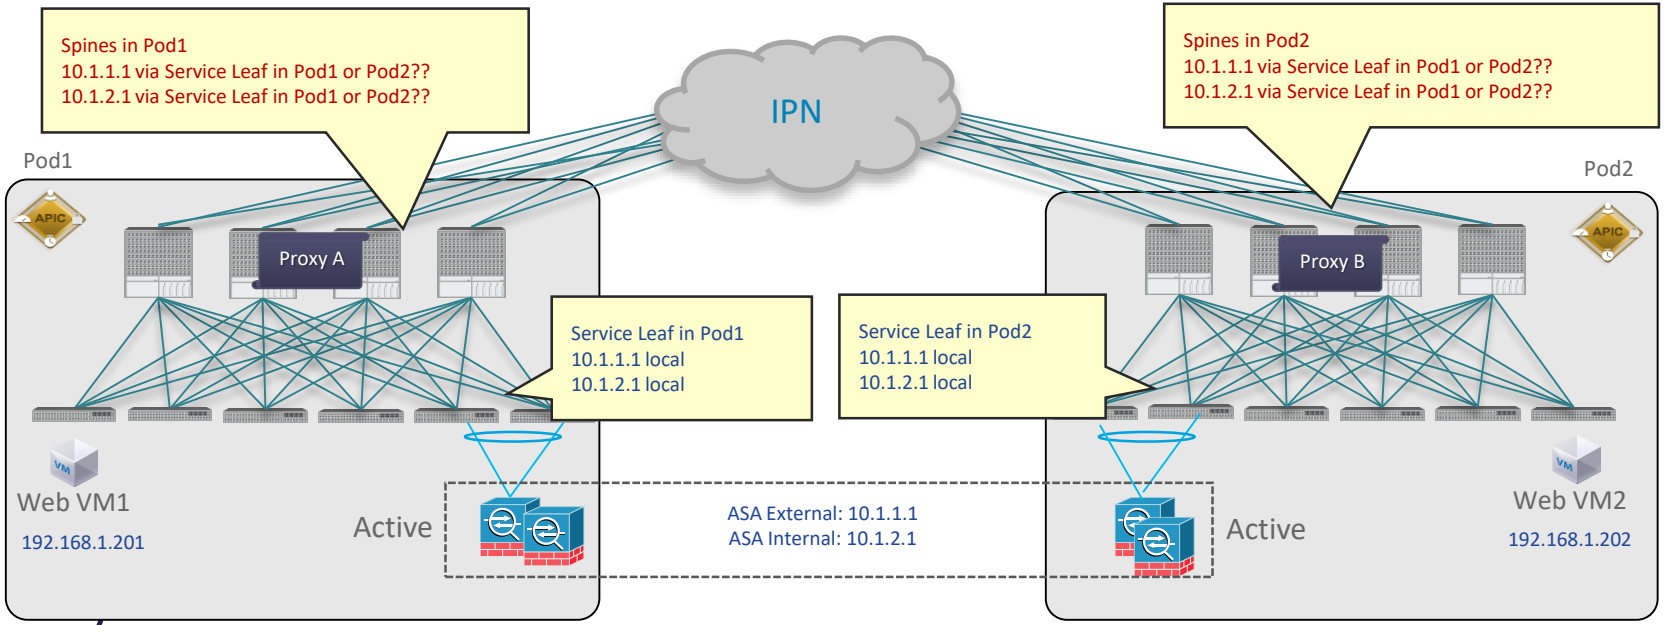

With Anycast IP/MAC Feature

Works with
Any Cast Service starting 3.2

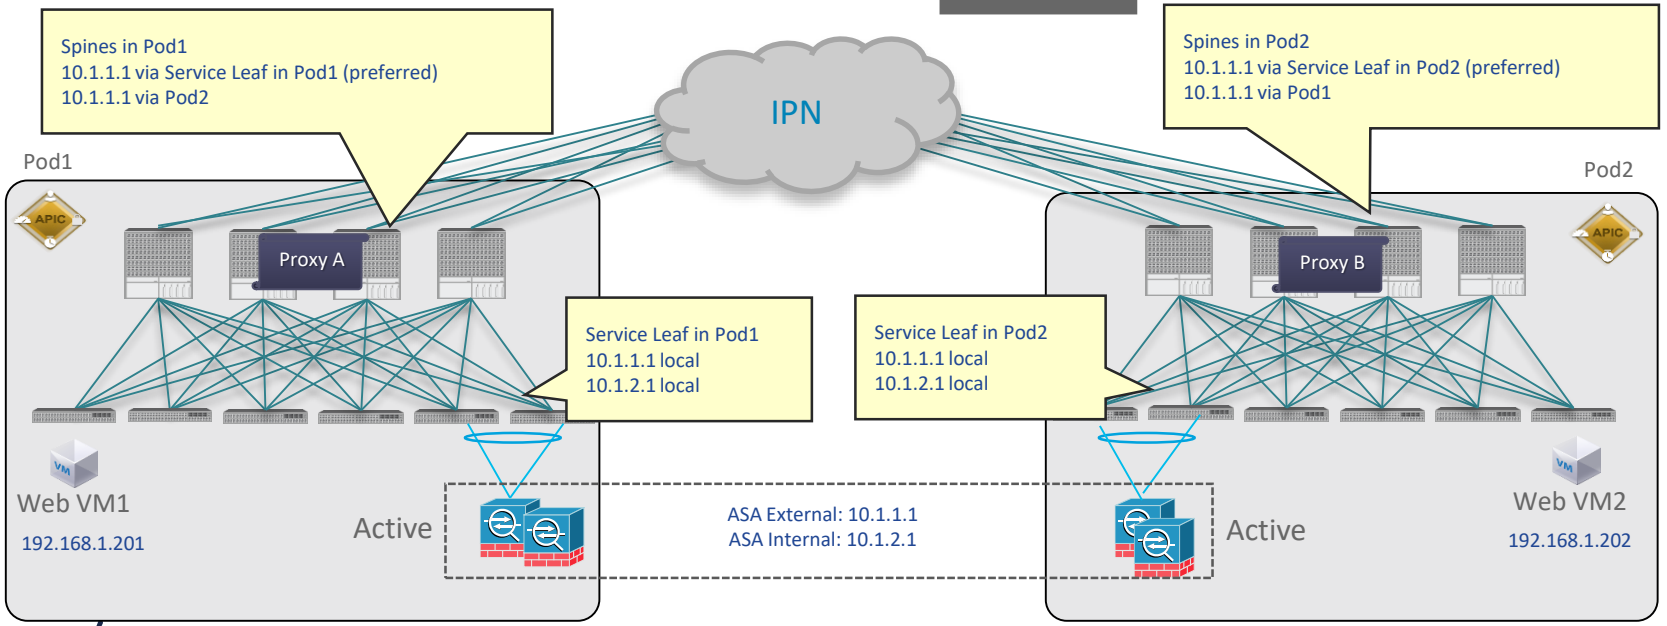

ACI Multi-Site and Network Services
Integration Models

For Your Reference

Deployment options fully supported with ACI Multi-Pod

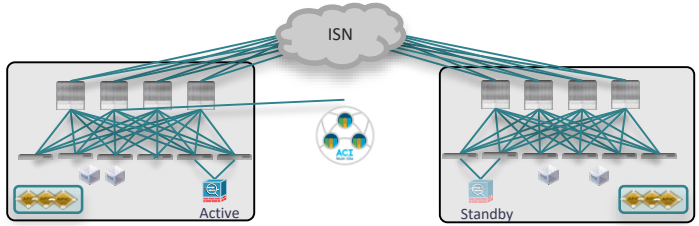

- Active and Standby pair deployed across Pods
- **Currently supported only if the FW is in L2 mode or in L3 mode but acting as default gateway for the endpoints**

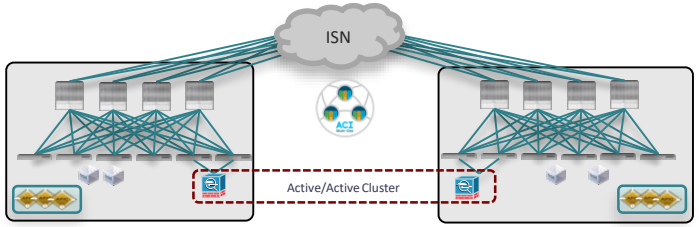

- Active/Active FW cluster nodes stretched across Sites (single logical FW)
- Requires the ability of discovering the same MAC/IP info in separate sites at the same time
- **Not supported**

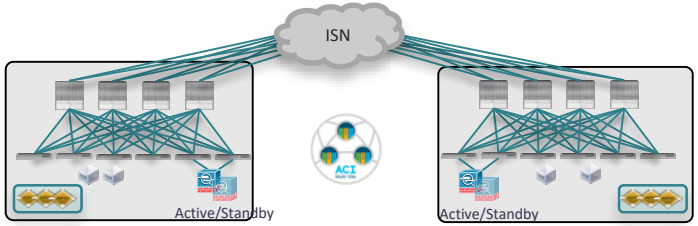

- **Recommended deployment model for ACI Multi-Site**
- Option 1: supported from 3.0 for N-S if the FW is connected in L3 mode to the fabric → mandates the deployment of traffic ingress optimization
- Option 2: supported from 3.2 release with the use of Service Graph with Policy Based Redirection (PBR)

Use of Service Graph and Policy Based Redirection
North-South Communication – Inbound Traffic

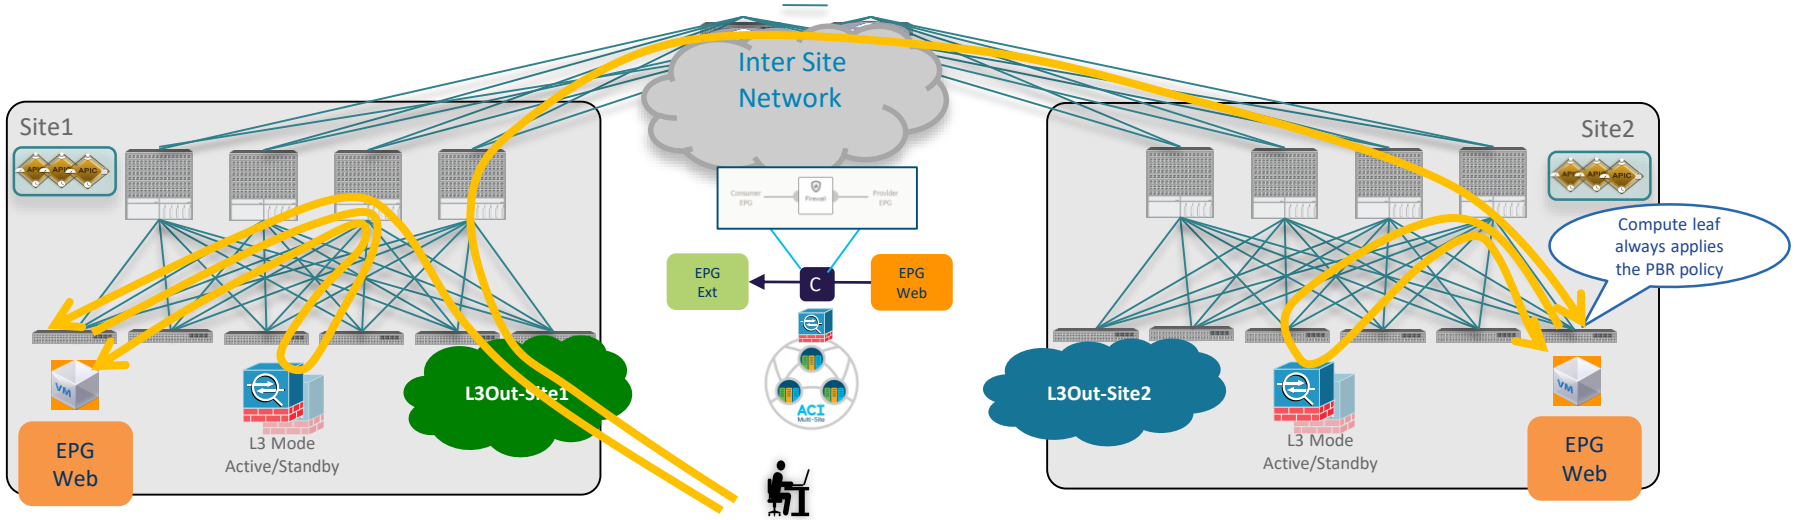

- Inbound traffic can enter any site when destined to a stretched subnet (if ingress optimization is not deployed or possible)
- PBR policy is **always applied on the compute leaf node** where the destination endpoint is connected
  - Requires the VRF to have the Ingress policy enforcement preference and direction
  - Supported only **intra-VRF** in ACI release 4.0.
  - Ext-EPG and Web EPG can indifferently be provider or consumer of the contract

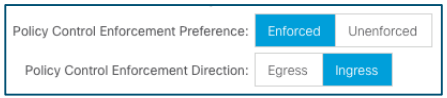

Use of Service Graph and Policy Based Redirection
North-South Communication – Inbound Traffic

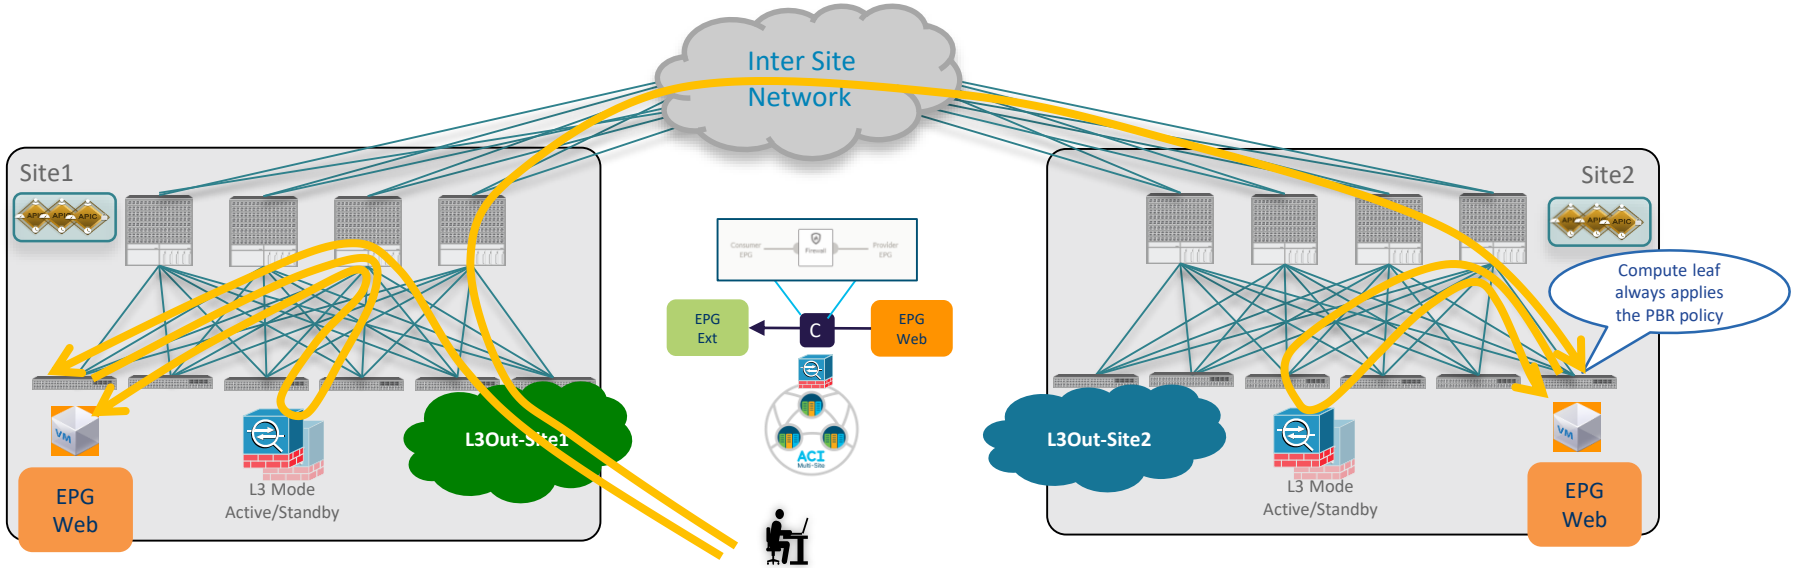

- Inbound traffic can enter any site when destined to a stretched subnet (if ingress optimization is not deployed or possible)
- PBR policy is **always applied on the compute leaf node** where the destination endpoint is connected
  - Requires the VRF to have the Ingress policy enforcement preference and direction
  - Supported only **intra-VRF** in ACI release 4.0.
  - Ext-EPG and Web EPG can indifferently be provider or consumer of the contract

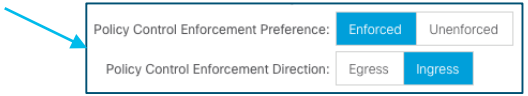

Use of Service Graph and Policy Based Redirection
North-South Communication – Outbound Traffic

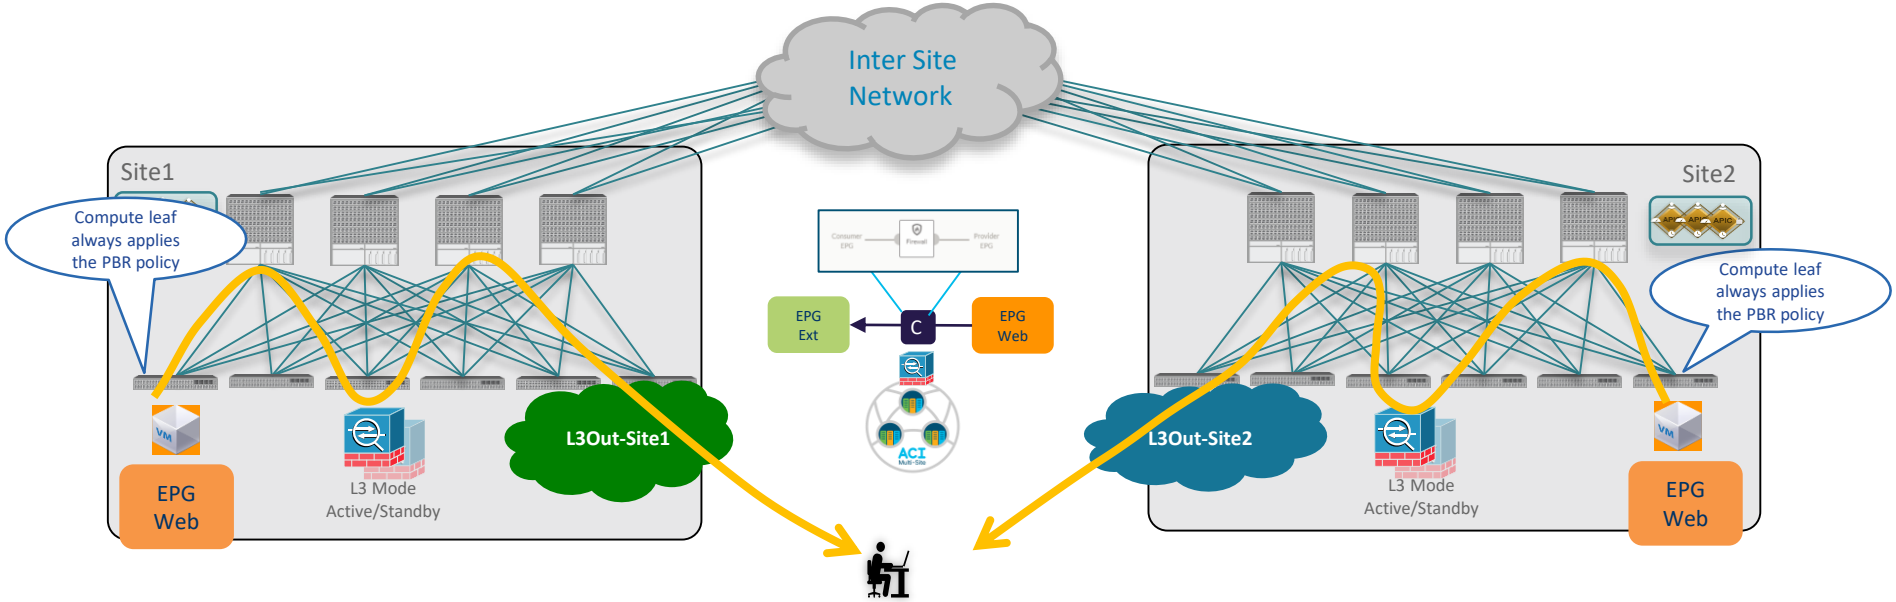

- PBR policy is **always applied on the same leaf** where it was applied for inbound traffic
- Ensures the same service node is selected for both legs of the flow
- Different L3Outs can be used for inbound and outbound directions of the same flow

Summary

- ACI Contract security
- ACI L4-L7 service integration
  - Firewall Design Options
    - Inline FW, FW as gateway, VRF sandwich or PBR
  - Load Balancer Design Options
    - LB as gateway, SNAT or PBR for return traffic
  - Multi-Pod/Multi-Site Design Options

Useful Links

- Service Graph Design with Cisco Application Centric Infrastructure White Paper

  https://www.cisco.com/c/en/us/solutions/collateral/data-center-virtualization/application-centric-infrastructure/white-paper-c11-734298.html

- Cisco Application Centric Infrastructure Policy-Based Redirect Service Graph Design White Paper

  https://www.cisco.com/c/en/us/solutions/data-center-virtualization/application-centric-infrastructure/white-paper-c11-739971.html

- ACI Fabric Endpoint Learning White Paper

  https://www.cisco.com/c/en/us/solutions/collateral/data-center-virtualization/application-centric-infrastructure/white-paper-c11-739989.html

Useful Links

- ACI Multi-pod and Service Node Integration White paper

  https://www.cisco.com/c/en/us/solutions/collateral/data-center-virtualization/application-centric-infrastructure/white-paper-c11-739571.html

- ACI Multi-site and Service Node Integration White paper

  https://www.cisco.com/c/en/us/solutions/collateral/data-center-virtualization/application-centric-infrastructure/white-paper-c11-743107.html

Thank you

You make **possible**