



The bridge to possible

Explore complexities and best practices for deploying applications in multi cluster service mesh

Sundar Srinivasaraghavan – Principal Architect
Ravi Jandyala – Product Management Architect
BRKCLD-2019

CISCO *Live!*

#CiscoLive

Cisco Webex App

Questions?

Use Cisco Webex App to chat with the speaker after the session

How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space


Webex spaces will be moderated by the speaker until June 7, 2024.



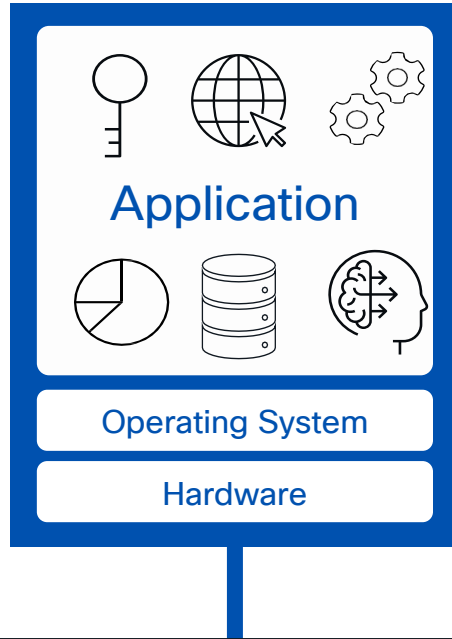
Agenda

- Why do we need Service Mesh ?
- What is Service Mesh ?
- Multi-cluster Service Mesh Deployment Models
- Service Mesh Deployment Challenges
- Application Deployment in Multi-cluster Best Practices
- Demo
- Conclusion


Why do we need Service Mesh ?



Monolithic Application Architecture

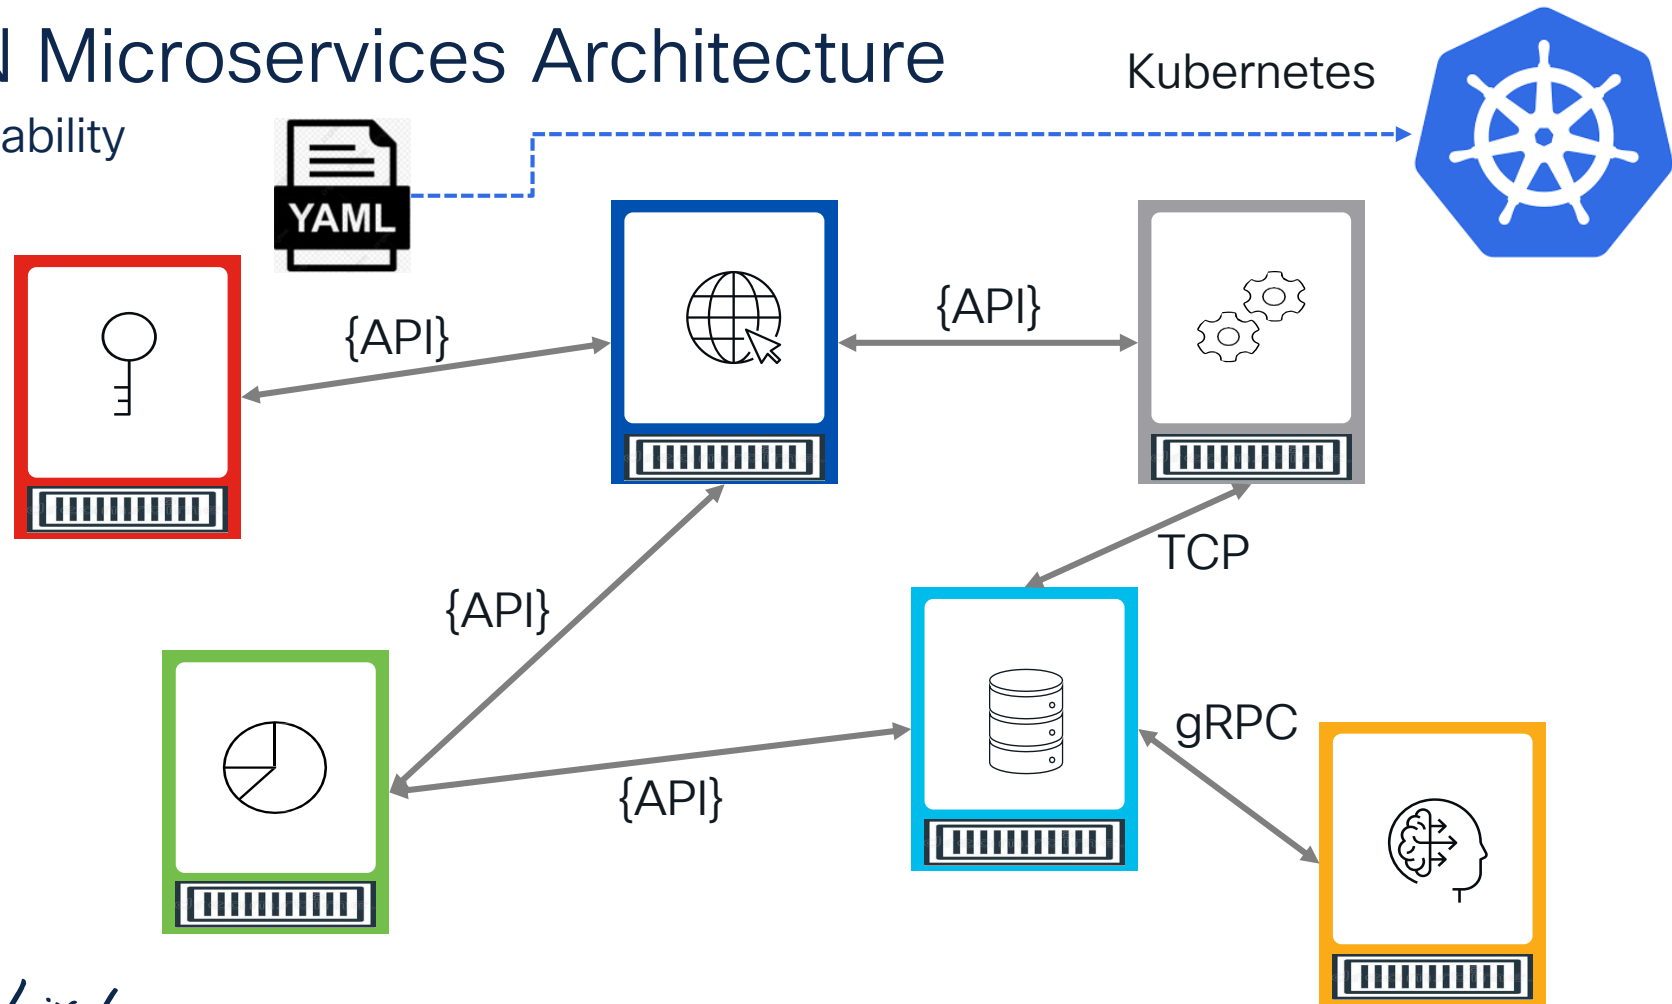


Cloud Native Microservices Architecture









CN Microservices Architecture

Scalability



CN Microservices Architecture

Scalability



Kubernetes




CN Microservices Architecture

Scalability

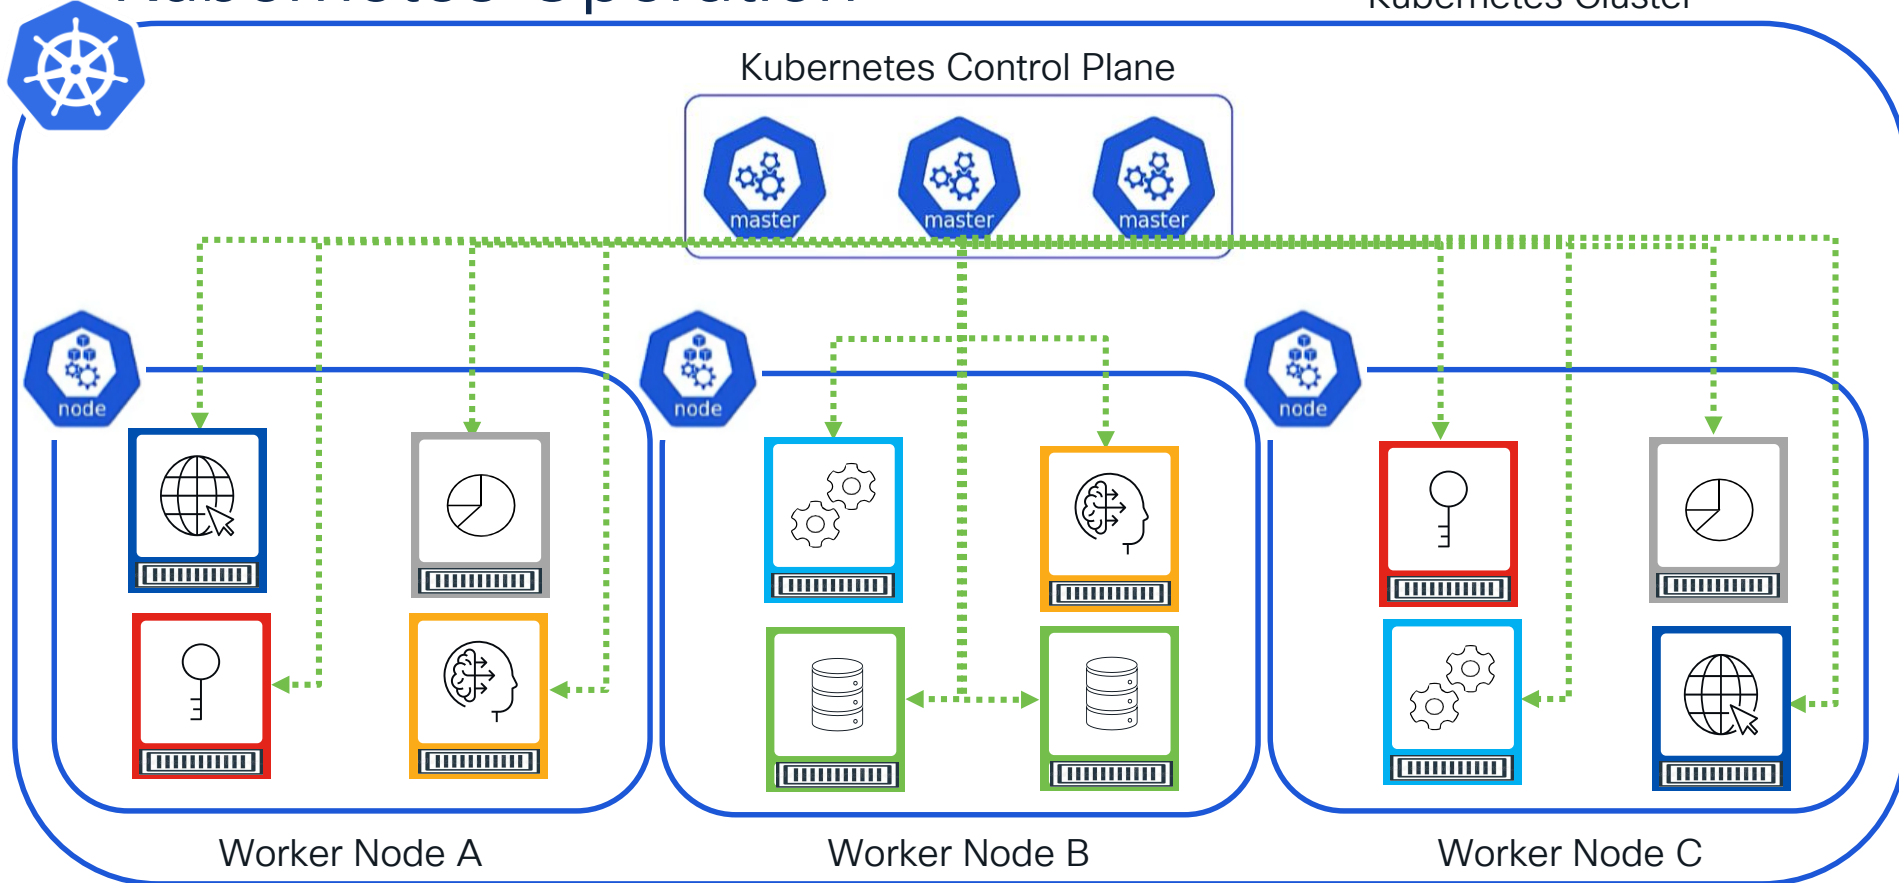


Kubernetes

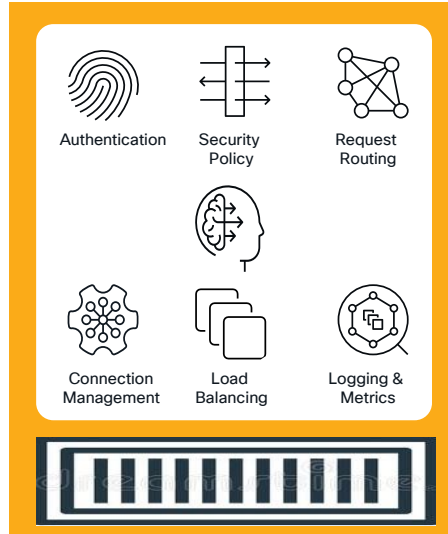



Kubernetes Operation



Kubernetes Cluster




Microservice Common Functions




Microservice Common Functions



CISCO *Live!*




#CiscoLive BRKCLD-2019



© 2024 Cisco and/or its affiliates. All rights reserved. Cisco Public




Service Mesh

A Service Mesh enables you to **connect**, **secure**, **control** and **observe** microservices



Benefits:


- **Consistent development**
- **Consistent deployment**
- **Consistent security** of microservices
- **Scalability** of microservice architecture



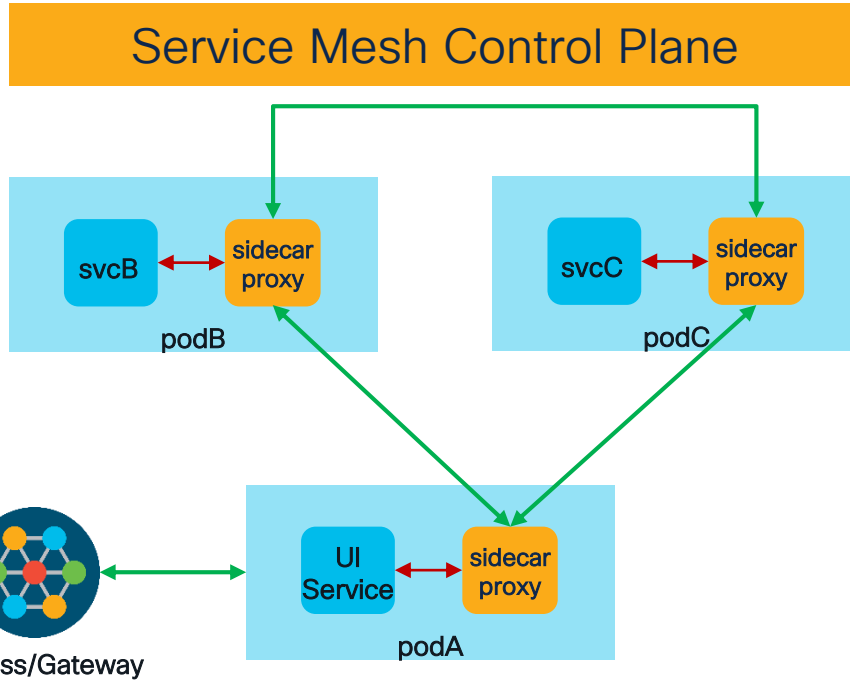
My Application Service Requirements

- I want to deploy a microservice
- I want to deploy using Kubernetes
- I have a bunch of requirements such as the need to handle:
 - Service failures
 - Retries
 - Circuit breaking
 - Topology changes
 - Monitoring
 - Tracing
 - Encryption between services
 - and more

What is a Service Mesh ?




What is a Service Mesh?



- Infrastructure layer for service-to-service communication
- Can use a mesh of sidecar proxies:
 - Can inspect API transactions at Layer 7 and 4 (TCP)
 - Intelligent routing rules can be applied between endpoints

Sidecar Proxies and Service Mesh



- In a generic Kubernetes environment, a containerized application microservice is usually assigned to a **dedicated pod**
- However, several **common service** functions (such as observability, access policy, encryption, load-balancing, traffic management, etc.) can be standardized and enabled by creating a **sidecar** within the pod
- These common services are in turn centrally controlled by the **service mesh** control plane

Istio Overview

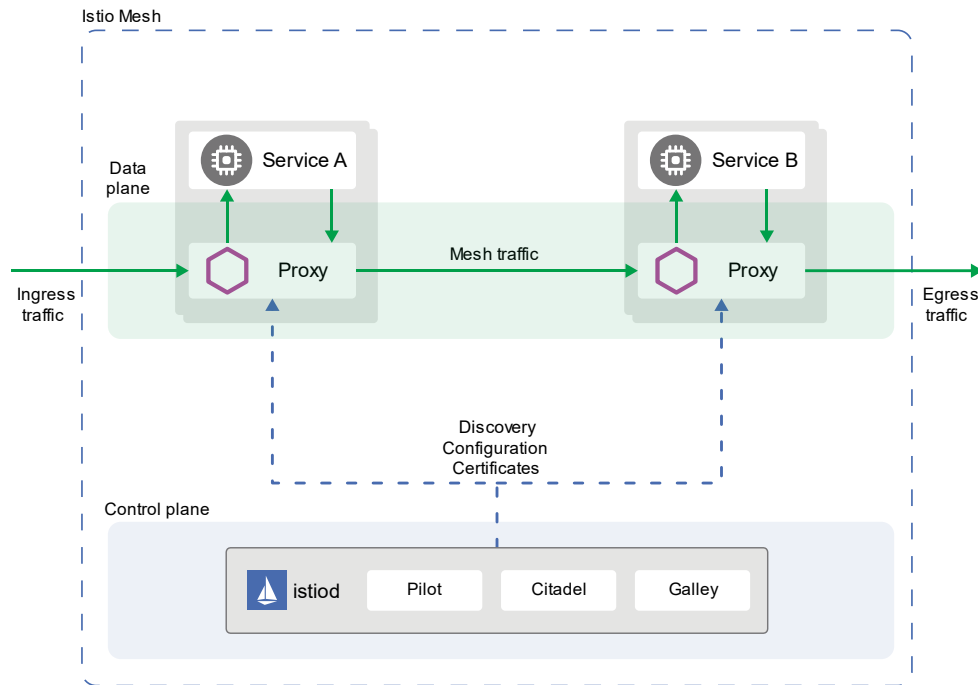
- An open-source project started by Google and IBM with help from the Envoy team at Lyft
 - <https://istio.io/>
 - <https://github.com/istio>
 - <https://www.envoyproxy.io/>
- Automatic **load balancing** for HTTP, gRPC, WebSocket, and TCP traffic
- Robust **multicluster connectivity**
- Fine-grained **control of traffic** behavior with rich routing rules, retries, failovers, and fault injection
- A pluggable policy layer and configuration API supporting access controls, rate limits and quotas
- Automatic **metrics, logs, and traces** for all traffic within a cluster, including cluster ingress and egress
- Secure service-to-service authentication with strong identity assertions between services in a cluster

gRPC - Cross-platform Remote, Open Source, High Performance Remote Procedure Calls


Istio Architecture

<https://istio.io/latest/docs/ops/deployment/architecture/>

- **istiod**
 - **Pilot**
 - Handles **service discovery** and config data
 - Provides the Envoy proxies with the mesh topology and **route rules**
 - **Galley**
 - **Validates** user authored Istio API configuration on behalf of other control plane components
 - Top-level **config ingestion**, processing and distribution
 - **Citadel**
 - Provides **certificates** to the Envoy proxies for authentication and authorization
- **Envoy**
 - A proxy attached to every microservice
 - The **connection point** for a microservice to attach to the mesh




Multi-cluster Service Mesh Deployment Models




Istio Multicluster

- There are several reasons for establishing connectivity between Kubernetes clusters to include:
 - Service load balancing
 - Data replication
 - Service dependencies
 - Partner-provided service connectivity
 - etc..
- <https://istio.io/latest/docs/ops/deployment/deployment-models/>
 - Primary-Remote – single network
 - Primary-Remote – multiple networks
 - Multi-Primary – single network
 - **Multi-Primary – multiple networks**




Single Cluster Deployment



Single Cluster

- Simplest Deployment
- Single Mesh/Control Plane
- Typically over same subnet
- End to end service visibility

Multi Cluster Deployment




- Multiple options
 - Single or Multiple Networks
 - Single or Multiple control planes
 - Zones or Regions
 - Distributed Applications
 - Loadbalancing and Istio Gateways

Single vs Multiple Networks


- Flat Networking
- Single subnet
- No overlapping IP
- Direct reachability between workloads without Istio gateways
- Overlapping IP or VIP ranges for **service endpoints**
- Crossing of administrative boundaries
- Fault tolerance
- Scaling of network addresses
- Compliance with standards that require network segmentation

Service Mesh Deployment Challenges




Service Mesh Deployment Challenges

- Lifecycle management
- Disparate/fragmented observability
- Multi-cluster challenges:
 - Availability
 - Cross-cluster service discovery
 - Inter-cluster traffic management policy
 - Multi-Tenancy



Service Mesh


Service Mesh Observability Challenges



Topology Console



Metrics Utility



Service Mesh




Logging Operator



- Repeat per cluster
- Aggregate & Correlate




Events Tool



Tracing System



Enabling Multi-cluster Service Mesh




- Service meshes can be extended across clusters, such as by extending the control plane from a **primary** cluster to a **remote** cluster
 - Stable IP
 - Expose Control Plane via Istio GW
- Deploying multiple control planes across clusters, which is called a **multi-primary control plane**

Pre-planning

- Network CIDR
- Service Naming
- Establish Trust
- Enable DNS Proxy
- External Load balancer
- Expose services

Application Deployment in Multi-cluster Best Practices



Establish Trust – MUST for mTLS

- Have a dedicated ROOT CA
 - Make sure to download istio binary in both clusters `curl -L https://istio.io/downloadIstio | sh -`
 - Add istio folder to \$PATH variable
- Setup Intermediate CA for each cluster
 - Copy cluster2 CA files to cluster 2
- Create secret in each cluster

```
sudo scp cluster2/*.pem administrator@172.40.140.22:~/istio-1.21.2/certs/cluster2/
```

Cluster 1

```
cd istio-1.21.2
mkdir -p certs
pushd certs
make -f ../tools/certs/Makefile.selfsigned.mk root-ca
make -f ../tools/certs/Makefile.selfsigned.mk cluster1-cacerts
!
make -f ../tools/certs/Makefile.selfsigned.mk cluster2-cacerts
```

Cluster 2

```
cd istio-1.21.2
mkdir -p certs
pushd certs
```

Cluster 1:

```
kubectl create secret generic cacerts -n istio-system --from-file=cluster1/ca-cert.pem --from-file=cluster1/ca-key.pem --from-file=cluster1/root-cert.pem --from-file=cluster1/cert-chain.pem
```

Cluster 2:

```
kubectl create secret generic cacerts -n istio-system --from-file=cluster2/ca-cert.pem --from-file=cluster2/ca-key.pem --from-file=cluster2/root-cert.pem --from-file=cluster2/cert-chain.pem
```

Enable DNS Proxy

- By default, Istio does not enable DNS proxy for services that are exposed to another cluster
- <https://istio.io/latest/docs/ops/configuration/traffic-management/dns-proxy/#getting-started>
- Without enabling DNS proxy, service in cluster 1 may not be resolvable on the 2nd cluster.

Add to the Istio Operator Config

```
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
spec:
  meshConfig:
    defaultConfig:
      proxyMetadata:
        # Enable basic DNS proxying
        ISTIO_META_DNS_CAPTURE: "true"
```

OR edit the config post-deployment

```
# kubectl edit istiocontrolplanes -n istio-system
meshConfig:
  defaultConfig:
    . . .<output_summarized>
    proxyMetadata:
      ISTIO_META_ALS_ENABLED: "true"
      ISTIO_META_DNS_CAPTURE: "true"
      PROXY_CONFIG_XDS_AGENT: "true"
```


Application Load Balancers

- For Ingress, Egress and Eastwest gateways
- Typically available in Public Clouds
- For on-prem clusters, make sure to deploy load balancer for each cluster. For example MetalLB

kubectl apply -f

<https://raw.githubusercontent.com/metallb/metallb/v0.14.5/config/manifests/metallb-native.yaml>

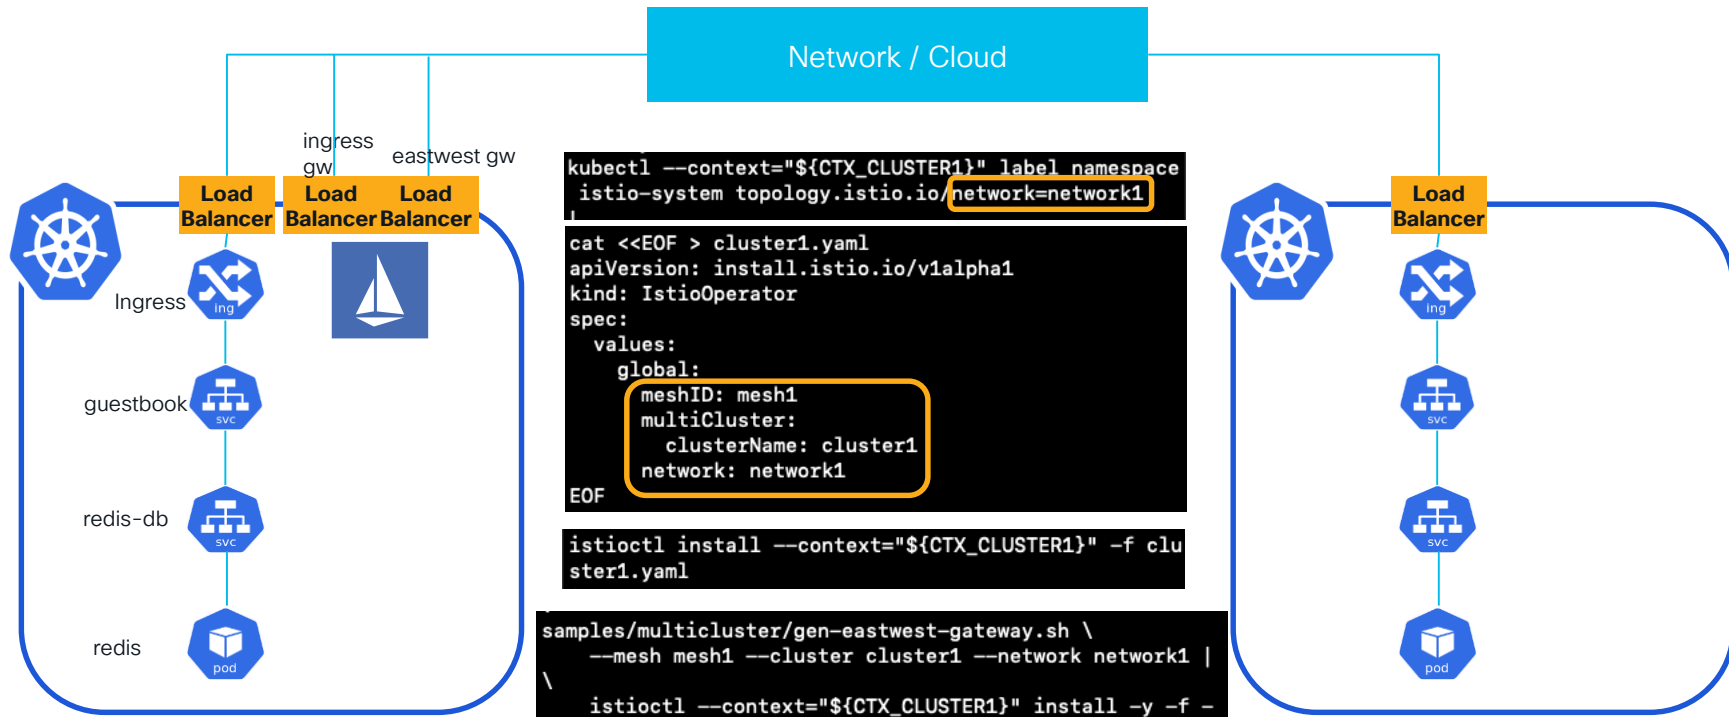
- Define IP Address Pools for each cluster
- Make sure these IPs are externally reachable for the GWs

```
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  name: lb-pool
  namespace: metallb-system
spec:
  addresses:
    - 172.40.143.181-172.40.143.190
```

```
apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
  name: example
  namespace: metallb-system
spec:
  ipAddressPools:
    - lb-pool
```

Istio Multi-Primary Deployment

Initialize Cluster1



cluster1

context: cluster1-admin@cluster1

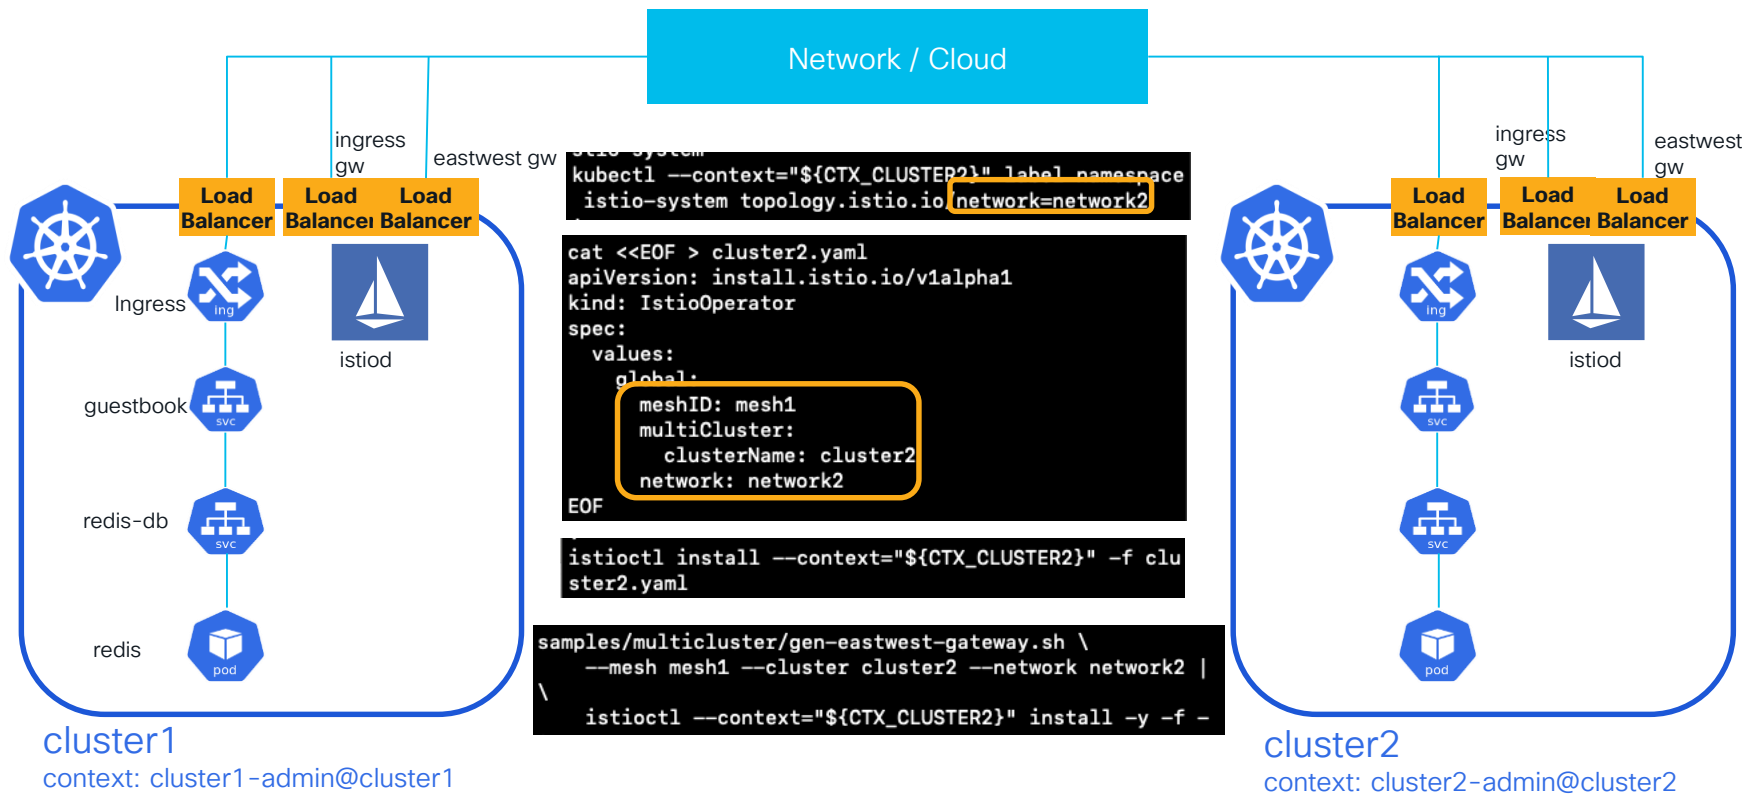
CISCO Live!

cluster2

context: cluster2-admin@cluster2

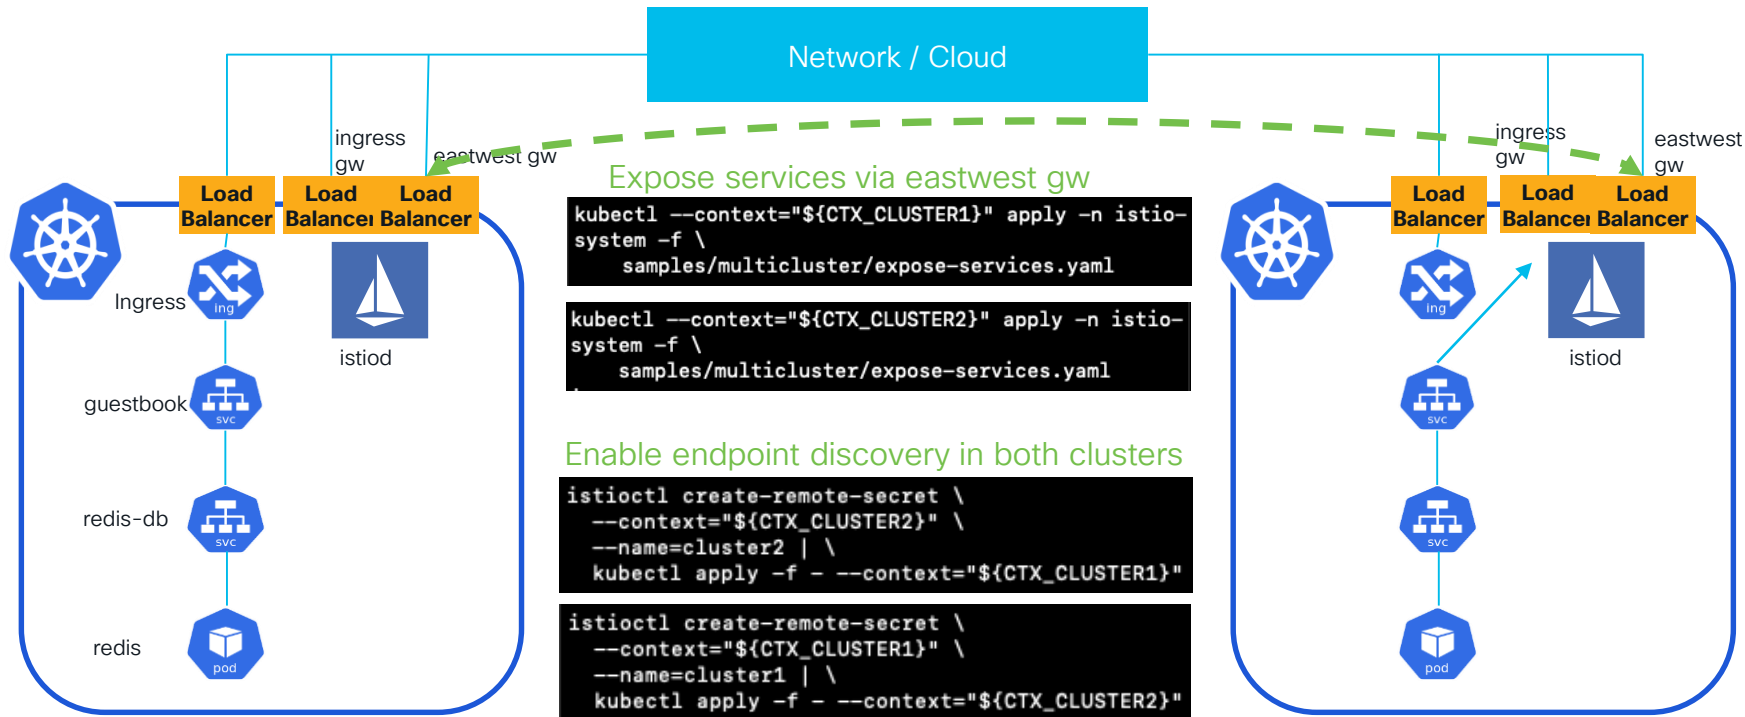
Istio Multi-Primary Deployment

Initialize Cluster2



Istio Multi-Primary Deployment

Expose services and Enable endpoint discovery



cluster1

context: cluster1-admin@cluster1

CISCO *Live!*


cluster2

context: cluster2-admin@cluster2

Mesh status – Cluster 1

```
administrator@cl1-istio-master:~/istio-1.21.2$ istioctl remote-clusters -c ~/.kube/config
```

NAME	SECRET	STATUS	ISTIOD
cluster2	istio-system/istio-remote-secret-cluster2	synced	istiod-6696b6844d-cwj44



```
administrator@cl1-istio-master:~/istio-1.21.2$ kubectl get pods -n istio-system
```

NAME	READY	STATUS	RESTARTS	AGE
istio-eastwestgateway-c89658c74-6sq8b	1/1	Running	0	103m
istio-ingressgateway-7b4fdf6d69-lj2ds	1/1	Running	0	104m
istiod-6696b6844d-cwj44	1/1	Running	0	104m

```
administrator@cl1-istio-master:~/istio-1.21.2$
```

```
administrator@cl1-istio-master:~/istio-1.21.2$
```


```
administrator@cl1-istio-master:~/istio-1.21.2$ kubectl get svc -n istio-system
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
istio-eastwestgateway	LoadBalancer	10.100.159.133	172.40.143.182	15021:31046/TCP,15443:32163/TCP,15012:30919/TCP,15017:32420/TCP	103m
istio-ingressgateway	LoadBalancer	10.102.23.142	172.40.143.181	15021:31555/TCP,80:30553/TCP,443:32092/TCP	105m
istiod	ClusterIP	10.105.208.143	<none>	15010/TCP,15012/TCP,443/TCP,15014/TCP	105m

```
administrator@cl1-istio-master:~/istio-1.21.2$
```

Mesh status – Cluster 2

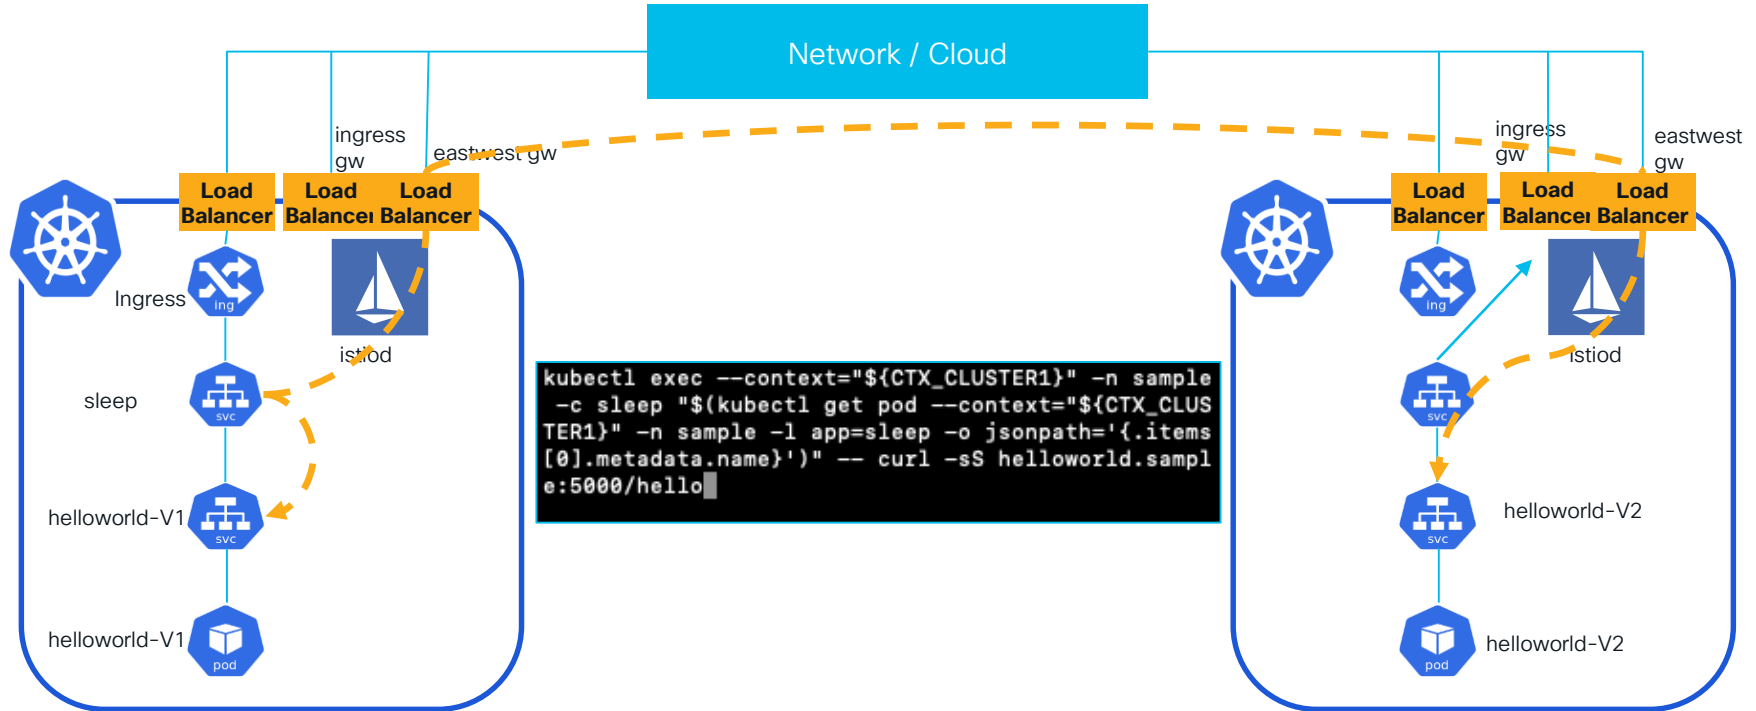
```
administrator@multi-primary-master:~/istio-1.21.2$ istioctl remote-clusters
NAME          SECRET                                STATUS    ISTIOD
cluster1      istio-system/istio-remote-secret-cluster1  synced    istiod-6c4d88c6d7-wpzz8
```



```
administrator@multi-primary-master:~/istio-1.21.2$ kubectl get pods -n istio-system
NAME                                READY   STATUS    RESTARTS   AGE
istio-eastwestgateway-7cbc7b94b8-pmkc7  1/1     Running   0           56m
istio-ingressgateway-6dbdfffc56f-v8vz8  1/1     Running   0           58m
istiod-6c4d88c6d7-wpzz8                1/1     Running   0           58m
administrator@multi-primary-master:~/istio-1.21.2$
administrator@multi-primary-master:~/istio-1.21.2$
administrator@multi-primary-master:~/istio-1.21.2$ kubectl get svc -n istio-system
NAME                                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)                                     AGE
istio-eastwestgateway               LoadBalancer  10.96.31.100     172.40.143.192   15021:31590/TCP,15443:31841/TCP,15012:31272/TCP,15017:32063/TCP  56m
istio-ingressgateway                LoadBalancer  10.108.147.230   172.40.143.191   15021:32147/TCP,80:31216/TCP,443:32293/TCP  58m
istiod                              ClusterIP      10.108.239.229   <none>           15010/TCP,15012/TCP,443/TCP,15014/TCP  58m
administrator@multi-primary-master:~/istio-1.21.2$
```

Deploy Application

Validate cross cluster connectivity



cluster1

context: cluster1-admin@cluster1

CISCO *Live!*

cluster2

context: cluster2-admin@cluster2

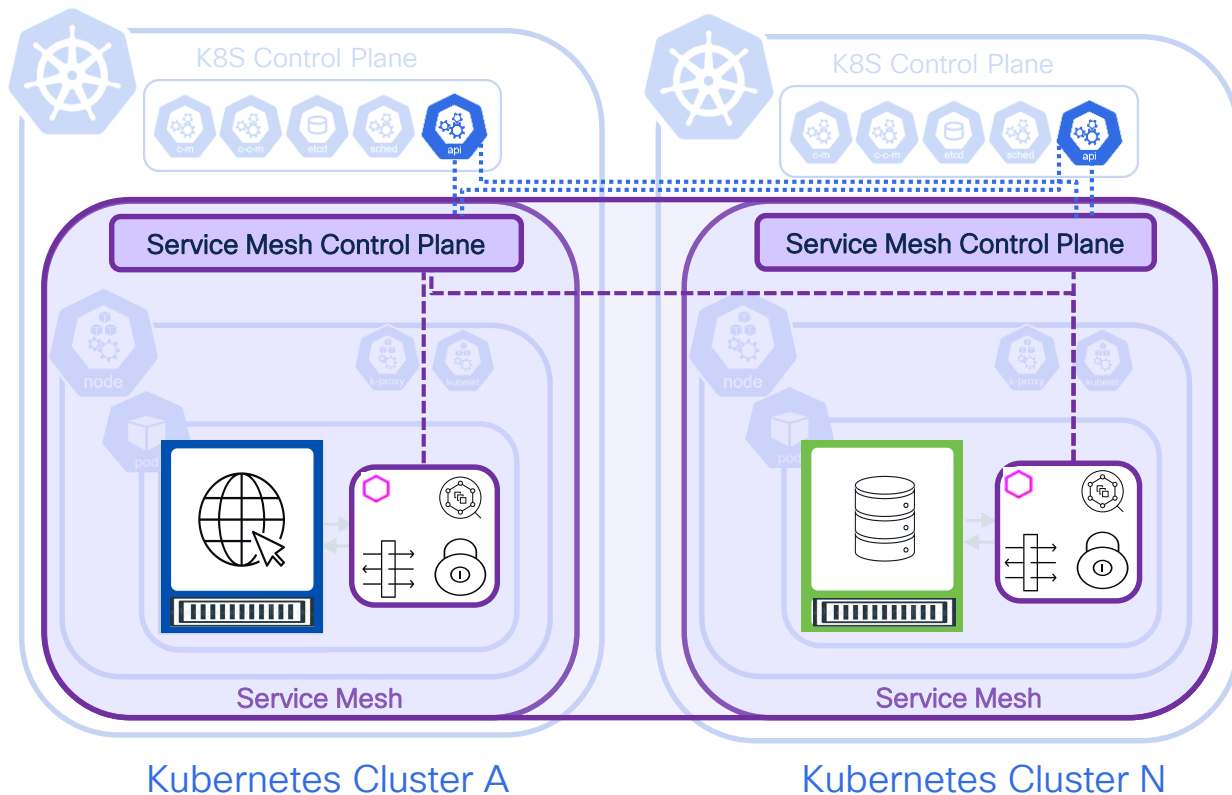
Deploy application in Multi-cluster

```
administrator@c11-istio-master:~/istio-1.21.2$ kubectl get pods -n sample -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE             NOMINATED NODE   READINESS GATES
helloworld-v1-867747c89-wzjch        2/2     Running   0           9m43s  192.168.184.139  c11-istio-work2   <none>           <none>
sleep-7656cf8794-qmx9p              2/2     Running   0           8m2s   192.168.138.131  c11-istio-work3   <none>           <none>

administrator@c11-istio-master:~/istio-1.21.2$
administrator@c11-istio-master:~/istio-1.21.2$
administrator@c11-istio-master:~/istio-1.21.2$ kubectl get svc -n sample
NAME      TYPE      CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
helloworld  ClusterIP  10.105.200.149 <none>         5000/TCP   11m
sleep      ClusterIP  10.96.147.41   <none>         80/TCP     8m15s

administrator@c11-istio-master:~/istio-1.21.2$
administrator@c11-istio-master:~/istio-1.21.2$
administrator@c11-istio-master:~/istio-1.21.2$ kubectl exec -n sample -c sleep "$(kubectl get pod --context="${CTX_CLUSTER1}" -n sample -l app=sleep -o jsonpath='{.items[0].metadata.name}')" -- curl -s helloworld.sample:5000/hello
Hello version: v1, instance: helloworld-v1-867747c89-wzjch
administrator@c11-istio-master:~/istio-1.21.2$ kubectl exec -n sample -c sleep "$(kubectl get pod --context="${CTX_CLUSTER1}" -n sample -l app=sleep -o jsonpath='{.items[0].metadata.name}')" -- curl -s helloworld.sample:5000/hello
Hello version: v1, instance: helloworld-v1-867747c89-wzjch
administrator@c11-istio-master:~/istio-1.21.2$ kubectl exec -n sample -c sleep "$(kubectl get pod --context="${CTX_CLUSTER1}" -n sample -l app=sleep -o jsonpath='{.items[0].metadata.name}')" -- curl -s helloworld.sample:5000/hello
Hello version: v2, instance: helloworld-v2-7f46498c69-fg5p7
administrator@c11-istio-master:~/istio-1.21.2$ kubectl exec -n sample -c sleep "$(kubectl get pod --context="${CTX_CLUSTER1}" -n sample -l app=sleep -o jsonpath='{.items[0].metadata.name}')" -- curl -s helloworld.sample:5000/hello
Hello version: v1, instance: helloworld-v1-867747c89-wzjch
administrator@c11-istio-master:~/istio-1.21.2$ kubectl exec -n sample -c sleep "$(kubectl get pod --context="${CTX_CLUSTER1}" -n sample -l app=sleep -o jsonpath='{.items[0].metadata.name}')" -- curl -s helloworld.sample:5000/hello
Hello version: v1, instance: helloworld-v1-867747c89-wzjch
administrator@c11-istio-master:~/istio-1.21.2$ kubectl exec -n sample -c sleep "$(kubectl get pod --context="${CTX_CLUSTER1}" -n sample -l app=sleep -o jsonpath='{.items[0].metadata.name}')" -- curl -s helloworld.sample:5000/hello
Hello version: v2, instance: helloworld-v2-7f46498c69-fg5p7
administrator@c11-istio-master:~/istio-1.21.2$
```


Enabling a Multi-Primary Control Plane

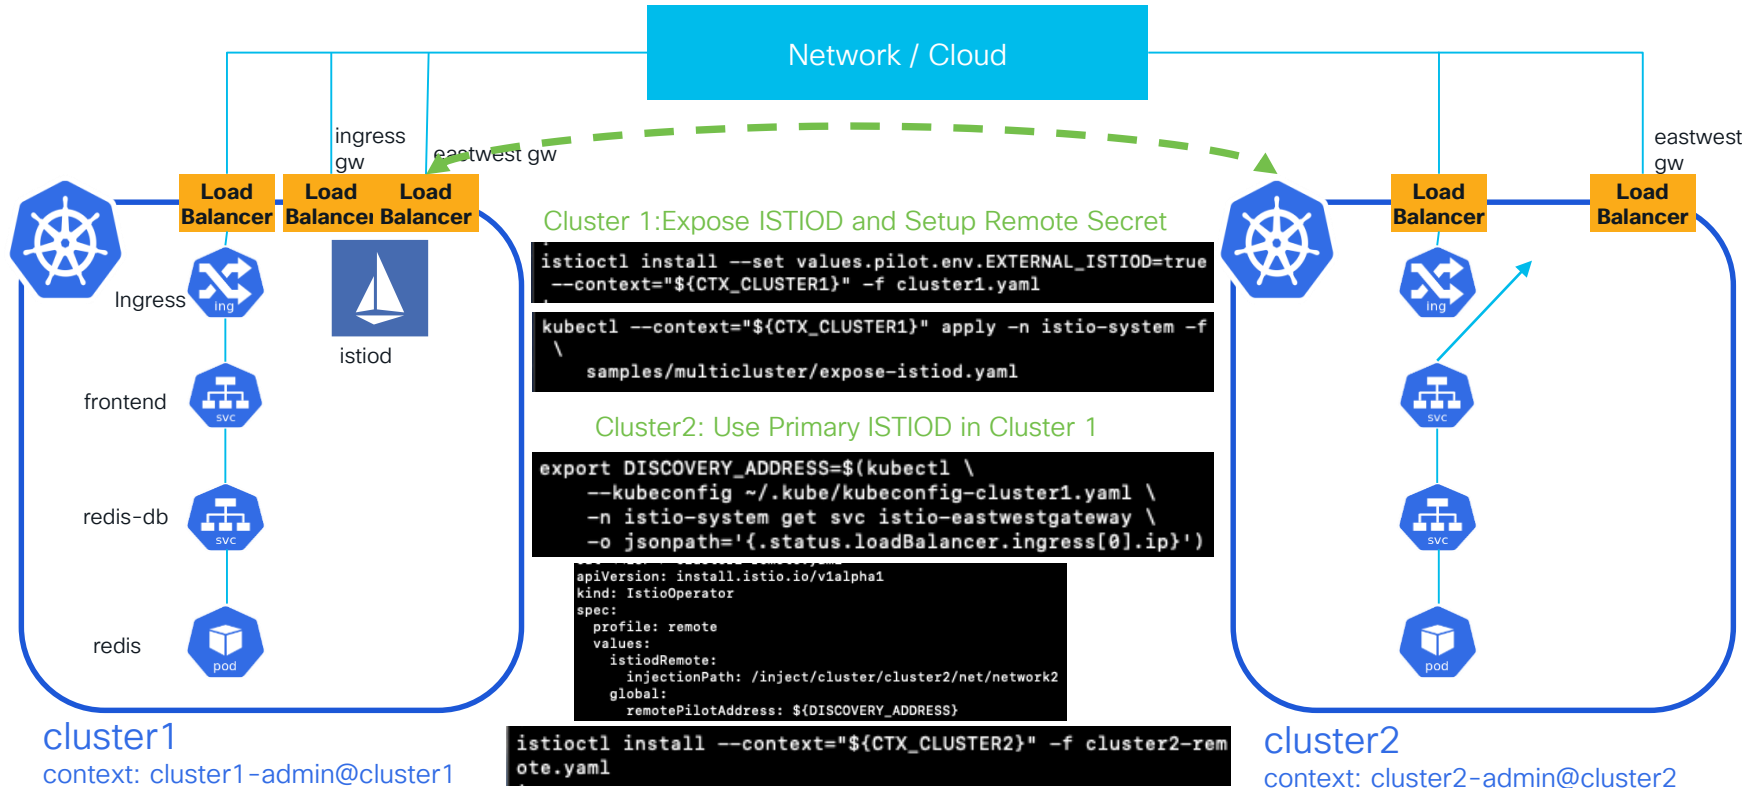


Benefits :


- Limited Scope
 - Cluster specific Configuration changes
 - Cluster specific impact if control plane is unavailable
 - Controlled Configuration rollout
- Service isolation/limited visibility
- High availability
- Cross-cluster endpoint/service discovery

Istio Primary-Remote Deployment

Key Differences



Demo



Conclusion

- Deployment options exist for Multi cluster Service Mesh
- Single vs Multi Network Deployment
- Certificate Setup
- DNS Proxy Setup
- Load balancers and Gateways
- Cross-cluster Service Discovery

Complete Your Session Evaluations



Complete a minimum of 4 session surveys and the Overall Event Survey to be entered in a drawing to **win 1 of 5 full conference passes** to Cisco Live 2025.



Earn 100 points per survey completed and compete on the Cisco Live Challenge leaderboard.



Level up and earn **exclusive prizes!**



Complete your surveys in the **Cisco Live mobile app**.

Continue your education

- All about Istio Service Mesh:
<https://istio.io/latest/docs/>
- Istio Examples:
<https://istio.io/latest/docs/examples/microservices-istio/>
- Attend the interactive education with DevNet, Capture the Flag, and Walk-in Labs
- Visit the On-Demand Library for more sessions at
www.CiscoLive.com/on-demand

Contact me at: Webex App -
<https://ciscolive.ciscoevents.com/ciscolivebot/#BRKCLD-2019>



The bridge to possible

Thank you

CISCO *Live!*

#CiscoLive

Verification of Endpoints in Cluster 1

```
[administrator@multi-primary-master:~/istio-1.21.2$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
loki-0	0/2	Pending	0	10d
ratings-v1-6484c4d9bb-25cnf	2/2	Running	0	79m
reviews-v3-5b9bd44f4-ztxsd	2/2	Running	0	79m

```
administrator@multi-primary-master:~/istio-1.21.2$ istioctl proxy-config endpoints ratings-v1-6484c4d9bb-25cnf --cluster "outbound|9080||productpage.default.svc.cluster.local"
```

ENDPOINT	STATUS	OUTLIER CHECK	CLUSTER
172.40.143.182:15443	HEALTHY	OK	outbound 9080 productpage.default.svc.cluster.local

```
administrator@multi-primary-master:~/istio-1.21.2$ istioctl proxy-config endpoints ratings-v1-6484c4d9bb-25cnf --cluster "outbound|9080||details.default.svc.cluster.local"
```

ENDPOINT	STATUS	OUTLIER CHECK	CLUSTER
172.40.143.182:15443	HEALTHY	OK	outbound 9080 details.default.svc.cluster.local

```
administrator@multi-primary-master:~/istio-1.21.2$ istioctl proxy-config endpoints ratings-v1-6484c4d9bb-25cnf --cluster "outbound|9080||reviews.default.svc.cluster.local"
```

ENDPOINT	STATUS	OUTLIER CHECK	CLUSTER
172.40.143.182:15443	HEALTHY	OK	outbound 9080 reviews.default.svc.cluster.local
192.168.179.215:9080	HEALTHY	OK	outbound 9080 reviews.default.svc.cluster.local