

The background features a vibrant, multi-colored abstract design. On the left, there are overlapping, wavy, organic shapes in shades of red, orange, and yellow. On the right, a bright white light source emits a series of sharp, radiating lines in various colors, including blue, green, and yellow, creating a sunburst or starburst effect. The overall color palette is a spectrum of rainbow colors.

cisco *Live!*

Let's go

#CiscoLive



The bridge to possible

Cloud Native Observability

Shannon McFarland- CCIE #5245
Distinguished Engineer, Emerging Technologies & Incubation
@eyepv6
BRKCLD-2158



Cisco Webex App

Questions?

Use Cisco Webex App to chat with the speaker after the session

How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space

Webex spaces will be moderated by the speaker until June 9, 2023.



<https://ciscolive.ciscoevents.com/ciscolivebot/#BRKCLD-2158>

Agenda

- What is Cloud Native Observability (CNO)?
- What is M.E.L.T?
 - Metrics
 - Events (and Alerts)
 - Logs
 - Traces
- Service Meshes – Built-in CNO
- Cisco Solutions for Observability
- Conclusion

What is Cloud Native Observability?

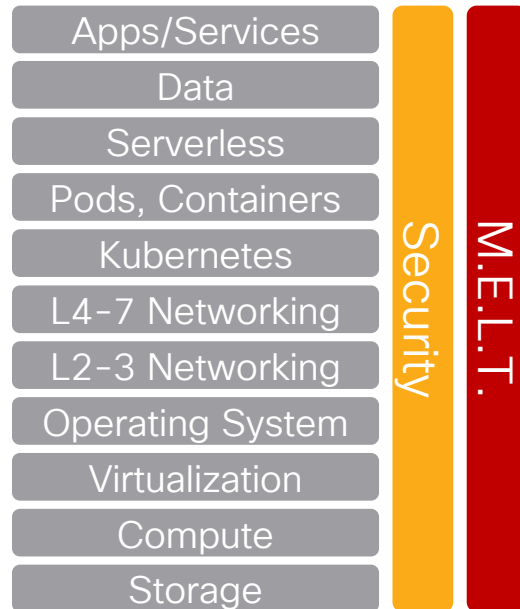
What is Cloud Native?

- “Cloud native technologies empower organizations to build and run scalable applications in modern, dynamic environments such as public, private, and hybrid clouds. **Containers**, service meshes, **microservices**, **immutable** infrastructure, and declarative **APIs** exemplify this approach.
- These techniques enable loosely coupled systems that are **resilient**, **manageable**, and **observable**. Combined with robust **automation**, they allow engineers to make **high-impact changes frequently and predictably with minimal toil.**” – CNCF
- <https://github.com/cncf/toc/blob/main/DEFINITION.md>
- Other Cloud Native criteria include:
 - Elasticity/Horizontal Scaling of Live Services
 - Leveraging Common Frameworks (Application service leverages a Service Mesh)

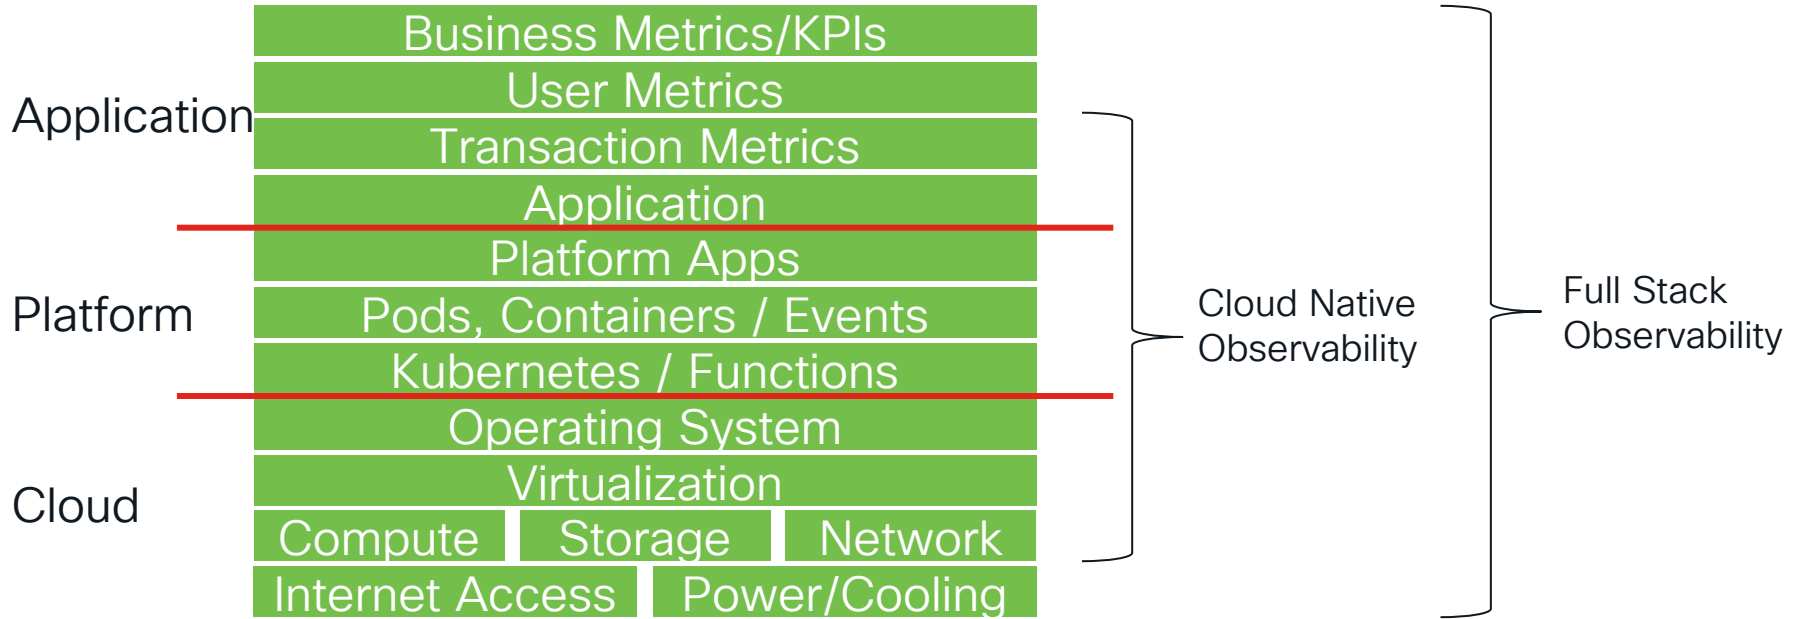
Persona/Role – Moving from Monitoring to Observability

- Platform Operator/Developer
 - They may want to see the application through a Layer 7 (HTTP/gRPC**) lens
 - What is the latency/RPS/memory/CPU for each service component?
 - Where is the bottleneck?
 - Does each component adhere to an SLO?
- Data Scientist/Data Engineer
 - They may want to see very specific parts of the streaming data pipeline that is sub-component of the overall application
- CISO/Security Architect/DevSecOps
 - They may want to see the same application view as the developer, but with a specific focus on CI/CD-centric security (image scanning, code scanning) and internal/external API security

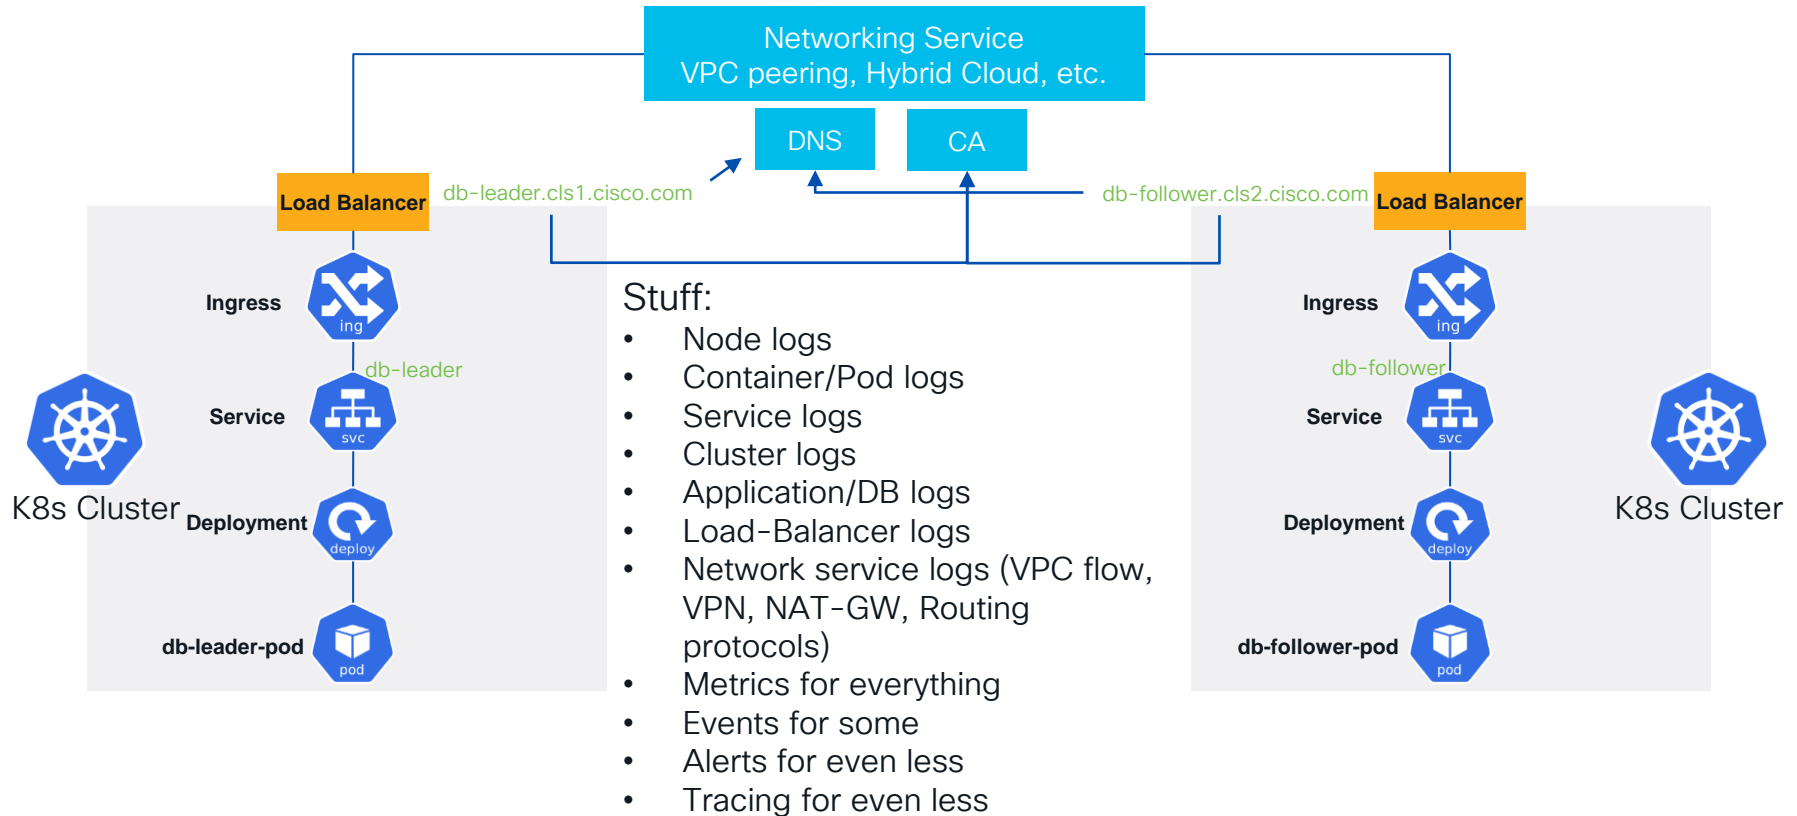
Cloud Native Observability



An Example: Cloud Native Stack



Let's look at a topology



“Full-Stack Observability” adds to traditional monitoring to support seamless digital experiences for modern architectures and teams

Monitoring

Passive detection of (sub)system
“health” issues



Full-Stack Observability

Seamless digital experiences for modern application
architectures and teams

- Encompasses full spectrum of **Visibility – Insights – Action** capabilities to actively understand issues and drive remediation
- Broad coverage across **application, infrastructure, networking, and security stack** with rich, real-time correlation across domains
 - Includes all systems impacting the digital experience for users
- Focused on **actively understanding and remediating issues**, enhanced with ML/AI to ultimately predict issues before they occur
- Facilitates **collaboration across modern teams** (e.g., DevOps / SRE) to achieve common objectives (e.g., SLOs)

What is M.E.L.T.?

Metrics

Metrics

Expose

Collect

Store

Query

Collect/Measure Data at Regular Intervals

- Expose
 - Infrastructure:
 - AWS Elastic Compute 2 (EC2) VM hosting Elastic Kubernetes Service (EKS) worker node - CPU, Memory, Storage, Network
 - Application:
 - NGINX, DB
- Collect
 - Scrape from exposed sources
- Store
 - Time-Series Database (TSDB)
- Query
 - PromQL, MQL (monitoring query language), MetricsQL (VictoriaMetrics)

Metrics

Expose

Collect

Store


Query

Collect/Measure Data at Regular Intervals

- Prometheus (<https://prometheus.io/>) – Open-source event monitoring and alerting tool
- Thanos (<https://thanos.io/>) – Adds high availability, long-term storage and global query capabilities for Prometheus
- Cortex metrics (<https://cortexmetrics.io/>) – Adds high availability, multi-tenant, horizontally scalable and long-term storage capabilities for Prometheus
- Grafana (<https://grafana.com/>) – Visualization of metrics, logs and events from MANY data sources to include Prometheus
- AWS CloudWatch Metrics
- Google Cloud Metrics
- Microsoft Azure Monitor Metrics

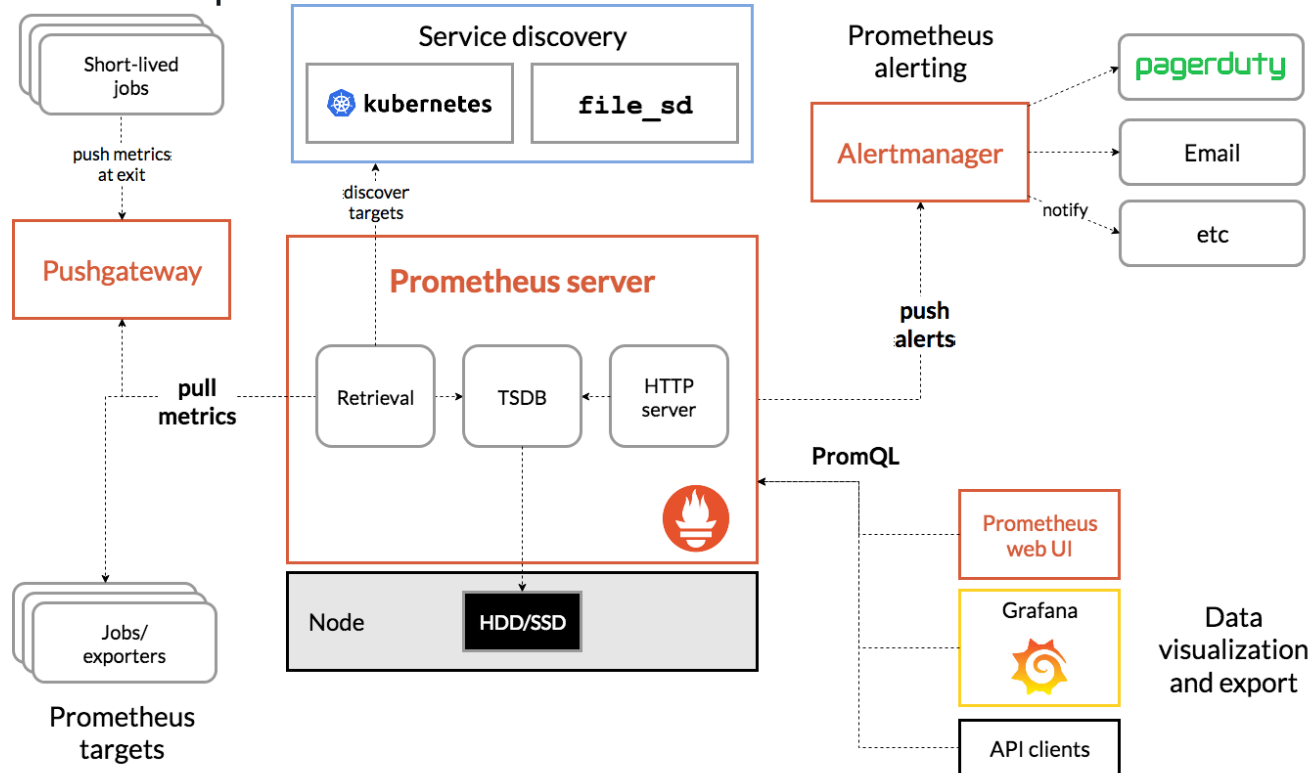
Metrics

Common Architectural Components



Metrics

Prometheus Example



Events and Alerts


Events

- Metrics and Events are two different data types:
 - Metrics = regular/predictable data
 - Events = irregular/unpredictable data
 - Scheduled or unscheduled state changes
- Kubernetes event example:

```
# kubectl get events --field-selector reason=NodeHasSufficientMemory
LAST SEEN   TYPE      REASON              OBJECT               MESSAGE
57m         Normal    NodeHasSufficientMemory node/<omitted>       Node <omitted> status is now: NodeHasSufficientMemory
```

Events

- Events can be paired with other toolsets to provide a robust Event<>Action framework
 - Metrics, Pub/Sub, AI/ML, DevOps, etc.
- Event-Driven Architecture (EDA):
 - KEDA – Kubernetes-Based Event-Driven Autoscaler: <https://keda.sh/>
 - AWS EventBridge
 - Many, many more



Events – AWS EventBridge

Amazon EventBridge



Getting started

Event buses

Rules

Archives

Replays

Integration

Partner event sources

API destinations

Schema registry

Schemas

Documentation

Amazon EventBridge > Rules

Rules

A rule watches for specific types of events. When a matching event occurs, the event is routed to the targets associated with the rule. A rule can be associated with one or more targets.

Select event bus

Event bus

Select or enter event bus name

default

Rules (2/2)



Edit

Delete

Enable

Create rule

Find rules

Any status ▼

< 1 ... >

	Name ▲	Status ▼	Type ▼	Description
<input type="radio"/>	CL_Event_State_Change	✔ Enabled	Standard	
<input type="radio"/>	Schemas-events-event-bus-default	✔ Enabled	Managed	This rule is used to route Events to Schema Discoverer

Events – AWS EventBridge

CloudWatch > Log groups > /aws/events/cl_event_lg > e19753db-94e2-377e-b6e5-ad02f4f4f82f

Log events

You can use the filter bar below to search for and match terms, phrases, or values in your log events. [Learn more about filter patterns](#)

☐ View as text



Actions ▼

Create Metric Filter

Filter events

Clear

1m

30m

1h

12h

Custom



Timestamp Message

No older events at this moment. [Retry](#)

▼ 2021-11-23T13:43:11.000-05:00 {"version":"0","id":"5b18f90d-8547-7c50-1d82-8ad503155334","detail-type":"EC2 Instance State-change Notification","so...

```
{
  "version": "0",
  "id": "5b18f90d-8547-7c50-1d82-8ad503155334",
  "detail-type": "EC2 Instance State-change Notification",
  "source": "aws.ec2",
  "account": "123456789012",
  "time": "2021-11-23T18:43:11Z",
  "region": "us-east-1",
  "resources": [
    "arn:aws:ec2:us-east-1:123456789012:instance/i-0f8b96fdcd607b54f"
  ],
  "detail": {
    "instance-id": "i-0f8b96fdcd607b54f",
    "state": "terminated"
  }
}
```

Copy

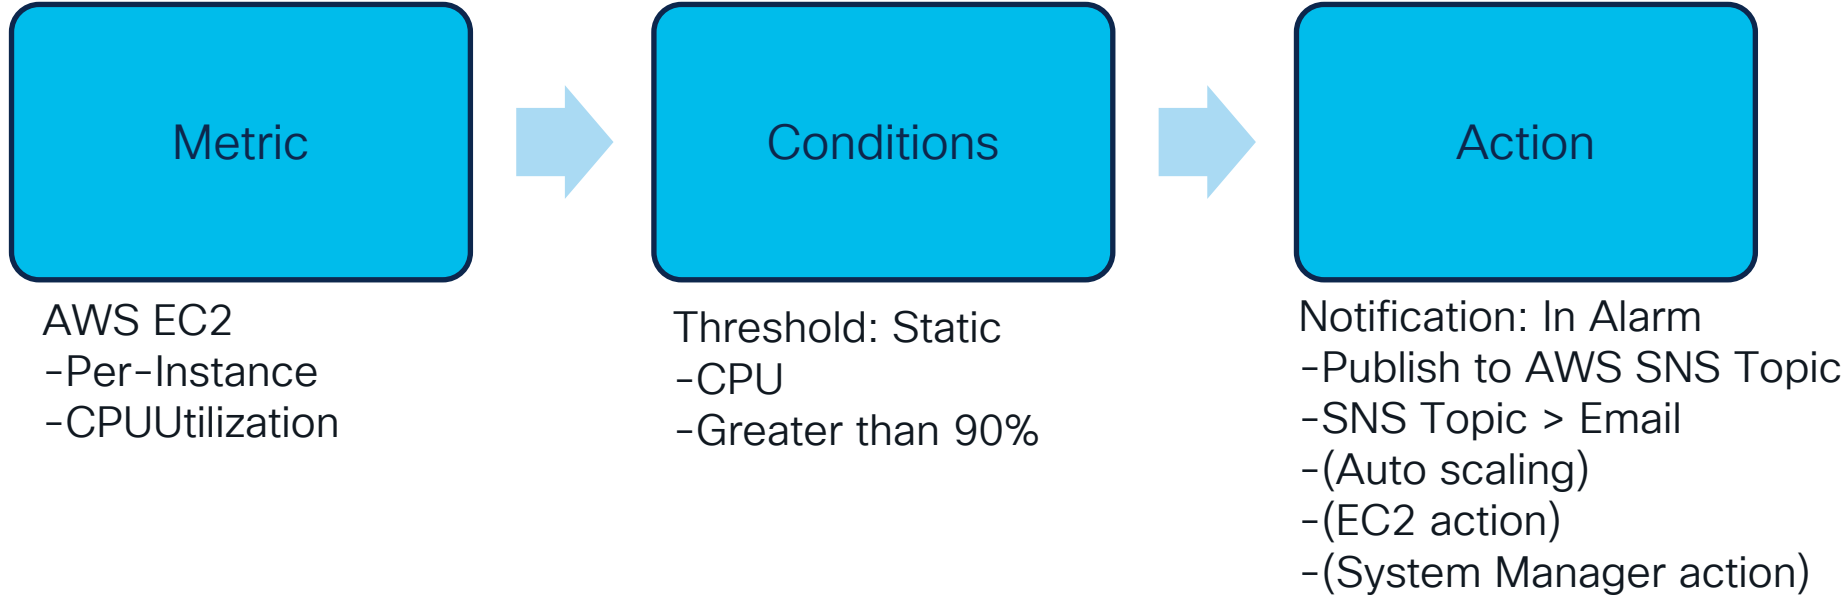
No newer events at this moment. Auto retry paused. [Resume](#)

Alerts

- Alerts – A predefined trigger based on a threshold or event
 - Static alert example:
 - HTTP Request Per Second (RPS) of 90% triggers alert to Slack
 - Kubernetes Node isn't ready for 1 minutes (Prometheus Example)

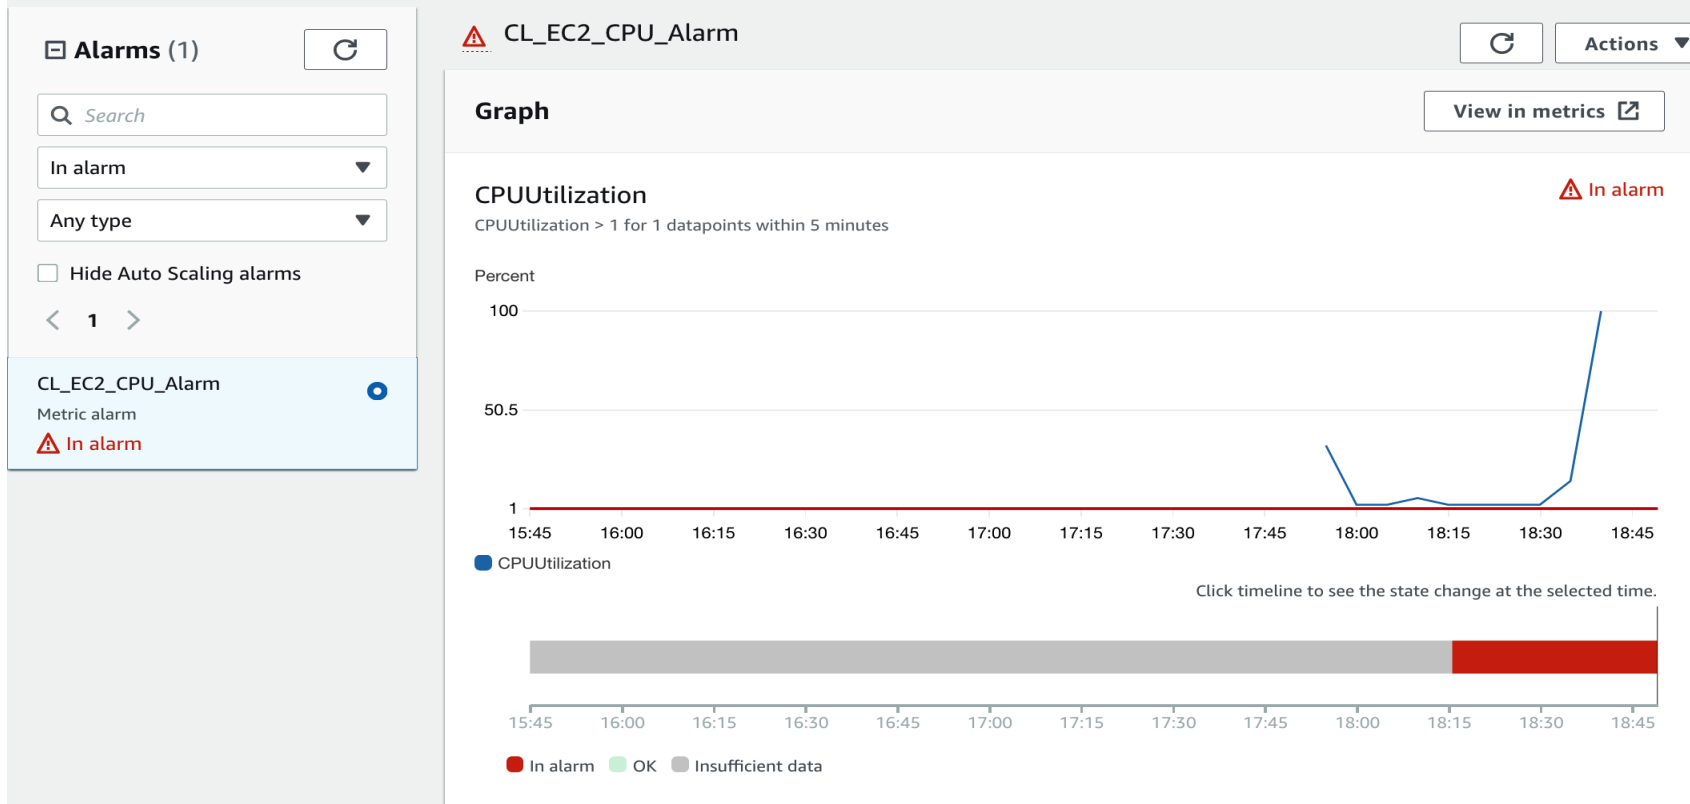
```
- alert: KubernetesNodeReady
  expr: kube_node_status_condition{condition="Ready",status="true"} == 0
  for: 1m
  labels:
    severity: critical
  annotations:
    summary: Kubernetes Node ready (instance {{ $labels.instance }})
    description: "Node {{ $labels.node }} unready \n  VALUE = {{ $value }}\n  LABELS = {{ $labels }}"
```

Alerts – AWS CloudWatch Alerts



Alerts – AWS CloudWatch Example

CloudWatch > Alarms > CL_EC2_CPU_Alarm



Alerts – AWS CloudWatch Example

Amazon SNS

Dashboard

Topics

Subscriptions

▼ Mobile

Push notifications

Text messaging (SMS)

Origination numbers

Amazon SNS > Topics > CL_Default_CloudWatch_Alarms_Topic

CL_Default_CloudWatch_Alarms_Topic

Edit

Delete

Publish message

Details

Name

CL_Default_CloudWatch_Alarms_Topic

ARN

arn:aws:sns:us-east-1:██████████:CL_Default_CloudWatch_Alarms_Topic

Type

Standard

Display name

-

Topic owner

██████████

Subscriptions

Access policy

Delivery retry policy (HTTP/S)

Delivery status logging

Encryption

Tags

Subscriptions (1)

Edit

Delete

Request confirmation

Confirm subscription

Create subscription

Search

<

1

>

⚙

ID

▼

Endpoint

▼

Status

▼

Protocol

▲



5294dbc9-0908-45ee-afad-526753dc0117



BRKCLD-2158



Confirmed

EMAIL

Alerts – AWS CloudWatch Example

You are receiving this email because your Amazon CloudWatch Alarm "CL_EC2_CPU_Alarm" in the US East (N. Virginia) region has entered the ALARM state, because "Threshold Crossed: 1 out of the last 1 datapoints [6.22950819672132 (23/11/21 18:10:00)] was greater than the threshold (1.0) (minimum 1 datapoint for OK -> ALARM transition)." at "Tuesday 23 November, 2021 18:15:36 UTC".

View this alarm in the AWS Management Console:

https://us-east-1.console.aws.amazon.com/cloudwatch/deeplink.js?region=us-east-1#alarmsV2:alarm/CL_EC2_CPU_Alarm

Alarm Details:

- Name: CL_EC2_CPU_Alarm

- Description:

- State Change: INSUFFICIENT_DATA -> ALARM

- Reason for State Change: Threshold Crossed: 1 out of the last 1 datapoints [6.22950819672132 (23/11/21 18:10:00)] was greater than the threshold (1.0) (minimum 1 datapoint for OK -> ALARM transition).

- Timestamp: Tuesday 23 November, 2021 18:15:36 UTC

- AWS Account: [REDACTED]

- Alarm Arn: arn:aws:cloudwatch:us-east-1:[REDACTED]:alarm:CL_EC2_CPU_Alarm

Threshold:

- The alarm is in the ALARM state when the metric is GreaterThanThreshold 1.0 for 300 seconds.

Monitored Metric:

- MetricNamespace: AWS/EC2

- MetricName: CPUUtilization

- Dimensions: [InstanceId = i-0f8b96fdcd607b54f]

- Period: 300 seconds

- Statistic: Maximum

- Unit: not specified

- TreatMissingData: missing

State Change Actions:

- OK:

- ALARM: [arn:aws:sns:us-east-1:[REDACTED]:CL_Default_CloudWatch_Alarms_Topic]

- INSUFFICIENT_DATA:

Logs

Logs

- Nearly everything in a Cloud Native (or other) environment produces logs in some form
- Logging has tremendous potential, but it is very complex to manage all the sources and then derive value out of what the logs say
- Collection and data formatting should be simple, but it isn't:
 - Currently, K8s doesn't enforce uniform structure for log messages*
 - You can't safely assume all log formats are in JSON
 - You may need to transform logs
- There are MANY gotchas on storage, forwarding, rotation – We don't have time for that today

[*https://github.com/kubernetes/enhancements/tree/master/keps/sig-instrumentation/1602-structured-logging](https://github.com/kubernetes/enhancements/tree/master/keps/sig-instrumentation/1602-structured-logging)

Logs

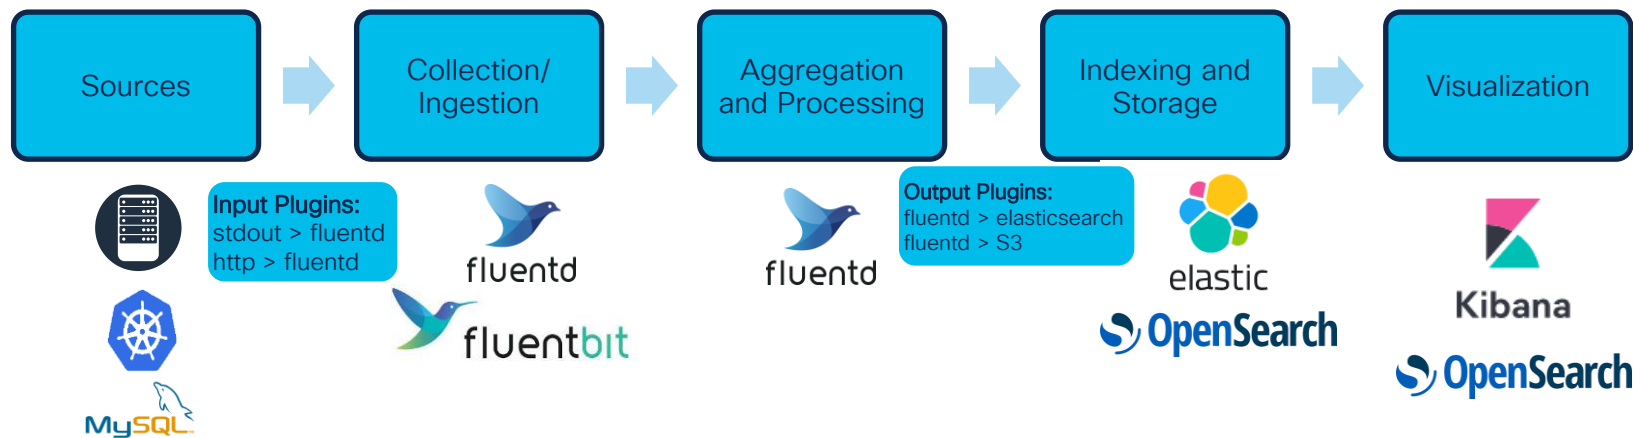
Common Components

- **Fluentd/Fluent Bit** – Open-Source log collection and processing (Bit is for highly resource-constrained environments)
- **Logstash** – Open-Source log collection and processing
- **Elasticsearch, OpenSearch, Grafana Loki**, etc – Aggregation, search and analytics
- **Kibana, OpenSearch Dashboard, Grafana Loki**, etc – Dashboard
- **AWS CloudWatch/CloudTrail, GCP Cloud Logging, Microsoft Azure Cloud Monitoring**
- Many more...

Common Stacks

- **EFK** – Elasticsearch, Fluentd/bit, Kibana
- **ELK** – Elasticsearch, Logstash, Kibana
- **OFO** – OpenSearch, Fluentd/bit, OpenSearch Dashboard
- **ENDLESS** combination of tools

Logs – Common Architecture Components

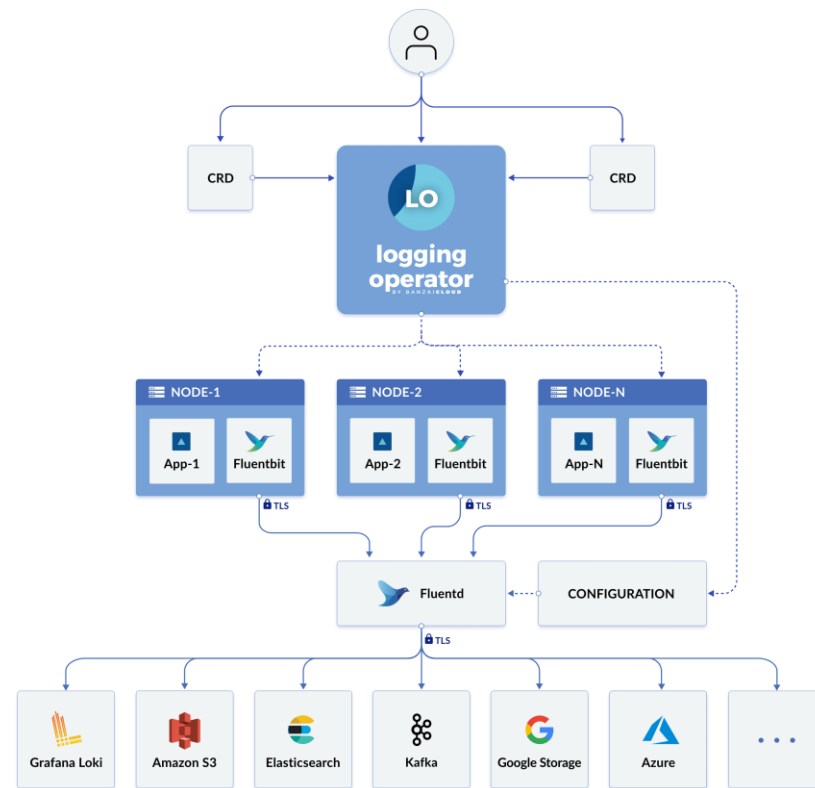


Building this by hand and as independent components is

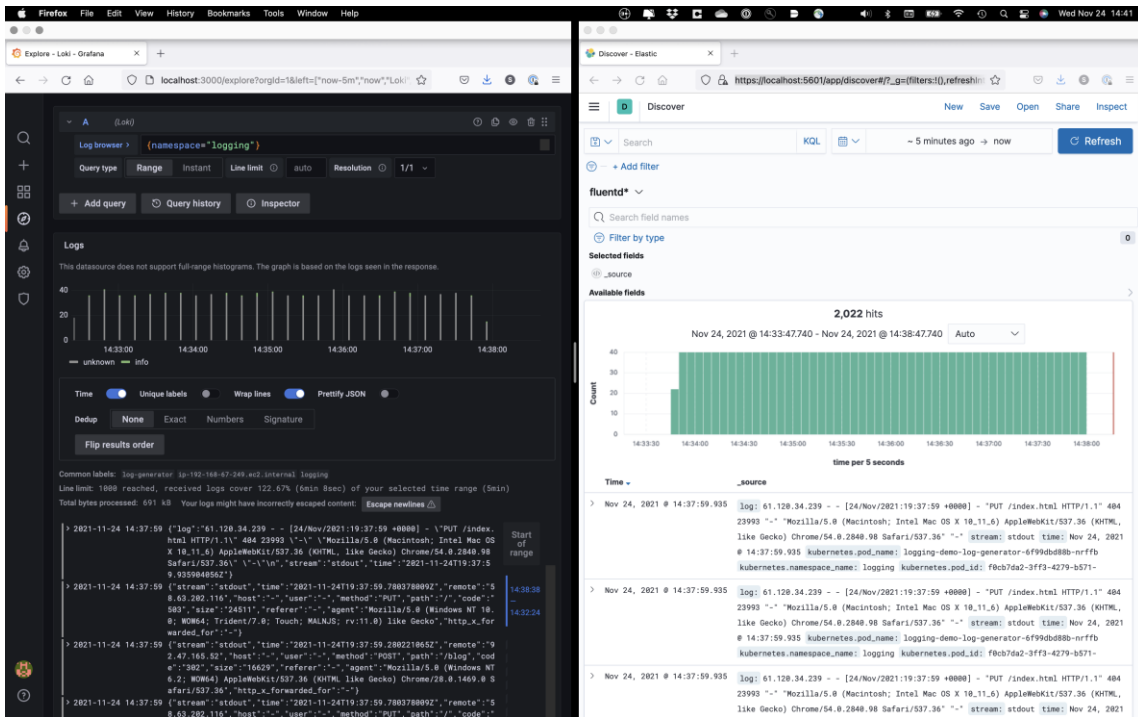


Logs

- Do things the easier way:
- Kubernetes:
 - Use Operators – Cisco Open-Source Logging Operator
 - <https://github.com/kube-logging/logging-operator>
- Fluentd/FluentBit and source configuration
- Security (TLS, RBAC, etc)
- Output configuration
 - AWS CloudWatch, S3, Azure Storage, GCP Storage, Elasticsearch, Grafana Loki, Kafka, etc.



Logging Operator – Example – One source, multi-outputs – NGINX to Elasticsearch/Kibana & Grafana Loki



Discover - Elastic

[←](#)
[→](#)
[↺](#)
[🏠](#)
[🔒](#)
[🌐](#)
[https://localhost:5601/app/discover/#/?_g=\(filters:!\(\),refreshInterval:5000,view:table,viewLayers:!\)](https://localhost:5601/app/discover/#/?_g=(filters:!(),refreshInterval:5000,view:table,viewLayers:!))
[🌟](#)
[🔔](#)
[⬇️](#)
[⚙️](#)
[👤](#)
[☰](#)

Elastic

Search Elastic

☰

D Discover

New Save Open Share Inspect

Table

JSON

_id	Cbh0U30BEpNoFPjK7bAq
_index	fluentd
_score	-
_type	_doc
kubernetes.annotations.kubernetes.io/psp	eks.privileged
kubernetes.container_hash	833498657557.dkr.ecr.us-east-2.amazonaws.com/banzaicloud/log-generator@sha256:c9facbdc427e4cf04f62580b1544479648b7361d93ca8a39c3e6ba66a32ea83b
kubernetes.container_image	833498657557.dkr.ecr.us-east-2.amazonaws.com/banzaicloud/log-generator:0.3.8
kubernetes.container_name	log-generator
kubernetes.docker_id	58357d3bc32b506d08fac9852ba517f13958e431430c2fbc458672ab639f3d7
kubernetes.host	ip-192-168-67-249.ec2.internal
kubernetes.labels.app.kubernetes.io/instance	logging-demo
kubernetes.labels.app.kubernetes.io/name	log-generator
kubernetes.labels.pod-template-hash	6f99dbd88b
kubernetes.namespace_name	logging
kubernetes.pod_id	f8cb7da2-3ff3-4279-b571-f39c85148447
kubernetes.pod_name	logging-demo-log-generator-6f99dbd88b-nrffb
log	61.120.34.239 - - [24/Nov/2021:19:37:59 +0000] - "PUT /index.html HTTP/1.1" 404 23993 "-" Mozilla/5.0 (Macintosh; Intel Mac OS X 10_11_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/54.0.2840.98 Safari/537.36 "-"
stream	stdout
time	Nov 24, 2021 @ 14:37:59.935

> Nov 24, 2021 @ 14:37:59.935

log: 61.120.34.239 - - [24/Nov/2021:19:37:59 +0000] - "PUT /index.html HTTP/1.1" 404

X

Find in page

☐ Highlight All
 ☐ Match Case
 ☐ Match Diacritics
 ☐ Whole Words

Traces

Traces

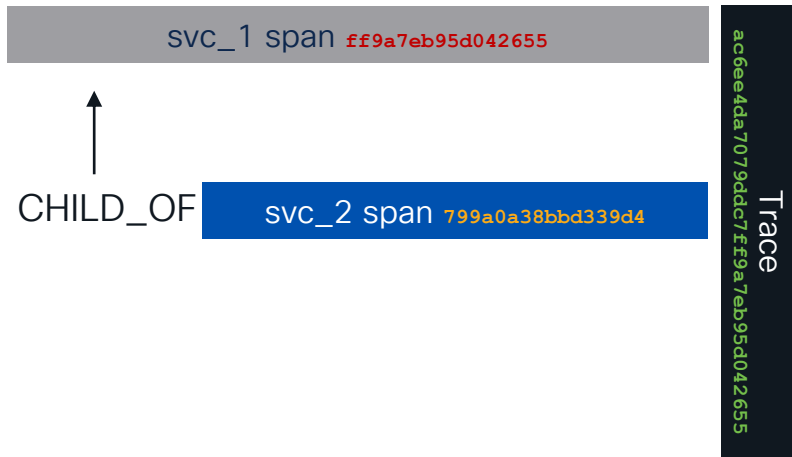
- Distributed tracing helps with:
 - Service mapping (topology)
 - Bottlenecks/Latency/Drops in a distributed architecture (network, microservice, etc.)
- Example projects/solutions:
 - OpenTelemetry - Combo of OpenCensus + OpenTracing - Library-based collection
 - Service Meshes - Istio, Linkerd, etc. - Sidecar-based collection
 - Jaeger - Visualize traces
 - W3C TraceContext/B3 TraceContext - Bringing some sanity to the format of a trace ID
 - AWS X-ray, GCP Cloud Trace, Azure Monitor (Application Insights) - Tracing libraries and visualization service

Traces

- Primer:

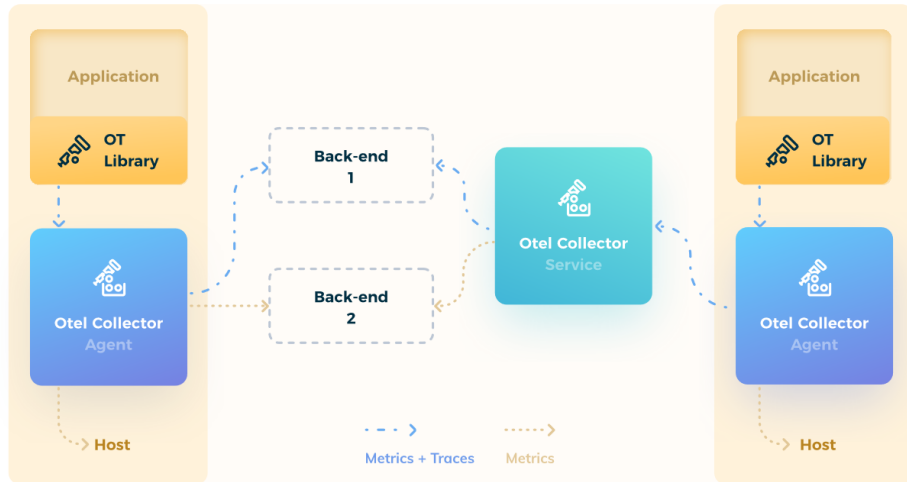
- A “span” is the foundational element of a distributed trace – it represents an individual unit of work
- A “span” can reference another span and when assembled, you have a “trace”
- “Context propagation” – Correlate trace metadata across service boundaries
- Not using a standard method for trace context propagation can lead to VERY painful deployments and VERY expensive workarounds

```
"traceID": "ac6ee4da7079ddc7ff9a7eb95d042655",
"spans": [
  {
    "traceID": "ac6ee4da7079ddc7ff9a7eb95d042655",
    "spanID": "ff9a7eb95d042655",
    "operationName": "frontpage.smm-demo.svc.cluster.local:8080/*",
    "references": [],
    "startTime": 1638308084933526,
    "duration": 803461,
  }
  {
    "traceID": "ac6ee4da7079ddc7ff9a7eb95d042655",
    "spanID": "799a0a38bbd339d4",
    "operationName": "frontpage.smm-demo.svc.cluster.local:8080/*",
    "references": [
      {
        "refType": "CHILD_OF",
        "traceID": "ac6ee4da7079ddc7ff9a7eb95d042655",
        "spanID": "ff9a7eb95d042655"
      }
    ]
  }
]
```



OpenTelemetry (OTel)

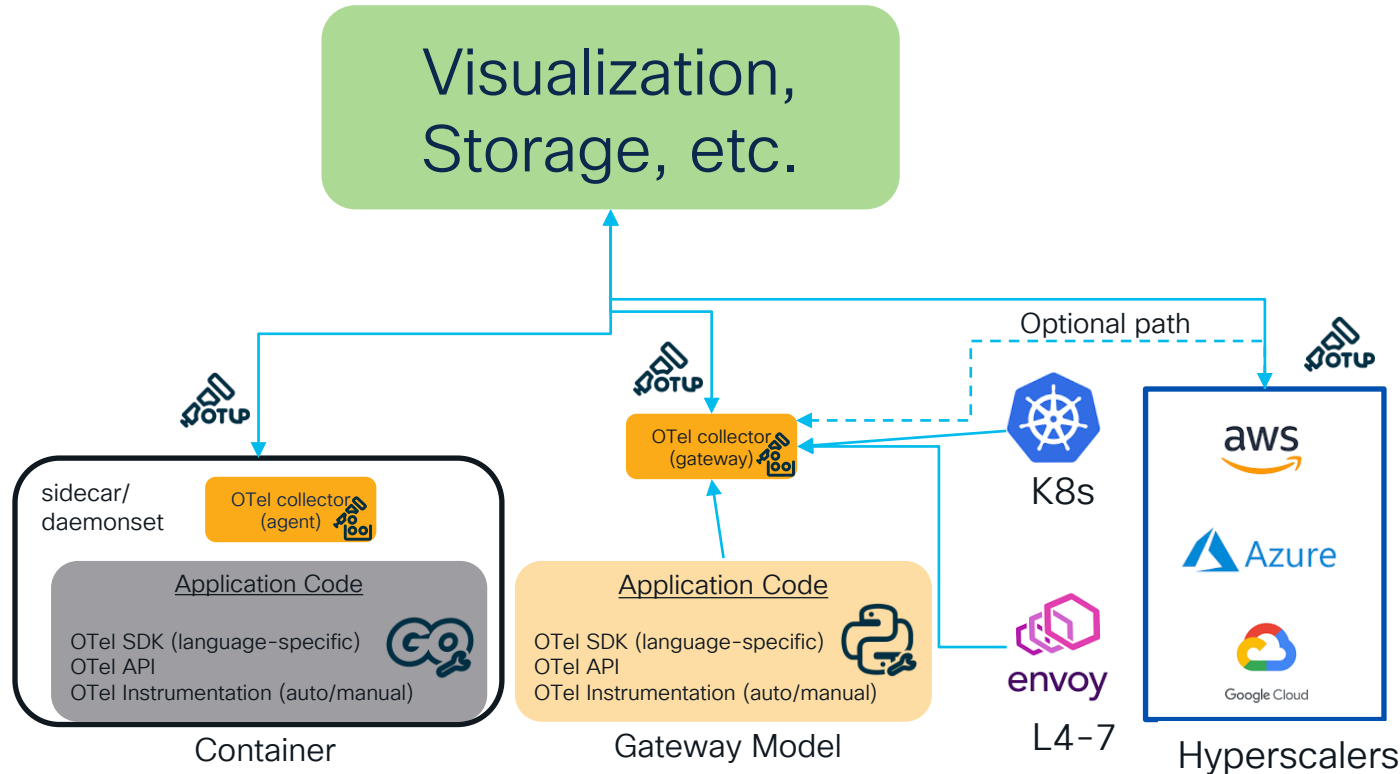
<https://opentelemetry.io/>



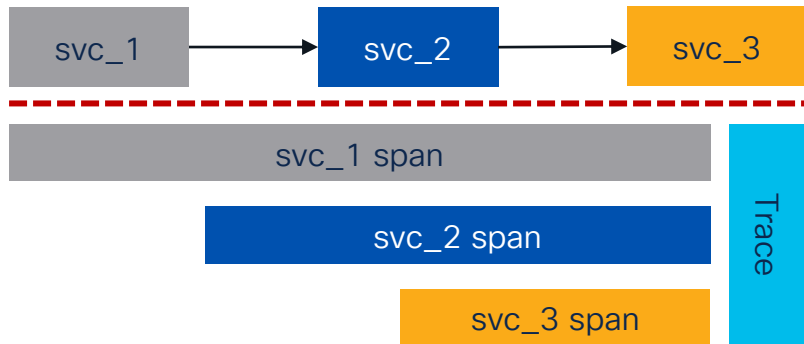
REFERENCE ARCHITECTURE

- Language-specific libraries
- Supports: Traces, Metrics, Logs
- The Collector recognizes multiple Trace Context formats
- Different form factors for the Collector
- <https://techblog.cisco.com/blog/getting-started-with-opentelemetry>

Example: OpenTelemetry Components in Action



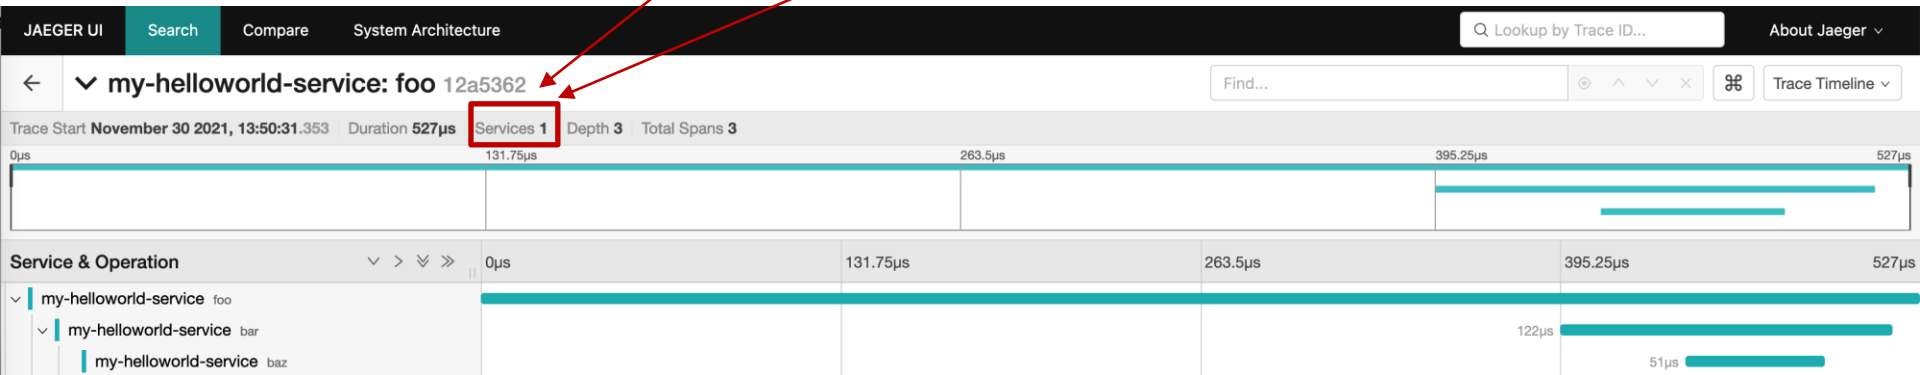
OpenTelemetry + Jaeger



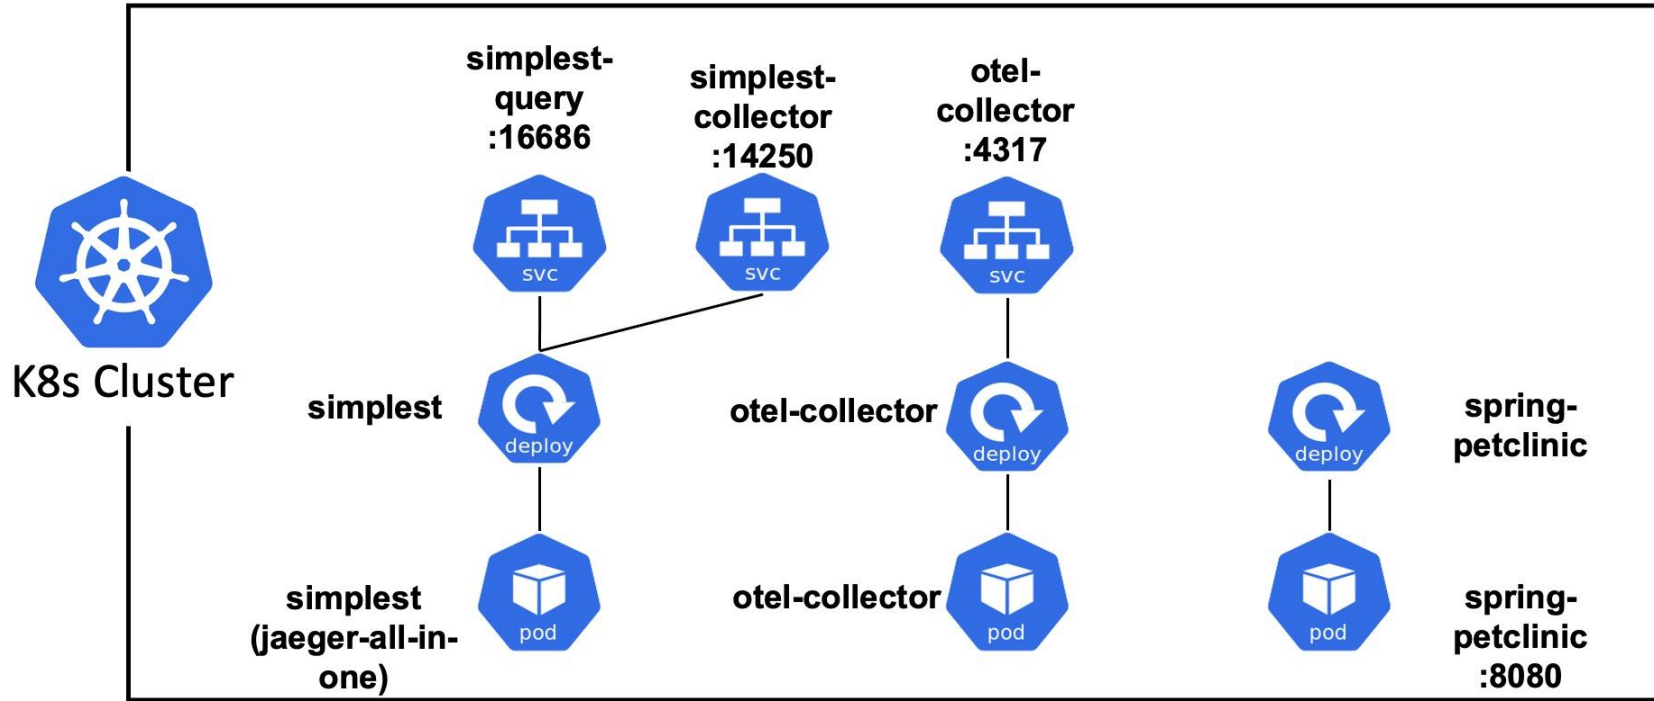
foo

Tags

internal.span.format	proto
otel.library.name	__main__
otel.library.version	
service.name	my-helloworld-service
span.kind	internal
telemetry.sdk.language	python
telemetry.sdk.name	opentelemetry
telemetry.sdk.version	1.7.1



Tracing Deployment Example



Example OpenTelemetry Deploy - 1

Deploy a test KinD cluster

```
# kind create cluster
```

Deploy Cert Manager

```
# kubectl apply -f https://github.com/cert-manager/cert-manager/releases/download/v1.11.0/cert-manager.yaml
```

Deploy the Jaeger Operator

```
# kubectl create namespace observability
# kubectl create -f https://github.com/jaegertracing/jaeger-operator/releases/download/v1.44.0/jaeger-operator.yaml -n observability
```

Deploy the Jaeger All-in-One Strategy

```
# kubectl apply -f - <<EOF
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: simplest
EOF
```

Deploy the OpenTelemetry Operator

```
# kubectl apply -f https://github.com/open-telemetry/opentelemetry-operator/releases/latest/download/opentelemetry-operator.yaml
```

Example OpenTelemetry Deploy - 2

Deploy the OTEL Collector

```
# kubectl apply -f - <<EOF
apiVersion: opentelemetry.io/v1alpha1
kind: OpenTelemetryCollector
metadata:
  name: otel
spec:
  config: |
    receivers:
      otlp:
        protocols:
          grpc:
          http:
    processors:
    exporters:
      logging:
      jaeger:
        endpoint: "simplest-collector:14250"
        tls:
          insecure: true
  service:
    pipelines:
      traces:
        receivers: [otlp]
        processors: []
        exporters: [jaeger]
EOF
```

Deploy the OTEL Java Auto-instrumentation CRD

```
# kubectl apply -f - <<EOF
apiVersion: opentelemetry.io/v1alpha1
kind: Instrumentation
metadata:
  name: my-instrumentation
spec:
  exporter:
    endpoint: http://otel-collector:4317
  propagators:
    - tracecontext
    - baggage
    - b3
  sampler:
    type: parentbased_traceidratio
    argument: "0.25"
  java:
    image: ghcr.io/open-telemetry/opentelemetry-
operator/autoinstrumentation-java:latest
  nodejs:
    image: ghcr.io/open-telemetry/opentelemetry-
operator/autoinstrumentation-nodejs:latest
  python:
    image: ghcr.io/open-telemetry/opentelemetry-
operator/autoinstrumentation-python:latest
EOF
```

Example OpenTelemetry Deploy - 3

Deploy the Spring Pet Clinic service

```
# kubectl apply -f - <<EOF
apiVersion: apps/v1
kind: Deployment
metadata:
  name: spring-petclinic
spec:
  selector:
    matchLabels:
      app: spring-petclinic
  replicas: 1
  template:
    metadata:
      labels:
        app: spring-petclinic
      annotations:
        sidecar.opentelemetry.io/inject: "true"
        instrumentation.opentelemetry.io/inject-java: "true"
    spec:
      containers:
        - name: app
          image: ghcr.io/pavolloffay/spring-petclinic:latest
EOF
```


```
# kubectl port-forward deployment.apps/spring-petclinic 8080:8080
```

```
# kubectl port-forward svc/simplest-query 16686:16686
```

Browser:

<http://localhost:8080>

<http://localhost:16686>



Example OpenTelemetry - Validation

```
# kubectl logs deployment.apps/otel-collector
. . . <output summarized>
builder/receivers_builder.go:73 Receiver started. {"kind": "receiver", "name": "otlp"}
. . .
jaegerexporter@v0.41.0/exporter.go:186 State of the connection with the Jaeger Collector backend {"kind": "exporter",
"name": "jaeger", "state": "READY"}
```

Browser-to-Jaeger Query Service (16686) and Otel Collector-to-Jaeger (14250)

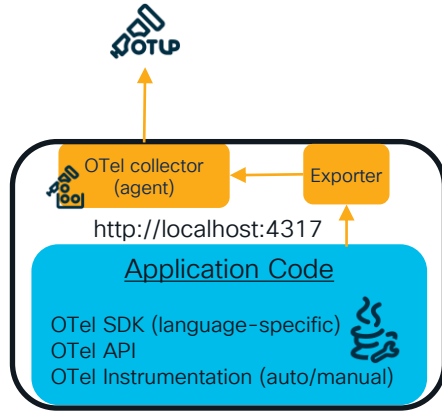
```
# kubectl exec -it simplest-797dd8fc67-619q4 -- netstat -at

Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
tcp        0      0 localhost:47280          localhost:16686         ESTABLISHED
tcp        0      0 simplest-797dd8fc67-619q4:14250 10-244-0-11.otel-collector.default.svc.cluster.local:49768 ESTABLISHED
```

Browser-to-Pet Clinic UI (8080) and Java Otel library-to-Otel Collector (4317)

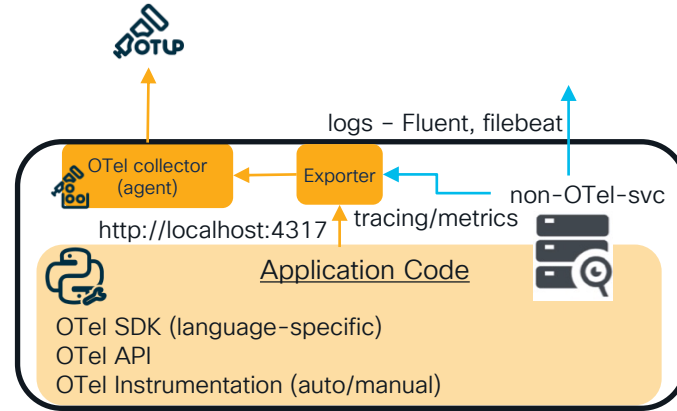
```
root@spring-petclinic-6d569df946-9f26m:/# netstat -at
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
tcp        0      0 localhost:55028          localhost:8080          ESTABLISHED
tcp        0      0 10.244.0.12:34208        otel-collector.def:4317 ESTABLISHED
```

Hybrid MELT Support



Java Agent - MLT over OTLP

```
java -javaagent:path/to/opentelemetry-  
javaagent-all.jar \ -jar myapp.jar
```



Python Agent - Tracing over OTLP

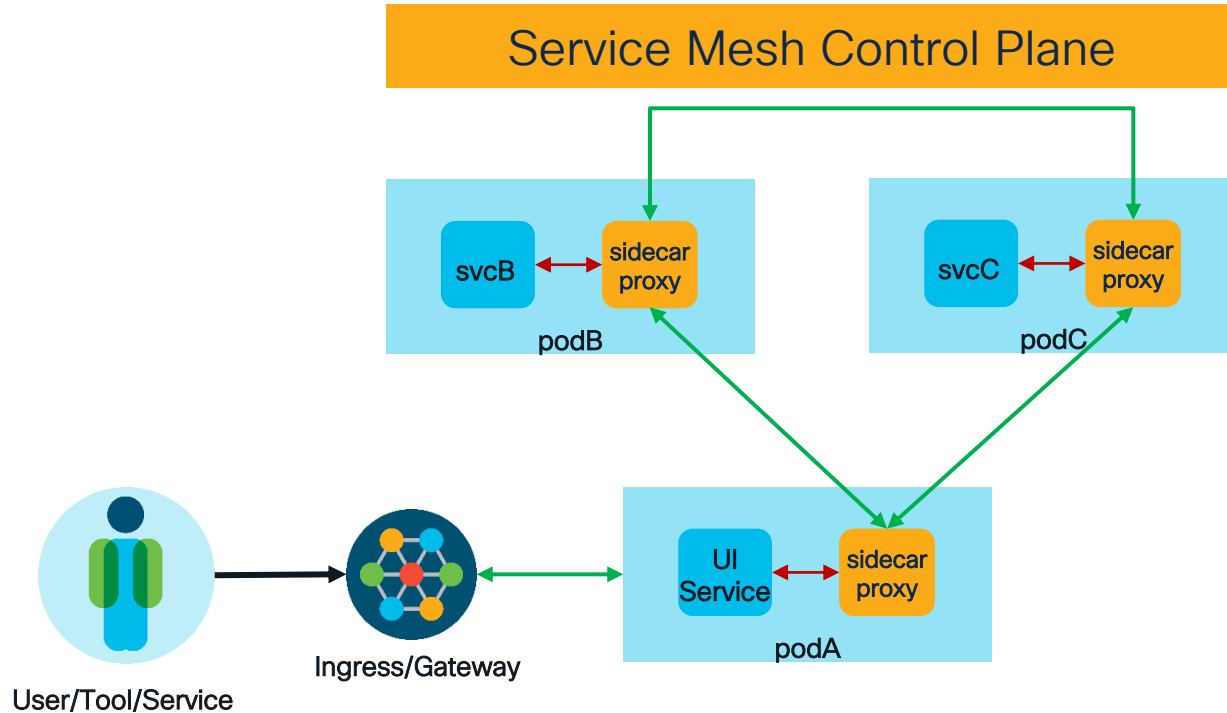
```
pip install opentelemetry-api  
pip install opentelemetry-sdk  
pip install opentelemetry-exporter-{exporter}  
pip install -U aws-xray-sdk
```

Application Instrumentation Options

- Library/SDK-based:
 - OpenTelemetry
 - Cisco/ApplicationDynamics
 - AWS X-ray, GCP Cloud Trace, Azure Monitor (Application Insights)
 - Many others
- Sidecar-based:
 - Service Meshes – Istio, Linkerd, Consul Connect, KongHQ, etc.

Service Mesh

What is a Service Mesh?



- Infrastructure layer for service-to-service communication
- Can use a mesh of sidecar proxies:
 - Can inspect API transactions at Layer 7 and 4 (TCP)
 - Intelligent routing rules can be applied between endpoints
- Allow for tracing and some application instrumentation without the need to add code/libraries/SDK to the application

Service Mesh Observability with Proxies

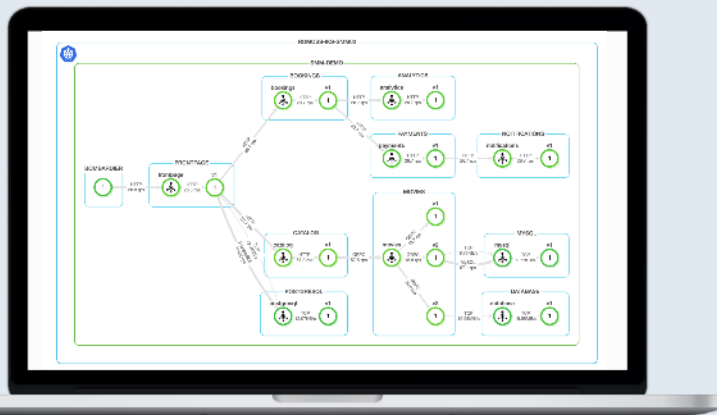
- Service Meshes provide observability via their sidecar proxies
- Observability info:
 - Mesh-specific metrics
 - Application-specific metrics
 - Distributed traces: Layer 4-7: TCP, HTTP, gRPC
 - Access logs (mesh and apps)
- In-mesh dashboards: Istio/Kiali, Linkerd/Viz, Cisco Calisti
- Out-of-mesh dashboards: Prometheus, Grafana, Jaeger, EFK, etc.

Cisco Solutions

<https://www.cisco.com/c/en/us/solutions/full-stack-observability.html>

Cisco Calisti

<https://calisti.app/>



Operationalize the Service Mesh

Multi-cloud, multi-cluster connectivity and observability
Connect any on-prem and public cloud together

Simplifies service mesh management
Single pane of glass, in depth metrics

Policy-based app networking & security
Policy management for DevOps teams

Apache Kafka on Kubernetes & Service Mesh
Lifecycle management of Apache Kafka and components

Traffic management ensures smooth app updates

Complete application and health observability

Security at all layers between clusters and clouds

Demo

Cisco Full-Stack Observability

Full Stack Observability (FSO)

is a requirement for business to deliver the most **optimal and secure experience** to users and applications.

Cisco Full-Stack Observability brings together data from multiple operations domains to provide **unified visibility**, derive **real-time insights** and recommend **actions** helping to:



Focus on what matters most: revenue, user experience, risk, costs



Reduce time to resolution of incidents and performance issues;



Minimize tool sprawl



Break down silos by reducing friction among teams, typically infrastructure, security, applications, networking and cloud

Cisco FSO Platform

Empowers an extensible observability ecosystem



Fully extensible entity-based data modeling



Natively supports OpenTelemetry



Anchored on MELT



Provides Unified Query Language (UQL)

FSO Exchange & Customers

Developer Ecosystem

OTEL Compliant

Developer Ready

Unified Query Language

MELT Data Foundation

Experience Foundation

Tenancy and Access Foundation

Summary

- There are a lot of things to keep track of and a lot of tools to help you do so:
 - Metrics, Events, Logs, Traces
 - Proprietary solutions, open-source solutions
 - Most solutions (vendor and OSS) do a handful of things well – most of the time up to you to ‘integrate’ them
- Next-gen solutions such as Cisco Full Stack Observability will reduce/remove the burden of you having to stitch together various tools to gain visibility to – derive value from and take action on your data
- Check out Cisco FSO: <https://www.cisco.com/c/en/us/solutions/full-stack-observability.html>
- Start working with Cisco Calisti!
 - <https://calisti.app/>

Fill out your session surveys!



Attendees who fill out a minimum of four session surveys and the overall event survey will get **Cisco Live-branded socks** (while supplies last)!



Attendees will also earn 100 points in the **Cisco Live Challenge** for every survey completed.



These points help you get on the leaderboard and increase your chances of winning daily and grand prizes

Continue your education



- Visit the Cisco Showcase for related demos
- Book your one-on-one Meet the Engineer meeting
- Attend the interactive education with DevNet, Capture the Flag, and Walk-in Labs
- Visit the On-Demand Library for more sessions at www.CiscoLive.com/on-demand

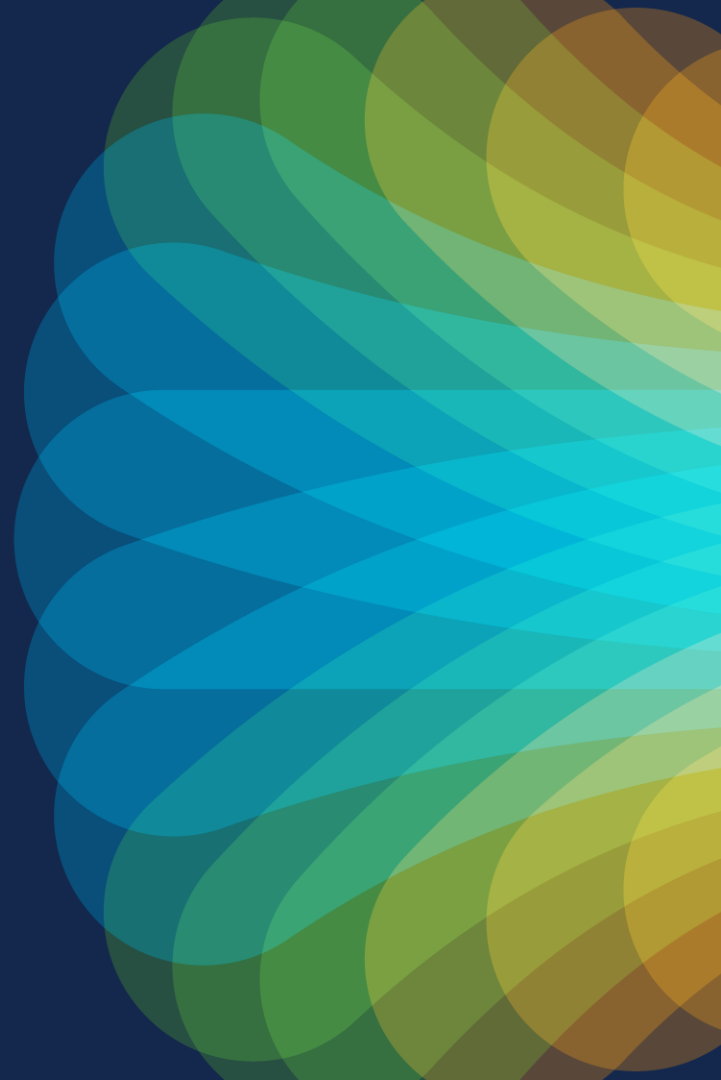


The bridge to possible

Thank you

CISCO *Live!*

#CiscoLive

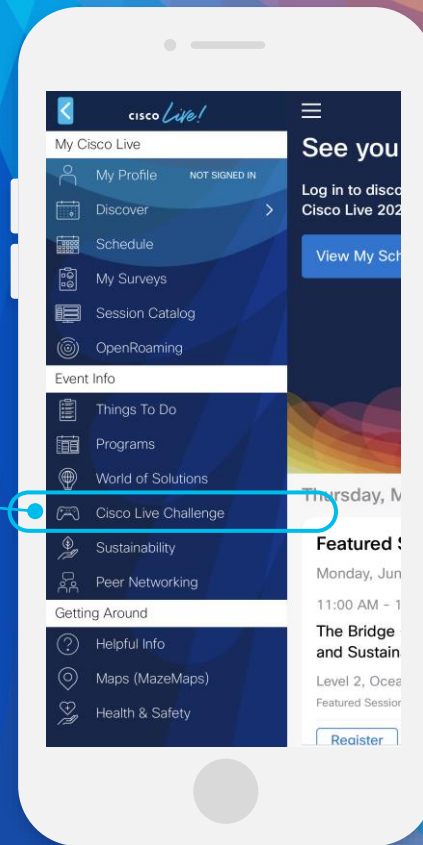
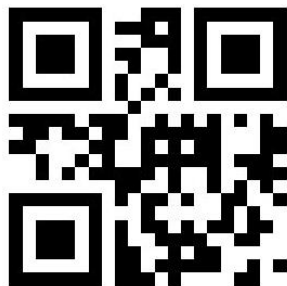


Cisco Live Challenge

Gamify your Cisco Live experience!
Get points for attending this session!

How:

- 1 Open the Cisco Events App.
- 2 Click on 'Cisco Live Challenge' in the side menu.
- 3 Click on View Your Badges at the top.
- 4 Click the + at the bottom of the screen and scan the QR code:



The background is a vibrant, abstract graphic. It features a central bright white light source from which numerous colorful rays emanate, creating a sunburst or starburst effect. The rays transition through a spectrum of colors including yellow, orange, red, and various shades of blue and green. Overlaid on this are large, flowing, wavy shapes in similar colors, giving the overall composition a sense of movement and energy.

cisco *Live!*

Let's go

#CiscoLive