



The bridge to possible

# But... why do I need a Service Mesh?

Traffic management with Istio

Julio Gómez, Principal Architect  
CCIE 9302, @juliodevops

# Cisco Webex App

## Questions?

Use Cisco Webex App to chat with the speaker after the session

## How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space

Webex spaces will be moderated until February 24, 2023.





# Agenda

- Introduction
- What is a Service Mesh?
- Istio Architecture
- How does it work?
- Capabilities
- Use cases demos
- Conclusion

DevOps

Microservices

Containers

Integration Server

Pipeline

Python

Kubernetes

Public Cloud

Helm

Docker

Git

Version Control

Repositories

Schedulers

Automation

Google Cloud

Grafana

AWS

Prometheus

API

YAML


Cloud Native Development

Drone

Istio

*“Service Meshes are  
the new **black**”*

Jessie Frazelle






**Dave Strebel @Hashiconf**  
@dave\_strebel



"Finally got Istio into production"



**"Information is not knowledge.  
The only source of knowledge  
is experience."**



# Introduction





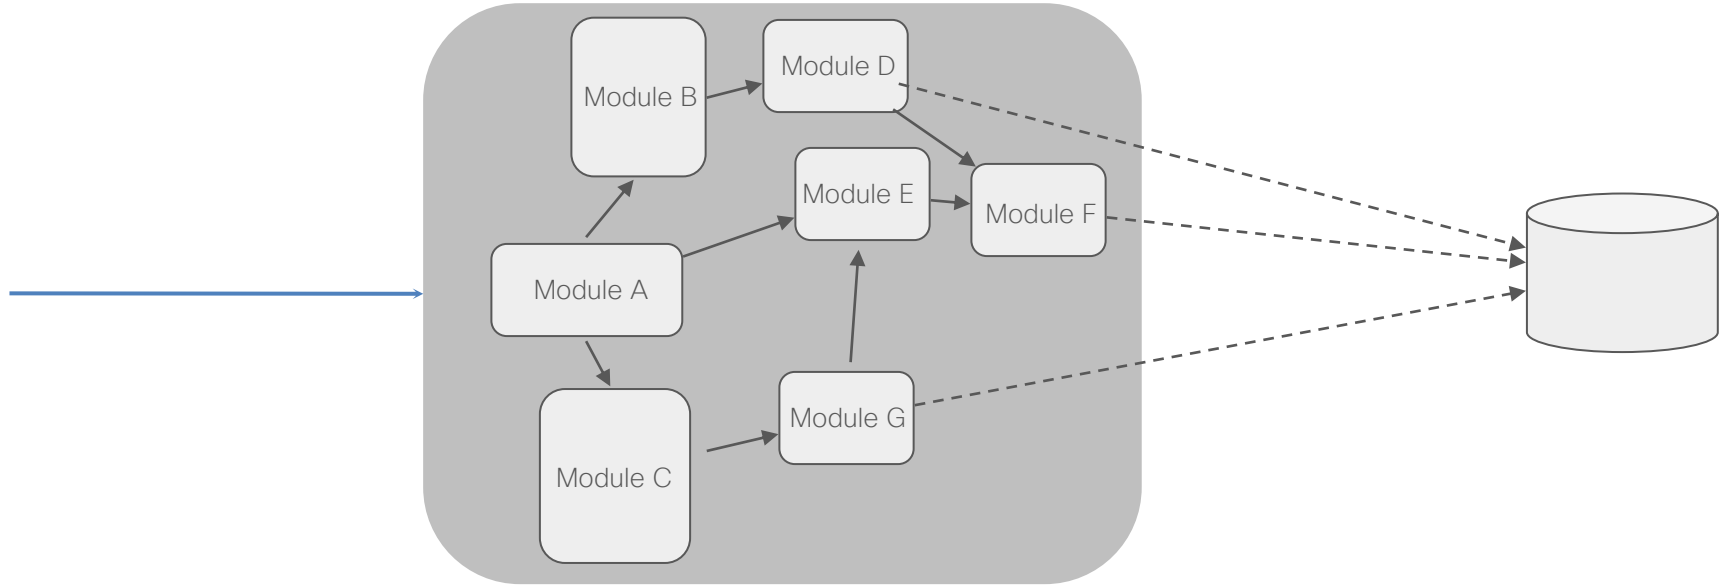
# Distributed world

The trends of containerization, microservices and hybrid/multi-cloud deployments have created more distributed applications than ever.


Developers, DevOps and SecOps personnel need modern tools to **secure, manage and monitor** distributed applications.



# Remember the Monolith?



# Microservices



# Microservices FTW

1. Gained development velocity
2. Easy testing because of abstractions
3. Scale services independently



**Problem**  
**SOLVED**

# But... what have we lost?

1. We replaced a reliable in-process call with an unreliable RPC
2. Trivial single stepping replaced by...?
3. Secure in-process communication is replaced by insecure network
4. Access control within process was a NOOP
5. Latency went up

wow, that abstraction was leaky...

# Can developers fix it?

1. Add retry logic to the application code
2. Add entry-exit traces
3. Secure inter-service connections with strong authentication

And while you are adding code... choose the RPC endpoint intelligently

- a) Endpoints with low latency
- b) Endpoints with warm caches

```
try {  
    HttpResponse response = httpClient.get(  
        "http://secretsauce.internal/recipe");  
    cook(response.body);  
} catch (NetworkError ne) {  
    fixmePleaseOMG(ne);  
}
```

```
try {  
    // Load balancing  
    IP ip =  
    DNS.lookupSRV("secretsauce.internal").pickOne();  
    HttpResponse response = httpClient.open(ip).get(  
        "http://secretsauce.internal/recipe");  
    cook(response.body);  
} catch (NetworkError ne) {  
    fixmePleaseOMG(ne);  
}
```

Load-balance  
Regionality



```
for (int i = 0; i < 3; i++) { // Retry
    try {
        IP ip =
        DNS.lookupSRV("secretsauce.internal").pickOne();
        HttpResponse response = httpClient.open(ip).get(
            "http://secretsauce.internal/recipe");
        cook(response.body);
    } catch (NetworkError ne) {
        if (i == 2) fixmePleaseOMG(ne);
        else Thread.sleep(random(5) * 1000);
    }
}
```

Retries  
Load-balance  
Regionality

```
Secret key = new Secret(new File("/somewhere/safe/key"));  
for (int i = 0; i < 3; i++) {  
    try {  
        IP ip =  
        DNS.lookupSRV("secretsauce.internal").pickOne();  
        HttpResponse response = httpClient.open(ip)  
            .setHeader("Authorization", key.toString())  
            .get("http://secretsauce.internal/recipe");  
        cook(response.body);  
    } catch (NetworkError ne) {  
        if (i == 2) fixmePleaseOMG(ne);  
        else Thread.sleep(random(5) * 1000);  
    }  
}
```

Authorization  
Retries  
Load-balance  
Regionality

```
Secret key = new Secret(new File("/somewhere/safe/key"));
for (int i = 0; i < 3; i++) {
    try {
        IP ip =
        DNS.lookupSRV("secretsauce.internal").pickOne();
        HttpResponse response = httpClient.open(ip)
            .setHeader("Authorization", key.toString())
            .get("http://secretsauce.internal/recipe");
        log("Success") ;
        cook(response.body) ;
    } catch (NetworkError ne) {
        log("Failed") ;
        if (i == 2) fixmePleaseOMG(ne);
        else Thread.sleep(random(5) * 1000);
    }
}
```

Monitoring?  
Health checks?  
Latency?  
Circuit breaking?  
Alerting?

Logging  
Authorization  
Retries  
Load-balance  
Regionality

# Problem: Too Much Infra Code in Services

- Too much work to make services production-ready
  - Load balancing, auto scaling, rate limiting, traffic routing...
- Done in different ways in every service
  - Retry, TLS, failover, deadlines, cancellation, etc, for each language, framework
- Service management across services is responsibility of each service
- Custom siloed implementations lead to fragmented implementation, no uniform policy application, difficult debugging
- Code bloat for each service

# Solution: Systemic Service Management

Core Idea: Deploy lightweight sidecar proxies to take over all ingress and egress traffic

- Transparently attach the proxies to service backends
- Mediate all inbound and outbound traffic between the application and the services it consumes
- Provides an array of infrastructure features

## Outbound features:

- ❖ Service authentication
- ❖ Load balancing
- ❖ Retry and failover
- ❖ Fine-grained routing
- ❖ Telemetry
- ❖ Tracing

## Inbound features:

- ❖ Service authentication
- ❖ Authorization
- ❖ Rate limits
- ❖ Load shedding
- ❖ Telemetry
- ❖ Tracing

# Service Mesh to democratize the solution

- Address service level concerns
- Unlock the full power of microservices
- Elevate the network to the needs of applications
- Uniform observability regardless of platform
- Security by default, everywhere



# What is a Service Mesh?

An open platform to  
manage service  
interactions among  
microservices in a  
consistent and  
secure way

Decouple operations from  
development with Istio anywhere


Enable customers to secure,  
monitor and manage services  
everywhere

Kubernetes first, but not  
Kubernetes only



# Why a service mesh?

A service mesh provides a transparent and language-independent way to flexibly and easily automate application network functions




# Value Proposition


(everybody needs something...  
but not everybody needs everything)

**À  
La Carte**

**CISCO** *Live!*



Observability



Control



Security

Why not Istio  
everything?





Why not Istio  
everything?



# Observability

Understanding services and their dependencies

Monitor uniform Service Level Indicators for every service

Collect telemetry, logs and traces

Improved understanding of applications at the *service* (not network) level

# Control

Scale by directing traffic to multiple versions

Roll out new versions without worrying about ops challenge

Apply access control, rate limiting policies to protect services from bad behavior

**TODAY**

# Security

**Secure by default** – new and existing applications

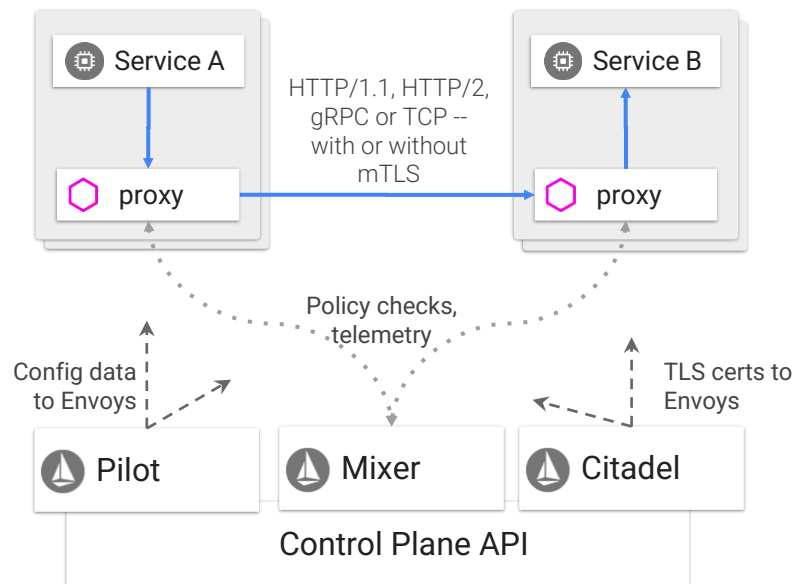
Meet compliance obligations by **encrypting data in transit**

mTLS assures a secure, proven **service-based identity** for every call

With strong identity, **authorization** can be explicitly required

# Istio architecture





**Pilot:** Control plane to configure and push service communication policies

**Envoy:** Network proxy to intercept communication and apply policies

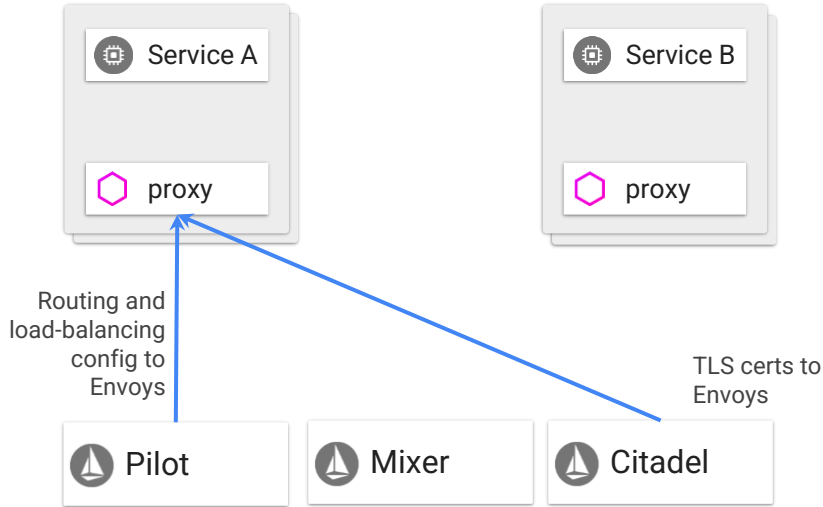
**Mixer:** Policy enforcement with a flexible plugin model for providers for a policy

**Citadel:** Service-to-service auth[n,z] using mutual TLS, with built-in identity and credential management

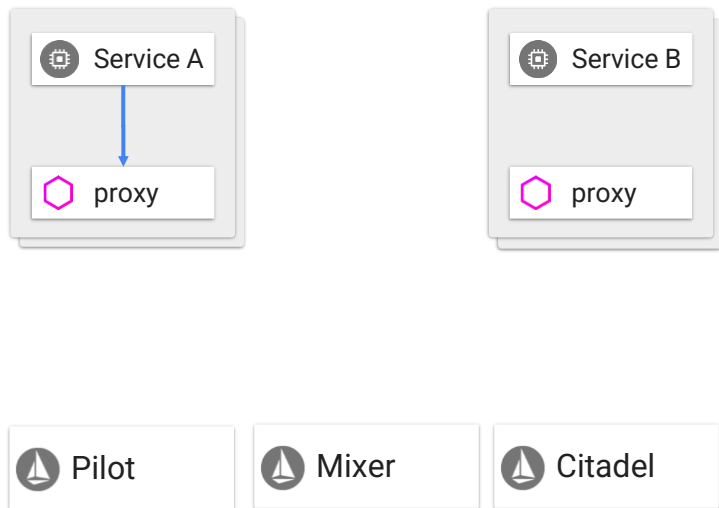
# How does it work?

Life of a request in the service mesh






Service A comes up. Envoy is deployed with it and fetches service information, routing and configuration policy from Pilot. If Citadel is being used, TLS certs are securely distributed as well.




Service A places a call to service B

Client-side Envoy intercepts the call

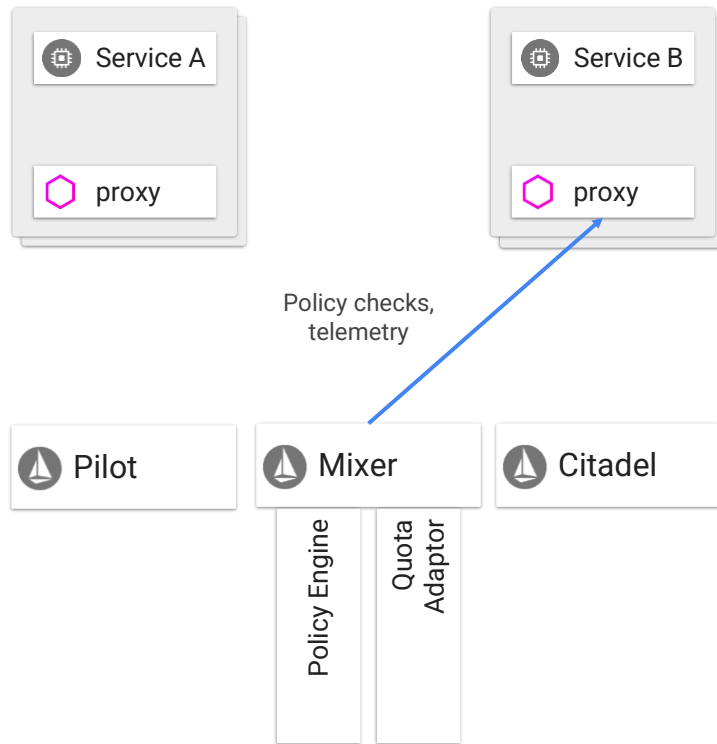
Envoy consults config to know  
how/where to route call to service B



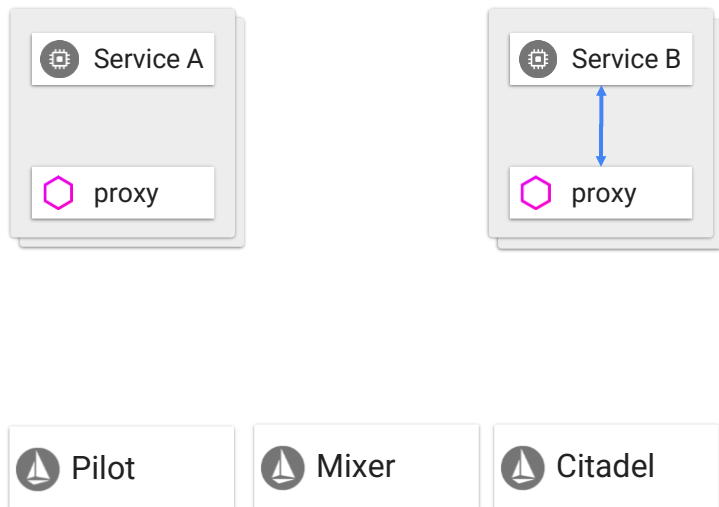
Envoy forwards request to appropriate instance of service B



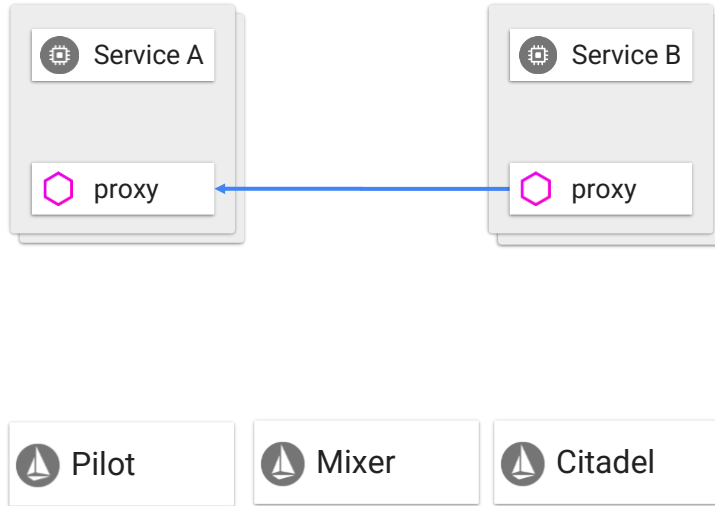
There, the Envoy proxy deployed with the service intercepts the call



Mixer checks with appropriate adaptors (policy engine, quota adaptor) to verify that the call can proceed and returns true/false to Envoy

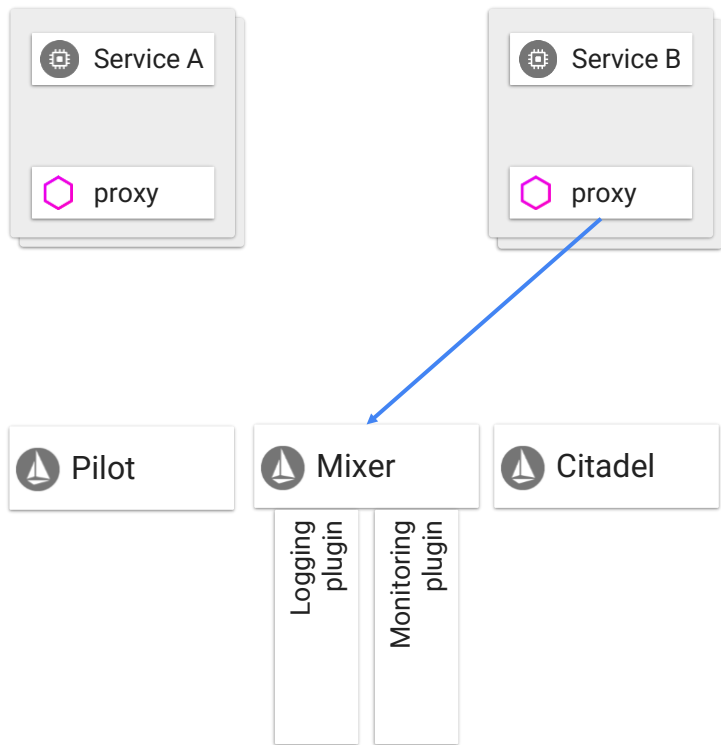


Server-side Envoy forwards request to service B, which processes request and returns response




Envoy forwards response to the original caller, where response is intercepted by Envoy on the caller side




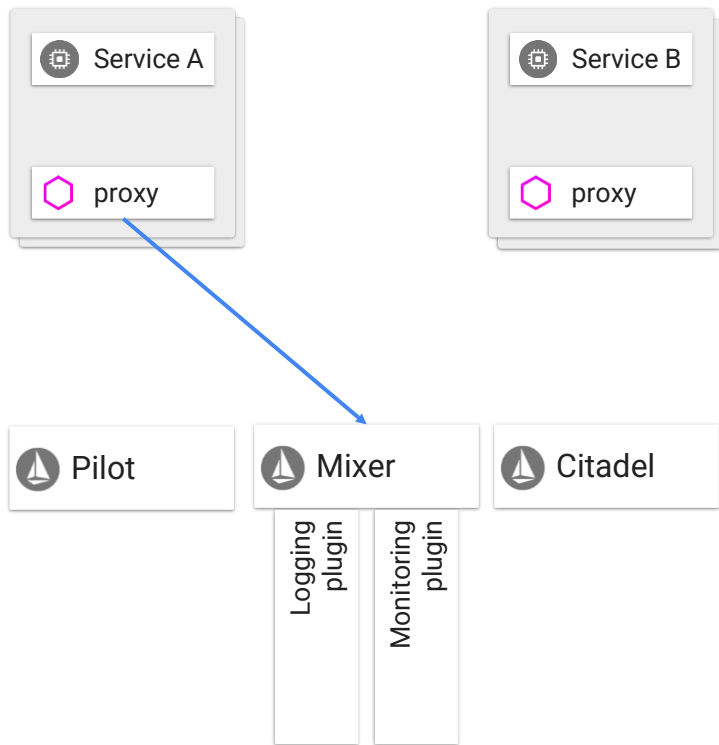


Envoy reports telemetry to Mixer,  
which in turn notifies appropriate  
plugins



Client-side Envoy forwards  
response to original caller






Client-side Envoy reports telemetry to Mixer (including client-perceived latency), which in turn notifies appropriate plugins

# Capabilities



# A network for services

- **Traffic Control**
- Visibility
- Resiliency & Efficiency
- Security



# Application Rollout

```
// A simple traffic splitting rule
```

```
destination: serviceB.example.cluster.local
```

```
match:
```

```
  source: serviceA.example.cluster.local
```

```
route:
```

```
- tags:
```

```
  version: v1.5
```

```
  env: us-prod
```

```
weight: 99
```

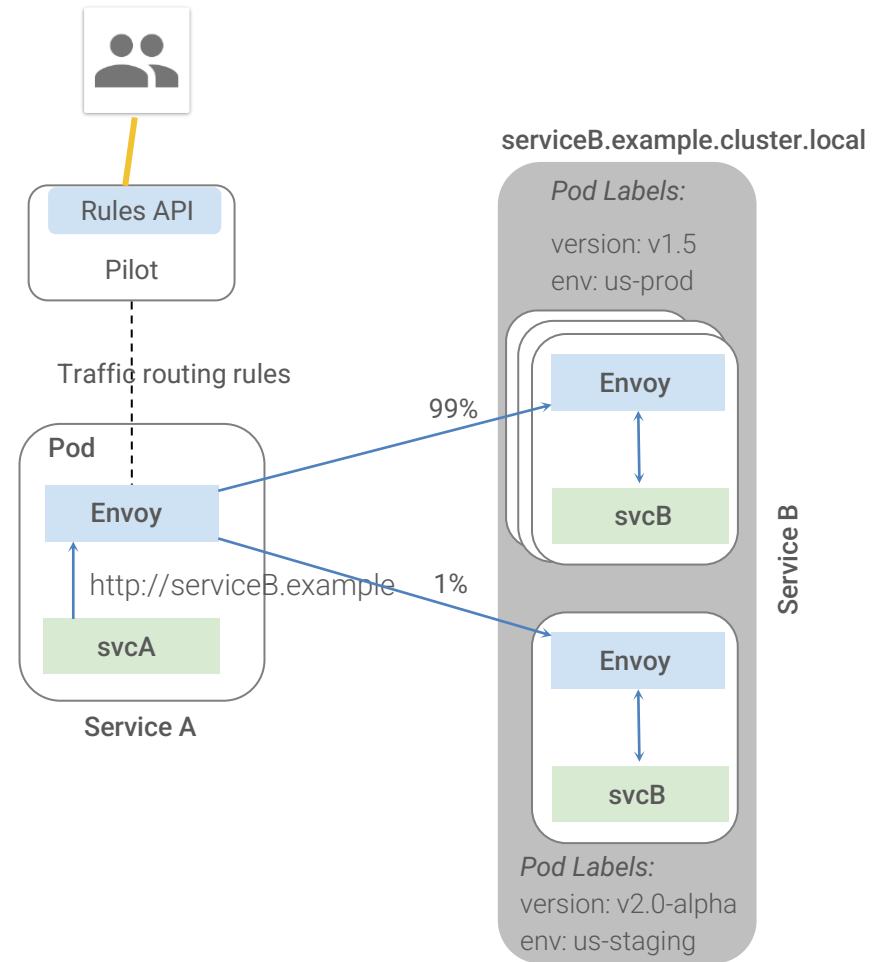
```
- tags:
```

```
  version: v2.0-alpha
```

```
  env: us-staging
```

```
weight: 1
```

Traffic control is decoupled  
from infrastructure scaling



# Traffic Steering

```
// Content-based traffic steering rule
```

```
destination: serviceB.example.cluster.local
```

```
match:
```

```
  httpHeaders:
```

```
    user-agent:
```

```
      regex: ^(.?*;)?(iPhone)(;.*)?$
```

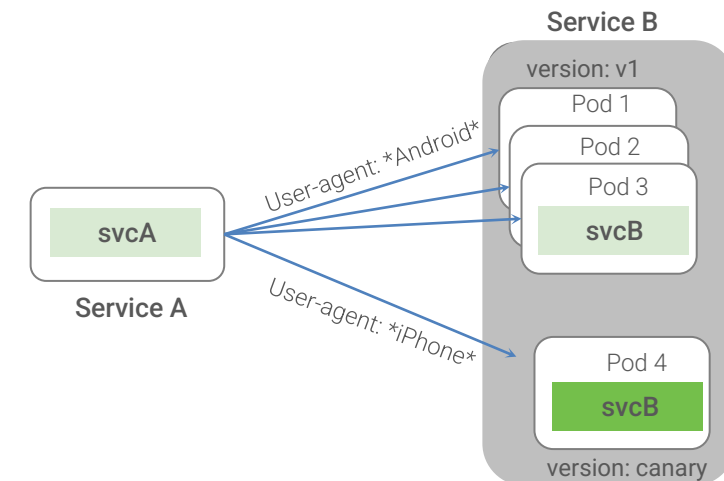

```
precedence: 2
```

```
route:
```

```
- tags:
```

```
  version: canary
```


## Content-based traffic steering





# A network for services

- Traffic Control
- **Visibility**
- Resiliency & Efficiency
- Security





- Admin
- Namespaces

Nodes

Persistent Volumes

Namespace

kube-system

- Workloads
- Deployments

Replica Sets



Replication Controllers

Daemon Sets

Stateful Sets

Jobs

Pods
- Services and discovery
- Services



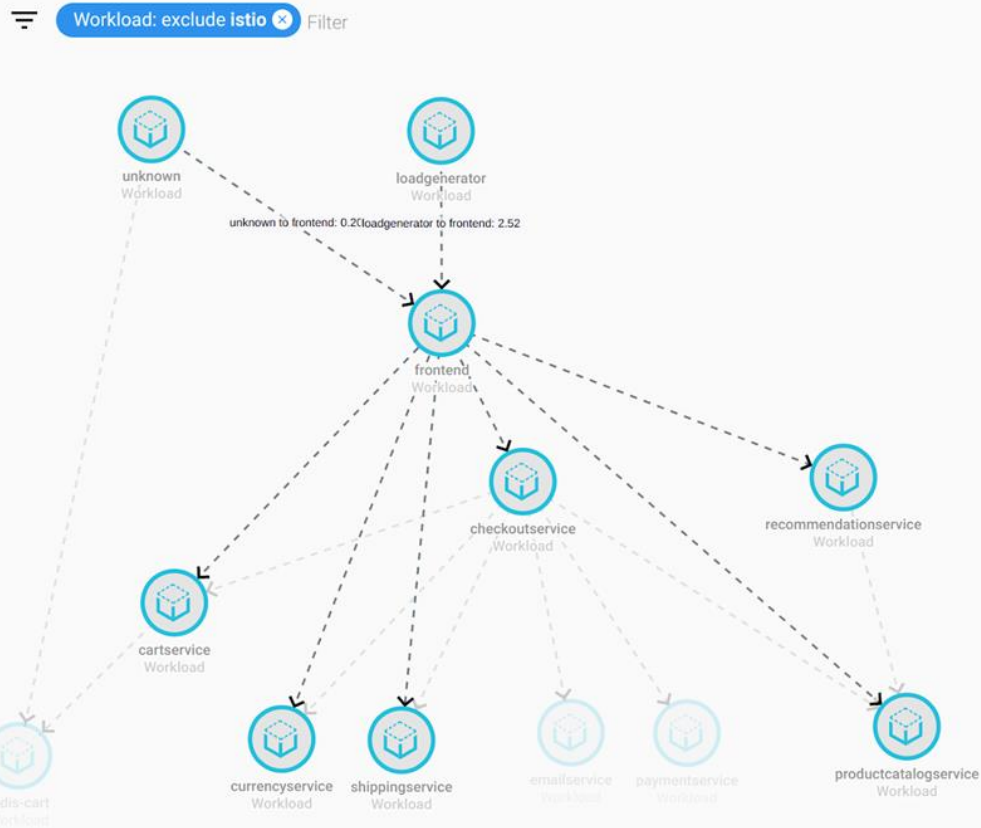
Pods1 - 10 of 18<<<>>>

Name	Status	Restarts	Age	CPU (cores)	Memory (bytes)
dashboard-same-i...	Running	0	14 minutes	0	5.4 Mi
fluentd-cloud-logg...	Running	0	2 months	0.01	120.8 Mi
fluentd-cloud-logg...	Running	0	2 months	0.009	94.3 Mi
fluentd-cloud-logg...	Running	0	2 months	0.014	174.3 Mi
heapster-v1.2.0-4...	Running	0	10 days	0.002	44.8 Mi

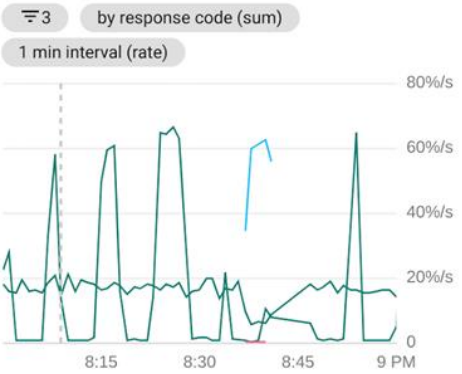


Filters

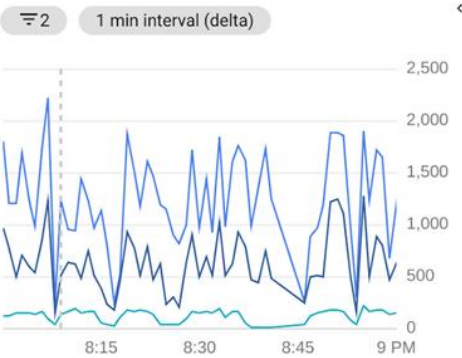
- Daemon Set
- Workload
- Pod
- Container
- Replica Set
- Deployment



Percent non 200 response codes



Response Latencies

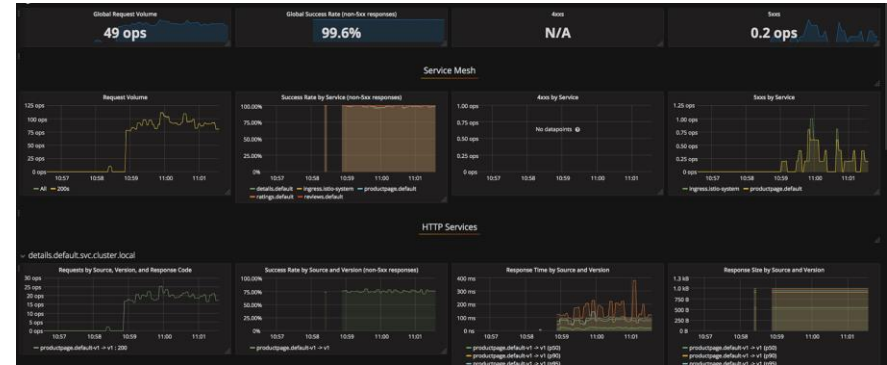


# Visibility


Monitoring & tracing should not be an afterthought in the infrastructure

## Goals

- Metrics without instrumenting apps
- Consistent metrics across fleet
- Trace flow of requests across services
- Portable across metric backend providers



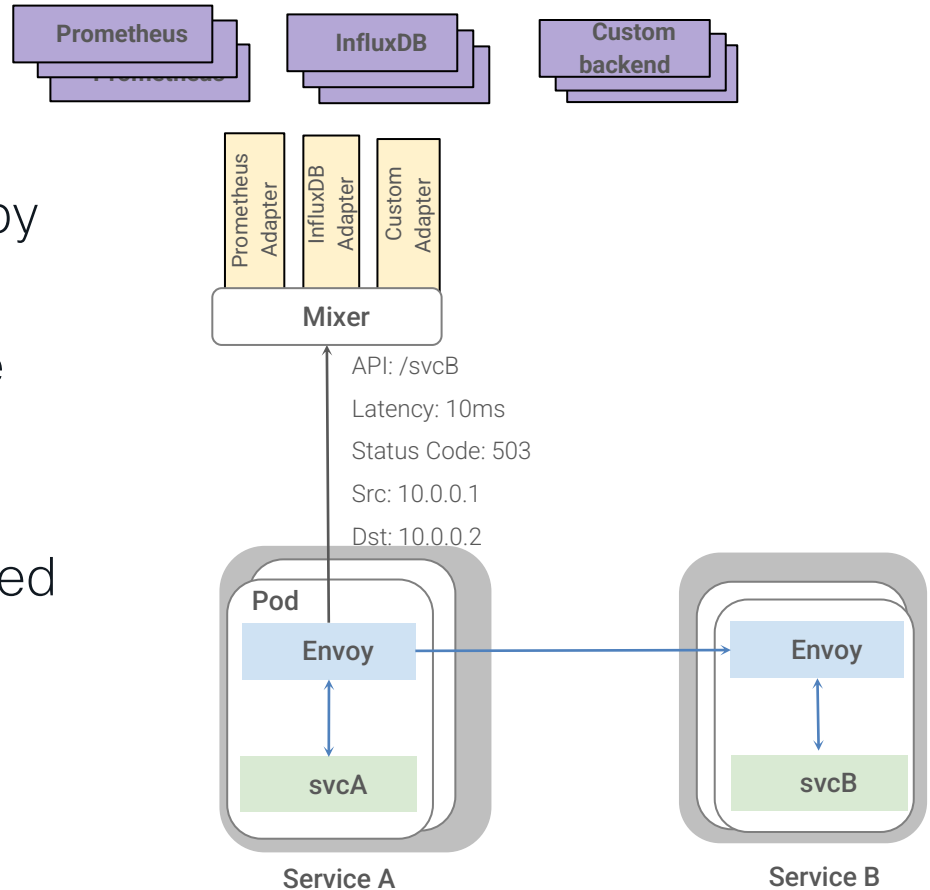
Istio - Grafana dashboard w/ Prometheus backend



Istio Zipkin tracing dashboard

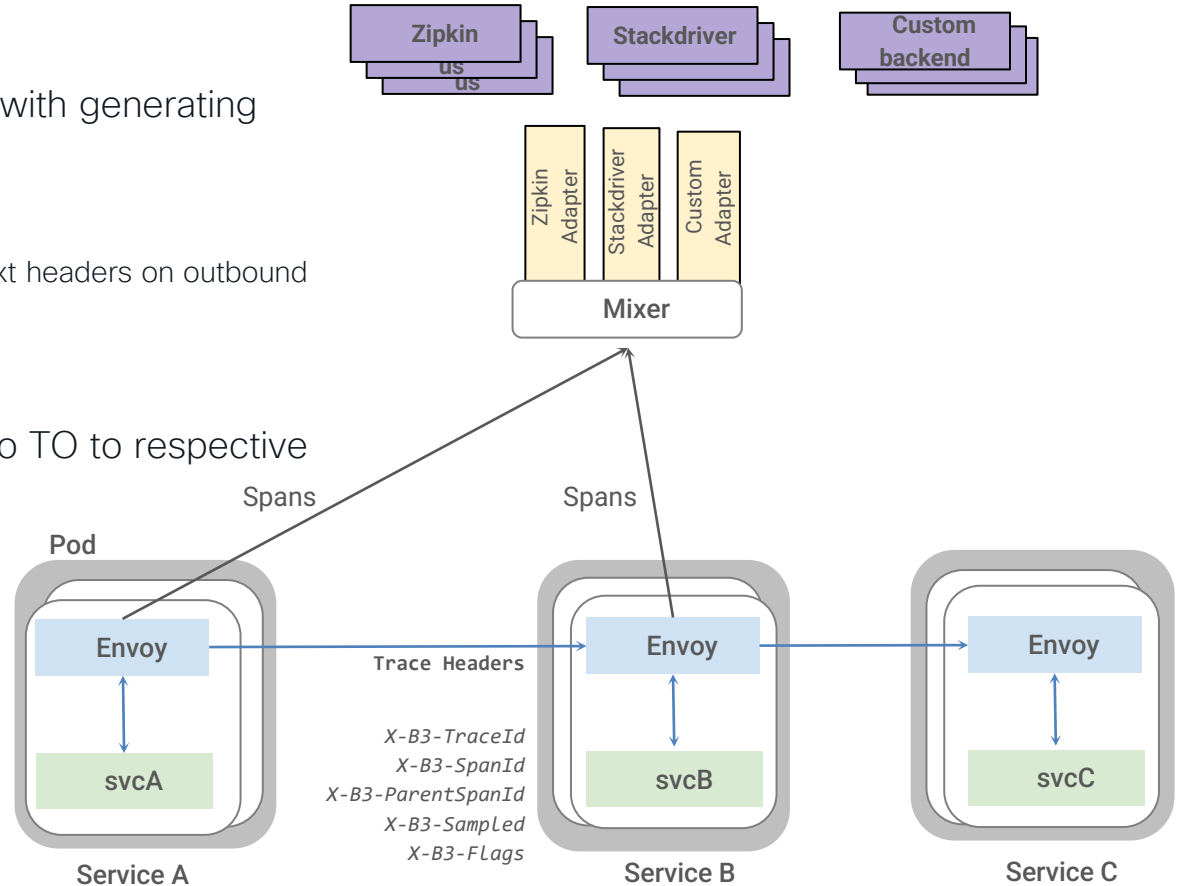
# Metrics flow

- Mixer collects metrics emitted by Envoys
- Adapters in the Mixer normalize and forward to monitoring backends
- Metrics backend can be swapped at runtime




# Tracing flow

- Application do not have to deal with generating spans or correlating causality
- Envoys generate spans
  - Applications need to *\*forward\** context headers on outbound calls
- Envoys send traces to Mixer
- Adapters at Mixer send traces to TO to respective backends



# A network for services

- Traffic Control
- Visibility
- **Resiliency & Efficiency**
- Security



# Resiliency

Istio adds fault tolerance to your application without any changes to code

```
// Circuit breakers

destination: serviceB.example.cluster.local
policy:
- tags:
    version: v1
  circuitBreaker:
    simpleCb:
      httpConsecutiveErrors: 7
      sleepWindow: 5m
      httpDetectionInterval: 1m
```

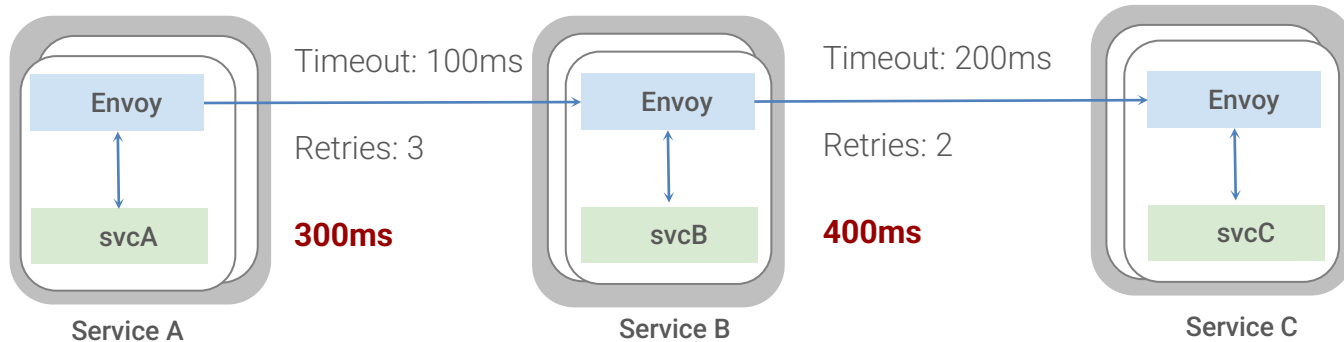
## Resilience features

- ❖ Timeouts
- ❖ Retries with timeout budget
- ❖ Circuit breakers
- ❖ Health checks
- ❖ AZ-aware load balancing w/ automatic failover
- ❖ Control connection pool size and request load
- ❖ Systematic fault injection

# Resiliency Testing

Systematic fault injection to identify weaknesses in failure recovery policies

- HTTP/gRPC error codes
- Delay injection






# Efficiency

- L7 load balancing
  - Passive/Active health checks, circuit breaks
  - Backend subsets
  - Affinity
- TLS offload
  - No more JSSE or stale SSL versions.
- HTTP/2 and gRPC proxying

# A network for services

- Traffic Control
- Visibility
- Resiliency & Efficiency
- **Security**

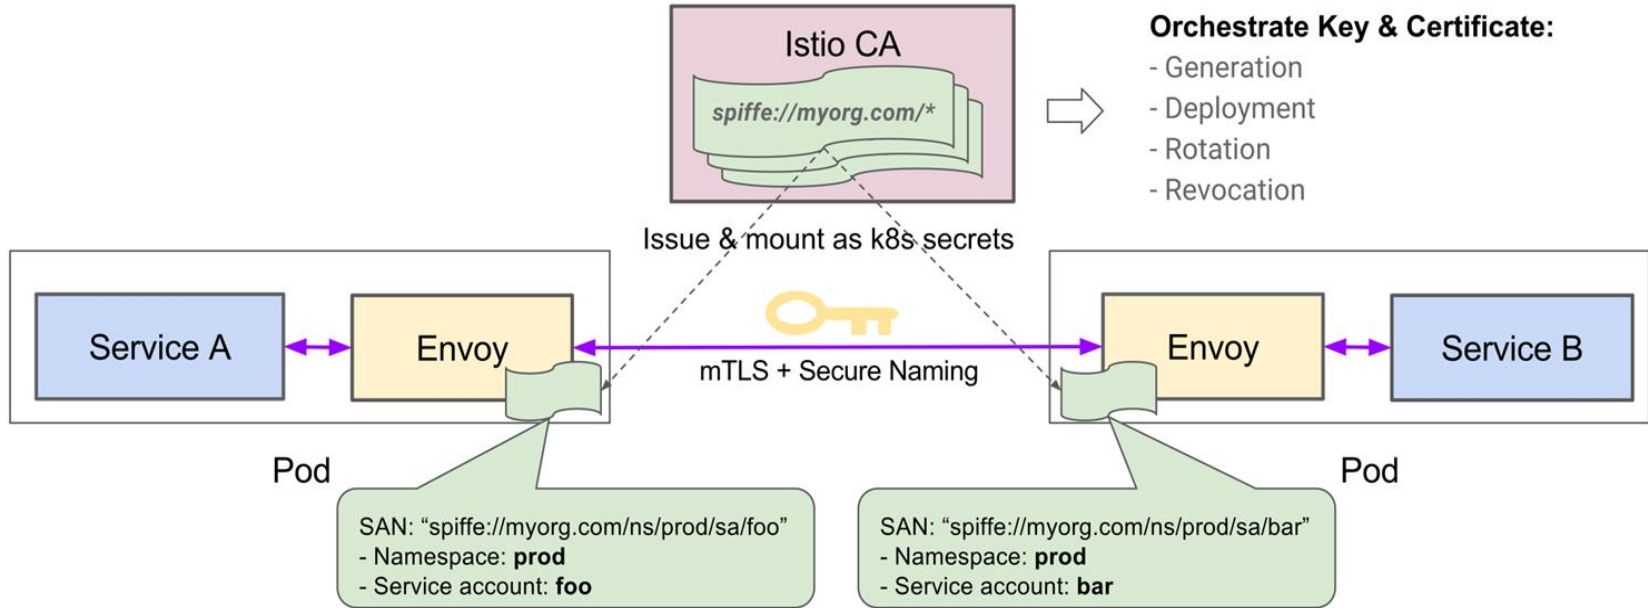


# Security

- Verifiable identity
- Secure naming / addressing
- Traffic encryption
- Revocation



# Security at Scale




# Istio vs Linkerd

	Istio	Linkerd
Source	Google + IBM + Lyft	CNCF
Support for VMs	Via Consul	No
Sidecar proxy	Envoy (mature and more features – ingress/egress and LB)	Own implementation (lightweight and fast)
Service relations	Included	Not included
Management console	Not included	Included
Policy Mgmt & Auth	Included	Not included
Ease of use	Harder	Easier

# Use cases demos


# Example microservices- based application





# 3 different HTTP front-end versions

```
kind: DestinationRule
spec:
  host: myhero-ui.myhero.svc.cluster.local
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
    - name: v3
      labels:
        version: v3
```






# Use case 1

## Routing to specific service version

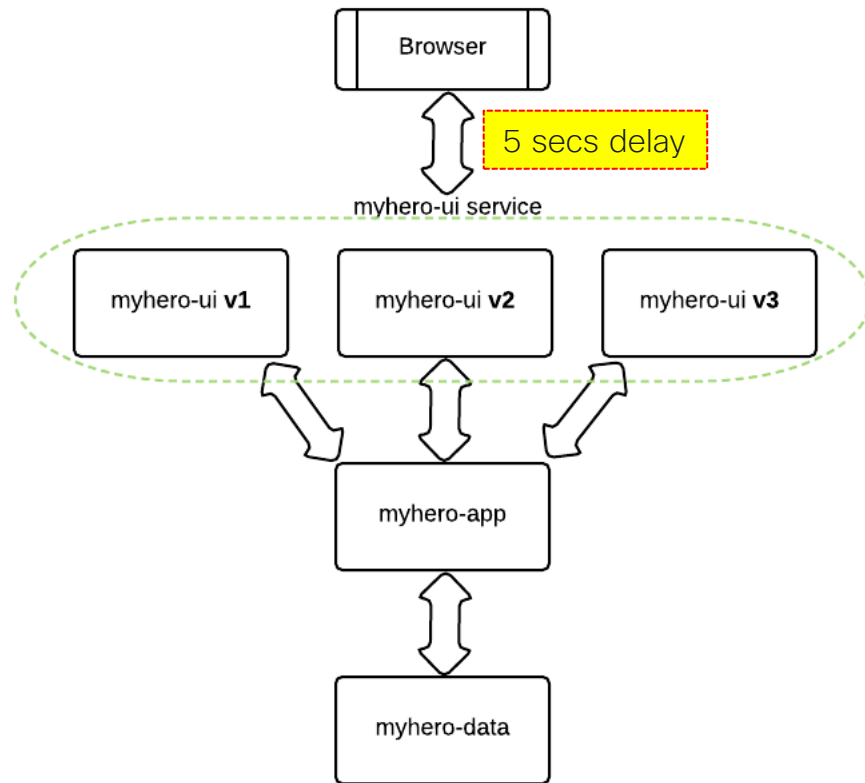
```
kind: VirtualService
spec:
  hosts:
  - "ui.xxx.com"
  gateways:
  - myhero-gateway
  http:
  - route:
    - destination:
        host: myhero-ui.myhero.svc.cluster.local
        subset: v1
```



# Use case 2

## Delay injection

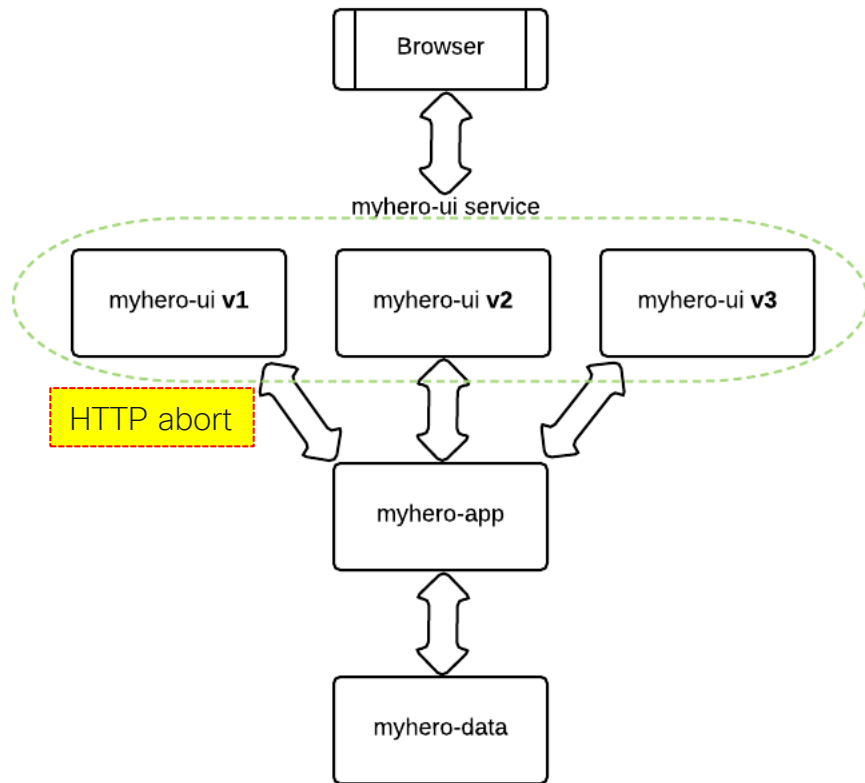
```
kind: VirtualService
spec:
  hosts:
  - "ui.xxx.com"
  gateways:
  - myhero-gateway
  http:
  - fault:
      delay:
        percent: 100
        fixedDelay: 5s
    route:
    - destination:
        host: myhero-ui.myhero.svc.cluster.local
        subset: v1
```



# Use case 3

## HTTP abort injection


```
kind: VirtualService
spec:
  hosts:
  - "ui.xxx.com"
  gateways:
  - myhero-gateway
  - mesh
  http:
  - fault:
      abort:
        percent: 100
        httpStatus: 500
    route:
    - destination:
        host: myhero-app.myhero.svc.cluster.local
```



# Use case 4

## Gradual migration


```
kind: VirtualService
spec:
  hosts:
  - "ui.xxx.com"
  gateways:
  - myhero-gateway
  http:
  - route:
    - destination:
        host: myhero-ui.myhero.svc.cluster.local
        subset: v1
      weight: 50
    - destination:
        host: myhero-ui.myhero.svc.cluster.local
        subset: v3
      weight: 50
```



# Use case 4

## Gradual migration


```
kind: VirtualService
spec:
  hosts:
  - "ui.xxx.com"
  gateways:
  - myhero-gateway
  http:
  - route:
    - destination:
        host: myhero-ui.myhero.svc.cluster.local
        subset: v3
      weight: 100
```



# Use case 5

## Traffic mirroring

```
kind: VirtualService
spec:
  hosts:
  - "ui.xxx.com"
  gateways:
  - myhero-gateway
  http:
  - route:
    - destination:
        host: myhero-ui.myhero.svc.cluster.local
        subset: v1
    mirror:
        host: myhero-ui.myhero.svc.cluster.local
        subset: v2
```



# Calisti

## The Cisco Service Mesh Manager

An enterprise ready Istio platform for DevOps and SREs that automates lifecycle management and simplifies connectivity, security & observability for microservice based applications.

# Calisti Core Value Propositions

Calisti – Simplifying Cloud Native Application Connectivity and Observability

## Turnkey solution for Kafka

Automatic mTLS encryption for all components

**Extend observability to event driven messaging applications**

Graceful upscale/downscale

Self-healing

Alert-based storage expansion

Automated authentication

## Integrated Service Observability

Provide actionable insights to operators; not just metrics

Detect outliers that affect service and workload health

Generate alerts to maintain application service level objectives

Speed up root cause analysis of issues with timelines and topology

Dive deep into communications with Traffic Tap, while maintaining end-to-end encryption

Visualize communications in context with distributed tracing

## Simplified Mesh Management

Automate Istio management to remove toil and risk from installs and updates

Validate configurations prior to deployment to maximize availability

Migrate virtual machines to your mesh to begin your cloud native journey

Add workloads flexibly to a mesh to support evolving business needs

Support complex multi-mesh architectures with ease

Integrate external services to provide additional features and functions

De-risk upgrades with canary deployments

Provide superior levels of application experience with circuit breakers



# Conclusion



# What did we learn?

- What is a service mesh
- Why you might need one
- Istio's architecture
- How does it work
- What are its capabilities
- How you can use Istio for Traffic Management



*“Service Meshes actually are the new *black*.”*

Questions?



[cs.co/julidevops](https://cs.co/julidevops)

[cs.co/julioblog](https://cs.co/julioblog)



# Complete your Session Survey

- Please complete your session survey after each session. Your feedback is important.
- Complete a minimum of 4 session surveys and the Overall Conference survey (open from Thursday) to receive your Cisco Live t-shirt.
- All surveys can be taken in the Cisco Events Mobile App or by logging in to the Session Catalog and clicking the "Attendee Dashboard" at <https://www.ciscolive.com/emea/learn/sessions/session-catalog.html>



# Continue Your Education



Visit the Cisco Showcase for related demos.



Book your one-on-one Meet the Engineer meeting.



Attend any of the related sessions at the DevNet, Capture the Flag, and Walk-in Labs zones.



Visit the On-Demand Library for more sessions at [ciscolive.com/on-demand](https://ciscolive.com/on-demand).



The bridge to possible

# Thank you

CISCO *Live!*



CISCO *Live!*

ALL IN