



Introduction to Hyperconverged with Cisco Hyperflex

Serge Charles TSA
BRKCOM-1110



Agenda

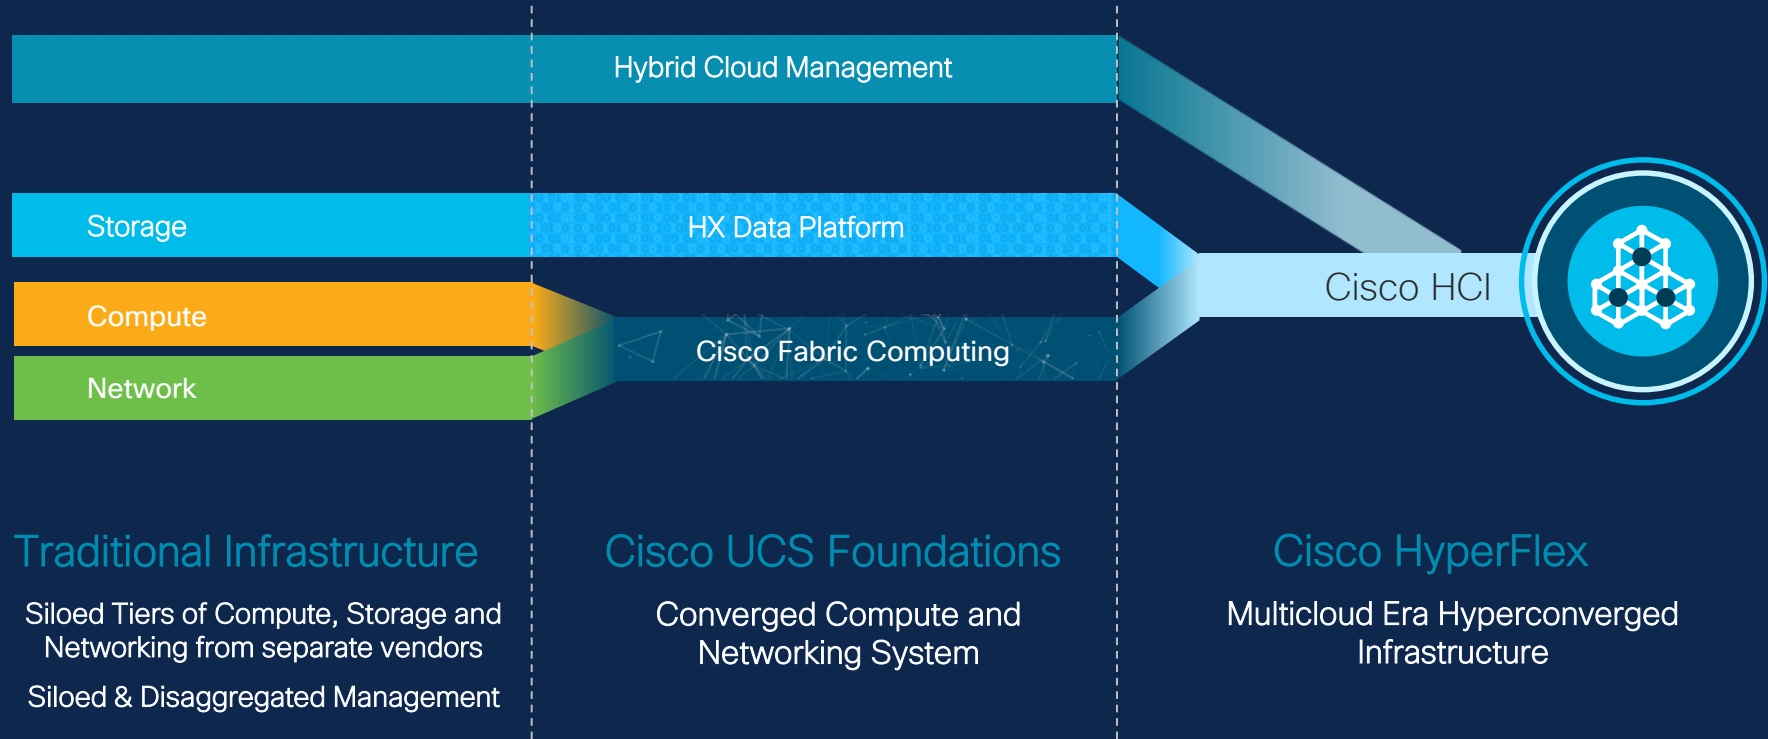
- Cisco Hyperflex Data Platform
- Platform Resiliency and scaling options
- Deployment and management options
- HX Connect demo



Cisco Hyperflex Data Platform

Complete Hyperconvergence

Compute, Network, HCI Software Engineered Together



The Cisco HX Data Platform

**HX LOG STRUCTURED FILE SYSTEM
DESIGNED SPECIFICALLY FOR HYPERCONVERGENCE**



DISTRIBUTED

Object-Based File System
Architected for Scale-Out,
Distributed Storage



ADVANCED DATA SERVICES

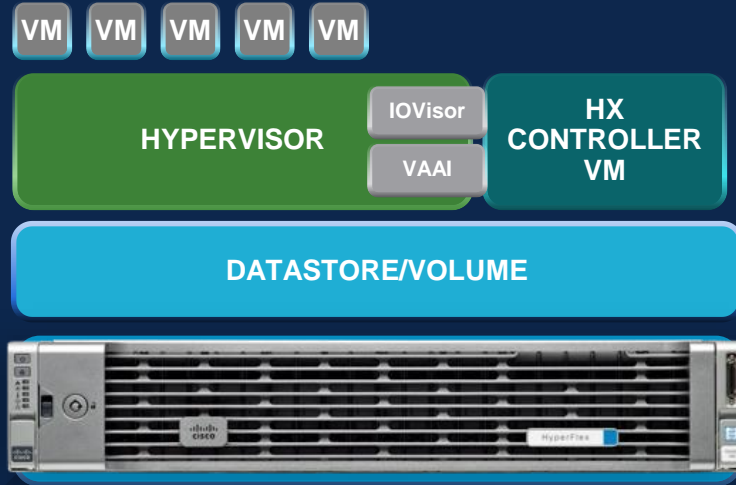
Built Into File
System Architecture



FUTURE READY

Designed for Containers
and Next-generation
Applications

Inside HX Data Platform Node

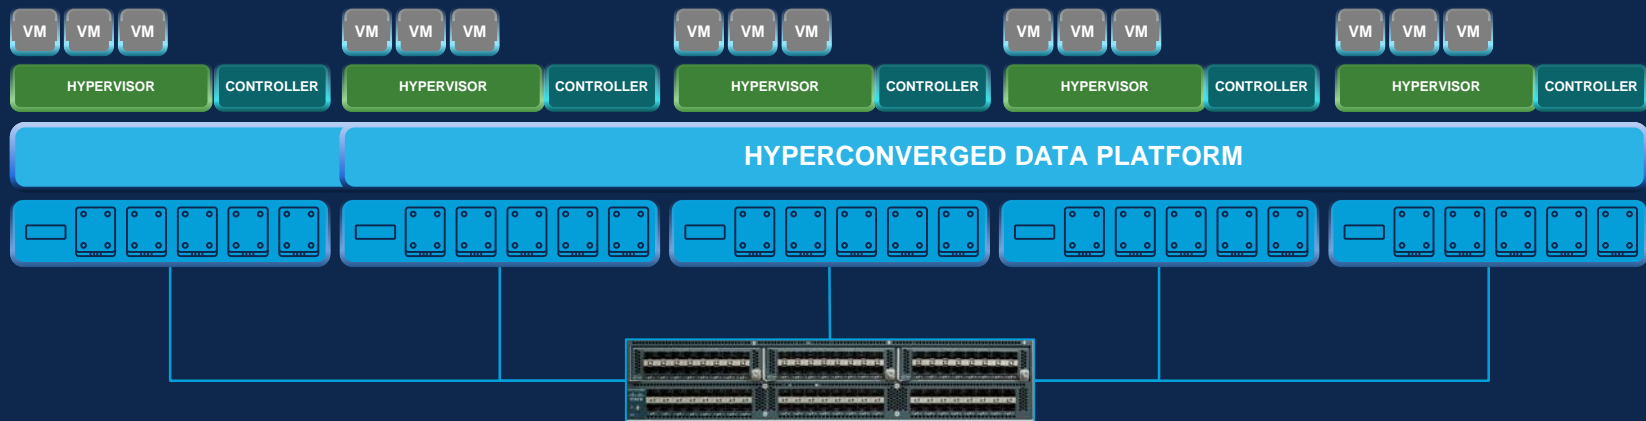


HX Controller VM Assumes
Direct Access of Local Storage

IOVisor Module Presents Pooled
Storage to HyperVisor and Stripes IO

Data Services are Offloaded
to HX Data Platform

Hyperconverged Scale Out and Distributed File System



Start with as Few
as Three Nodes

Hyperconverged
Data Platform
Installs in Minutes

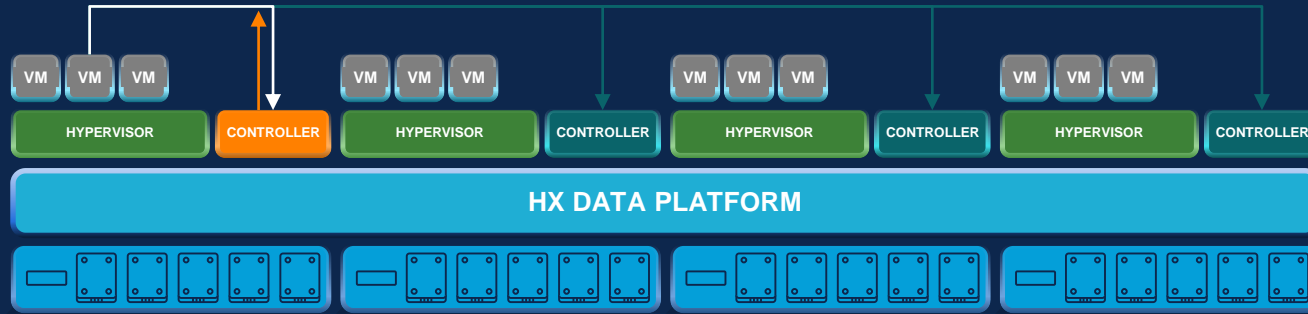
Network Fabric
Policy Configures
QoS Settings

Add Servers, One
or More at a Time

Distribute and
Rebalance Data
Across Servers
Automatically

Retire Older
Servers

Dynamic Data Distribution




Systems Built on Conventional File Systems Write Locally, Then Replicate, Creating Performance Hotspots

HX Data Platform Stripes Data Across All Nodes Simultaneously, Leveraging Cache Across all SSDs for Fast Writes

Balanced Space Utilization: No Data Migration Required Following a VM Migration

Capacity and Network Utilization

- HX balances space utilization: no data migration required following a VM migration




Balanced Space Utilization


No Data Migration on VM Migration
Less Stress on Network

HyperFlex All-Flash Performance Consistency

Each dot represents a Virtual Machine and its average IOPS over an hour of load



**Consistent Performance for
All VMs on the Cluster**




**Wildly Inconsistent VM
Performance Across the Cluster**




Continuous Data Optimization

Log-Structured File System Yields More Efficient Data Optimization




Inline Compression

30–50% space savings



Inline Deduplication

20–30% space savings



No Special Hardware
No Performance Impact
No Config lock-in
No Additional License

Lower Cost

UCS: Embedded Automation

Rapid Deployment of HyperFlex with Service Profiles



Server SME



Network SME



Storage SME

Server Policy
Storage Policy
Network Policy
Virtualization Policy
Application Profiles

Service Profile Templates Pre-Defined at the Factory

Uplink port configuration,
VLAN, VSAN, QoS, and
EtherChannels

Server port configuration
including LAN and SAN
settings
Network interface card (NIC)
configuration: MAC address,
VLAN, and QoS settings;
host bus adapter HBA configuration:
worldwide names (WWNs), VSANs,
and bandwidth constraints;
and firmware revisions


Unique user ID (UUID),
firmware revisions,
and RAID controller settings

Service profile assigned to server,
chassis slot, or pool

Quick Deployment



HX Ready



1

Subject Matter Expert
Define Policies



2

Policies Used to Create
Service Profile Templates



3

Service Profile Templates
Create Service Profiles



4

Associating Service
Profiles with Hardware
Configures Servers
Automatically

Fabric Centric Design



High Performance

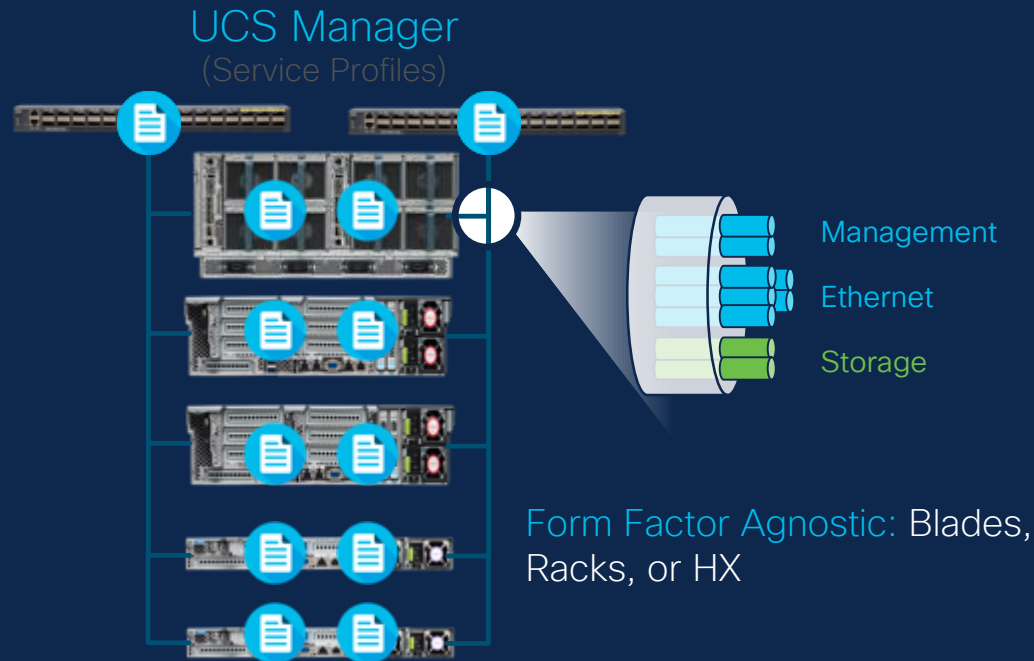
40 GB/s Ethernet; 320 GB/s per Chassis

Unified Fabric

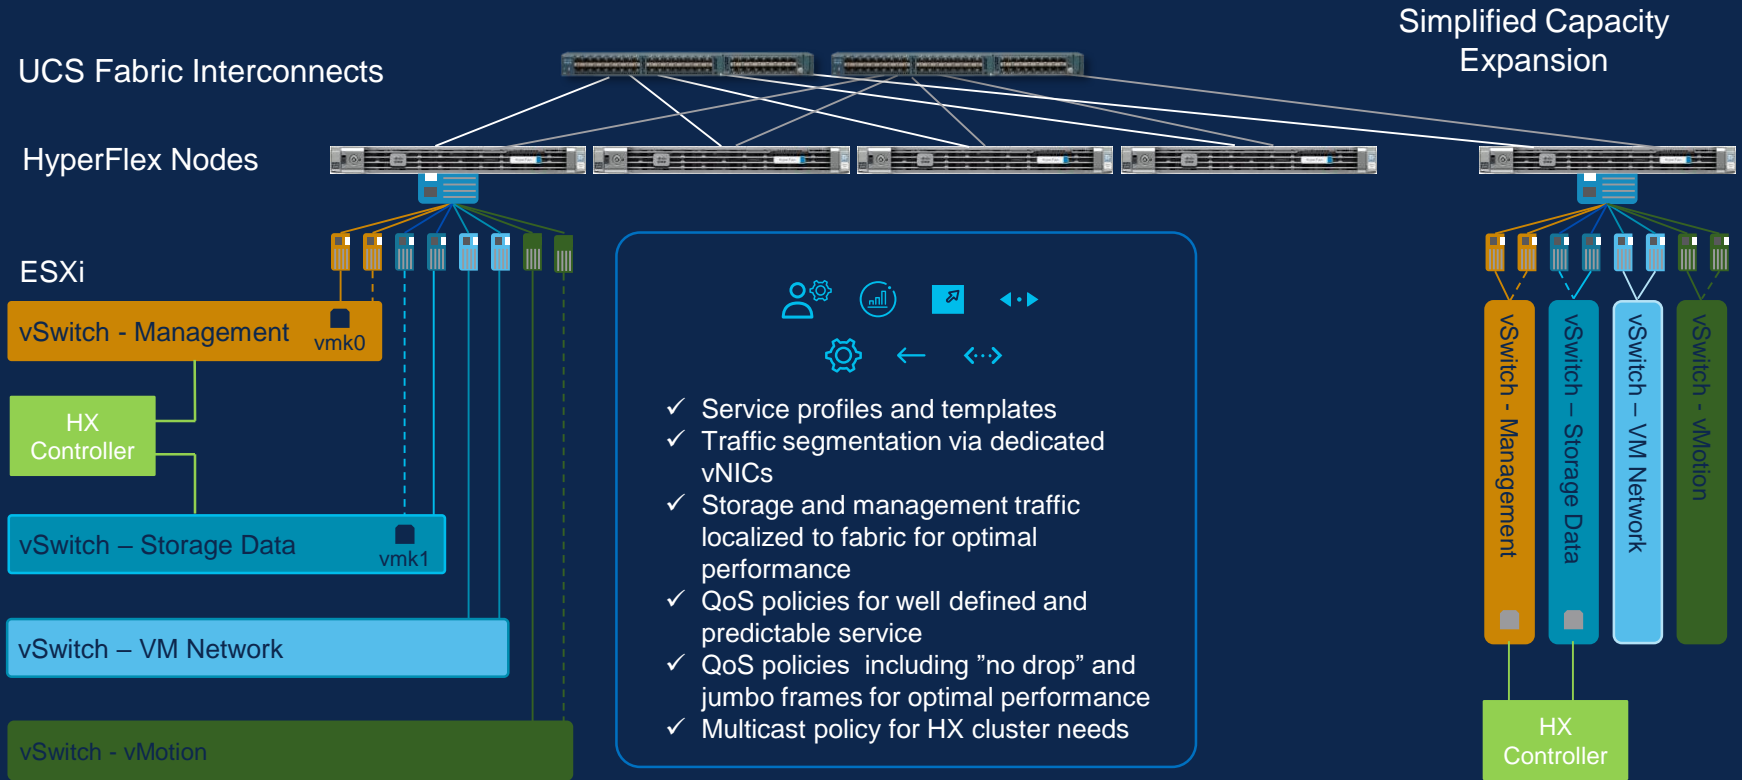
Single Cable for Network, Storage, and Management Traffic

Easy to Scale

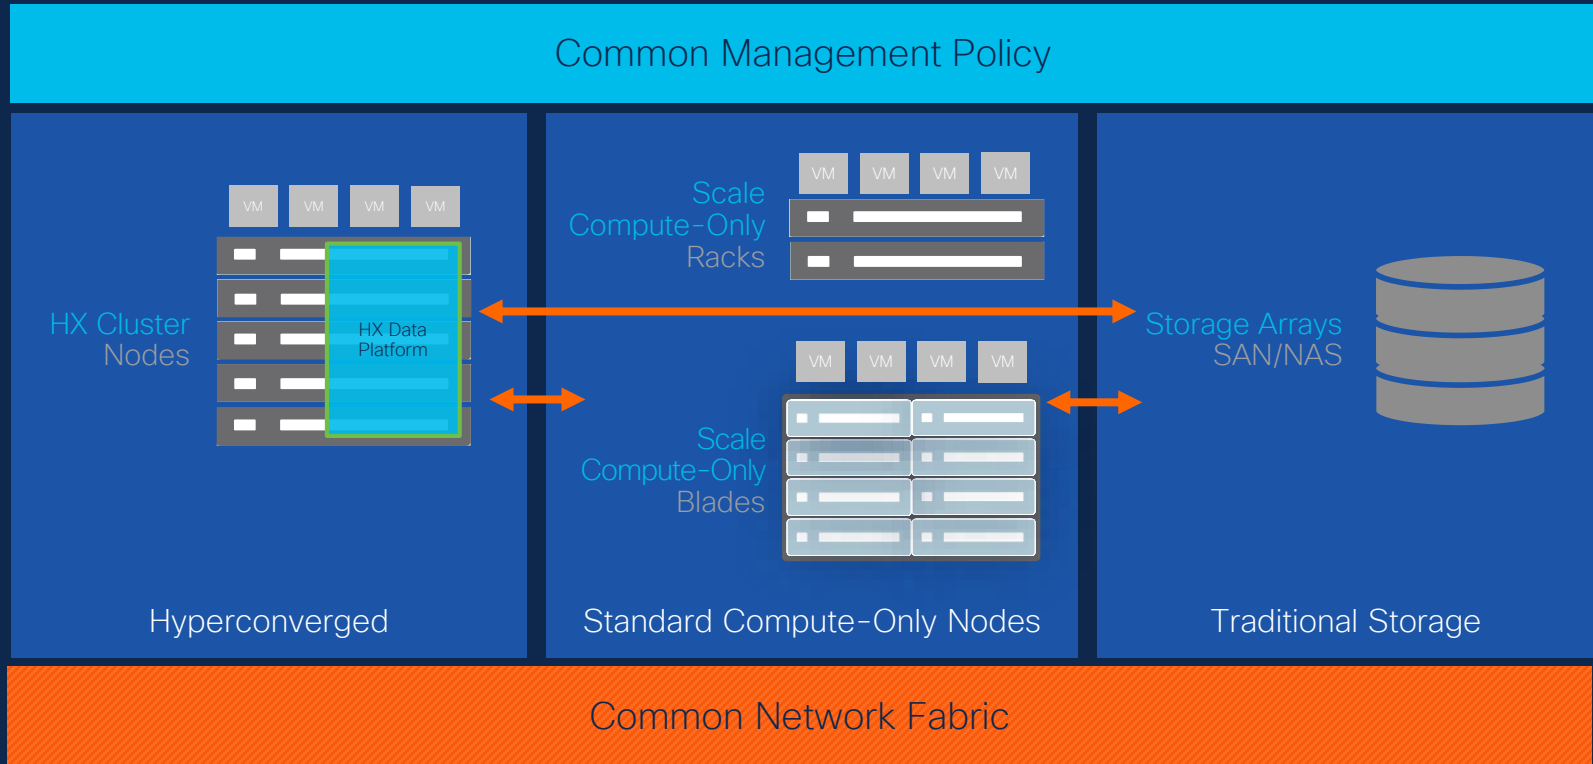
Single Point of Management:
Add Cables for Bandwidth
vs. Fabric Type



Current Networking model for HX



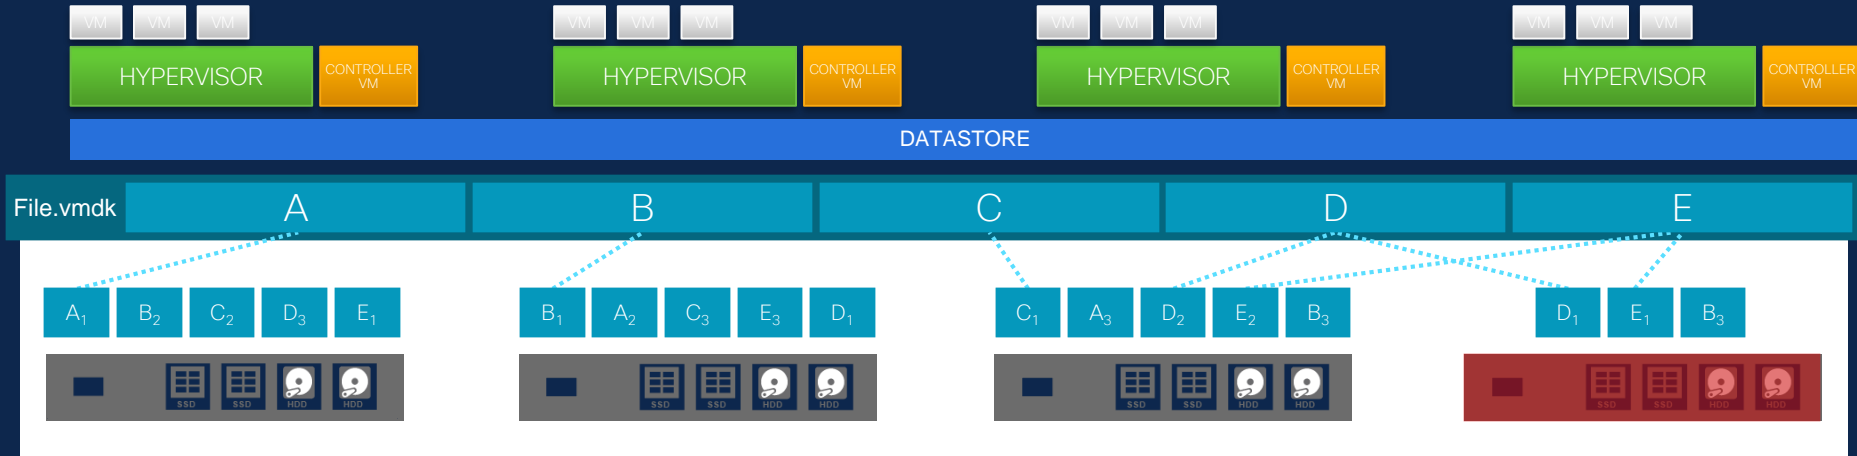
Integrate Into Your Existing Infrastructure





Platform Resiliency and scaling options

Non-Disruptive Operations



- Stripe blocks of a file across servers
- Replicate one or two additional copies to other servers
- Handle entire server or disk failures
- Restore back to original number of copies
- Rebalance VMs and data post replacement
- Rolling software upgrades

Data Protection and High Availability

Data Protected by Replication of Data Across the Cluster Nodes


Replication Factor 3 (RF3)

Default and Recommended is Replication Factor = 3

Every block is written to 3 different nodes in the cluster

Higher availability to survive multi-point failures;
Higher device protection

Reduces raw disk capacity to 33%



Note: RF3 is strongly recommended for high availability

Replication Factor	3 or 4 Node Cluster	5+ Node Cluster
3	Simultaneous Failures Supported: 1 node / 2 drives*	Simultaneous Failures Supported: 2 nodes / 2 drives*

*drives across different nodes

Data Protection and High Availability

Data Protected by Replication of Data Across the Cluster Nodes


Replication Factor 2 (RF2)

Default and Recommended is Replication Factor = 3

Every block is written to 2 different nodes in the cluster

Lower availability to survive failures; Lower device protection

Reduces raw disk capacity to 50%



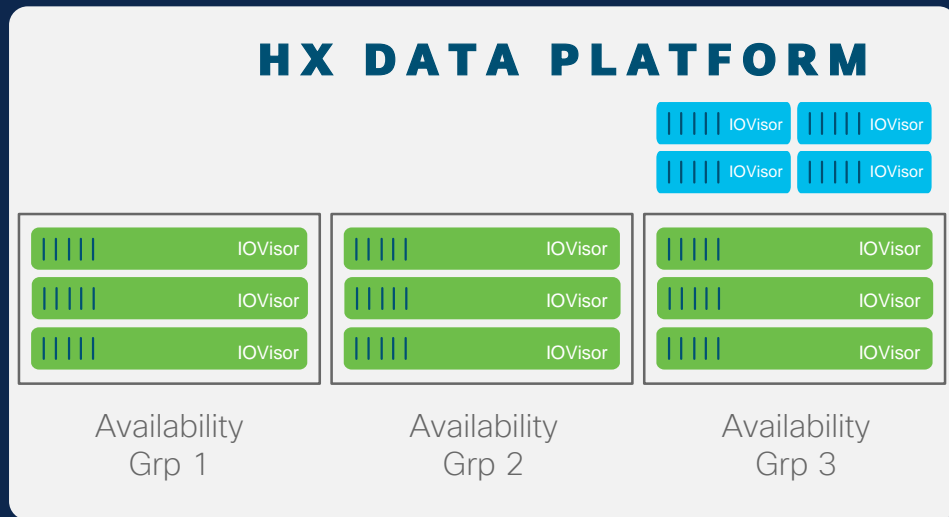
Note: RF3 is strongly recommended for high availability

Replication Factor	3 or 4 Node Cluster	5+ Node Cluster
2	Simultaneous Failures Supported: 1 node / 1 drive	Simultaneous Failures Supported: 1 node / 1 drive

*drives across different nodes

Logical Availability Zones (LAZ)

- Cluster Scale With High Availability
- Increased resiliency without added manageability overhead
- How does it work?
 - HX nodes grouped into logical “availability groups” (N/A for compute nodes)
 - HXDP never places 2 copies of the data in the same availability group
 - Clusters with LAZ can survive > 2 simultaneous node failures without data loss or loss of availability
 - Tolerate more independent failures



LAZ Failure Scenario

Cluster State: **Online**



LAZ: Off



LAZ Failure Scenario

Cluster State: **Online**



LAZ: **On**



HyperFlex Data Protection

Fast and Flexible Native Snapshots



- Pointer-based snapshots
 - Space-efficient with no performance penalty vs. VMware Redo Log Snaps
 - Fast creations and deletions
- Fine-grained or coarse-grained
 - VM-level or VM folder-level
- VAAI-integrated
 - Quiesced and crash-consistent
- Use vCenter Snapshot Manager
- Policy-based schedules and retention

Schedule Snapshot for workload_vm_a1

☒ Enable Hourly Snapshot

Hourly Snapshot

Schedule

Start At: 10:00 AM

Until: 05:00 PM

On: ☐ Sunday ☒ Monday ☒ Tuesday ☒ Wednesday ☒ Thursday ☒ Friday ☐ Saturday

Retention

☒ A maximum of 10 snapshot per VM

☐ Never Expires

☒ Enable Daily Snapshot

Daily Snapshot

Schedule

Start At: 09:00 PM

On: ☐ Sunday ☒ Monday ☒ Tuesday ☒ Wednesday ☒ Thursday ☒ Friday ☐ Saturday

Retention

☒ A maximum of 7 snapshot per VM

☐ Never Expires

☒ Enable Weekly Snapshot

Weekly Snapshot

Schedule

Start At: 05:00 AM

On: ☐ Sunday ☐ Monday ☐ Tuesday ☐ Wednesday ☐ Thursday ☐ Friday ☒ Saturday

Retention

☒ A maximum of 4 snapshot per VM

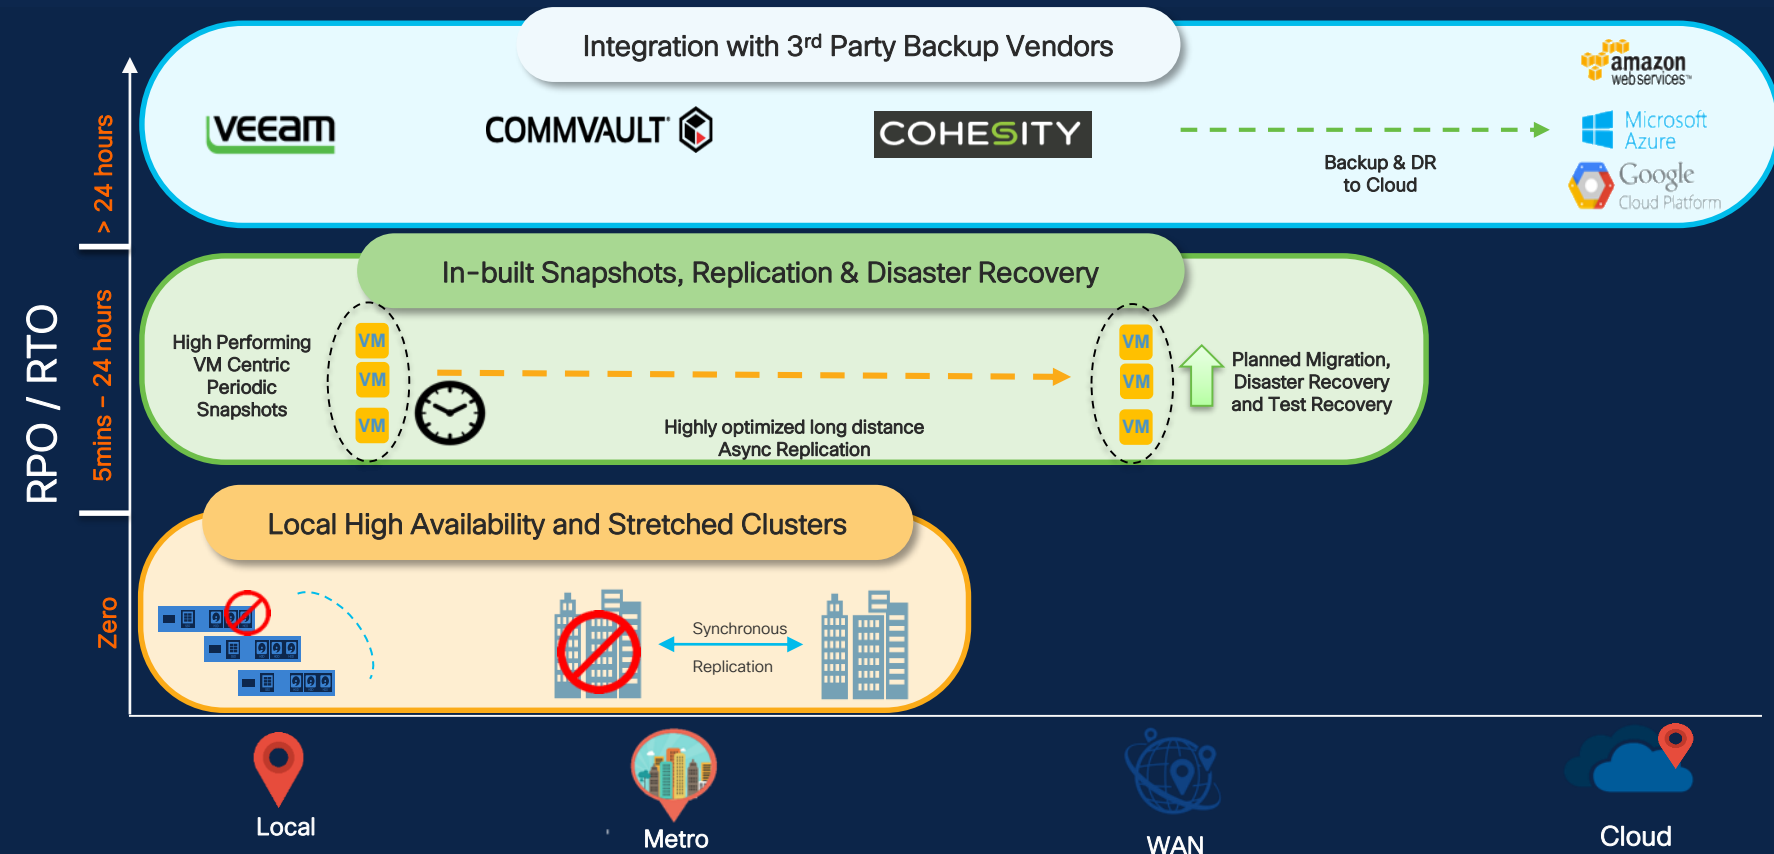
☐ Never Expires

OK Cancel

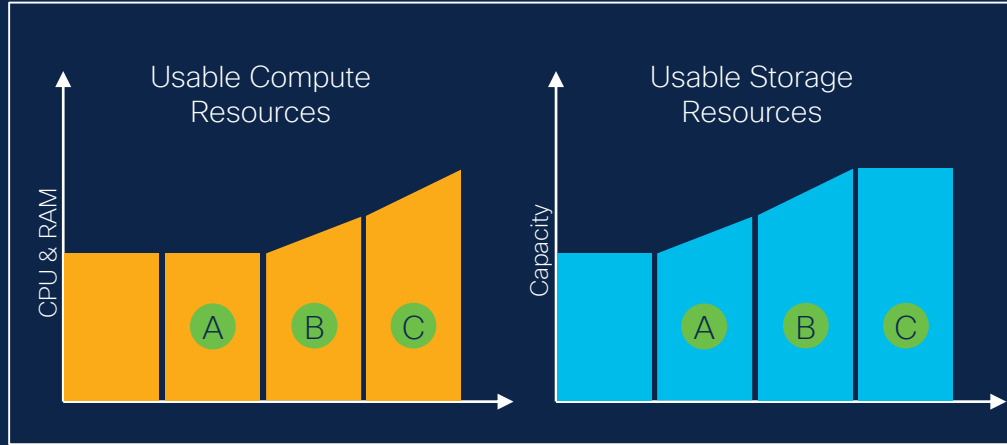
All This Functionality Enabled with the HX Data Platform Filesystem

HyperFlex Data Protection

Flexibility to meet business needs





Independent Scaling of Compute and Capacity





C Add Compute-Only Nodes


Non-HyperFlex hosts can connect to storage with IOVisor



A Scale Capacity Within Converged Nodes



B Add Converged Nodes



Scaling Options in HXDP 4.0.1a

VMware

SFF

(excluding All-NVMe)

Converged Nodes

3-32



Compute-Only Nodes

0-32

✓ Expansion Supported

2:1

Max ratio **Compute** to
Converged Nodes*

Max Cluster Size

64

LFF

Converged Nodes

3-16



Compute-Only Nodes

0-32

✓ Expansion Supported

2:1

Max ratio **Compute** to
Converged Nodes*

Max Cluster Size

48

! Hybrid Only

All-NVMe

Converged Nodes

3-16



Compute-Only Nodes

0-16

✓ Expansion Supported

1:1

Max ratio **Compute** to
Converged Nodes*

Max Cluster Size

32

! Requires Enterprise License

HyperV

SFF

Converged Nodes

3-16



Compute-Only Nodes

0-16

✓ Expansion Supported

1:1

Max ratio **Compute** to
Converged Nodes*

Max Cluster Size

32

LFF

Converged Nodes

3-16**



Compute-Only Nodes

0-16

✓ Expansion Supported

1:1

Max ratio **Compute** to
Converged Nodes*

Max Cluster Size

32

! Hybrid Only

* 2:1 - Enterprise license (HXDP-P) if # of **Compute** > # of **Converged** Nodes.

* 1:1 - Standard license (HXDP-S) if # of **Compute** <= # of **Converged** Nodes.

** Max LFF cluster size currently limited to 8 **Converged** nodes when using new LFF 12TB drives.

Scaling Options in HXDP 4.0.1a (cont.)

Edge

SFF
(excluding All-NVMe)

Converged Nodes
2-4



Expansion Not
Currently Supported

1GbE and 10GbE
Network Options
Available

Single and Dual
Switch Support

Deploy and Manage
via Cisco Intersight

! HX*220 Nodes Only

Stretched Cluster**

SFF

Converged Nodes
2-16 / Site (4-32 / Cluster)



Compute-Only Nodes
0-21 / Site (0-42 / Cluster)
[Limited by Max Cluster size]

✓ Expansion Supported***

2:1

Max ratio **Compute** to
Converged Nodes*

Max Cluster Size
32/Site | **64**/Cluster

LFF

Converged Nodes
2-8 / Site (4-16 / Cluster)



Compute-Only Nodes
0-16 / Site (0-32 / Cluster)

✓ Expansion Supported***

2:1

Max ratio **Compute** to
Converged Nodes*

Max Cluster Size
24/Site | **48**/Cluster

! Hybrid Only

* 2:1 – Enterprise license (HXDP-P) if # of **Compute** > # of **Converged** Nodes.

* 1:1 – Standard license (HXDP-S) if # of **Compute** <= # of **Converged** Nodes.

** Stretched cluster requires Enterprise license (HXDP-P)

*** Requires uniform expansion of converged nodes across both sites



Deployment and management options

Cluster Deployment

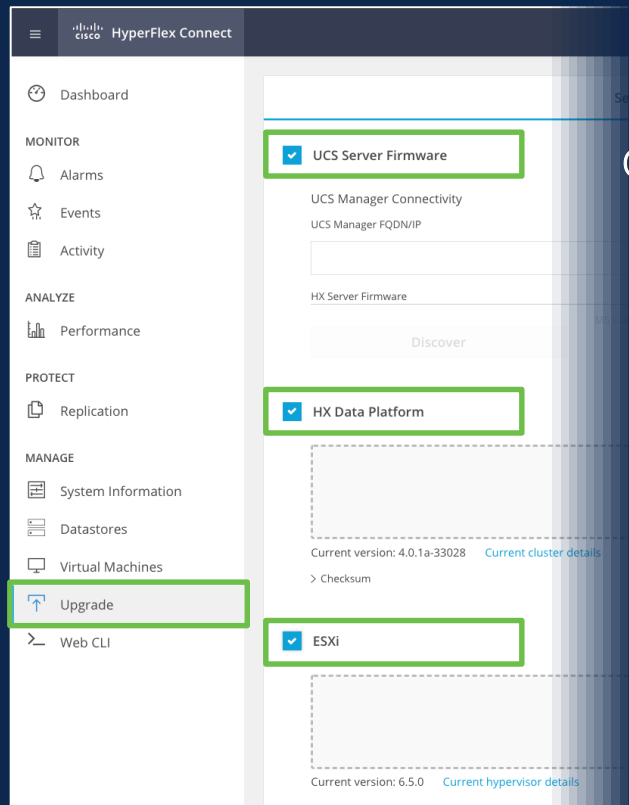


- Download latest OVA from Cisco.com
- Deploy on existing infrastructure
- Run single cluster deployment



- Deploy from anywhere!
- No need for existing infrastructure or OVA setup
- Latest version always available
- Reusable policy for rapid & consistent deployment
- Simple ramp-up of large HX projects with simultaneous background deployment


One-Click Full Stack HyperFlex Upgrades



One-click, fully integrated stack upgrades via HX Connect.



*Non-Disruptive
online upgrades!*





HX Connect demo



Enterprise Application Ready

Enterprises Run Mission Critical Apps on HyperFlex



#1 in Performance



1

3x higher
VM density

2

3x reduced
read/write latency

3

7:1 reduction
IOPS variability



3x better
TCO

More workloads
On Hyperconverged

Predictable
End User Experience

Enterprise Grade HCI

+35%

Clusters Running
Databases(ASUP data)

25+%

Enterprise scale
deployments

10+

New CVDs/Solution
Guides for DB Apps



Thank you





You make **possible**