

The background is a vibrant, abstract graphic. It features a central bright white light source from which numerous colorful rays emanate, creating a sunburst or starburst effect. The rays transition through a spectrum of colors including yellow, orange, red, and various shades of blue and green. Overlaid on this are large, flowing, wavy shapes in similar colors, giving the impression of liquid or smoke being illuminated by the light. The overall effect is dynamic and energetic.

cisco *Live!*

Let's go

#CiscoLive



The bridge to possible

Routing IPv6 In VXLAN BGP EVPN Fabrics

Venkat Kirishnamurthy , Principal Architect, CX
BRKDCN-2682



Cisco Webex App

Questions?

Use Cisco Webex App to chat with the speaker after the session

How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space

Webex spaces will be moderated by the speaker until June 9, 2023.



https://ciscolive.ciscoevents.com/ciscolivebot/#_BRKDCN-2682

Agenda

- Introduction
- Why IPv6 in the Data Center?
- Introduction to Underlay and Overlay Networks
- IPv6 Unicast Routing in the Overlay with IPv4 Underlay
- IPv6 Unicast Routing in the Overlay with IPv6 Underlay
- Conclusion

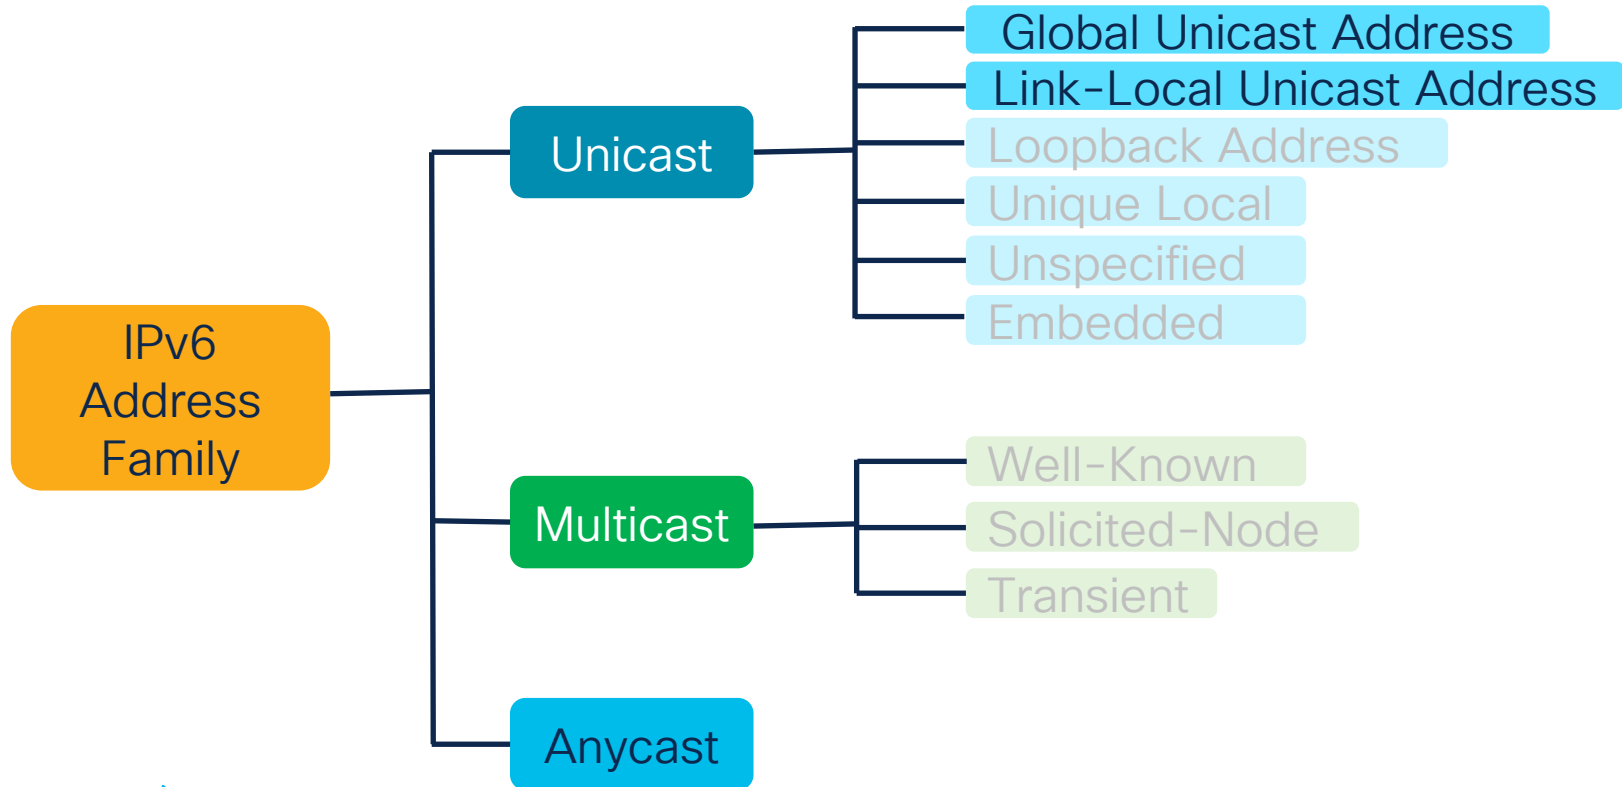
IPv6 in The Data Center

Why IPv6 in the DC?

Key IPv6 Drivers in the Data Center

- We are running out of IPv4 addresses
- Explosion of IP address requirements in the Data center due to:
 - Disaggregation of application into micro services
 - Proliferation of Kubernetes and Containers in the Data center
- Massive scale service provider /Cloud Data center deployments
- Federal Government Mandate / Requirements
 - DOD requirements to transition to IPv6 by 2025
 - FedRAMP Mandatory IPv6 only requirement

IPv6 Address Hierarchy




IPv6 Unicast Address

Global Unicast Address

Global unicast addresses (GUAs) are:


- Routable globally and reachable in the IPv6 Internet
- Same as routable IPv4 addresses



IPv6 Unicast Address

Link-Local Address

- Link-Local address : unicast address confined to a subnet/segment; Typically, auto created
- Non-Routable address
- Used to communicate with routers, DHCPv6 servers in the same subnet/segment




IPv6 Multicast Address

Well-Known Multicast Address

Well-Known Multicast Address assigned by IANA :

- Uses the prefix **ff00::/12**
- Are predefined or reserved multicast addresses.
 - ff02::1 : All IPv6 Devices
 - ff02::2 : All IPv6 routers




T- Flag: 0 - Well-defined; 1 - Transient/ Dynamically defined

IPv6 Multicast Address

Solicited-node multicast Address

Solicited-mode Multicast Address:

- Used by Neighbor Discovery Protocol for Address resolutions and Duplicate Address Detection (DAD)
- Auto created using a special mapping of the device unicast address with the solicited-node multicast prefix **ff02:0:0:0:0:1:ff00::/104**
- Auto created for every unicast IPv6 address on a node




HW that supported VXLAN with IPv6 in underlay

Cisco Nexus 9332C	Cisco Nexus 9364C	Cisco Nexus 9300-EX
Cisco Nexus 9300-FX	Cisco Nexus 9300-FX2	Cisco Nexus 9300-FX3
Cisco Nexus 9300-FXP	Cisco Nexus 9300-GX\	Cisco Nexus 9300-GX2

Introduction to Underlay and Overlay Networks



Underlay Network Architecture



Unicast Protocols



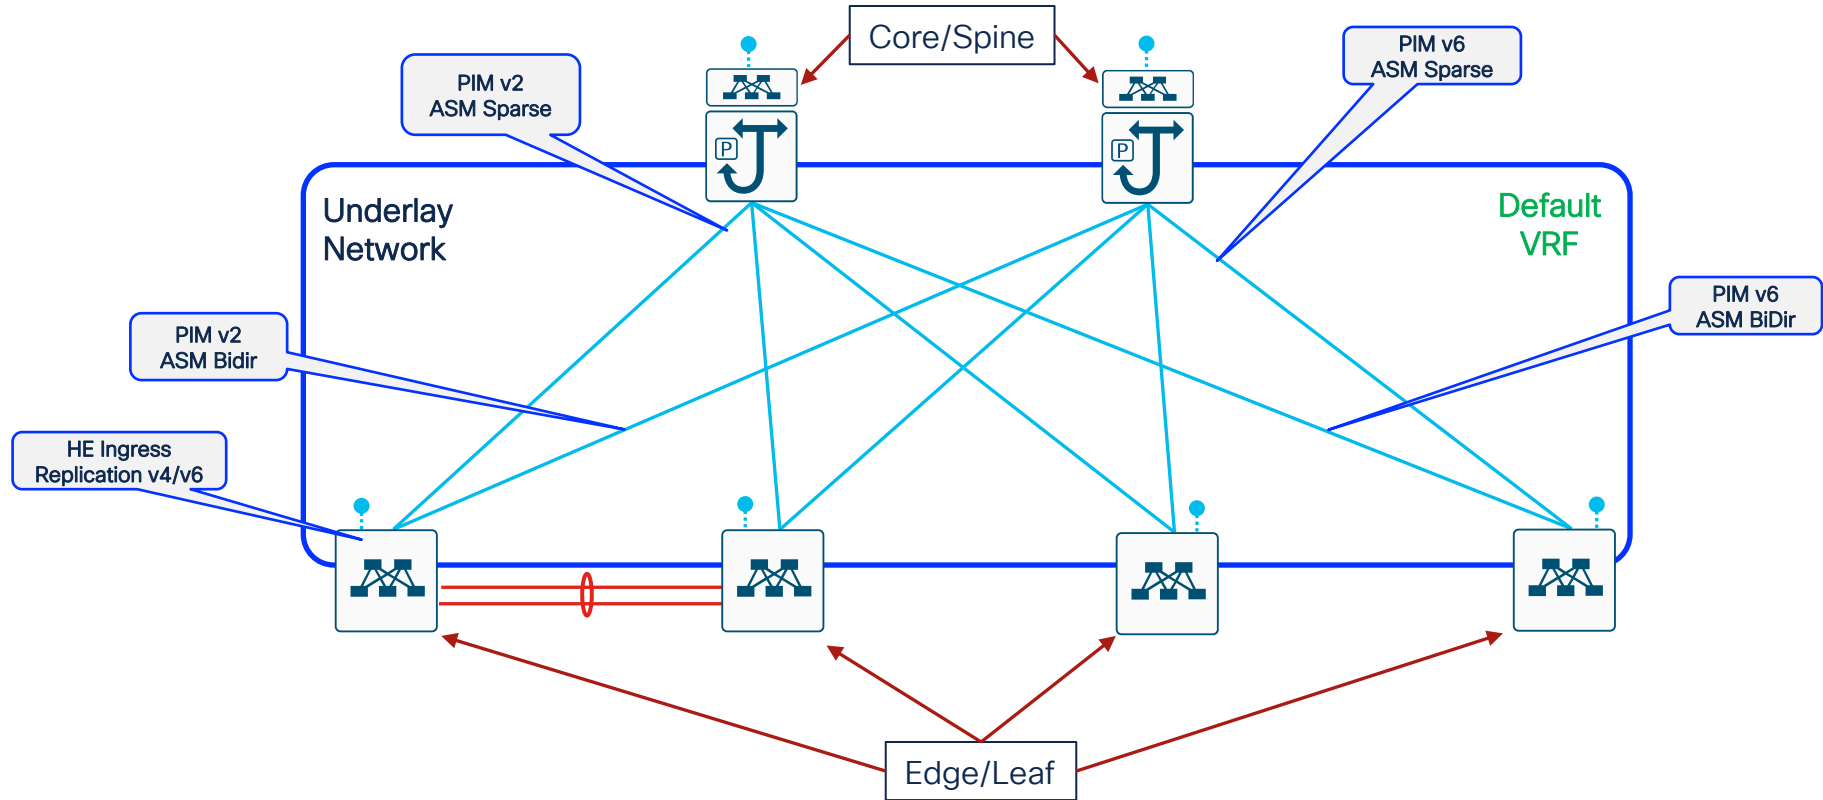
Underlay Network Protocols

Unicast Protocol


Device Role	Interface	Protocol	Function
Spine, Leaf	Loopback0	OSPF/ISIS/BGP	Router ID
Leaf	Loopback1	OSPF/ISIS/BGP	NVE Source Int
Spine, Leaf	Fabric Interface	OSPF/ISIS/BGP	Switch Peering
Leaf (vPC)	vPC backup SVI	OSPF/ISIS/BGP	vPC peer link routing

Underlay Network Protocols


Multicast Protocols



Overlay Network Architecture



Unicast




Overlay Network Protocols

Unicast Protocol

Device Role	Interface	Protocol	Function
Leaf	L2 VNI SVI	BGP EVPN	Anycast Gateway Subnet
Leaf	L3 VNI SVI	BGP EVPN	VRF Routing
Spine	N/A	BGP EVPN	Route Reflector
Border Leaf/Gateway	External Interface	OSPF/ISIS/eBGP	External Routing

Overlay Network Protocols

Multicast



Overlay Network Protocols

Multicast Protocol

Device Role	Interface	Protocol	Function
Leaf	L2 VNI SVI	PIM v2/PIM v6 mVPN	Multicast Routing
Leaf	L3 VNI SVI	PIM v2/PIM / mVPN	Multicast Routing
Leaf	L2 VNI SVI	IGMPv2/v3	Host Join/Leave
Leaf	L3 VNI SVI	IGMPv2/v3	Host Join/Leave
Border Leaf/Gateway	External Interface	PIM v2/PIM v6	External Multicast Routing

VXLAN IPv6 Deployment modes in the Data Center

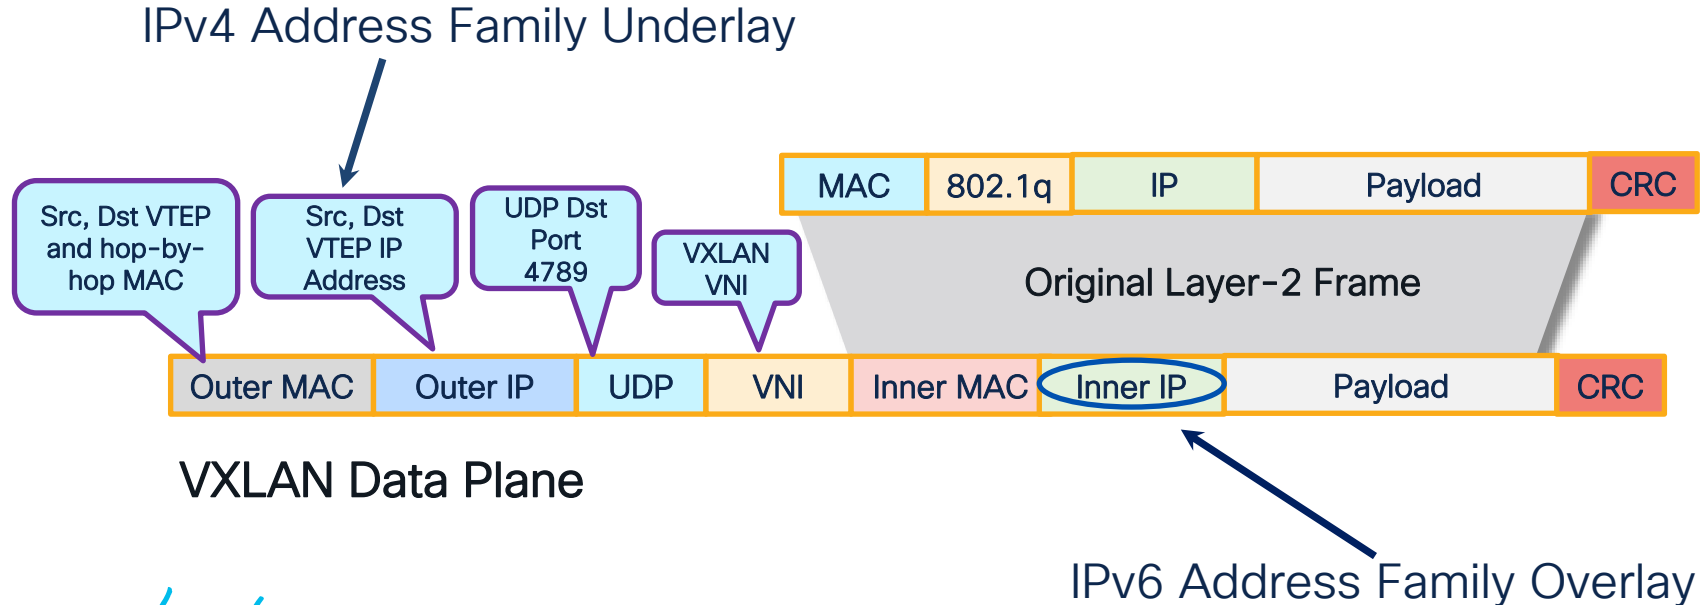
VXLAN IPv6 routing in the Data Center

- IPv6 routing in a IPv4 Underlay
 - IPv6 specific configs needed in a IPv4 Underlay
 - IPv6 routing in the overlay with IPv4 in the Underlay
- IPv6 routing in a IPv6 Underlay
 - VXLAN IPv6 deployment in the underlay
 - IPv6 routing in the overlay with IPv6 in the Underlay

Adding configs to IPv4 underlay to support IPv6 in the Overlay

VXLAN Headers IPv6 transport thru IPv4 Underlay

The IPv6 packet is encapsulated in the overlay as IPv4 UDP packets and uses IPv4 routing to transport the VXLAN encapsulated traffic.



IPv6 Overlay Configuration Steps

- Synchronize IPv6 Neighbor discovery in the vPC domain
- Configure IPv6 gateway address on the L2VNI SVI and host
- Enable IPV6 Forwarding for L3VNI VLANs
- The IPv6 Address Family must be enabled in the VRF and BGP EVPN Configurations

IPv6 Configuration on SVIs

- Configure IPv6 ND synchronize between vpc domain.
- Assign IPv6 GW address on L2VNI SVI and enable IPv6 forwarding on SVIs for L3 VNI VLANs.
- The IPv6 address “**use-link-local-only**” serves the same purpose as “**ip forward**” for IPv4

L2VNI AND VPC CONFIG

```
vpc domain 1
<snip>
ipv6 nd synchronize
<snip>
interface Vlan13
  no shutdown
  vrf member tenant-1
  ip address 192.168.13.254/24 tag 12345
  ipv6 address 2023:192:168:13::254/64 tag 12345
  fabric forwarding mode anycast-gateway
```

L3VNI CONFIG

```
interface Vlan3000
  no shutdown
  mtu 9216
  vrf member tenant-1
  no ip redirects
  ip forward
  ipv6 address use-link-local-only
  no ipv6 redirects
```

Enable IPv6 AF in VRFs

- Add configuration to the VRF definition
- Add configuration to the VRF definition under BGP

VRF

```
vrf context tenant-1
  vni 50000
  rd auto
  address-family ipv4 unicast
    route-target both auto
  route-target both auto evpn
  address-family ipv6 unicast
    route-target both auto
  route-target both auto evpn
```

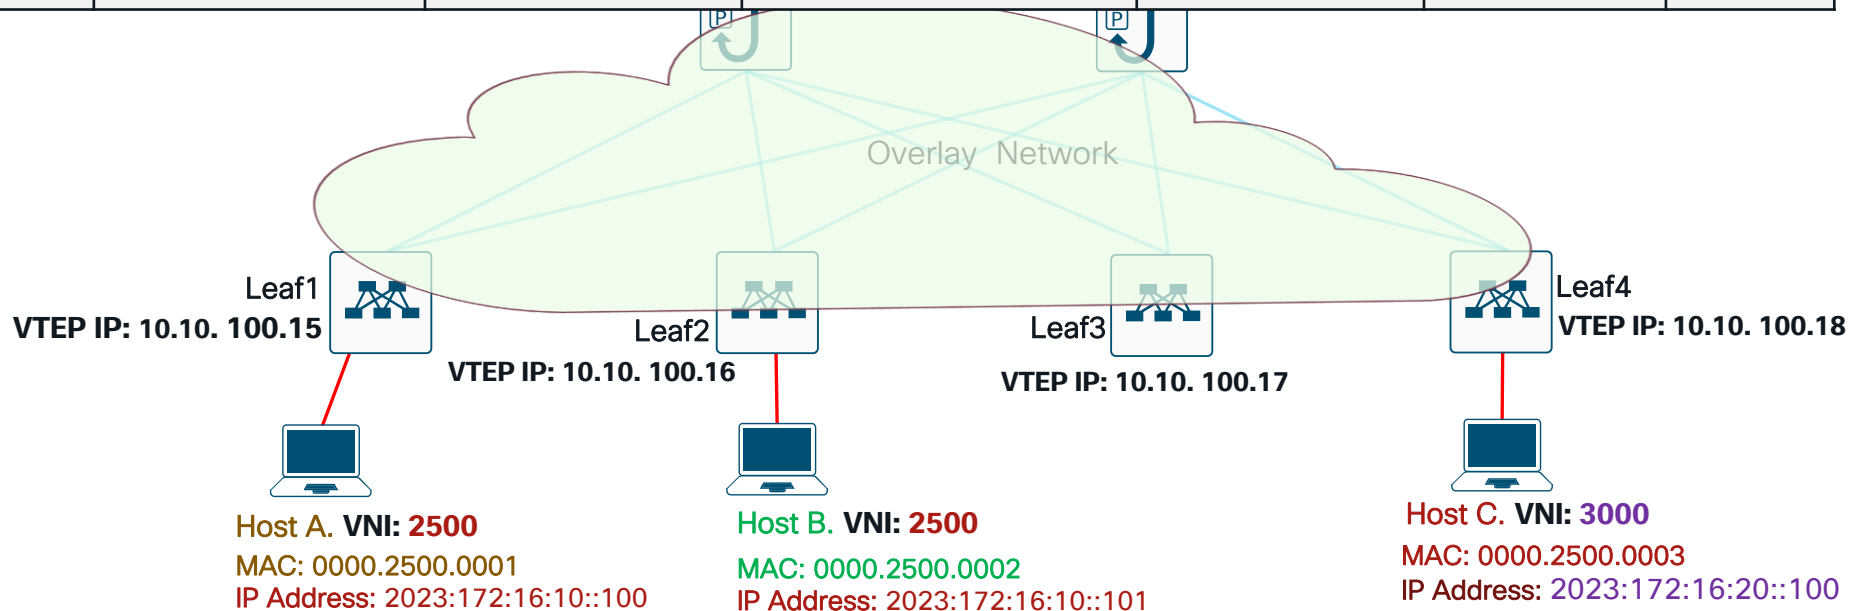
BGP VRF

```
router bgp 65501
  vrf tenant-1
    address-family ipv4 unicast
      advertise l2vpn evpn
      redistribute <..> fabric-rmap-redist-subnet
      maximum-paths ibgp 2
    address-family ipv6 unicast
      advertise l2vpn evpn
      redistribute <..> fabric-rmap-redist-subnet
      maximum-paths ibgp 2
```


VXLAN IPv6 unicast routing in the Overlay with IPv4 Underlay

IPv6 overlay host route advertisement

Type	MAC/Length	L2VNI/RT	IP/Length	L3VNI /RT	Next-Hop	SEQ.
2	0000.2500.0001 /48	2500, 65500:2500	2023:172:16:10::100/128	5000,65500:5000	10.10.100.15	
2	0000.2500.0002 /48	2500, 65500:2500	2023:172:16:10::101/128	5000,65500:5000	10.10.100.16	
2	0000.3000.0003 /48	3000, 65500:3000	2023:172:16:20::100/128	5000,65500:5000	10.10.100.18	



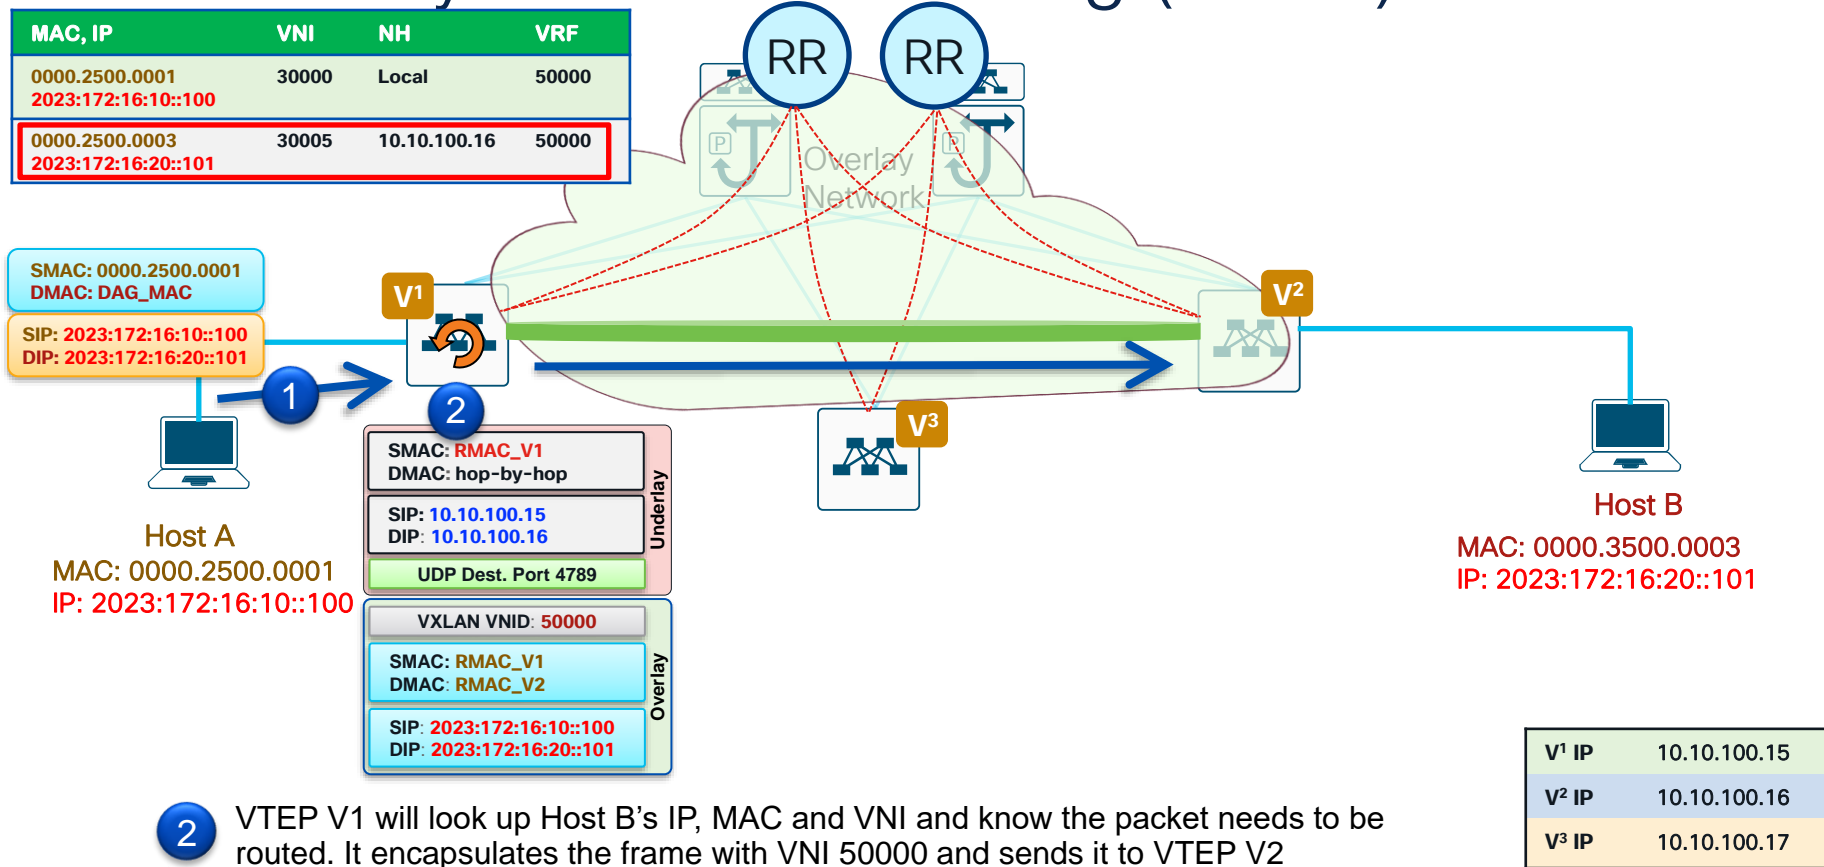
IPv6 overlay Packet Forwarding (Route)



1 Host A needs to send packets to Host B sends it to VTEP V1

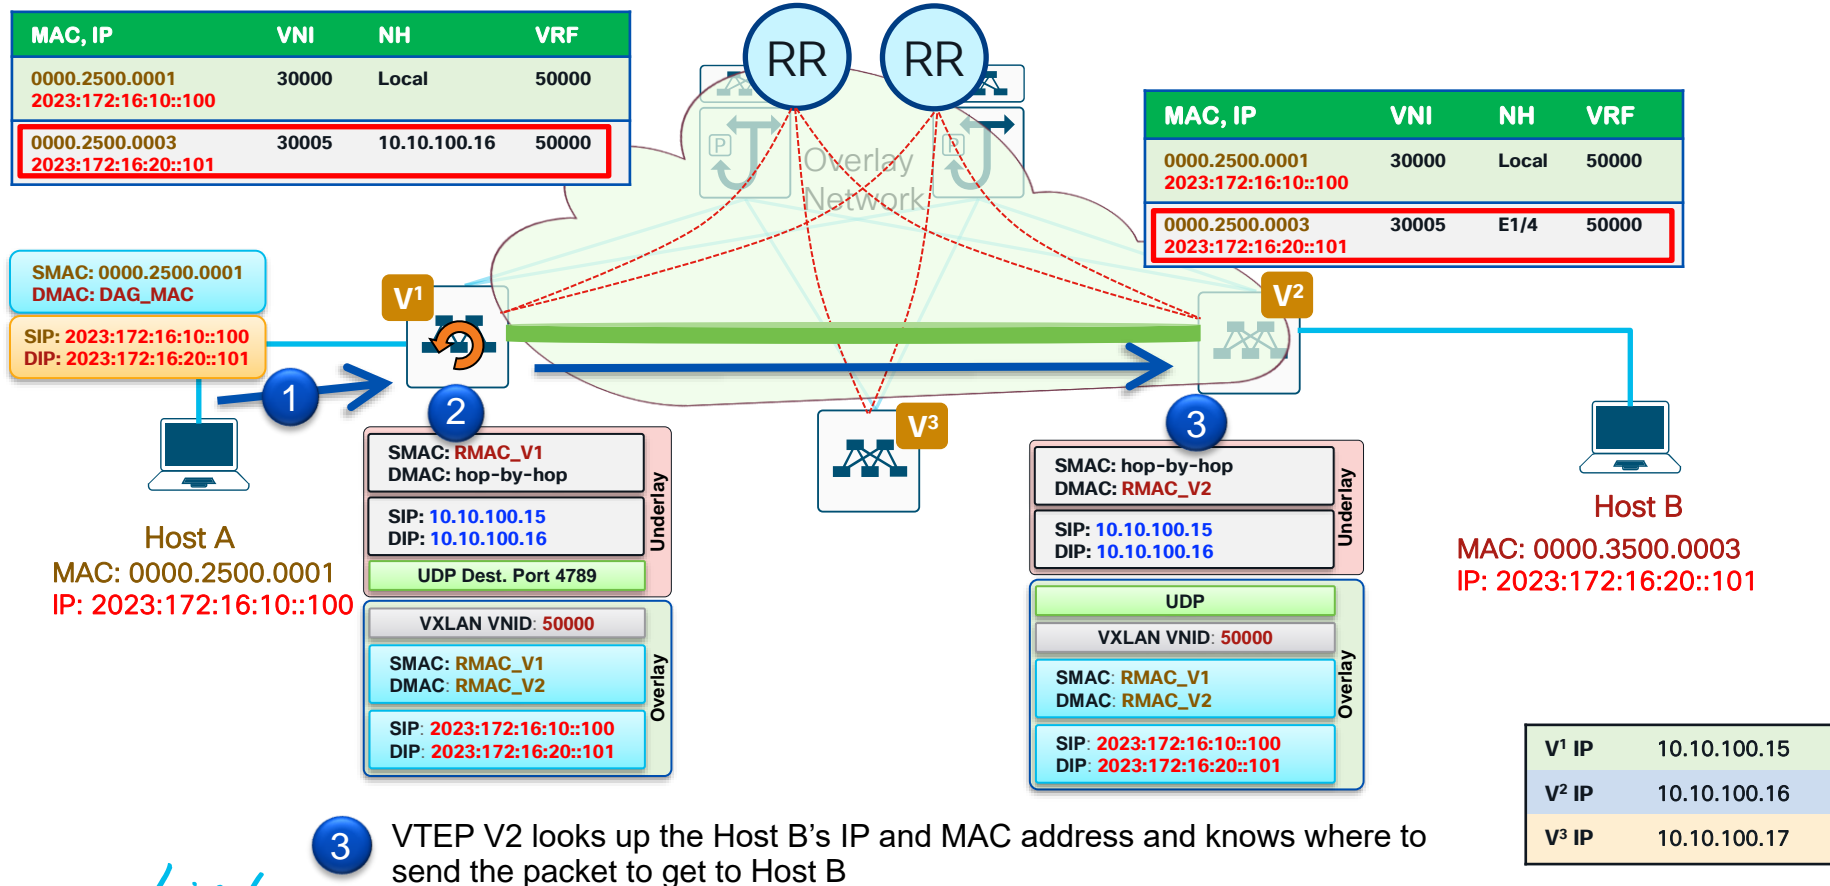
V ¹ IP	10.10.100.15
V ² IP	10.10.100.16
V ³ IP	10.10.100.17

IPv6 overlay Packet Forwarding (Route)

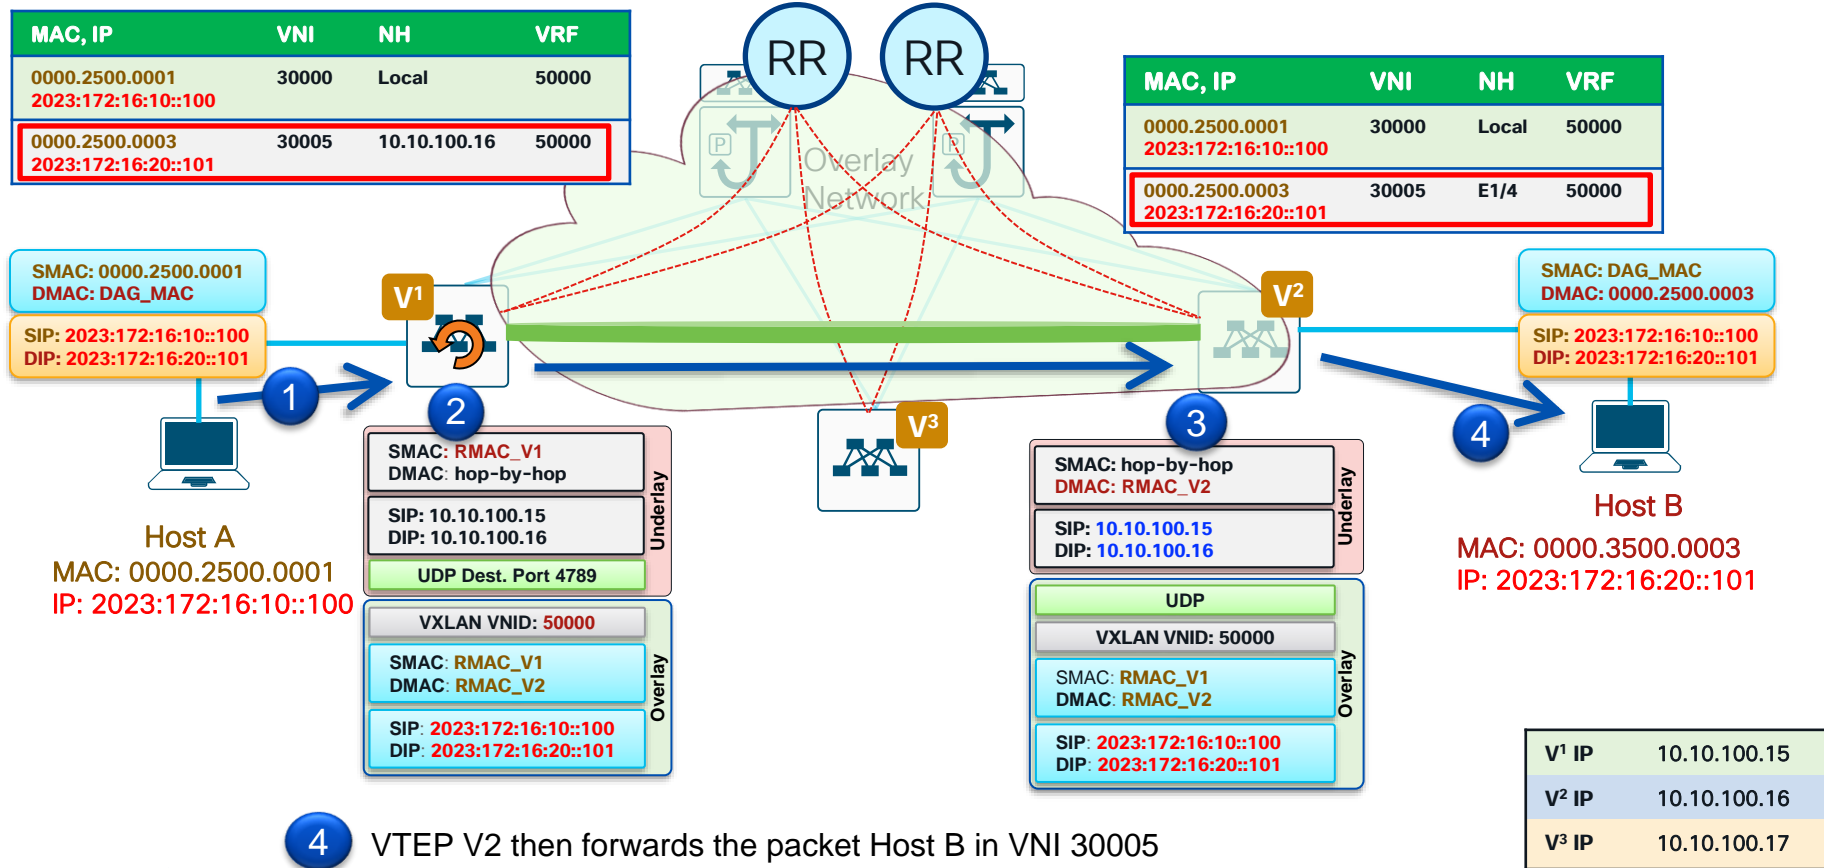


2 VTEP V1 will look up Host B's IP, MAC and VNI and know the packet needs to be routed. It encapsulates the frame with VNI 50000 and sends it to VTEP V2

IPv6 overlay Packet Forwarding (Route)



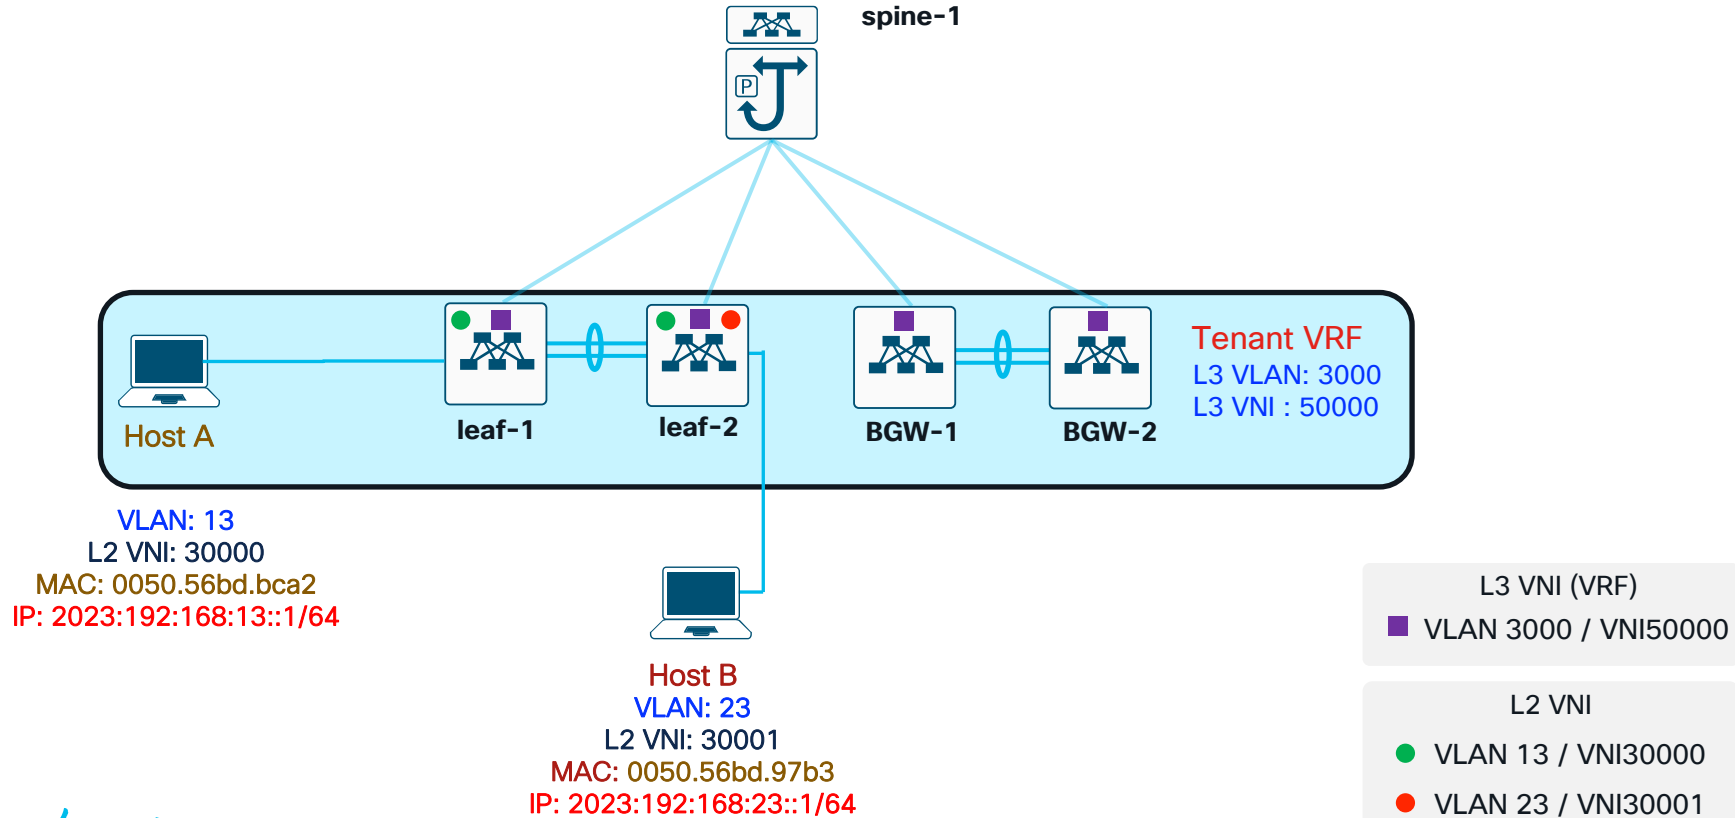
IPv6 overlay Packet Forwarding (Route)



Validating Routing Configs:

Underlay : IPv4
Overlay : IPv6

IPv6 in the Overlay with IPv4 in the underlay



IPv6 Prefix Verification

Type -2 Route Validation

```
Leaf1# show l2route evpn mac-ip all detail
```

Topology	Mac Address	Host IP	Prod	Flags	Seq No	Next-Hops
-----	-----	-----	-----	-----	-----	-----
13	0050.56bd.bca2 L3-Info: 50000 Sent To: BGP	192.168.13.1	HMM	L,	0	Local
13	0050.56bd.bca2 L3-Info: 50000 Sent To: BGP	2023:192:168:13::1	HMM	L,	0	Local

IPv6 Prefix Verification

Type -2 Route Validation

Leaf2# `show l2route evpn mac-ip all detail`

Topology	Mac Address	Host IP	Prod	Flags	Seq No	Next-Hops
13	0050.56bd.bca2 encap-type:1	192.168.13.1	BGP	--	0	10.4.1.1 (Label: 30000)
13	0050.56bd.bca2 encap-type:1	2023:192:168:13::1	BGP	--	0	10.4.1.1 (Label: 30000)
23	0050.56bd.97b3 L3-Info: 50000 Sent To: BGP	192.168.23.1	HMM	L,	0	Local
23	0050.56bd.97b3 L3-Info: 50000 Sent To: BGP	2023:192:168:23::1	HMM	L,	0	Local

leaf-1
VTEP (IPv4)

IPv6 Prefix verification

BGP EVPN Routes – Host A (VNI 30000)

```
Leaf2# show bgp l2vpn evpn vni-id 30000
```

BGP routing table information for VRF default, address family L2VPN EVPN
BGP table version is 39, Local Router ID is 10.4.0.5

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 10.4.0.5:32780 (L2VNI 30000)					
*>i[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[0]:[0.0.0.0]/216					
	10.4.1.1		100	0	i
*>i[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[32]:[192.168.13.1]/272					
	10.4.1.1		100	0	i
*>i[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[128]:[2023:192:168:13::1]/368					
	10.4.1.1		100	0	i

leaf2
RID

32767+
13 (vlan)

Leaf1 VTEP

IPv6 Prefix verification

BGP EVPN Routes – Host B (VNI 30001)

```
Leaf2# show bgp l2vpn evpn vni-id 30001
```

BGP routing table information for VRF default, address family L2VPN EVPN

BGP table version is 39, Local Router ID is 10.4.0.5

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 10.4.0.5:32790 (L2VNI 30001)					
*>1[2]:[0]:[0]:[48]:[0050.56bd.97b3]:[0]:[0.0.0.0]/216					
	10.4.1.4		100	32768	i
*>1[2]:[0]:[0]:[48]:[0050.56bd.97b3]:[32]:[192.168.23.1]/272					
	10.4.1.4		100	32768	i
*>1[2]:[0]:[0]:[48]:[0050.56bd.97b3]:[128]:[2023:192:168:23::1]/368					
	10.4.1.4		100	32768	i

leaf2
RID

32767+
23 (vlan)

Leaf2 VTEP

IPv6 Prefix Verification

VRF Tenant-1 – L3VNI (50000)

Leaf2# **show bgp l2vpn evpn vrf tenant-1**

BGP routing table information for VRF default, address family L2VPN EVPN
BGP table version is 39, Local Router ID is 10.4.0.5

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 10.4.0.5:5 (L3VNI 50000)					
*>i[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[32]:[192.168.13.1]/272					
	10.4.1.1		100	0	i
*>i[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[128]:[2023:192:168:13::1]/368					
	10.4.1.1		100	0	i
* i[5]:[0]:[0]:[24]:[192.168.13.0]/224					
	10.4.1.1	0	100	0	?
*>1					
	10.4.1.4	0	100	32768	?
*>1[5]:[0]:[0]:[24]:[192.168.23.0]/224					
	10.4.1.4	0	100	32768	?
*>1[5]:[0]:[0]:[64]:[2023:192:168:13::]/416					
	10.4.1.4	0	100	32768	?
* i					
	10.4.1.1	0	100	0	?
*>1[5]:[0]:[0]:[64]:[2023:192:168:23::]/416					
	10.4.1.4	0	100	32768	?

IPv6 Prefix Verification

Host Route Details – Host A (VNI 30000)

```
Leaf1# show bgp l2vpn evpn 2023:192:168:13::1
BGP routing table information for VRF default, address family L2VPN EVPN
Route Distinguisher: 10.4.0.1:32780 (L2VNI 30000)
BGP routing table entry for [2]:[0]:[0]:[48]:[0050.56bd.bca2]:[128]:[2023:192:168:13::1]/368,
version 30
Paths: (1 available, best #1)
Flags: (0x000102) (high32 00000000) on xmit-list, is not in l2rib/evpn
```

Advertised path-id 1

Path type: local, path is valid, is best path, no labeled nexthop

AS-Path: NONE, path locally originated

10.4.1.1 (metric 0) from 0.0.0.0 (10.4.0.1)

Origin IGP, MED not set, localpref 100, weight 32768

Received label 30000 50000

Extcommunity: RT:65501:30000 RT:65501:50000 ENCAP:8 Router MAC:ac7a.56fc.ffd7

Path-id 1 advertised to peers:

10.4.0.3

Who gets this
information

IPv6 Prefix Verification -1

Host Route Details - Host B (VNI 30001)

```
Leaf1# show bgp l2vpn evpn 2023:192:168:23::1
```

```
BGP routing table information for VRF default, address family L2VPN EVPN
```

```
Route Distinguisher: 10.4.0.5:32790
```

```
BGP routing table entry for [2]:[0]:[0]:[48]:[0050.56bd.97b3]:[128]:[2023:192:168:23::1]/368,  
version 33
```

```
Paths: (1 available, best #1)
```

```
Flags: (0x000202) (high32 00000000) on xmit-list, is not in l2rib/evpn, is not in HW
```

```
Advertised path-id 1
```

```
Path type: internal, path is valid, is best path, no labeled nexthop
```

```
Imported to 2 destination(s)
```

```
Imported paths list: tenant-1 L3-50000
```

```
AS-Path: NONE, path sourced internal to AS
```

```
10.4.1.4 (metric 9) from 10.4.0.3 (10.4.0.3)
```

```
Origin IGP, MED not set, localpref 100, weight 0
```

```
Received label 30001 50000
```

```
Extcommunity: RT:65501:30001 RT:65501:50000 ENCAP:8 Router MAC:ac7a.56fd.24f7
```

```
Originator: 10.4.0.5 Cluster list: 10.4.0.3
```

Who gets this information

IPv6 Prefix Verification -2

Host Route Details – Host B (VNI 30001)

Route Distinguisher: 10.4.0.1:5 (L3VNI 50000)

BGP routing table entry for [2]:[0]:[0]:[48]:[0050.56bd.97b3]:[128]:[2023:192:168:23::1]/368, version 34

Paths: (1 available, best #1)

Flags: (0x000202) (high32 00000000) on xmit-list, is not in l2rib/evpn, is not in HW

Advertised path-id 1

Path type: internal, path is valid, is best path, no labeled nexthop

Imported from

10.4.0.5:32790:[2]:[0]:[0]:[48]:[0050.56bd.97b3]:[128]:[2023:192:168:23::1]/368

AS-Path: NONE, path sourced internal to AS

10.4.1.4 (metric 9) from 10.4.0.3 (10.4.0.3)

Origin IGP, MED not set, localpref 100, weight 0

Received label 30001 50000

Extcommunity: RT:65501:30001 RT:65501:50000 ENCAP:8 Router MAC:ac7a.56fd.24f7

Originator: 10.4.0.5 Cluster list: 10.4.0.3

Who gets this information

IPv6 Prefix Verification

MAC Address Table

```
Leaf1# show mac address-table
```

Legend:

* - primary entry, G - Gateway MAC, (R) - Routed MAC, O - Overlay MAC
age - seconds since last seen, + - primary entry using vPC Peer-Link,
(T) - True, (F) - False, C - ControlPlane MAC, ~ - vsan,
(NA) - Not Applicable

	VLAN	MAC Address	Type	age	Secure	NTFY	Ports
*	13	0050.56bd.bca2	dynamic	NA	F	F	Eth1/2
G	-	2020.0000.00aa	static	-	F	F	sup-eth1 (R)
G	-	ac7a.56fc.ffd7	static	-	F	F	sup-eth1 (R)
G	13	ac7a.56fc.ffd7	static	-	F	F	sup-eth1 (R)
G	3000	ac7a.56fc.ffd7	static	-	F	F	sup-eth1 (R)

IPv6 Prefix Verification

MAC Address Table

Leaf2# show mac address-table

egend:

* - primary entry, G - Gateway MAC, (R) - Routed MAC, O - Overlay MAC
age - seconds since last seen, + - primary entry using vPC Peer-Link,
(T) - True, (F) - False, C - **ControlPlane** MAC, ~ - vsan,
(NA) - Not Applicable

VLAN	MAC Address	Type	age	Secure	NTFY	Ports
-----+-----+-----+-----+-----+-----						
C 13	0050.56bd.bca2	dynamic	NA	F	F	nve1(10.4.1.1)
* 23	0050.56bd.97b3	dynamic	NA	F	F	Eth1/2
G -	2020.0000.00aa	static	-	F	F	sup-eth1 (R)
G -	ac7a.56fd.24f7	static	-	F	F	sup-eth1 (R)
G 13	ac7a.56fd.24f7	static	-	F	F	sup-eth1 (R)
G 23	ac7a.56fd.24f7	static	-	F	F	sup-eth1 (R)
G 3000	ac7a.56fd.24f7	static	-	F	F	sup-eth1 (R)

IPv6 Prefix Verification

IPv6 Neighbor Discovery (~ARP)

Leaf1# `show ipv6 neighbor vrf tenant-1` ← (`show ip arp` for `ipv6`)

IPv6 Adjacency Table for VRF tenant-1

Total number of entries: 2

Address	Age	MAC Address	Pref	Source	Interface	Mobility Flags
2023:192:168:13::1	02:54:41	0050.56bd.bca2	50	icmpv6	Vlan13	
fe80::b6e0:1e7e:40f2:7b7c	02:54:35	0050.56bd.bca2	50	icmpv6	Vlan13	

Link Local IPv6

Leaf2# `show ipv6 neighbor vrf tenant-1`

IPv6 Adjacency Table for VRF tenant-1

Total number of entries: 2

Address	Age	MAC Address	Pref	Source	Interface	Mobility Flags
2023:192:168:23::1	02:51:37	0050.56bd.97b3	50	icmpv6	Vlan23	
fe80::b6e0:1e7e:40f2:7b7c	02:51:32	0050.56bd.97b3	50	icmpv6	Vlan23	

Link Local IPv6

Considerations for deploying IPv6 Underlay



VXLAN IPv6 Guidelines

- Dual Stack (IPv4 and IPv6) is not supported for VXLAN underlay. It should either be IPv4 or IPv6, not both.
- NVE Source interface loopback for VTEP can either be IPv4 (VXLANv4) or IPv6 (VXLANv6), and not both
- For VXLAN with IPv6 in the Underlay supports IS-IS, OSPFv3 and eBGP
- Multicast in the underlay is not supported yet, instead Ingress replication is supported
- Next hop address in overlay (in bgp l2vpn evpn address family updates) should be resolved in underlay to the same address family. E.g.:

VTEP Address Family	BGP l2vpn evpn peering
IPv4 Address	Only over IPv4 Addresses
IPv6 Address	Only over IPv6 Addresses

vPC requirements for IPv6 Underlay

2 Loopbacks

- VXLAN with IPv4 underlay leveraged the concept of secondary IP address for the purpose of VIP. However, there is no concept of secondary IP address in IPv6.
- Primary IP address (PIP) and VIP need to be two separate loopback interfaces for IPv6 underlay.
- NVE CLI in IPv6 underlay to specify the loopback interface that carries the VIP in case of VXLANv6 vPC

```
interface loopback1 ← PIP
  ipv6 address 2001:db8:0:0:1:0:0:1/128
interface loopback10 ← VIP
  ipv6 address 2001:db8:0:0:1:0:0:2/128
interface nve1
  source-interface loopback1 anycast loopback10
```

- VTEP IP address cannot be a link local IPv6 address.
- Loopback1 and loopback10 should be valid /128 IPv6 **Global Unicast Address**

vPC requirements for IPv6 Underlay

vMAC

- vPC VTEPs use vMAC (virtual MAC) with the VIP/PIP feature
- vMAC is used with VIP and the system MAC is used with PIP
- By default, in a IPv6 underlay the vMAC is autogenerated by picking the last 48 bits from the IPv6 VIP
 - Autogenerated vMAC = 0x06 + the last 4 bytes of the IPv6 VIP address
- vMAC can also be manually configured and will be given precedence over autogenerated

```
interface nve1
    virtual-rmac <Unique 48-bit mac address>
```

- vMAC needs to be unique in the fabric

vPC requirements for IPv6 Underlay

Peer Keepalive

- vPC needs to allow IPv6 addresses to be used peer-keepalive links
- peer-keepalive must be a global unicast address
- The keepalive link becomes operational when when the peers have correctly configured IPv6 address, and they are reachable from each peer
- Peer-keepalive can be configured on the in-band or out-of-band interfaces

```
vpc domain 1
  peer-keepalive destination 2023:172:16:10::100 source 2023:172:16:10::101
vrf default
```

BGP requirements for IPv6 Underlay

BGP Underlay

- BGP IPv6 neighbor must support L2VPN EVPN address-family session
- BGP requires a 32-bit router id to be configured to identify the instance. this can any 32-bit integer. It is not used for routing
- With IPv6 in the undelay for VXLAN fabric, none of the interfaces have IPv4 addresses, the router ID will need to be manually configured to a 32-bit integer value

```
Router bgp 65001
  ! IPV4 Router ID
  router-id 2.2.2.2
  address-family ipv6 unicast
    redistribute direct route-map allow
```

Implementing IPv6 in the Underlay

IPv6 Underlay Configuration Steps

- Enable OSPFv3 /IS-IS/ feature for underlay routing
- Configure **OSPFv3**/ IS-IS in the underlay on all loopback and directly connected switch interfaces
- Configure IPv6 gateway address on the SVI and host.
- Configure IPv6 forwarding on the L3VNI VLANs
- Configure IPv6 for VRF and for VRF under BGP

IPv6 Underlay Configuration Steps

- Configure two separate loopback on each of the vPC switches; **loopback1 (PIP)** and **Loopback10 (VIP)**
- Ingress replication used instead of underlay Multicast under NVE interface config.
- Configure peer-keepalive configuration to support IPv6 addresses in the vPC domain (use loopback)

IPv6 Underlay Configuration - 1

- Enable **OSPFv3** / IS-IS feature for underlay routing
- Configure **OSPFv3** / IS-IS and IPv6 routing in the underlay on all loopback and fabric switch interfaces

FEATURE OSPFV3

```
feature OSPFv3
ipv6 switch-packets lla
router ospfv3 UNDERLAY
  router-id 10.4.1.1
interface Ethernet1/1
ipv6 link-local use-bia
  ipv6 router ospfv3 UNDERLAY area 0.0.0.0
no shutdown
```

LOOPBACK

```
interface loopback0
  description Routing loopback interface
  ipv6 address fd00::a02:0/128
  ipv6 router ospfv3 UNDERLAY area 0.0.0.0
interface loopback1
  description VTEP loopback interface
  ipv6 address fd00::a03:0/128
  ipv6 router ospfv3 UNDERLAY area 0.0.0.0
```

IPv6 Underlay Configuration - 1

- Assign IPv6 address on L2VNI VLAN SVI and enable IPv6 forwarding on SVIs for L2/L3 VNI VLANs
- The IPv6 address “**use-link-local-only**” serves the same purpose as “**ip forward**” for IPv4

L2VNI

```
interface Vlan13
  no shutdown
  vrf member tenant-1
  ip address 192.168.13.254/24 tag 12345
  ipv6 address 2023:192:168:13::254/64 tag 12345
  fabric forwarding mode anycast-gateway
```

L3VNI

```
interface Vlan3000
  no shutdown
  mtu 9216
  vrf member tenant-1
  no ip redirects
  ip forward
  ipv6 address use-link-local-only
  no ipv6 redirects
```

IPv6 Underlay Configuration -2

- Configure IPv6 for VRF and for VRF under BGP

VRF CONFIG

```
vrf context tenant-1
  vni 50000
  rd auto
  address-family ipv4 unicast
    route-target both auto
    route-target both auto evpn
  address-family ipv6 unicast
    route-target both auto
    route-target both auto evpn
```

VRF CONFIG UNDER BGP

```
router bgp 65503
  vrf tenant-1
    address-family ipv4 unicast
      advertise l2vpn evpn
      redistribute <..> fabric-rmap-redist-subnet
      maximum-paths ibgp 2
    address-family ipv6 unicast
      advertise l2vpn evpn
      redistribute <..> fabric-rmap-redist-subnet
      maximum-paths ibgp 2
```

IPv6 Underlay Configuration -3

- Configure two separate loopback on each of the vPC switches; **loopback1 (PIP)** and **Loopback10 (VIP)**. **loopback10** shares the same IP on both vPC pair
- Configure Ingress replication for the L2 VNI that are on multiple leaf(s)
- Configure the **loopback10** interfaces under NVE interface

VP VIP LOOPBACK

```
interface loopback10
  description Anycast loopback interface
  ipv6 address fd00::a03:1/128
  ipv6 router ospfv3 UNDERLAY area 0.0.0.0
```

NVE INTERFACE & LO10 IN L2 VNI

```
interface nve1
  no shutdown
  host-reachability protocol bgp
  source-interface loopback1 anycast loopback10
  member vni 30000
  ingress-replication protocol bgp
  member vni 50000 associate-vrf
```

IPv6 Underlay Configuration -4

- Configure peer-keepalive configuration to support IPv6 addresses in the vPC domain
- Synchronize neighbor discovery (nd) between the vPC peers

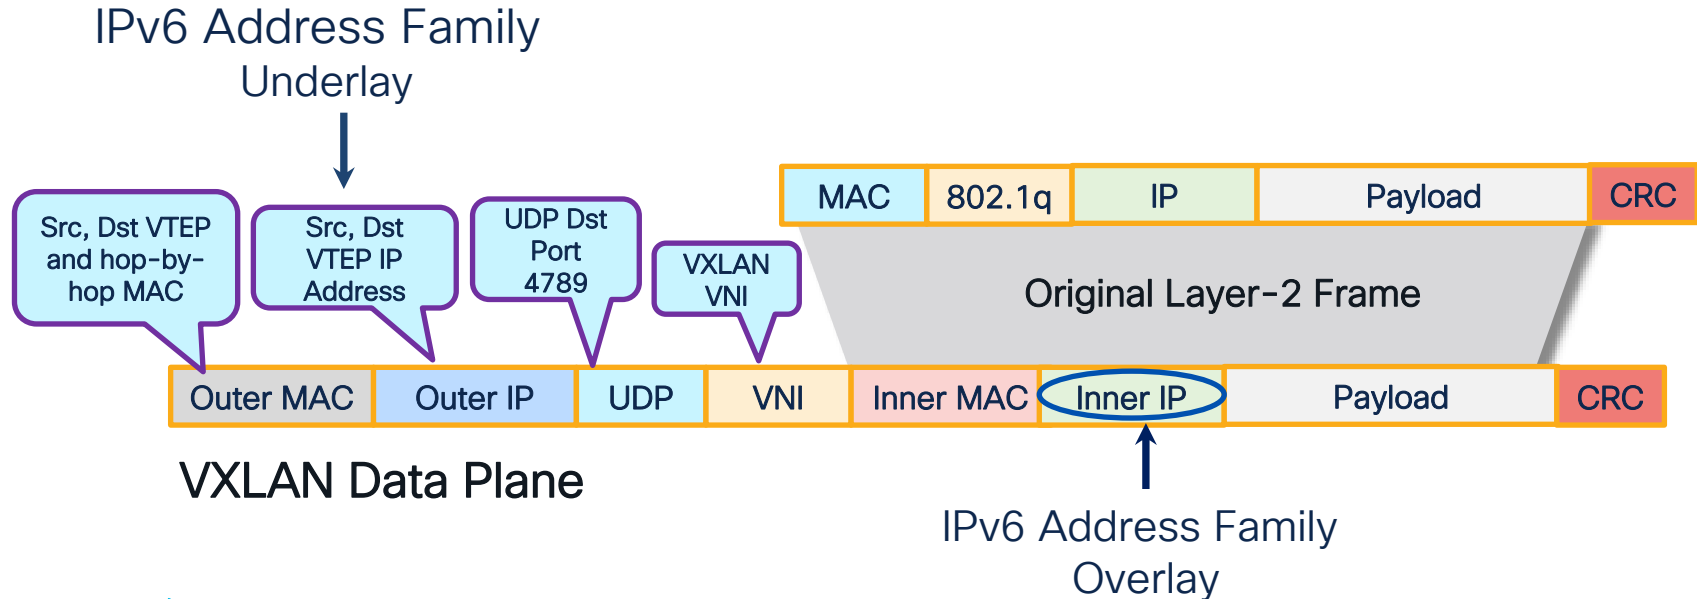
VPC DOMAIN

```
vpc domain 1
  peer-switch
  peer-keepalive destination fd00::a02:0 source fd00::a02:2 vrf default
  delay restore 150
  peer-gateway
  auto-recovery reload-delay 360
  ipv6 nd synchronize
  ip arp synchronize
```

VXLAN IPv6 unicast routing in the Overlay with IPv6 Underlay

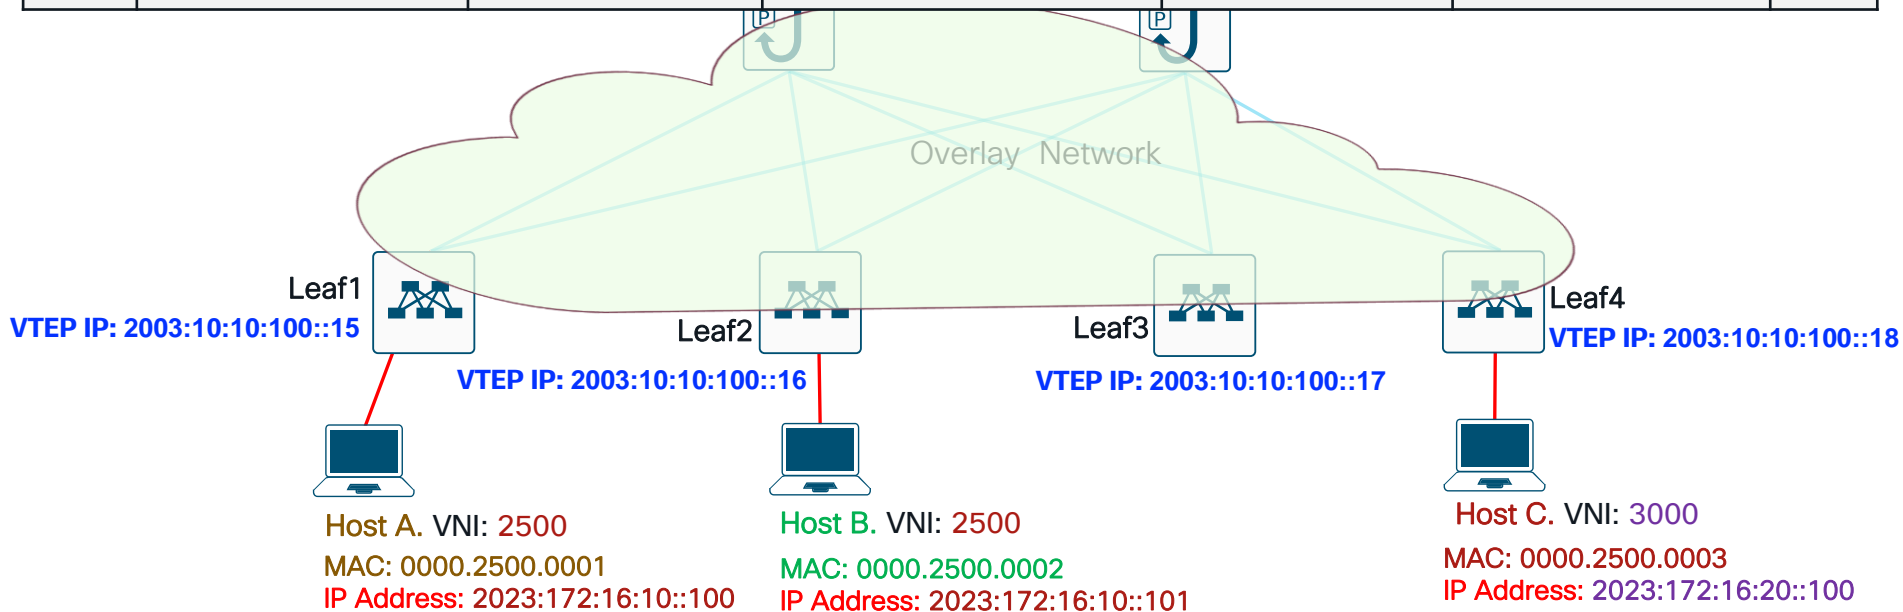
IPv6 VXLAN deployment the Data Center

The IPv6 packet is encapsulated in the overlay as IPv6 packets and uses IPv6 routing to transport the VXLAN encapsulated traffic.




IPv6 overlay host route advertisement

Type	MAC/Length	L2VNI/RT	IP/Length	L3VNI /RT	Netx-Hop	SEQ.
2	0000.2500.0001 /48	2500, 65500:2500	2023:172:16:10::100/128	5000,65500:5000	2003:10:10:100::15	
2	0000.2500.0002 /48	2500, 65500:2500	2023:172:16:10::101/128	5000,65500:5000	2003:10:10:100::16	
2	0000.3000.0003 /48	3000, 65500:3000	2023:172:16:20::100/128	5000,65500:5000	2003:10:10:100::18	



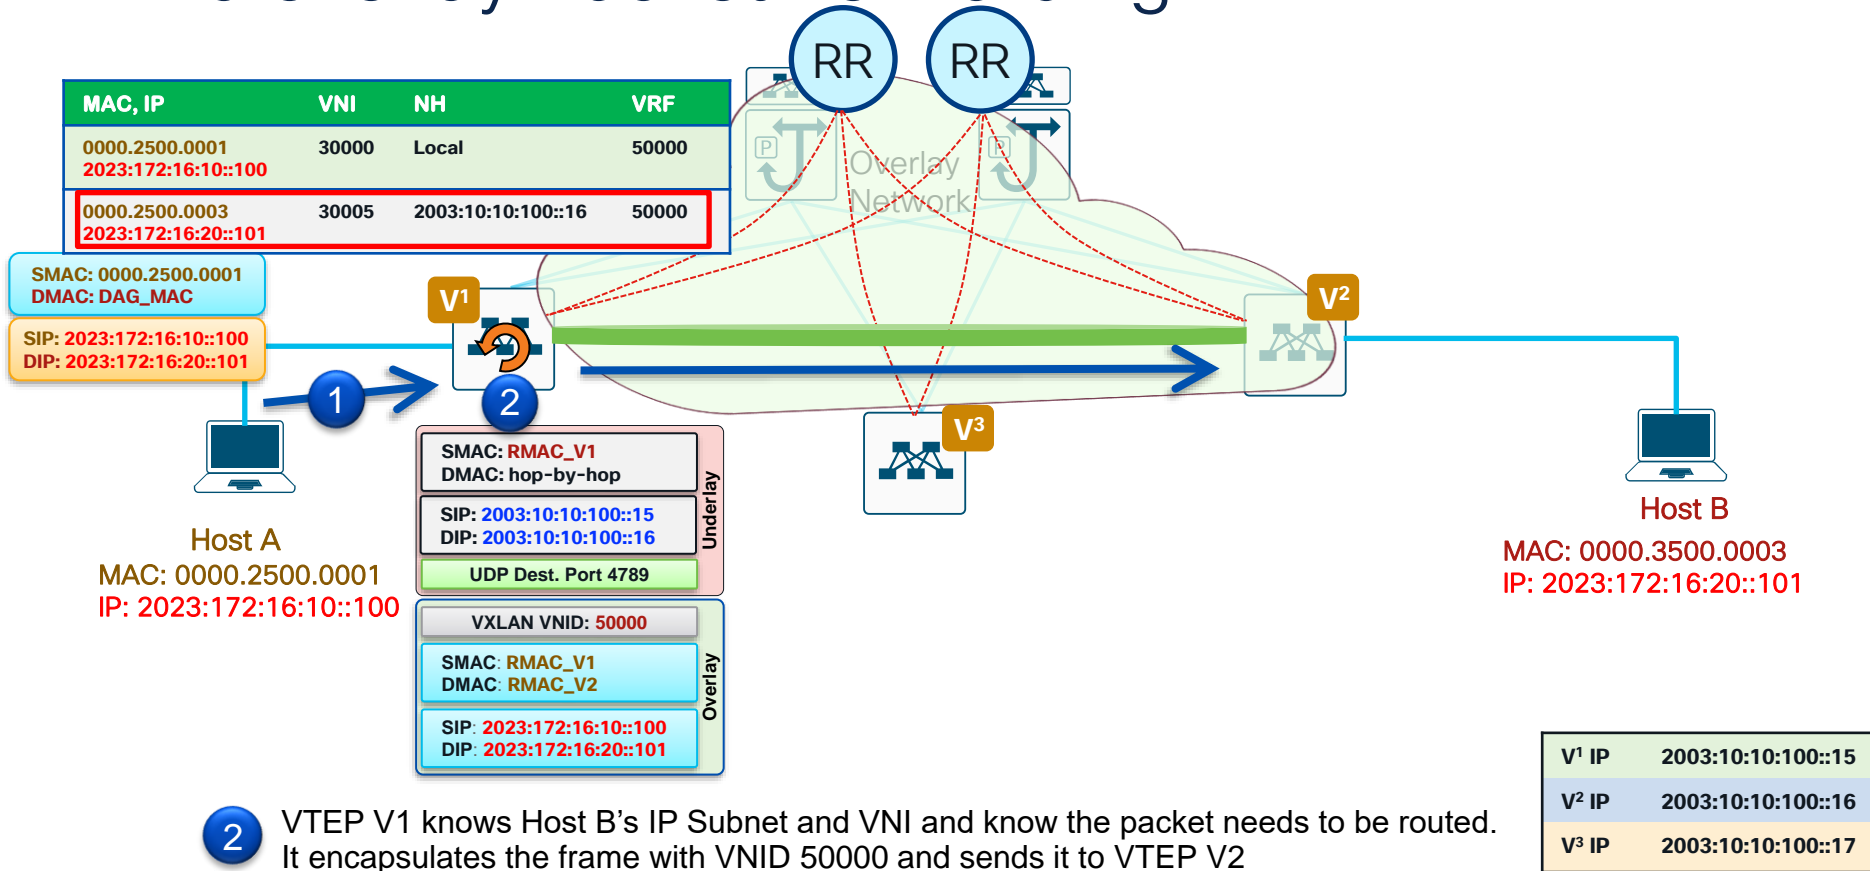
IPv6 overlay Packet Forwarding (Route)



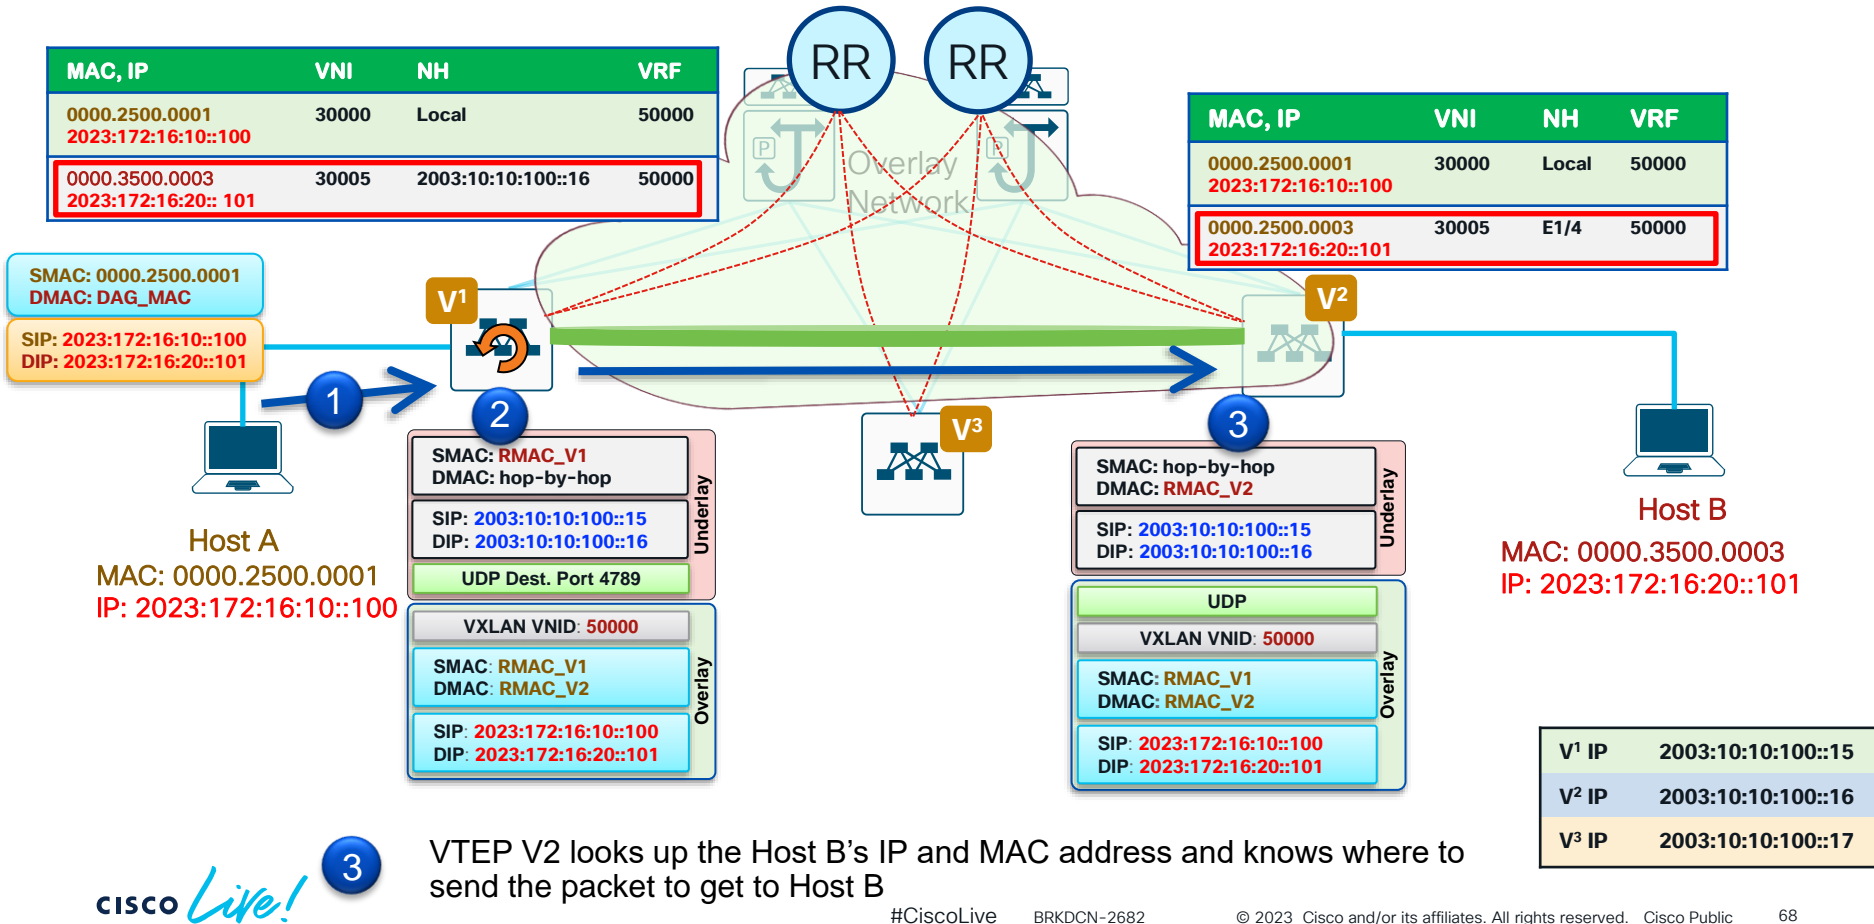
1 Host A needs to send packets to Host B sends it to VTEP V1

V ¹ IP	2003:10:10:100::15
V ² IP	2003:10:10:100::16
V ³ IP	2003:10:10:100::17

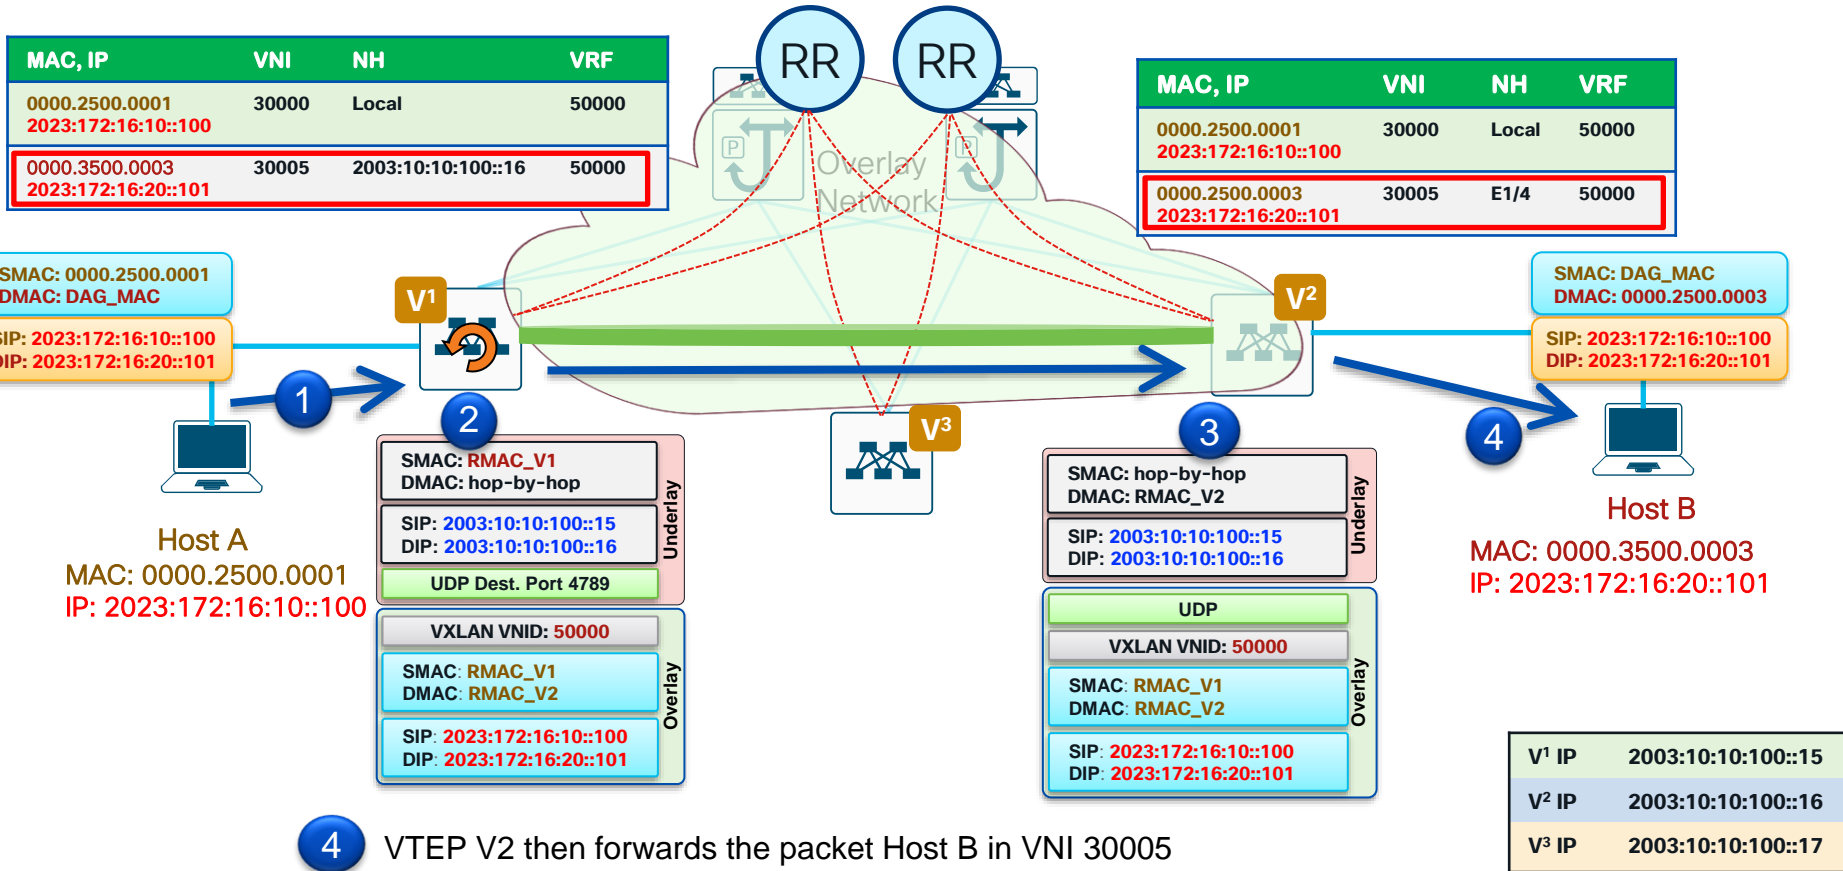
IPv6 overlay Packet Forwarding



IPv6 overlay Packet Forwarding (Route)



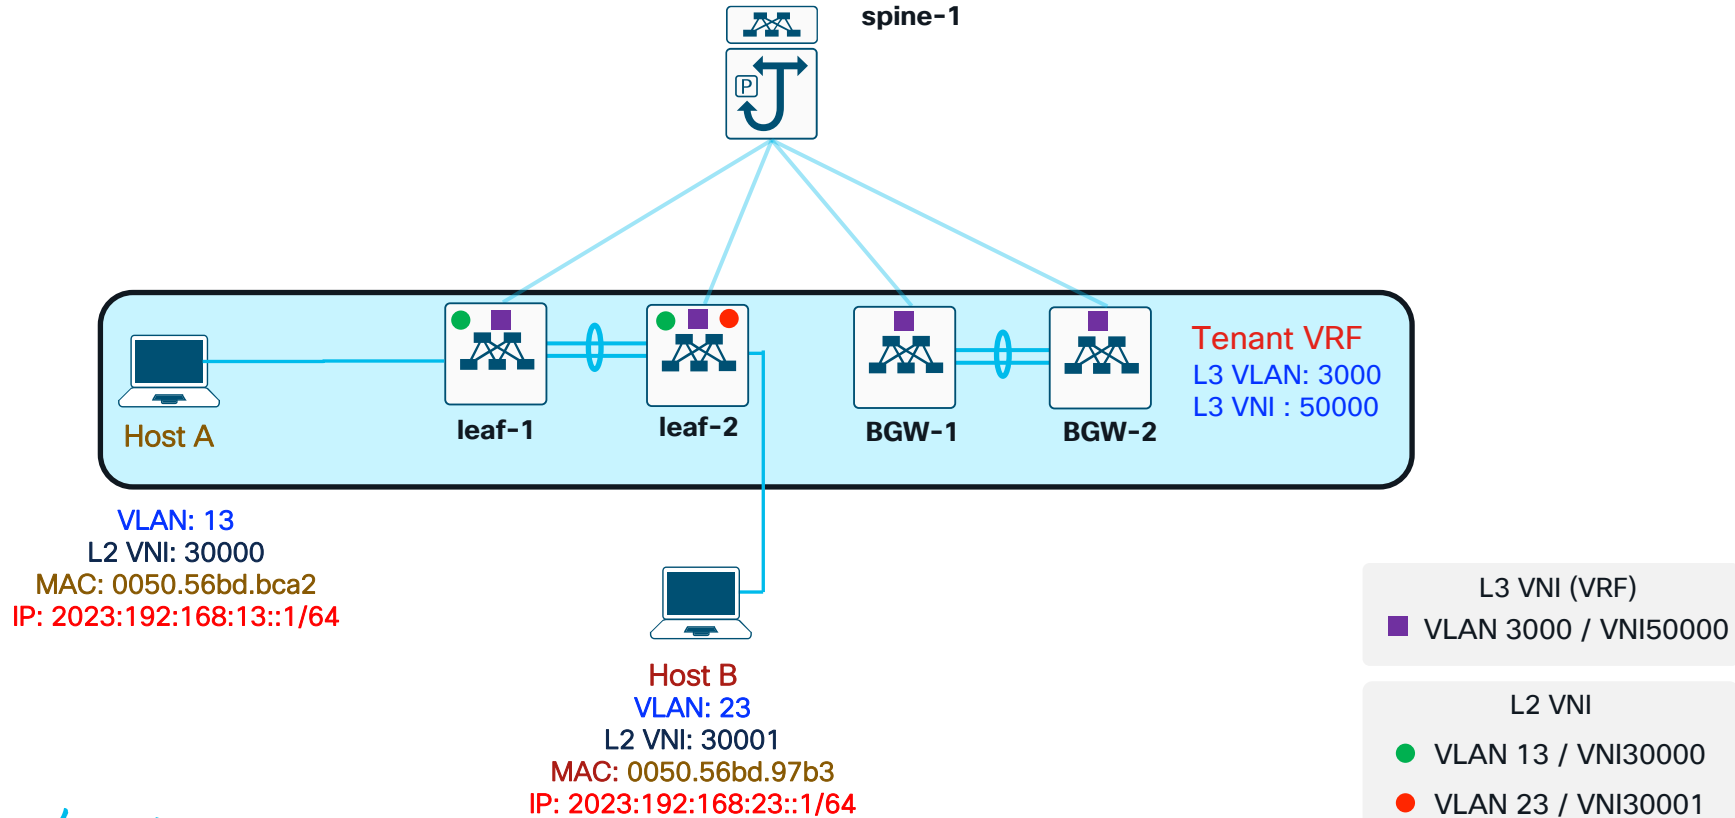
IPv6 overlay Packet Forwarding (Route)



Validating IPv6 routing in the Overlay



IPv6 in the overlay with IPV6 in the underlay



IPv6 Prefix Verification

Type -2 Route Validation

```
Leaf1# show l2route evpn mac-ip all detail
```

Topology	Mac Address	Host IP	Prod	Flags	Seq No	Next-Hops
-----	-----	-----	-----	-----	-----	-----
13	0050.56bd.bca2 L3-Info: 50000 Sent To: BGP	192.168.13.1	HMM	L,	0	Local
13	0050.56bd.bca2 L3-Info: 50000 Sent To: BGP	2023:192:168:13::1	HMM	L,	0	Local

IPv6 Prefix Verification

Type -2 Route Validation

Leaf2# show l2route evpn mac-ip all detail

Topology	Mac Address	Host IP	Prod	Flags	Seq No	Next-Hops
13	0050.56bd.bca2 encap-type:1	192.168.13.1	BGP	--	0	fd00::a03:0 (Label: 30000)
13	0050.56bd.bca2 encap-type:1	2023:192:168:13::1	BGP	--	0	fd00::a03:0 (Label: 30000)
23	0050.56bd.97b3 L3-Info: 50000 Sent To: BGP	192.168.23.1	HMM	L,	0	Local
23	0050.56bd.97b3 L3-Info: 50000 Sent To: BGP	2023:192:168:23::1	HMM	L,	0	Local

Leaf-1
VTEP (IPv6)

IPv6 Prefix Verification

Ingress Replication - Route Type 3 on Leaf1

```
Leaf1# show bgp l2vpn evpn vni-id 30000
```

BGP routing table information for VRF default, address family L2VPN EVPN

BGP table version is 20, Local Router ID is 10.4.1.1

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher:	10.4.1.1:32780	(L2VNI 30000)			
*>1[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[0]:[0.0.0.0]/216	fd00::a03:0		100	32768	i
*>1[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[128]:[2023:192:168:13::1]/368					
IR Route Type	fd00::a03:0		100	32768	i
*>1[3]:[0]:[128]:[fd00::a03:0]/184	fd00::a03:0		100	32768	i
*>i[3]:[0]:[128]:[fd00::a03:3]/184	fd00::a03:3		100	0	i

Leaf-1
VTEP

Leaf-2
VTEP

IPv6 Prefix Verification

Ingress Replication - Route Type 3 on Leaf2

```
Leaf2# show bgp l2vpn evpn vni-id 30000
```

```
BGP routing table information for VRF default, address family L2VPN EVPN
```

```
BGP table version is 22, Local Router ID is 10.4.1.5
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher:	10.4.1.5:32780	(L2VNI 30000)			

```
*>i[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[0]:[0.0.0.0]/216
```

fd00::a03:0	100	0	i
-------------	-----	---	---

```
*>i[2]:[0]:[0]:[48]:[0050.56bd.bca2]:[128]:[2023:192:168:13::1]/368
```

IR Route Type

fd00::a03:0	100	0	i
-------------	-----	---	---

```
*>i[3]:[0]:[128]:[fd00::a03:0]/184
```

fd00::a03:0	100	0	i
-------------	-----	---	---

```
*>l[3]:[0]:[128]:[fd00::a03:3]/184
```

fd00::a03:3	100	32768	i
-------------	-----	-------	---

Leaf-1
VTEP

Leaf-2
VTEP

IPv6 Prefix Verification

VRF Tenant-1 – L3VNI (50000)

```
Leaf1# show bgp l2vpn evpn vrf tenant-1
```

BGP routing table information for VRF default, address family L2VPN EVPN
BGP table version is 20, Local Router ID is 10.4.1.1

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 10.4.1.1:7 (L3VNI 50000)					
*>i[2]:[0]:[0]:[48]:[0050.56bd.97b3]:[128]:[2023:192:168:23::1]/368	fd00::a03:3		100	0	i
* i[5]:[0]:[0]:[24]:[192.168.13.0]/224	fd00::a03:3	0	100	0	?
*>1	fd00::a03:0	0	100	32768	?
*>i[5]:[0]:[0]:[24]:[192.168.23.0]/224	fd00::a03:3	0	100	0	?
* i[5]:[0]:[0]:[64]:[2023:192:168:13::]/416	fd00::a03:3	0	100	0	?
*>1	fd00::a03:0	0	100	32768	?
*>i[5]:[0]:[0]:[64]:[2023:192:168:23::]/416	fd00::a03:3	0	100	0	?

Diagram annotations:

- Leaf-2 VTEP (green oval) points to the first instance of fd00::a03:3.
- Leaf-1 VTEP (blue oval) points to the second instance of fd00::a03:3.

IPv6 Prefix Verification – 1

Ingress Replication – Route Type 3 Detailed

```
Leaf1# show bgp l2vpn evpn route-type 3
```

```
BGP routing table information for VRF default, address family L2VPN EVPN
```

```
Route Distinguisher: 10.4.1.1:32780 (L2VNI 30000)
```

```
BGP routing table entry for [3]:[0]:[128]:[fd00::a03:0]/184, version 14
```

```
Paths: (1 available, best #1)
```

```
Flags: (0x000002) (high32 00000000) on xmit-list, is not in l2rib/evpn
```

```
Advertised path-id 1
```

```
Path type: local, path is valid, is best path, no labeled nexthop, is extd
```

```
AS-Path: NONE, path locally originated
```

```
fd00::a03:0 (metric 0) from 0:: (10.4.1.1)
```

```
Origin IGP, MED not set, localpref 100, weight 32768
```

```
Extcommunity: RT:65503:30000 ENCAP:8
```

```
PMSI Tunnel Attribute:
```

```
flags: 0x00, Tunnel type: Ingress Replication
```

```
Label: 30000, Tunnel Id: fd00::a03:0
```

Who gets this information

IPv6 Prefix Verification – 2

Ingress Replication – Route Type 3 Detailed

BGP routing table entry for [3]:[0]:[128]:[fd00::a03:3]/184, version 16
Paths: (1 available, best #1)
Flags: (0x000012) (high32 00000000) on xmit-list, is in l2rib/evpn, is not in HW

Advertised path-id 1

Path type: internal, path is valid, is best path, no labeled nexthop

Imported from 10.4.1.5:32780:[3]:[0]:[128]:[fd00::a03:3]/184

AS-Path: NONE, path sourced internal to AS

fd00::a03:3 (metric 8) from fd00::a02:2 (10.4.1.3)

Origin IGP, MED not set, localpref 100, weight 0

Extcommunity: RT:65503:30000 ENCAP:8

Who gets this information

Originator: 10.4.1.5 Cluster list: 10.4.1.3

PMSI Tunnel Attribute:

flags: 0x00, Tunnel type: Ingress Replication

Label: 30000, Tunnel Id: fd00::a03:3

IPv6 Prefix Verification

Host Route Details – Host A (VNI 30000)

```
Leaf1# show bgp l2vpn evpn 2023:192:168:13::1
```

```
BGP routing table information for VRF default, address family L2VPN EVPN
```

```
Route Distinguisher: 10.4.1.1:32780 (L2VNI 30000)
```

```
BGP routing table entry for [2]:[0]:[0]:[48]:[0050.56bd.bca2]:[128]:[2023:192:168:13::1]/368,  
version 19
```

```
Paths: (1 available, best #1)
```

```
Flags: (0x000102) (high32 00000000) on xmit-list, is not in l2rib/evpn
```

```
Advertised path-id 1
```

```
Path type: local, path is valid, is best path, no labeled nexthop
```

```
AS-Path: NONE, path locally originated
```

```
fd00::a03:0 (metric 0) from 0:: (10.4.1.1)
```

```
Origin IGP, MED not set, localpref 100, weight 32768
```

```
Received label 30000 50000
```

```
Extcommunity: RT:65503:30000 RT:65503:50000 ENCAP:8 Router MAC:ac7a.56fc.ffd7
```

IPv6 Prefix Verification

Host Route Details – Host B (VNI 30001)

```
Leaf1# show bgp l2vpn evpn 2023:192:168:23::1
```

```
BGP routing table information for VRF default, address family L2VPN EVPN
```

```
Route Distinguisher: 10.4.1.5:32790
```

```
BGP routing table entry for [2]:[0]:[0]:[48]:[0050.56bd.97b3]:[128]:[2023:192:168:23::1]/368,  
version 20
```

```
Paths: (1 available, best #1)
```

```
Flags: (0x000202) (high32 00000000) on xmit-list, is not in l2rib/evpn, is not in HW
```

```
Advertised path-id 1
```

```
Path type: internal, path is valid, is best path, no labeled nexthop
```

```
Imported to 2 destination(s)
```

```
Imported paths list: tenant-1 L3-50000
```

```
AS-Path: NONE, path sourced internal to AS
```

```
fd00::a03:3 (metric 8) from fd00::a02:2 (10.4.1.3)
```

```
Origin IGP, MED not set, localpref 100, weight 0
```

```
Received label 30001 50000
```

```
Extcommunity: RT:65503:30001 RT:65503:50000 ENCAP:8 Router MAC:ac7a.56fd.24f7
```

```
Originator: 10.4.1.5 Cluster list: 10.4.1.3
```

Source

Leaf-2
VTEP

Spine-1
RID

Who Gets this information

Route Reflector (Spine)

IPv6 Prefix Verification

IPv6 Route table for VRF Tenant-1

```
Leaf1# sh ipv6 route vrf tenant-1
```

```
IPv6 Routing Table for VRF "tenant-1"
```

```
'*' denotes best ucast next-hop
```

```
'**' denotes best mcast next-hop
```

```
'[x/y]' denotes [preference/metric]
```

```
2023:192:168:13::/64, ubest/mbest: 1/0, attached
```

```
    *via 2023:192:168:13::254, Vlan13, [0/0], 04:19:43, direct, , tag 12345
```

```
2023:192:168:13::1/128, ubest/mbest: 1/0, attached
```

```
    *via 2023:192:168:13::1, Vlan13, [190/0], 04:16:49, hmm
```

```
2023:192:168:13::254/128, ubest/mbest: 1/0, attached
```

```
    *via 2023:192:168:13::254, Vlan13, [0/0], 04:19:43, local, tag 12345
```

```
2023:192:168:23::/64, ubest/mbest: 1/0
```

```
    *via fd00::a03:3/128%default, [200/0], 04:21:27, bgp-65503, internal, tag 65503,  
segid 50000 VTEP:(fd00::a03:3, underlay_vrf: 1) encap: VXLAN
```

```
2023:192:168:23::1/128, ubest/mbest: 1/0
```

```
    *via fd00::a03:3/128%default, [200/0], 04:16:18, bgp-65503, internal, tag 65503,  
segid 50000 VTEP:(fd00::a03:3, underlay_vrf: 1) encap: VXLAN
```

IPv6 Prefix Verification

MAC Address Table

```
Leaf1# show mac address-table
```

Legend:

* - primary entry, G - Gateway MAC, (R) - Routed MAC, O - Overlay MAC
age - seconds since last seen, + - primary entry using vPC Peer-Link,
(T) - True, (F) - False, C - ControlPlane MAC, ~ - vsan,
(NA) - Not Applicable

	VLAN	MAC Address	Type	age	Secure NTFY Ports		
*	13	0050.56bd.bca2	dynamic	NA	F	F	Eth1/2
G	-	2020.0000.00aa	static	-	F	F	sup-eth1 (R)
G	-	ac7a.56fc.fffd7	static	-	F	F	sup-eth1 (R)
G	13	ac7a.56fc.fffd7	static	-	F	F	sup-eth1 (R)
G	3000	ac7a.56fc.fffd7	static	-	F	F	sup-eth1 (R)

IPv6 Prefix Verification

MAC Address Table

Leaf2# **show mac address-table**

Legend:

* - primary entry, G - Gateway MAC, (R) - Routed MAC, O - Overlay MAC
age - seconds since last seen, + - primary entry using vPC Peer-Link,
(T) - True, (F) - False, C - ControlPlane MAC, ~ - vsan,
(NA) - Not Applicable

VLAN	MAC Address	Type	age	Secure	NTFY	Ports
C 13	0050.56bd.bca2	dynamic	NA	F	F	nve1 (fd00::a03:0)
* 23	0050.56bd.97b3	dynamic	NA	F	F	Eth1/2
G -	2020.0000.00aa	static	-	F	F	sup-eth1 (R)
G -	ac7a.56fd.24f7	static	-	F	F	sup-eth1 (R)
G 13	ac7a.56fd.24f7	static	-	F	F	sup-eth1 (R)
G 23	ac7a.56fd.24f7	static	-	F	F	sup-eth1 (R)
G 3000	ac7a.56fd.24f7	static	-	F	F	sup-eth1 (R)

IPv6 Prefix Verification

IPv6 Neighbor Discovery (~ARP)

Leaf1# `show ipv6 neighbor vrf tenant-1` ← (`show ip arp for ipv6`)

IPv6 Adjacency Table for VRF tenant-1

Total number of entries: 2

Address	Age	MAC Address	Pref	Source	Interface	Mobility Flags
2023:192:168:13::1	02:54:41	0050.56bd.bca2	50	icmpv6	Vlan13	
fe80::b6e0:1e7e:40f2:7b7c	02:54:35	0050.56bd.bca2	50	icmpv6	Vlan13	

Link Local IPv6

Leaf2# `show ipv6 neighbor vrf tenant-1`

IPv6 Adjacency Table for VRF tenant-1

Total number of entries: 2

Address	Age	MAC Address	Pref	Source	Interface	Mobility Flags
2023:192:168:23::1	02:51:37	0050.56bd.97b3	50	icmpv6	Vlan23	
fe80::b6e0:1e7e:40f2:7b7c	02:51:32	0050.56bd.97b3	50	icmpv6	Vlan23	

Link Local IPv6

Conclusion

Fill out your session surveys!



Attendees who fill out a minimum of four session surveys and the overall event survey will get **Cisco Live-branded socks** (while supplies last)!



Attendees will also earn 100 points in the **Cisco Challenge** for every survey completed.



These points help you get on the leaderboard and increase your chances of winning daily and grand prizes

Continue your education



- Visit the Cisco Showcase for related demos
- Book your one-on-one Meet the Engineer meeting
- Attend the interactive education with DevNet, Capture the Flag, and Walk-in Labs
- Visit the On-Demand Library for more sessions at www.CiscoLive.com/on-demand




The bridge to possible

Thank you



#CiscoLive




Cisco Live Challenge

Gamify your Cisco Live experience!
Get points for attending this session!

How:

- 1 Open the Cisco Events App.
- 2 Click on 'Cisco Live Challenge' in the side menu.
- 3 Click on View Your Badges at the top.
- 4 Click the + at the bottom of the screen and scan the QR code:



The background is a vibrant, abstract graphic. It features a central bright white light source from which numerous colorful rays emanate, creating a sunburst or starburst effect. The rays transition through a spectrum of colors including yellow, orange, red, and various shades of blue and green. Overlaid on this are several large, semi-transparent, wavy shapes in similar color tones, giving the overall image a sense of motion and energy.

cisco *Live!*

Let's go

#CiscoLive

Apendix



Features supported with IPv6v In Underlay

- Cisco Nexus Dashboard Fabric Controller (NDFC) integration
- Address Resolution Protocol (ARP) suppression in the overlay
- Access Control List (ACL) Quality of Service (QoS)
- Border Node with VRF-Lite
- Dynamic Host Configuration Protocol (DHCP)
- Internet Group Management Protocol (IGMP) Snooping in the overlay
- Virtual Extensible Local Area Network (VXLAN) Operation, Administration, and Maintenance (OAM)

Features supported with IPv6v In Underlay

- Virtual Port Channel (vPC) with VIP and PIP support
- VXLAN Policy-Based Routing (PBR)
- vPC Fabric Peering
- Storm Control for host ports (Access Side)
- VXLAN Access Features

Private VLAN (PVLAN)

802.1x

Port Security

Port VLAN Translation

QinVNI

SelQinVNI

QinQ QinVNI

Features not supported (yet) with IPv6 Underlay

- EVPN Multi-homing with Ethernet Segment (ES)
- VXLAN Flood and Learn
- Multicast underlay (PIM-BiDir, Protocol Independent Multicast (PIM) Any Source Multicast (ASM), Snooping)
- Tenant Routed Multicast (TRM)
- VXLAN Multi-Site
- Downstream VNI

Features not supported (yet) with IPv6 Underlay

- Bidirectional Forwarding Detection (BFD)
- Centralized Route Leak
- NetFlow
- Peer vtep command
- Virtual Network Functions (VNF) Multipath

Underlay Network Protocols

Multicast Protocol

Device Role	Interface	Protocol	Function
Spine	Loopback254	PIM v2/PIM v6	Router RP
Spine, Leaf	Loopback0	PIM v2/PIM v6	Router ID
Leaf	Loopback1	PIM v2/PIM v6	NVE Source Int
Spine, Leaf	Fabric Interface	PIM v2/PIM v6	Switch Peering
Leaf (vPC)	vPC backup SVI	PIM v2/PIM v6	vPC peer link routing