



The bridge to possible

Model-Driven Streaming Telemetry and Programmability on Cisco NX-OS

Nick Mortari
Technical Marketing Engineer
Cloud Networking Team

BRKDCN-2904

CISCO *Live!*

#CiscoLive

Cisco Webex App

Questions?

Use Cisco Webex App to chat with the speaker after the session

How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space

Webex spaces will be moderated by the speaker until June 7, 2024.






Agenda

- Why Are We Here Today?
- What Does “Model-Driven” Mean?
- Building Blocks of Programmability
- Building Blocks of Telemetry
- Open-Source Telemetry Stack
- Demonstrations Throughout

Why Are We Here Today?




Why Are We Here Today?

- I need to change port VLANs
- I need to change MTU values
- I need to configure routes

You



- Are my BGP neighbors up?
- Are my links fully utilized?
- Are my devices overheating?



Why Are We Here Today?

Data

- Can't access the information I need
- Can't push or pull from different devices

Performance

- My tools don't scale well
- Can't push or pull data fast enough

Why Are We Here Today?


Model-driven methods save the day!




Why Are We Here Today?



Access
Detailed
Information

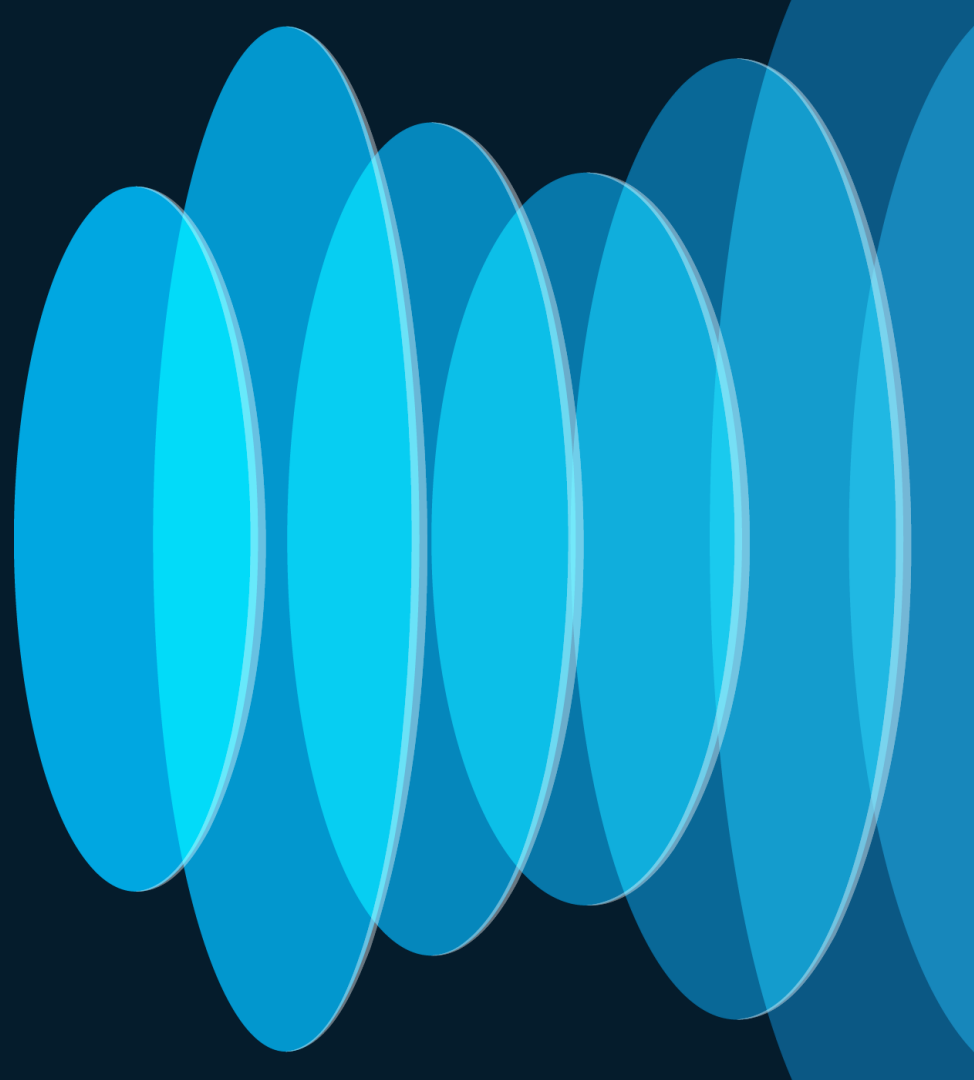


Scalability




Performance

What Does “Model-Driven” Mean?



What Is Model-Driven?




Using a mutually agreed upon method for structuring information



Example: Standard address format must be used when sending mail



What Is Model-Driven Programmability?



Sending configuration information to a device using a common model

I understand this!



I understand this!



I understand this!



010110
110010
001011



I know what to send!



Example: Set CDP hold time on my switches

What Is Model-Driven Telemetry?



Receiving state information from a device using a common model

I know what to send!



I know what to send!




I know what to send!



010110
110010
001011




I understand this!



Example: Collect current power consumption from my devices

Model-Driven Programmability




Building Blocks of Programmability




- **Data Structure**
- Data Encoding
- Data Transport

Data Structure

Two Options




DME (Data Management Engine)



YANG Models

Data Structure

DME - Programmability




- Configuration and operational data is stored in DME
- Tree data structure
- DN (Distinguished Name) is in .../.../.../... format
- Configuration data can be accessed with the DN as a sensor path

- `sys/bgp/inst` represents configuration and state data for BGP process

Data Structure

DME - Programmability



- Almost entire OS is available
- As of 10.4(2)F, almost all commands are DMEized
- We can use DME paths for programmability

Data Structure

DME – Programmability Paths

Visore is a built-in DME browser of NX-OS, navigate to [https://\[switch_ip\]/visore.html](https://[switch_ip]/visore.html)

rmonIfIn	
broadcastPkts	199779
clearTs	never
discards	0
dn	sys/intf/phys-[eth1/27]/dbgIfIn < > [1] [34]
errors	0
modTs	2022-03-28T16:45:11.658+00:00
multicastPkts	345290
nUcastPkts	545069
noBuffer	0
octetRate	3657496
octets	11346525403646
packetRate	3438
rateInterval	300
ucastPkts	3777158007
unknownEtype	0
unknownProtos	0

API reference is also available:

[https://developer.cisco.com/site/nxapi-dme-model-reference-api/?version=10.3\(2\)](https://developer.cisco.com/site/nxapi-dme-model-reference-api/?version=10.3(2))

lACP:If

The LACP information that is operated at an interface (member port of the port channel) level.

Telemetry Sensor Path(s)

- `sys/lacp/inst/if-[id]`

Configurable Properties

PROPERTY NAME	DATA TYPE	DESCRIPTION	PERMITTED VALUES
adminSt	nw:IfAdminSt (nw:AdminSt)	The administrative state of the object or policy.	SELECTION: 1 - enabled 2 - disabled
descr	naming:Descr1024 (string:Basic)	Description	MAX SIZE: 1024
name	naming:Name256 (string:Basic)	The name of the object.	MAX SIZE: 63
prio	lacp:PortPrio (scalar:UInt16)	Specifies the LACP interface port priority. When there is a limitation that prevents all compatible ports from aggregating then the port priority should be set to standby mode. A higher port priority value means a lower priority for LACP.	RANGE: [1, 65535] DEFAULT: 32768
txRate	lacp:TxRate (scalar:Enum8)	Specifies the rate at which the LACP packets are transmitted.	SELECTION: 1 - normal 2 - fast DEFAULT: normal



Data Structure


YANG Models

- YANG (Yet Another Next Generation) is a data modeling language
- Defines the data structure and data type for the model we use



← Data Structure

↗ Data Type





Data Structure

YANG Models - Programmability

- NX-OS supports two YANG models for telemetry
 - OpenConfig YANG model
 - Cisco native YANG model
- Configuration elements can be accessed with YANG model paths



Change MTU

OpenConfig path:

[/interfaces/interface/config/mtu](#)

Cisco Native path:

[/System/intf-items/phys-items/PhysIf-list/mtu](#)




Data Structure

YANG Models – Choosing Your Model



Cisco Native Model

- Vendor specific
- Created by Cisco
- Supports almost every feature on NX-OS



You can
use both!



OPENCONFIG

OpenConfig Model

- Vendor agnostic
- Created by many networking companies (open-source)
- Does not support every feature on NX-OS

NX-OS

Cisco Model

OpenConfig Model

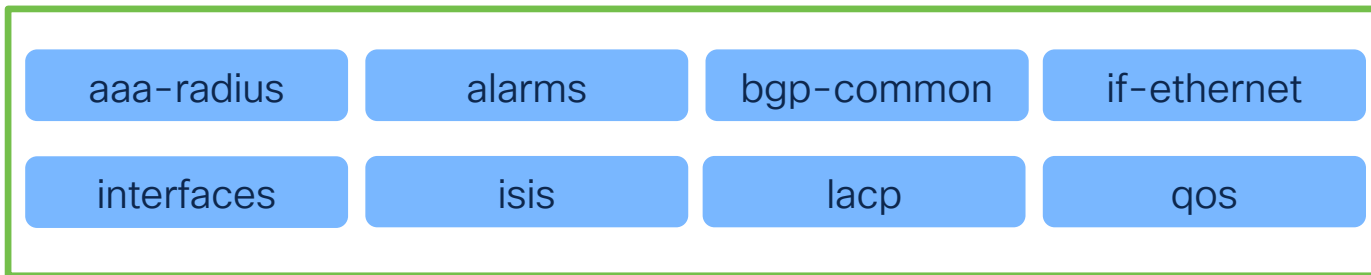


Data Structure

YANG Models – Modules

- YANG models are broken down into modules
- Allows for easier changes when working with others
- The Cisco model is technically “one large module”

OpenConfig



Data Structure

YANG Models – OpenConfig



- openconfig-bgp.yang
- openconfig-evpn-types.yang
- openconfig-evpn.yang
- openconfig-extensions.yang
- openconfig-if-aggregate.yang
- openconfig-if-ethernet.yang
- openconfig-if-ip-ext.yang
- openconfig-if-ip.yang
- openconfig-if-types.yang
- openconfig-igmp-types.yang
- openconfig-igmp.yang
- openconfig-inet-types.yang
- openconfig-interfaces.yang
- openconfig-isis-lsdb-types.yang
- openconfig-isis-lsp.yang
- openconfig-isis-policy.yang

- 100+ OpenConfig modules supported
- To enable OC YANG Model:
 - Before 10.2(2), mtz-openconfig-all rpm needs to be installed on switch
 - After 10.2(2), run feature openconfig

Data Structure

YANG Models – OpenConfig

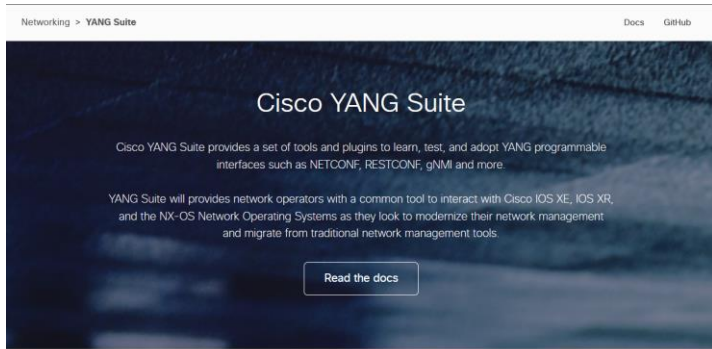


- cisco-nx-openconfig-acl-deviations.yang
- cisco-nx-openconfig-bfd-deviations.yang
- cisco-nx-openconfig-bgp-policy-deviations.yang
- cisco-nx-openconfig-if-aggregate-deviations.yang
- cisco-nx-openconfig-if-ethernet-deviations.yang
- cisco-nx-openconfig-if-ip-deviations.yang
- cisco-nx-openconfig-if-ip-ext-deviations.yang
- cisco-nx-openconfig-interfaces-deviations.yang
- cisco-nx-openconfig-lacp-deviations.yang
- cisco-nx-openconfig-lldp-deviations.yang
- cisco-nx-openconfig-macsec-deviations.yang
- cisco-nx-openconfig-network-instance-deviations.yang

- Be aware of deviations, modules can be partially supported
- Deviation can mean:
 - Path does not follow the original YANG module definition
 - Path is not supported
- A full list of supported modules and deviations can be found at:
<https://github.com/YangModels/yang/tree/master/vendor/cisco/nx>

Data Structure

YANG Suite

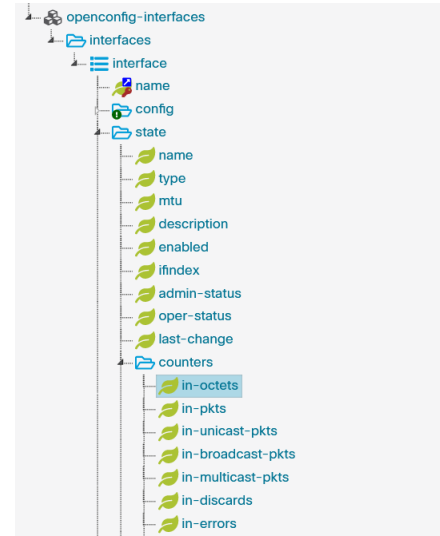


YANG Suite In Your Network

Network automation and programmability capabilities include browsing YANG modules in a graphical interface, creating RPC payload messages to interact with devices, and a gRPC Dial-Out model driven telemetry collector for streaming telemetry. The user-interface is updated with HTML5 and provides flexible deployment options with Docker containers.



- Tool to assist with YANG model exploration and testing
- Includes YANG browser for both models



<https://developer.cisco.com/yangsuite>


Building Blocks of Programmability




- Data Structure
- **Data Encoding**
- Data Transport

Data Encoding

Two Options



XML (eXtensible Markup Language)



JSON (JavaScript Object Notation)

Data Encoding

XML

```
<interface>
  <name>eth1/49</name>
  <config>
    <access-vlan>200</access-vlan>
    <interface-mode>ACCESS</interface-mode>
    <description>To server</description>
    <duplex>auto</duplex>
    <admin-state>up</admin-state>
    <speed>50000</speed>
    <mtu>1500</mtu>
  </config>
</interface>
```

- Every element must have an opening and closing tag
- Value is placed between the tags
- Low transfer efficiency

 Tag

 Value

Data Encoding

JSON

```
{  
  "interface": {  
    "name": "eth1/49",  
    "config": {  
      "access-vlan": "200",  
      "interface-mode": "ACCESS",  
      "description": "To server",  
      "duplex": "auto",  
      "admin-state": "up",  
      "speed": "50000",  
      "mtu": "1500"  
    }  
  }  
}
```

- Every element must have a key
- Value is placed after the key
- Better transfer efficiency

 Key


 Value

Building Blocks of Programmability


- Data Structure
- Data Encoding
- **Data Transport**

Data Transport


Four Options




NETCONF (Network Configuration)



RESTCONF (RESTful Configuration)



NX-API REST




gNMI (Google Remote Procedure

Call Network Management Interface)

Data Transport

NETCONF

- Protocol that connects over SSH to leverage a set of RPCs
- Supports configuration management with `edit-config`
- Supports YANG as the data structure
- Supports XML for data encoding
- Enabled with feature `netconf` command



Data Transport

RESTCONF

- Protocol that connects over HTTP/HTTPS to pass data
- Supports configuration management through RESTful methods
- Supports YANG as the data structure
- Supports XML and JSON for data encoding
- Enabled with feature `restconf` command



Data Transport

NX-API REST


- Protocol that connects over HTTP/HTTPS to pass data
- Supports configuration management through RESTful methods
- Supports DME as the data structure
- Supports XML and JSON for data encoding
- Enabled with feature `nxapi` command



Data Transport

gNMI

- Protocol that connects over HTTP/2 to leverage a set of RPCs
- Supports configuration management with `gnmi set`
- Supports DME and YANG as the data structure
- Supports XML and JSON for data encoding
- Enabled with feature `grpc` command



Data Transport

NX-OS Sandbox

The screenshot displays the NX-API Sandbox web interface. At the top, the Cisco logo is on the left, followed by 'NX-API Sandbox'. To the right, user information is shown: 'Username:admin Role:network-admin Version:10.4(1) Management ip:10.195.233.164 Hostname:RU30'. Further right are links for 'NX-API References', 'Yang Browser', and 'Yang Models', and a 'Logout' button.

The main area is divided into two sections. The top section shows a CLI configuration snippet: `router bgp 65221`, `router-id 205.0.0.14`, and `address-family ipv4 unicast`. The bottom section shows a RESTCONF request. The 'Method' is set to 'RESTCONF (Yang)' and the 'Message format' is 'json'. The request URL is `restconf/data/Cisco-NX-OS-device:System/`. Below the URL bar are 'Send', 'Reset', and 'Convert' buttons. The 'Request' tab is selected, showing a JSON payload:

```
{
  "bgp-items": {
    "inst-items": {
      "asn": "65221",
      "dom-items": {
        "Dom-list": [
          {
            "name": "default",
            "rtrId": "205.0.0.14",
            "rtrIdAuto": "disabled",
            "af-items": {
              "DomAf-list": [
                {
```

 A 'Copy' button is visible next to the JSON response.

- Tool to construct payloads and scripts from CLI
- Navigate to [https://\[switch_ip\]/](https://[switch_ip]/)

Data Transport

gNMI - Firewalls

gNMI Client

Firewall


```
permit tcp any any eq 22
permit tcp any any eq 443
deny ip any any
```

NX-OS

gRPC
Server

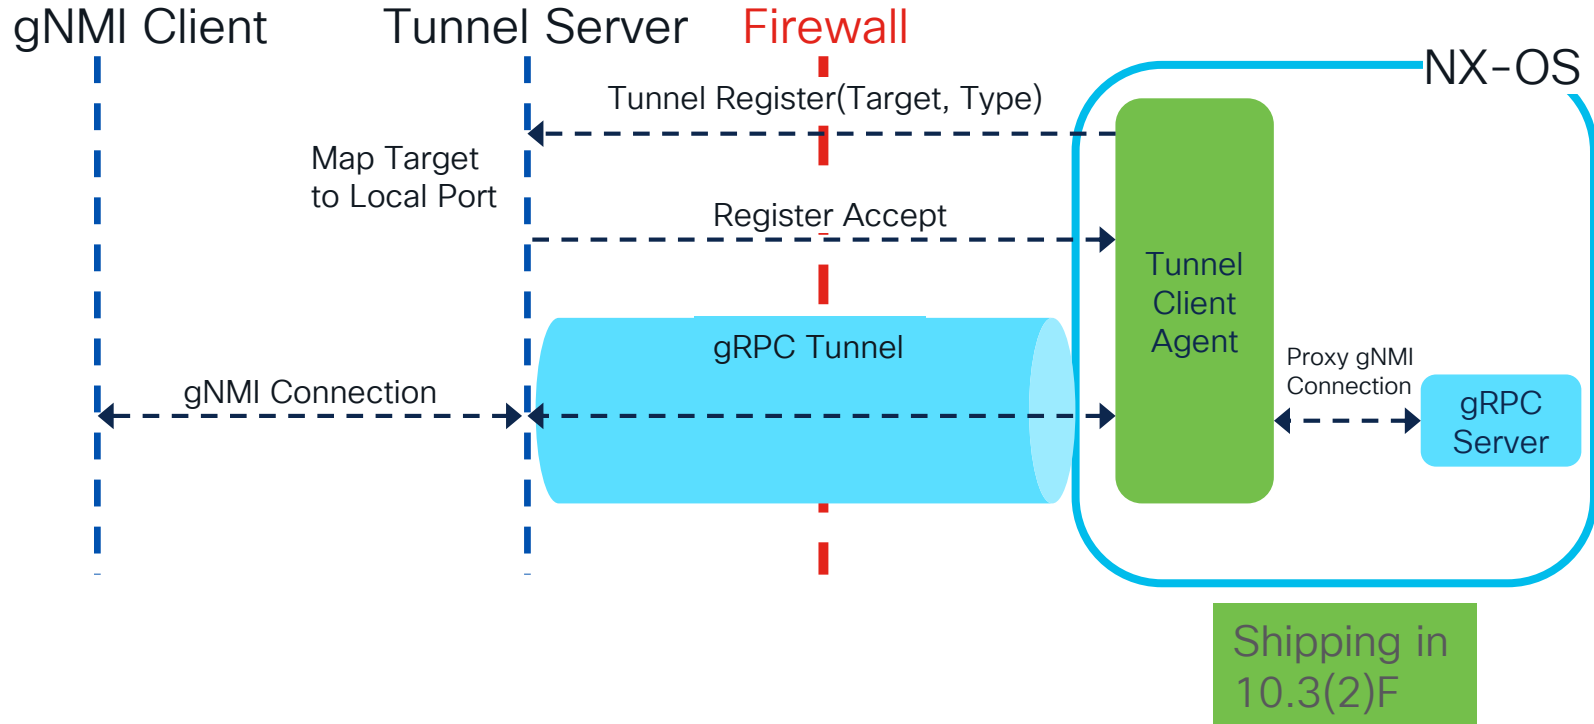
Data Transport

gNMI – NGINX Proxy



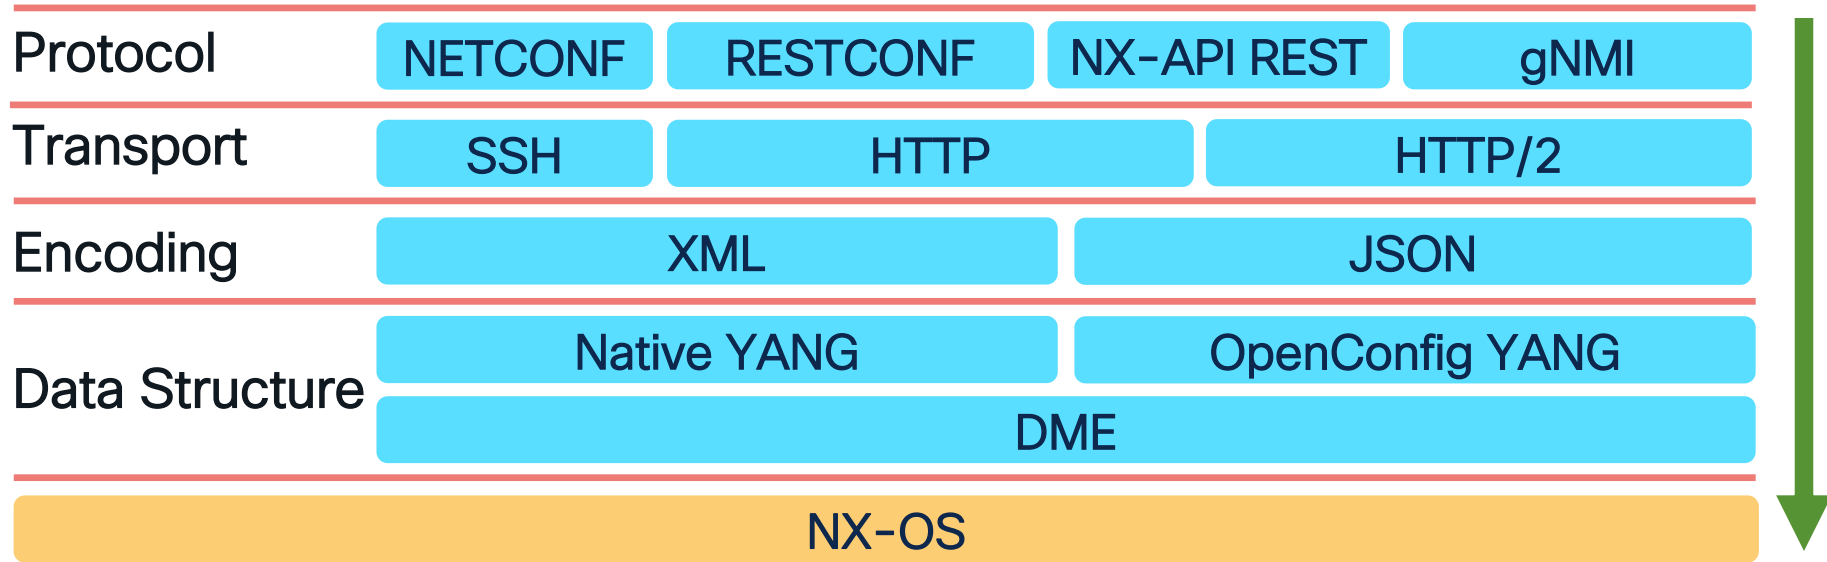
Data Transport

gNMI – gRPC Tunnel



Data Transport

Programmability Overview




Data Transport

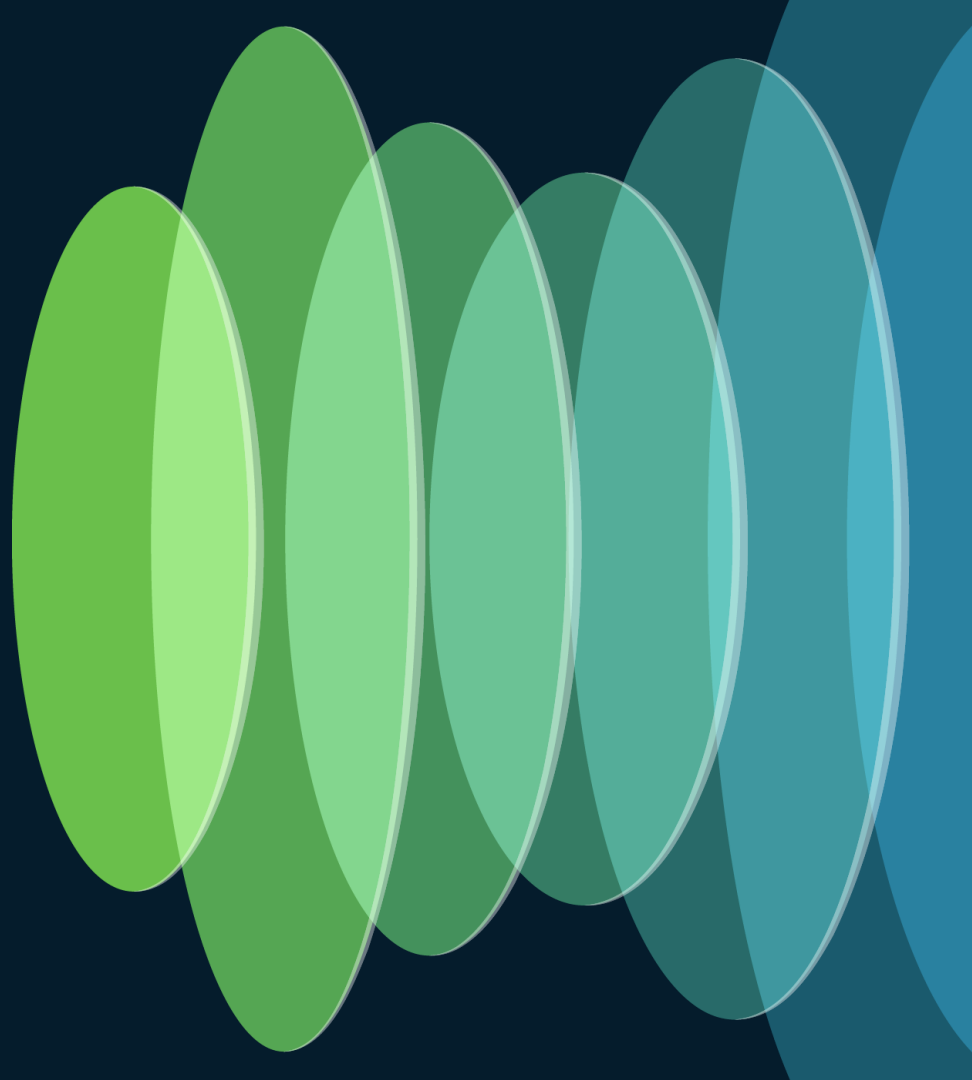
Programmability Options Compared

	NETCONF	RESTCONF	NX-API REST	gNMI
Data Model	YANG	YANG	DME	YANG / DME
Encoding	XML	JSON / XML	JSON / XML	JSON / XML
Transport Method	SSH	HTTP / HTTPS	HTTP / HTTPS	HTTP/2

Demo




Model-Driven Streaming Telemetry




Model-Driven Streaming Telemetry

Performance



~~Telemetry~~ The Lemon Tree




Building Blocks of Streaming Telemetry




- **Data Structure**
- Data Frequency
- Data Encoding
- Data Transport

Data Structure


Three Options (well...actually two)



DME



YANG

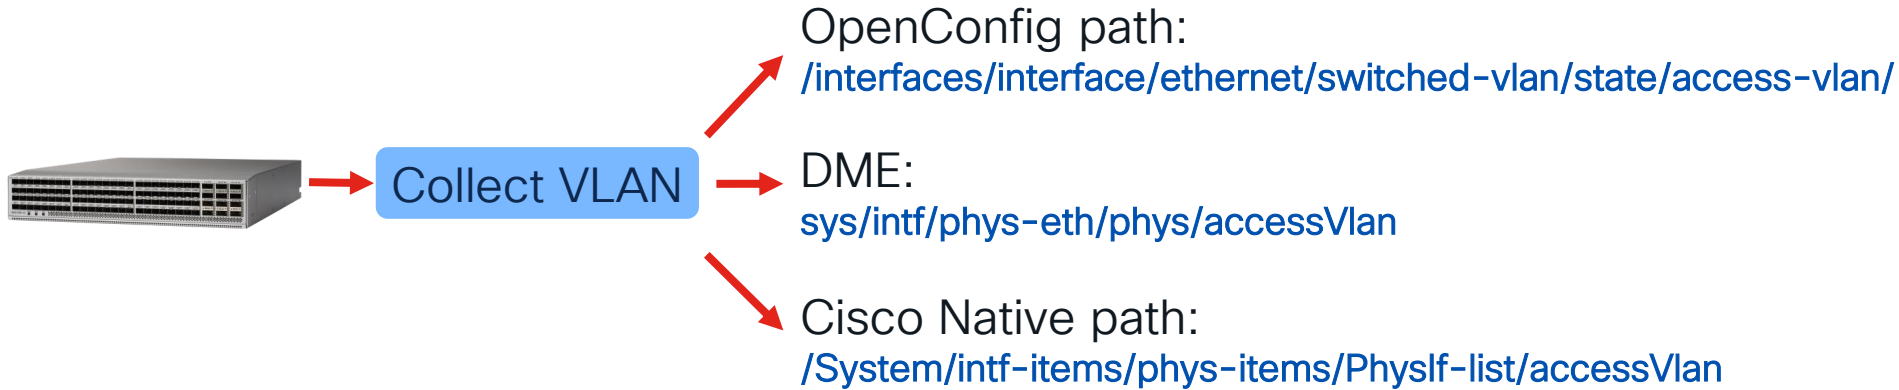


NX-API CLI

Data Structure

DME and YANG - Telemetry

- DME and YANG paths can also be used for telemetry
- We can “subscribe” to a path and receive a stream of structured data



Data Structure

NX-API CLI

```
93240YC-FX2-L02-S4# show nve vni | json-pretty
{
  "TABLE_nve_vni": {
    "ROW_nve_vni": [
      {
        "if-name": "nve1",
        "vni": "30000",
        "mcast": "239.1.1.1",
        "vni-state": "Up",
        "mode": "CP",
        "type": "L2 [2300]",
        "flags": null,
        "dci-mcast": "Unconfigured"
      },
      ...
    ]
  }
}
```



- 100% of customer-facing show commands of NX-OS have output



- Only supports sample-based telemetry
- CLI doesn't have data types, all values are strings
 - The collector will need to parse the result and “guess” data type

Data Structure

Platform Support

Nexus Platform	DME	CLI/NX-API	YANG	Release
3000 with 8G+ RAM	✓	✓	✓ *	7.0(3)I7(1)
9300	✓	✓	✓ *	7.0(3)I5(1)
9500/9400/9800	✓	✓	✓ *	7.0(3)I7(1)
7000/7700	✗	✓	✗	8.3(1)

* Streaming YANG models starting from 9.2(1)

Building Blocks of Streaming Telemetry




- Data Structure
- **Data Frequency**
- Data Encoding
- Data Transport

Data Frequency

Two Options

Sample-Based Collection



Collect information at
every sample interval

Event-Based Collection




Collect information at
every event

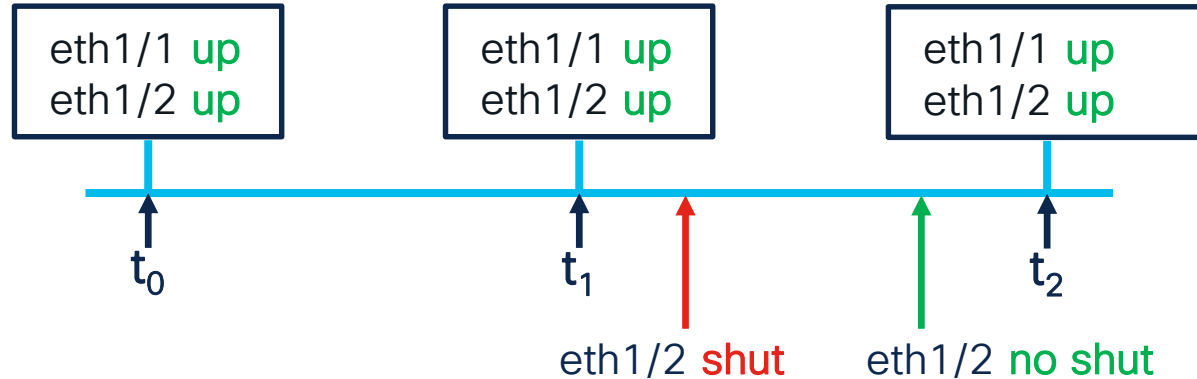
Data Frequency

Sample-Based Collection

Scenario 1



Scenario 2



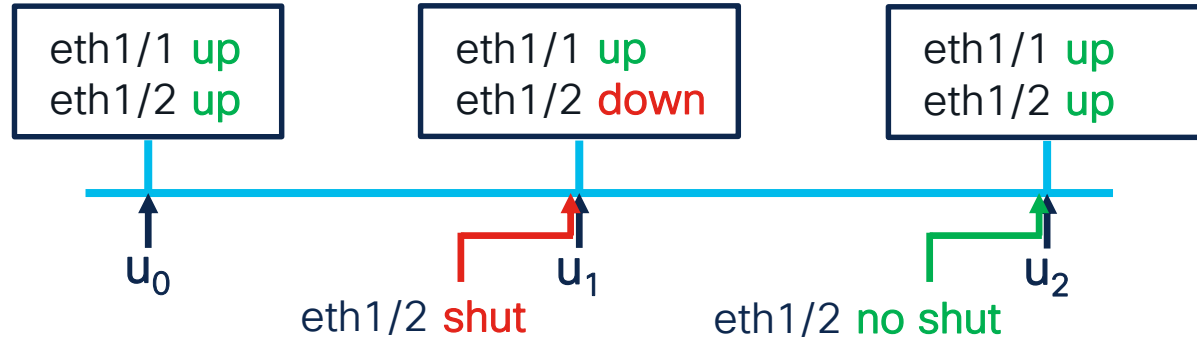
Data Frequency

Event-Based Collection

Scenario 1



Scenario 2




Building Blocks of Streaming Telemetry


- Data Structure
- Data Frequency
- **Data Encoding**
- Data Transport

Data Encoding


Three Options



XML (eXtensible Markup Language)



JSON (JavaScript Object Notation)



GPB (Google Protocol Buffers)

Data Encoding

GPB - The Need for Speed

XML and JSON are great
but...

You



Data Encoding

GPB (Google Protocol Buffers)

```
<interface>
  <name>eth1/49</name>
  <state>
    <counters>
      <in-broadcast-pkts>2</in-broadcast-pkts>
      <in-discards>0</in-discards>
      <in-errors>0</in-errors>
      <in-fcs-errors>0</in-fcs-errors>
      <in-multicast-pkts>30543</in-multicast-pkts>
      <in-octets>13320913920</in-octets>
      <in-unicast-pkts>5406026</in-unicast-pkts>
      <in-unknown-protos>0</in-unknown-protos>
      <out-broadcast-pkts>3</out-broadcast-pkts>
      <out-discards>0</out-discards>
      <out-errors>0</out-errors>
      <out-multicast-pkts>26070</out-multicast-pkts>
      <out-octets>143144868</out-octets>
      <out-unicast-pkts>1424051</out-unicast-pkts>
    </counters>
  </state>
</interface>
```



Variable Length Integers!
(varints)

```
1:"eth1/49"
2:{
  1:{
    1:2
    2:0
    3:0
    4:0
    5:30543
    6:13320913920
    7:5406026
    8:0
    9:3
    10:0
    11:0
    12:26070
    13:143144868
    14:1424051
  }
}
```

Data Encoding

GPB (Google Protocol Buffers)

```
<interface>
  <name>eth1/49</name>
  <state>
    <counters>
      <in-broadcast-pkts>2</in-broadcast-pkts>
      <in-discards>0</in-discards>
      <in-errors>0</in-errors>
      <in-fcs-errors>0</in-fcs-errors>
      <in-multicast-pkts>30543</in-multicast-pkts>
      <in-octets>13320913920</in-octets>
      <in-unicast-pkts>5406026</in-unicast-pkts>
      <in-unknown-protos>0</in-unknown-protos>
      <out-broadcast-pkts>3</out-broadcast-pkts>
      <out-discards>0</out-discards>
      <out-errors>0</out-errors>
      <out-multicast-pkts>26070</out-multicast-pkts>
      <out-octets>143144868</out-octets>
      <out-unicast-pkts>1424051</out-unicast-pkts>
    </counters>
  </state>
</interface>
```



```
1:"eth1/49"
2:{
  1:{
    1:2
    2:0
    3:0
    4:0
    5:30543
    6:13320913920
    7:5406026
    8:0
    9:3
    10:0
    11:0
    12:26070
    13:143144868
    14:1424051
```

High wire efficiency,
but hard to develop the encoder and decoder

Data Encoding

GPB-KV (Key-Value)

```
"counters":{  
  "in-octets": 13320913920,  
  "out-octets": 143144868  
}
```

```
message TelemetryField {  
  uint64      timestamp = 1;  
  string      name = 2;  
  oneof value_by_type {  
    bytes      bytes_value = 4;  
    string     string_value = 5;  
    bool       bool_value = 6;  
    uint32     uint32_value = 7;  
    uint64     uint64_value = 8;  
    sint32     sint32_value = 9;  
    sint64     sint64_value = 10;  
    double     double_value = 11;  
    float      float_value = 12;  
  }  
  repeated TelemetryField fields = 15;  
}
```

```
{  
  2:"in-octets"  
  8:0x319FD0400  
},  
{  
  2:"out-octets"  
  8:0x88837A4  
}
```



Building Blocks of Streaming Telemetry

- Data Structure
- Data Frequency
- Data Encoding
- **Data Transport**

Data Transport

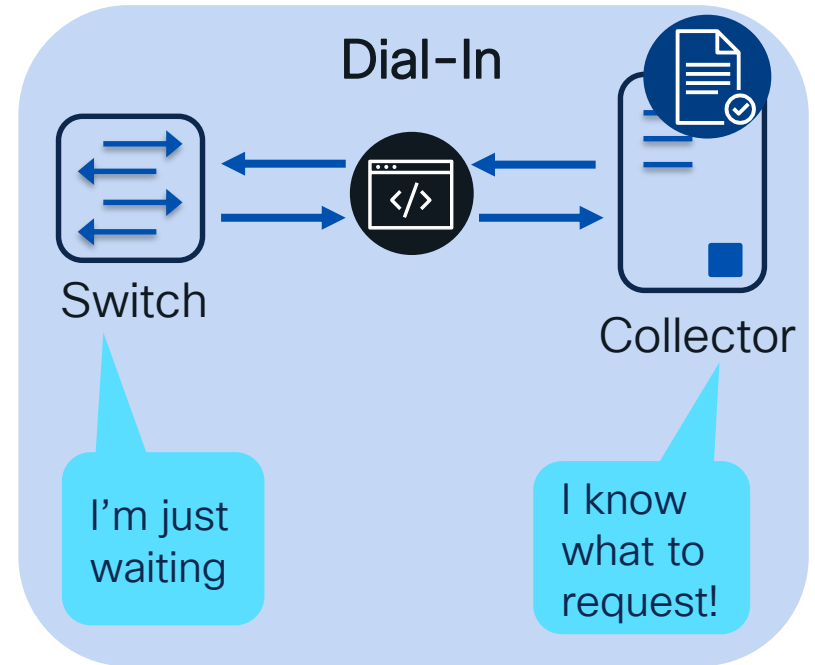
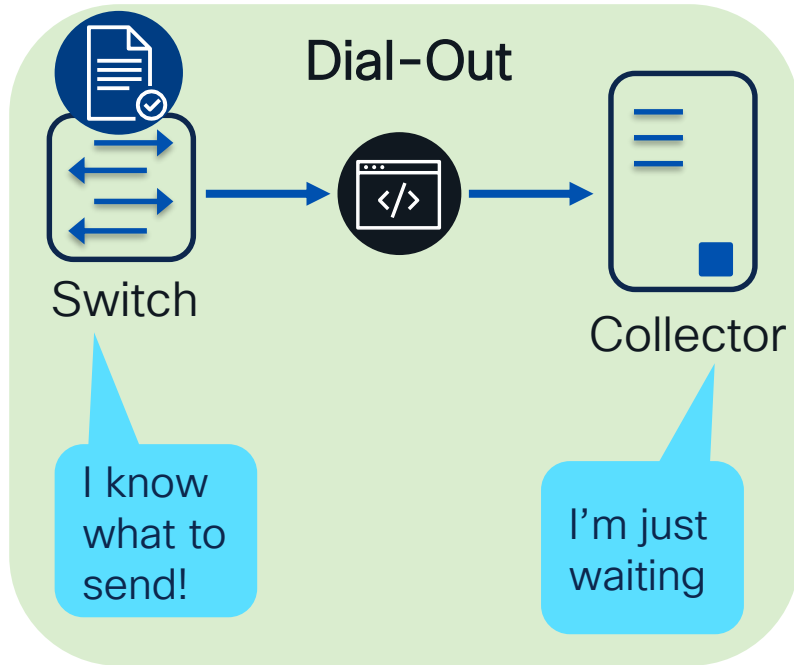
Dial-Out vs Dial-In

- TCP connection is always persistent in telemetry
- The difference is which part initializes the connection



Data Transport


Dial-Out vs Dial-In: Configuration



Data Transport

Design Considerations


Dial-Out



Collectors can be set up behind load balancer, all switches stream to the same VIP of collector

CISCO *Live!*

Dial-In



To distribute workload, the collectors need to dial-in to different switch groups, need to keep the sensor configuration synchronized across the cluster

Data Transport

Dial-Out vs Dial-In

Dial-Out

Protocol

Supports gRPC, HTTP, UDP as transport protocols

Configuration

Telemetry configuration must be setup on the switch

Access

No need to open a specific port to the management interface

Load Balancing

Easier to load balance with collectors behind VIP

Dial-In

Only gNMI protocol supported

Single channel for both subscription and data transport

External firewalls must allow ingress connection to the gNMI service

gNMI clients need to be distributed to handle entire fabric

Data Transport

gNMI - Features



Capabilities

Collect capabilities from device



Get

Collect current data values from device



Set

Modify data on the device



Subscribe

Subscribe to a data stream for a path in the data model

Data Transport

gNMI - Features



sample

Sample-based collection



on_change

Event-based collection



target_defined

Switch decides which frequency type



suppress_redundant

Don't send redundant data



heartbeat_interval

Every X interval, give me an update

Data Transport

gNMI – NX-OS Implementation

RPC

- gNMI fully supported from NX-OS 9.3(5)
- NX-OS 10.4(x) updated to gNMI version 0.8.0

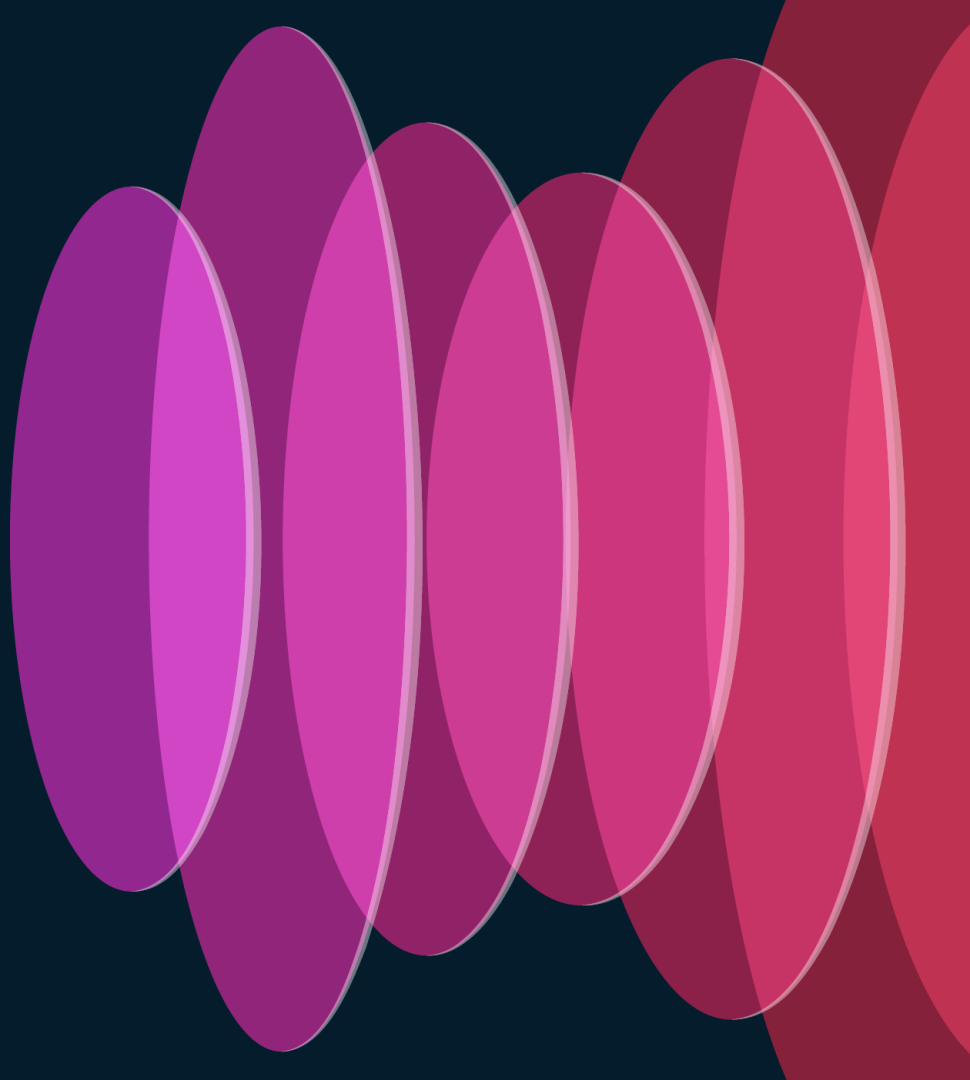
Security

- gNMI leverages TLS
- Mutual TLS supported

Data Encoding

- Native and OpenConfig models supported
- Supports Google Protobuf and JSON as encoding
- Wild card supported from NX-OS 10.2(2)

Open-Source Telemetry Stack



Open-Source Telemetry Stack

Requires Three Pieces

1

Collection Agent

A service that understands the data collected from the device

2

Time Series Database


A database with very precise time stamping that stores the collected data

3

Using Stored Data


Integrating the data with an automation system, or graphically displaying the data

Open-Source Telemetry Stack



gNMI
Dial-In

Telemetry
Dial-Out



Write Data



https://github.com/dsx1123/telemetry_collector

Takeaways

- NX-OS has several choices for the data model, encoding, and transport options.
 - Customers can choose based on business requirements
- Most customers are interested in gNMI dial-in
 - Pros and cons between dial-out and dial-in
- Use OpenConfig model first, fall back to native model and DME

Continue Your Journey

- DEVNET-2135: Mastering Network Automation: Unleashing the Power of gNMI (Wednesday)
- DEVNET-2235: Industry Standard Streaming Telemetry with Cisco NX-OS (Thursday)
- Explore the NX-OS sandbox
- Install YANG suite to explore the models
- Install the TIG stack to get started with telemetry

Links

- API guide

[https://developer.cisco.com/site/nxapi-dme-model-reference-api/?version=10.3\(2\)](https://developer.cisco.com/site/nxapi-dme-model-reference-api/?version=10.3(2))

- Supported YANG modules

<https://github.com/YangModels/yang/tree/main/vendor/cisco/nx>

- Cisco YANG Suite

<https://developer.cisco.com/yangsuite/>

- TIG dashboard

https://github.com/dsx1123/telemetry_collector

- BGP Python script

<https://github.com/nmortari/gNMI-Programming/blob/main/gNMI-BGP-Setup.py>

Continue your education



- Visit the Cisco Showcase for related demos
- Book your one-on-one Meet the Engineer meeting
- Attend the interactive education with DevNet, Capture the Flag, and Walk-in Labs
- Visit the On-Demand Library for more sessions at www.CiscoLive.com/on-demand



The bridge to possible

Thank you

CISCO *Live!*

#CiscoLive