



The bridge to possible

# A Tale of Two Ships

SD-Access and BGP EVPN Designs and Deployments  
on Cruise Ships

Cheeho Yan, Sr. Technical Leader  
Yianni Thallas, Product Management Architect  
BRKENS-2822

CISCO *Live!*

#CiscoLive

# Cisco Webex App


## Questions?

Use Cisco Webex App to chat with the speaker after the session

## How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space

Webex spaces will be moderated by the speaker until June 7, 2024.



# How Did the Tale of Two Ships Start?




- Network implementation started from around 2022-23
- SD-Access LISP (SDA) & BGP EVPN deployed on two ships from two different cruise lines
- Several additional ships might implement the same solutions across both fleets
- We will get to how the tale ended

# Agenda


- Fabric Choices: SD-Access vs. BGP EVPN
- Business Requirement
- Ship #1: SD-Access Solution
- Ship #2: BGP EVPN Solution
- Plan, Design, Implementation, and Support Best Practices
- Sneak Preview: Catalyst Center Orchestrated BGP EVPN

# Fabric Choices: SD-Access vs. BGP EVPN



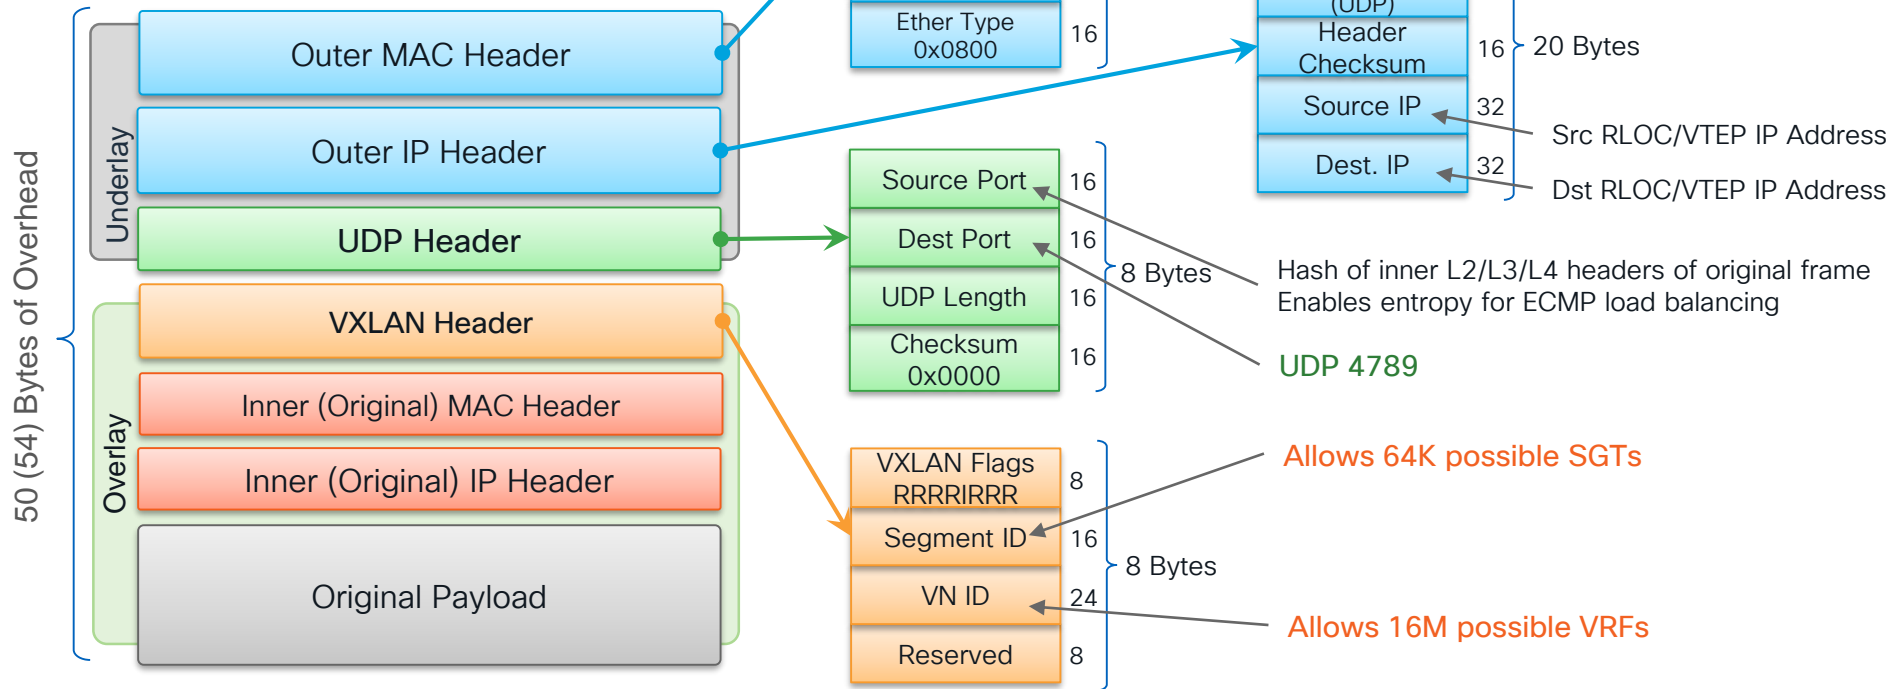
# A Fabric is an Overlay

- A logical topology used to virtually connect devices, built on top of physical underlay topology
- Provides additional services not provided by the underlay




# VXLAN Header

## MAC-in-IP Encapsulation




# Fabrics with VXLAN Encapsulation

VXLAN Provides a Network with Segmentation & Scale



SDA LISP

- **Network Simplification** Lightweight, extensible, massive scale with rapid convergence. Single overlay for wired/wireless
- **Mobility First Requirement** Fabric Integrated Wireless, L2 mobility, enhanced wireless performance
- **Segmentation** Zero-Trust Architecture with Micro and Macro Segmentation. Unified Wired + Wireless policy



BGP EVPN

- **One Fabric Architecture (Campus & DC)** Operational ease with a single familiar protocol
- **Multi-vendor interoperability** Vendor-agnostic solution
- **Flexibility** Customizable overlay network types and topologies



**CISCO** *Live!*



© 2024 Cisco and/or its affiliates. All rights reserved. Cisco Public

- 9

# Enterprise Campus BGP EVPN Drivers



Industry Standard



Multi-vendor IT strategy



One Fabric Architecture



Unified operation across network infrastructure



Proven and Scalable



BGP protocol history. Minimum new learning curve



Hierarchical Fabric Domain



Multi-tier overlay network architecture



Flexible Overlay



Use-case driven. Customized overlay networks types & topologies

# BGP EVPN System Roles



## BORDER-GATEWAY:

A gateway point of between two or more BGP EVPN administrative domain boundary



## BORDER :

A gateway point of between EVPN fabric and external network domain



## INTERMEDIATE :

A Layer 2 or Layer 3 (IP/MPLS) underlay network system providing basic transport and forwarding plane



## SPINE :

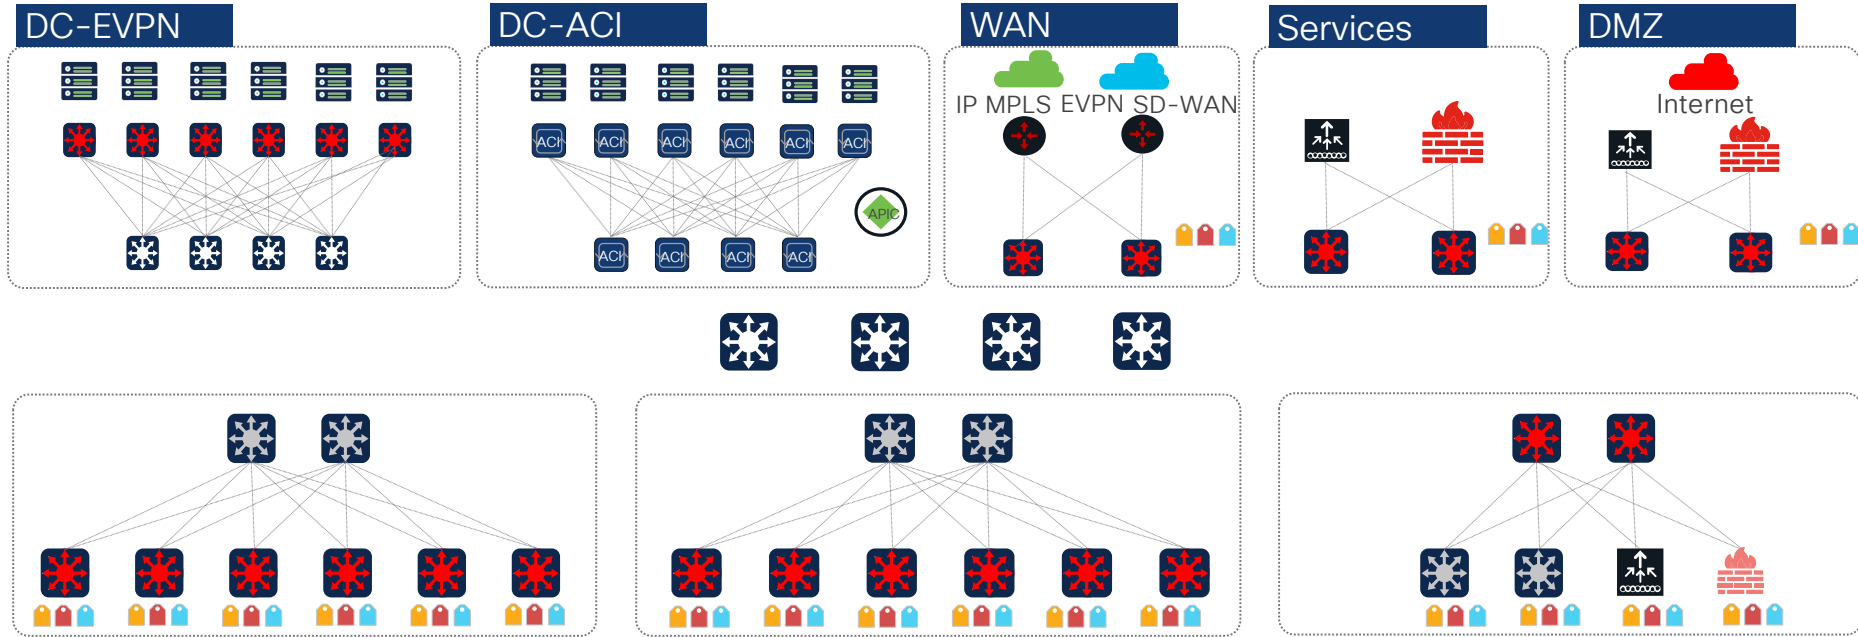
A BGP EVPN role that reflects the L2/L3 VPN prefixes providing hierarchical neighbor peering, learning and distribution point



## VTEP (LEAF) :

An origination and termination point of VXLAN enabled overlay network

# Enterprise BGP EVPN Reference Architecture



Industry Standard



Unified Fabric



Proven



Hierarchical



Flexible

**CISCO** Live!

# SDA-LISP: Recommended by Cisco for Campus

3,600+ Campus Deployments  
30% Year to Year Growth

## Simple

- Industry standard-based and optimized for enterprise campus
- Configurations automated based on business intent

## Efficient

- Scales to a large number of endpoints
- Rapid converge


## Extensible

- Simple directory structure to map endpoints to locations
- Highly extensible to address new use cases

## Wired + Wireless

- Centralized control plane with distributed data plane
- Seamless campus wide mobility

# Business Requirement



# Responses to Meet Business Needs

## Requirement


Business Continuity

Services to End Users

Ease of Network Deployment ,  
Operation & Support



SD-Access



BGP EVPN

# Business Continuity



- Network Downtime impacts revenue



- Network redundancy is required



- Limited network change windows



- Ships are offshore most of the time



Note:



Specific to cruise line industry



# Services to End Users



- Most endpoints are connected using wireless



- Extensive use of IP based IOT devices



- Seamless endpoint onboarding



- Ease of problem identification/troubleshooting/resolution



- Large number of APs and cabin switches



# Ease of Network Deployment / Operation / Support



- Standard based network architecture



- Tight timeline for network implementation



- Secured endpoint admission control



- Security segmentation




- Onboard + Onshore IT support model




# Two Cruise Ships: Similar Physical Configuration

- Passengers & Crews: 3,000 – 4,500
- Wired & Wireless Endpoints: up to 10,000
- Network Domains: Data Centers (servers) + Campus (passengers, crews)
- Network Infrastructure: Green field implementation



# Ship #1: SD- Access Solution



# Challenges & Solutions




Automation, Assurance, TrustSec,  
Wire/Wireless integrated in Fabric

Catalyst Center / SDA



Limited Satellite Bandwidth

Catalyst Center Air-Gap



Extended Nodes Exceeding Limit

Fabric Zone



Need L2 Connectivity to Servers

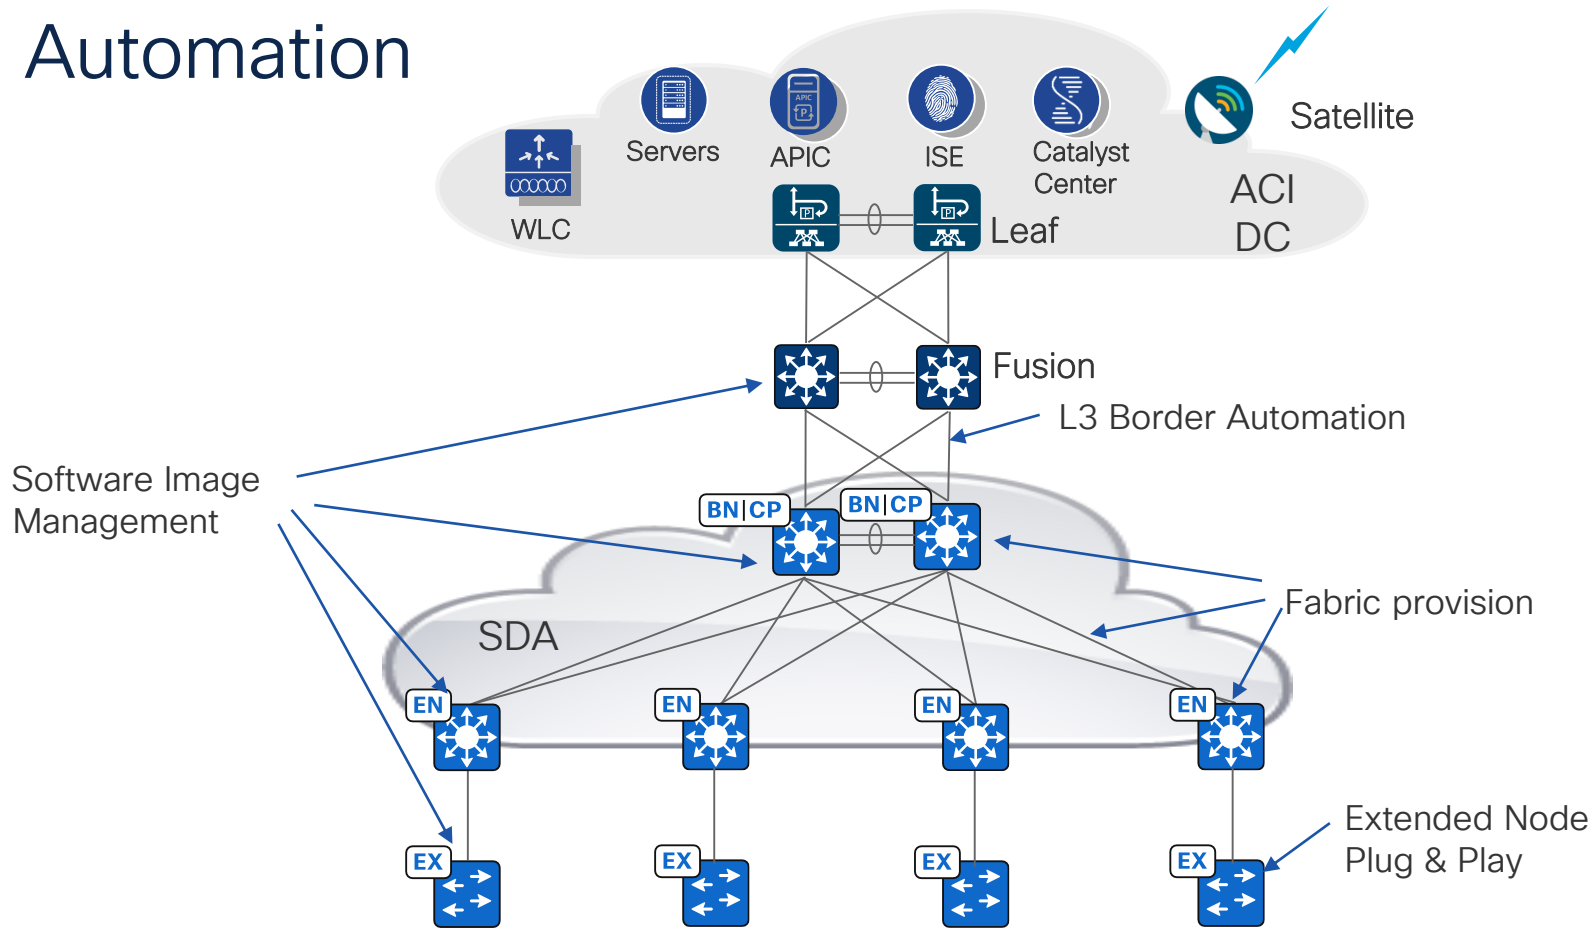
Layer 2 Virtual Network



Survivability for Critical Services

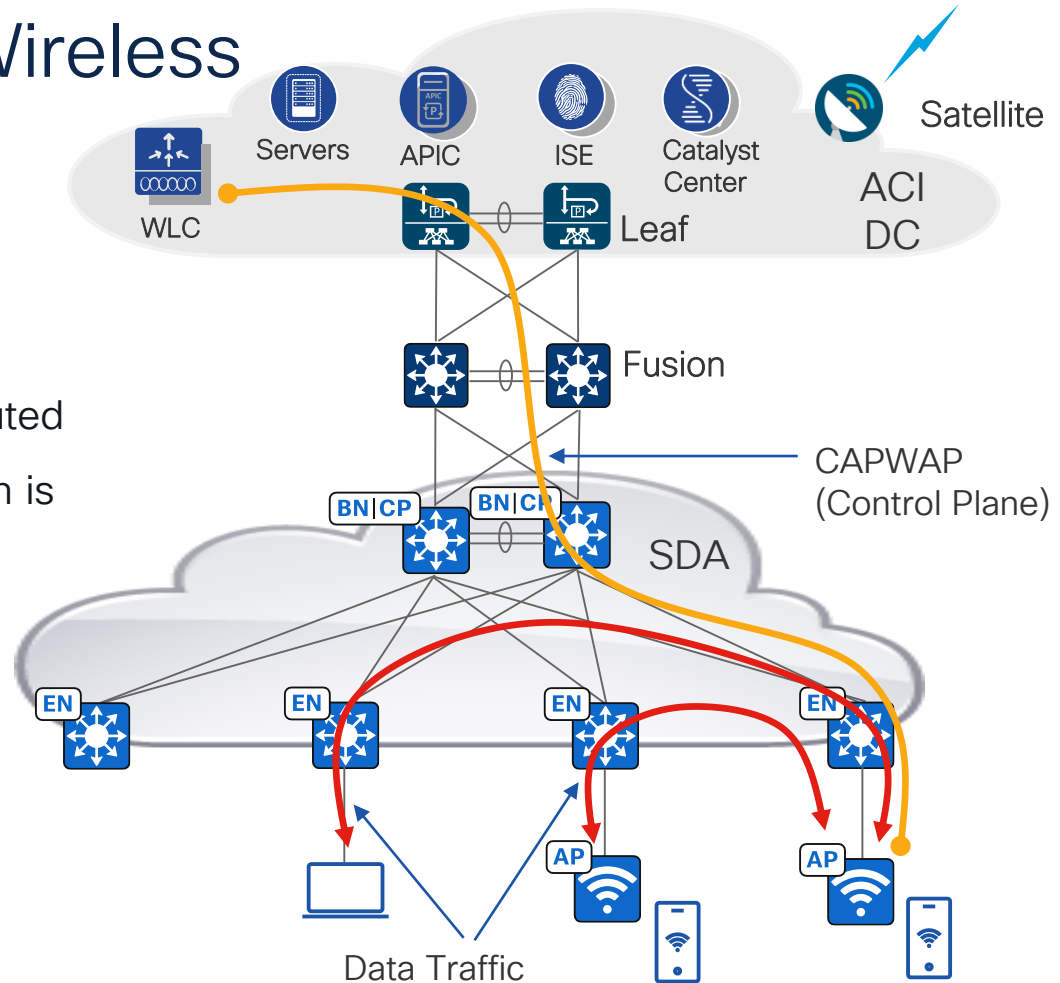
Spread Network Devices Across Fire  
Zones

# Fabric Automation



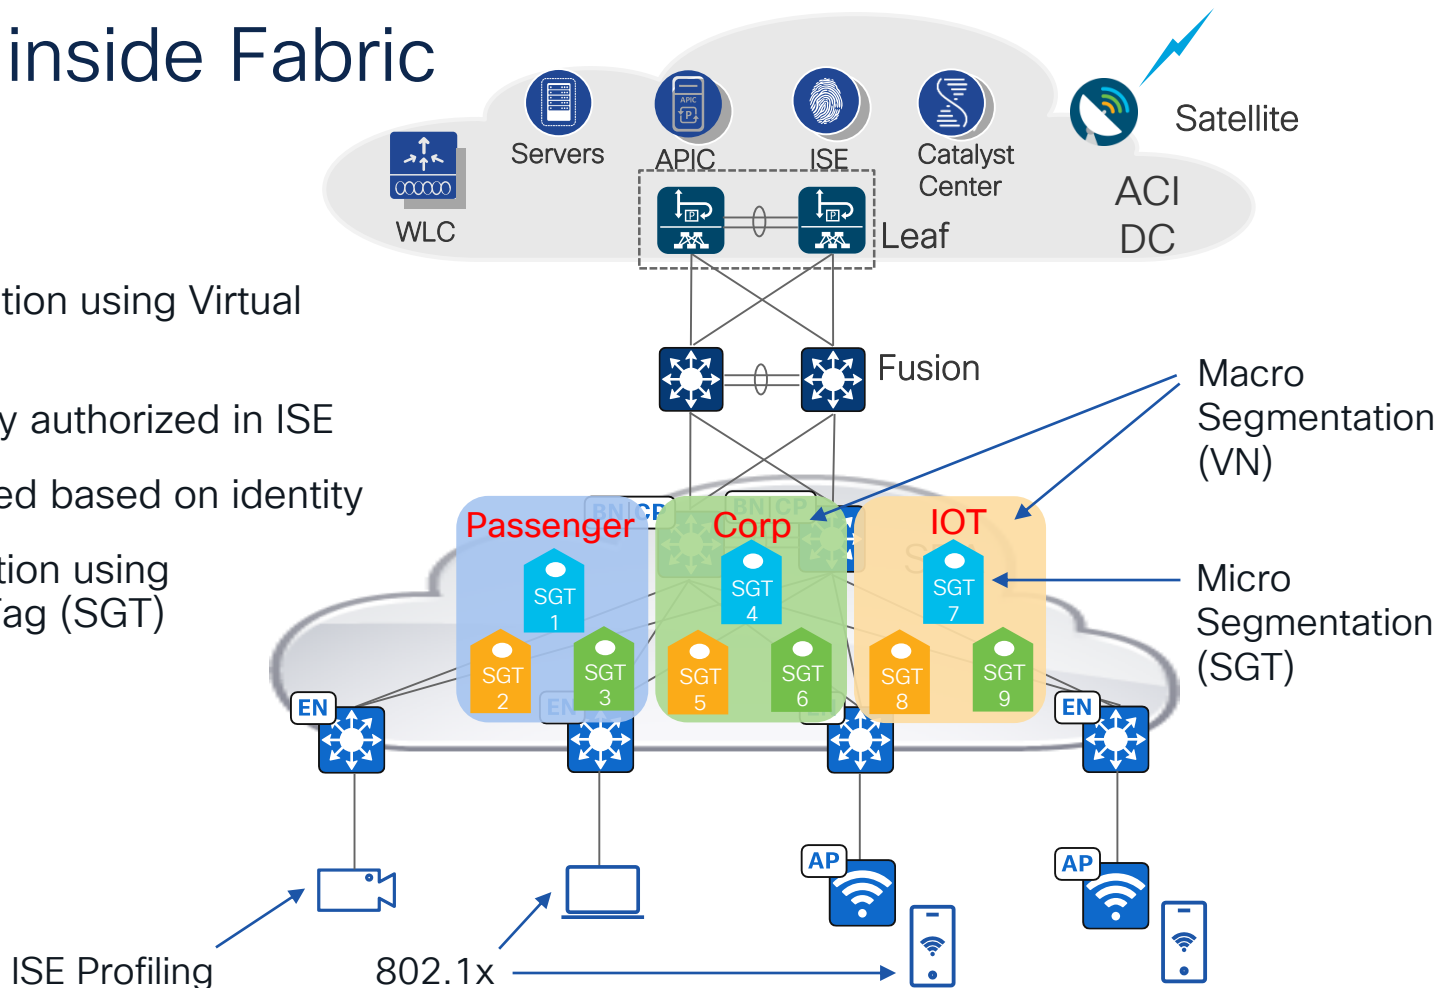
# Fabric Enabled Wireless

- No air-pinning to centralized controller
- Wireless client traffic is distributed
- Wired-Wireless communication is directly through Fabric



# TrustSec inside Fabric

- Macro segmentation using Virtual Networks (VN)
- Centralized policy authorized in ISE
- Endpoints grouped based on identity
- Micro Segmentation using Security Group Tag (SGT)





# Fabric Assurance

## Assurance Dashboard

- Health & critical issues summary
- Fabric, device & client health scores
- Wireless trends

## Device & Client 360 Views

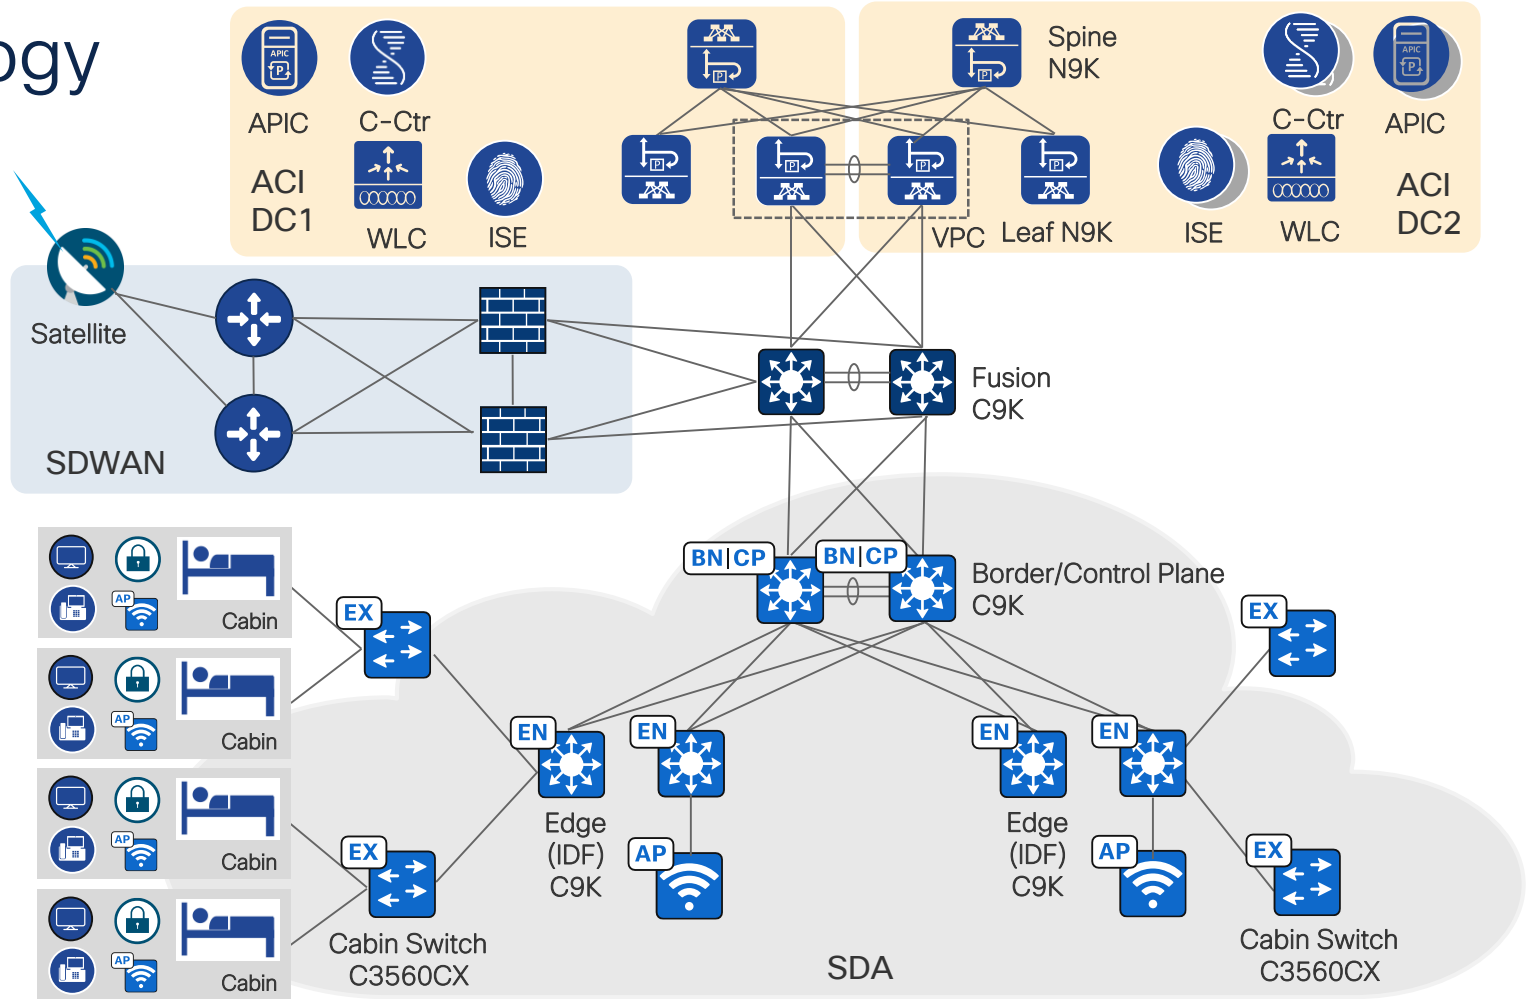
- Time travel view
- View connected neighbors & clients
- Universal search for elements in the network

## Troubleshooting

- Issues with explanations
- Suggested actions



# Topology



# Satellite Bandwidth

## Challenge:

- Limited Internet bandwidth through Satellite
- Bandwidth priority is given to the paid passengers



## Solution: Catalyst Center in Air-Gap Mode



### Typical Use Cases

- Department of Defence
- Government



# Airgap Implementation Considerations

Supported	Not Supported
Catalyst Center Software Offline Updates	Geo Maps Update
Assurance (Except AI Network Analytics)	License Manager (Devices that do not support SLR/PLR)
Application Policy, Topology, Third Party SDK, Audit Log, Global Search, Home Page, Bonjour, Plug and Play, PSIRTs	AI Network Analytics
SWIM: Manual Image Import, Manual KGV Import, Addons (Manual import required) – SMU, Sub package, APSP, APDP	SWIM Addons – ROMMON; Automated Image Download from cisco.com
SDA (with DHCP server in Airgap)	Integrations outside Airgap
License Manager (SLR/PLR Supported)	Make a Wish



Ensure Catalyst Center does not have backdoors to Internet to download online packages automatically, otherwise it would cause upgrade issues

# Scale

## Challenge:

- Supported fabric devices per SDA Fabric site is 1,200 for DN2-HW-APL-XL
- Cruise line required 2000+ Extended Nodes to service all the cabins
- Same IP pools stretching multi-Fabric site was not an option

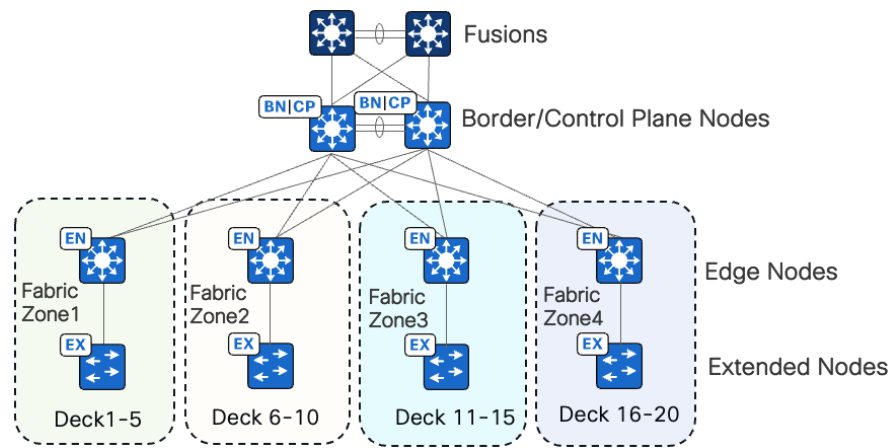


## Solution: Fabric Zone



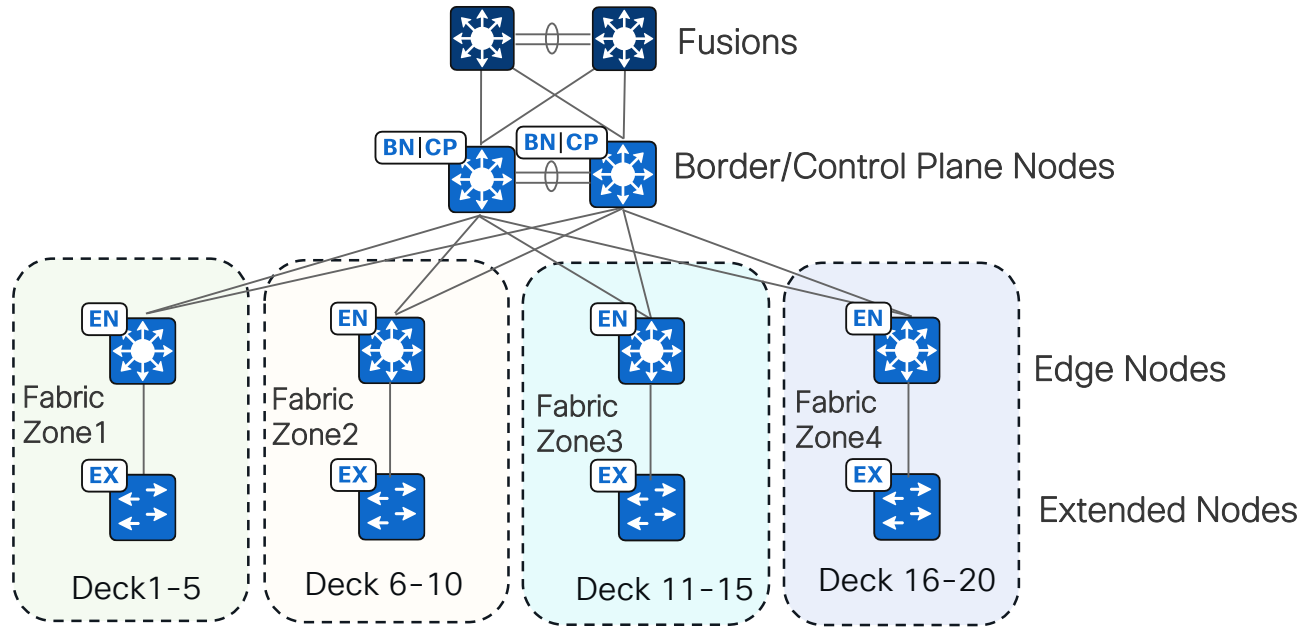
### Typical Use Cases


- Manufacturing



# Fabric Zones on the Ship

- 4 Zones with 500 Extended Nodes each
- No common changes to all Fabric Zones at the same time
- Onboard 50 Extended Nodes at a time



 Review design with Cisco first. Multiple Fabric Site is the preferred choice in other use cases

# L2 Connectivity to DC

## Challenge:

- Certain legacy applications require servers at DC to be on the same L2 segment as the endpoints at the Fabric
- Security requirement for Gateway of a subnet to be outside the Fabric

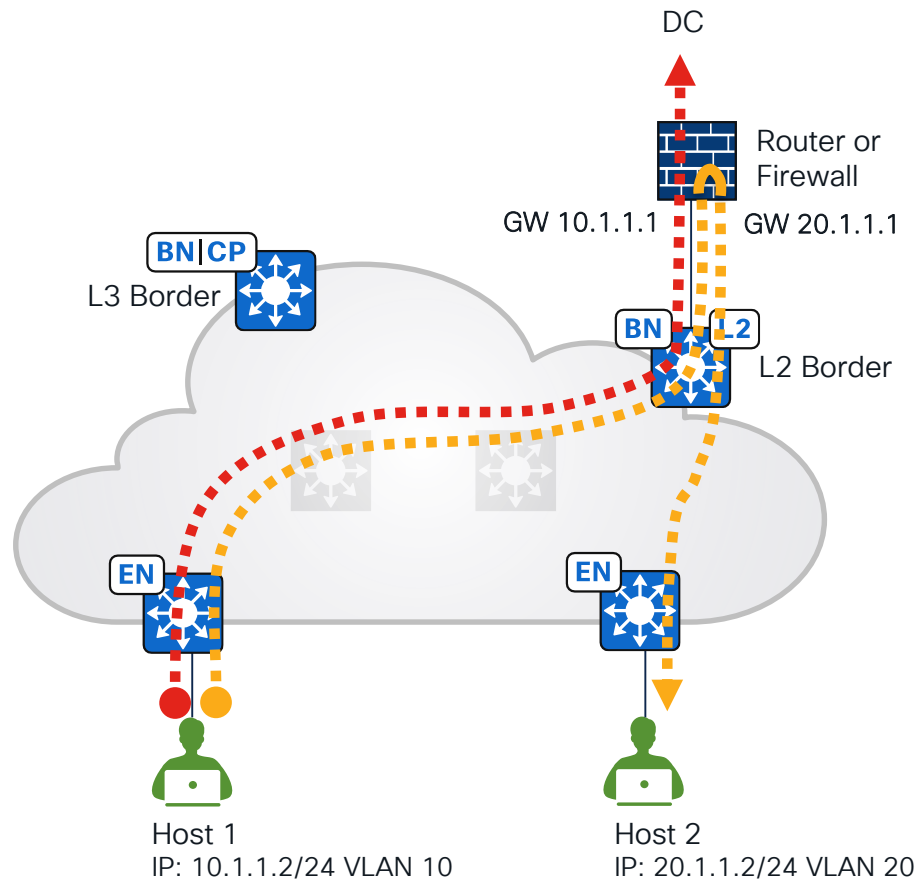


## Solution: Layer 2 Virtual Network



### Typical Use Cases

- Manufacturing
- IoT segments
- BMS



# Location Redundancy


## Challenge:

- DCs and critical network devices need to have location redundancy

Solution: Spread them across Fire Zones and locations


## Typical Use Cases

- Transportation
- Financial





# Ship #2: BGP EVPN Solution



# Challenges & Solutions



Automate Fabric Overlay from DC to Campus

Use Nexus Dashboard Fabric Controller (NDFC)



Reduce Fabric Fault Domain

Multisite VXLAN BGP EVPN



L2 Connectivity to Servers

Layer 2 overlay tunnels



Survivability for Critical Services

Spread network devices across fire zones

# Fabric Automation

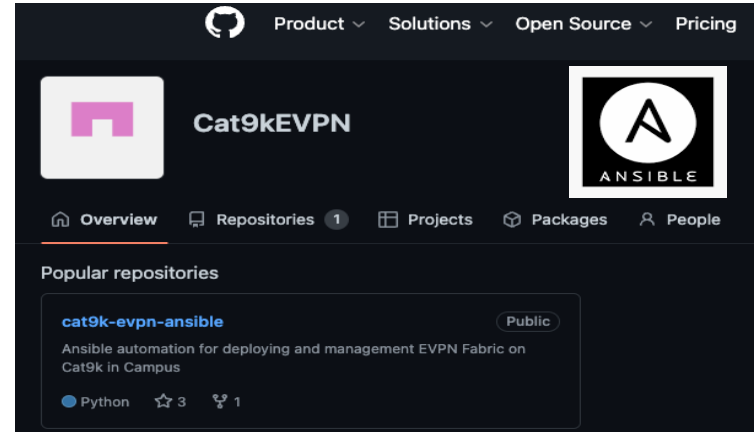
## Challenge:

- Need automation for Nexus and Catalyst switches with BGP EVPN deployment with a single pane of glass



## Solution:

- NDFC with built in python & CLI policies for configuration generation
- Cat 9000 programmability




# Multisite

## Challenge:

- Single fabric fault domain from DC to campus is not scalable

## Solution:

- Implement DC site and Campus site
- Back-to-back multi-site anycast BGW

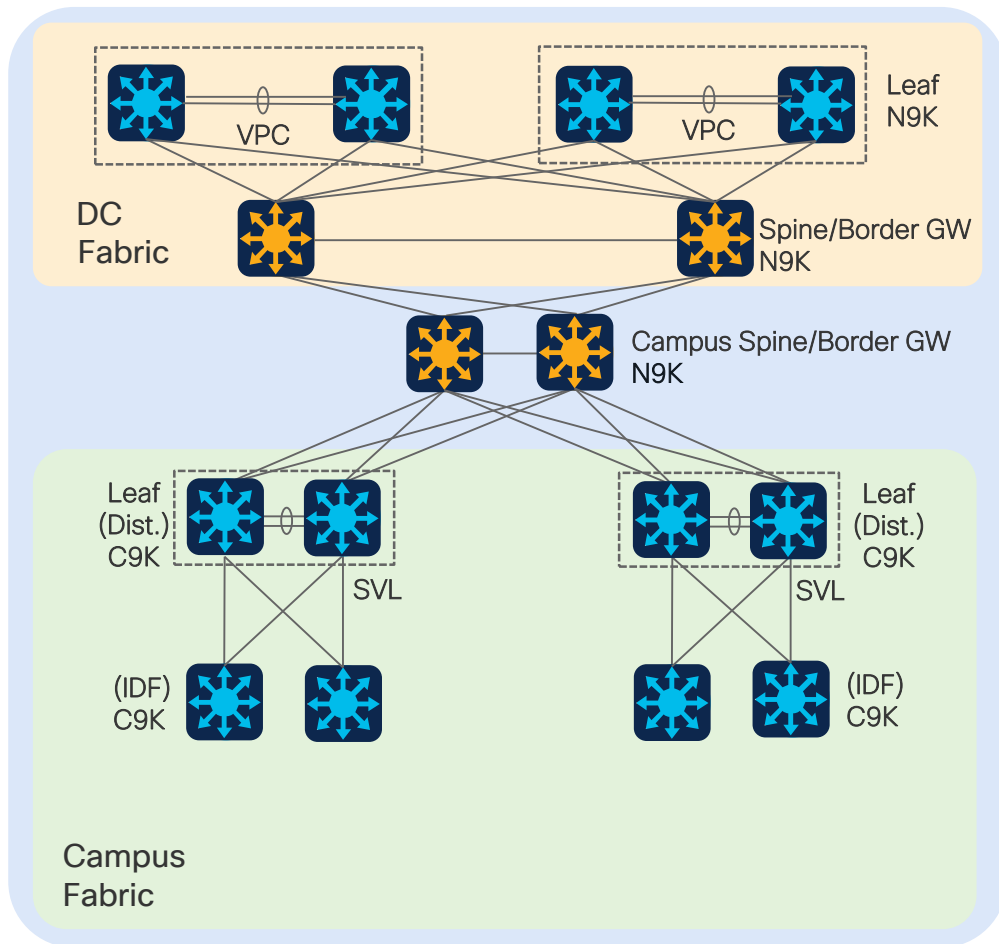


# Multisite

- Separate replication domains for BUM (Broadcast, Unknown Unicast, and Multicast) traffic
- Granular control on cross-site L2 and L3 communication
- Better overall scalability
- Smaller fault domain

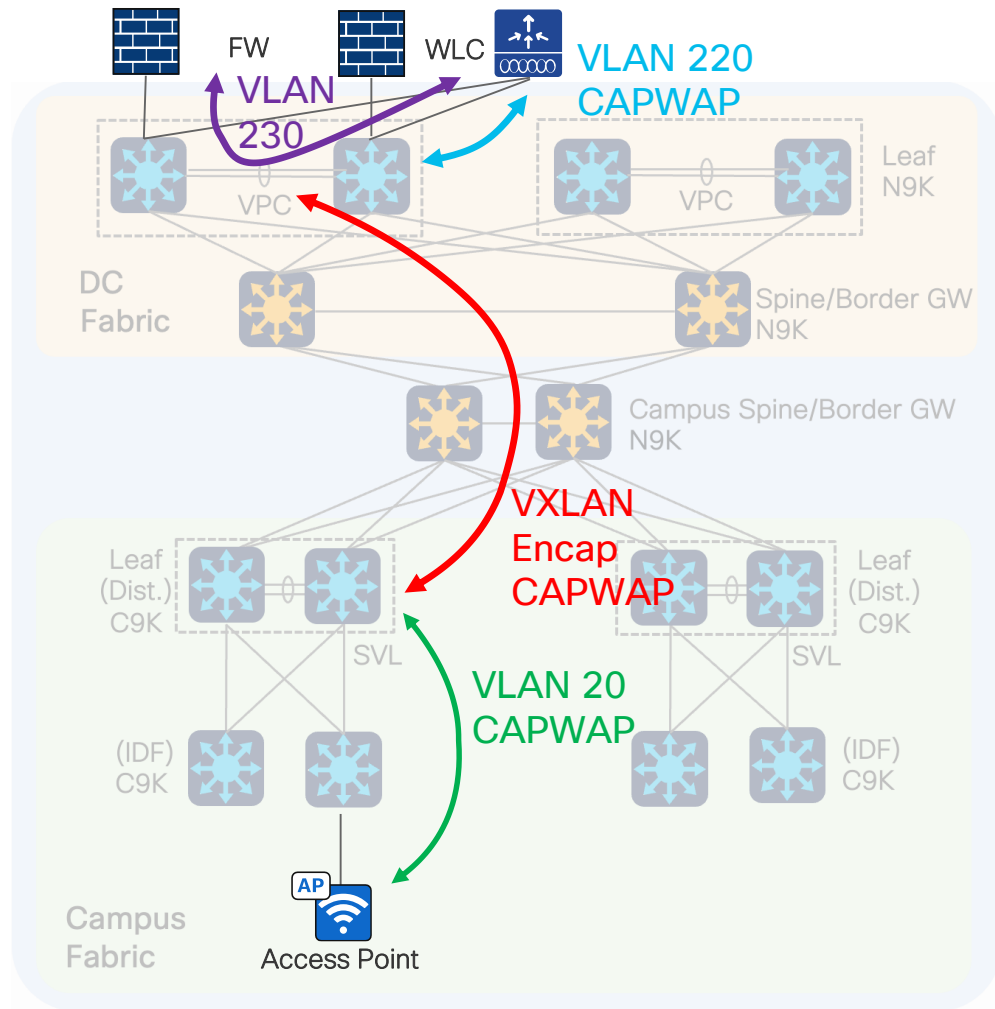


## Micro segmentation across campus and DC is not supported




# Wireless Inter-Site Traffic

- CAPWAP from Access Point reaches Leaf over VLAN 20
- CAPWAP encapsulated in VXLAN reaches DC Leaf
- CAPWAP traffic from service Leaf sent to WLC over VLAN 220 as CAPWAP
- Traffic sent as Ethernet from WLC to FW over VLAN 230

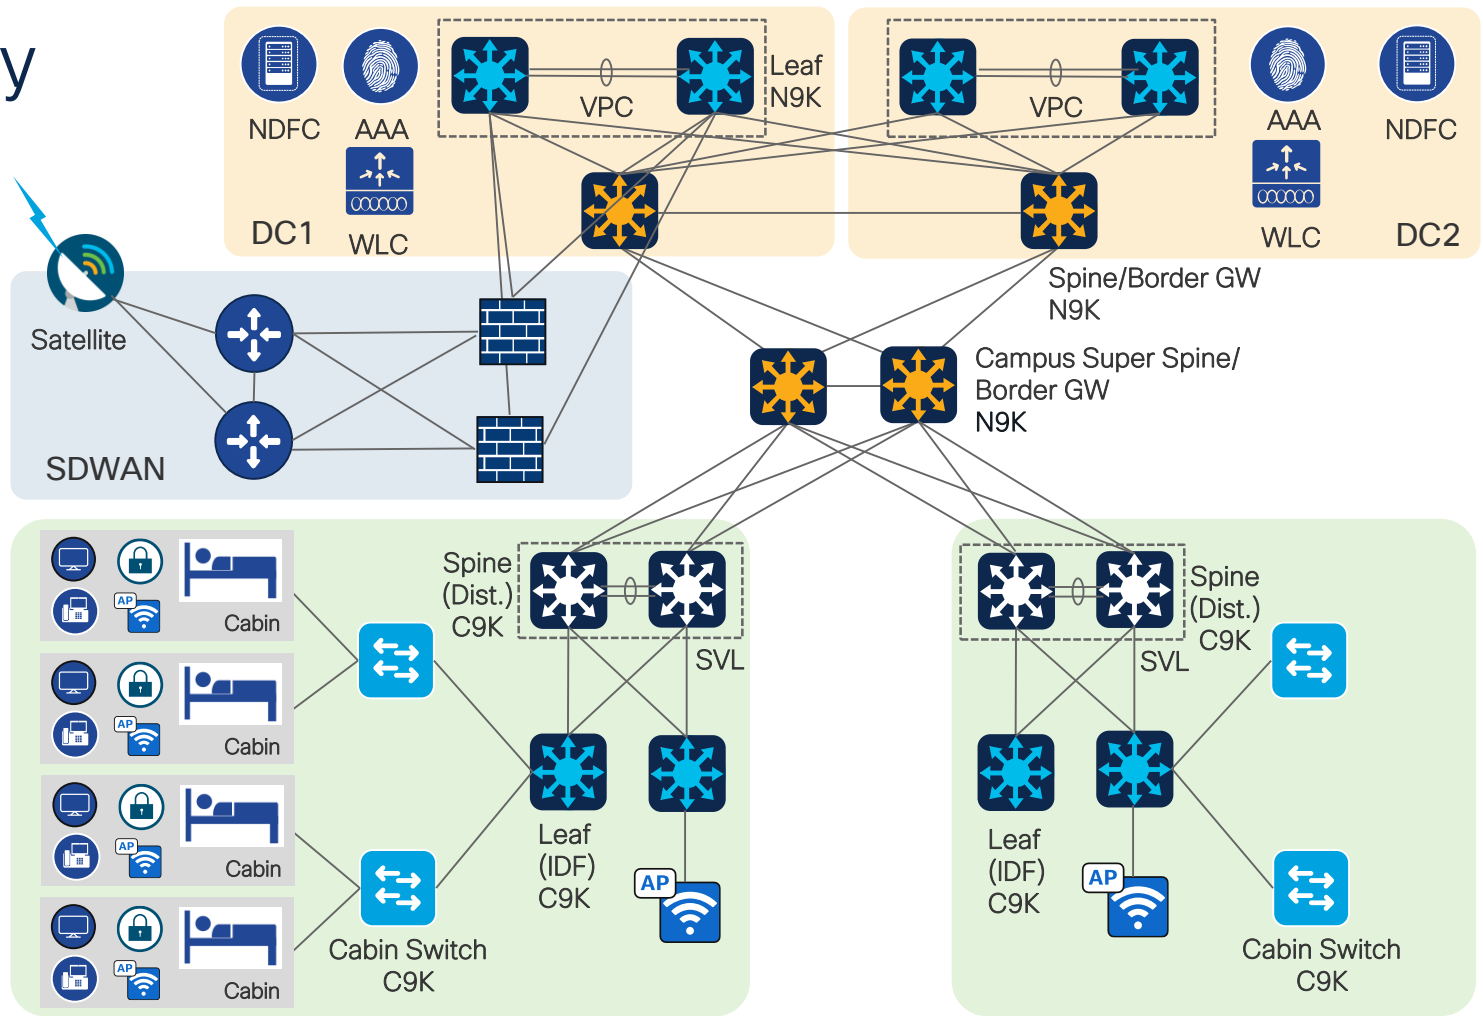


# L3 Inter-Site Traffic

- Anycast Gateway required for each subnet
- Anycast Gateway for server resides in DC
- Anycast Gateway for IDF resides in Campus
- Traffic routed between different VLANs/subnets with the same L3 VNI



# Topology





# L2 Connectivity to DC

## Challenge:

- Certain legacy applications require endpoints at the Campus to be on the same L2 segment as servers at DC

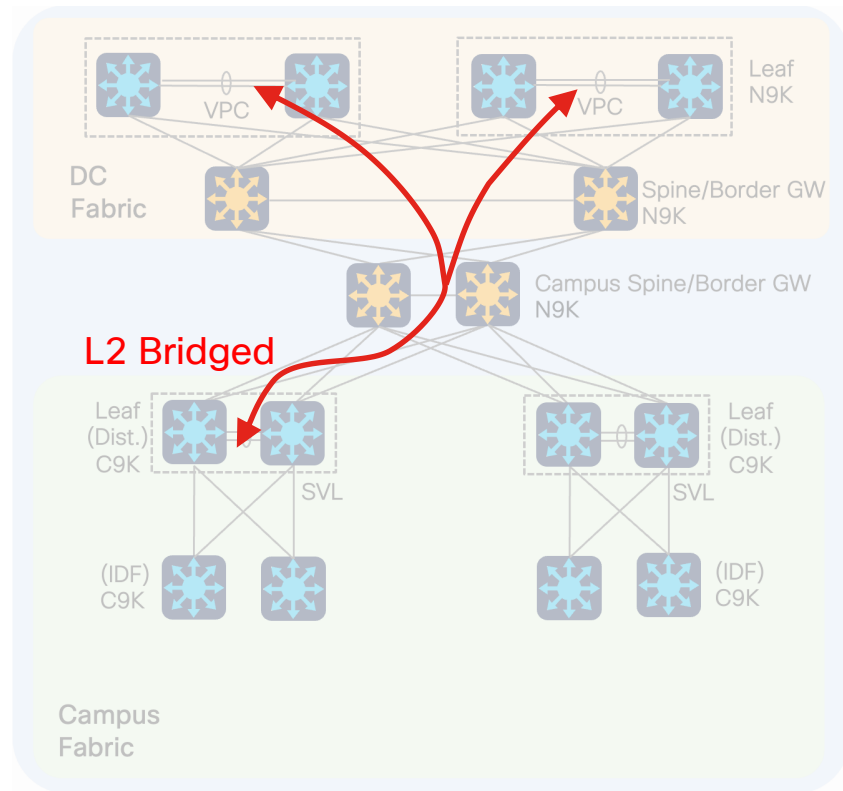


## Solution: Layer 2 Overlay Tunnels



### Typical Use Cases

- Manufacturing
- IoT



# Location Redundancy

## Challenge:


- DCs and critical network devices need to have location redundancy




**Solution:** Spread them across Fire Zones and locations

## Typical Use Cases

- Transportation
- Financial



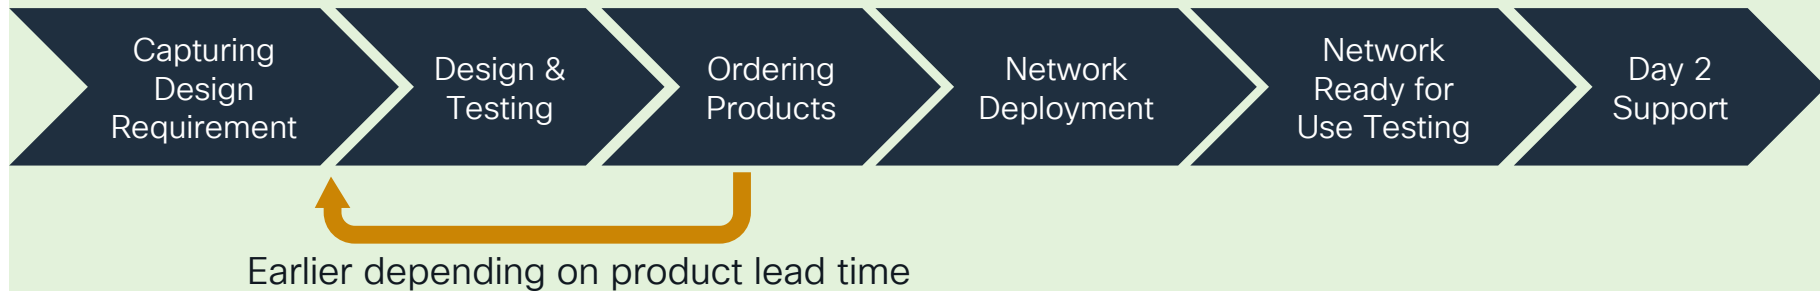
# Plan, Design, Implementation, and Support Best Practices



# Plan

## Project Planning

- Plan according to contractually binding ship sail date
- Set major milestones



- Engage all involved parties relevant to deployment phases
- Track progress and take actions as needed

# Plan

## Capture Design Requirement

- Current operation model
- Priorities and approaches to meet them
- Traffic Flow: Campus – DC – Internet
- Application requirement

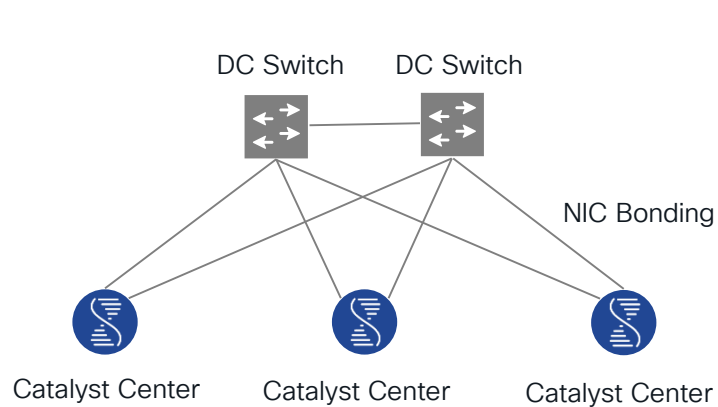
## Security Policies

- Define security policies based on business requirement
- Macro and micro segmentation
- Group Based Policies
- Endpoint authentication and authorization
- Further developed in Design and fine-tuned in lab and onsite


# Design

## Redundancy

- Link, controller and hardware redundancy
- Put critical roles on dedicated devices
- Routing if possible, bridging only if you have to



Catalyst Center 3-Node Cluster HA

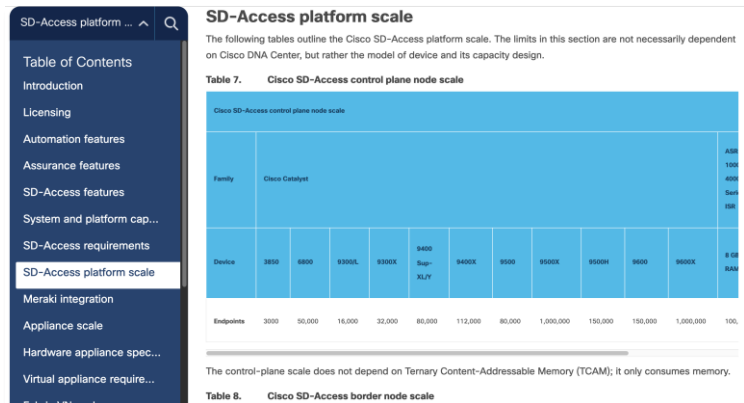


SDA L2 Border

# Design

## Sale & Feature Support

- Consider platform & controller scales
- Older hardware affects Fabric wide feature support
- Use recommended controller & device software versions
- Standardize network site hierarchy, scalable IP address schemes & device naming



## New Deployment

Release **2.3.5.5 (recommended)**

Device Role **Fabric Edge**

Submit Query


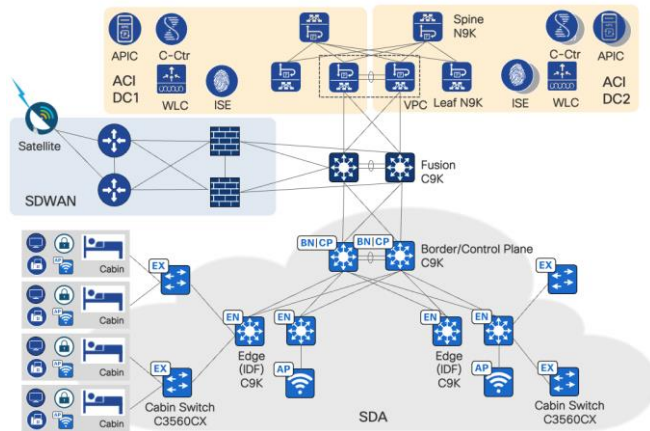
## SD-Access Compatibility Matrix for Cisco Catalyst Center 2.3.5.5 (recommended release)

Device Role	Device Series	Device Model	Recommended Release	Supported Release
Fabric Edge	Cisco Catalyst 9200 Series Switches	C9200-24T	IOS XE 17.9.4a	IOS XE 17.13.x
		C9200-24P		IOS XE 17.12.x
		C9200-24PXG		IOS XE 17.11.x

# Design

## Lab Environment

- Set up a scaled down proof of concept (POC) lab if feasible
- The lab assists in developing implementation procedures
- Software version and hardware certification
- Ongoing lab testing prior to implementing new features to production network






# Deployment

## Coordination

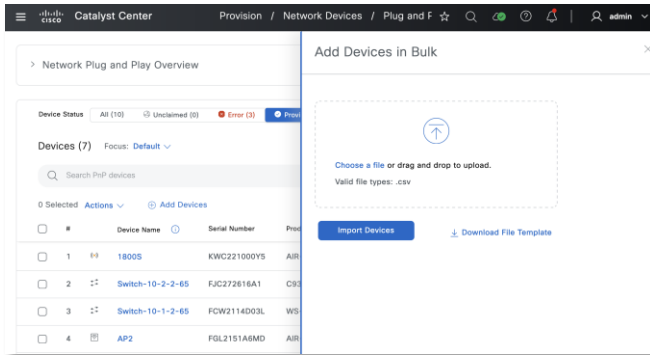
- Consider international shipyard construction constrains
- Address issues with shipyard construction timely:
  - Cabling
  - Unaccountable APs and Cabin switches
- Conduct post implementation WiFi RF survey before ships are in service



# Deployment

## SDA Automation

- 100s of APs and cabin switches can be connected by shipyard construction at different times
- Automate AP onboarding using a CSV template
- Large numbers of Extended Node onboarding can be automated using templates



Serial Number	Model	AP Name	Site	RF Profile
FJC253786GH	C9115AXE-B	AP-9-2-9016	Global/ADVENTURE/FABRIC_ZONE_03/DECK_09	PASSENGER
FJC253716MV	C9115AXE-B	AP-9-2-9018	Global/ADVENTURE/FABRIC_ZONE_03/DECK_09	CREW

# Extended Node Onboarding

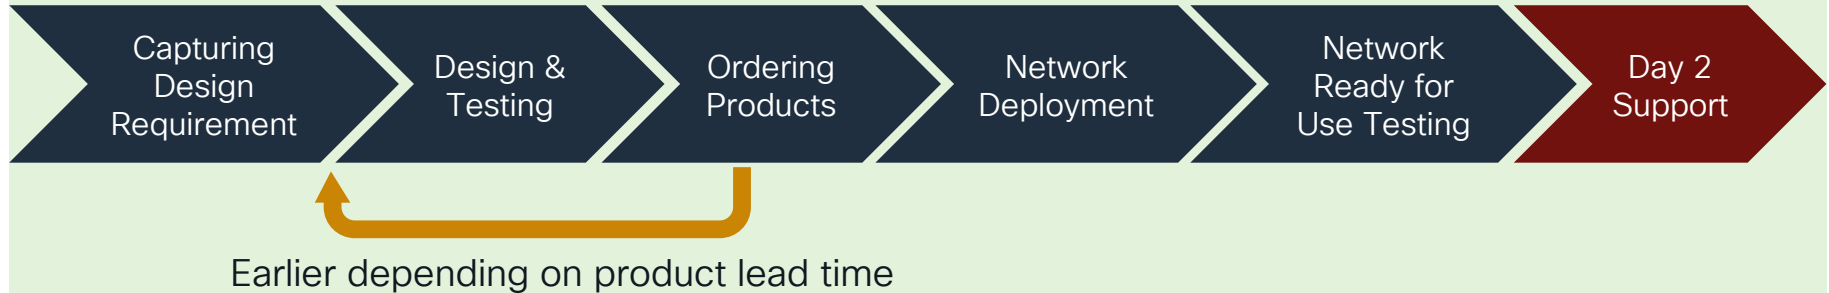
- Configure an EEM script in a template
- Push the template into upstream Edge switches by provisioning
- The Edge creates a port-channel and trunk when it sees a C3560CX as its CDP neighbor
- The C3560CX is onboarded as an Extended Node automatically
- Once all the C3560CX are onboarded, remove the EEM script from the Edge

```
event manager applet detect-3560CX authorization bypass
event neighbor-discovery interface regexp Ethernet.* cdp add
action 1.0 regexp "(C3560CX)" "$_nd_cdp_platform"
action 2.0 if $_regexp_result eq "1"
action 2.1 cli command "enable"
action 3.0 cli command "show running-config interface $_nd_local_intf_name"
action 4.0 regexp "(channel-group)" "$_cli_result"
action 5.0 if $_regexp_result eq "0"
action 6.0 cli command "enable"
action 7.0 cli command "config t"
action 8.0 cli command "default interface $_nd_local_intf_name"
action 8.1 cli command "int $_nd_local_intf_name"
action 8.3 cli command "switchport mode trunk"
action 8.4 syslog msg "3560CX detected, no port-channel present"
action 8.5 break
action 8.6 else
action 8.7 syslog msg "3560CX detected, port-channel present"
action 9.1 end
action 9.2 end
```


# Post Implementation Support

## Day 2 Support





- Design/Implementation team provides support immediately after ship launch
- Transfer network As-Built documents to Day-2 support
- Train Day-2 support team on features and technology
- Changes made to production should be reviewed, approved and implemented during change windows



# Sneak Preview: Catalyst Center Orchestrated BGP EVPN









# SDA-EVPN Supported Hardware

Catalyst Center		
Physical Appliance	Virtual Appliance	Cloud
DN2-HW-APL DN2-HW-APL-L DN2-HW-APL-XL 	DN2-SW-APL 	Amazon Web Services 
	Catalyst Center can support either SDA-LISP or SDA-EVPN. Co-existence of both control-planes on single Catalyst Center is not supported	


Catalyst Switches (IOS-XE 17.12.x and above)	
Core	Access
Catalyst 9500 Non-High-Performance Catalyst 9500 High-Performance Catalyst 9600 Sup-1	Catalyst 9300L/LM Catalyst 9300/9300B Catalyst 9300-X Catalyst 9400 Sup-1 Catalyst 9400-X Sup-2

# Fabric Roles and Access Modes




# Flexible Underlay Options

## LAN Automation



## Brownfield



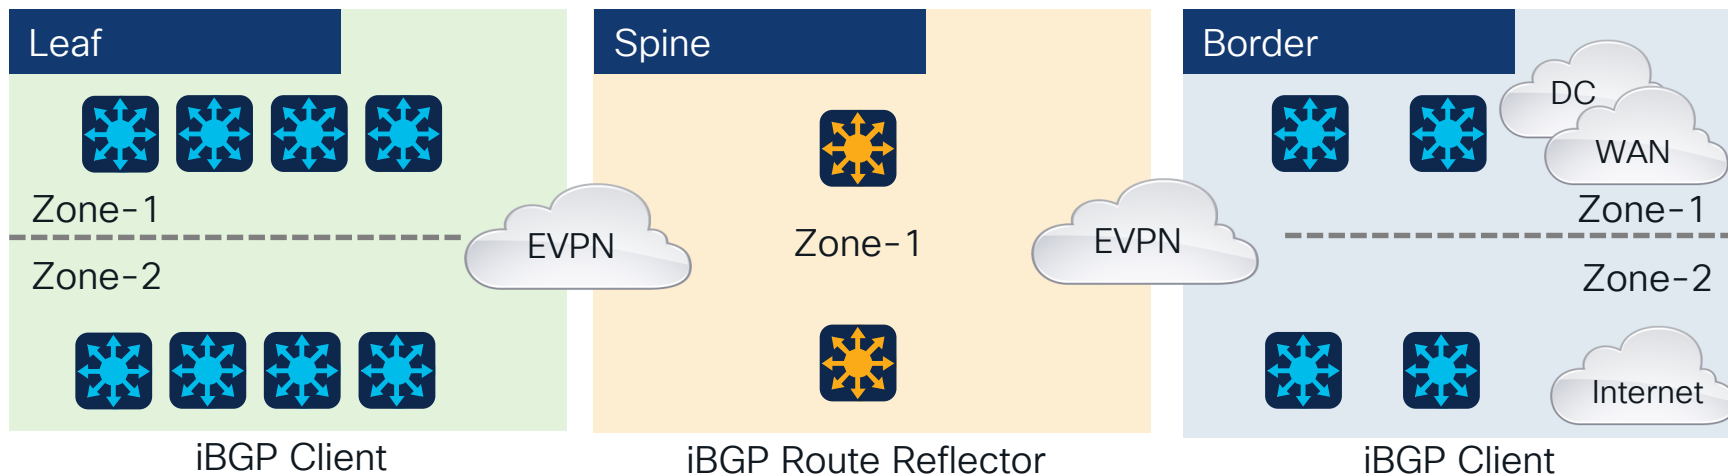
- Network virtualization does not mandate change to existing network
- Customers can build SDA-EVPN Fabric while managing brownfield underlay independently



SDA-EVPN does not support fabric automation on network switches with pre-configured BGP



# Control Plane & Fabric Zone



## Control Plane

- Underlay IGPs can be used to build iBGP sessions over loopbacks

## Fabric Zone

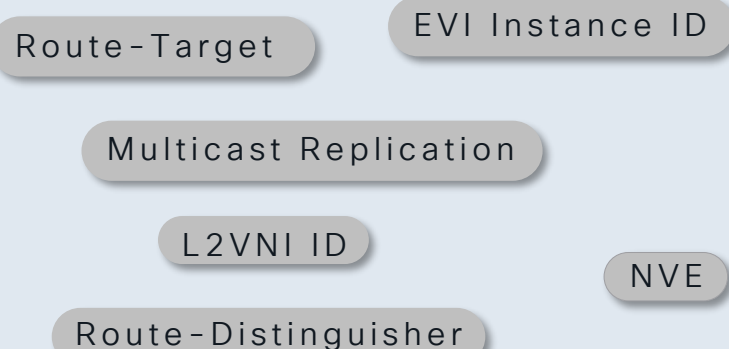
- A device can only be associated to single Fabric Zone
- A Fabric site can have multiple Leaf and Border zones but one Spine zone only

# Simplified Addressing

## L3VN Address Pool



## L2VN Address Pool

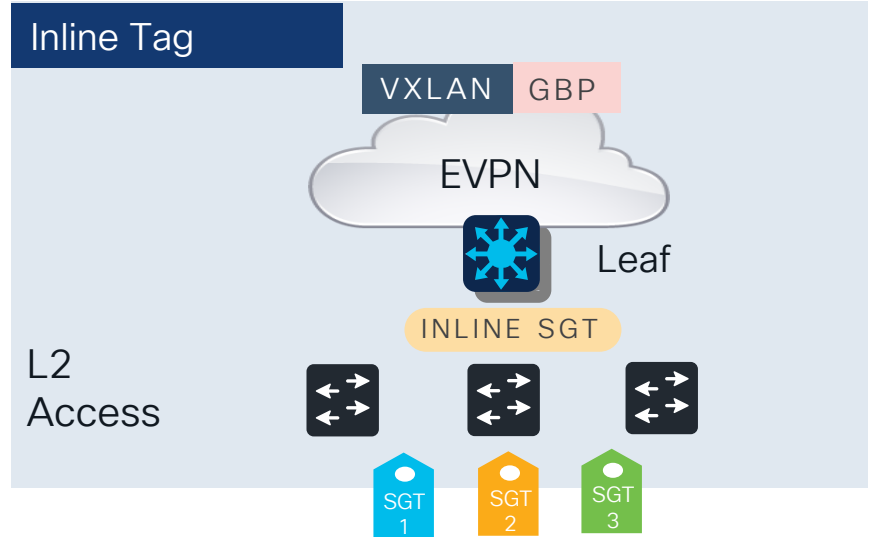



- Fabric Resource Pools auto-generated
- Automation workflows dynamically reserve & release addresses to pools
- Reduces operational complexity



Fabric Resource Pool is a one-time initial fabric site configuration step. It cannot be expanded or modified once fabric site automation is completed

# Context Aware Fabric – Micro Segmentation




- VXLAN with Group Based Policy (GBP) extension uses SGT for endpoint contexts
- VXLAN data-plane with GBP extents policy-plane across fabric

# Wireless Integration

## Centralized



## Distributed



- Centralized/Distributed wireless integrated with SDA-EVPN overlay
- AP to WLC Control-plane preserved in underlay



SDA-EVPN does not automate wireless network and switch ports connected to WLC & APs

# Cisco Campus BGP EVPN Solution Evolution

	Do It Yourself CLI	Programmable Ansible, Terraform	NDFC	SDA-EVPN Catalyst Center
Fabric Configuration	✓ CLI templates	✓ Automated	✓ Automated	✓ Intent-Based
Software Life Cycle Management		✓	✓	✓
Config Compliance			✓	✓
Assurance & Analytics (Endpoint, Policy & Trust)				✓
Zero-Trust Architecture				✓

Enterprise Ready Catalyst 9000 Foundation

# How Did the Tale of Two Ships End?



- Network implementation started from around 2022-23
- SD-Access LISP (SDA) & BGP EVPN deployed on two ships from two different cruise lines
- Around ten vessels have been deployed using these two solutions, with more to come
- Both solutions have been very stable

# Complete Your Session Evaluations



Complete a minimum of 4 session surveys and the Overall Event Survey to be entered in a drawing to **win 1 of 5 full conference passes** to Cisco Live 2025.

---



**Earn 100 points** per survey completed and compete on the Cisco Live Challenge leaderboard.

---



Level up and earn **exclusive prizes!**

---



Complete your surveys in the **Cisco Live mobile app**.

# Continue your education

- Visit the Cisco Showcase for related demos
- Book your one-on-one Meet the Engineer meeting
- Attend the interactive education with DevNet, Capture the Flag, and Walk-in Labs
- Visit the On-Demand Library for more sessions at [www.CiscoLive.com/on-demand](http://www.CiscoLive.com/on-demand)

Contact me at: [chyan@cisco.com](mailto:chyan@cisco.com),  
[ythallas@cisco.com@cisco.com](mailto:ythallas@cisco.com@cisco.com)





The bridge to possible

# Thank you

CISCO *Live!*

#CiscoLive