

CISCO *Live!*



#CiscoLive



The bridge to possible

SD Access: Troubleshooting the Fabric

Michel Peters
Technical Leader Engineering
BRKENT-3820



#CiscoLive

Cisco Webex App

Questions?

Use Cisco Webex App to chat with the speaker after the session

How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space

Webex spaces will be moderated by the speaker until June 17, 2022.



<https://ciscolive.ciscoevents.com/ciscolivebot/#BRKENT-3820>




Agenda

- Fabric overview
- Endpoints Registrations
- Reaching Remote Endpoints


Fabric overview



Overlay Technology

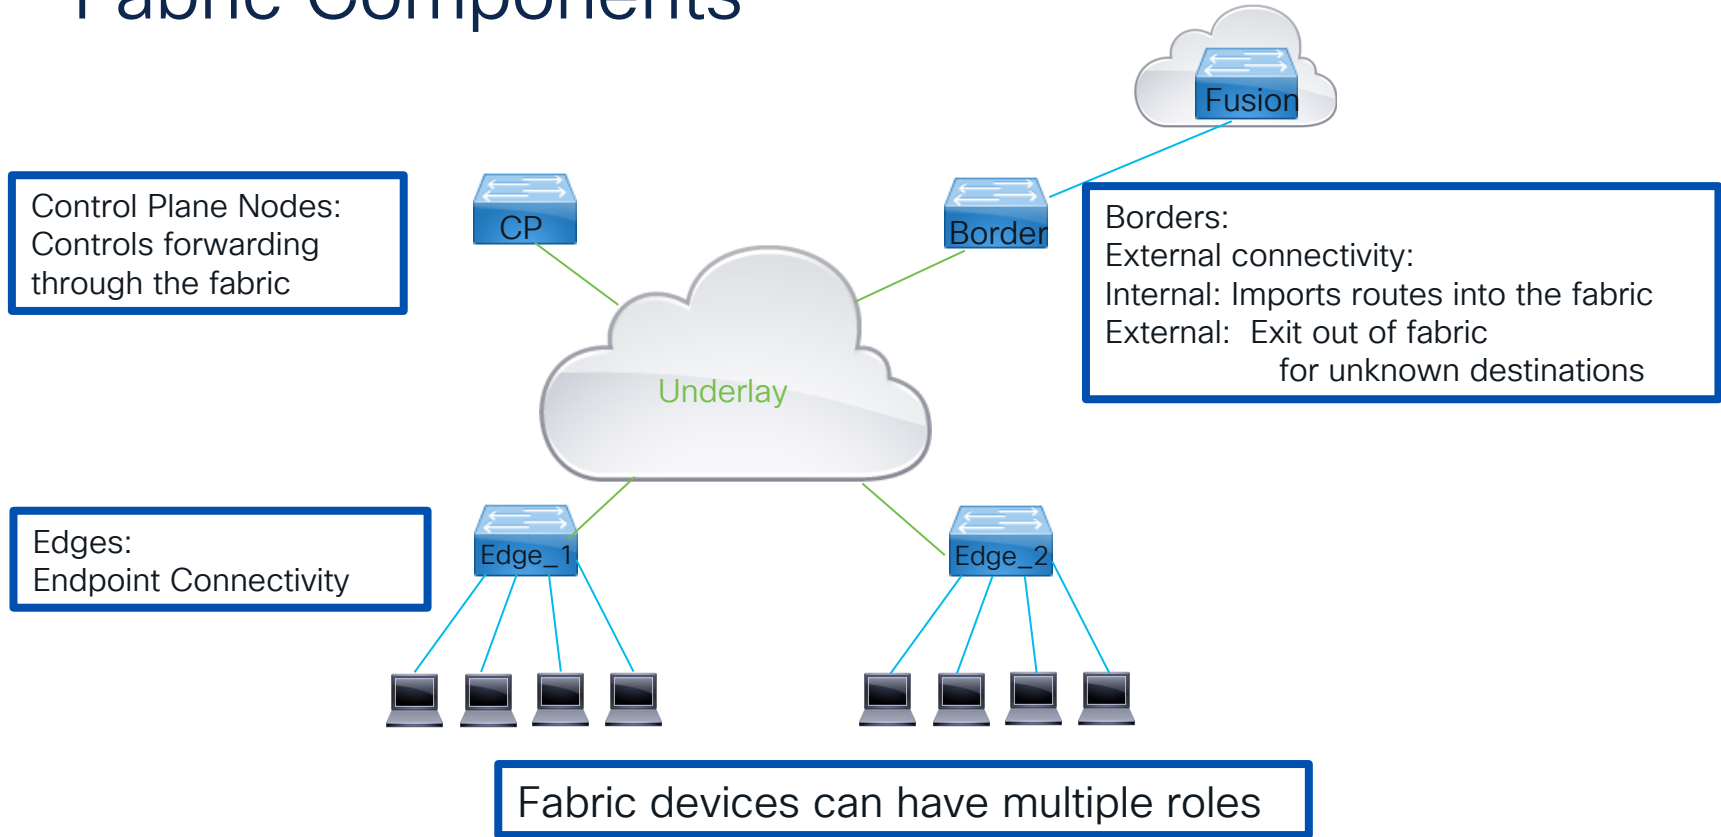


One Underlay
Layer 3 links



Multiple Virtual
Networks as Overlays

Fabric Components



SD Access Fabric Key Technologies

- Locator/ID Separation Protocol,
Control plane protocol used inside the fabric for the Overlay networks
- VXLAN,
Used for encapsulating all Dataplane traffic through the underlay to form the overlay networks
- Cisco TrustSec,
Assigning of a Policy label to all packets and enforcing security policies
- Authentication,
Assigns endpoints using Dot1x/MAB with their respective authorization profiles and associated pools

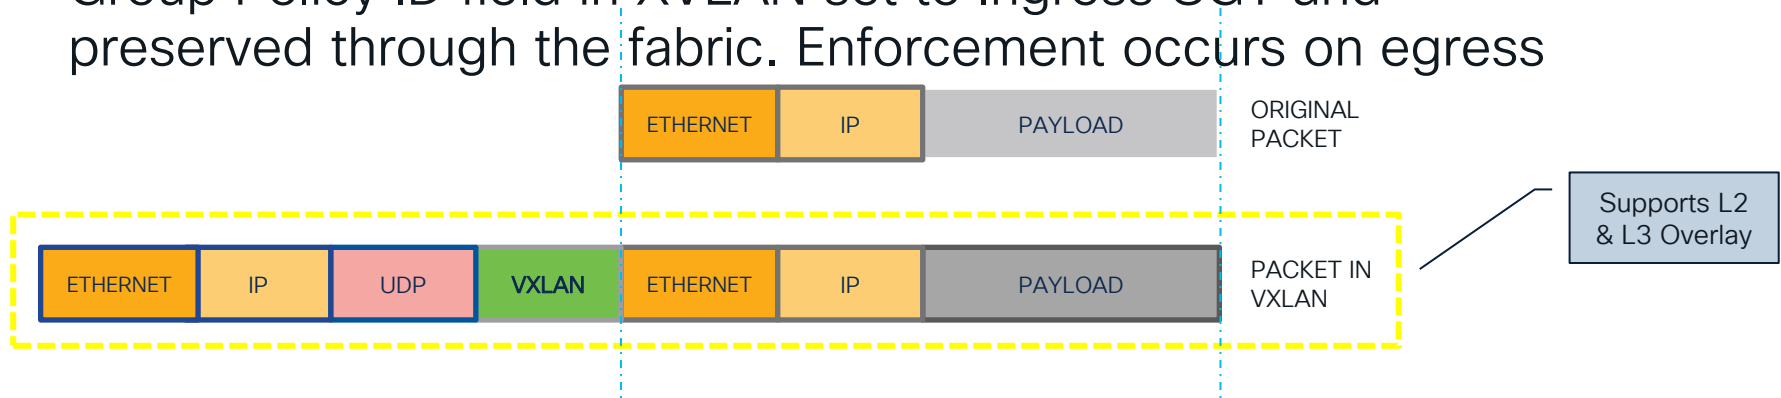
LISP Basic operation

- LISP is a routing architecture, not just a routing protocol
- LISP creates a level of indirection by using two spaces: “locators” (RLOC) and “endpoints” (EID)
- Advertise “locators” in core routing. Removes “hosts” from routing tables. Host prefixes moved to an alternative system database
- Routers in Underlay only need routing information to RLOC space, simplifies Underlay network
- To get path information to end hosts, routers query locator-end host map servers. Mapping analogous to DNS.
- Routers hold map-cache of locator-hosts.

LISP Device	SD Access	Function
RLOC (Routing Locator)	Fabric Devices	Routing Locator. Exists in global routing tables. Authoritative to reach EID space.
EID (Endpoint ID)	IP pools/End Points	Endpoint Identifier. IP addresses. Hidden from core network routing table. RLOC acts next-hop to reach EID space.
ETR (Egress Tunnel Router)& PETR (Proxy ETR)	Edge Device & Border node	Connects a LISP site to a LISP capable core network. Registers EID prefixes with Map Server (MS). Decapsulates LISP packets received from LISP core. PETR works on behalf of non-LISP domain and provides LISP-non-LISP connectivity.
ITR (Ingress Tunnel Router) & PITR (Proxy Ingress Tunnel Router)	Edge Device and Border node	Responsible for forwarding local traffic to external destinations. Resolves RLOC for a given destination by sending Map-request to Map Resolver. Encapsulates traffic and send to fabric. Typically, this is a Access Layer Switch. PITR works on behalf of non-LISP domain and provides LISP-non-LISP connectivity.
XTR (X Tunnel Router)	Edge Device	When both ITR and ETR functions are handled by one router, it is called XTR. This is typical in practice.
MR (Map Resolver)	Control Plane Node	Responds to Map-requests from ITR. Map-requests will be replied with a (Negative) Map-reply or forwarded to appropriate ETR
MS (Map Server)	Control Plane Node	Registers EID space upon receiving Map-register messages from ETR. Updates Map Resolver with EID and RLOC data.
MSMR (Map Server Map Resolver)	Control Plane Node	When a device acts as both Map Server and Map Resolver, it is called MS MR. This is typical in practice.

Data Plane

- In SD Access the entire Layer 2 packet is encapsulated
- VXLAN encapsulation used. IP address in new header set to RLOC
- VXLAN Network Identifier field set to LISP instance ID
- Group Policy ID field in VXLAN set to Ingress SGT and preserved through the fabric. Enforcement occurs on egress



Endpoints Registrations

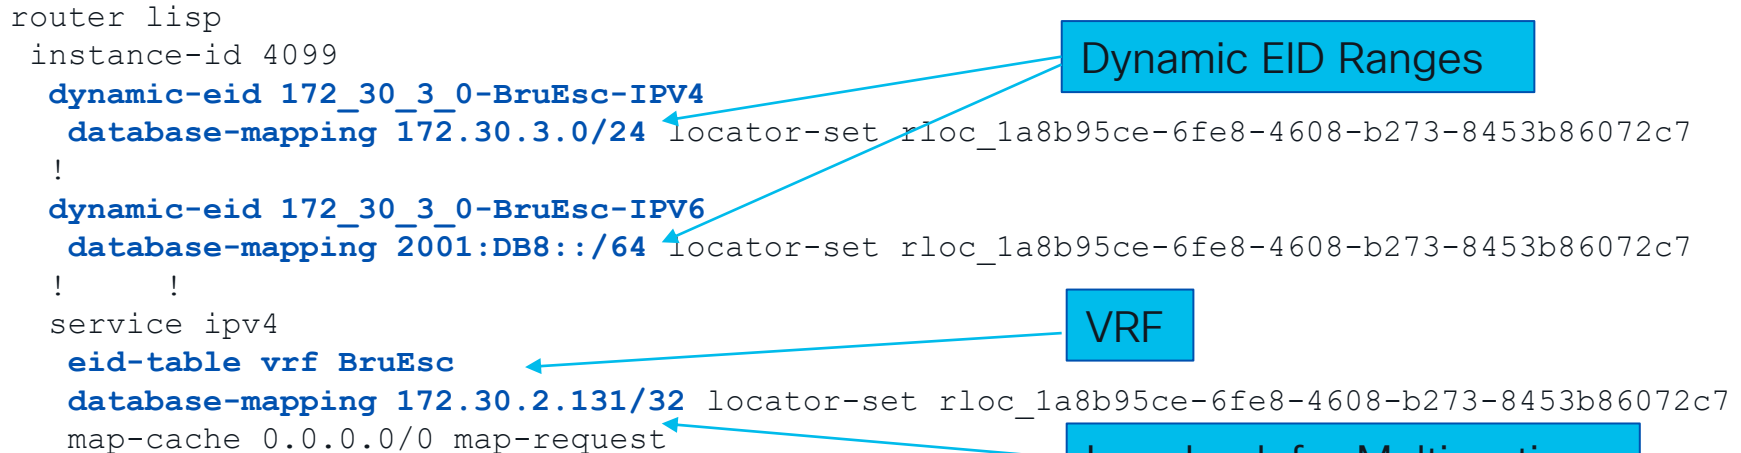


Virtual Networks and LISP instances

- VRF's correlate to Virtual Networks
- Dynamic EID ranges specify which EID are learned
- When Multicast enabled Loopback interfaces are created inside VRF

```
Edge_1#sh ip vrf BruEsc
Name                               Interfaces
BruEsc                             Lo4099
                                   V11021
                                   LI0.4099
                                   Tu2
                                   V11022
```

```
router lisp
instance-id 4099
dynamic-eid 172_30_3_0-BruEsc-IPV4
  database-mapping 172.30.3.0/24 locator-set rloc_1a8b95ce-6fe8-4608-b273-8453b86072c7
!
dynamic-eid 172_30_3_0-BruEsc-IPV6
  database-mapping 2001:DB8::/64 locator-set rloc_1a8b95ce-6fe8-4608-b273-8453b86072c7
!
service ipv4
  eid-table vrf BruEsc
  database-mapping 172.30.2.131/32 locator-set rloc_1a8b95ce-6fe8-4608-b273-8453b86072c7
  map-cache 0.0.0.0/0 map-request
```



Edge Configuration: SVI/VLAN Configuration

- Layer 3 Subnets and Layer 2 Pools deployed to all Edges is consistent throughout a fabric site
- Same IP address and mac are used on all edges(IP Anycast)
- Connections between edges should be L3 to avoid mac-learning issues


```
Edge_1#sh run int vlan 1021
interface Vlan1021
  mac-address 0000.0c9f.f377
  vrf forwarding BruEsc
  ip address 172.30.3.1 255.255.255.0
  ip helper-address 10.48.91.148
  no lisp mobility liveness test
  lisp mobility 172_30_3_0-BruEsc-IPV4
```

```
Edge_2#sh run int vlan 1021
interface Vlan1021
  mac-address 0000.0c9f.f377
  vrf forwarding BruEsc
  ip address 172.30.3.1 255.255.255.0
  ip helper-address 10.48.91.148
  no lisp mobility liveness test
  lisp mobility 172_30_3_0-BruEsc-IPV4
```

LISP operation, registering with Map Server

Instance	RLOC	EID (mac address)
8189	Edge_1	10f9.206d.e5b7
8189	Edge_2	10f9.206d.e5b6
4099	Edge_1	172.30.3.3/32
4099	Edge_2	172.30.3.2/32
4099	Border	10.48.91.128/25

- Fabric devices dynamically learn the IP and Mac addresses of attached devices to register with control plane node using map-register messages
- Layer 2 Instances start with 8xxx, 1 instance per Vlan
- 4xxx Instances used for Layer 3 Virtual Networks, 1 per Virtual Network, IPv4 and IPv6 share same instance
- Control Plane nodes maintain central database mapping



LISP Database

```
Edge_1#sh ip arp vrf BruEsc 172.30.3.3
Protocol  Address          Age (min)  Hardware Addr  Type   Interface
Internet  172.30.3.3        171       10f9.206d.e5b7  ARPA   Vlan1021
Edge_1#sh lisp instance-id 4099 ipv4 database 172.30.3.3/32
LISP ETR IPv4 Mapping Database for EID-table vrf BruEsc (IID 4099), LSBs: 0x1
172.30.3.3/32, dynamic-eid 172_30_3_0-BruEsc-IPV4,..
  Uptime: 3d15h, Last-change: 3d15h
  Locator          Pri/Wgt  Source      State
  172.30.233.6     10/10   cfg-intf    site-self, reachable
```

- LISP Database registers Learned Endpoints that are inside the EID ranges
- Endpoints can be learned via ARP/DHCP Snooping/Device Tracking
- Locator IP address is typically Loopback0 of Fabric Device in the Underlay network, needs to be reachable inside routing tables
- Wildcard (*) when used will show all instances with lisp commands

Registration of Endpoints with Map Server (CP)

- LISP Reliable Transport used with SDA. Using TCP in stead of UDP
- LISP Session Down can be due to failed communication or no EID's to be registered
- Registration only succeeds when LISP key matches with CP node
- Map register messages send to CP nodes to register all EID's

```
Edge_1#sh lisp session
```

Peer	State	Up/Down	In/Out	Users
172.31.255.182:4342	Up	00:00:25	54/22	12

```
Edge_1#sh tcp brief
```

TCB	Local Address	Foreign Address	(state)
7EFDC4E8BA90	172.30.233.6.43136	172.31.255.182.4342	ESTAB

```
Edge_1#sh lisp instance-id 4100 ipv4 statistics | sec Map-Register
```

```
Map-Register records in/out: 0/28
```

```
Map-Server AF disabled: 0
```

```
Authentication failures: 0
```

Layer 2 Endpoints

- Entries with Type CP_LEARN are remote entries
- Only endpoint macaddresses are registered with Control Plane

```
Edge_1#sh mac ad vlan 1021
```

Vlan	Mac Address	Type	Ports
----	-----	-----	-----
1021	0000.0c9f.f377	STATIC	Vl1021
1021	10f9.206d.e5b7	STATIC	Te1/0/11
1021	701f.539b.0a75	STATIC	Vl1021
1021	10f9.206d.e5b6	CP_LEARN	L2LI0

Total Mac Addresses installed by LISP: REMOTE: 1

```
Edge_1#sh lisp instance-id 8189 ethernet database
```

```
LISP ETR MAC Mapping Database for EID-table Vlan 1021 (IID 8189), LSBs: 0x1
0000.0c9f.f377/48, dynamic-eid Auto-L2-group-8189, do not register, inherited from
default locator-set rloc_1a8b95
  Uptime: 3d23h, Last-change: 3d23h
    Locator      Pri/Wgt  Source      State
    172.30.233.6  10/10   cfg-intf    site-self, reachable
10f9.206d.e5b7/48, dynamic-eid Auto-L2-group-8189, inherited from default locator-set
rloc_1a8b95
  Uptime: 3d23h, Last-change: 3d23h
    Locator      Pri/Wgt  Source      State
    172.30.233.6  10/10   cfg-intf    site-self, reachable
```

Control Plane Node (MSMR)

- Control Plane Node acts as both Map Server and Map resolver (MSMR)
- Keeps database of all EID registrations for all AF(Ethernet/IPv4/IPV6)
- No synchronization between Control Plane nodes
- Show lisp site command gives overview of all IPv4/IPv6 registrations

```
Border_CP_1#sh lisp site instance-id 4099
```

```
LISP Site Registration Information
```

Site Name	Last Register	Up	Who Last Registered	Inst ID	EID Prefix
site_uci	never	no	--	4099	0.0.0.0/0
	never	no	--	4099	172.30.2.128/25
	05:17:04	yes#	172.30.233.6:43136	4099	172.30.2.131/32
	00:00:07	yes#	172.30.233.1:4342	4099	172.30.2.132/32
	never	no	--	4099	172.30.3.0/24
	00:00:07	yes#	172.30.233.1:4342	4099	172.30.3.2/32
	05:17:04	yes#	172.30.233.6:43136	4099	172.30.3.3/32
	never	no	--	4099	172.30.4.0/24

Control Plane Node (MSMR) details on EID

```
Border_CP_1#sh lisp site 172.30.3.2/32 instance-id 4099
```

```
EID-prefix: 172.30.3.2/32 instance-id 4099
```

```
First registered: 4d23h
```

```
Last registered: 00:00:01
```

```
Origin: Dynamic, more specific of 172.30.3.0/24
```

```
Proxy reply: Yes
```

```
TTL: 1d00h
```

```
State: complete
```

```
Extranet IID: Unspecified
```

```
Registration errors:
```

```
Authentication failures: 0
```

```
Allowed locators mismatch: 0
```

```
ETR 172.30.233.1, last registered 00:00:01, proxy-reply, map-notify
```

```
TTL 1d00h, no merge, hash-function sha1, nonce 0x768..
```

```
state complete, no security-capability
```

```
xTR-ID 0x41DCA445-0xF8480845-0x4E7EB2E4-0xFA8E33CF
```

```
site-ID unspecified
```

Locator	Local	State	Pri/Wgt	Scope
172.30.233.1	yes	up	10/10	IPv4 none

Age of EID

With proxy set Control plane responds on behalf of XTR

ETR Information

RLOC Information

Layer 2 Control Plane

```
Border_CP_1#sh lisp instance-id 8189 ethernet server
```

```
LISP Site Registration Information
```

Site Name	Last Register	Up	Who	Last Registered	Inst ID	EID Prefix
site_uci	never	no	--		8189	any-mac
	03:57:06	yes#	172.30.233.1	51300	8189	10f9.206d.e5b6/48
	10:12:16	yes#	172.30.233.6	43136	8189	10f9.206d.e5b7/48

- Mac registrations on CP shown using show lisp instance-id * ethernet server

```
Border_CP_1#sh lisp inst 8189 ethernet server 10f9.206d.e5b6 registration-history
```

```
Roam = Did host move to a new location?
```

```
WLC = Did registration come from a Wireless Controller?
```

```
Prefix qualifier: + = Register Event, - = Deregister Event, * = AR register event
```

Timestamp (UTC)	Instance	Proto	Roam	WLC	Source	EID prefix
Jun 6 02:51:41.699	8189	TCP	No	No	172.30.233.1	+ 10f9.206d.e5b6/48
Jun 6 03:51:49.913	8189	TCP	No	No	172.30.233.1	- 10f9.206d.e5b6/48
Jun 6 03:52:06.392	8189	TCP	No	No	172.30.233.1	+ 10f9.206d.e5b6/48


Reaching Remote Endpoints



LISP basic operation, resolving

Instance	RLOC	EID (mac address)
8189	Edge_1	10f9.206d.e5b7
8189	Edge_2	10f9.206d.e5b6
4099	Edge_1	172.30.3.3/32
4099	Edge_2	172.30.3.2/32
4099	Border	10.48.91.128/25

- Endpoint 1 sends packet towards Endpoint 2
- Edge_1 initiates map request to CP node
- CP responds to Edge_2 with map-response containing RLOC information
- Edge_1 creates map-cache entry and is ready to forward traffic



Map Cache

- Map-requests triggered by hitting an Entry with send-map-request action
map-cache 0.0.0.0/0 map-request
- Responses from Control Plane Nodes are cached on fabric devices to build the map cache.
- Successful map-requests are cached with a default TTL of 1 day
Time to Live set by registering device with “etr map-cache-ttl”
- Negative map-requests have TTL of 15 minutes
- Control plane node returns largest possible block containing requested EID when sending Negative Map Reply

Resolving Remote Destinations

Triggers map-request

```
Edge_1#sh lisp instance-id 4099 ipv4 map-cache
LISP IPv4 Mapping Cache for EID-table vrf BruEsc (IID 4099), 7 entries
0.0.0.0/0, uptime: 5d05h, expires: never, via static-send-map-request
    Encapsulating to proxy ETR
0.0.0.0/1, uptime: 11:28:43, expires: 00:10:14, via map-reply, forward-native
    Encapsulating to proxy ETR
172.30.2.129/32, uptime: 11:30:36, expires: 00:29:39, via map-reply, complete
    Locator      Uptime      State  Pri/Wgt      Encap-IID
    172.31.255.182 11:30:36  up      10/10        -
172.30.3.0/24, uptime: 5d05h, expires: never, via dynamic-EID, send-map-request
    Encapsulating to proxy ETR
172.30.3.2/32, uptime: 00:16:31, expires: 23:43:28, via map-reply, complete
    Locator      Uptime      State  Pri/Wgt      Encap-IID
    172.30.233.1  00:16:31  up      10/10        -
172.30.4.0/24, uptime: 5d05h, expires: never, via dynamic-EID, send-map-request
    Encapsulating to proxy ETR
```

NMR, send to petr


Encapsulate to RLOC

Map Cache shows EID range, source of cache entry and action to be taken.

LISP basic operation, packet forwarding

Instance	RLOC	EID (mac address)
8189	Edge_1	10f9.206d.e5b7
8189	Edge_2	10f9.206d.e5b6
4099	Edge_1	172.30.3.3/32
4099	Edge_2	172.30.3.2/32
4099	Border	10.48.91.128/25

- Overlay traffic in SD Access is encapsulated in Vxlan and send between RLOC addresses
- Underlay network unaware of overlay topology
- Reachability to RLOC should exist in Route table
 - ipv4 locator reachability minimum-mask-length 32
 - ipv4 locator reachability exclude-default



Layer 2 or Layer 3 forwarding

- SDA supports both layer 2 and Layer 3 forwarding inside fabric
- Traffic inside IP pool will be encapsulated using Layer 2 instance
- Traffic destined outside IP pool send using Layer 3 instance
- Layer 2 forwards traffic based on Destination Mac Address and L2 Map-cache
- Optional flooding of BUM traffic using Multicast group in underlay
- Traffic outside IP pool will be routed using an Layer 3 LISP Instance

LISP Remote forwarding

- Show ip route does not show full detail on forwarding
- Default route and remote entries not showing on edges. As Null0 routes on Border

```
Edge_1#sh ip route vrf BruEsc
```

```
..
```

```
Gateway of last resort is not set
```

```
172.30.0.0/16 is variably subnetted, 7 subnets, 2 masks
```

```
C 172.30.2.131/32 is directly connected, Loopback4099
```

```
C 172.30.3.0/24 is directly connected, Vlan1021
```

```
L 172.30.3.1/32 is directly connected, Vlan1021
```

```
1 172.30.3.3/32 [10/1] via 172.30.3.3, 4d07h, Vlan1021
```

```
Border_CP_1#sh ip route vrf BruEsc
```

```
..
```

```
Gateway of last resort is not set
```

```
172.30.3.0/24 [200/0], 6w4d, Null0
```

```
C 172.30.3.1/32 is directly connected, Loopback1021
```

```
1 172.30.3.2/32 [250/1], 07:20:46, Null0
```

```
1 172.30.3.3/32 [250/1], 13:35:56, Null0
```

LISP Remote forwarding, more detail

```
Edge_1#sh ip cef vrf BruEsc 172.30.3.2 detail
172.30.3.2/32, epoch 1, flags [subtree context, check lisp eligibility]
  SC owned,sourced: LISP remote EID - locator status bits 0x00000001
  LISP remote EID: 2 packets 1152 bytes fwd action encap, dynamic EID need encap
  SC inherited: LISP cfg dyn-EID - LISP configured dynamic-EID
  LISP EID attributes: localeID No, c-dynEID Yes, d-dynEID No
  SC inherited: LISP generalised SMR - [enabled, inheriting, 0x7EFDC4E7A0A8 locks: 4]
  LISP source path list
    nexthop 172.30.233.1 LISP0.4099
  2 IPL sources [no flags]
  nexthop 172.30.233.1 LISP0.4099
```

- CEF gives accurate view of forwarding inside fabric device
- LISP subinterface is Instance-id , nexthop IP Address is RLOC of destination
- Show ip cef <nexthop> gives egress interface information in underlay for next hop.

Packet Encapsulation

Apply a display filter ... <%%/>

No.	Protocol	Source	Destination	Time	Info
→ 3	ICMP	172.30.3.2	172.30.3.3	0.116267	Echo (ping) request id=0x069b, seq=9688/55333, ttl=64 (reply in 4)
← 4	ICMP	172.30.3.3	172.30.3.2	0.116365	Echo (ping) reply id=0x069b, seq=9688/55333, ttl=64 (request in 3)
5	ICMP	172.30.3.3	172.30.2.2	1.023982	Echo (ping) request id=0x0659, seq=97/24832, ttl=63 (reply in 6)
6	ICMP	172.30.2.2	172.30.3.3	1.024255	Echo (ping) reply id=0x0659, seq=97/24832, ttl=252 (request in 5)
7	ICMP	172.30.3.2	172.30.3.3	1.140294	Echo (ping) request id=0x069b, seq=9689/55589, ttl=64 (reply in 8)
8	ICMP	172.30.3.3	172.30.3.2	1.140385	Echo (ping) reply id=0x069b, seq=9689/55589, ttl=64 (request in 7)
9	ICMP	172.30.3.3	172.30.2.2	2.047999	Echo (ping) request id=0x0659, seq=98/25088, ttl=63 (reply in 10)
10	ICMP	172.30.2.2	172.30.3.3	2.048247	Echo (ping) reply id=0x0659, seq=98/25088, ttl=252 (request in 9)
11	ICMP	172.30.3.2	172.30.3.3	2.164316	Echo (ping) request id=0x069b, seq=9690/55845, ttl=64 (reply in 12)
12	ICMP	172.30.3.3	172.30.3.2	2.164408	Echo (ping) reply id=0x069b, seq=9690/55845, ttl=64 (request in 11)

Frame 3: 148 bytes on wire (1184 bits), 148 bytes captured (1184 bits) on interface 0
Ethernet II, Src: Cisco_9b:0b:40 (70:1f:53:9b:0b:40), Dst: Cisco_1c:49:d8 (2c:5a:0f:1c:49:d8)
Internet Protocol Version 4, Src: 172.30.233.1, Dst: 172.30.233.6
User Datagram Protocol, Src Port: 65472, Dst Port: 4789

New Header

Virtual eXtensible Local Area Network
Flags: 0x8800, GBP Extension, VXLAN Network ID (VNI)
Group Policy ID: 4
VXLAN Network Identifier (VNI): 8189
Reserved: 0

SGT

VXLAN Header

LISP Instance ID

Ethernet II, Src: 10:f9:20:6d:e5:b6 (10:f9:20:6d:e5:b6), Dst: 10:f9:20:6d:e5:b7 (10:f9:20:6d:e5:b7)
Destination: 10:f9:20:6d:e5:b7 (10:f9:20:6d:e5:b7)
Source: 10:f9:20:6d:e5:b6 (10:f9:20:6d:e5:b6)
Type: IPv4 (0x0800)
Internet Protocol Version 4, Src: 172.30.3.2, Dst: 172.30.3.3
Internet Control Message Protocol

Encapsulated packet

Questions?



Related sessions:

- BRKTRS-3010 : SD Access: Advanced Fabric Troubleshooting
- BRKTRS-3090 : Troubleshooting the Cisco Catalyst 9000 Series Switches
- BRKTRS-2811a & BRKTRS-2811b : Overview of Packet Capturing Tools in Cisco Switches and Routers
- LABTRS-2391 Packet Capturing Tools in Enterprise Switching Environments

Cisco Networking Bot

<https://cnb.cisco.com/>

Empowering User by Digitizing Cisco Product & Adoption Experiences



EoS/EoL



Product Migration/Adoption



Config / Tshoot Guides



Security Advisories



Release Recommendations



HW-SW Compatibility

[Transforming Customer Experience with Cisco AI Chatbots](#)

Got a Question
on Enterprise Products?

Chat with CN Bot now



Technical Session Surveys

- Attendees who fill out a minimum of four session surveys and the overall event survey will get Cisco Live branded socks!
- Attendees will also earn 100 points in the Cisco Live Game for every survey completed.
- These points help you get on the leaderboard and increase your chances of winning daily and grand prizes.



Cisco Learning and Certifications

From technology training and team development to Cisco certifications and learning plans, let us help you empower your business and career. www.cisco.com/go/certs

Pay for Learning with Cisco Learning Credits

(CLCs) are prepaid training vouchers redeemed directly with Cisco.



Learn

Cisco U.

IT learning hub that guides teams and learners toward their goals

Cisco Digital Learning

Subscription-based product, technology, and certification training

Cisco Modeling Labs

Network simulation platform for design, testing, and troubleshooting

Cisco Learning Network

Resource community portal for certifications and learning



Train

Cisco Training Bootcamps

Intensive team & individual automation and technology training programs

Cisco Learning Partner Program

Authorized training partners supporting Cisco technology and career certifications

Cisco Instructor-led and Virtual Instructor-led training

Accelerated curriculum of product, technology, and certification courses



Certify

Cisco Certifications and Specialist Certifications

Award-winning certification program empowers students and IT Professionals to advance their technical careers

Cisco Guided Study Groups

180-day certification prep program with learning and support

Cisco Continuing Education Program

Recertification training options for Cisco certified individuals

Here at the event? Visit us at **The Learning and Certifications lounge at the World of Solutions**



Continue your education

- Visit the Cisco Showcase for related demos
- Book your one-on-one Meet the Engineer meeting
- Attend the interactive education with DevNet, Capture the Flag, and Walk-in Labs
- Visit the On-Demand Library for more sessions at www.CiscoLive.com/on-demand



The bridge to possible

Thank you

CISCO *Live!*



#CiscoLive