



The bridge to possible

CiscoLive: Brought to You By Cisco NSO and Cisco Modeling Labs

Joe Clarke, Distinguished Engineer

Cisco Webex App

Questions?

Use Cisco Webex App to chat with the speaker after the session

How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space

Webex spaces will be moderated until February 24, 2023.



Abstract

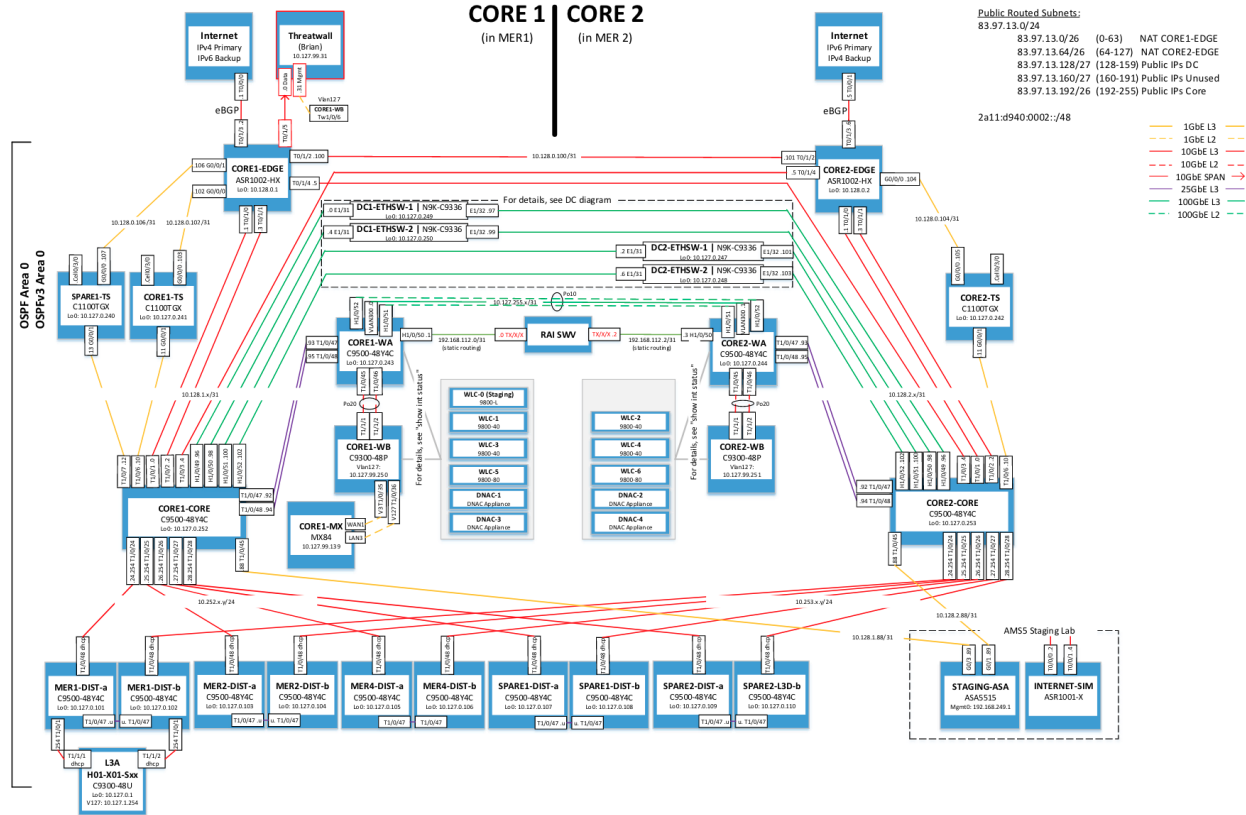
In order to deliver a world-class event, we need a world-class network. The network that runs the event is staged ahead of time, but only in two weeks. It has to provide robust, stable services for thousands of attendees, and it only has a small window to make the best possible impression. Oh, and like all dynamic, agile environments, requests for changes come in beyond the last minute. So how do you build a network that you know will fit the bill? This session will deep-dive into the solution used to design, build, and test the data centre network here at CiscoLive. This DC infrastructure is driven by Cisco NSO service definitions and modeled ahead of time using Cisco Modeling Labs. Attendees will see how it was built and why.

What We've Been Looking For



We're Getting There...

The Cisco Live Europe Network



CISCO *Live!*



So Why Are You Here?

While this session is focused on what was done for the CiscoLive data centre, I'll wager...



[Insert Critical Service Here]



Agenda


- Introduction
 - Our Problem
 - Cisco NSO
 - Cisco Modeling Labs
- Building the Data Centre Model
- Testing Things Out
- Looking Forward to 2024

Our Problem



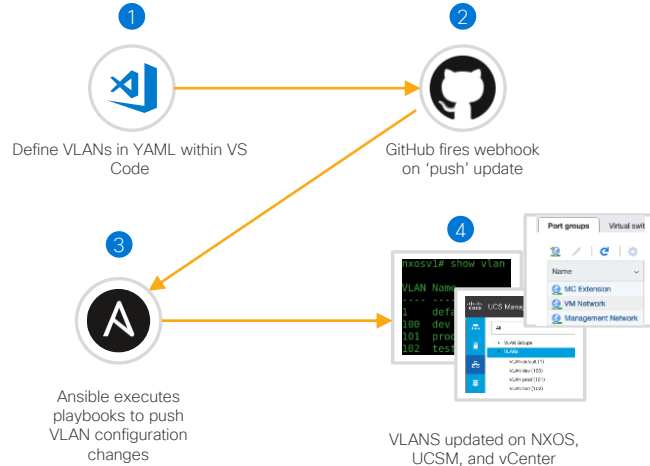
Event (All?) Networks Are Dynamic

No plan survives
contact with
production...



Solve It With Automation

Approach From 2020



```
2 vlans:
3   - id: 3
4     name: PUBLIC_INTERNET
5     vm_name: Public Internet
6     is_stretched: true
7     standard_svi_v4: true
8     standard_svi_v6: true
9   - id: 100
10     name: CROSS_DC_VMs
11     vm_name: CROSS DC VMs
12     is_stretched: true
13     standard_svi_v4: true
14     standard_svi_v6: true
```



Code available from
<https://marcuscom.com/git/MarcusCom/ciscolive/src/master/automation/cleu-ansible-n9k>

Drawbacks With This First Approach


- While it worked well, it wasn't so intuitive for network engineers, even through the CLI

```
➔ ./add_vlan.py
usage: ./add_vlan.py [-h] --vlan-name <VLAN_NAME> --vlan-id <VLAN_ID> [--vm-vlan-name <VM_VLAN_NAME>] [--svi-v4-network <SVI_NETWORK>]
                    [--svi-subnet-len <SVI_PREFIX_LEN>] [--svi-standard-v4] [--svi-v6-network <SVI_NETWORK>] [--svi-standard-v6]
                    [--svi-descr <SVI_DESCRIPTION>] [--no-add-acl] [--mtu <MTU>] [--is-stretched] [--no-hsrp]
                    [--no-passive-interface] [--v6-link-local] [--ospf-broadcast] [--interface <INTF>] [--generate-iflist]
                    [--vmware-cluster <CLUSTER>] --username <USERNAME> [--limit <HOSTS_OR_GROUP_NAMES>] [--tags <TAG_LIST>]
                    [--list-tags] [--test-only]
./add_vlan.py: error: the following arguments are required: --vlan-name/-n, --vlan-id/-i, --username/-u
```

- Required additional work for each dynamic feature that was needed (i.e., we needed playbooks and scripts for static routes, switchports, credentials, etc.)
- Lacked state for easy iteration, backouts, and re-deployments

Something New For 2021...er, 2022...er, 2023, yeah!

- The time away provided ample opportunity to think through something different
- Wanted to **iterate** and **test** as I went
- Began work in late 2021, settling on NSO for service development and a common interface to the data centre with CML as my initial testbed
- Is it perfect? No, but don't let perfect be the enemy of good.



Cisco Network Services Orchestrator (NSO)



NSO will help a lot now, and even more later



Today:
Network CLI

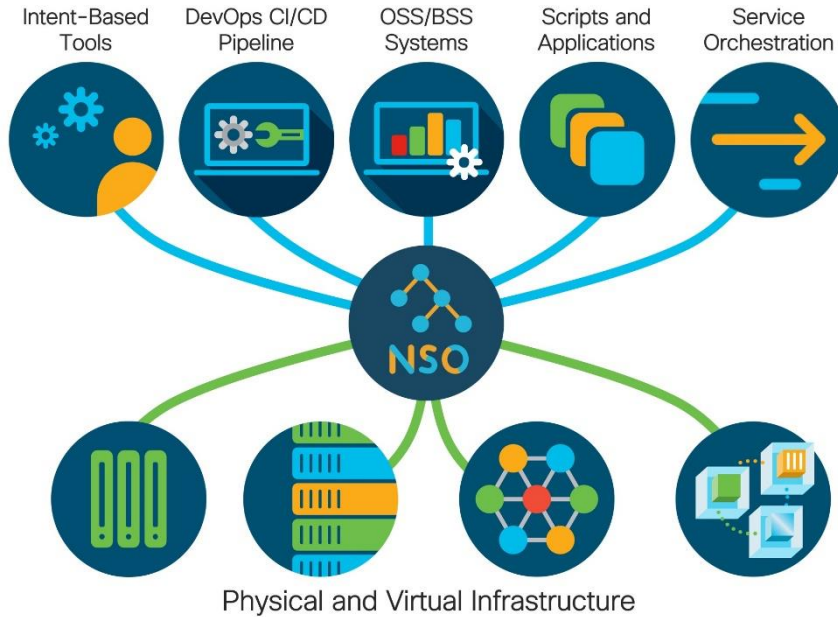
- Immediate benefits with familiar approach
- 2 fully functional CLI options
- Strict separation of operational and config data
- Perform operations on groups of devices
- Full AAA integration
- Database-style two-phase commit on changes



Tomorrow:
Network API

- Adopt over time as comfort and skills grow
- Python is a good place to start
- Multiple language bindings available
- Eventually integrate your own toolchain

Cisco NSO is a bridge

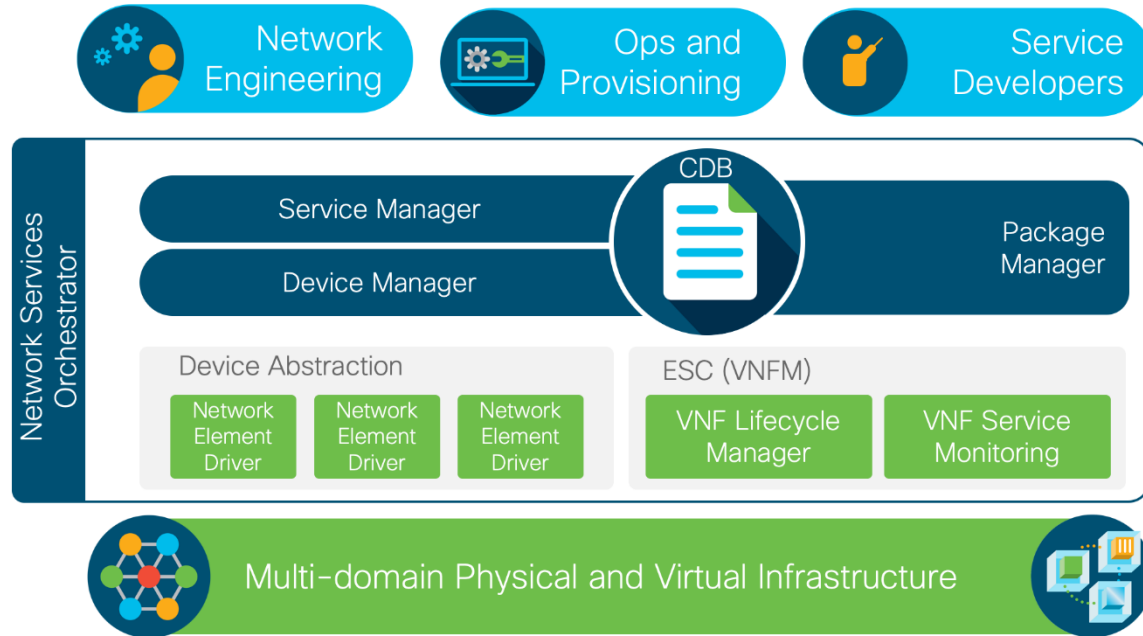


Between people that build services and ones that operate infrastructure

Across different **domains** and vendors


Over both **physical** and **virtual** infrastructure

NSO Architecture



- Multivendor abstraction through Network Element Drivers (NEDs)
- Single datastore for all network elements under management
- Multiple interfaces including CLI, REST, Java Python
- Templates and compliance reporting
- YANG-based configuration schema

A Few Words About YANG



```
list interface {  
    key "name";  
    unique "type location";  
  
    leaf name {  
        type string;  
        reference  
            "RFC 2863: The Interfaces Group MIB - ifName";  
    }  
  
    leaf description {  
        type string;  
    }  
  
    ...  
    container statistics {  
        config false;  
        leaf discontinuity-time {  
            type yang:date-and-time;  
        }  
  
        leaf in-octets {  
            type yang:counter64;  
            reference  
                "RFC 2863: The Interfaces Group MIB - ifHCInOctets";  
        }  
    }  
    ...  
}
```

Data modeling language built for network configuration

Human readable

Hierarchical configuration

Extensible through augmentation

Reusable types and groupings

Modular and expressive

Standards-based (defined in RFC7950)

Getting NSO Files

- Available FREE on DevNet for non-production use
 - <https://developer.cisco.com/docs/nso/#!getting-nso>
- What you'll get
 - NSO installation file for Intel Mac or Linux
 - Latest NEDs for Cisco platforms (IOS, XR, NX-OS, ASA)

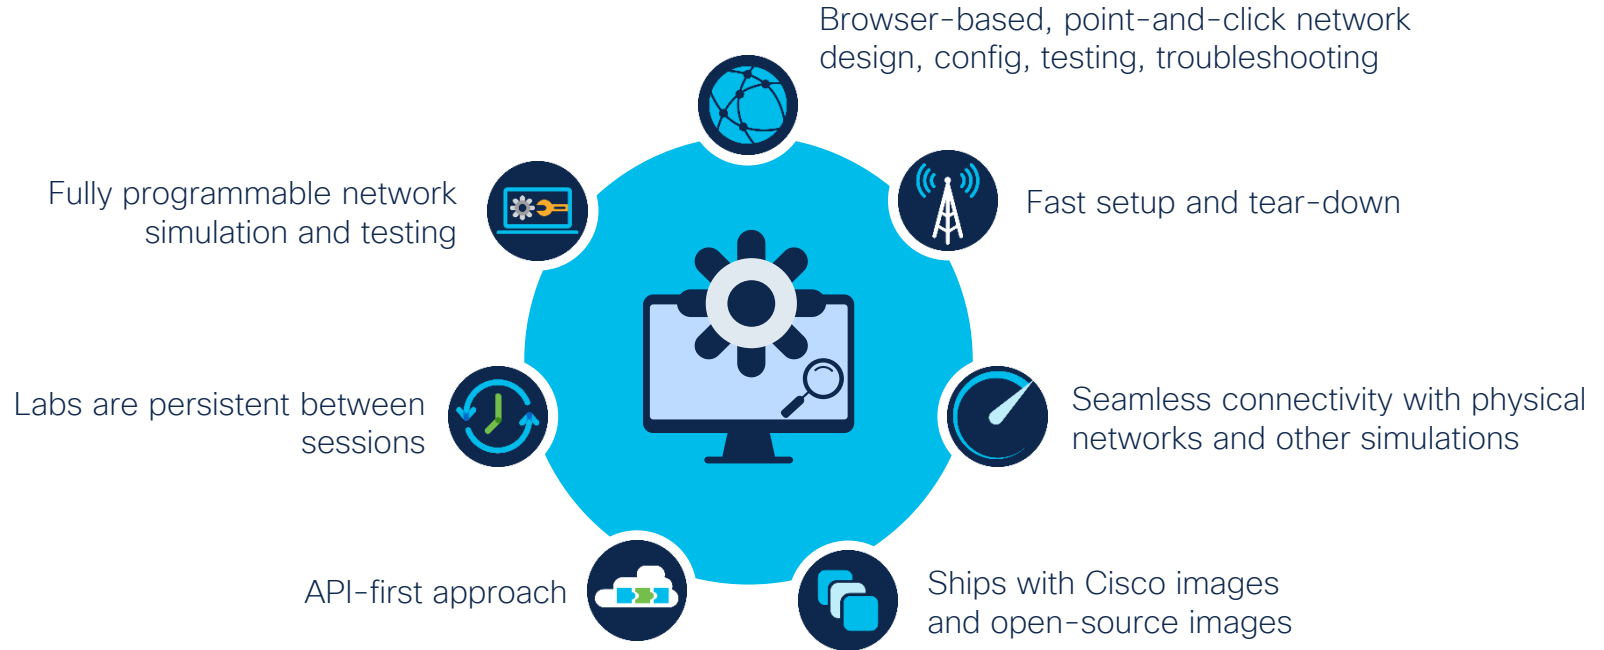


Cisco Modeling Labs

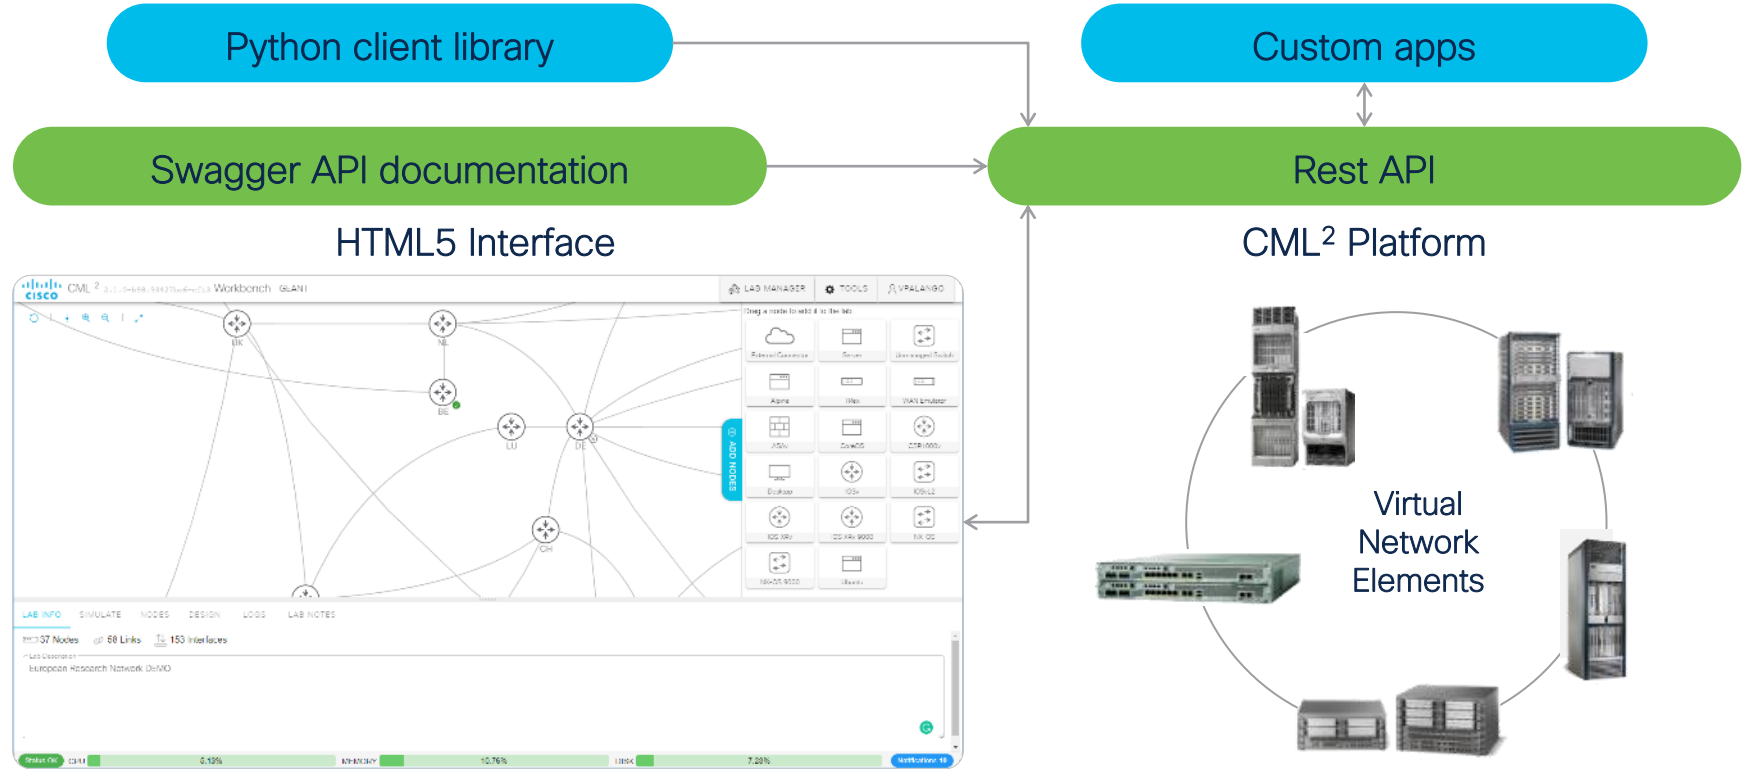


Cisco Modeling Labs

Your lab anywhere



Architecture Overview



Key Features



Labs = Isolated virtual networks

Configurable network connectivity

Console multiplexing

Flexible console access



Live topology modification –
drag, drop, and wire networks in a running
simulation

Device persistence –
just like shutting down a real router

Admin control of all labs

Ability to define and import
third-party device types that support the
Linux
KVM hypervisor



Server-side labs
with an easy-to-use
HTML5 front end


{ REST:API }

API First Design

Swagger interface

Python client library

Getting the CML Bits

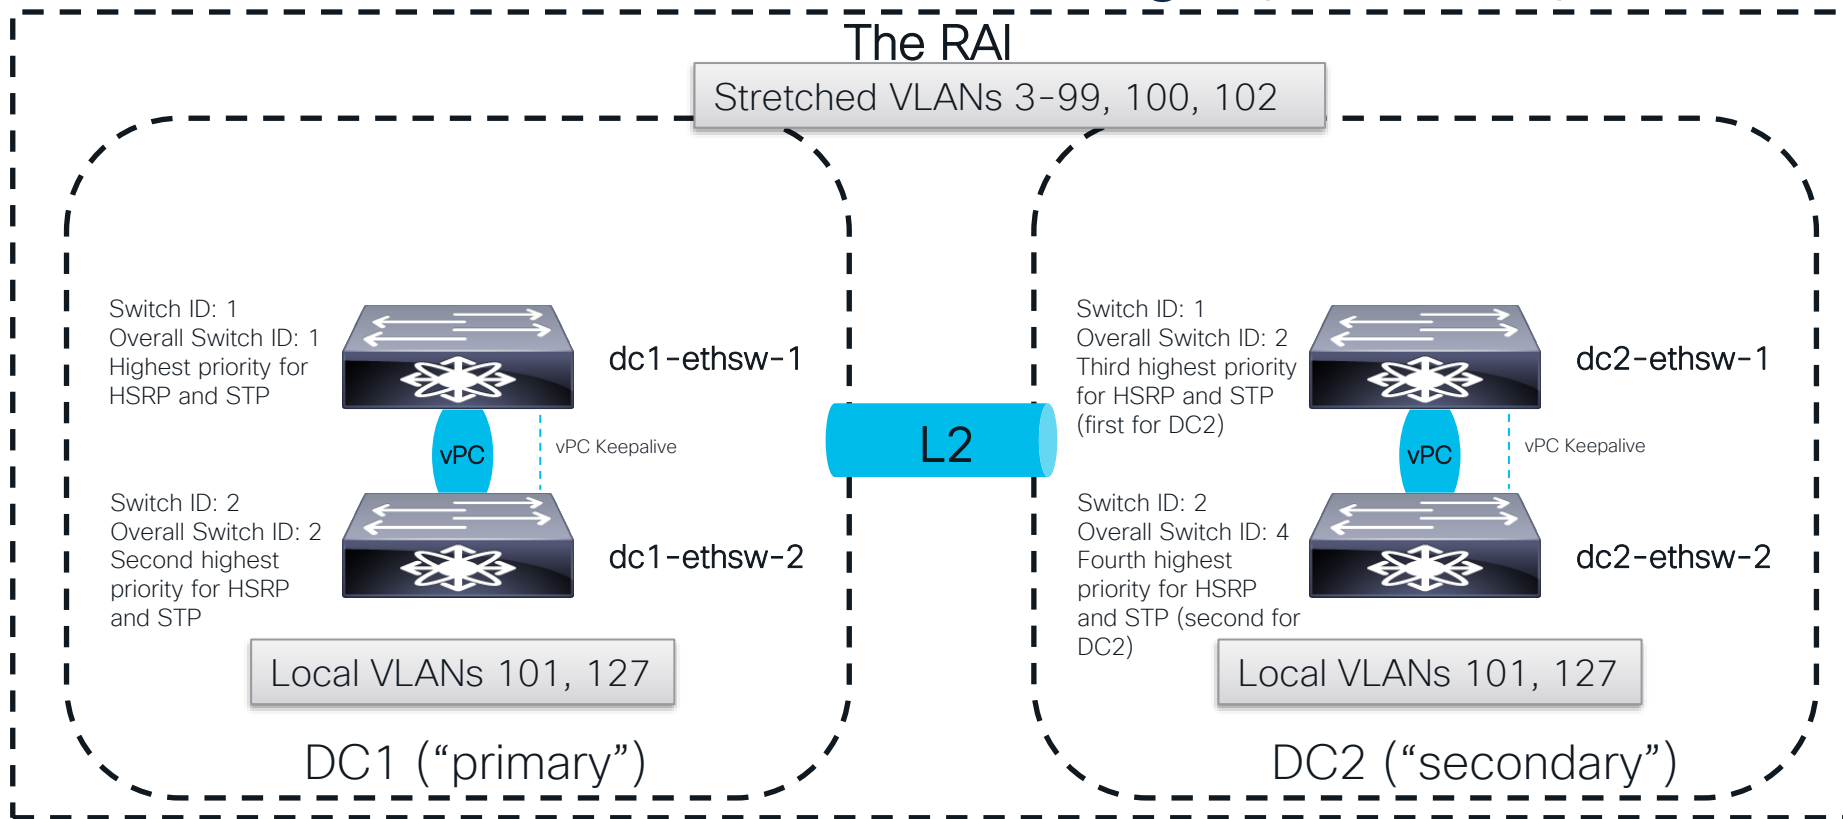


- Obtaining:
<https://developer.cisco.com/docs/modeling-labs/#!downloading-files-for-cml-installation>
- Docs:
<https://developer.cisco.com/docs/modeling-labs/#!introduction>
- Sandbox:
<https://devnetsandbox.cisco.com/RM/Topology> (Reserve “Cisco Modeling Labs Enterprise”)

Building the Data Centre Model



Some More About Our DC Design (Network)



Some More About Our DC Design (Compute)

The RAI

Stretched VLANs 3-99, 100, 102

Fabric Interconnects
for both UCS-Mini
and HyperFlex

vNIC Templates in
use

ESXi/vCenter with
distributed virtual
switches in use:

- dVS for DC-specific VLANs
- dVS for stretched VLANs



dc1-ucsm



ESXi clusters

Local VLANs 101, 127

DC1 (“primary”)

Fabric Interconnects
for both UCS-Mini
and HyperFlex

vNIC Templates in
use

ESXi/vCenter with
distributed virtual
switches in use:

- dVS for DC-specific VLANs
- dVS for stretched VLANs



dc2-ucsm



ESXi clusters

Local VLANs 101, 127

DC2 (“secondary”)

Choosing the Service Design Approach

Top-Down



Bottom-Up

Ultimately, this is choosing the user experience, too.

Top-Down Approach

- Good for green-field services
- One thinks of the overall model and flow first
- From that the templates (i.e., southbound* is developed)

* Southbound: meaning towards the end devices




Bottom-Up Approach

- Work from the templates or the CLI and derive the model from that
- Effective when you have a configuration or service that you want to automate



CiscoLive Met in the Middle



- Had an idea of what I wanted to be for the experience
 - Knew which parameters I wanted to be variable and which I wanted to be static
 - Allowed for some new service elements (e.g., HSRP security)
-
- Used existing configs from previous years to inform the model
 - Provided checks and balances to ensure nothing was left out
 - Able to identify technical debt





NSO Components Used

- Cisco NSO 5.6.0 ← Started here
- Ultimately landed on **NSO 6.0.0** for CLEU 2023
- NEDs
 - Cisco-NX 5.23.7 : For N9K switches
 - Cisco-UCS 3.4.4 : For UCS Manager
 - VMware-vSphere 3.3.5 : For vCenter
- **NOTE:** Initially lacked some IPv6 features (OSPFv3 and neighbor discovery config features); thanks to NSO TAC/dev for quick turn-around 🍷

Getting Started

- The service *package* for CiscoLive uses multiple templates (think interface config, VLAN config, router config, etc.) so Python is required

```
ncs-make-package --service-skeleton python-and-template  
ciscolive
```

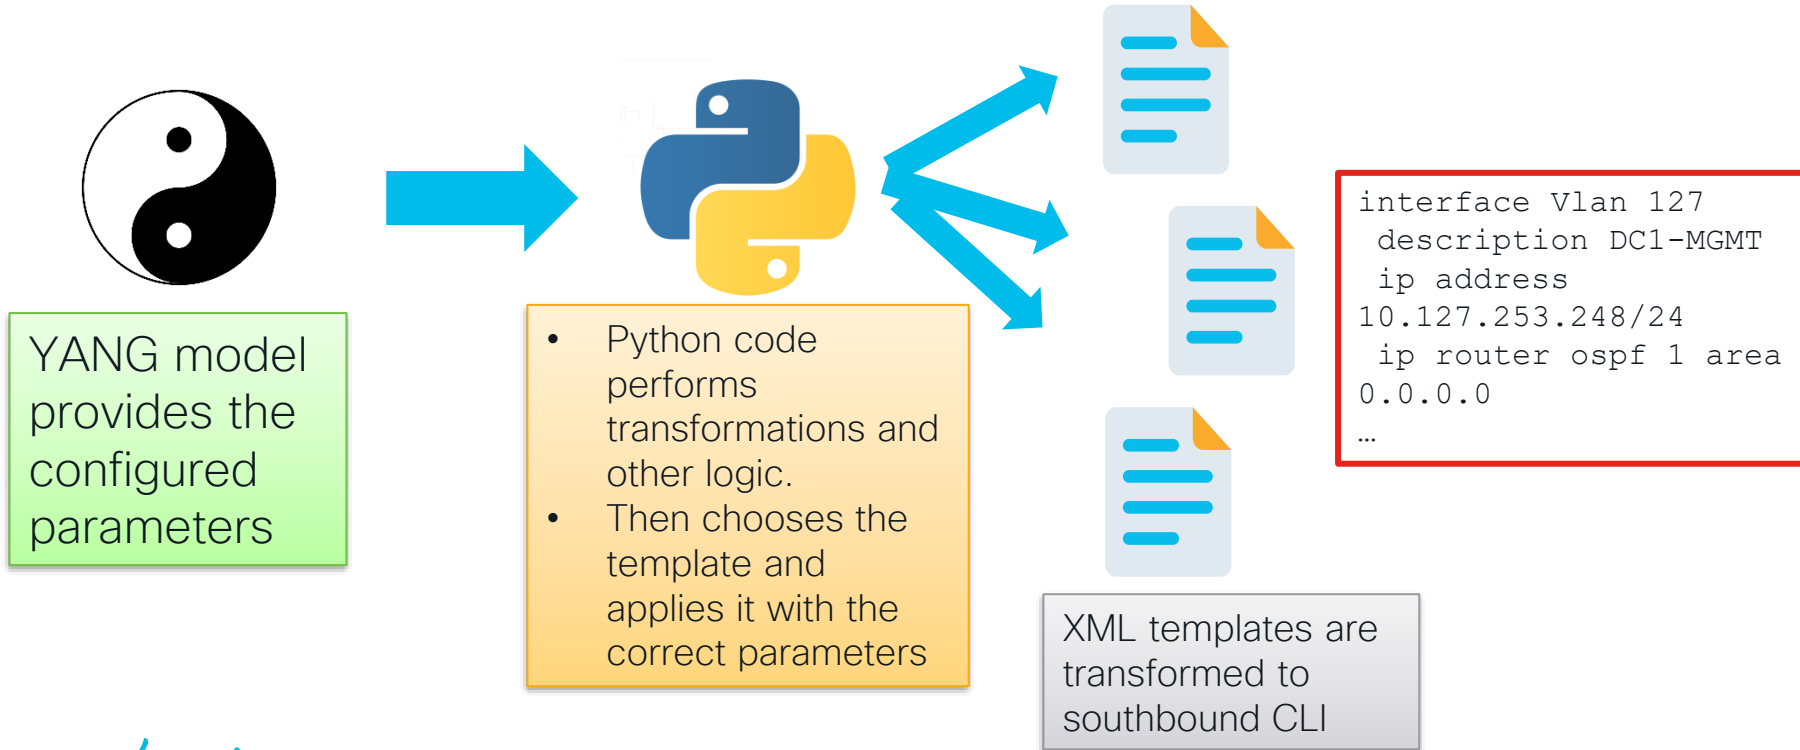
- Iteration is **key** so did all work in git (plus you can have the code 😊)
- For my `ncs.conf`, I added one bit (for Cisco-style CLI):

```
<cli>
```

```
  <style>c</style>
```

```
</cli>
```

How The Python Package Works



The Model Tree

```
module: ciscoive
  +--rw ciscoive* [location year]
  +--rw location string
  +--rw year uint16
  +--rw contact string
  +--rw bandwidth? uint64
  +--rw management
  | +--rw interface
  | | +--rw v4-subnet inet:ipv4-prefix
  | | +--rw v6-prefix inet:ipv6-prefix
  | | +--rw v4-network* inet:ipv4-prefix
  | | +--rw v6-network* inet:ipv6-prefix
  +--rw routing
  | +--rw ip* [prefix next-hop]
  | | +--rw prefix inet:ipv4-prefix
  | | +--rw next-hop inet:ipv4-address
  | | +--rw redistribute? boolean
  | | +--rw data-center? -> /ciscoive/data-center/id
  | +--rw ipv6* [prefix next-hop]
  | | +--rw prefix inet:ipv6-prefix
  | | +--rw next-hop inet:ipv6-address
  | | +--rw redistribute? boolean
  | | +--rw data-center? -> /ciscoive/data-center/id
  +--rw pim
  | +--rw rp? inet:ipv4-address
  | +--rw ssm-range inet:ipv4-prefix
  +--rw dns
  | +--rw server* inet:ipv4-address
  | +--rw domain inet:domain-name
  | +--rw v6-server* inet:ipv6-address
  | +--rw search-list* inet:domain-name
  +--rw dhcp
  | +--rw relay* inet:ipv4-address
  +--rw security
  | +--rw ospf-key talif:aes-cfb-128-encrypted-string
  | +--rw hsp-key talif:aes-cfb-128-encrypted-string
  | +--rw user* [name]
  | | +--rw name string
  | | +--rw role enumeration
  | | +--rw password talif:aes-cfb-128-encrypted-string
  | +--rw oaa
  | | +--rw server* inet:ipv4-address
  | | +--rw tacplus-key talif:aes-cfb-128-encrypted-string
  +--rw ntp
  | +--rw server* inet:ipv4-address
  +--rw snmp
  | +--rw community? talif:aes-cfb-128-encrypted-string
  | +--rw user string
  | +--rw password talif:aes-cfb-128-encrypted-string
  +--rw logging
  | +--rw server* inet:ipv4-address
  +--rw netflow
  | +--rw exporter inet:ipv4-address
  +--rw vcenter* [device]
  | +--rw device -> /ncs:devices/device/name
  | +--rw datacenter
  | | +--rw name -> deref(../device)/../ncs:config/vmm/vcenter/datacenter/name
  | | +--rw cross-dc-switch -> deref(../name)/../vmm:switch/name
  | | +--rw dc-switch* -> deref(../name)/../vmm:switch/name
```

- Generated with pyang -f tree
- We'll discuss some of these elements

```
+--rw vlan* [id]
  +--rw id uint16
  +--rw name string
  +--rw category? segment-category
  +--rw cross-dc? boolean
  +--rw natmed? boolean
  +--rw dhcp boolean
  +--rw native? boolean
  +--rw ip
  | +--rw (prefix-or-dc)?
  | | +--(prefix)
  | | | +--rw prefix inet:ipv4-prefix
  | | | +--(dc)
  | | | +--rw data-center* [id]
  | | | | +--rw id -> /ciscoive/data-center/id
  | | | | +--rw prefix inet:ipv4-prefix
  | | | +--rw access-group? string
  | | +--rw ipid
  | | +--rw (prefix-or-dc)?
  | | | +--(prefix)
  | | | | +--rw prefix inet:ipv4-prefix
  | | | +--(dc)
  | | | +--rw data-center* [id]
  | | | | +--rw id -> /ciscoive/data-center/id
  | | | | +--rw prefix inet:ipv4-prefix
  | | | +--rw traffic-filter? string
  | | +--rw link-local-only? boolean
  +--rw up-down
  | +--rw peer-port-channel uint16
  +--rw member-interface* string
  +--rw knapsack-interface? union
  | +--rw Ethernet* [name]
  | | +--rw name string
  | | +--rw speed? enumeration
  | | +--rw ip
  | | | +--rw address* inet:ipv4-address
  | | | +--rw ipid
  | | | +--rw address* inet:ipv4-address
  | | +--rw (single-description-or-list)?
  | | | +--(single-description)
  | | | | +--rw description? string
  | | | | +--(description-list)
  | | | | +--rw description-list* string
  | | | +--rw mode? enumeration
  | | +--rw category* segment-category
  | | +--rw vlan -> /ciscoive/vlan/id
  | | +--rw all-land-vlan* -> /ciscoive/vlan/id
  | +--rw port-channel* [name]
  | | +--rw name uint16
  | | +--rw protocol? enumeration
  | | +--rw member-interface* [name]
  | | | +--rw name string
  | | | +--rw (single-description-or-list)?
  | | | | +--(single-description)
  | | | | +--rw description? string
  | | | | +--(description-list)
  | | | | +--rw description-list* string
  | | | +--rw (single-description-or-list)?
  | | | | +--(single-description)
  | | | | +--rw description? string
  | | | | +--(description-list)
  | | | | +--rw description-list* string
  | | +--rw mode? enumeration
  | | +--rw category* segment-category
  | | +--rw vlan -> /ciscoive/vlan/id
  | | +--rw all-land-vlan* -> /ciscoive/vlan/id
  +--rw data-center* [id]
  +--rw id uint8
  +--rw location string
  +--rw tenant* [id]
  | +--rw id uint8
  | +--rw device -> /ncs:devices/device/name
  | +--rw knapsack-ip inet:ipv4-address
  +--rw fabric-interconnect* [device]
  +--rw device -> /ncs:devices/device/name
  +--rw unic-template-trunk* [name]
  +--rw name string
  +--rw (org-or-nas)
  | +--(org)
  | | +--rw org? -> /ncs:devices/device/ncs-name-current(Cr/../device)/conf/ucs/org/org/name
  | | +--(nas)
  | | | +--rw nas? empty
  | | | +--rw (org-temp)
  | | | | +--rw (org-temp) -> /ncs:devices/device/ncs-name-current(Cr/../device)/conf/ucs/org/org/name-current(Cr/../org)/nsc-temp/name
  | | | +--rw (nas-temp)
  | | | | +--rw (nas-temp) -> /ncs:devices/device/ncs-name-current(Cr/../device)/conf/ucs/org/org/nsc-temp/name
  | | +--rw description? string
  | | +--rw needs-native? boolean
```

Spinning Up a New "ciscolive"

```
module: ciscolive
  |--rw ciscolive* [location year]
    |--rw location      string
    |--rw year          uint16
```

Tree

```
list ciscolive {
  tailf:info "An instance of the CiscoLive network for a given
year and location.";
  tailf:cli-full-command;
  uses ncs:service-data;
  ncs:servicepoint "ciscolive-servicepoint";
  key "location year";
  leaf location {
    tailf:info "Location of this instance of the
CiscoLive network.";
    type string;
  }
  leaf year {
    tailf:info "Year for this instance of the CiscoLive DC
network.";
    type uint16;
  }
}
```

YANG

- NSO expects service *instances* to be a list
- Each “ciscolive” is indexed using its location and year
- The `tailf:cli-full-command` extension is a Cisco nicety which lets one enter “ciscolive” sub-mode with a command like `ciscolive Amsterdam 2023`
- The `tailf:info` extension provides CLI context-sensitive help

Choosing IDs and Octets

- Some small Python snippets are used to determine its position in a list object (also used for calculating HSRP and STP priorities)
- The core architects like(d) .254/::fe for the [HSRP] gateways, so also grab high values for per-switch octets
- sid == switch ID (1 or 2)
- did == DC ID (1 or 2)
 - dc1-ethsw-1: $[1 - 1 + (1 * 2 - 2) = 0]$
 - $[253 - ((6 - 0)) = 247]$

Python

```
def get_switch_index(did, sid):
```

```
    """
```

```
    Get the switch's flattened list index given the DC ID and the switch ID.
```

```
    """
```

```
    return int(sid) - 1 + (int(did) * 2 - 2)
```

```
def get_switch_octet(did, sid):
```

```
    """
```

```
    Get the IP octet that represents the switch.
```

```
    """
```

```
    return 253 - ((6 - get_switch_index(did, sid)))
```

HSRP and STP

- dc1-ethsw-1 should have lowest STP priority and highest HSRP priority
- Each subsequent switch should follow in suit
- STP:
 - dc1-ethsw-1: $0 * 4096 = 0$
 - dc2-ethsw-1: $2 * 4096 = 8192$
- HSRP:
 - dc1-ethsw-2: $105 - 1 = 104$
 - dc2-ethsw-2: $105 - 3 = 102$

```
# Set a spanning-tree priority based on switch and DC IDs.  
stp_prio = get_switch_index(dc.id, switch.id) * 4096  
...  
  
# Determine HSRP priority based on switch ID and DC ID  
svi_vars.add("HSRP_PRIORITY", 105 - get_switch_index(dc.id,  
switch.id))
```

Python

With Respect to Passwords...

```
show running-config
ciscolive Amsterdam 2023
snmp
```


```
container snmp {
  leaf community {
    tailf:info "SNMP community string (read-only)";
    type tailf:aes-cfb-128-encrypted-string;
  }
  leaf password {
    tailf:info "SNMPv3 password.";
    type tailf:aes-cfb-128-encrypted-string;
    mandatory true;
  }
}
```

YANG

```
from _ncs import decrypt
...
trans = ncs.maagic.get_trans(root)
trans.maapi.install_crypto_keys()
...
snmp_vars.add("COMMUNITY",
decrypt(self.service.snmp.community))
```

Python

```
ciscolive Amsterdam 2023
snmp
  community not-public
  user      CLEUR
  password really-secret
!
```



```
ciscolive Amsterdam 2023
snmp
  community $8$Ykq...
  user      CLEUR
  password $8$5IQ...
!
```

Modeling VLANs

```

+--rw vlan* [id]
  |   +--rw id          uint16
  |   +--rw name        string
  |   +--rw category?   segment-
category
  |   +--rw cross-dc?   boolean
  |   +--rw routed?    boolean
  |   +--rw dhcp        boolean
  |   +--rw native?     boolean
  |   +--rw ip
...
  |   +--rw ipv6
...
```

Tree

- Each VLAN can be L2 only or routed (i.e., having an SVI)
- A VLAN can either be local to a DC or stretched (this mainly affects IP space)
- The `category` leaf is used to automatically group allowed VLANs on trunks

Segment Categories

- An enumeration typedef to better describe VLANs
- Tied to interfaces (interfaces can have multiple categories)

```
typedef segment-category {  
    tailf:info "Usage category of a given  
segment.";  
    type enumeration {  
        enum vm {  
            tailf:info "Segment is used for virtual  
machines.";  
        }  
        enum netapp {  
            tailf:info "Segment is used for NetApp-  
only traffic.";  
        }  
        enum fabric-interconnect {  
            tailf:info "Segment is used for FI-only  
traffic.";  
        }  
        enum peer {  
            tailf:info "Segment is used for route  
peering.";  
        }  
    }  
}
```

YANG

Constructing the L2 VLAN Parameters for Switches and FIs

- Parameterize the VLAN ID
- When creating the “name” parameter, substitute “{dc}” and “{peer_dc}” strings with the DC ID and peer DC ID respectively
- Convenient for those DC-local VLANs
- Allows for a name like “MGMT-DC1” vs. simply “MGMT”

Python

```
for vlan in self.service.vlan:  
    l2_vars.add("VLAN_ID", vlan.id)  
    l2_vars.add("VLAN_NAME",  
vlan.name.format(dc=dc.id,  
peer_dc=get_peer_id(dc.id)))
```

Constructing the L2 VLAN Parameters for vCenter

- Like the network and compute sides
- PG or Port Group is more appropriate as a parameter name
- The correct dVS also needs to be specified
- Note: the category leaf is used here so that we only add VM-supporting VLANs

Python

```
if vlan.category != "vm":  
    continue  
  
vc_vars.add("VLAN_ID", vlan.id)  
vc_vars.add("PG_NAME", vlan.name.format(dc=dc.id,  
peer_dc=get_peer_id(dc.id)))  
  
if vlan.cross_dc:  
    vc_vars.add("VSWITCH",  
vcenter.datacenter.cross_dc_vswitch)  
else:  
    vc_vars.add("VSWITCH",  
list(vcenter.datacenter.dc_vswitch)[int(dc.id) - 1])
```

Constructing the L3 SVI Parameters

If the VLAN is a stretched VLAN get its single prefix

Calculate the last octet for the SVI and for the HSRP VIP as the hex version of the v4 octet

```
# Add IPv6 parameters
if ("ciscolive:prefix" in vlan.ipv6 and vlan.ipv6.prefix and vlan.ipv6.prefix != "") or (
    "ciscolive:data-center" in vlan.ipv6 and not vlan.ipv6.link_local_only
):
    if vlan.cross_dc and vlan.category != "peer":
        ipv6_prefix = ipaddress.ip_network(vlan.ipv6.prefix)
    else:
        dc_prefix = next(d for d in vlan.ipv6.data_center if d["id"] == dc.id)
        ipv6_prefix = ipaddress.ip_network(dc_prefix.prefix)

vip_octet = format(int(str(v4_vip).split(".")[3]), "x")
v6_octet = format(int(str(v4_addr).split(".")[3]), "x")
svi_vars.add(
    "HSRP_V6_VIP",
    f"{ipv6_prefix.network_address}{vip_octet}",
)
svi_vars.add(
    "SVI_V6",
    f"{ipv6_prefix.network_address}{v6_octet}/{ipv6_prefix.prefixlen}",
)
svi_vars.add("LINK_LOCAL", "False")
if vlan.ipv6.traffic_filter and vlan.ipv6.traffic_filter != "":
    svi_vars.add("TRAFFIC_FILTER", vlan.ipv6.traffic_filter)
else:
    svi_vars.add("TRAFFIC_FILTER", "")
```

Python

Note: for IPv6, an SVI doesn't need a global address. We can use link-local only.

...Else get its per-DC prefix

The model lets us specify a traffic-filter (or ACL), but those are not defined in the model

Modeling Ethernet Interfaces

- Each ethernet interface has the same config on all switches
- The must constraints ensure an interface isn't already used for vPC peer or part of a port-channel (those are handled differently in the model)
- The `tailf:cli-add-mode` extension allows a sub-mode for interface Ethernet 1/1
- The `tailf:cli-allow-join-with-key` extension allows for interface Ethernet1/1
- ☹️ Why no leafref for name?

YANG

```
container interface {
  tailf:info "Physical interface configuration.";
  tailf:cli-add-mode;
  list Ethernet {
    tailf:info "List of ethernet interfaces.";
    tailf:cli-allow-join-with-key;
    key "name";
    leaf name {
      tailf:info "Ethernet interface name.";
      type string {
        pattern
          '[0-9]+/[0-9]+(/[0-9]+)?';
      }
      must "count(/cisco/vpc-peer/member-interface[. = current()]) < 1" {
        error-message
          "Interface cannot be used as a vPC peer link.";
      }
      must "count(..../port-channel/member-interface/name[. = current()]) < 1" {
        error-message
          "Interface cannot be used as it is already a port-channel member.";
      }
      must "count(/cisco/vpc-peer/keepalive-interface[. = current()]) < 1" {
        error-message
          "Interface cannot be used as it is already a vPC keepalive interface.";
      }
    }
  }
}
```

The Power of Constraints

```
port-channel 66
  member-interface 1/35
  description-list [ "-> HX-FI-A_e1/35" "-> HX-FI-A_e1/36" ]
  !
  description-list [ "-> DC1-HX-FI-A" "-> DC2-HX-FI-A" ]
  mode                trunk
  category             [ vm ]
```

```
jclarke@ncs (config-ciscolive-Amsterdam/2023) # interface Ethernet1/35
jclarke@ncs (config-Ethernet-1/35) # ip address [ 1.1.1.1 1.1.1.2 1.1.1.3 1.1.1.4 ]
jclarke@ncs (config-Ethernet-1/35) # commit
Aborted: 'ciscolive Amsterdam 2023 interface Ethernet 1/35 name' (value "1")
Interface cannot be used as it is already a port-channel member.
```



Ethernet Interface IPs

- An ethernet interface can be used for direct peering (i.e., to the network core)
- Note: mode here can be edge (DC edge), access, trunk, or cross-dc-link
- Therefore, each switch will have a different IP
- I chose a list approach using the switch's index

```
container ip {  
    tailf:info "IPv4 commands for edge  
interfaces.";  
    leaf-list address {  
        tailf:info "IPv4 address for  
interface on each switch (used with  
/31).";  
        ordered-by user;  
        type inet:ipv4-address;  
        min-elements 4;  
        max-elements 4;  
        when '.../.../mode = "edge"';  
    }  
}
```

YANG

Ethernet 1/31

```
ip address [ 10.128.1.97 10.128.1.99  
10.128.1.101 10.128.1.103 ]
```

dc1-ethsw-1

dc1-ethsw-2

dc2-ethsw-1

dc2-ethsw-2

Port-channel Interfaces

- NX-OS makes you specify an interface is part of a port-channel
- I found it easier, especially for constraints, to have port-channels specify their members
- It's your model. You have the flexibility

```
list member-interface {  
    tailf:info "Member interfaces of this  
port-channel.";  
    key "name";  
    leaf name {  
        type string {  
            pattern '[0-9]+/[0-9]+(/[0-9]  
9]+)?';  
        }  
    }  
    uses description-details;  
    min-elements 1;  
}
```

YANG

A Trade-Off Between Variability and Ease of Use

- Not everything needs to be a YANG element
- You know your service intent, so feel free to have more static parameters in the XML templates
- Or, put them in a separate variable Python file; or add them as config false nodes

XML

```
<ospf>
  <?if {$CATEGORY = "peer"}?>
    <authentication>
      <authentication-type>message-
digest</authentication-type>
      <key-chain>OSPF_KEY</key-chain>
    </authentication>
    <passive-interface>>false</passive-interface>
    <?end?>
    <network>point-to-point</network>
  </ospf>
<redirects>>false</redirects>
<router>
  <ospf>
    <name>1</name>
    <area>0.0.0.0</area>
  </ospf>
</router>
```

Example of Static Parameters as config false

- Create a top-level OSPF area parameter
- Use `show cisco live Amsterdam 2023 ospf-area` to display it
- Accessible inside Python as `service.ospf_area`
- Doesn't clutter the `show running-config` output

```
leaf ospf-area {  
    tailf:info "Static OSPF area ID";  
    type inet:ipv4-address;  
    default "0.0.0.0";  
    config false;  
}
```

Final Service Config



Almost...

- Final config at https://github.com/CiscoLearning/ciscolive-brkops-2040/blob/main/cl_2023.cfg
- Number of non-comment lines: **263**
- Number of resulting southbound config commands: **3467**

13x Config Compression!

Getting the Code

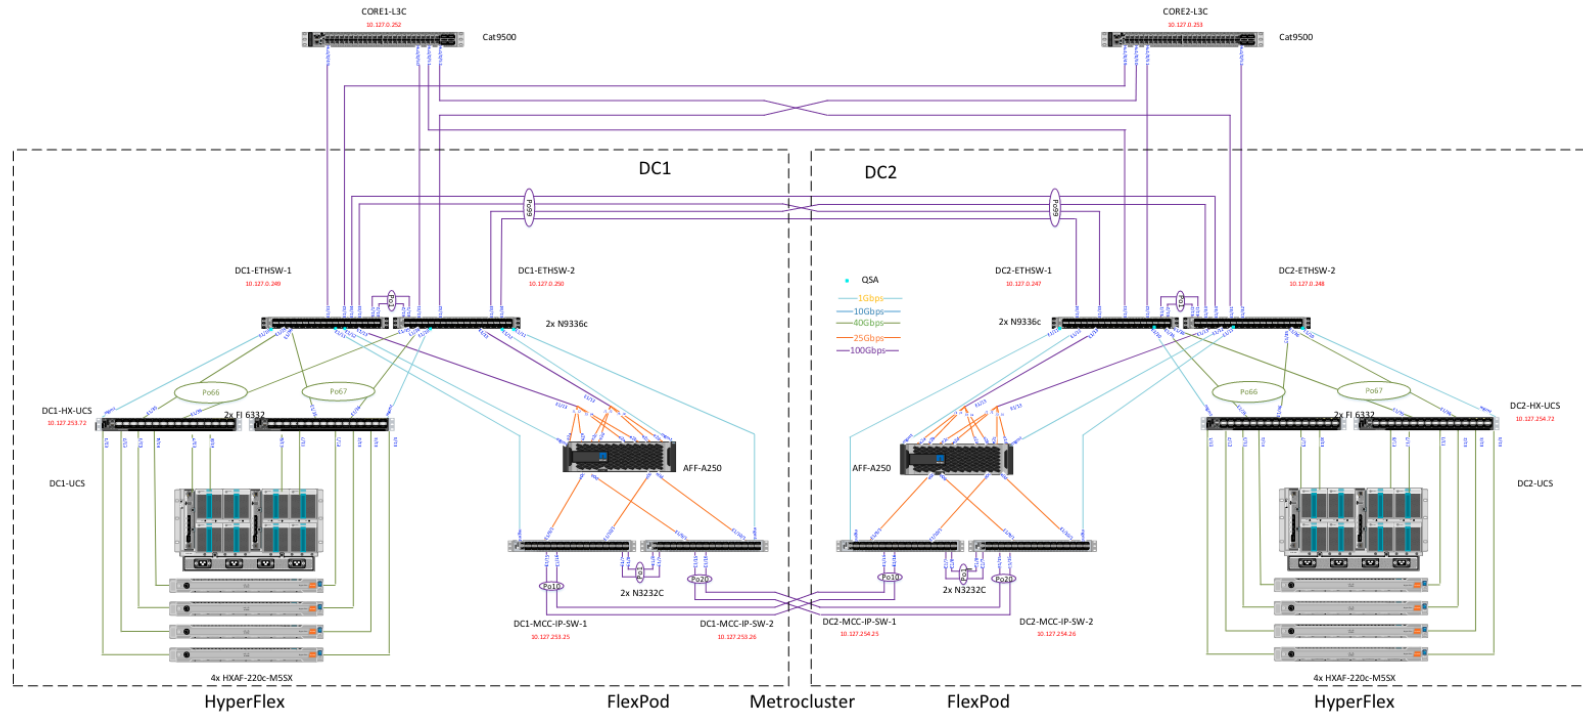
- Code is in GitHub at <https://github.com/CiscoLearnin/g/ciscolive-brkops-2040>
- YANG module is in `src/yang/`
- XML templates are in `templates/`
- Python code is in `python/ciscolive/`




Testing Things Out




Our Actual Network



Our CML Test Network

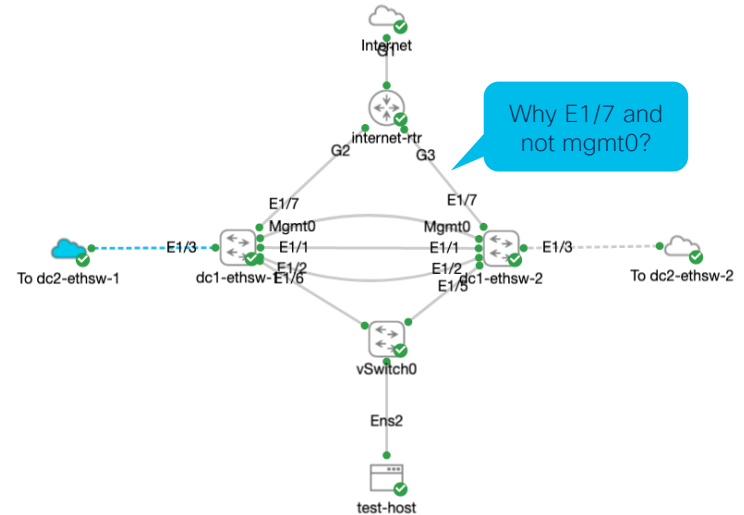


Exports are in
cm1/ directory in
the git repo



Connection Woes

- The current DC layout uses mgmt0 for the vPC keepalive, but connected back-to-back
- This is old technical debt that has not yet been removed
- It required more initial config on the switch and re-deploy reconcile to assume ownership of config by the service



Some Q&A

1. Why CML?
 2. Why use two labs/topologies?
 3. What about UCSM and vCenter?
1. While NSO Netsim is nice for strict model testing, CML let me confirm ultimate functionality.
 2. Using two labs was a bit cleaner, and easier to iterate as the model expanded.
 3. The networking pieces were the biggest lift. While I could have added the UCS simulator or a test vCenter, it wasn't as critical.

Testing Tips



- Commit frequently; use `un-deploy` if you want to back everything out
- As you iterate, use `re-deploy` to resync the network with the model
- Use `outformat` to aid in template development and understand your southbound changes
- The `dry-run` argument is your friend!

```
admin@ncs(config-ipv6)# commit dry-run outformat native
```

```
...
```

```
data vlan 143
  name "CL_Test"
  exit
interface Vlan143
  no shutdown
  description CL_Test
  ip address 10.143.252.247/24
  no ip arp gratuitous hsrp duplicate
  ip ospf network point-to-point
  no ip redirects
  ip router ospf 1 area 0.0.0.0
  ip flow monitor CL-ipv4-nb input
  ipv6 address 2a11:d940:2:8ffc::f7/64
```

```
admin@ncs(config-ipv6)# commit dry-run outformat xml
```

```
...
```

```
data <ciscolive xmlns="http://example.com/ciscolive">
  <location>Amsterdam</location>
  <year>2023</year>
  <vlan>
    <id>143</id>
    <name>CL_TEST</name>
    <routed>true</routed>
    <dhcp>true</dhcp>
    <ip>
      <prefix>10.143.252.0/24</prefix>
    </ip>
    <ipv6>
      <prefix>2a11:d940:2:8ffc::/64</prefix>
    </ipv6>
  </vlan>
</data>
```

Looking Forward to 2024



Things I'd Do Differently

While the NSO solution worked quite well, there's always room for improvement. Remember, it's about iteration. Don't let perfect be the enemy of good.



Model

Use `config false`

Fold static data into the model as `config false nodes`

Compute IP[v6] prefixes

Prefixes are based on VLAN, so they need not be specified explicitly



Network

Dedicated out-of-band network

Will be much easier to provision if the mgmt0 interfaces are reachable by default



Integration

Incorporate source of truth more

We use NetBox for inventory management and IPAM. Use it for the authoritative list of VLANs so they only need to be added once

ACL Management

ACLs come from "The Tool". Either integrate with this as with NetBox or consider modeling them in NSO for the DC

What To Learn More?

- Come to our panel discussion on Friday, **PNLNMS-1035**
- Each part of the CiscoLive network will be presented
- We talk architecture, lessons learned and stats
- We know it will have been a long week; we make this talk fun!



What Will You Automate?

We have all come here to learn. What will you take back?



Brought to you by **Cisco** *live!*

Complete your Session Survey

- Please complete your session survey after each session. Your feedback is important.
- Complete a minimum of 4 session surveys and the Overall Conference survey (open from Thursday) to receive your Cisco Live t-shirt.
- All surveys can be taken in the Cisco Events Mobile App or by logging in to the Session Catalog and clicking the "Attendee Dashboard" at <https://www.ciscolive.com/emea/learn/sessions/session-catalog.html>



Continue Your Education



Visit the Cisco Showcase for related demos.



Book your one-on-one Meet the Engineer meeting.



Attend any of the related sessions at the DevNet, Capture the Flag, and Walk-in Labs zones.



Visit the On-Demand Library for more sessions at ciscolive.com/on-demand.



The bridge to possible

Thank you

CISCO *Live!*

CISCO *Live!*

ALL IN