



You make **possible**



Traffic Engineering with Segment Routing

Leonir Hoxha – Senior Consulting Engineer
BRKSPG-2021



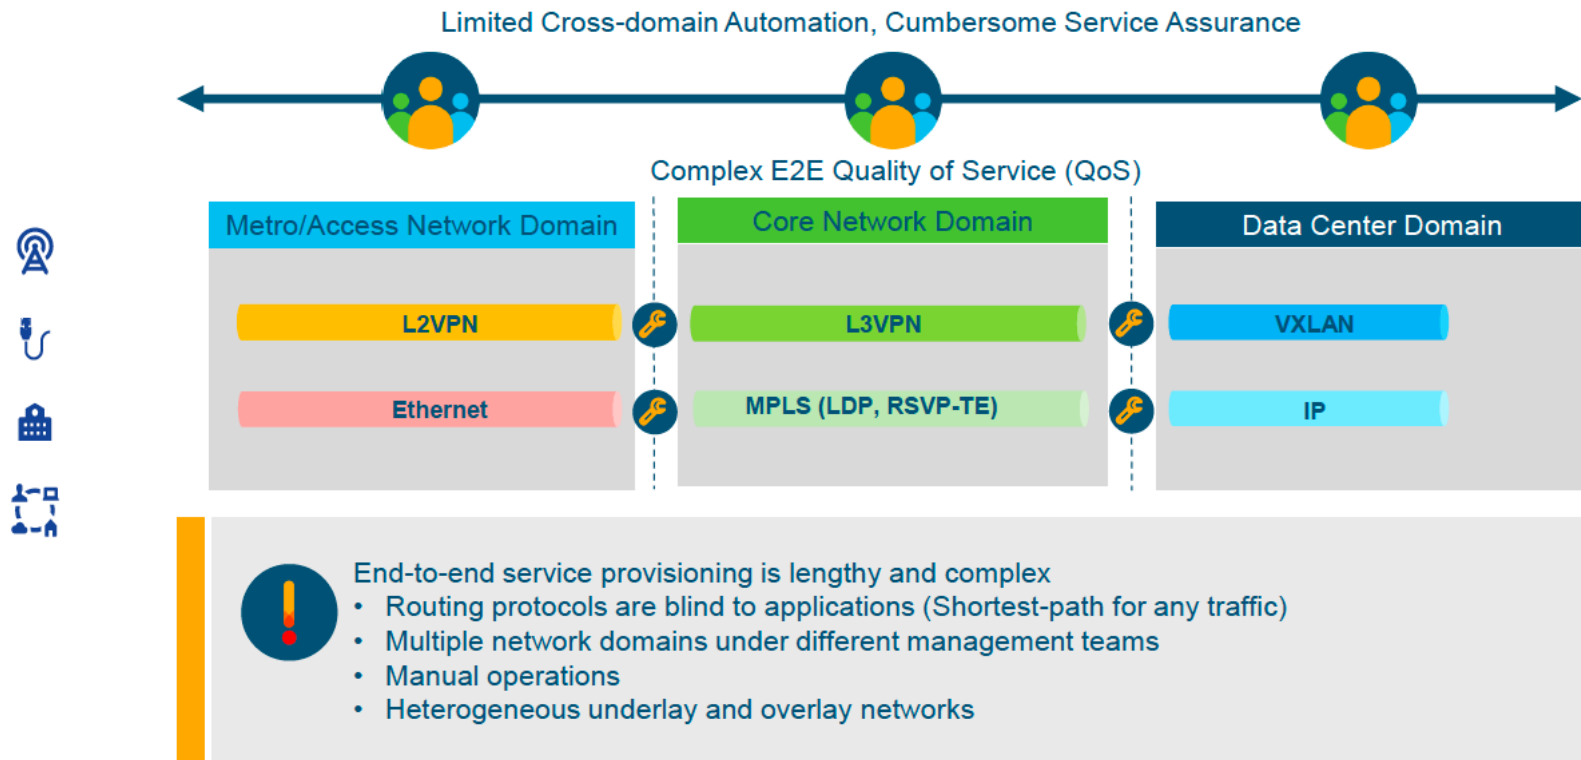
Agenda

- Segment Routing Overview
- Control & Data Plane
- Segment Routing Traffic Engineering
- Automated Steering
- On-Demand Next-Hop
- SR-PCE
- Flexible Algorithm (Flex-Algo)

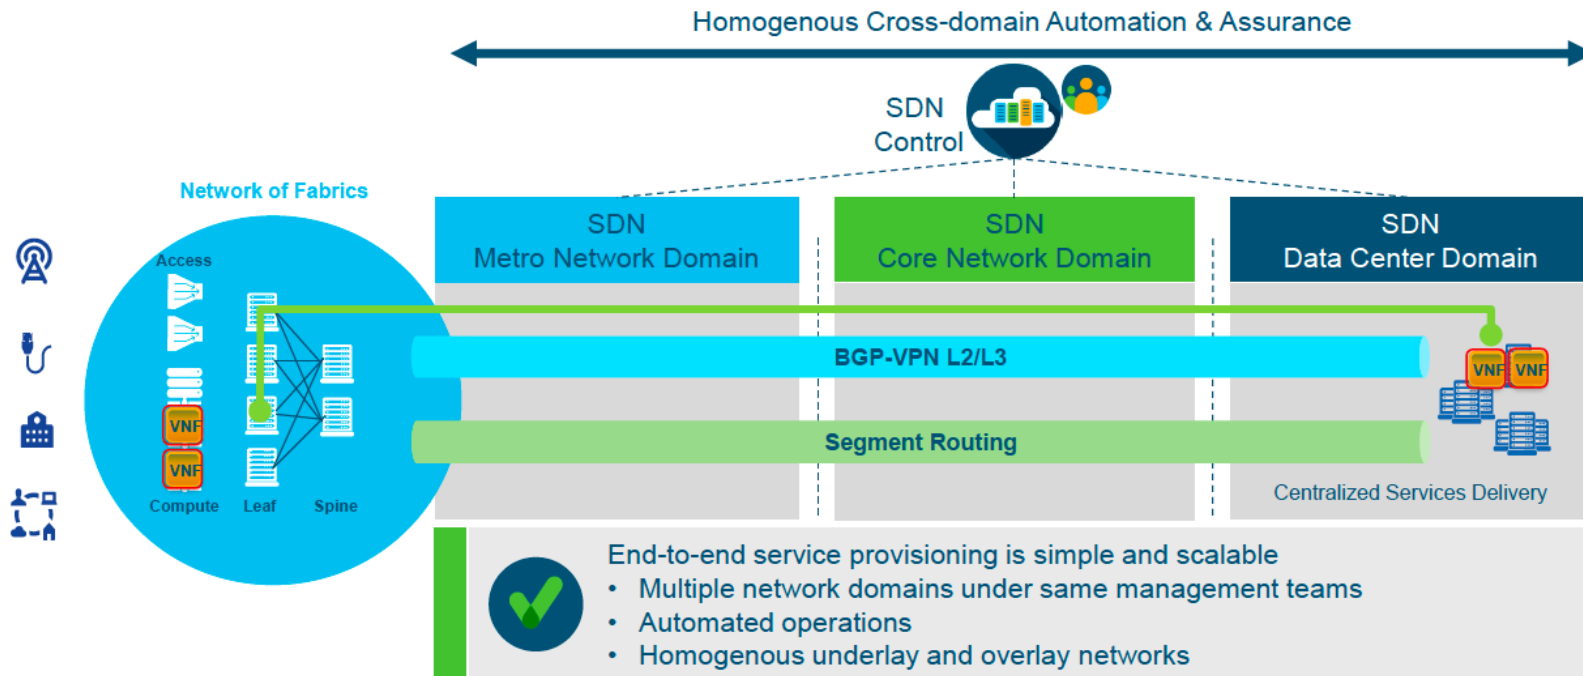


Segment... what?

Problem Statement: Today's Service Creation



Segment Routing Unified Fabric Vision



Segment Routing: Value Proposition



Multi-vendor consensus – Designed and built with network operators

Segment Routing Overview



You make networking **possible**

Segment Routing

- **Source Routing**
 - the source chooses a path and encodes it in the packet header as an ordered list of segments
 - the rest of the network executes the encoded instructions
- **Segment**: an identifier for any type of instruction
 - forwarding or service

Segment Routing – Forwarding Plane


- **MPLS**: an ordered list of segments is represented as a stack of labels
- **IPv6**: an ordered list of segments is encoded in a routing extension header
- This session: **MPLS data plane**

Segment Routing – IGP Segments

- Two basic building blocks distributed by IGP
 - Prefix Segments
 - Adjacency Segments

IGP Prefix Segment


- Shortest-path to the IGP prefix
 - Equal Cost Multipath (ECMP) Aware
- Global Segment
- Label = 16000 + Index
 - Advertised as index
- Distributed by ISIS/OSPF
- Manually assigned



All nodes use default SRGB
16,000 – 23,999

IGP Prefix Segment

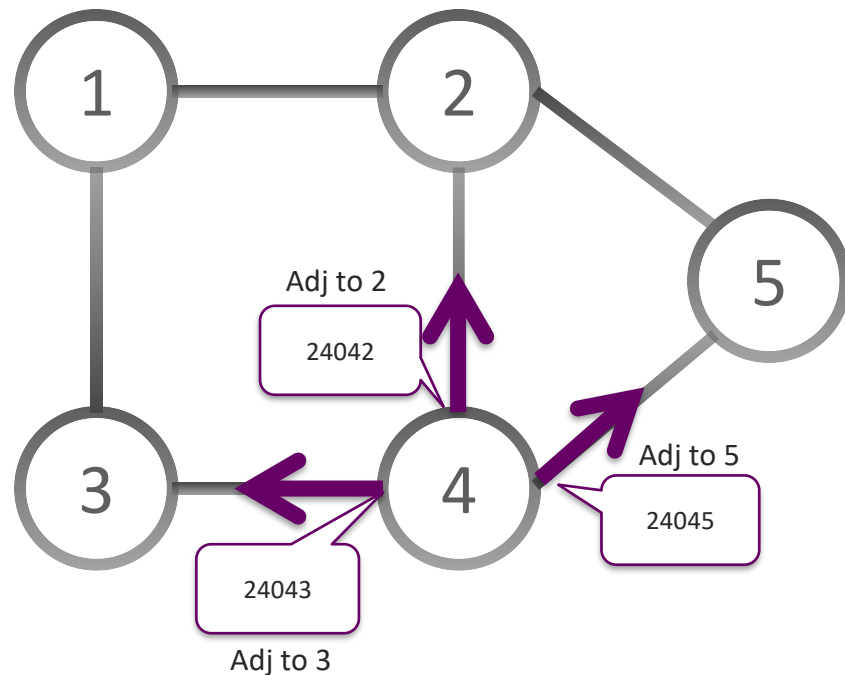
- Shortest-path to the IGP prefix
 - Equal Cost Multipath (ECMP) Aware
- Global Segment
- Label = 16000 + Index
 - Advertised as index
- Distributed by ISIS/OSPF
- Manually assigned



All nodes use default SRGB
16,000 – 23,999


IGP Adjacency Segment

- Forward on the IGP adjacency
- Local Segment
 - Dynamically Allocated
- Distributed by ISIS/OSPF
- Advertised as a label value



Combining IGP Segments

- Steer traffic on any path through the network
- Path is specified by a stack of labels
- No path is signaled
- No per-flow state is created
- Single protocol: IS-IS or OSPF





Control & Data Plane



You make security **possible**



MPLS Control and Forwarding Operation with Segment Routing

Services



No changes to control or forwarding plane

Packet Transport



IGP or BGP label distribution for IPv4 and IPv6. Forwarding plane remains the same (MPLS)

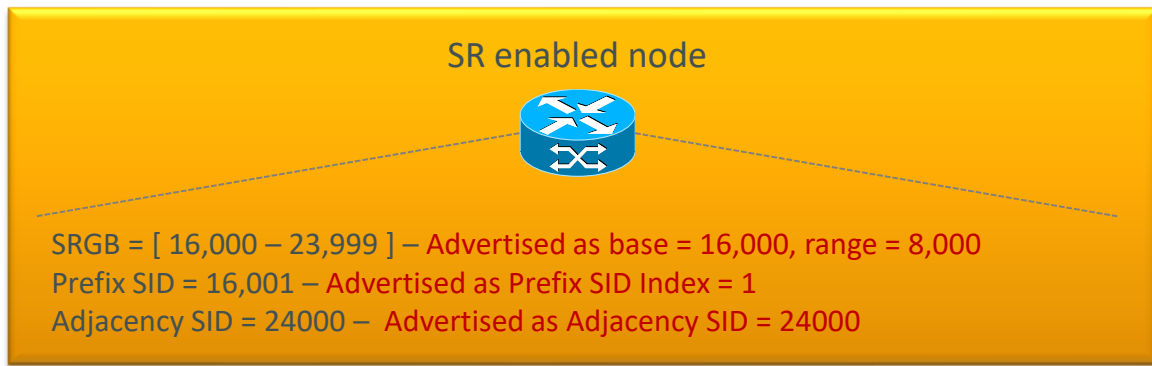
SID Encoding

- **Prefix SID**

- **Label** form SR Global Block (SRGB)
- **SRGB** advertised within IGP via TLV
- In the configuration, Prefix-SID can be configured as an **absolute** value or an **index**
- In the protocol advertisement, Prefix-SID is always encoded as a **globally unique index**
Index represents an **offset** from SRGB base, zero-based numbering, i.e. 0 is 1st index
E.g. index **1** → SID is $16,000 + 1 = 16,001$

- **Adjacency SID**

- **Locally significant**
- **Automatically allocated** by the IGP for each adjacency
- Always encoded as an absolute (i.e. not indexed) value



SR IS-IS Control Plane Summary


- IPv4 and IPv6 control plane
- Level 1, level 2 and multi-level routing
- Prefix Segment ID (Prefix-SID) for host prefixes on loopback interfaces
- Adjacency Segment IDs (Adj-SIDs) for adjacencies
- Prefix-to-SID mapping advertisements (mapping server)
- MPLS penultimate hop popping (PHP) and explicit-null signaling

SR OSPF Control Plane Summary

- OSPFv2 control plane
- Multi-area
- IPv4 Prefix Segment ID (Prefix-SID) for host prefixes on loopback interfaces
- Adjacency Segment ID (Adj-SIDs) for adjacencies
- Prefix-to-SID mapping advertisements (mapping server)
- MPLS penultimate hop popping (PHP) and explicit-null signaling


MPLS Data Plane Operation (Labeled)

Prefix SID



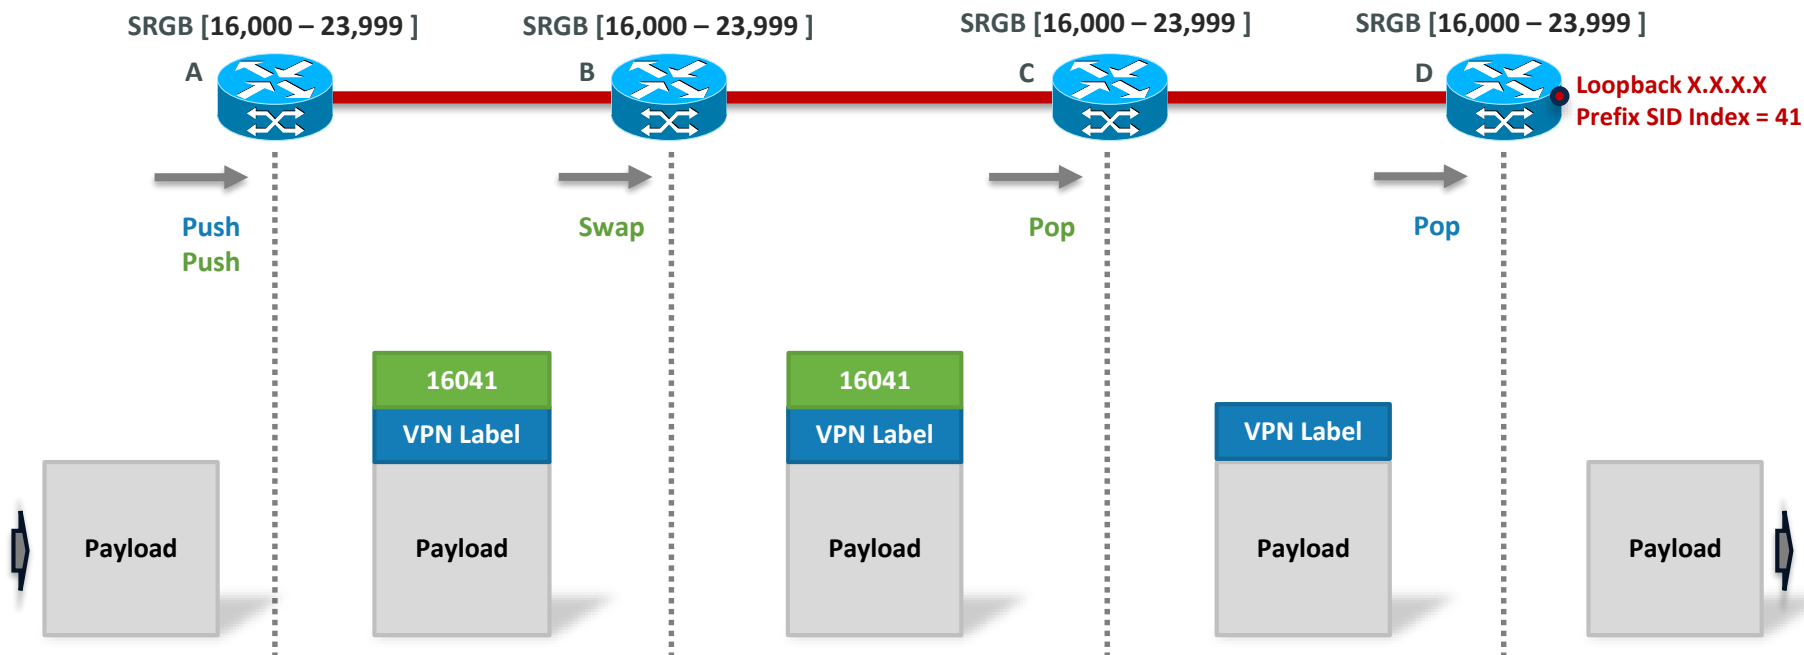
- Packet forwarded along IGP shortest path (ECMP)
- Swap operation performed on input label
- Same top label if same/similar SRGB
- PHP if signaled by egress LSR

Adjacency SID

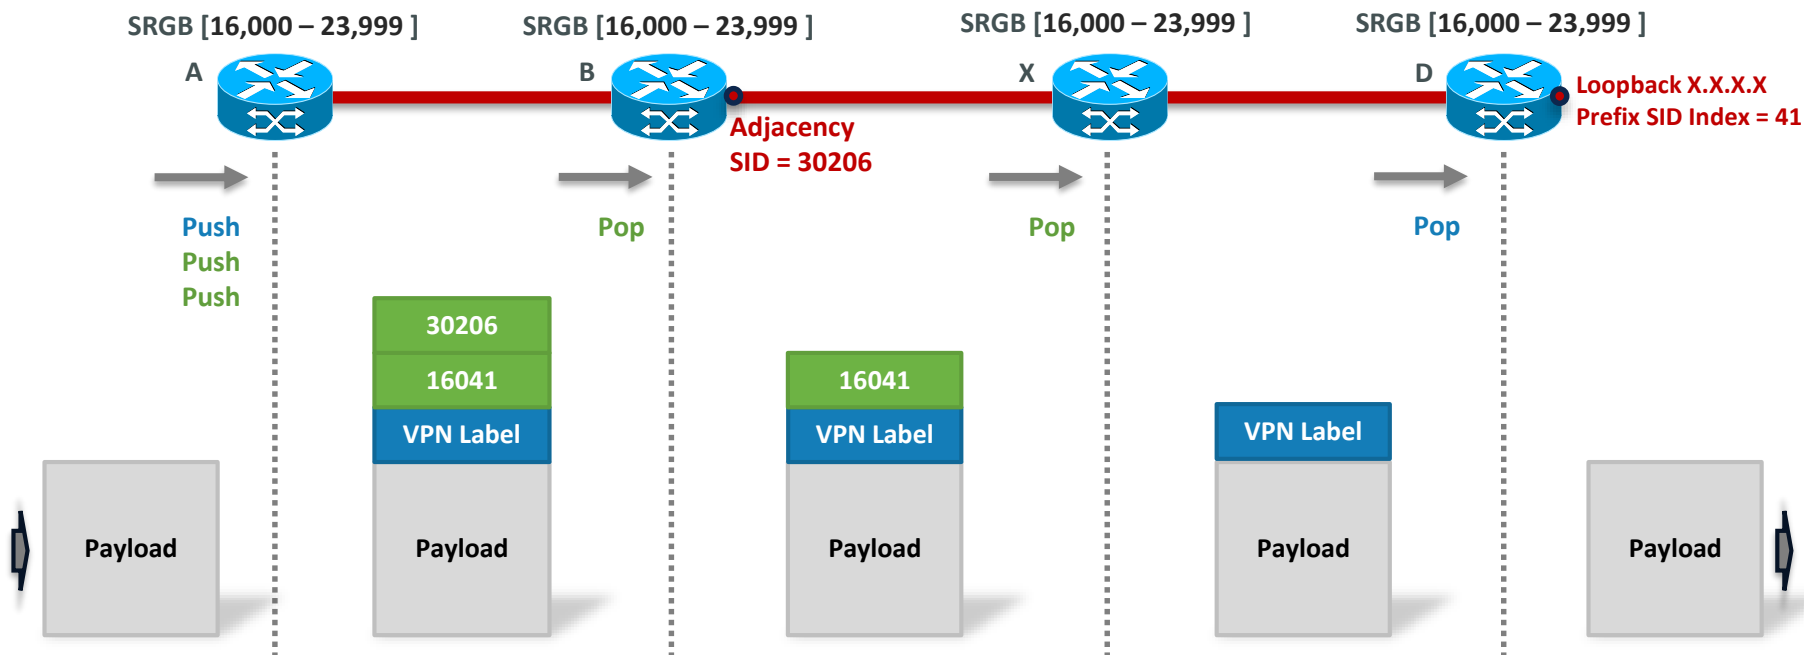


- Packet forwarded along IGP adjacency
- Pop operation performed on input label
- Top labels will likely differ
- Penultimate hop always pops last adjacency SID

MPLS Data Plane Operation (Prefix SID)



MPLS Data Plane Operation (Adjacency SID)




Segment Routing Traffic Engineering



You make the power of data **possible**


Traffic Engineering with SR

- Source based routing
 - State only at the Ingress PE
 - Supports constraint-based routing
- Uses existing ISIS / OSPF extensions
 - To advertise link attributes
- No RSVP-TE to establish LSPs
- ECMP Aware




SR Policy

- SR-Policy is a Traffic Engineering intent by means of a solution SID-list.
- An SR-Policy is uniquely identified by a tuple:
 - **Head-end:** where the SR Policy has been instantiated (configured)
 - **End-point:** the destination of the SR Policy
 - **Color:** a numerical value to differentiate multiple SRTE Policies between the same pair of nodes



SR Policy Color

- Each SR Policy has a color
 - Color can be used to indicate a certain treatment (SLA, policy) provided by an SR Policy
- Only one SR Policy with a given color C can exist between a given node pair (head-end (H), end-point (E))
 - In other words: each SR Policy triplet (H, C, E) is unique
- Example:
 - Low-cost="blue", Low-delay="green"
 - steer traffic to 1.1.1.0/24 via Node4 into Low-cost SR Policy (1, blue, 4)
 - steer traffic to 2.2.2.0/24 via Node4 into Low-delay SR Policy (1, green, 4)




SR Policy on IOS-XR

- SR Policy Configuration on IOS-XR

```
segment-routing
traffic-eng
policy POLICY1
  color 20 end-point ipv4 1.1.1.4
  candidate-paths
    preference 100
    dynamic
    metric
    type igp
```

Binding SID

- The **Binding SID** is a fundamental building block of SR-TE solution
- Explicitly configured or dynamically allocated
- A Binding SID identifies a SRTE Policy
 - Packet received with Binding-SID as Top Label is steered into the SRTE Policy associated with the Binding-SID
 - Binding-SID label is popped, SRTE Policy's SID list is pushed



Binding SID Allocation

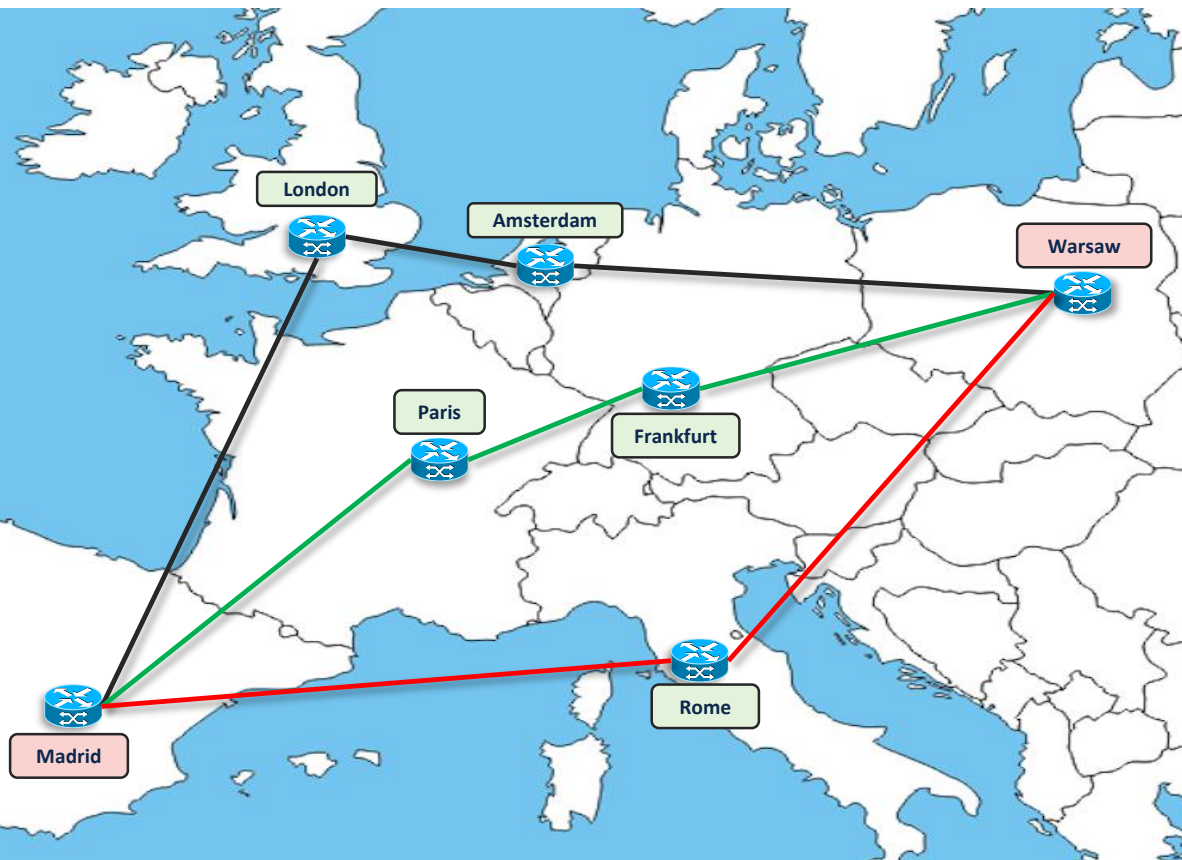
- Explicit allocation

```
segment-routing
traffic-eng
policy POLICY1
  color 20 end-point ipv4 1.1.1.6
  binding-sid mpls 50100
  candidate-paths
  preference 100
  dynamic
  metric
  type igp
```

- Dynamic allocation is the default

Practical Use-Cases (Optimization Based SR-TE Path)

- TE Affinity-bit (coloring)
- SRLG
- Delay
- Disjoint Path



High Bandwidth Path:

Madrid – London – Amsterdam – Warsaw

Secure / Encrypted Path (MacSec enabled):

Madrid – Paris – Frankfurt – Warsaw

Low Delay Path:

Madrid – Rome – Warsaw

Disjoint Path:

1. MAD – PAR – FFM – WAW
2. MAD – RME – WAW
3. MAD – LON – AMS – WAW

Practical Use-Cases (SR Policy Configuration)

MacSec enabled Path

```
segment-routing
traffic-eng
affinity-map
!! 32-bit maps
name macsec bit-position 0
!
interface HundredGig0/0/0/0
affinity name macsec
!
policy POLICY1
color 20 end-point ipv4 1.1.1.4
candidate-paths
preference 100
dynamic
metric type igp
constraints
affinity
include-all name macsec
```

Low Delay Path

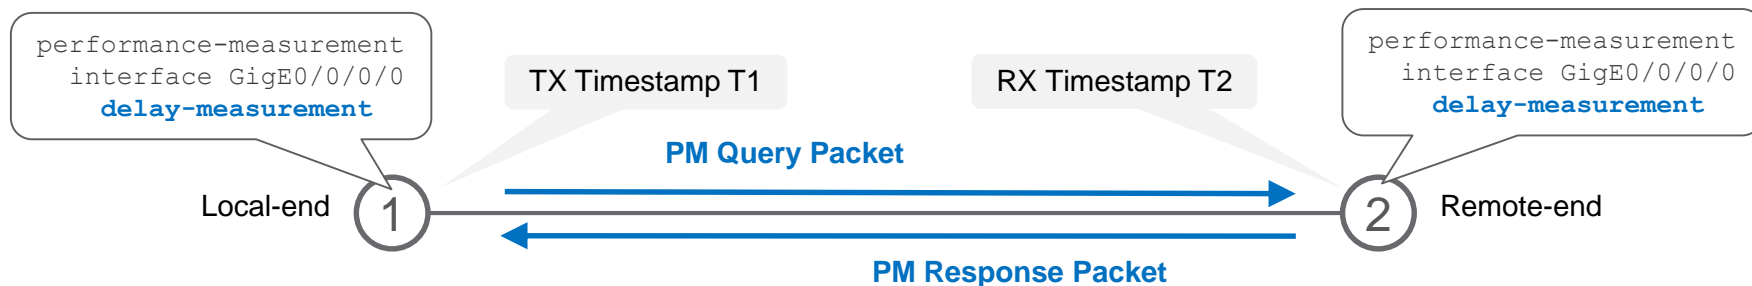
```
segment-routing
traffic-eng
!
policy POLICY2
color 30 end-point ipv4 1.1.1.4
candidate-paths
preference 110
dynamic
metric
type delay
```

Disjoint Path

```
segment-routing
traffic-eng
!
policy POLICY3
color 10 end-point 1.1.1.4
candidate-paths
preference 120
dynamic
pcep
metric type te
constraints
disjoint-path group-id 1 type node
```


Delay Measurement

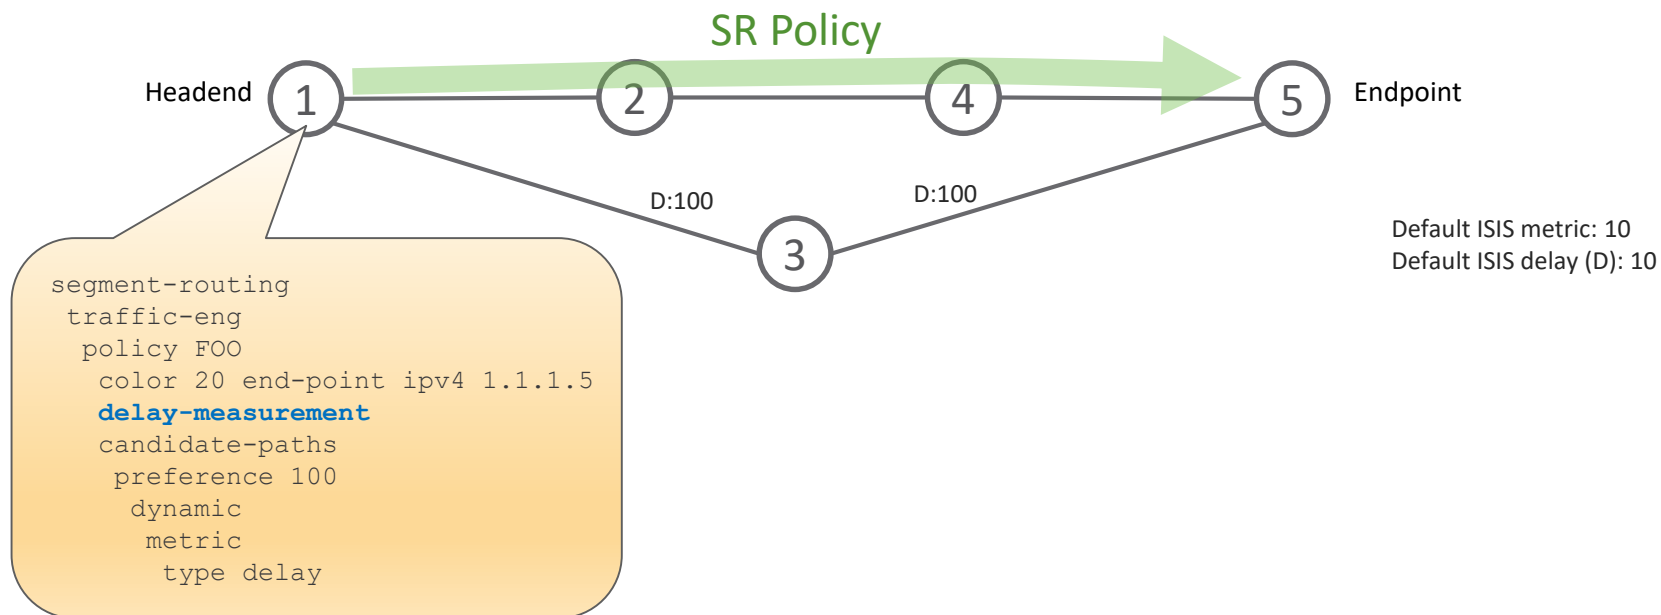
- Link Delay (using Probes) Measurement feature




- One Way Delay = $(T2 - T1)$
- Timestamps added in hardware
- PM Query format: RFC 6374 (MPLS/GAL) or RFC 5357 (IP/UDP/TWAMP)

Default: every 3 sec

Per SR-Policy Delay Measurement



Probe Measurement



- One Way Delay = $(T2 - T1)$
 - Requires clock synchronization
- Default: Send Query every 3 sec

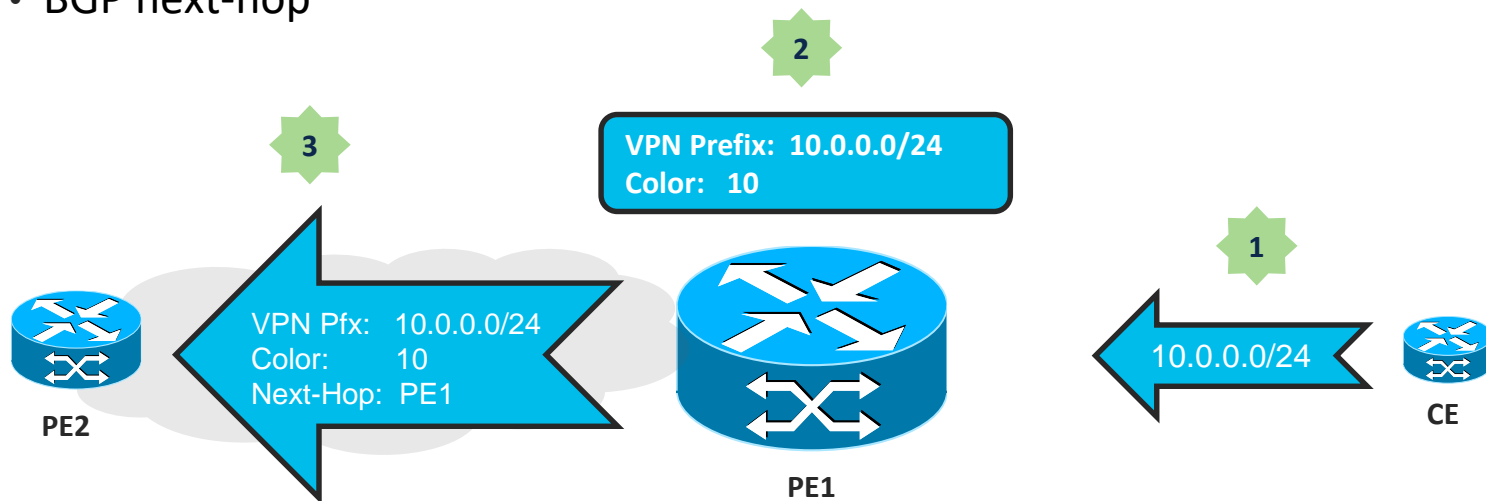
Automated Steering



You make customer experience **possible**


Automated Steering

- Consists of:
 - Prefix / Customer routes (L2/L3 VPN)
 - Color (BGP extended community attribute)
 - BGP next-hop



Automated Steering

- BGP can automatically steer traffic into an SR Policy based on **BGP next-hop** and **color** of a route
 - Color of a route is specified by its color extended community attribute
- When a **BGP next-hop** and **color** of a route match the end-point and color of an SR Policy, then BGP installs the route resolving on the BSID of the SR Policy




Automated Steering – Setting Color of a Route


- The color extended community is specified in **RFC 5512** and updated in *draft-previdi-idr-segment-routing-te-policy*
- The color of a BGP route is typically set at the egress PE by adding a color extended community to the route
 - The color extended community is propagated to the ingress PE
 - Traffic steering on the ingress PE is then done automatically based on the color

Automated Steering – Color Assignment on PE

- Node1 has two SR Policies with end-point Node3:
 - POL10 with color 10 (blue) via Node2
 - POL20 with color 20 (red) via Node4
- Node3 advertises two prefixes with next-hop 1.1.1.3 in BGP:
 - 10.10.10.0/24 with color 10 (blue)
 - 10.20.20.0/24 with color 20 (red)



Automated Steering – Color Assignment Egress PE



Node3

```
extcommunity-set opaque BLUE
  10
end-set
!
extcommunity-set opaque RED
  20
end-set
!
route-policy SET_COLOR
  if destination in (10.10.10.0/24) then
    set extcommunity color BLUE
  endif
  if destination in (10.20.20.0/24) then
    set extcommunity color RED
  endif
end-policy
!
router bgp 1
  neighbor 1.1.1.1
  remote-as 1
  update-source Loopback0
  address-family ipv4 unicast
  route-policy SET_COLOR out
```

On-Demand Next-Hop (ODN)



You make multi-cloud **possible**

On-Demand Next-Hop (ODN)

- A head-end router **automatically instantiates an SR Policy** to a BGP next-hop when required (on-demand)
- Color community is used as SLA indicator
- Reminder: an SR Policy is defined (color, end-point)

BGP Color
Community

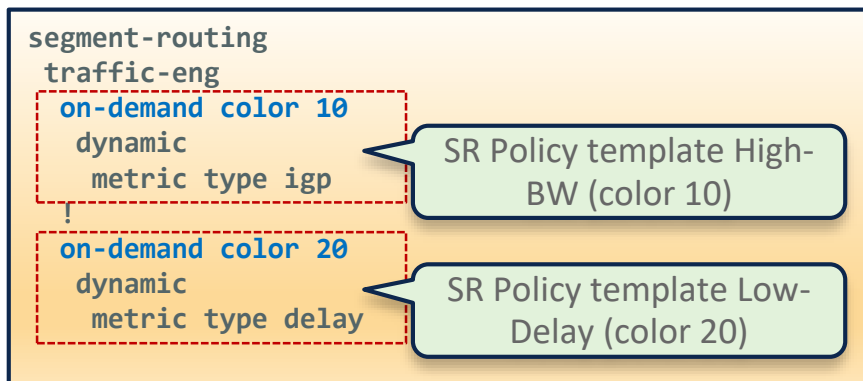


BGP
Next-hop

- **Automated Steering (AS)** automatically steers the BGP traffic into this SR Policy, also based on next-hop and color.

On-Demand SR Policy

- Configure an SR Policy template for each color for which on-demand SR Policy instantiation is desired
- An example with two color templates configured:
 - color 10 for high bandwidth (optimize IGP metric)
 - color 20 for low-delay (optimize link-delay metric)




ODN Benefits

- SLA-aware BGP service
- No a-priori full-mesh of SR Policy configuration
 - 3 to 4 common optimization templates are used throughout the network
 - color → optimization objective
- No complex steering configuration
 - Automated Steering of BGP routes on the right SLA path
 - Data plane performant
 - BGP PIC FRR data plane protection is preserved
 - BGP NHT fast control plane convergence is preserved

ODN Summary

The ODN functionality is only triggered when receiving a **service route** with an **authorized color**. A color is authorized when an on-demand template is configured for that color.



SR-PCE

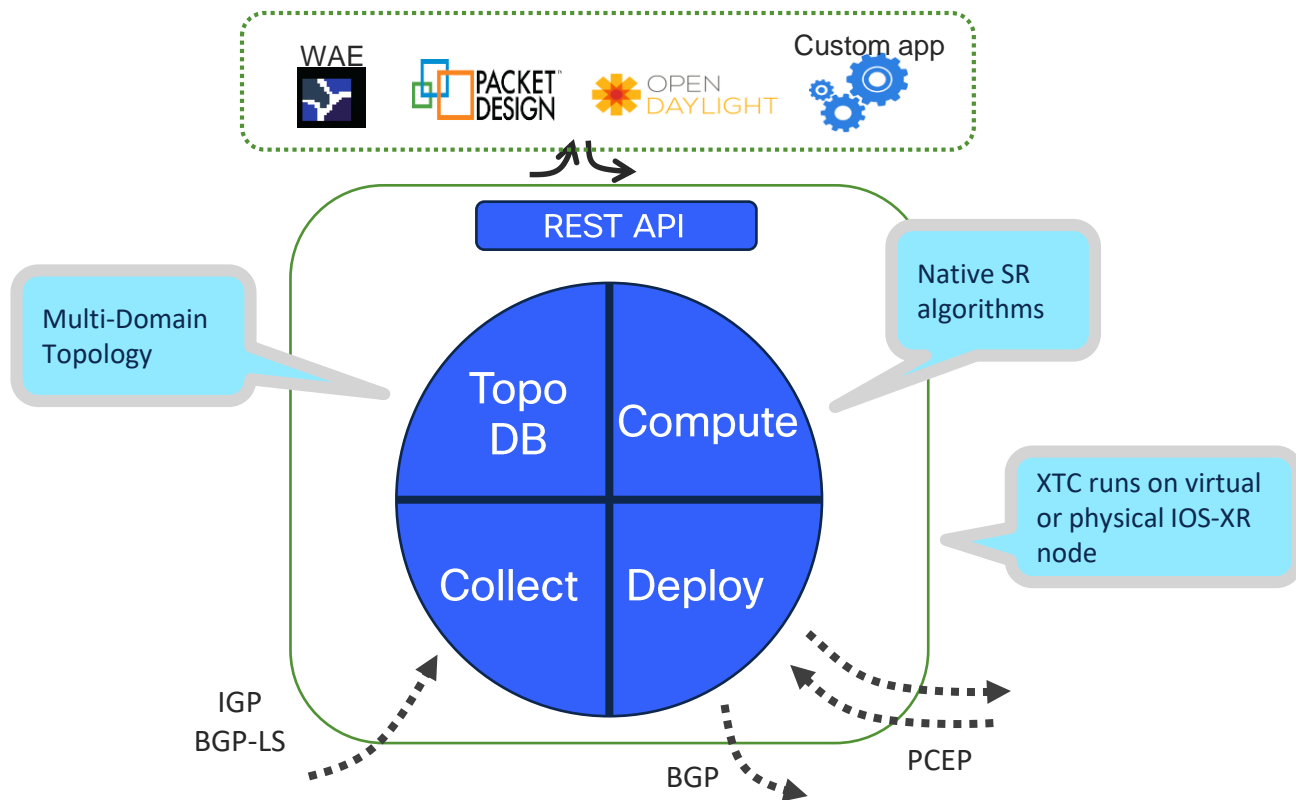


You make customer experience **possible**

What about inter-domain path calculation?

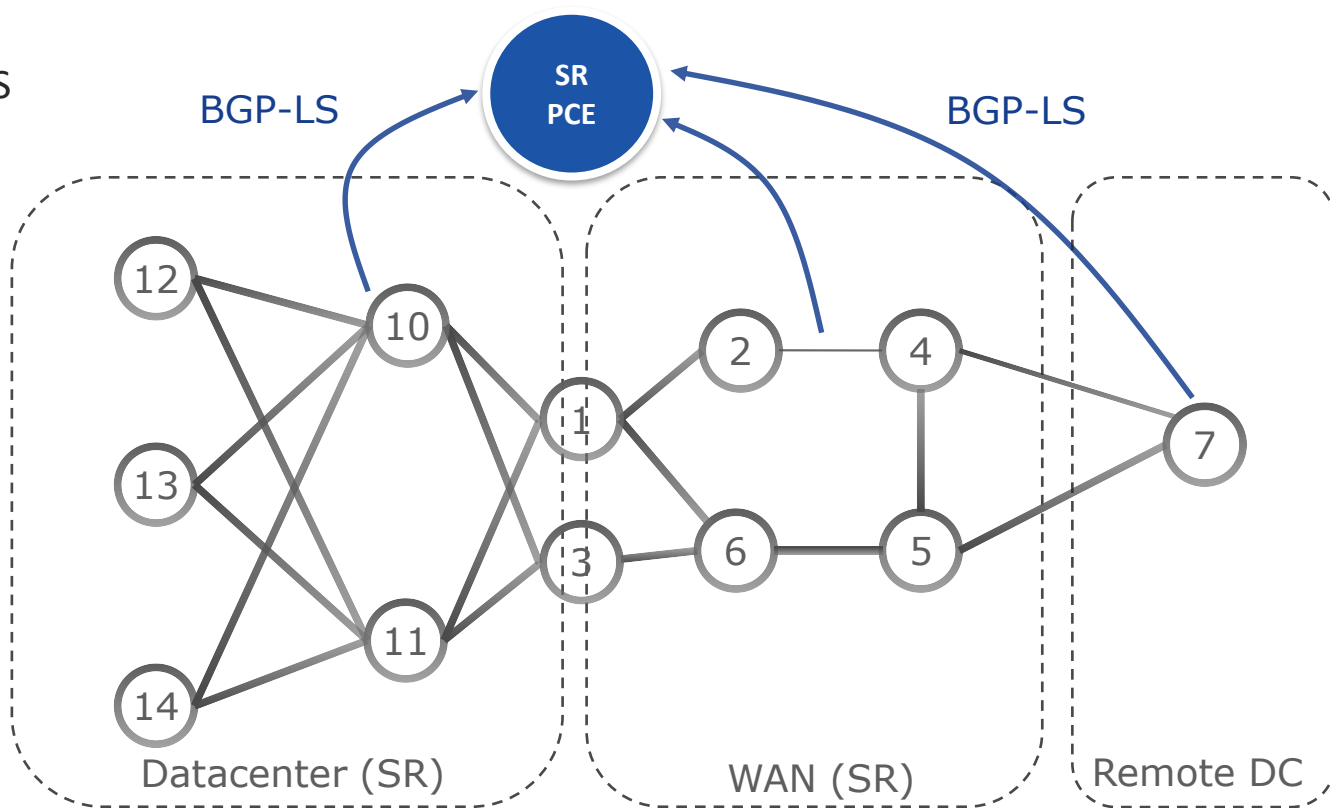
- SR-PCE is an IOS XR multi-domain stateful SR Path Computation Element (PCE)
 - **IOS XR:** XTC functionality is available on any physical or virtual IOS XR node, activated with a single configuration command
 - **Multi-domain:** Real-time reactive feed via BGP-LS/ISIS/OSPF from multiple domains; computes inter-area/domain/AS paths
 - **Stateful:** takes control of SRTE Policies, updates them when required

SR-PCE Building Blocks



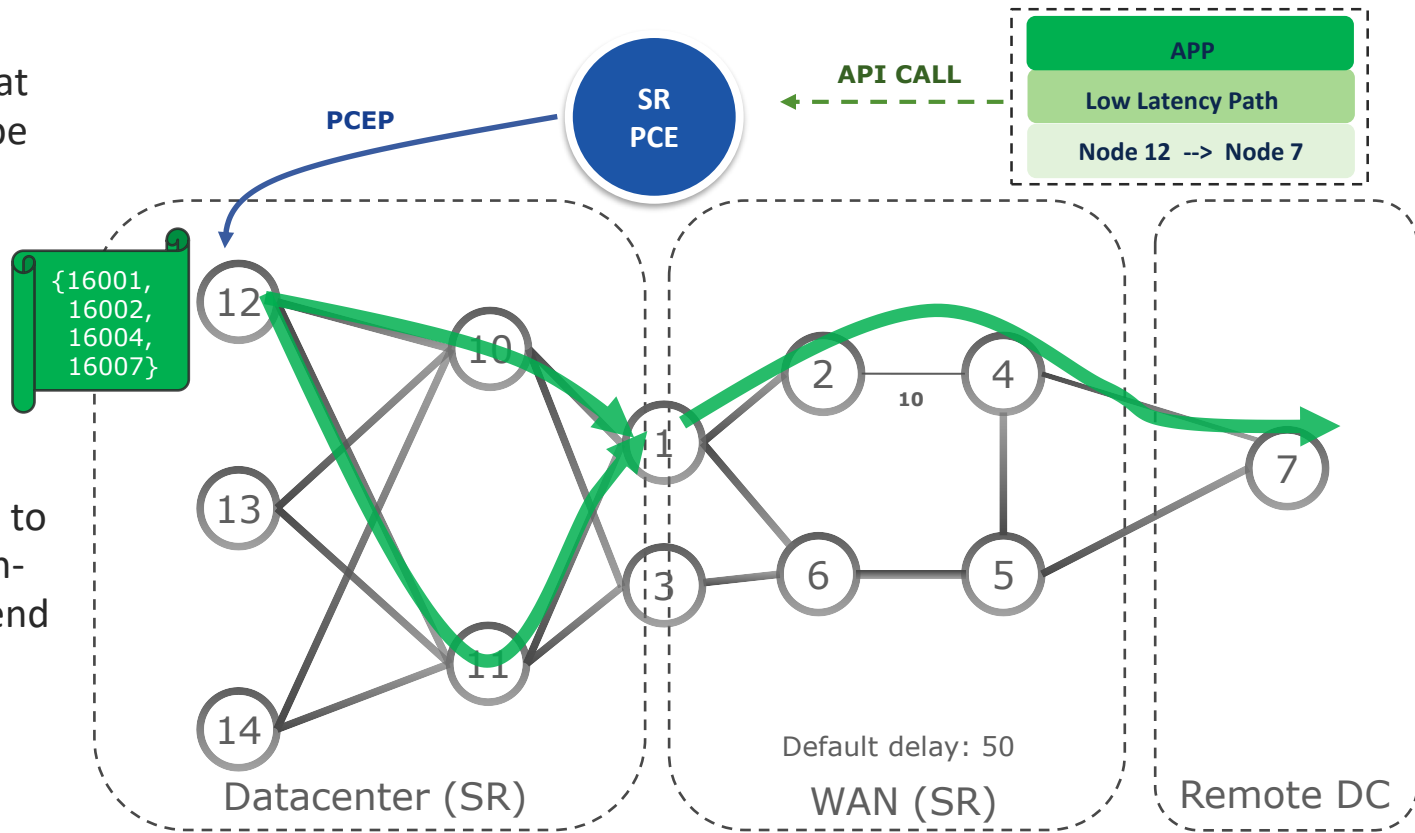
SR-PCE Controller

- SR PCE collects via BGP-LS
 - IGP segments
 - BGP segments
 - Topology



An End-to-end Path as a List of Segments (1)

- SR PCE computes that the green path can be encoded as
 - 16001
 - 16002
 - 16004
 - 16007
- SR PCE programs a single per-flow state to create an application-engineered end-to-end policy



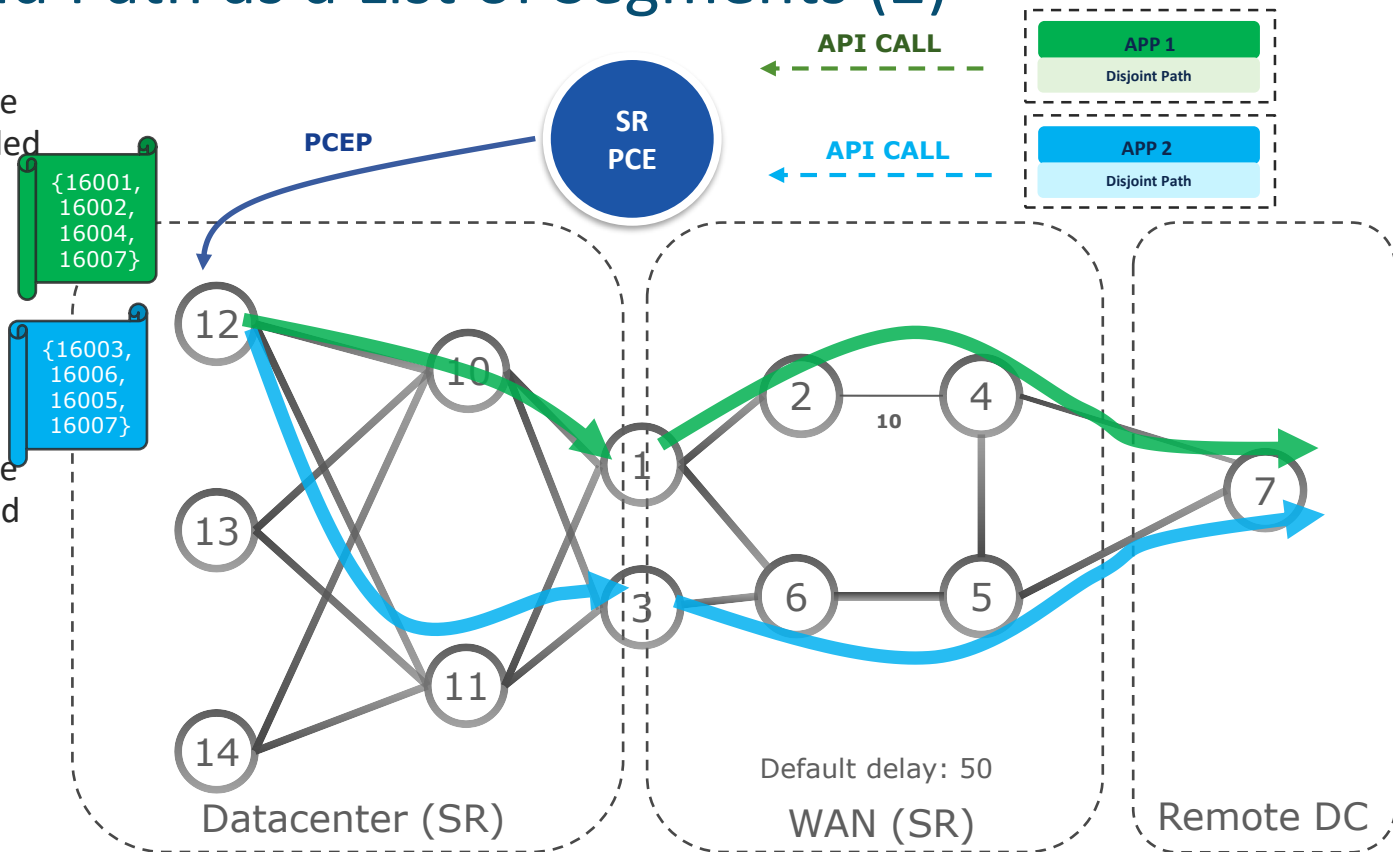
An End-to-end Path as a List of Segments (2)

- SR PCE computes that the **green** path can be encoded as

- 16001
- 16002
- 16004
- 16007

- SR PCE computes that the **blue** path can be encoded as

- 16003
- 16006
- 16005
- 16007



Flexible Algorithm (Flex-Algo)



You make networking **possible**

Flex-Algo

- **Complements** the SRTE solution by adding new Prefix-Segments with specific optimization objective and constraints
 - Minimize igp-metric or delay or te-metric
 - Avoid SRLG or affinity
- **Leverages** the SRTE benefits of simplicity and automation:
 - Automated sub-50msec FRR (TILFA)
 - On-Demand Policy (ODN)
 - Automated Steering (AS)

Flex-Algo Definition



- The definition of an “**Algorithm**” is defined by the operator, on a per-deployment basis
- Flex-Algo ‘**K**’ is defined as
 - The specification of minimal metric: IGP, delay, ...
 - The exclusion of certain link properties: link-affinity, SRLG, ...
- Each node **MUST** advertise Flex-Algo(s) that it is participating in

Currently Defined Algorithms


- **0: Shortest Path First (SPF) algorithm based on link metric.**
 - This is the well-known shortest path algorithm as computed by the IS-IS Decision process. Consistent with the deployed practice for link-state protocols, algorithm 0 permits any node to overwrite the SPF path with a different path based on local policy
- **1: Strict Shortest Path First (SPF) algorithm based on link metric.**
 - The algorithm is identical to algorithm 0 but algorithm 1 requires that all nodes along the path will honor the SPF routing decision. **Local policy MUST NOT alter the forwarding decision computed by algorithm 1** at the node claiming to support algorithm

Flex-Algo Definition

- Example:
 - Operator-1 defines Flex-Algo 128 as “**minimize IGP metric and avoid link-affinity ‘green’**”
 - Operator-2 defines Flex-Algo 128 as “**minimize delay metric**”




Use Case: Multi-Plane Networks




- All the nodes support Algo 0: minimize IGP metric
- Green nodes also support 128: minimize IGP metric
- Red nodes also support 129: minimize Delay

Use Case: Multi-Plane Networks (cont)




- Path to Node 9 across Algo 0
- ECMP based forwarding across all Planes

Use Case: Multi-Plane Networks (cont)



- Path to Node 9 across **Flex-Algo 128**
- ECMP based forwarding WITHIN **green** Plane




Use Case: Multi-Plane Networks (cont)



- Path to Node 9 across **Flex-Algo 129**
- ECMP based forwarding **WITHIN red Plane**




Use Case: Delay vs Cost of Transport

- All nodes support Algo 0 & 128
- ISIS link metric 10
- Algo 128: minimize delay metric
- Per-link measurement of delay and advertisement as delay metric via ISIS
- Delay metric at that time shown in green



Use Case: SRTE for Secure Paths

- ISIS link metric 10
- Link colors shown **Unencrypted** / **Encrypted**
- All nodes support Algo 0 & 128
- Algo 128: minimize IGP while traversing links with encryption enabled (**exclude brown**)
- Per-link colors flooded in IGP



Flex-Algo Benefits

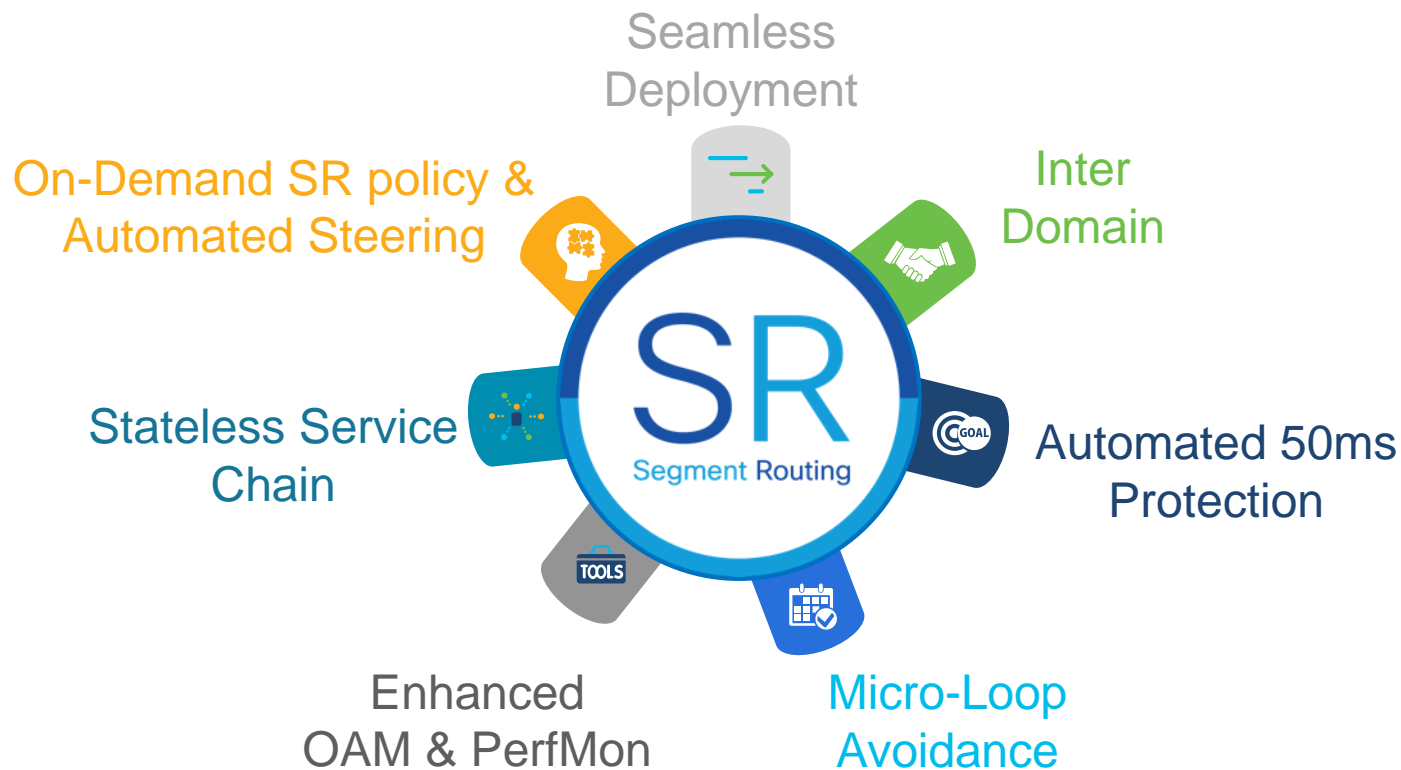
- **Minimize label stack, and maximize ECMP at midpoints:**
 - Single SID needed to enforce traffic towards a specific Path/Plane
 - ECMP if possible (cost related) on that specific Path/Plane
- **Slicing of the network on a per:**
 - Latency
 - Bandwidth
 - Secure links – MACSec
- **TI-LFA aware**
 - Protected path stays in Flex-Algo virtual topology

Conclusion



You make networking **possible**

SR Unified Fabric Attributes



Industry at Large Backs up SR



Strong customer adoption

WEB, SP, DC, Metro,
Enterprise



De-facto SDN Architecture



Standardization IETF



Multi-vendor Consensus



Open Source Linux, VPP



Thank you





You make **possible**