



You make **possible**



My Machine is an Honor Student at Cisco!

Quality of machine learning experience

Mike Mikhail, Architect

mamikhai@cisco.com

@MikeMikhail

DEVLIT-4002



Barcelona | January 27-31, 2020



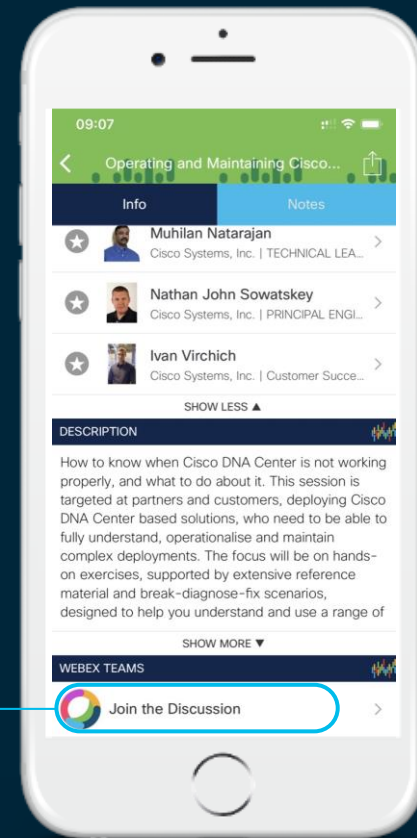
Cisco Webex Teams

Questions?

Use Cisco Webex Teams to chat with the speaker after the session

How

- 1 Find this session in the Cisco Events Mobile App
- 2 Click “Join the Discussion”
- 3 Install Webex Teams or go directly to the team space
- 4 Enter messages/questions in the team space



Agenda

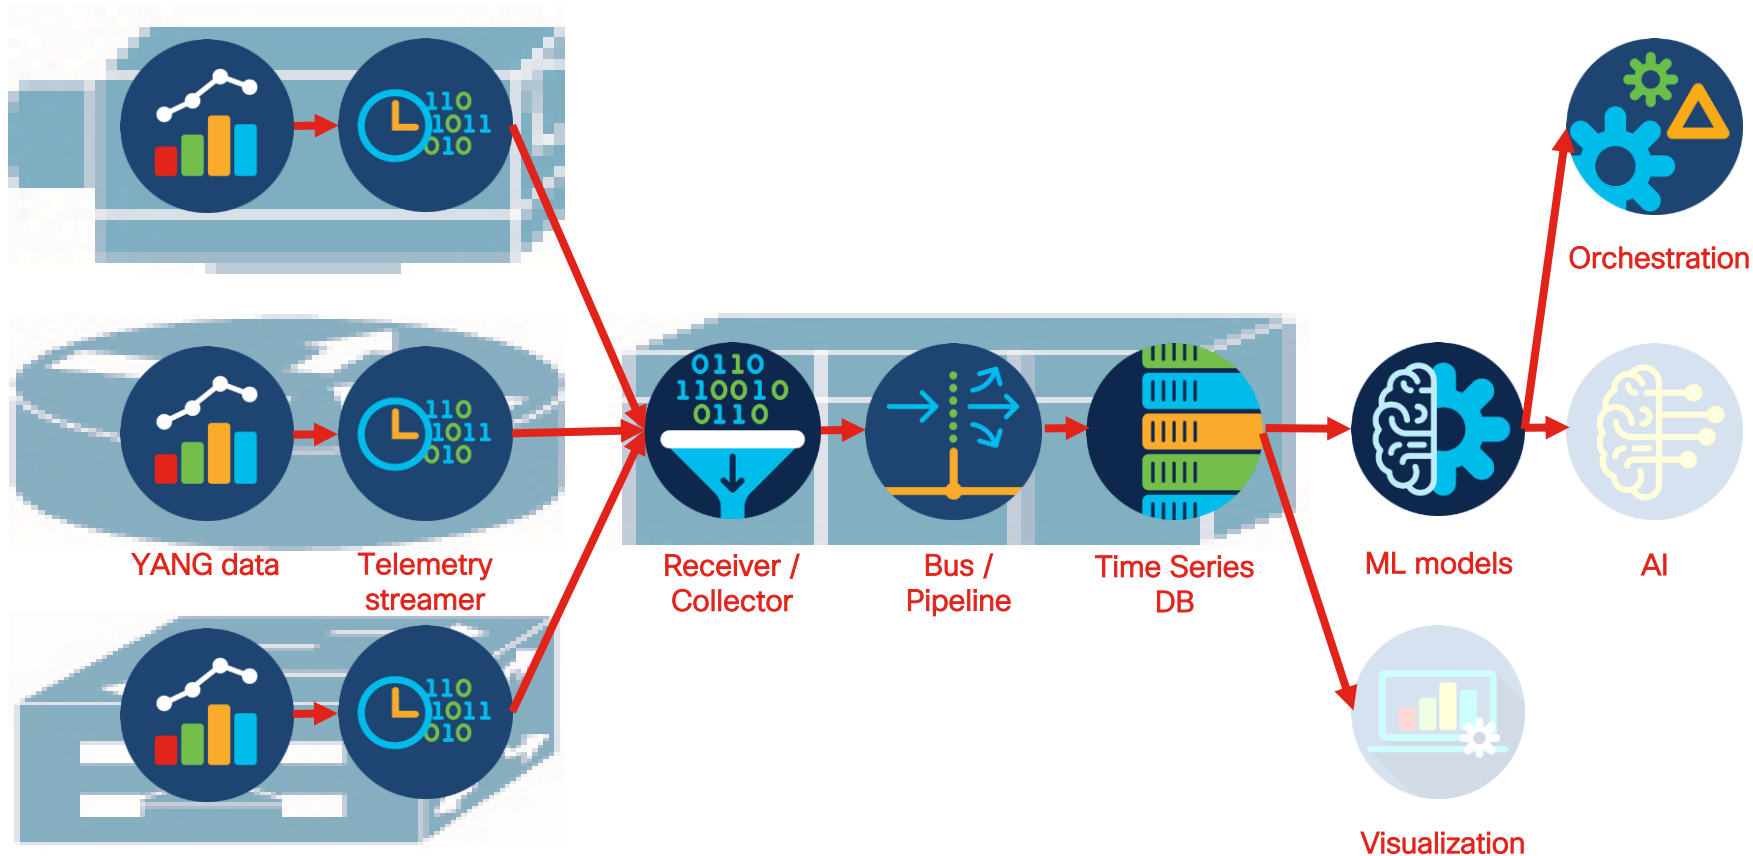
- Data analysis and quality
- The process
- Data conditioning
- Visualization
- Quality assurance






Initiate ML model, initial training...

Production Line

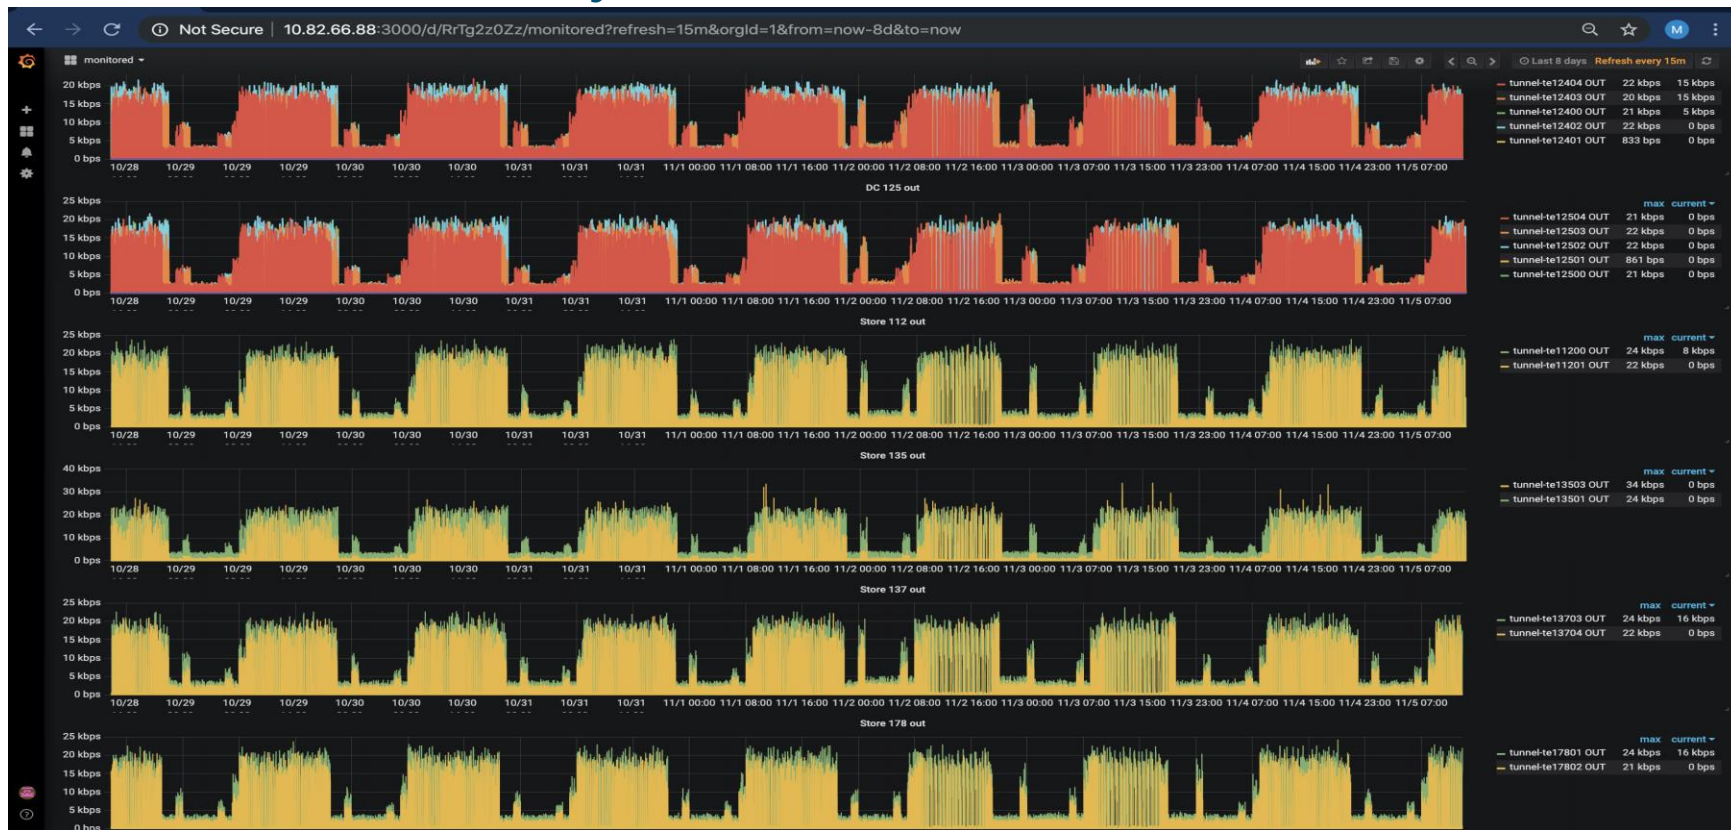


Data Source

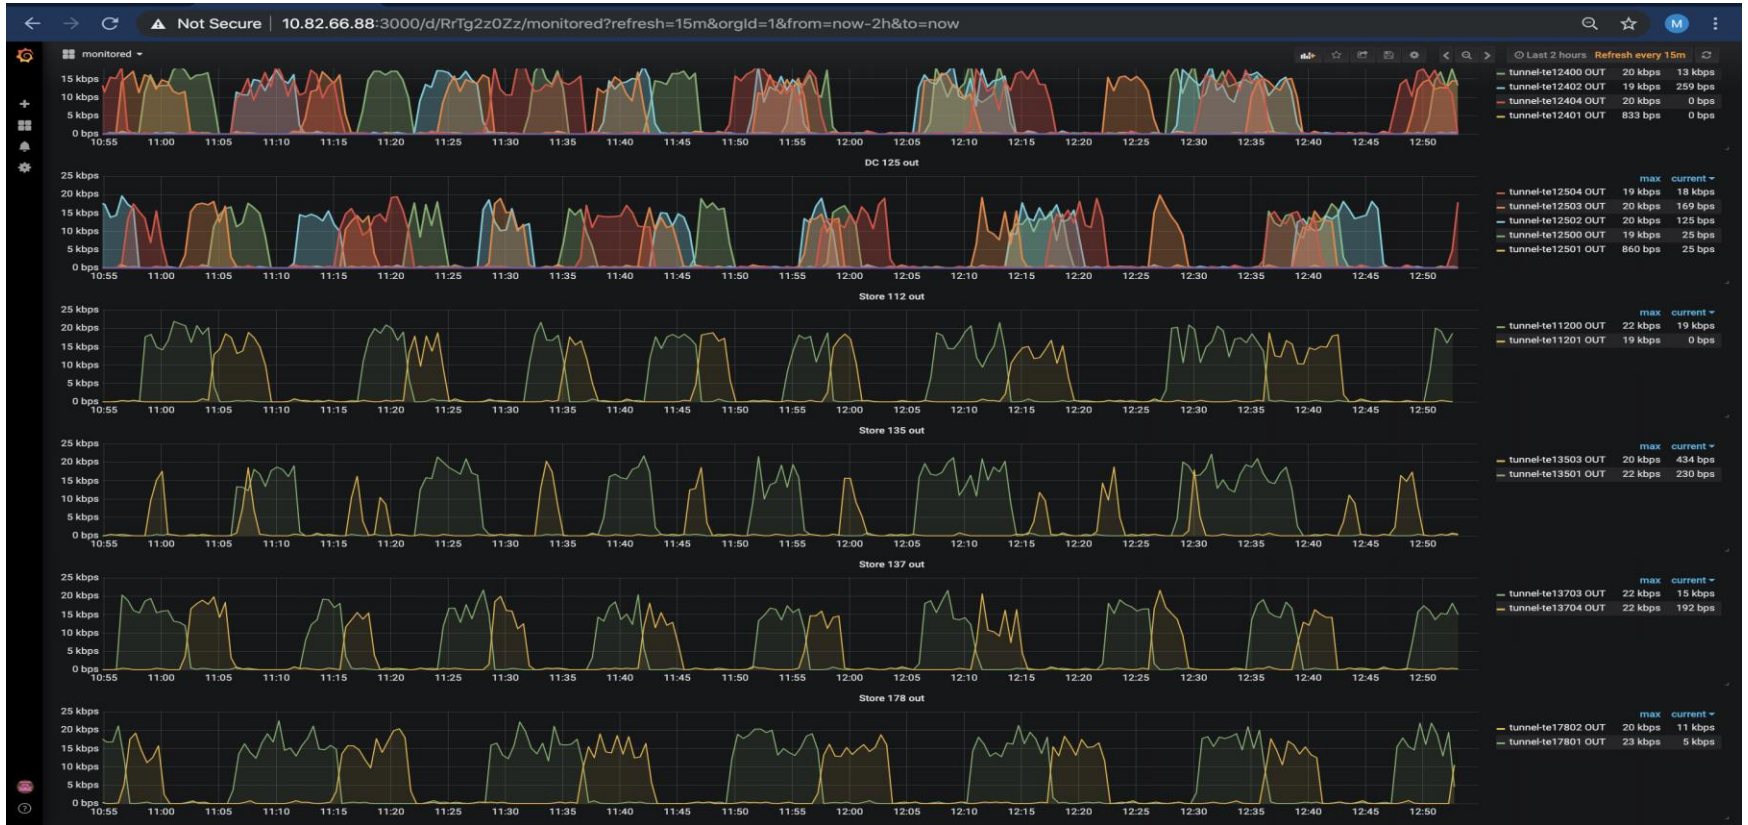
Business app data



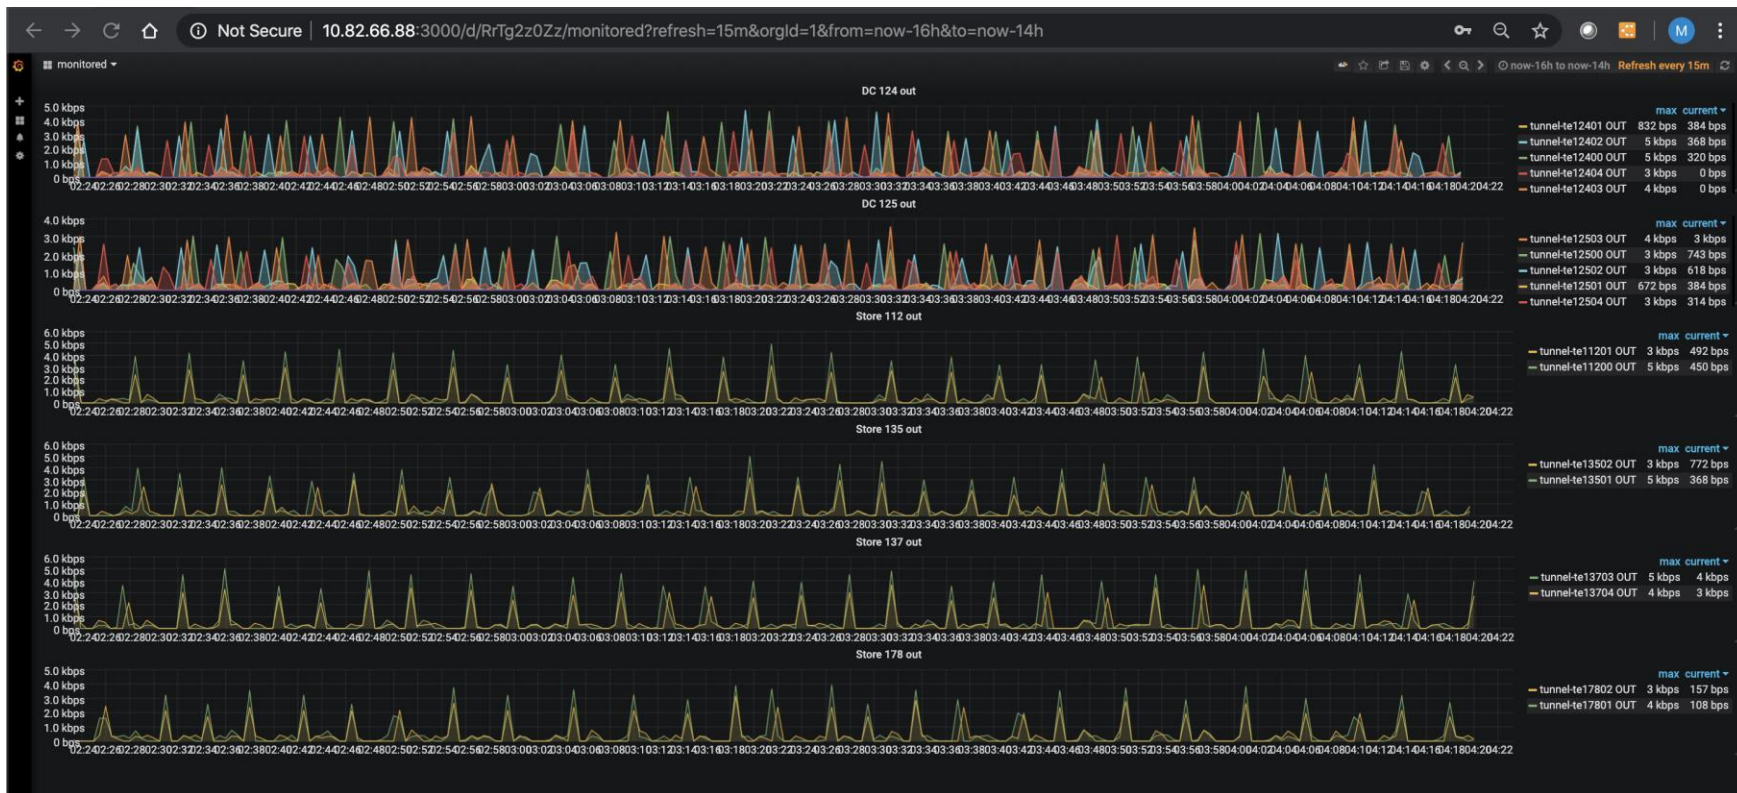
Visualized – 8days



Visualized – 2 business hours

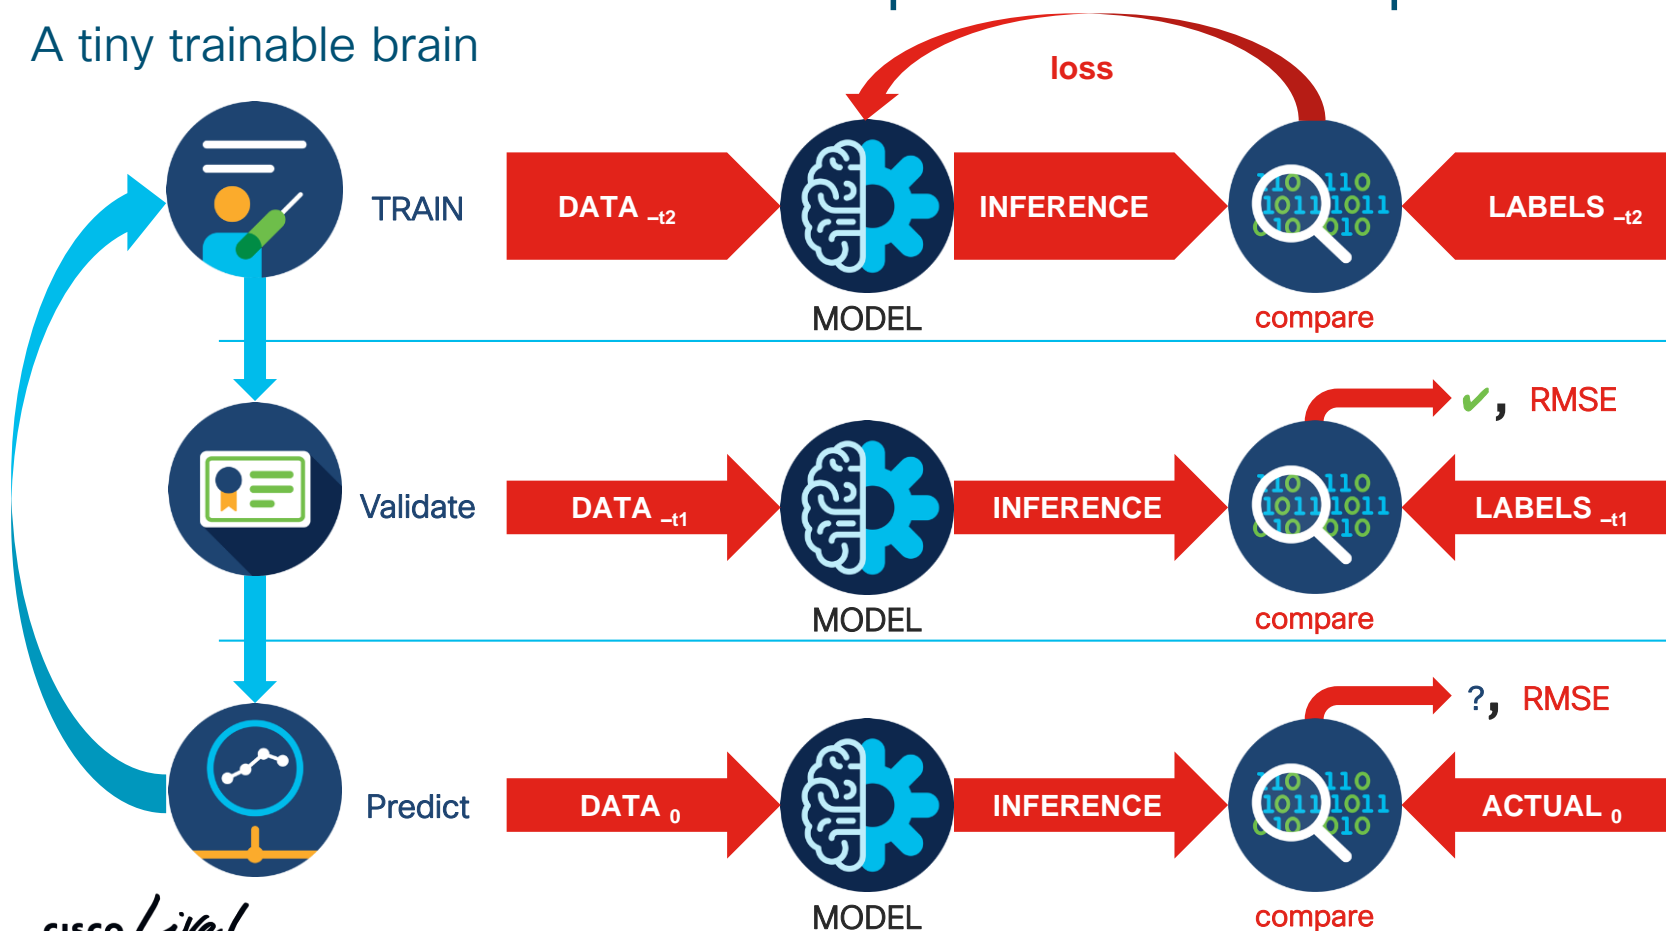


Visualized – 2 off hours




What's an ML Model? Supervised example

A tiny trainable brain



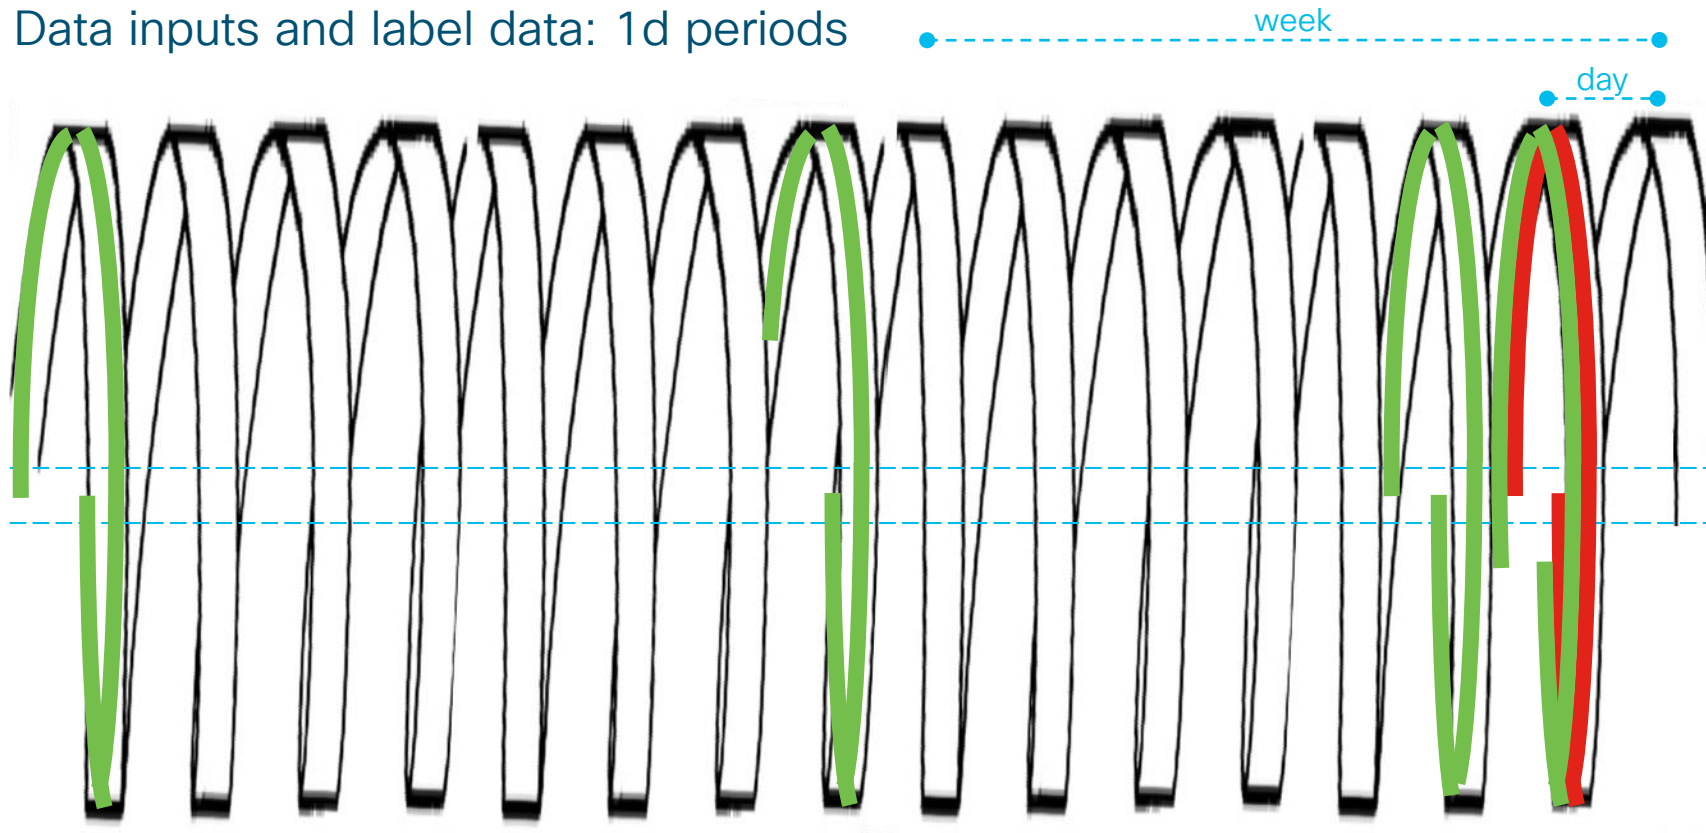
Basis for Inference

Data inputs and validation data




Training: Larger Datasets

Data inputs and label data: 1d periods




Training

Data inputs and label data: -1d



Training: Initial, and Periodic; Predict

“Monitor” flowchart



The Data in 4x 24-hour DataFrame, [4x18]x2880



4 24-hour_period x 18 path x 2880 30-second_field = 72 column x 2880 row DataFrame

DataFrame

```
def read_train_long(record_count, label_prefix, verbose=True):
```

```
    print('\ntraining long data')
    print(train.describe())
```

training long data

	d_te11200_previous	d_te11200_1d	d_te11200_1w	d_te11200_2w
count	2.880000e+03	2.880000e+03	2.880000e+03	2.880000e+03
mean	1.093179e+07	1.210950e+07	1.215980e+07	2.038959e+07
std	1.072837e+07	1.140074e+07	1.141777e+07	1.141292e+07
min	0.000000e+00	1.365000e+03	1.087200e+04	2.841000e+03
25%	1.663094e+06	1.665956e+06	1.689180e+06	1.086874e+07
50%	5.952412e+06	8.211694e+06	8.303540e+06	2.399129e+07
75%	1.959832e+07	2.178496e+07	2.212936e+07	3.107957e+07
max	3.269652e+07	3.305397e+07	3.308679e+07	3.262037e+07

Column "label"

df is 72 x 2880

Math description
of column

	d_te11201_previous	d_te11201_1d	d_te11201_1w	d_te11201_2w
count	2.880000e+03	2.880000e+03	2.880000e+03	2.880000e+03
mean	7.527995e+06	8.306467e+06	8.337991e+06	1.384803e+07
std	7.289889e+06	7.747967e+06	7.757351e+06	7.804774e+06
min	0.000000e+00	1.380000e+03	8.182000e+03	1.001000e+03

DataFrame: same format, smaller size

validation data

	d_te11200_previous	d_te11200_1d	d_te11200_1w	d_te11200_2w \
count	1.200000e+02	1.200000e+02	1.200000e+02	1.200000e+02
mean	1.358485e+06	9.795289e+05	1.179847e+06	1.246457e+06
std	8.000545e+05	6.292864e+05	6.797141e+05	7.001280e+05
min	1.380000e+03	7.731200e+04	9.600000e+01	1.928000e+03
25%	4.905140e+05	3.478915e+05	6.696790e+05	8.413410e+05
50%	1.362629e+06	1.087073e+06	1.240100e+06	1.377682e+06
75%	2.190353e+06	1.559045e+06	1.827315e+06	1.777683e+06
max	2.840542e+06	2.067730e+06	2.336623e+06	2.262407e+06

df is 72 x 120

	d_te11201_previous	d_te11201_1d	d_te11201_1w	d_te11201_2w \
count	1.200000e+02	1.200000e+02	1.200000e+02	1.200000e+02
mean	7.821397e+05	1.026548e+06	7.012429e+05	7.430594e+05
std	5.078806e+05	4.558566e+05	4.654541e+05	5.078973e+05
min	0.000000e+00	0.000000e+00	0.000000e+00	1.350000e+03
25%	3.321160e+05	5.965380e+05	4.511570e+05	3.695710e+05
50%	9.229315e+05	9.878900e+05	8.389290e+05	7.320545e+05
75%	1.225480e+06	1.511388e+06	1.236286e+06	1.205447e+06
max	1.499793e+06	1.857222e+06	1.556056e+06	1.526029e+06

Fetching, Formatting, Conditioning the Data

```
def read_data(field_key, measurement_name, condition1, condition2, condition3, limit,
label):
    query_db = str('SELECT "%s" FROM "%s" WHERE %s AND %s AND %s LIMIT %d ' % (field_key,
measurement_name, condition1, condition2, condition3, limit+1))
    data_db = client.query(query_db)
    print('\ndata_db:\n', data_db)
    data_df = pd.DataFrame(data_db[str(measurement_name)])
    print('\ndata_df:\n', data_df)
    print('\ndata_df description:\n', data_df.describe())
    data_df.columns = [label]
    data_df.reset_index(drop=True, inplace=True)
    data_df.fillna(method='ffill', inplace=True)
    data_df.fillna(method='bfill', inplace=True)
    data_df -= data_df.min()
    data_df.drop(data_df.index[0], inplace=True)
    print('\ndata_df:\n', data_df)
    print('\ndata_df description:\n', data_df.describe())
    sys.exit()
    # data_df = data_df.sub(data_df.shift(fill_value=0))
    # print('\n', query_db, '\n', data_df.describe())
    return data_df
```

Raw Data

```
data_db:
  defaultdict(<class 'list'>, {'Cisco-IOS-XR-infra-statsd-oper:infra-
statistics/interfaces/interface/latest/generic-
counters':
                                bytes-sent
2020-01-03 06:34:04.032000+00:00 2269840208
2020-01-03 06:34:34.031000+00:00 2269840208
2020-01-03 06:35:04.027000+00:00 2269840208
2020-01-03 06:35:34.031000+00:00 2269840208
2020-01-03 06:36:04.032000+00:00 2269840820
2020-01-03 06:36:34.041000+00:00 2269842350
2020-01-03 06:37:04.032000+00:00 2269842554
2020-01-03 06:37:34.036000+00:00 2269842554
2020-01-03 06:38:04.032000+00:00 2269842554
2020-01-03 06:38:34.033000+00:00 2269842554
2020-01-03 06:39:04.090000+00:00 2269842554
2020-01-03 06:39:34.035000+00:00 2269853583
2020-01-03 06:40:04.036000+00:00 2269854129
2020-01-03 06:40:34.033000+00:00 2269854129
2020-01-03 06:41:04.035000+00:00 2269854228
2020-01-03 06:41:34.033000+00:00 2269855713
2020-01-03 06:42:04.035000+00:00 2269856604
```

In Single Column DataFrame

data_df:

		bytes-sent
2020-01-03	06:34:04.032000+00:00	2269840208
2020-01-03	06:34:34.031000+00:00	2269840208
2020-01-03	06:35:04.027000+00:00	2269840208
2020-01-03	06:35:34.031000+00:00	2269840208
2020-01-03	06:36:04.032000+00:00	2269840820
2020-01-03	06:36:34.041000+00:00	2269842350
2020-01-03	06:37:04.032000+00:00	2269842554
2020-01-03	06:37:34.036000+00:00	2269842554
2020-01-03	06:38:04.032000+00:00	2269842554
2020-01-03	06:38:34.033000+00:00	2269842554
2020-01-03	06:39:04.090000+00:00	2269842554
2020-01-03	06:39:34.035000+00:00	2269853583
2020-01-03	06:40:04.036000+00:00	2269854129
2020-01-03	06:40:34.033000+00:00	2269854129
2020-01-03	06:41:04.035000+00:00	2269854228
2020-01-03	06:41:34.033000+00:00	2269855713
2020-01-03	06:42:04.035000+00:00	2269856604
2020-01-03	06:42:34.034000+00:00	2269866268
2020-01-03	06:43:04.034000+00:00	2269866916

data_df description:

	bytes-sent
count	2.881000e+03
mean	2.284429e+09
std	1.233722e+07
min	2.269840e+09
25%	2.271348e+09
50%	2.283434e+09
75%	2.296978e+09
max	2.303451e+09

Column Label Changed, Baselined, Conditioned

data_df:

d_te11200_previous

1	0
2	0
3	0
4	612
5	2142
6	2346
7	2346
8	2346
9	2346
10	2346
11	13375
12	13921
13	13921
14	14020
15	15505
16	16396
17	26060
18	26708
19	28328

data_df description:

d_te11200_previous

count	2.880000e+03
mean	1.459358e+07
std	1.233636e+07
min	0.000000e+00
25%	1.507463e+06
50%	1.362955e+07
75%	2.713798e+07
max	3.361109e+07

Constructing Multi-Column DataFrame


```
def read_train(record_count, label_prefix, verbose=True):
    for interface in tunnel_ifs:
        query_if = str('("interface-name" = \'%s\')' % (interface))
        label = str(label_prefix + interface[-7:] + "_previous")
        read_if = read_data('bytes-sent', 'Cisco-IOS-XR-infra-statsd-oper:infra-
statistics/interfaces/interface/latest/generic-counters', query_if, 'time >= now() - {} - 2h -
1m'.format(previous), 'time <= now()', record_count, label)
        if interface == tunnel_ifs[0]:
            train = read_if
        else:
            train = pd.concat([train, read_if], axis=1, sort=False)
            label = str(label_prefix + interface[-7:] + "_1d")
            read_if = read_data('bytes-sent', 'Cisco-IOS-XR-infra-statsd-oper:infra-
statistics/interfaces/interface/latest/generic-counters', query_if, 'time >= now() - {} - 1d - 1h -
1m'.format(previous), 'time <= now()', record_count, label)
            train = pd.concat([train, read_if], axis=1, sort=False)
            label = str(label_prefix + interface[-7:] + "_2w")
            read_if = read_data('bytes-sent', 'Cisco-IOS-XR-infra-statsd-oper:infra-
statistics/interfaces/interface/latest/generic-counters', query_if, 'time >= now() - 3w - {} - 1h -
1m'.format(previous), 'time <= now()', record_count, label)
            train = pd.concat([train, read_if], axis=1, sort=False)

    train.fillna(method='ffill', inplace=True)
    return train
```

Unique column label

Concatenate columns

The Model



The Neural Network, and Sample Cycle

```
. my_optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)
my_optimizer = tf.contrib.estimator.clip_gradients_by_norm(my_optimizer, 5.0)
dnn_regressor = tf.estimator.DNNRegressor(
    feature_columns=construct_feature_columns(training_examples),
    hidden_units=hidden_units,
    optimizer=my_optimizer,
    model_dir= model_directory,
    label_dimension= len(tunnel_ifs) + len(physical_ifs)
.
hidden_units = [72, 36, 18] # probably an overkill for our small scale
.
cycle += 1
print('cycle number ', cycle)
dnn_regressor = train_nn_regression_model(
    learning_rate = 0.0003,
    steps = 1000,
    batch_size = 120,
    hidden_units = hidden_units,
    training_examples = read_train(120, 'd_'),
    training_targets = read_train_target(120, 'l_'),
    validation_examples = read_validate(120, 'd_'),
    validation_targets = read_last_target(120, 'v_'),
    prediction = True
.
```

Set once, for
a new model

Can change
every call

+ Normalization Layer

```
def read_train_long(record_count, label_prefix, verbose=True):
    global feature_mean
    global feature_std
    global feature_max
    .
    if feature_mean == 0:
        feature_mean = train.mean().mean()
        print('feature mean: ', feature_mean)
    if feature_std == 0:
        feature_std = train.std().mean()
        print('feature std: ', feature_std)
    if feature_max == 0:
        feature_max = train.max().mean() / 24 # The mean max per 1 hour
        print('feature max: ', feature_max)
    .
def construct_feature_columns(input_features):
    .
    # epsilon = 0.000001
    .
    # choose best normalization of input data
    .
    return set([tf.feature_column.numeric_column(my_feature, normalizer_fn=lambda val: (val) /
(feature_max))
                for my_feature in input_features])
```

The data - inputs

feature mean: 7692395.787210648
feature std: 7084487.678199986
feature max: 899857.4594907407

Data indicators,
Useful for normalization

training long data

	d_te11200_previous	d_te11200_1d	d_te11200_1w	d_te11200_2w
count	2.880000e+03	2.880000e+03	2.880000e+03	2.880000e+03
mean	1.460392e+07	1.148071e+07	1.377143e+07	1.286884e+07
std	1.208010e+07	1.079281e+07	1.288531e+07	1.193284e+07
min	5.194400e+04	8.214400e+04	7.731200e+04	1.128000e+03
25%	5.464170e+06	2.660394e+06	3.144108e+06	2.952690e+06
50%	7.655800e+06	5.238823e+06	6.525733e+06	6.685706e+06
75%	2.390421e+07	2.051300e+07	2.478536e+07	2.295742e+07
max	4.010818e+07	3.499195e+07	3.808687e+07	3.528133e+07

Column key

Record count

Min value

Max value

	d_te11201_previous	d_te11201_1d	d_te11201_1w	d_te11201_2w
count	2.880000e+03	2.880000e+03	2.880000e+03	2.880000e+03
mean	9.872998e+06	7.983878e+06	9.340476e+06	8.721320e+06
std	8.096162e+06	7.326994e+06	8.648336e+06	8.025967e+06
min	0.000000e+00	1.080000e+03	0.000000e+00	0.000000e+00
25%	3.781685e+06	2.003244e+06	2.256103e+06	2.102746e+06
50%	5.334513e+06	3.862060e+06	4.561317e+06	4.507384e+06

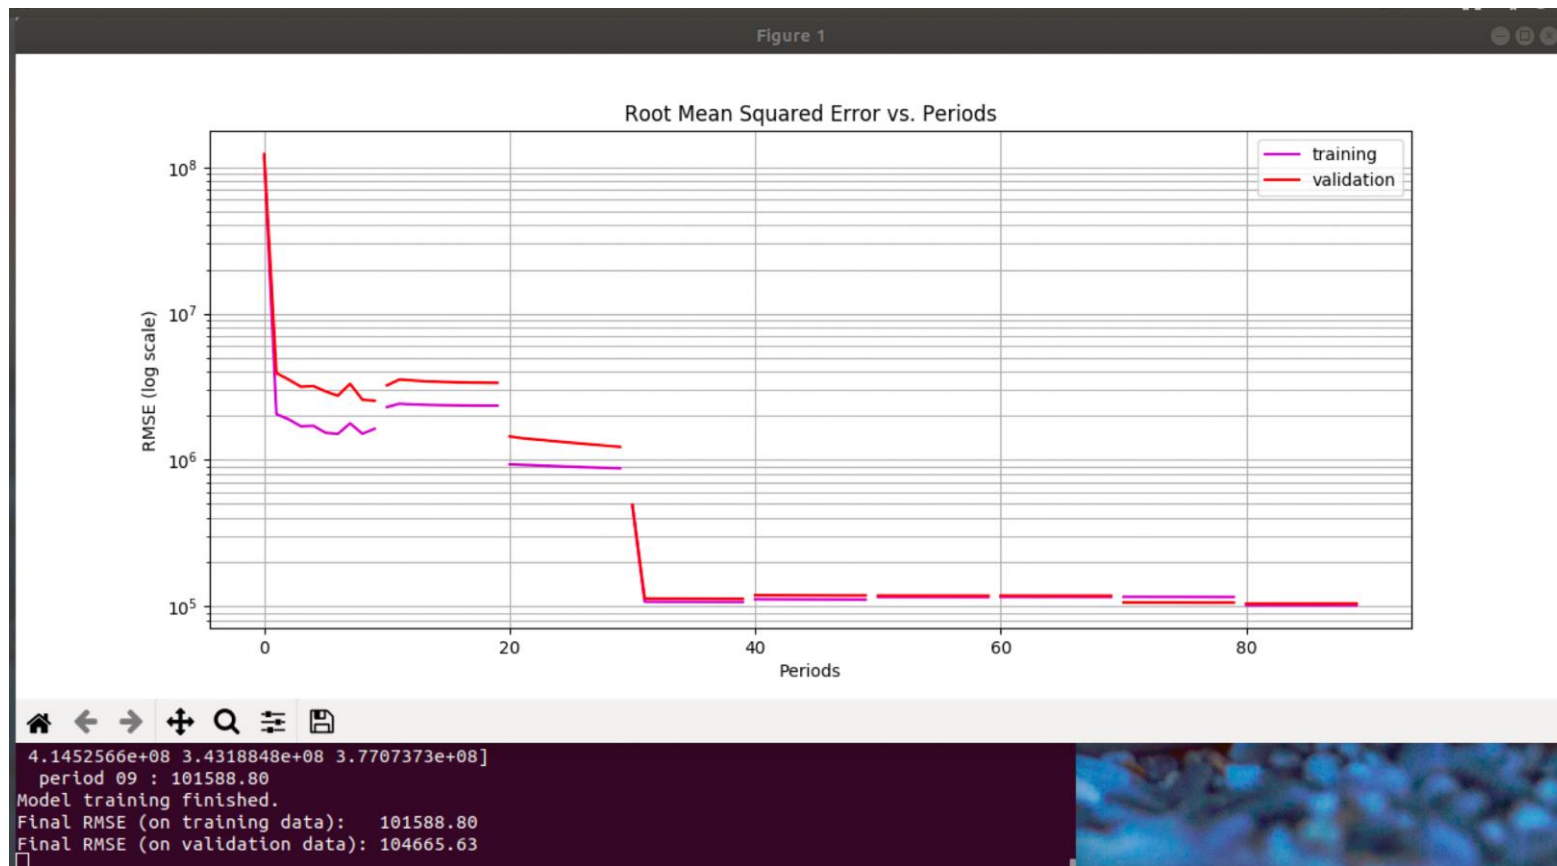
Check the Data: Model Output

```
.
training_predictions baoundaries
[ -9676.203      43134.777      30126.502      49435.23      -32687.848
  53458.71      -41498.527      2386.6956      -7947.478      841.848
  25334.666      -29982.818      -35702.715      39189.33      822.5124
  43476.05      48315.46      459.85422]
[2567748.5      1706653.8      1980904.5      2473544.2      2354506.8      2123372.
 2704984.5      1812864.2      2203519.      67233.766      1694452.8      1959683.9
 2268782.5      1549306.6      66137.98      2231304.2      1900274.9      1607212.8 ]
validation_predictions boundaries
[ 9514.773      37849.348      3103.1733      43236.957      -5364.3784
 22101.646      -17330.65      16654.14      8049.155      765.1315
 3570.572      -6731.2876      -15120.497      33285.02      766.89764
 37848.816      19647.463      13101.8      ]
[2846600.      2264547.2      3210291.      2034950.      3650229.2      1717540.4
 1242460.5      1751227.8      2441805.8      68177.71      2717810.5      3045913.5
 1010453.4      2042300.1      66662.13      1811712.9      1519572.6      1552536.8 ]
  period 08 : 55037.89
.
```

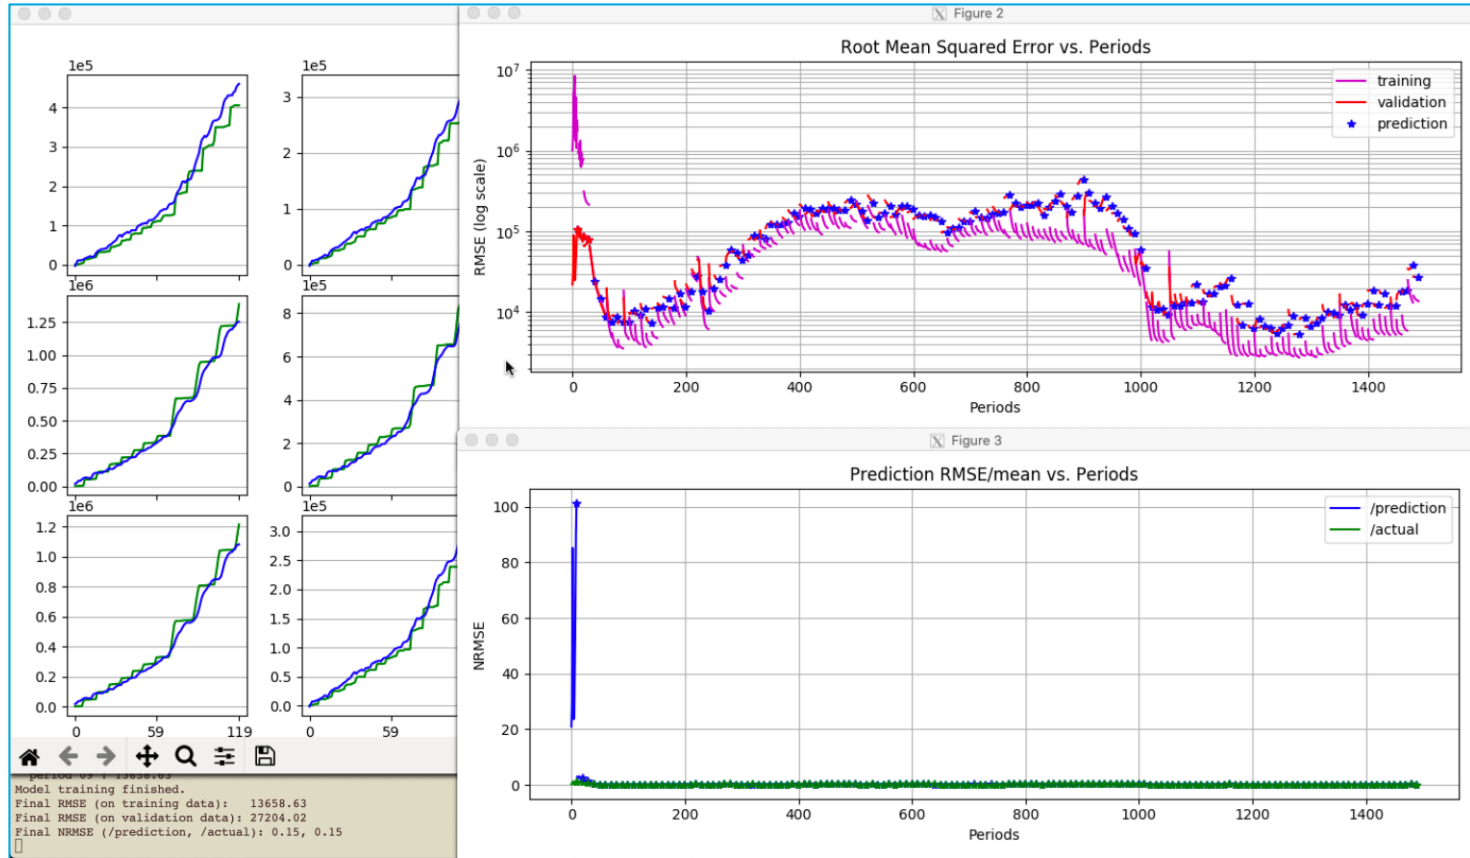
The Training

```
.
period 08 : 3582.59
training_predictions boundaries
[ 7208.3516    3251.541    7443.031    4104.243    4939.107    4307.9517
 -4889.6      -2828.935    5941.4014    726.0451    7325.6025    4610.5356
 -3598.1191    3449.1812    635.07806   4949.      4667.8926    -1990.8337 ]
[228423.4   177642.34 234673.56 181417.75 247542.95 188667.36 242869.14
 185446.6   211784.83 67781.5   213718.23 222849.58 222998.84 172893.55
 67602.49 174859.56 182980.92 178175.17]
validation_predictions boundaries
[ 6515.7734    3770.0752    8726.353    3760.6396    13467.836    3967.8716
 -5080.483    -3837.6772    5243.997    539.8233    8261.776    11545.85
 -3959.705    3852.7988    432.32617   4614.4375    4369.189    -3105.7751 ]
[193788.67   157614.84 150174.69 216238.94 187304.67 213302.8
 288468.47   197036.98 181182.31 51306.69 143677.8 166325.1
 265046.75   152532.97 51454.254 206665.23 205994.08 187957.39 ]
period 09 : 3548.03
Model training finished.
Final RMSE (on training data): 3548.03
Final RMSE (on validation data): 22812.22
Final NRMSE (/prediction, /actual): 0.25, 0.25
cycle number 788
```

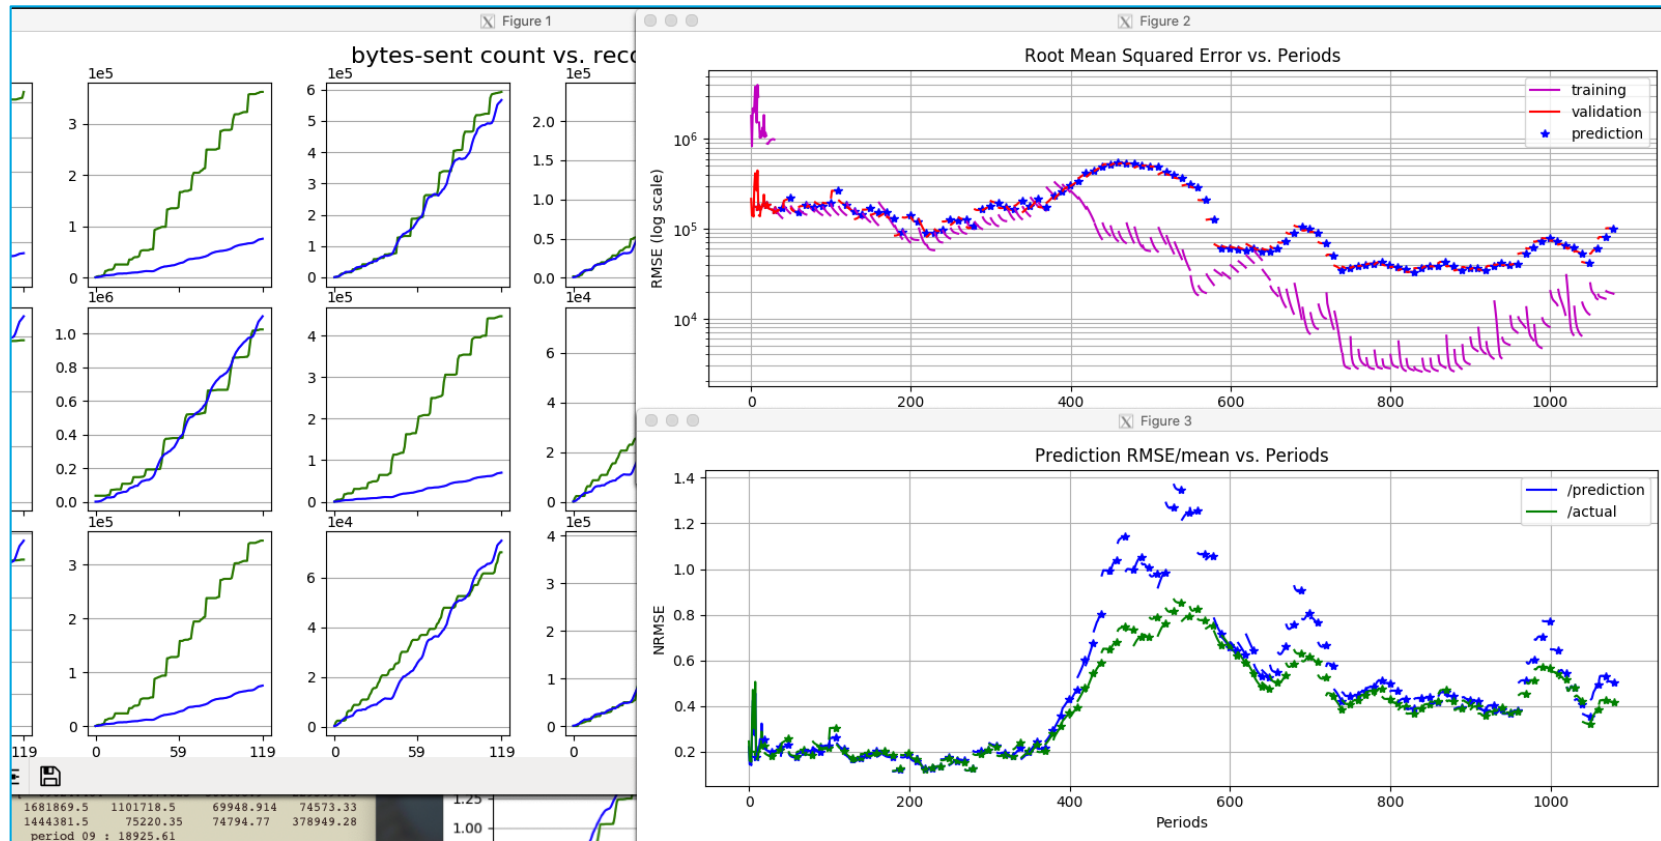
Training Makes Perfect!



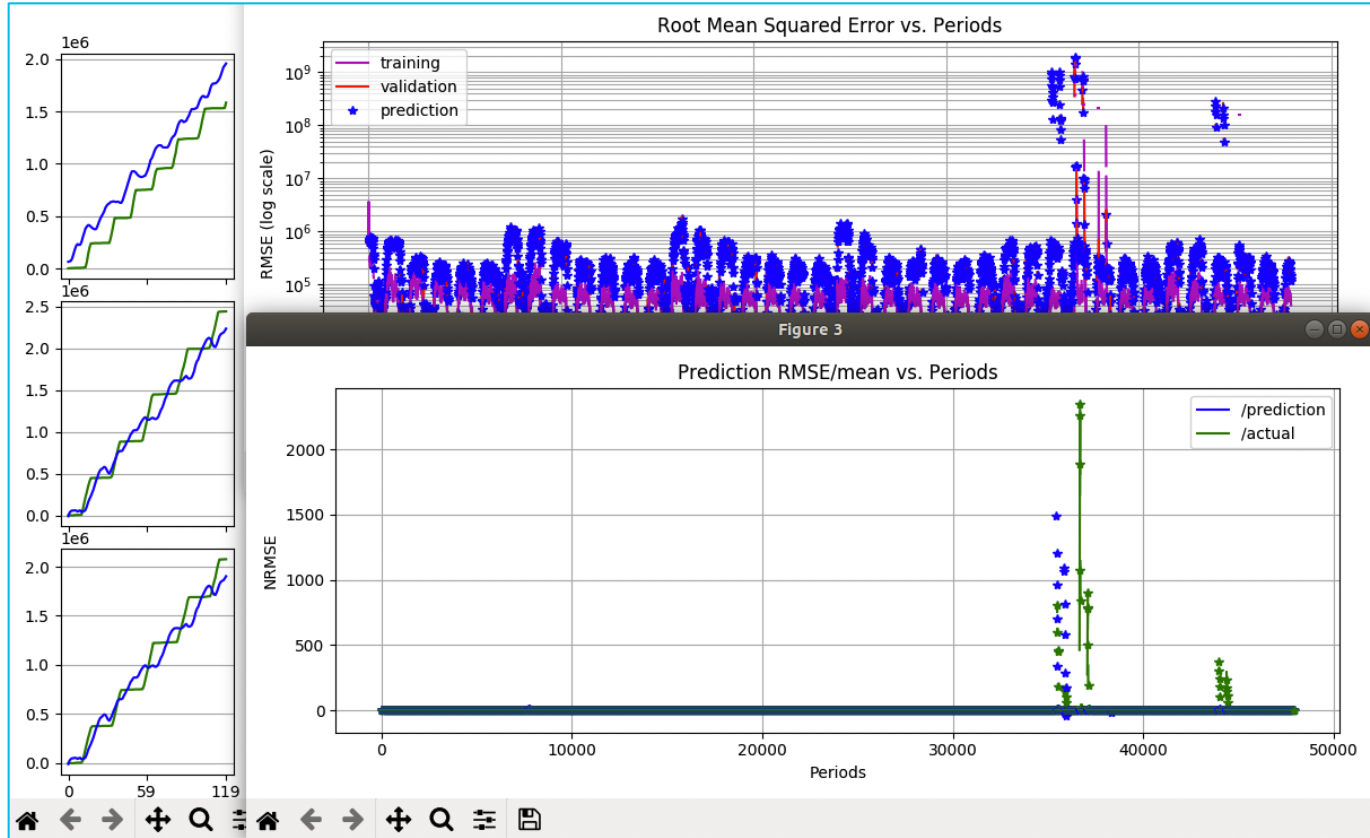
Best Indicator: NRMSE



A Short Term, Small Issue, & the Morning After




A Severe Issue



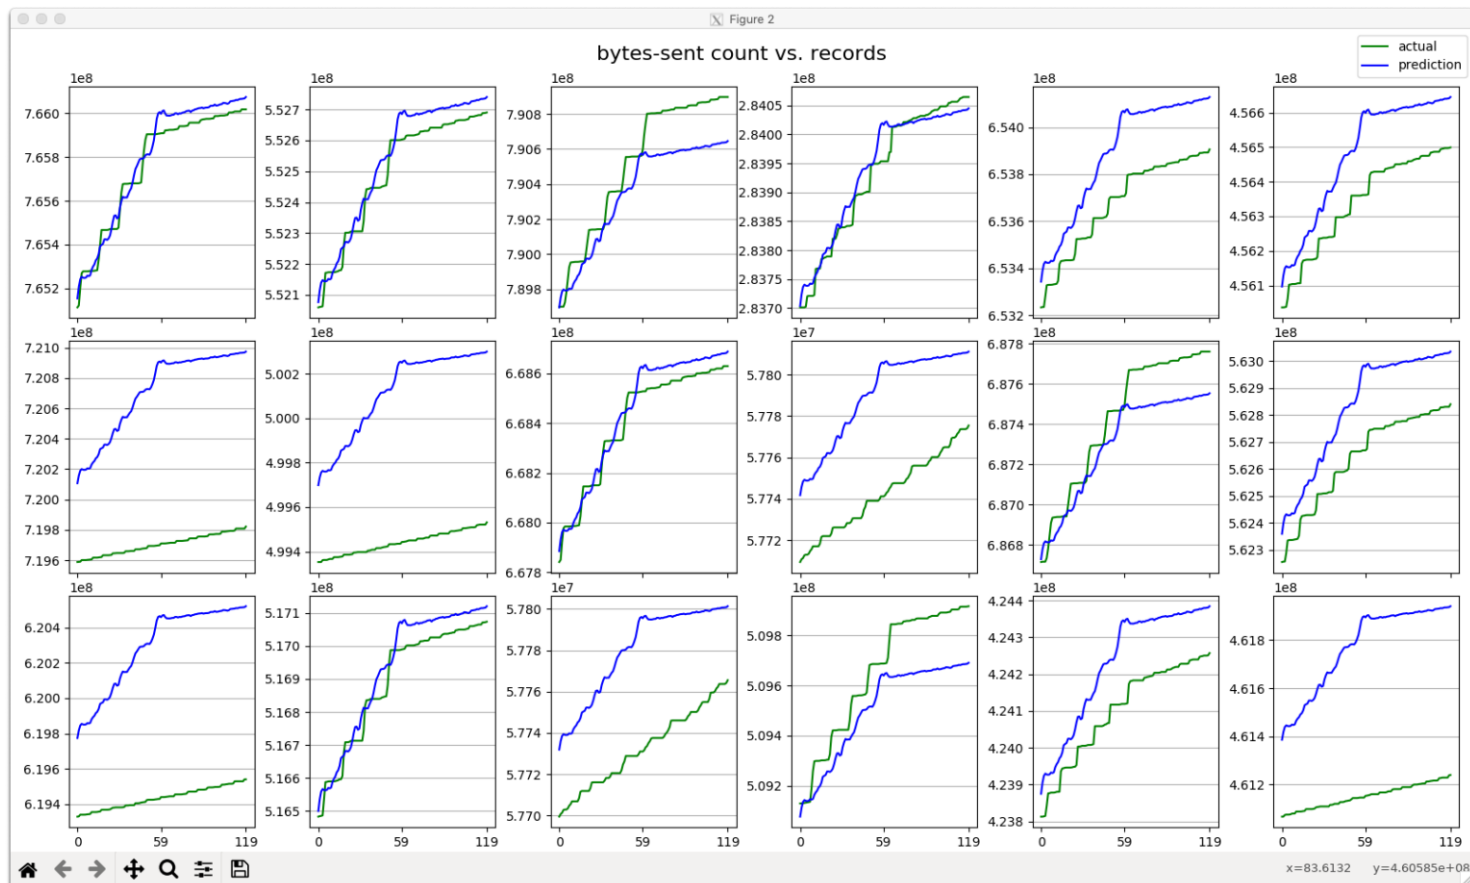
Visualize Progress: RMSE

```
.
for period in range(0, periods):
    # Train the model, starting from the prior state.
    dnn_regressor.train(
        input_fn=training_input_fn,
        steps=steps_per_period
    )
    # Take a break and compute predictions.
    training_predictions = dnn_regressor.predict(input_fn=predict_training_input_fn)
    training_predictions = np.array([[item['predictions'][i] for i in range(0,
len(tunnel_ifs) + len(physical_ifs))] for item in training_predictions])
.
    validation_predictions =
dnn_regressor.predict(input_fn=predict_validation_input_fn)
    validation_predictions = np.array([[item['predictions'][i] for i in range(0,
len(tunnel_ifs) + len(physical_ifs))] for item in validation_predictions])
.
    if if_plot:
        # RMSE values and graphs
        global x_periods
        x_periods += periods
        plt.ion()
.
```

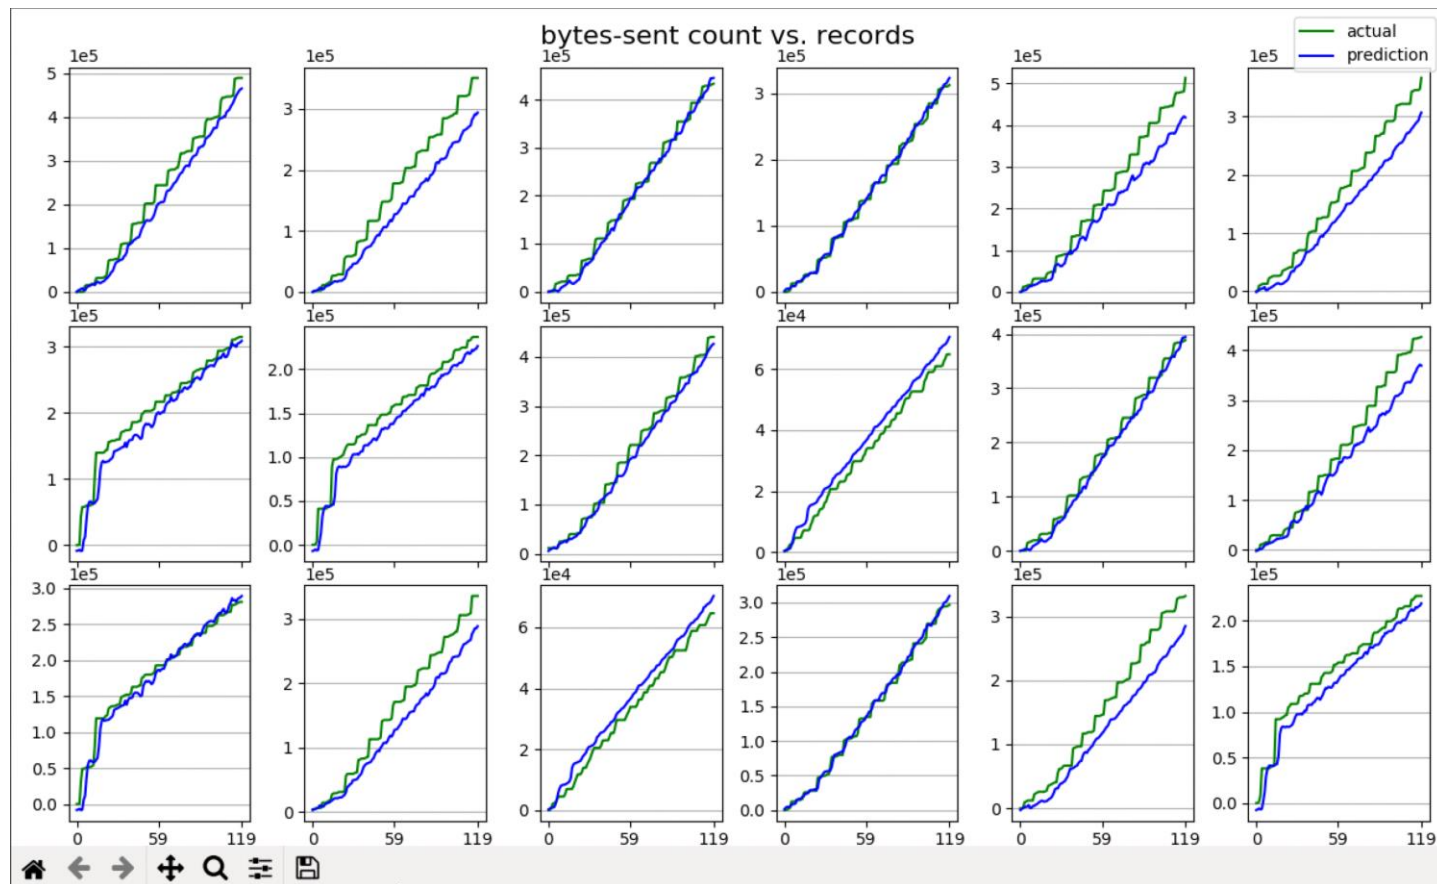
Visualize Progress: RMSE



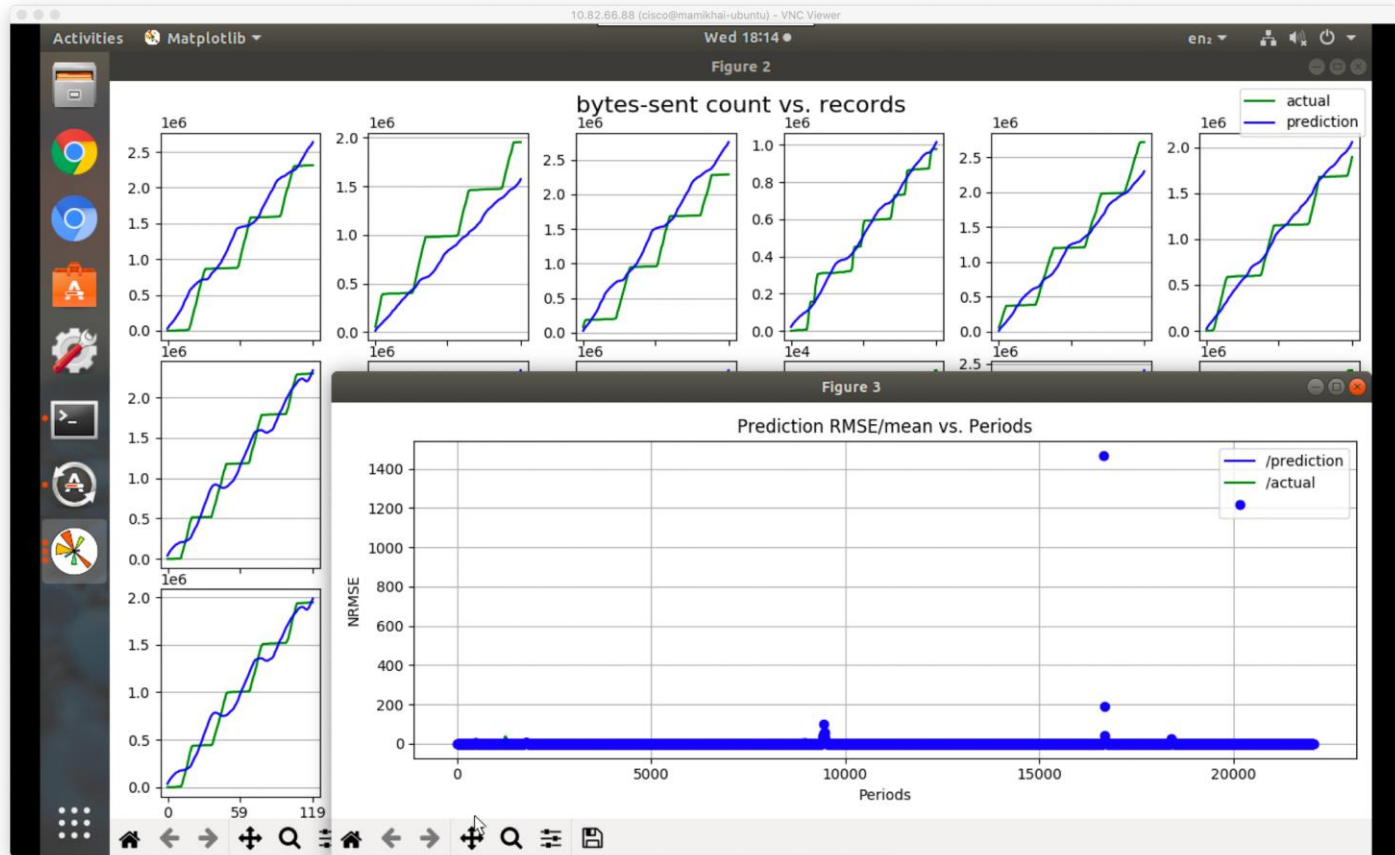
Validate: Visualize Prediction vs. Actual



Validate: Visualize Prediction vs. Actual



Eyes are Open?



Complete your online session survey



- Please complete your session survey after each session. Your feedback is very important.
- Complete a minimum of 4 session surveys and the Overall Conference survey (starting on Thursday) to receive your Cisco Live t-shirt.
- All surveys can be taken in the Cisco Events Mobile App or by logging in to the Content Catalog on ciscolive.com/emea.

Cisco Live sessions will be available for viewing on demand after the event at ciscolive.com.

Continue your education



Demos in the
Cisco Showcase



Walk-In Labs



Meet the Engineer
1:1 meetings



Related sessions



Thank you





You make **possible**