



The bridge to possible

Learning and Using Kubernetes with a Cluster of Raspberry Pis

Part 2 of 2

Jason Davis, Distinguished Engineer
@SNMPguy



Agenda

- Introduction
- Some Basic Kubernetes commands
- Installing MetalLB
- Layer-2 network integration
- Layer-3 network integration
- Other Uses
- Conclusion

Recap

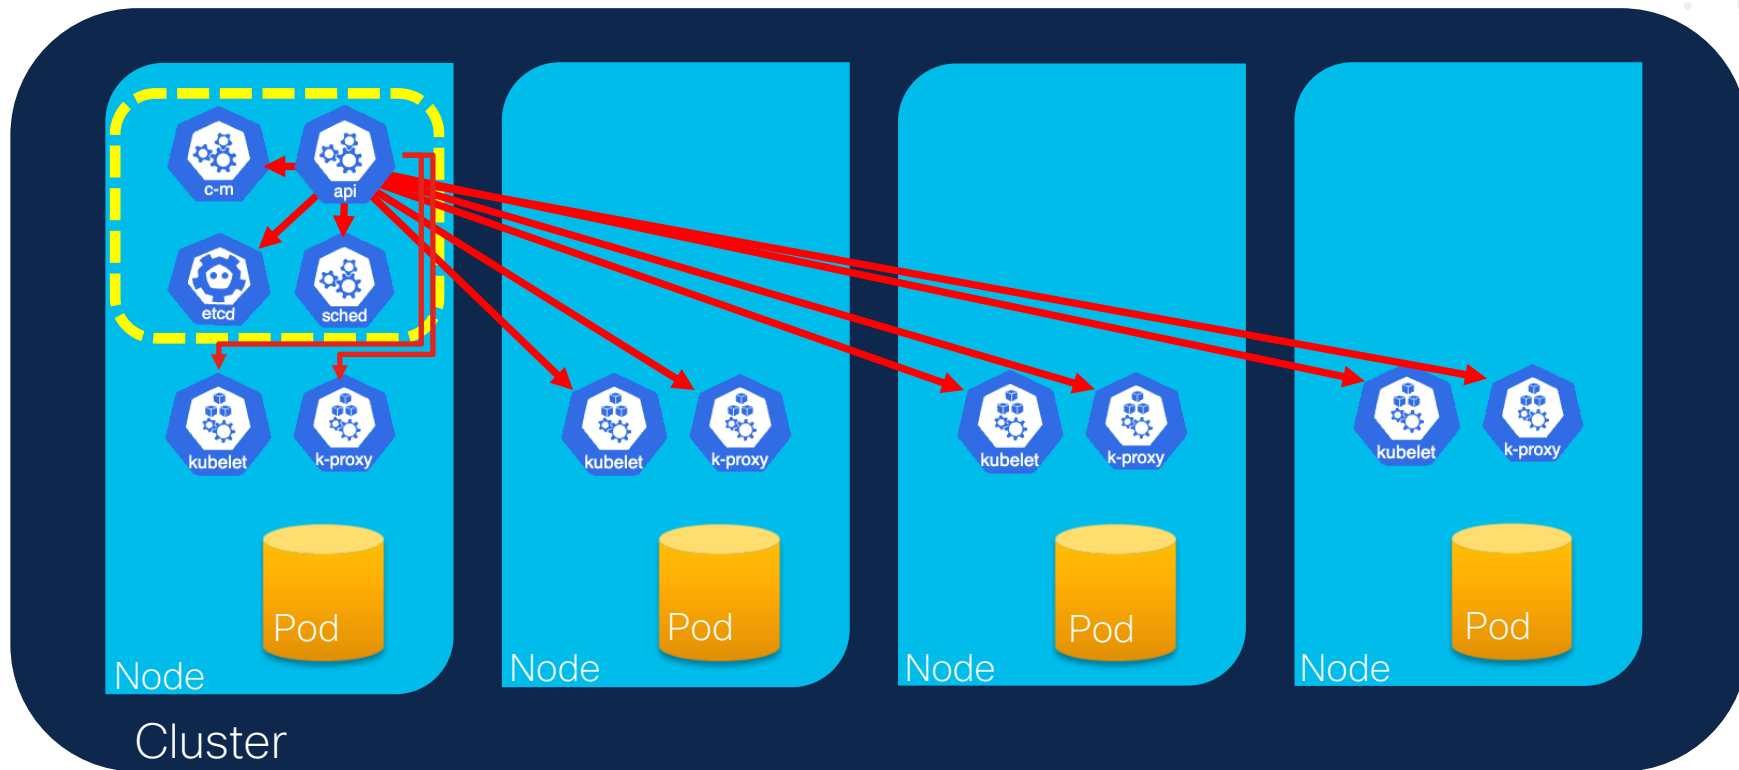
- Part 1, DEVNET-1279, we built a cluster of Raspberry Pis, installed microK8s and joined the nodes to a cluster
- In Part 2, DEVNET-2279 (THIS session), we will learn and use some microK8s administrative commands and configure MetalLB to interoperate with a routed network using BGP
- Join Part 2 directly if you have an existing microK8s cluster or Linux VMs already provisioned



kubernetes

NOTE: Cloud-based K8s clusters do not use MetalLB, as there are other mechanisms for network integration


Component View of Cluster



Get microK8s status

microk8s status


Run from controller node



```
jason@pi-cluster-node1:~$ microk8s status
microk8s is running
high-availability: no
addons:
  enabled:
    dns                # (core) CoreDNS
    helm               # (core) Helm - the package manager for Kubernetes
    helm3              # (core) Helm 3 - the package manager for Kubernetes
    hostpath-storage   # (core) Storage class; allocates storage from host directory
    metallb            # (core) Loadbalancer for your Kubernetes cluster
    storage             # (core) Alias to hostpath-storage add-on, deprecated
  disabled:
    cert-manager        # (core) Cloud native certificate management
    community           # (core) The community addons repository
    dashboard           # (core) The Kubernetes dashboard
    ha-cluster          # (core) Configure high availability on the current node
    host-access         # (core) Allow Pods connecting to Host services smoothly
    ingress             # (core) Ingress controller for external access
    kube-ovn            # (core) An advanced network fabric for Kubernetes
    mayastor            # (core) OpenEBS MayaStor
    metrics-server      # (core) K8s Metrics Server for API access to service metrics
    . . .
```

Get K8s nodes

```
microk8s kubectl get nodes -o wide
```




```
jason@pi-cluster-node1:~$ microk8s kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE
pi-cluster-node1	Ready	<none>	23d	v1.26.0	192.168.1.161	<none>	Ubuntu 22.04.1 LTS
5.15.0-1021-raspi	containerd://1.6.8						
pi-cluster-node2	Ready	<none>	22d	v1.26.0	192.168.1.162	<none>	Ubuntu 22.04.1 LTS
5.15.0-1021-raspi	containerd://1.6.8						
pi-cluster-node3	Ready	<none>	21d	v1.26.0	192.168.1.163	<none>	Ubuntu 22.04.1 LTS
5.15.0-1021-raspi	containerd://1.6.8						
pi-cluster-node4	Ready	<none>	21d	v1.26.0	192.168.1.164	<none>	Ubuntu 22.04.1 LTS
5.15.0-1021-raspi	containerd://1.6.8						


Determine node function by label

```
microk8s kubectl get nodes -l 'node.kubernetes.io/microk8s-controlplane'
microk8s kubectl get nodes -l '!node.kubernetes.io/microk8s-controlplane'
```



```
jason@pi-cluster-node1:~$ microk8s kubectl get nodes -l 'node.kubernetes.io/microk8s-controlplane'
```

NAME	STATUS	ROLES	AGE	VERSION
pi-cluster-node1	Ready	<none>	139d	v1.26.0



```
jason@pi-cluster-node1:~$ microk8s kubectl get nodes -l '!node.kubernetes.io/microk8s-controlplane'
```

NAME	STATUS	ROLES	AGE	VERSION
pi-cluster-node2	Ready	<none>	30m	v1.26.0
pi-cluster-node3	Ready	<none>	24m	v1.26.0
pi-cluster-node4	Ready	<none>	23m	v1.26.0


Labeling for custom querying

```
microk8s kubectl label node <node> node-role.kubernetes.io/<label_info>
```

```
jason@pi-cluster-node1:~$ microk8s kubectl label node pi-cluster-node1 node-  
role.kubernetes.io/controller=controller  
node/pi-cluster-node1 labeled  
jason@pi-cluster-node1:~$ microk8s kubectl label node pi-cluster-node2 node-  
role.kubernetes.io/worker=worker  
node/pi-cluster-node2 labeled  
jason@pi-cluster-node1:~$ microk8s kubectl label node pi-cluster-node3 node-  
role.kubernetes.io/worker=worker  
node/pi-cluster-node3 labeled  
jason@pi-cluster-node1:~$ microk8s kubectl label node pi-cluster-node4 node-  
role.kubernetes.io/worker=worker  
node/pi-cluster-node4 labeled
```


Get K8s services

```
microk8s kubectl get services
```



```
jason@pi-cluster-node1:~$ microk8s kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.152.183.1	<none>	443/TCP	139d

Create a deployment



Container hostname: microbot-7dc9ffb6cf-mpfz7

```
jason@pi-cluster-node1:~$ microk8s kubectl create deployment microbot \
--image=cdkbot/microbot-arm64:latest
deployment.apps/microbot created
jason@pi-cluster-node1:~$ microk8s kubectl scale deployment microbot --replicas=4
deployment.apps/microbot scaled
```

```
jason@pi-cluster-node1:~$ microk8s kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS
default	microbot-7dc9ffb6cf-5kxhp	1/1	Running	
default	microbot-7dc9ffb6cf-8qlsb	1/1	Running	
default	microbot-7dc9ffb6cf-stgqd	1/1	Running	
default	microbot-7dc9ffb6cf-vnxc1	1/1	Running	
kube-system	coredns-6f5f9b5d74-ggj94	1/1	Running	
kube-system	hostpath-provisioner-69cd9ff5b8-mvffv	1/1	Running	


microbot is a convenient docker image to demonstrate K8s functions

It is VERY small; based on Alpine Linux with NGINX web-server showing a static image and the K8s node identifier

<https://hub.docker.com/r/cdkbot/microbot-arm64>

Expose the service

- Initially we'll define a service as a NodePort
- Each node will expose the same port number individually
- Directed queries to the node's IP with the identified port will get that pod instance's output



```
jason@pi-cluster-node1:~$ microk8s kubectl expose deployment microbot --type=NodePort --port=80 --name=microbot-service
service/microbot-service exposed
```

View the deployment

jason@pi-cluster-node1:~\$ microk8s kubectl get all -n default

NAME	READY	STATUS	RESTARTS	AGE
pod/microbot-7dc9ffb6cf-jsj9b	1/1	Running	0	4m15s
pod/microbot-7dc9ffb6cf-mpfz7	1/1	Running	0	4m15s
pod/microbot-7dc9ffb6cf-xs8cj	1/1	Running	0	4m15s
pod/microbot-7dc9ffb6cf-zcxt4	1/1	Running	0	4m15s

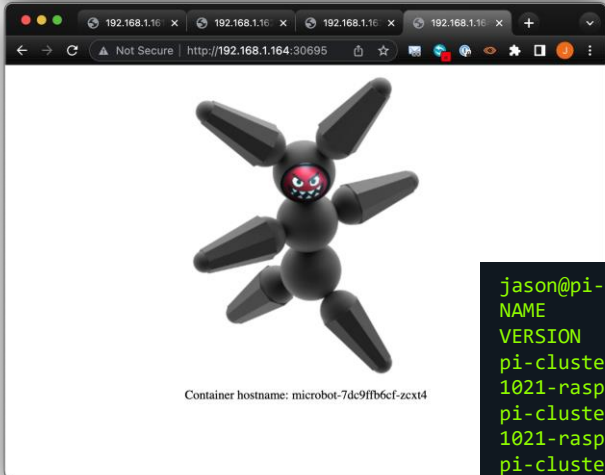
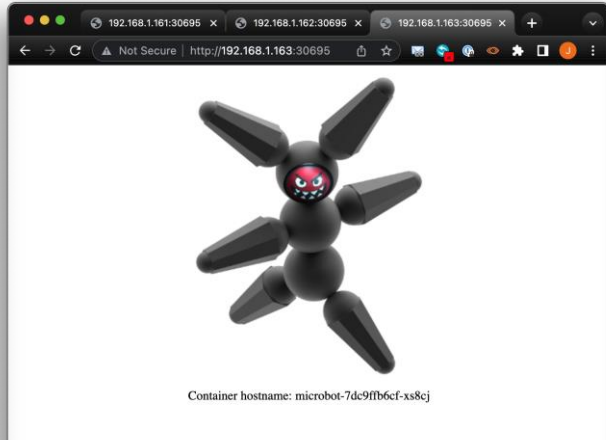
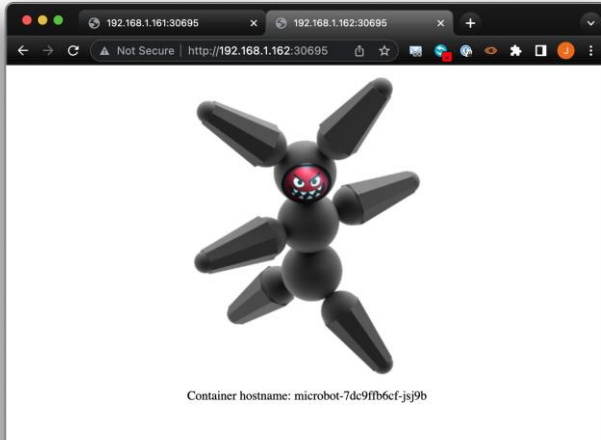
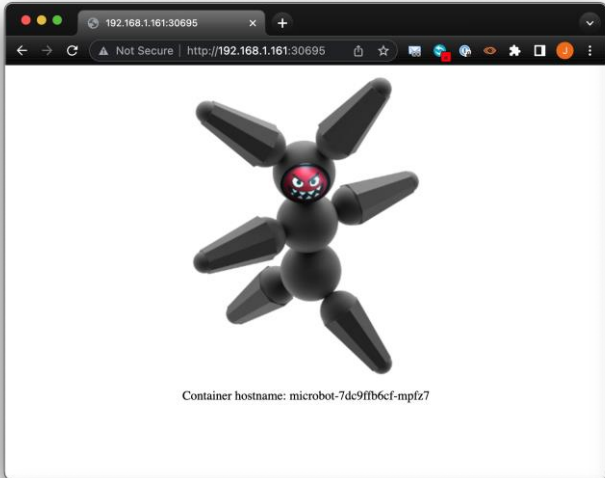
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	
service/kubernetes	ClusterIP	10.152.183.1	<none>	443/TCP	22d
service/microbot-service	NodePort	10.152.183.98	<none>	80:30695/TCP	5s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/microbot	4/4	4	4	13m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/microbot-7dc9ffb6cf	4	4	4	13m

Great!

Now we can access via
NodePort externally



```
jason@pi-cluster-node1:~$ microk8s kubectl get pods,svc -o wide
```

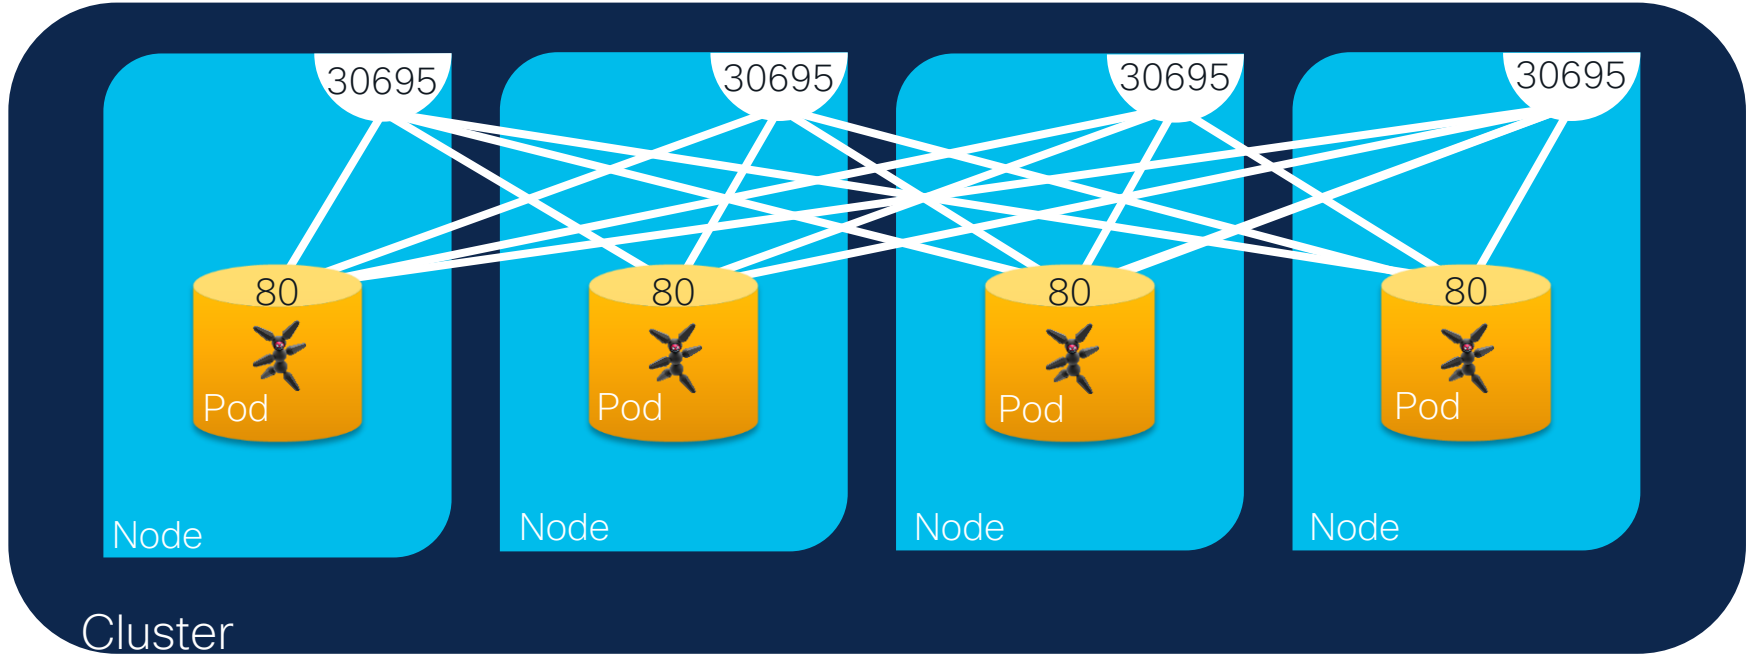
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
pod/microbot-7dc9ffb6cf-jsj9b	1/1	Running	0	32m	10.1.27.5	pi-cluster-node2
pod/microbot-7dc9ffb6cf-mpfz7	1/1	Running	0	32m	10.1.1.7	pi-cluster-node1
pod/microbot-7dc9ffb6cf-xs8cj	1/1	Running	0	32m	10.1.81.3	pi-cluster-node3
pod/microbot-7dc9ffb6cf-zcxt4	1/1	Running	0	32m	10.1.100.6	pi-cluster-node4

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
service/kubernetes	ClusterIP	10.152.183.1	<none>	443/TCP	22d	<none>
service/microbot-service	NodePort	10.152.183.98	<none>	80:30695/TCP	28m	

app=microbot

```
jason@pi-cluster-node1:~$ microk8s kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KERNEL-
pi-cluster-node1	Ready	<none>	23d	v1.26.0	192.168.1.161	<none>	Ubuntu 22.04.1 LTS	5.15.0-
1021-raspi	containerd://1.6.8							
pi-cluster-node2	Ready	<none>	22d	v1.26.0	192.168.1.162	<none>	Ubuntu 22.04.1 LTS	5.15.0-
1021-raspi	containerd://1.6.8							
pi-cluster-node3	Ready	<none>	21d	v1.26.0	192.168.1.163	<none>	Ubuntu 22.04.1 LTS	5.15.0-
1021-raspi	containerd://1.6.8							
pi-cluster-node4	Ready	<none>	21d	v1.26.0	192.168.1.164	<none>	Ubuntu 22.04.1 LTS	5.15.0-
1021-raspi	containerd://1.6.8							



MetalLB

Connect Externally with LoadBalancer

- Get MetalLB as a K8s manifest
- <https://metallb.universe.tf/installation/#installation-by-manifest>


```
microk8s kubectl apply -f  
https://raw.githubusercontent.com/metallb/metallb/v0.13.7/config/manifests/metallb-  
native.yaml
```


Apply MetalLB manifest



```
jason@pi-cluster-node1:~$ microk8s kubectl apply -f
https://raw.githubusercontent.com/metallb/metallb/v0.13.7/config/manifests/metallb-native.yaml
namespace/metallb-system created
customresourcedefinition.apiextensions.k8s.io/addresspools.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bfdprofiles.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bgpadvertisements.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bgppeers.metallb.io created
customresourcedefinition.apiextensions.k8s.io/communities.metallb.io created
customresourcedefinition.apiextensions.k8s.io/ipaddresspools.metallb.io created
customresourcedefinition.apiextensions.k8s.io/l2advertisements.metallb.io created
serviceaccount/controller created
serviceaccount/speaker created
role.rbac.authorization.k8s.io/controller created
role.rbac.authorization.k8s.io/pod-lister created
clusterrole.rbac.authorization.k8s.io/metallb-system:controller created
clusterrole.rbac.authorization.k8s.io/metallb-system:speaker created
rolebinding.rbac.authorization.k8s.io/controller created
rolebinding.rbac.authorization.k8s.io/pod-lister created
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:controller created
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:speaker created
secret/webhook-server-cert created
service/webhook-service created
deployment.apps/controller created
daemonset.apps/speaker created
validatingwebhookconfiguration.admissionregistration.k8s.io/metallb-webhook-configuration created
jason@pi-cluster-node1:~$
```

Review the metallb-system namespace created



```
jason@pi-cluster-node1:~$ microk8s kubectl get all -n metallb-system
```

NAME	READY	STATUS	RESTARTS	AGE
pod/controller-577b5bdfcc-7fflq	1/1	Running	0	31s
pod/speaker-bsr9f	1/1	Running	0	31s
pod/speaker-k4txn	1/1	Running	0	31s
pod/speaker-mk999	1/1	Running	0	31s
pod/speaker-nmgdw	1/1	Running	0	31s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/webhook-service	ClusterIP	10.152.183.227	<none>	443/TCP	31s

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
daemonset.apps/speaker	4	4	4	4	4	kubernetes.io/os=linux	31s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/controller	1/1	1	1	31s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/controller-577b5bdfcc	1	1	1	31s

```
jason@pi-cluster-node1:~$
```

Define and deploy a Layer-2 configmap

<https://metallb.universe.tf/configuration/#layer-2-configuration>


```
---
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  name: default
  namespace: metallb-system
spec:
  addresses:
    - 192.168.1.170-192.168.1.179
---
apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
  name: default
  namespace: metallb-system
```

Use an available address range in the same subnet as the INTERNAL-IPs of your nodes

Check with:
microk8s kubectl get nodes -o wide

save as 'cr-metallb-l2arp.yaml'

Apply the YAML definition



```
jason@pi-cluster-node1:~$ microk8s kubectl apply -f cr-metallb-l2arp.yaml
ipaddresspool.metallb.io/default created
l2advertisement.metallb.io/default created
```

Apply service and get status

jason@pi-cluster-node1:~\$ microk8s kubectl get services -A

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	kubernetes	ClusterIP	10.152.183.1	<none>	443/TCP	22d
default	microbot-service	NodePort	10.152.183.98	<none>	80:30695/TCP	92m
kube-system	kube-dns	ClusterIP	10.152.183.10	<none>	53/UDP,53/TCP,9153/TCP	21d

jason@pi-cluster-node1:~\$ microk8s kubectl delete service microbot-service

service "microbot-service" deleted

jason@pi-cluster-node1:~\$ microk8s kubectl expose deployment microbot --type=LoadBalancer --port=80 --name=microbot-service

service/microbot-service exposed

jason@pi-cluster-node1:~\$ microk8s kubectl get svc -A

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	kubernetes	ClusterIP	10.152.183.1	<none>	443/TCP	140d
metallb-system	webhook-service	ClusterIP	10.152.183.198	<none>	443/TCP	18h
kube-system	kube-dns	ClusterIP	10.152.183.10	<none>	53/UDP,53/TCP,9153/TCP	18h
default	microbot-service	LoadBalancer	10.152.183.22	192.168.1.170	80:31182/TCP	72s

jason@pi-cluster-node1:~\$

Apply service and get status

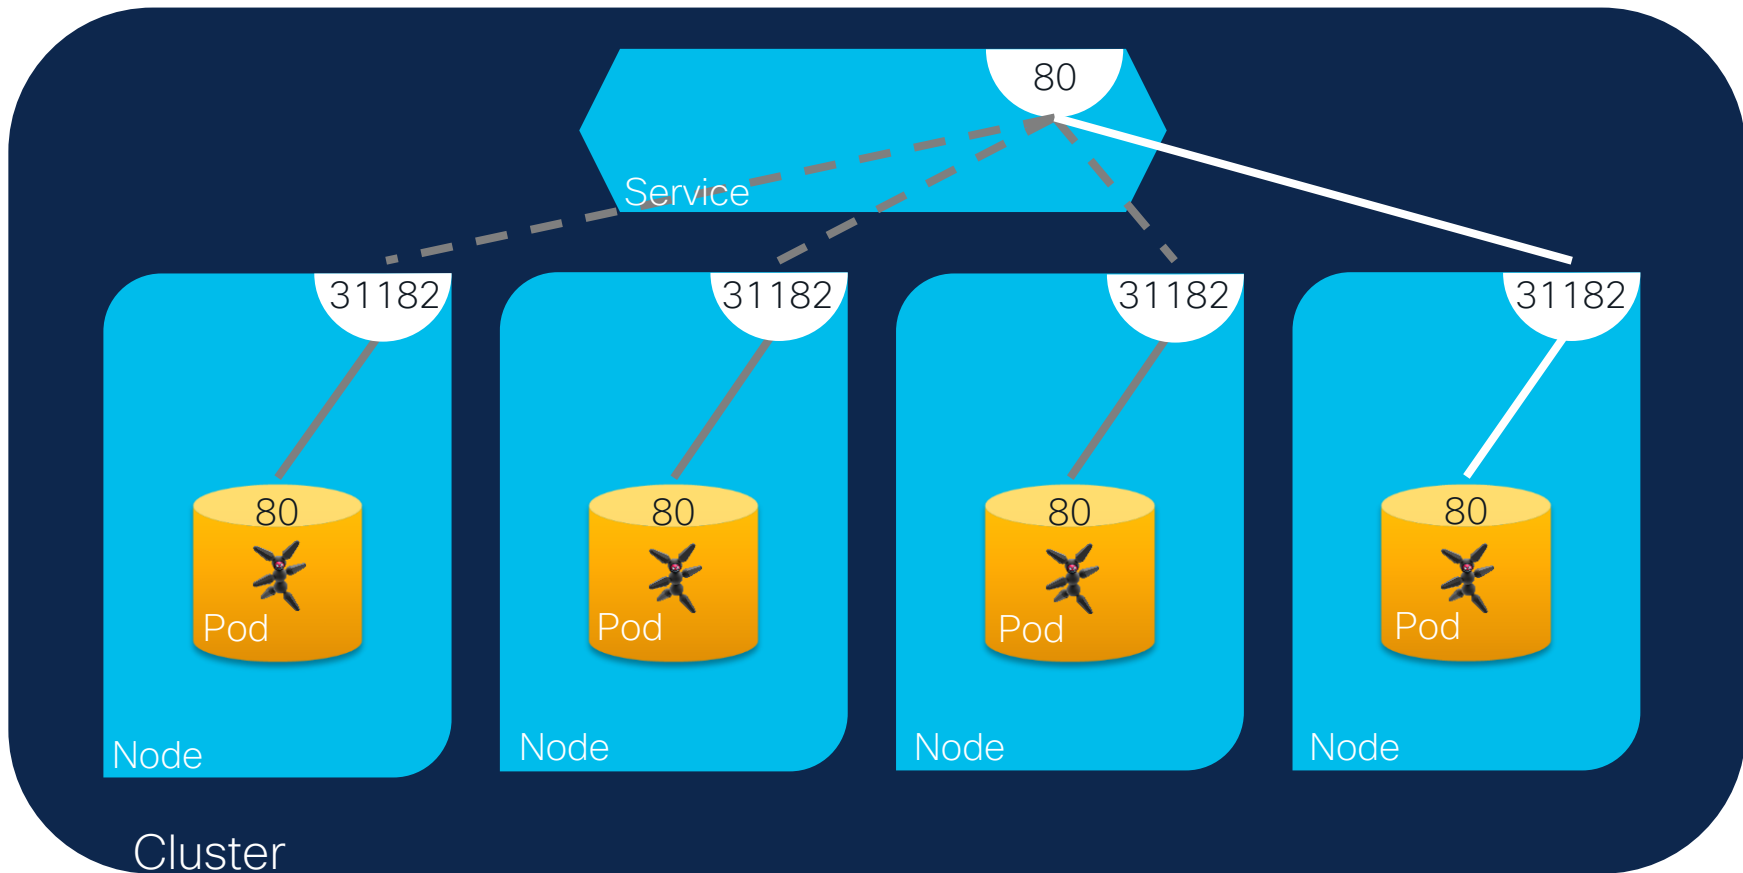
```
jason@pi-cluster-node1:~$ microk8s kubectl describe service/microbot-service
```

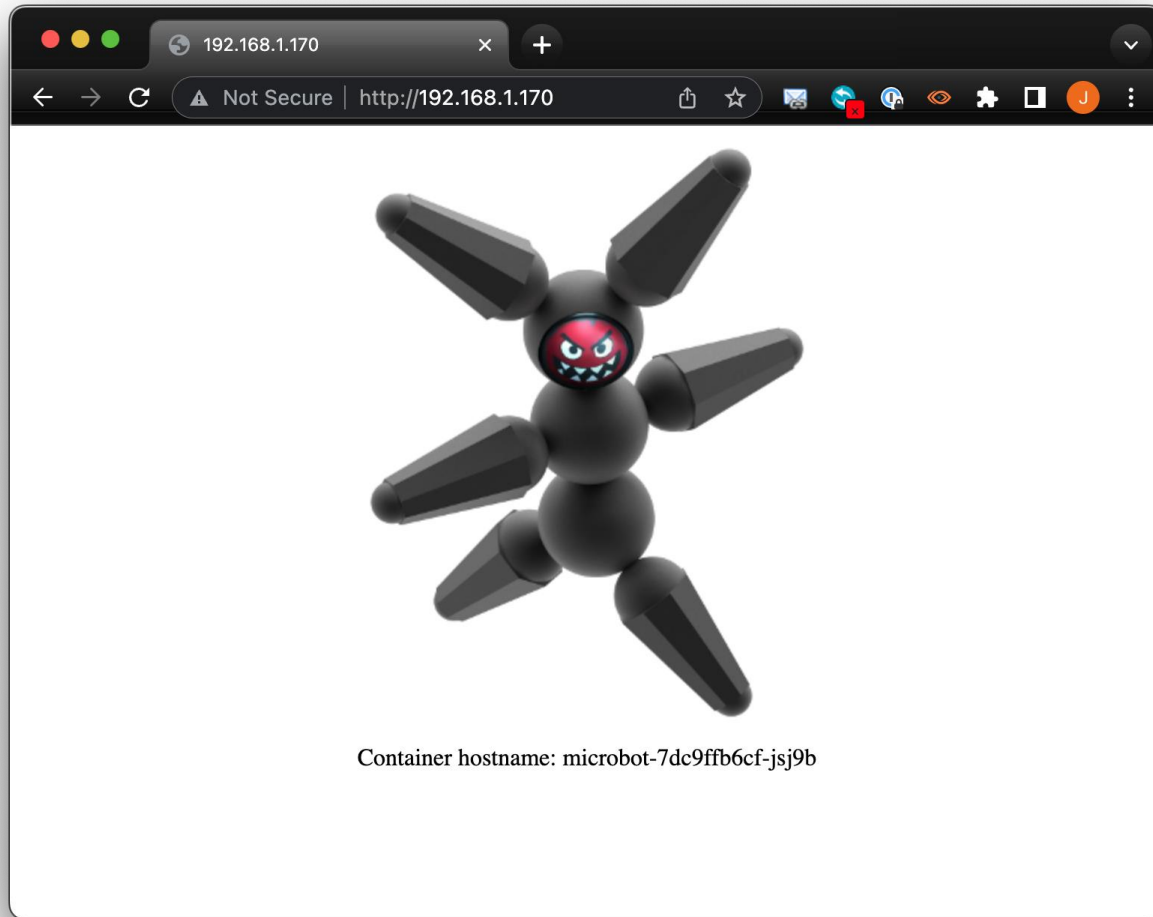
```
Name: microbot-service
Namespace: default
Labels: app=microbot
Annotations: <none>
Selector: app=microbot
Type: LoadBalancer
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.152.183.41
IPs: 10.152.183.41
LoadBalancer Ingress: 192.168.1.170
Port: <unset> 80/TCP
TargetPort: 80/TCP
NodePort: <unset> 31182/TCP
Endpoints: 10.1.1.7:80,10.1.100.6:80,10.1.27.5:80 + 1 more...
```

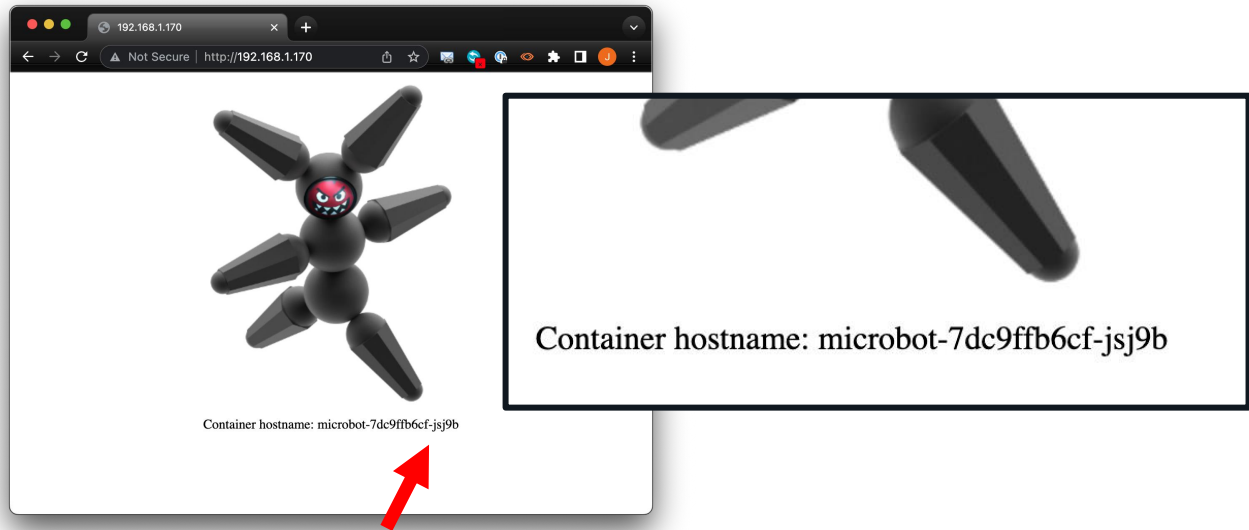
```
Session Affinity: None
External Traffic Policy: Cluster
```

```
Events:
```

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	IPAllocated	3m3s	metallb-controller	Assigned IP ["192.168.1.170"]
Normal	nodeAssigned	3m3s	metallb-speaker	announcing from node "pi-cluster-node4" with protocol "layer2"







Now delete the pod hosting this microbot instance and see another takeover

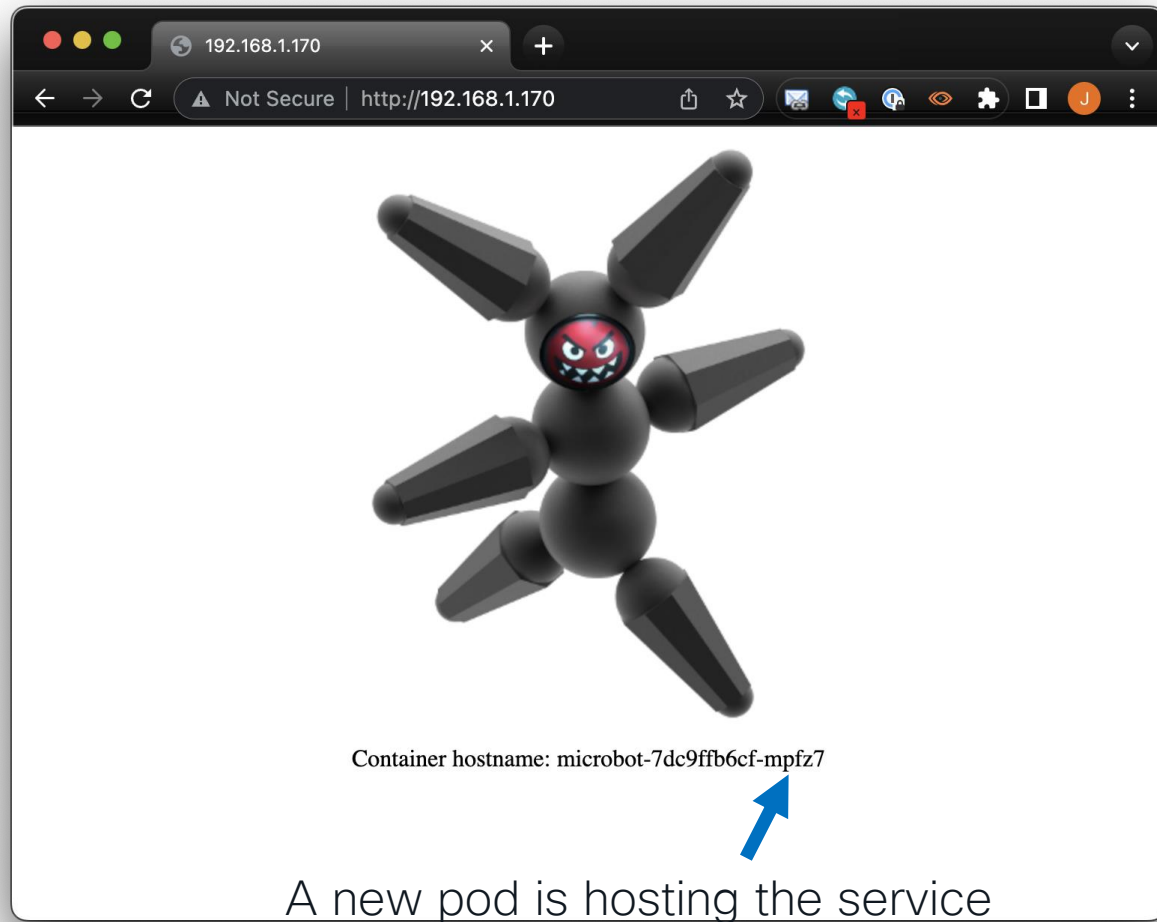
```
jason@pi-cluster-node1:~$ microk8s kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
microbot-7dc9ffb6cf-mpfz7	1/1	Running	0	25h	10.1.1.7	pi-cluster-node1
microbot-7dc9ffb6cf-ffpf2	1/1	Running	0	22h	10.1.27.8	pi-cluster-node2
microbot-7dc9ffb6cf-xs8cj	1/1	Running	0	25h	10.1.81.3	pi-cluster-node3
microbot-7dc9ffb6cf-jsj9b	1/1	Running	0	25h	10.1.100.6	pi-cluster-node4

```

→ jason@pi-cluster-node1:~$ microk8s kubectl delete pod/microbot-7dc9ffb6cf-jsj9b
pod "microbot-7dc9ffb6cf-jsj9b" deleted
jason@pi-cluster-node1:~$
jason@pi-cluster-node1:~$ microk8s kubectl get pods -o wide
NAME                                READY   STATUS              RESTARTS   AGE   IP            NODE
microbot-7dc9ffb6cf-mpfz7           1/1     Running             0          25h   10.1.1.7      pi-cluster-node1
microbot-7dc9ffb6cf-ffpf2           1/1     Running             0          22h   10.1.27.8     pi-cluster-node2
microbot-7dc9ffb6cf-xs8cj           1/1     Running             0          25h   10.1.81.3     pi-cluster-node3
→ microbot-7dc9ffb6cf-zcxt4         1/1     ContainerCreating   0          5s    10.1.100.6    pi-cluster-node4
jason@pi-cluster-node1:~$
jason@pi-cluster-node1:~$ microk8s kubectl get pods -o wide
NAME                                READY   STATUS   RESTARTS   AGE   IP            NODE
→ microbot-7dc9ffb6cf-mpfz7           1/1     Running   0          25h   10.1.1.7      pi-cluster-node1
microbot-7dc9ffb6cf-ffpf2           1/1     Running   0          22h   10.1.27.8     pi-cluster-node2
microbot-7dc9ffb6cf-xs8cj           1/1     Running   0          25h   10.1.81.3     pi-cluster-node3
→ microbot-7dc9ffb6cf-zcxt4           1/1     Running   0          15s   10.1.100.6    pi-cluster-node4

```



Limitations on MetalLB Layer-2

- Depends on ARP/ND
- Single-node Bottleneck
- Slowish fail-over performance

Now for Magic...

Integration with BGP Routing




Reconfigure microk8s

- Remove earlier L2/ARP config

```
jason@pi-cluster-node1:~$ microk8s kubectl delete -f ./cr-metallb-l2arp.yaml  
ipaddresspool.metallb.io "default" deleted  
l2advertisement.metallb.io "default" deleted
```

Reconfigure microk8s

- Create new L3/BGP config/CR



```
---
apiVersion: metallb.io/v1beta2
kind: BGPPeer
metadata:
  name: bgppeering
  namespace: metallb-system
spec:
  myASN: 64512
  peerASN: 64512
  peerAddress: 192.168.1.180
---
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  name: address-pool
  namespace: metallb-system
spec:
  addresses:
    - 192.168.1.170-192.168.1.179
---
apiVersion: metallb.io/v1beta1
kind: BGPAdvertisement
metadata:
  name: bgpadvertisement
  namespace: metallb-system
```

save as 'cr-metallb-l3bgp.yaml'

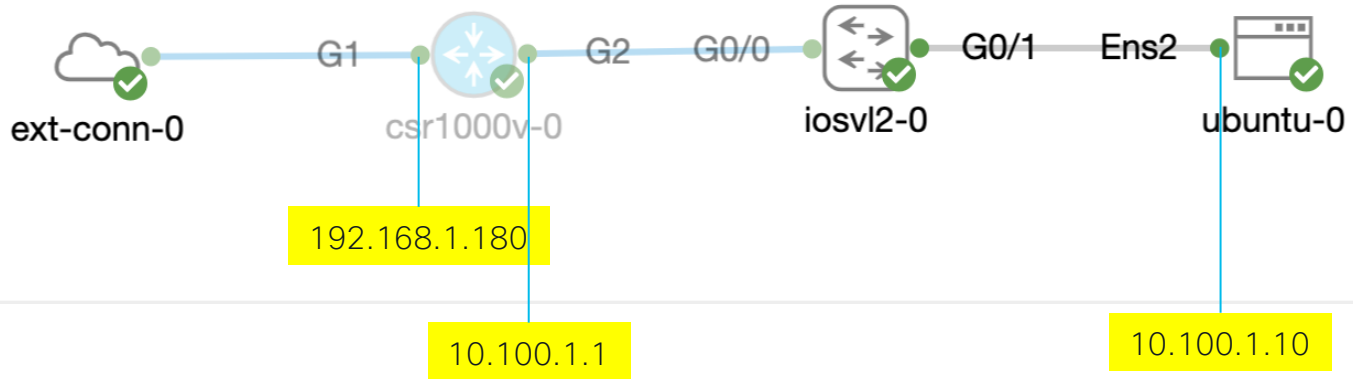
Reconfigure microk8s

- Apply new L3/BGP config/CR

```
jason@pi-cluster-node1:~$ microk8s kubectl apply -f ./cr-metallb-l3bgp.yaml
bgppeer.metallb.io/bgppeering created
ipaddresspool.metallb.io/address-pool created
bgpadvertisement.metallb.io/bgpadvertisement created
jason@pi-cluster-node1:~$
```


Use Cisco Modeling Labs (CML) to build up BGP Connectivity

Cisco Modeling Labs Workbench RPi K8s MetalLB



Configure CSR1000v for BGP routing

```
router bgp 64512
  bgp log-neighbor-changes
  redistribute connected
  neighbor 192.168.1.161 remote-as 64512
  neighbor 192.168.1.161 description rpi-cluster-node-1
  neighbor 192.168.1.162 remote-as 64512
  neighbor 192.168.1.162 description rpi-cluster-node-2
  neighbor 192.168.1.163 remote-as 64512
  neighbor 192.168.1.163 description rpi-cluster-node-3
  neighbor 192.168.1.164 remote-as 64512
  neighbor 192.168.1.164 description rpi-cluster-node-4
!
```



ADD NODES

NODE INFO SIMULATE CONNECTIVITY **CONSOLE** VNC EDIT CONFIG INTERFACES

SERIAL LINE 0 DOWNLOAD HISTORY SETTINGS

```
*Dec 5 21:32:55.871: %SYS-6-TTY_EXPIRE_TIMER: (exec timer expired, tty 0 (0.0.0.0)), user
inserthostname here>show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
        D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
        N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
        E1 - OSPF external type 1, E2 - OSPF external type 2, m - OMP
        n - NAT, Ni - NAT inside, No - NAT outside, Nd - NAT DIA
        i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
        ia - IS-IS inter area, * - candidate default, U - per-user static route
        H - NHRP, G - NHRP registered, g - NHRP registration summary
        o - ODR, P - periodic downloaded static route, l - LISP
        a - application route
        + - replicated route, % - next hop override, p - overrides from PfR
        % - replicated local route overrides by connected

Gateway of last resort is 192.168.1.1 to network 0.0.0.0

S*   0.0.0.0/0 [254/0] via 192.168.1.1
     10.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
C     10.100.1.0/24 is directly connected, GigabitEthernet2
L     10.100.1.1/32 is directly connected, GigabitEthernet2
     192.168.1.0/24 is variably subnetted, 3 subnets, 2 masks
C     192.168.1.0/24 is directly connected, GigabitEthernet1
B     192.168.1.170/32 [20/0] via 192.168.1.162, 00:57:24
L     192.168.1.180/32 is directly connected, GigabitEthernet1
inserthostname here>
```

CPU 26.60% MEMORY 28.97% DISK 6.83% Notifications 32 Status OK

Use Ubuntu node to do Console curl

```
cisco@inserthostnamehere:~$ curl http://192.168.1.170
<!DOCTYPE html>
<html>
  <style type="text/css">
    .centered
    {
      text-align:center;
      margin-top:0px;
      margin-bottom:0px;
      padding:0px;
    }
  </style>
  <body>
    <p class="centered"></p>
    <p class="centered">Container hostname: microbot-66c45859f5-4gxwm</p>
  </body>
</html>
cisco@inserthostnamehere:~$
```



Container hostname: microbot-66c45859f5-4gxwm

. . . Or Desktop node with VNC

Cisco Modeling Labs Workbench RPi K8s MetalLB

DASHBOARD TOOLS ADMIN

ext-conn-0 csr1000v-0 iosvl2-0 G0/1 Ens2 G0/2 ubuntu-0 E0 desktop-0

ADD NODES


NODE INFO SIMULATE CONNECTIVITY CONSOLE VNC EDIT CONFIG INTERFACES

CTRL+ALT+DEL FULLSCREEN

Applications: Mozilla Firefox

192.168.10.170/

192.168.10.170



Container hostname: microbot-7dc9ffb6cf-cccqt

CPU 34.50% MEMORY 32.58% DISK 7.28%

Notifications 7 Status OK

Other Uses

Another idea for the cluster – Network testing

- Network Multitool project
- <https://github.com/wbitt/Network-MultiTool>
- Includes useful tools like curl, mtr (my traceroute), tcptraceroute, arping, iperf3, tshark, rsync, nc (netcat), socat and ab (ApacheBench)
- All in a 220 Mbyte container

```
jason@pi-cluster-node1:~$ microk8s kubectl create deployment multitool --  
image=praqma/network-multitool:alpine-extra  
jason@pi-cluster-node1:~$  
jason@pi-cluster-node1:~$ microk8s kubectl exec -it pod/multitool-<pod_instance>  
-- /bin/bash
```

Network-Multitool / mtr

```
bash-5.1# which mtr
/usr/sbin/mtr
bash-5.1# mtr -c 100 developer.cisco.com
bash-5.1#
```

```
multitool-6d5cffb4d7-ndq85 (10.1.100.8) -> developer.cisco.com
Keys: Help  Display mode  Restart statistics  Order of fields  quit

My traceroute  [v0.94]
2023-01-09T16:55:42+0000
```

Host	Packets	Pings
	Loss% Snt	Last Avg Best Wrst StDev
1. 10.1.100.1	0.0% 100	0.5 0.4 0.3 0.6 0.1
2. gateway.1	0.0% 100	0.6 0.6 0.4 0.8 0.1
3. 192.168.1.254	0.0% 100	1.0 1.1 0.7 1.5 0.2
4. 99-178-168-1.lightspeed.rlghnc.sbcglobal.net	0.0% 100	4.1 4.4 1.9 7.5 1.4
5. 99.173.76.21	5.0% 100	2.7 4.5 2.1 7.2 1.4
6. (waiting for reply)		
7. (waiting for reply)		
8. (waiting for reply)		
9. 32.130.24.165	0.0% 100	19.5 16.6 11.4 22.0 2.6
10. cr1.cl2oh.ip.att.net	0.0% 100	21.1 16.9 10.9 23.1 2.6
11. 12.123.10.37	0.0% 100	12.5 12.7 9.8 16.2 1.4
12. 12.117.216.210	0.0% 100	11.5 14.8 11.0 50.0 5.0
13. (waiting for reply)		
14. (waiting for reply)		
15. (waiting for reply)		
16. (waiting for reply)		
17. (waiting for reply)		
18. (waiting for reply)		
19. server-18-67-65-49.iad89.r.cloudfront.net	0.0% 100	15.0 13.5 11.0 16.7 1.4

Network-Multitool / mtr

```
bash-5.1# mtr -z -j -c 10 developer.cisco.com
```

```
{  
  "report": {  
    "mtr": {  
      "src": "multitool-6d5cffb4d7-ndq85",  
      "dst": "developer.cisco.com",  
      "tos": 0,  
      "tests": 10,  
      "psize": "64",  
      "bitpattern": "0x00"  
    },  
    "hubs": [  
      {  
        "count": 1,  
        "host": "10.1.100.1",  
        "ASN": "AS???",  
        "Loss%": 0.0,  
        "Snt": 10,  
        "Last": 0.381,  
        "Avg": 0.426,  
        "Best": 0.348,  
        "Wrst": 0.521,  
        "StDev": 0.047  
      },  
      {  
        "count": 2,  
        "host": "gateway-  
        "ASN": "AS???",  
        "Loss%": 0.0,  
        "Snt": 10,  
        "Last": 0.566,  
        "Avg": 0.535,  
        "Best": 0.407,  
        "Wrst": 0.61,
```

Secure Shell VT220 Xmodem ZOC2301.LOG

01:22:53 10/33

Network-Multitool / ab (ApacheBench)

```
bash-5.1# ab -n 1000 -c 20 https://developer.cisco.com/
This is ApacheBench, Version 2.3 <$Revision: 1879490 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/
```

```
Benchmarking developer.cisco.com (be patient)
```

```
Completed 100 requests
Completed 200 requests
Completed 300 requests
Completed 400 requests
Completed 500 requests
Completed 600 requests
Completed 700 requests
Completed 800 requests
Completed 900 requests
Completed 1000 requests
Finished 1000 requests
```

```
Server Software:      CloudFront
Server Hostname:      developer.cisco.com
Server Port:          443
SSL/TLS Protocol:     TLSv1.2,ECDHE-RSA-AES128-GCM-SHA256,2048,128
Server Temp Key:       X25519 253 bits
TLS Server Name:       developer.cisco.com
```

```
Document Path:        /
Document Length:       6681 bytes
```

Network-Multitool / ab (ApacheBench)

```
Concurrency Level:      20
Time taken for tests:    5.313 seconds
Complete requests:      1000
Failed requests:         0
Total transferred:      7828012 bytes
HTML transferred:       6681000 bytes
Requests per second:    188.22 [#/sec] (mean)
Time per request:       106.260 [ms] (mean)
Time per request:       5.313 [ms] (mean, across all concurrent requests)
Transfer rate:          1438.84 [Kbytes/sec] received
```

Connection Times (ms)

	min	mean[+/-sd]	median	max
Connect:	51	76 5.9	76	97
Processing:	12	25 7.2	25	175
Waiting:	12	22 6.5	21	163
Total:	78	101 8.8	101	226

Percentage of the requests served within a certain time (ms)

50%	101
66%	104
75%	105
80%	107
90%	110
95%	113
98%	118
99%	121
100%	226 (longest request)

Other ideas

- Pi-Hole to block DNS and web advertisements in your browser
- Development of microservices using Flask containers with simple Python scripts
- Learning and practical experience development with Load balancing and how to swap out/in updated docker containers of microservices

In Conclusion – What did we learn?

- How to provision Raspberry Pis for Kubernetes use
- Basic Kubernetes (and microK8s) commands
- Integrating a Kubernetes cluster in an L2/ARP network with MetalLB
- Integrating a Kubernetes cluster in an L3/BGP network with MetalLB
- Hosting basic container services as replicated deployments
- A few fun practical tool use cases with K8s images

Cisco Webex App

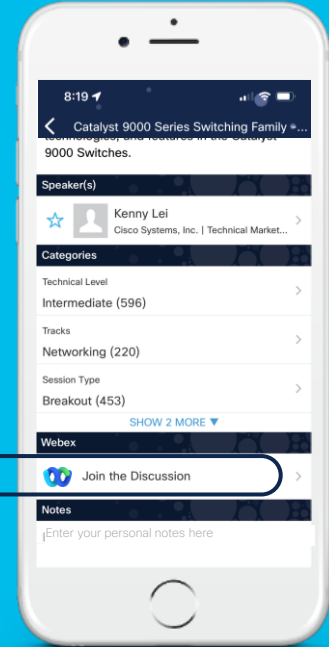
Questions?

Use Cisco Webex App to chat with the speaker after the session

How

- 1 Find this session in the Cisco Live Mobile App
- 2 Click “Join the Discussion”
- 3 Install the Webex App or go directly to the Webex space
- 4 Enter messages/questions in the Webex space

Webex spaces will be moderated until February 24, 2023.



Complete your Session Survey

- Please complete your session survey after each session. Your feedback is important.
- All surveys can be taken in the Cisco Events Mobile App or by logging in to the Session Catalog and clicking the "Attendee Dashboard" at <https://www.ciscolive.com/emea/learn/sessions/session-catalog.html>



Continue Your Education



Visit the Cisco Showcase for related demos.



Book your one-on-one Meet the Engineer meeting.



Attend any of the related sessions at the DevNet, Capture the Flag, and Walk-in Labs zones.



Visit the On-Demand Library for more sessions at ciscolive.com/on-demand.



The bridge to possible

Thank you

CISCO *Live!*

CISCO *Live!*

ALL IN