



Possibilities

#CiscoLive


Catalyst 9k Troubleshooting

Michel Peters,
Technical Leader Engineering
DGTL-BRKARC-3190



#CiscoLive





Agenda

Agenda

- Catalyst 9k Platform & Software Architecture
- Troubleshooting tools
- Platform Issues
- Packet Drops
- Forwarding Issues

The background is a dark blue field filled with numerous small, semi-transparent squares and dots in various colors including light blue, teal, yellow, orange, and red. These elements are scattered across the frame, with a higher concentration of yellow and orange squares forming a diagonal streak from the top right towards the bottom right.

Platform & Software Architecture

Catalyst 9k family

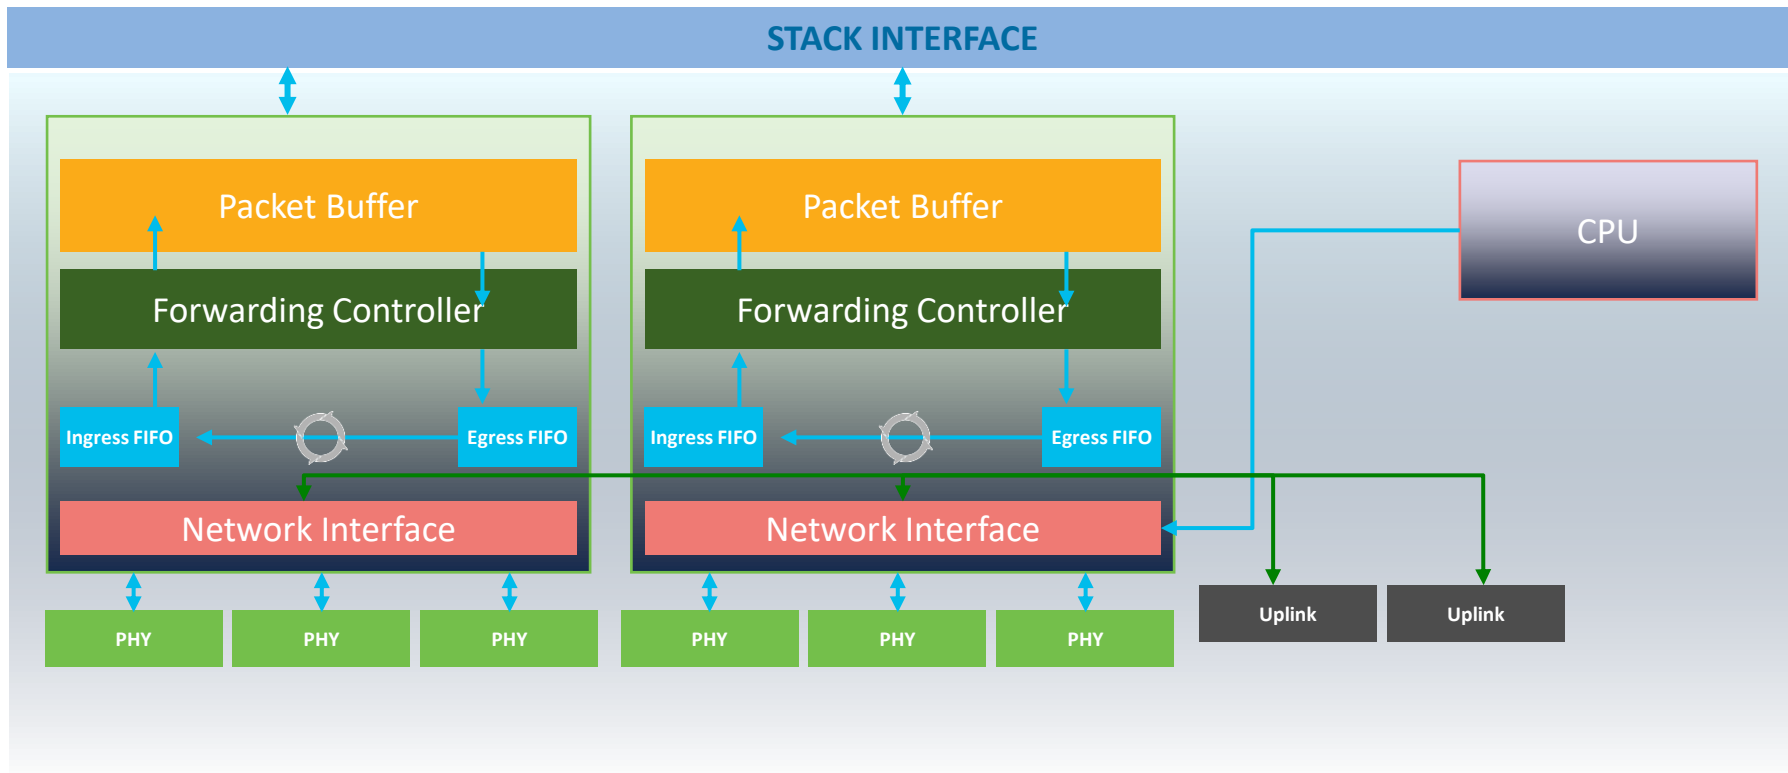


	9200	9300	9400	9500	9600
UADP	2.0 mini	2.0	2.0 XL	2.0 XL & 3.0	3.0
Format	Stackable	Stackable	Chassis	Standalone	Chassis

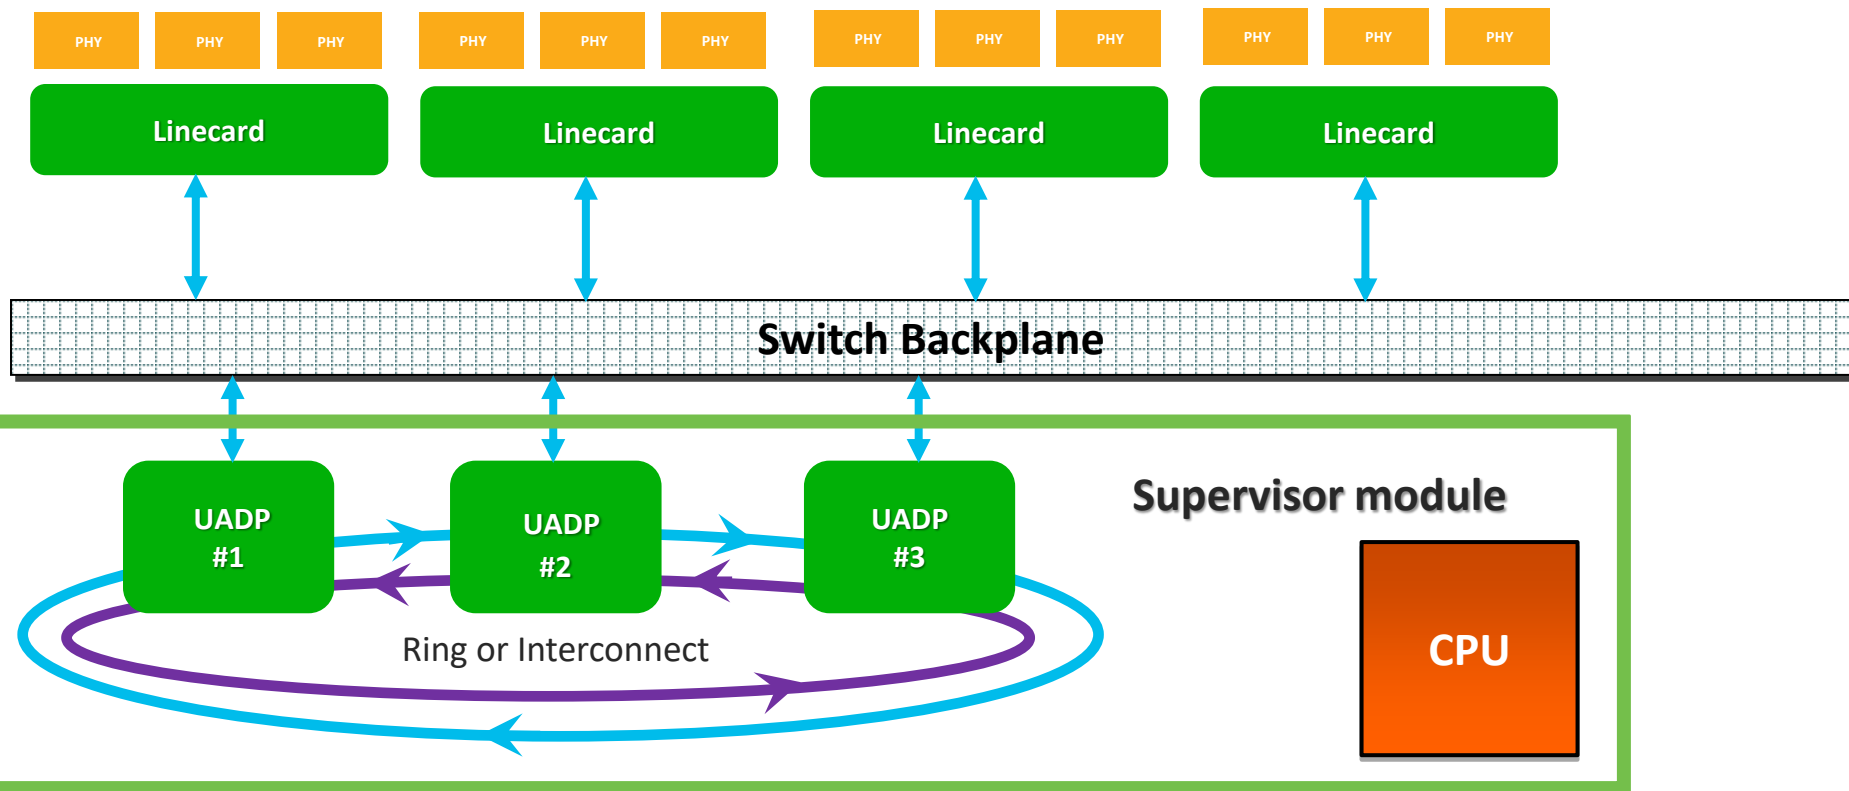
- UADP Asics are the foundation of all Catalyst 9K switches
- 3850/3650 switches originally launches with UADP 1.x
- Port Asics (UADP) responsible for data plan forwarding
- CPU's handle control plane traffic and forwarding tables



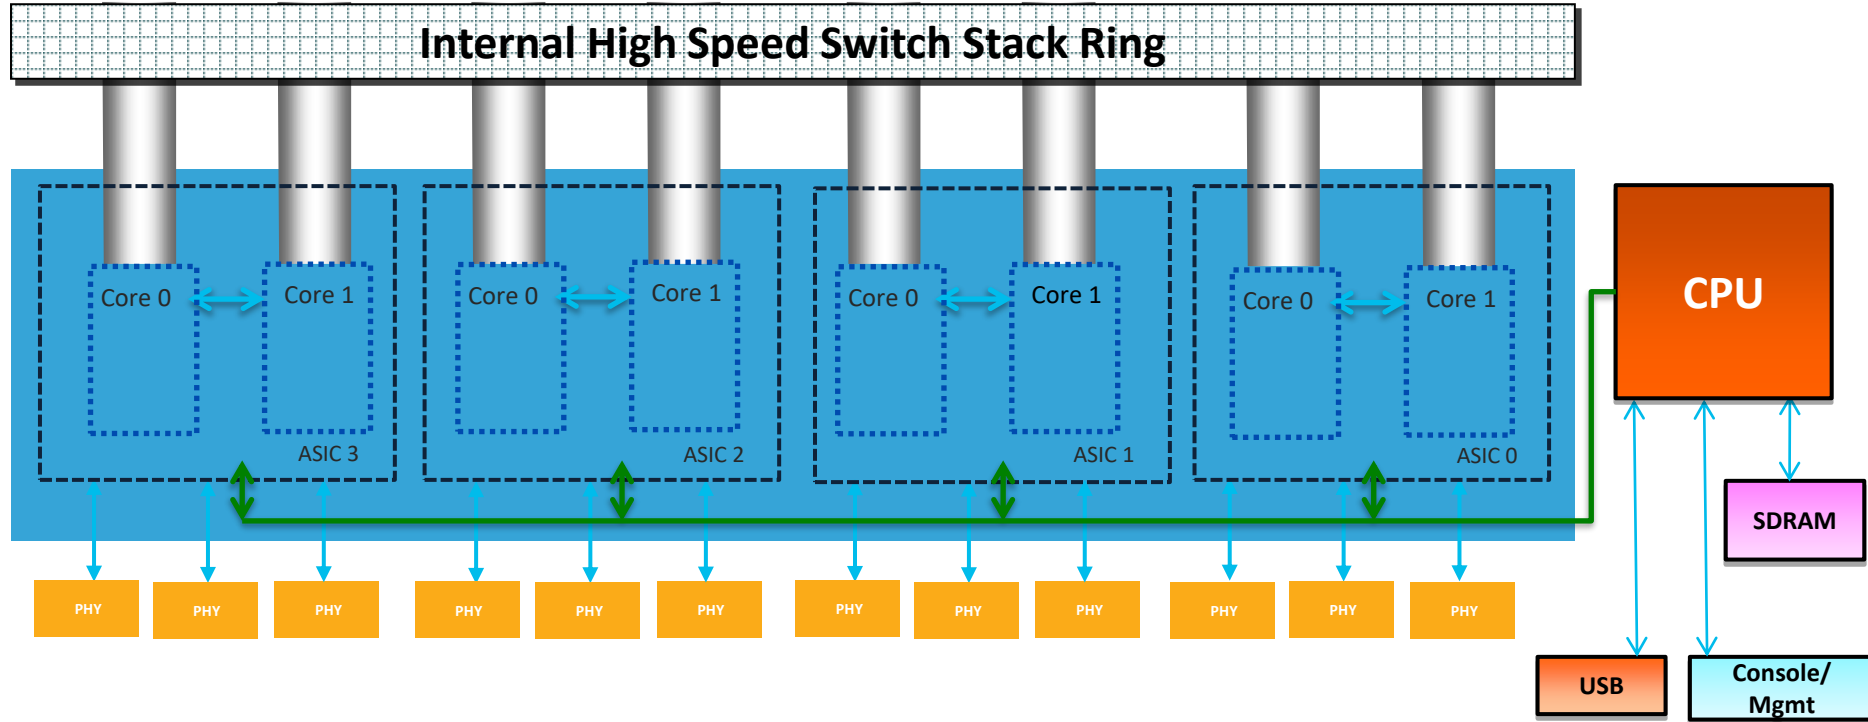
Catalyst 9200/9300 Stackable Switches



Catalyst 9400/9600 Chassis based



Catalyst 9500 Fixed Switches



Interface Internal Mappings

Interface to ASIC mapping important to understand data flows

```
9300_1#show platform software fed switch active ifm mappings
```

Interface	IF_ID	Inst	Asic	Core	Port	SubPort	Mac	Cntx	LPN	GPN	Type	Active
GigabitEthernet1/0/1	0x8	1	0	1	0	0	26	6	1	1	NIF	Y
GigabitEthernet1/0/2	0x4c	1	0	1	1	0	6	7	2	2	NIF	Y
GigabitEthernet1/0/3	0x4d	1	0	1	2	0	28	8	3	3	NIF	Y

Internally used interface addressing:

- LPN : Local Port Number
- GPN : Global Port Number
- IF_ID : Interface Identification, used for many fed CLI
- Type : Type of interface, NIF = Network Interface
- Inst : Instance : ASIC + Core
- Port : Asic Ports
- Active : Is Interface Active , multipurpose ports

IFM Mappings 9400/9500/9600

9500_1#sh platform software fed switch active ifm mappings | inc 1/./[1] |Int

Interface	IF_ID	Inst	Asic	Core	Port	SubPort	Mac	Cntx	LPN	GPN	Type	Active
TenGigabitEthernet1/0/1	0x2c	1	0	1	0	0	11	0	1	1	NIF	Y
TenGigabitEthernet1/1/1	0x18	0	0	0	16	0	0	19	17	17	NIF	N
FortyGigabitEthernet1/1/1	0x3c	0	0	0	24	0	4	4	25	25	NIF	Y

9600_1#sh platform software fed active ifm mappings | inc /0/[12][5] |Int

Interface	IF_ID	Inst	Asic	Core	Port	SubPort	Mac	Cntx	LPN	GPN	Type	Active
FortyGigabitEthernet1/0/15	0x17	3	1	1	24	0	30					
HundredGigE1/0/25	0x21	0	0	0	0	0	0					
TwentyFiveGigE2/0/15	0x47	1	0	1	21	0	21					
TwentyFiveGigE2/0/25	0x51	3	1	1	7	0	7					

Different mappings on different platforms

C9407R#sh platform software fed active ifm mappings | inc /0/1 |Int

Interface	IF_ID	Inst	Asic	Core	Port	SubPort	Mac	Cntx	LPN	GPN	Type	Active
GigabitEthernet2/0/1	0x7	1	0	1	0	0	7	7	1	201	NIF	Y
TenGigabitEthernet3/0/1	0x7f	4	2	0	19	0	19	3	1	301	NIF	Y
TenGigabitEthernet4/0/1	0x7e	5	2	1	3	0	11	3	1	401	NIF	Y
GigabitEthernet5/0/1	0x4b	3	1	1	0	0	4	4	1	501	NIF	Y

More detail IF_ID mappings

```
9300_1#sh platform software fed switch active ifm interfaces ?
```

access-tunnel	Access Tunnel Interface information
detail	Information for all type interfaces
efp	EFP Interface Information
ethernet	Interface ethernet information
internal	Internal Interface information
l2-lisp	Layer 2 LISP Interface information
lisp	LISP Interface information
loopback	Interface loopback information
nve	NVE Interface Information
svi	Interface SVI information
sw-subif	SW Sub-interface information
tunnel	Layer 3 Tunnel Interface Information
vfi-segment	L2VPN VFI Segment Interface Information
virtualportgroup	Interface vpg information
vlan	Interface VLAN information
wired-client	Interface wired client information

```
9300_1#sh plat soft fed sw act ifm int svi
```

Interface	IF_ID	State

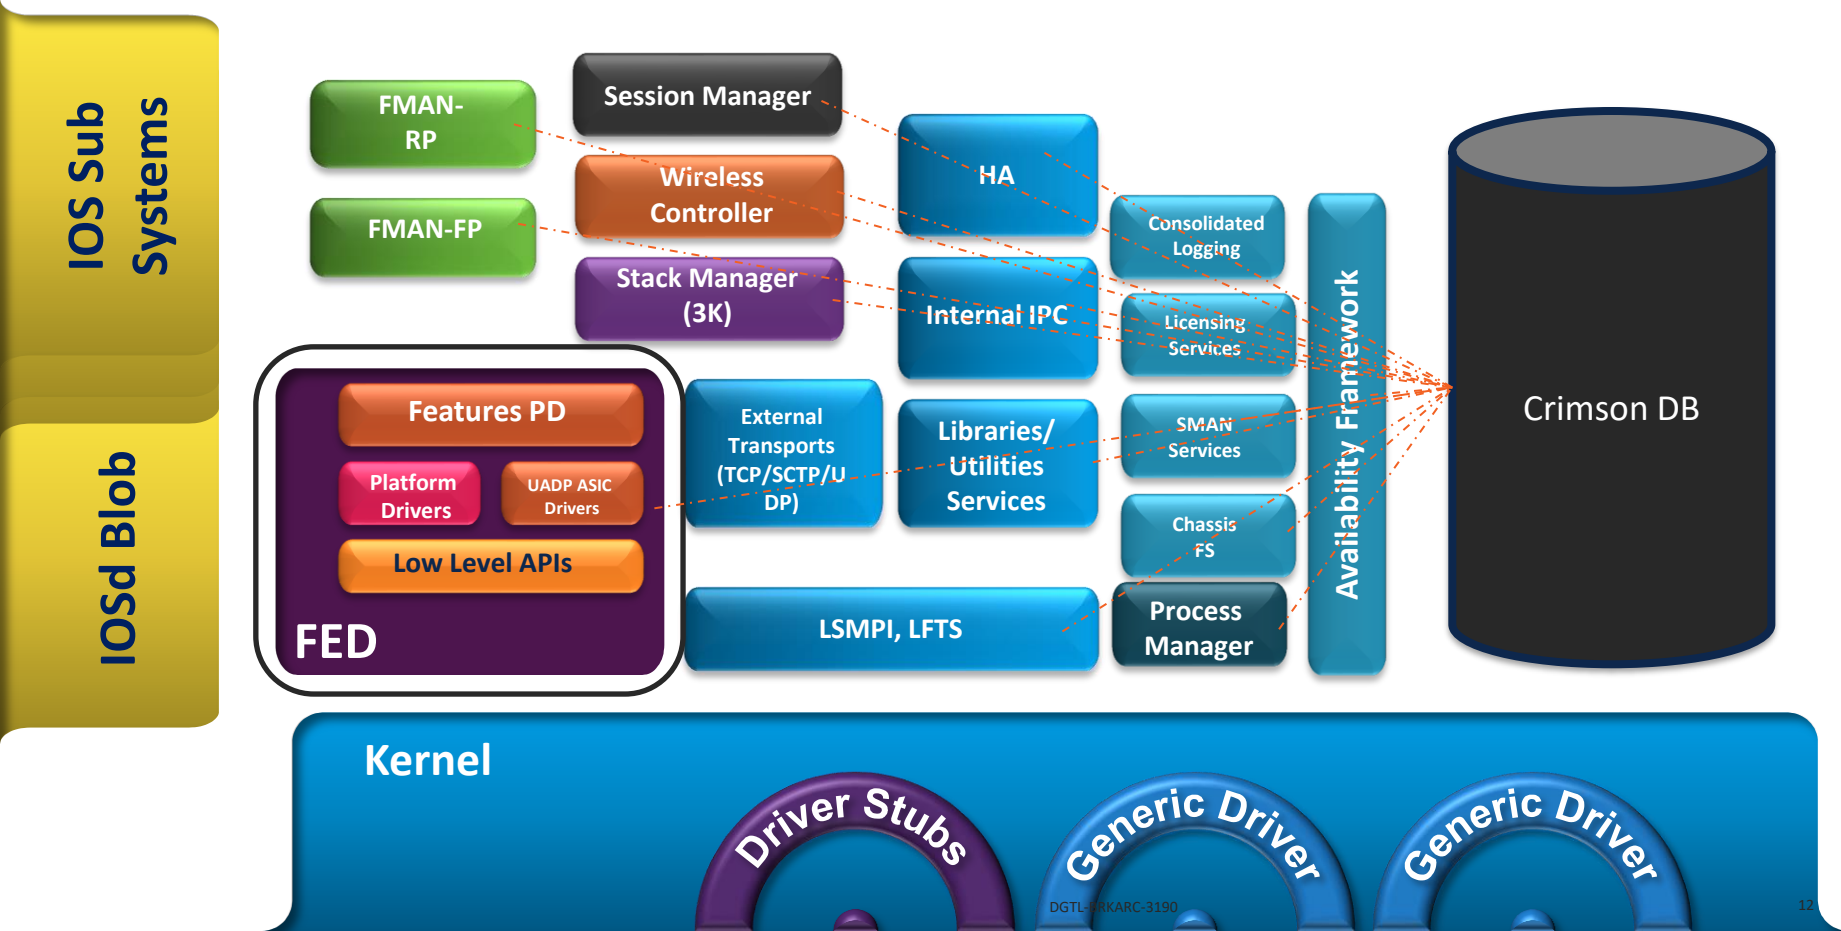
Vlan1	0x0000004b	READY
Vlan100	0x0000005a	READY

All physical and logical
Interfaces have an IF_ID

More detail about interfaces can be gathered with

show platform software fed switch <switch> ifm if-id <id>

IOS-XE 16 & IOS-XE 17 graphical overview



Kernel CPU information

```
9300_1#sh proc cpu platform sorted location switch active R0
CPU utilization for five seconds: 1%, one minute: 1%, five minutes: 1%
Core 0: CPU utilization for five seconds: 3%, one minute: 2%, five minutes: 2%
Core 7: CPU utilization for five seconds: 1%, one minute: 1%, five minutes: 0%

          5Sec    1Min    5Min  Status      Size  Name
-----
14416    13034     6%     6%     6%  S          223452  fed main event
10014     9623     1%     1%     1%  S           52212  sif_mgr
 9738     9215     1%     1%     1%  S          818660  linux_iosd-imag
```

- Underlying kernel shows Multi Core Architecture
- IOSd runs as process on kernel-> linux_iosd-image
- Not all processes run on all switches

```
Switch_1#sh processes cpu platform sorted location switch 1 R0 | in fman
27777    26990     0%     0%     0%  S          314179584  fman_rp
19145    17642     0%     0%     0%  S          296591360  fman_fp_image
Switch_1#sh processes cpu platform sorted location switch 3 R0 | in fman
20643    19400     0%     0%     0%  S          296599552  fman_fp_image
```

IOSd CPU utilization

```
9300_1#show processes cpu sorted
CPU utilization for five seconds: 1/0% one minute: 0%; five minutes: 0%
PID Runtime(ms)      Invoked      uSecs      5Sec      1Min      5Min  TTY Process
434      8197      1230039          6    0.07%    0.01%    0.00%    0 MMON MENG
203     10890      614953         17    0.07%    0.03%    0.01%    0 VRRS Main thread
287          8          46        173    0.07%    0.00%    0.00%    0 Exec
221     12377     1226864         10    0.07%    0.02%    0.00%    0 IP ARP Retry Age
113     11806      20043        589    0.07%    0.03%    0.02%    0 Crimson flush tr
218     12527     1226864         10    0.07%    0.02%    0.01%    0 IPAM Manager
238     11425      393615         29    0.07%    0.02%    0.01%    0 UDLD
```

- Many process like still running as processes inside IOSd process (BGP, RIP, CEF, ARP, UDLD, CDP , etc)
- CPU utilization 1%/0% showing Utilization for Processes and Utilization for Interrupt
- IOSd does not provide multicore architecture

Platform Memory

```
9300_1#sh processes memory platform sorted location switch 1 R0
```

```
System memory: 7711304K total, 2388036K used, 5323268K free,
```

```
Lowest: 5321604K
```

Pid	Text	Data	Stack	Dynamic	RSS	Name
9738	151604	817392	136	408	817392	linux_iosd-imag
14416	152	224856	136	85308	224856	fed main event
21595	238	168844	136	2780	168844	dbm
20109	176	128988	136	5404	128988	sessmgrd
21225	6825	117488	136	3004	117488	fman_rp
22061	260	93396	136	124	93396	cli_agent
22864	600	91152	136	16444	91152	smand

IOSd

FED

SMD

Forwarding
manager

- Kernel memory utilization is available per switch
- linux_iosd-image process is IOSd
- Resident Set Size(RSS), memory occupied by each Process

IOSd Memory

```
9300_1#sh processes memory sorted
```

```
Processor Pool Total: 1445417856 Used: 290878080 Free: 1154539776
```

```
lsmapi_io Pool Total: 6295128 Used: 6294296 Free: 832
```

PID	TTY	Allocated	Freed	Holding	Getbufs	Retbufs	Process
0	0	288949984	53306504	214634744	0	0	*Init*
4	0	22511680	100104	22277344	0	0	RF Slave Main Th
81	0	24107152	2420648	13717584	0	0	IOSD ipc task
472	0	4133424	105760	4069608	849828	0	EEM ED Syslog
0	0	62739512	58194512	2978824	23259559	382788	*Dead*
609	0	6717728	3968088	2803904	0	0	ISIS Upd
490	0	1719800	90880	1659432	0	0	EEM Server

- IOSd runs as a process , but does still provides some kernel features like memory management for all processes running inside IOSd
- Processor Pool: Pool for Processes on IOSd
- lsmapi_io : Linux Shared Memory Punt Interface memory , IO buffers

Catalyst IOS-XE Software release schedule

	Everest 16.6	Fuji 16.9	Gibraltar 16.12	Amsterdam 17.1	Amsterdam 17.2
Next planned release	16.6.9	16.9.6	16.12.4	17.1.2	17.2.2
9200	No	No	Yes	Yes	Yes
9300/9400/9500	Yes	Yes	Yes	Yes	Yes
9600	No	No	Yes	Yes	Yes
Maintenance Throttle	Yes	Yes	Yes	No	No

Maintenance throttles receive more rebuilds and thus would be recommended over feature releases

Not all SKU support all IOS-XE versions, above table indicative

Recommended releases: <https://www.cisco.com/c/en/us/support/docs/switches/catalyst-9300-series-switches/214814-recommended-releases-for-catalyst-9200-9.html>

Troubleshooting Tools

Platform Specific Show commands

- Platform Independent Show commands are similar like on any IOS platform.
Ex. Show cdp, show bgp, show udd
- IOS-XE specific troubleshooting commands are under show platform

```
Switch#show platform software ip switch active R0 cef  
ASR_1k#show platform software ip rp active cef
```

- *Show platform [hardware/software] fed* contain Platform Dependent troubleshooting for the Catalyst 3k and 9k (FED layer)

```
9300#show platform software fed switch active ifm mappings  
9300#show platform software fed switch 5 ifm mappings  
9400#show platform software fed active ifm mappings
```

- Outputs might vary depending on if it executed on active/standby and or/on Forwarding Processor or Route Processor

Debugging IOSd processes

- Features running inside IOSd can be debugged in traditional methods/debugs

```
9300_1#debug arp
ARP packet debugging is on
9300_1#debug ip icmp
ICMP packet debugging is on
```

- Some processes on IOSd support additional event-tracing inside IOSd

```
9300_1#sh monitor event-trace arp all
*Jan 20 03:52:48.977: RCV PKT: S: 10.100.10.100 T: 10.100.10.1 IF: Vlan100 OP: rep
*Jan 20 03:52:52.740: RCV PKT: S: 172.17.11.2 T: 172.17.11.1 IF: Vlan1 OP: req
9300_1#sh monitor event-trace bgp all
Tracing currently disabled, from exec command
```

Global Config:
monitor event-trace ..

Debugging, non-IOSd Processes

- IOSd cannot do real time debugging on processes outside IOSd
- To facilitate debugging/logging trace logs are available per process
- Tracing levels set with granularity (default notice). Tracing always on
- Common processes: smd , fed, forwarding-manager

```
Switch#set platform software trace smd switch active R0 dot1x-all debug
Switch#show platform software trace level smd switch active R0 | inc dot1x
dot1x          Notice
dot1x-all      Debug
dot1x-redun     Notice
```

- Processes *can* run on any switch in the stack, not just on active

Always on Tracing usage

- Tracelog files are stored in crashinfo:/logs in binary format.
- Traces can be displayed using show platform software trace command
- Archive of traces can be created using the command
“request platform software trace archive”
- Archives contain binary files, not readable with text viewer

Example fed logs from switch 1.

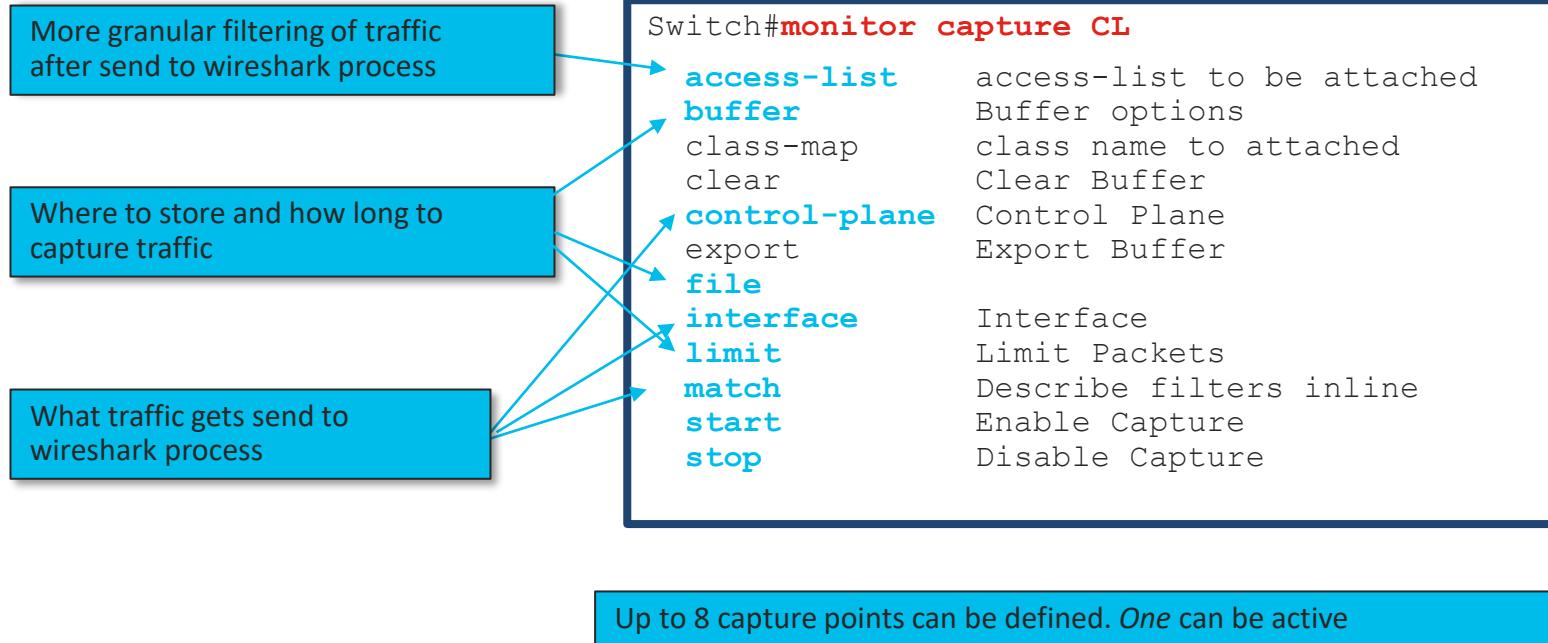
‘| inc dot1x’ filters out dot1x logs

```
Switch#show platform software trace message smd switch active | include dot1x
2020/01/29 21:49:20.540 {smd_R0-0}{1}: [dot1x] [19244]: UUID: 0, ra: 0, TID: 0 (debug): [0000.0000.0000:Gi1/0/1]
Received pkt saddr = 70e4.22e5.c4f7 , daddr = 0180.c200.0003, pae-ether-type = 0x888E
2020/01/29 21:49:20.540 {smd_R0-0}{1}: [dot1x] [19244]: UUID: 0, ra: 0, TID: 0 (debug): [0000.0000.0000:Gi1/0/1]
Queuing an EAPOL pkt on Authenticator Q
2020/01/29 21:49:20.540 {smd_R0-0}{1}: [dot1x] [19244]: UUID: 0, ra: 0, TID: 0 (debug): [0000.0000.0000:unknown]
Pkt body: 02 a3 00 16 04 10 12 ee 6c c4 9e 43 29 ef dd f4 0e 9c ae 75 34 0a
2020/01/29 21:49:20.540 {smd_R0-0}{1}: [dot1x] [19244]: UUID: 0, ra: 0, TID: 0 (info): [70e4.22e5.c4f7:Gi1/0/1]
EAP Packet - RESPONSE, ID : 0xa3
2020/01/29 21:49:20.540 {smd_R0-0}{1}: [dot1x] [19244]: UUID: 0, ra: 0, TID: 0 (info): [70e4.22e5.c4f7:Gi1/0/1]
Received EAPOL packet - Version : 1,EAPOL Type : EAP, Payload Length : 22, EAP-Type = MD5-Challenge
2020/01/29 21:49:20.540 {smd_R0-0}{1}: [dot1x] [19244]: UUID: 0, ra: 0, TID: 0 (debug): [0000.0000.0000:unknown]
Received EAPOL pkt (size=92u) on 12 socket
```

Embedded Packet Capture (EPC)

- EPC provides insight into both Data Plane and Control Plane traffic
- Captures can be done on Interfaces or Control Plane
- Data can be buffered and exported or stored directly in flash
- Data capture implemented on Port Asics, traffic copied to EPC process.
- EPC process provides deeper packet capture and display filtering
- Analysis can be done off-box or on box using included packet dissectors

Running Embedded Packet Capture



Running Embedded Packet Capture

```
Switch#monitor capture CL interface GigabitEthernet 1/0/2 both
Switch#monitor capture CL match ipv4 any any
Switch#monitor capture CL limit duration 60
Switch#monitor capture CL file location flash:cl.cap
Switch#monitor capture CL start display brief
```

Where and what to capture

Starting the packet display Press Ctrl + Shift + 6 to exit

```
 1  0.000000 10.254.111.100 -> 10.254.254.1 TCP 74 734 b^F^R 2049 [SYN]
Seq=0 Win=29200 Len=0 MSS=1460 SACK_PERM=1 TSval=445826583 TSecr=0 WS=128
 2  0.000501 10.254.254.1 -> 10.254.111.100 TCP 60 2049 b^F^R
Seq=1 Ack=1 Win=0 Len=0
 3  1.001299 10.254.111.100 -> 10.254.254.1 TCP 74 711 b^F^R
Seq=0 Win=29200 Len=0 MSS=1460 SACK_PERM=1 TSval=445826833 TSecr=0 WS=128
 4  1.001582 10.254.254.1 -> 10.254.111.100 TCP 60 2049 b^F^R 711 [RST, ACK]
Seq=1 Ack=1 Win=0 Len=0
```

Storage location and duration

Capture CLUS stopped - Capture duration limit reached

Starts the capture and enables live capture

Displaying packet captures

On Box Analysis of saved captures

```
Switch#show monitor capture file flash:cl.cap brief
```

```
Starting the packet display ..... Press Ctrl + Shift + 6 to exit
```

```
1 0.000000 10.200.10.100 -> 10.200.10.200 TCP 66 59498 b^F^R 80 [ACK] Seq=1 Ack=1
```

```
Switch#show moni capture file flash:cl.cap packet-number 1 detailed | be Transmission
```

```
Transmission Control Protocol, Src Port: 59498 (59498), Dst Port: 80 (80), Seq: 1, Ack:
```

```
1, Source Port: 59498
```

```
Destination Port: 80
```

```
Sequence number: 1 (relative sequence number)
```

```
Acknowledgment number: 1 (relative ack number)
```

```
Header Length: 32 bytes
```

```
Flags: 0x010 (ACK)
```

```
000. .... = Reserved: Not set
```

```
...0 .... = Nonce: Not set
```

```
.... 0... = Congestion Window Reduced (CWR): Not set
```

```
.... .0.. = ECN-Echo: Not set
```

```
.... ..0. = Urgent: Not set
```

```
.... ...1 .... = Acknowledgment: Set
```

```
.... .... 0... = Push: Not set
```

```
.... .... .0.. = Reset: Not set
```

```
.... .... ..0. = Syn: Not set
```

```
.... .... ...0 = Fin: Not set
```

```
[TCP Flags: *****A****]
```

```
Window size value: 24464
```

Details packet decodes

Off Box Analysis using
Wireshark also possible

SPAN

- Local SPAN, local traffic copied to interface on switch (Local)
- Remote SPAN , traffic copied to/from special rspan VLAN (Layer 2 domain)
- Encapsulated RSPAN, traffic encap/decap'd using GRE (Layer 3 domain)

```
monitor session 1 type erspan-source
source interface Gi2/0/3
destination
  erspan-id 123
  ip address 10.48.91.180
origin ip address 172.31.255.12
```

Using ERSPAN to send traffic
from Gi2/0/3 to PC

```
Frame 25670: 103 bytes on wire (824 bits), 103 bytes captured (824 bits) on interface 0
Ethernet II, Src: Cisco_5b:4b:00 (6c:41:6a:5b:4b:00), Dst: Vmware_b5:c6:36 (00:50:56:b5:c6:36)
Internet Protocol Version 4, Src: 172.31.255.12, Dst: 10.48.91.180
  0100 .... = Version: 4
  .... 0101 = Header Length: 20 bytes (5)
  Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
  Total Length: 89
  Identification: 0x006b (107)
  Flags: 0x0000
  ...0 0000 0000 0000 = Fragment offset: 0
  Time to live: 248
  Protocol: Generic Routing Encapsulation (47)
  Header checksum: 0xb0fa [validation disabled]
  [Header checksum status: Unverified]
  Source: 172.31.255.12
  Destination: 10.48.91.180
Generic Routing Encapsulation (ERSPAN)
Encapsulated Remote Switch Packet Analysis Type II
  0001 .... = Version: Type II (1)
  .... 0000 0000 0000 = Vlan: 0
  000. .... = COS: 0
  ...1 1... .... = Encap: VLAN tag preserved in frame (3)
  .... 0... .... = Truncated: Not truncated (0)
  .... ..00 0111 1011 = SpanID: 123
  0000 0000 0000 .... = Reserved: 0
  .... 0000 0000 1100 1000 0000 = Index: 3200
IEEE 802.3 Ethernet
  Destination: Spanning-tree-(for-bridges)_00 (01:80:c2:00:00:00)
  Source: Cisco_59:e5:7e (e4:aa:5d:59:e5:7e)
  Length: 39
Logical-Link Control
Spanning Tree Protocol
```

Embedded Event Manager, reacting to events

Definition

```
event manager applet RouteChange authorization bypass
```

```
event routing network 10.48.91.0/24 type all vrf Bru
```

```
action 0.1 cli command "enable"
```

```
action 1.0 cli command "show ip bgp vpnv4 all | append flash:rclog.txt"
```

```
action 1.1 cli command "show ip route vrf Bru | append flash:rclog.txt"
```

```
action 2 syslog msg "Route change detected, flash:rclog.txt appended"
```

Trigger

Actions

- EEM framework allows actions to be taken based on a number of triggers
- Assists with information gathering and possibly corrective actions
- Automated capture allows information to be gathered close to an event

Show Tech Enhancements

- Show tech contains lot of generic information, not feature specific
- For more focused information gathering show tech <keyword>
- Scripted command generation based on provided parameters
- Examples:

```
show tech-support cts  
show tech-support port  
show tech-support acl  
show tech-support fabric
```

Show tech can be large,
redirect to flash

```
9300_1#sh tech identity mac 0001.0001.0001 interface Gi 1/0/1 | redirect flash:shtech.log  
9300_1#dir flash:shtech.log  
Directory of flash:/shtech.log  
671754  -rw-          1504931  Jun 10 2019 00:07:47 +00:00  shtech.log  
11353194496 bytes total (9337597952 bytes free)
```

Show tech enhancement example

- Show tech enhancements go further then just show commands
- Results of outputs are used to generate and execute follow up commands

```
9300_1#sh tech-support cef ipv4 10.48.91.128/25 | re flash:shtechcef.log
9300_1#more flash:shtechcef.log | inc ---- show
----- show clock -----
----- show version -----
----- show running-config -----
----- show cef state -----
----- show cef state capabilities -----
----- show cef table internal -----
----- show ip cef vrf Default 10.48.91.128/25 internal -----
----- show ip route 10.48.91.128 255.255.255.128 -----
----- show interface TenGigabitEthernet1/1/6 -----
----- show cef interface TenGigabitEthernet1/1/6 internal -----
----- show adj link IP TenGigabit1/1/6 172.31.250.30 connectionid 0 internal -----
----- show arp 172.31.250.30 TenGigabitEthernet1/1/6 detail ---
----- show monitor event-trace cef ipv4 vrf Default 10.48.91.128 all ----
```

Programmability & Troubleshooting

- IOSXE Programmability helps gathering information and detecting failures
- Ability to execute CLI on IOSd, parse and possibly react to events
- Python (Guest shell)
Python on guestshell allows interaction with IOSd, config, cli etc

```
9300_1#guestshell enable
9300_1#guestshell run python3
Python 3.6.3 (default, Nov 1 2018, 15:47:26)
[GCC 4.8.5 20150623 (Red Hat 4.8.5-36)] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> import cli
>>> print (cli.execute("show run | inc app-iox" ))
iox
app-hosting appid guestshell
app-vnic management guest-interface 0
```

Programmability & Troubleshooting

- TCL included in IOS-XE.
- Does not require guestshell to be enabled
- TCL scripts can be sourced from filesystem or inputted

```
9300_1#tclsh  
9300_1(tcl)#puts [ exec "show ver" ]  
Cisco IOS XE Software, Version 17.01.01  
Cisco IOS Software [Amsterdam], Catalyst L3 Switch Software (CAT9K_IOSXE), Version 17.1.1, RELEASE  
SOFTWARE (fc3)  
Technical Support: http://www.cisco.com/techsupport
```

```
9300_1#tclsh tftp://10.48.91.151/cl.tcl  
Loading cl.tcl from 10.48.91.151 (via TenGigabitEthernet1/1/6): !  
[OK - 418 bytes]  
Number of Up CLNS neighbours 4 Number of up BFD sessions 4  
Number of sessions match
```


Platform Issues

Platform Health overview

Overview of Memory & CPU usages

```
9300_1#show platform resources
```

```
**State Acronym: H - Healthy, W - Warning, C - Critical
```

Resource	Usage	Max	Warning	Critical	State
Control Processor	1.39%	100%	5%	10%	H
DRAM	2336MB (31%)	7530MB	90%	95%	H

```
9300_1#show environment all
```

```
Switch 1 FAN 1 is OK
```

```
Switch 1 FAN 2 is OK
```

```
Switch 1 FAN 3 is OK
```

```
FAN PS-1 is OK
```

```
FAN PS-2 is NOT PRESENT
```

```
Switch 1: SYSTEM TEMPERATURE is OK
```

```
Inlet Temperature Value: 37 Degree Celsius
```

```
Temperature State: GREEN
```

```
Yellow Threshold : 46 Degree Celsius
```

```
Red Threshold : 56 Degree Celsius
```

```
Hotspot Temperature Value: 56 Degree Celsius
```

```
Temperature State: GREEN
```

```
Yellow Threshold : 105 Degree Celsius.
```

```
Red Threshold : 125 Degree Celsius
```

SW	PID	Serial#	Status	Sys Pwr	PoE Pwr	Watts
----	-----	---------	--------	---------	---------	-------

1A	PWR-C1-715WAC	DCA2120G1D4	OK	Good	Good	715
----	---------------	-------------	----	------	------	-----

```
1B Not Present
```

Yellow threshold and Fan Fail action configurable.

CPU Inband Traffic

- Dataplane forwarding occurring on UADP without CPU interaction
- CPU Destined traffic sent to corresponding CPU queue by UADP
- CPU queues:
 - Routing Protocol, ex EIGRP, OSPF, PIM, etc
 - Layer 2 Control , ex Spanning Tree, UDLD , etc
 - Software Forwarding, Traffic hitting exception, TCAM overload
 - Host Traffic (For us), ex SNMP, SSH, NETCONF, ICMP
 - ICMP, traffic needing ICMP generation
 - ..
- Control Plane policing default *enabled* to protect CPU resources

Inband Packet Statistics

```
9300_1#sh controllers cpu-interface
```

queue	retrieved	dropped	invalid	hol-block

Routing Protocol	42856	0	0	0
L2 Protocol	52577	0	0	0
sw forwarding	1	0	0	0
broadcast	11360	0	0	0
icmp gen	0	0	0	0
icmp redirect	0	0	0	0
logging	0	0	0	0
rpf-fail	0	0	0	0
DOT1X authentication	0	0	0	0
Forus Traffic	120	0	0	0
Forus Resolution	46750	0	0	0
Inter FED	0	0	0	0
L2 LVX control	0	0	0	0
EWLC control	0	0	0	0
EWLC data	0	0	0	0
L2 LVX data	0	0	0	0
Openflow	0	0	0	0
Topology control	85166	0	0	0
Proto snooping	0	0	0	0
DHCP snooping	0	0	0	0

- Counts packets received at CPU
- Different Queues to handle different classes of traffic for CPU

Control Plane Policing Statistics

```
9300_1#show policy-map control-plane
Control Plane
Service-policy input: system-cpp-policy
Class-map: system-cpp-police-ios-routing (match-any)
  0 packets, 0 bytes
  5 minute offered rate 0000 bps, drop rate 0000 bps
  Match: none
  police:
    rate 13000 pps, burst 3173 packets
    conformed 379638519 bytes; actions:
      transmit
    exceeded 0 bytes; actions:
      drop
Class-map: system-cpp-police-ios-feature (match-any)
  0 packets, 0 bytes
  5 minute offered rate 0000 bps, drop rate 0000 bps
  Match: none
  police:
    rate 6000 pps, burst 1464 packets
    conformed 20422413 bytes; actions:
      transmit
    exceeded 0 bytes; actions:
      drop
```

- Multiple Queues might map to one class map
- COPP policy map configurable (not recommended)
- Drops in critical queues can lead to instability

Control Plane Policing HW stats

```
9300_1#show plat hardware fed switch active qos queue stats internal cpu policer
```

CPU Queue Statistics

QId	PlcIdx	Queue Name	Enabled	(default) Rate	(set) Rate	Queue Drop (Bytes)	Queue Drop (Frames)
0	11	DOT1X Auth	Yes	1000	1000	0	0
1	1	L2 Control	Yes	2000	2000	0	0
2	14	Forus traffic	Yes	4000	4000	0	0

Per queue stats

CPU Queue Policer Statistics

Policer Index	Policer Accept Bytes	Policer Accept Frames	Policer Drop Bytes	Policer Drop Frames
0	17261371	11408	0	0
1	17682901	52775	0	0
2	357304765	43037	0	0

Policer stats

CPP Classes to queue map

PlcIdx	CPP Class	Queues
0	system-cpp-police-data	ICMP GEN/ BROADCAST/ ICMP Redirect/
10	system-cpp-police-sys-data	Openflow/ Exception/ EGR Exception/ NFL SAMPLED DATA/ RPF
13	system-cpp-police-sw-forward	Sw forwarding/ LOGGING/ L2 LVX Data Pack/ Transit Traffic

Mappings

Determining where CPU traffic originates from

```
9500_1#show plat software fed switch active punt rates interfaces
```

Punt Rate on Interfaces Statistics

Packets per second averaged over 10 seconds, 1 min and 5 mins

Active Interfaces
sending to CPU

=====							
Interface Name	IF_ID	Recv 10s	Recv 1min	Recv 5min	Drop 10s	Drop 1min	Drop 5min
=====							
TenGigabitEthernet1/0/1	0x0000002e	1	1	1	0	0	0
FortyGigabitEthernet1/1/1	0x00000032	1	1	18	0	0	0

```
9500_1#show platform software fed switch active punt cpuq rates
```

Punt Rate CPU Q Statistics

Packets per second averaged over 10 seconds, 1 min and 5 mins

Per Queue Statistics

=====								
Q	Queue	Rx	Rx	Rx	Drop	Drop	Drop	
no	Name	10s	1min	5min	10s	1min	5min	
=====								
2	CPU_Q_FORUS_TRAFFIC	0	0	17	0	0	0	
4	CPU_Q_ROUTING_CONTROL	3	3	3	0	0	0	

Ternary Content-Addressable Memory

- UADP asics use TCAM memory for storing a majority of the forwarding related tables
- Normal Memory has 0 or 1. TCAM has 0, 1 and “don’t care”
- Provides ability to do fast lookups needed for forwarding
- TCAM stores tables with Values , Mask and Result (VMR)
- TCAM memory shared by various features. Allocation of TCAM space done using Switching Database Manager (SDM)
- Different SDM templates available depending on platform
- SDM changes require reload
- Exceeding TCAM space impacts ability to do Hardware forwarding. Depending on feature SW forwarding, drops or flooding

```
9500#show sdm prefer ?
core          Core Template
distribution  Distribution Template
nat           NAT Template
sda           SDA Template
```


Monitoring TCAM Utilization

```
9300_1#show platform hardware fed switch active fwd-asic resource tcam utilization  
CAM Utilization for ASIC [0]
```

Table	Max Values	Used Values
-----	-----	-----
Unicast MAC addresses	32768/1024	81/21
L3 Multicast entries	8192/512	0/7
L2 Multicast entries	8192/512	0/9
Directly or indirectly connected routes	24576/8192	24/37
QoS Access Control Entries	5120	40
Security Access Control Entries	5120	125
Ingress Netflow ACEs	256	8
Policy Based Routing ACEs	1024	20
Egress Netflow ACEs	768	8
..		

- Not all Port Asics get programmed equally.
- Exceeding TCAM resources could lead to performance issues, traffic might be software forwarded, flooded or dropped

Dumping TCAM entries

- TCAM information can be displayed per switch and per ASIC
- For forwarding check on both ingress Asic and egress Asic
- TCAM dumps show raw information, content might not be easily readable

```
9300_1#sh plat hard fed sw act fwd-asic resource tcam table ?
acl                L3 IPv4 ACL Table
client_le          Client LE Table
l2                 L2 Lookup Table
l2_multicast       L2 IPv4 Multicast Forwarding Table
l3_multicast       L3 IPv4 Multicast Forwarding Table
qos                L3 IPv4 QoS Table
sghash             SGT/DGT Lookup Table
spd_lookup         SPD Lookup Table
tunnel             Tunnel Termination Table
unicast            L3 Unicast Table
```

```
9300_1#sh platform hardware fed switch active fwd-asic resource tcam table l2
```

```
Printing entries for region L2_EXACT_MATCH (0) type 1 asic 0
```

```
=====
```

```
HashTable 0 Left bucket 593 Index 0 (A:0,C:0) Valid
```

```
KEY - vlan:53 mac:0xa0f849104867 l3_if:1 gpn:19 epoch:0 static:0 flood_en:0 vlan_lead_wless_flood_en: 0
client_home_asic: 0 learning_peerid 0, learning_peerid_valid 0
```


```
MASK - vlan:0 mac:0x0 l3_if:1 gpn:0 epoch:0 static:0 flood_en:0 vlan_lead_wless_flood_en: 0 client_home_asic: 0
learning_peerid 0, learning_peerid_valid 0
```

```
SRC_AD - need_to_learn:0 lrn_v:0 catchall:0 static_mac:0 chain_ptr_v:0 chain_ptr: 0 static_entry_v:0 auth_state:0
auth_mode:0 auth_behavior_tag:0 traf_m:0 is_src_ce:0
```


```
DST_AD - si:0x2d bridge:0 replicate:0 blk_fwd_o:0 v4_rmac:1 v6_rmac:1 catchall:0 ign_src_lrn:0 port_mask_o:0
afd_cli_f:0 afd_lbl:0 prio:3 dest_mod_idx:0 destined_to_us:1 pv_trunk:0 smr:0
```

Catalyst 9200, 9300 Stacking

- Stack connections are crucial to system stability
- 9200/9300 stackable switches deploy ring topology to form a stack
- One failure in stack recoverable. Dual Failure fatal
- Ensure proper fitting and tightening of stack connectors.



Dual Rings



Single Ring

Stack Port Overview

```
Switch#show switch stack-ports summary
```

Sw#/Port#	Port Status	Neighbor	Cable Length	Link OK	Link Active	Sync OK	#Changes to LinkOK	In Loopback
1/1	OK	3	50cm	Yes	Yes	Yes	1	No
1/2	OK	2	50cm	Yes	Yes	Yes	1	No
2/1	OK	1	50cm	Yes	Yes	Yes	1	No
2/2	OK	3	50cm	Yes	Yes	Yes	1	No
3/1	OK	2	50cm	Yes	Yes	Yes	1	No
3/2	OK	1	50cm	Yes	Yes	Yes	1	No

Neighbor switch on
port

Cable length should
be showing

All stack links should be Active and in sync.
Changes to LinkOK should be low

Stack numbers are NOT allocated based on physical presense.
Renumber switches or bring up sequentially

Switch Stacks

Switch# **show switch detail**

Switch/Stack Mac Address : **58bf.eab6.7f80** - Local Mac Address

Mac persistency wait time: Indefinite

			H/W	Current	
Switch#	Role	Mac Address	Priority	Version	State

1	Member	58bf.eab6.4200	1	V01	Ready
*2	Active	58bf.eab6.7f80	4	V01	Ready
3	Standby	58bf.eab7.5800	3	V01	Ready

Stack Port Status			Neighbors	
Switch#	Port 1	Port 2	Port 1	Port 2

1	OK	OK	3	2
2	OK	OK	1	3
3	OK	OK	2	1

Every stack has 1 Active switch. Standby and members are optional

Set Priority to designate Active/Standby Switches with switch x priority x

SVI interfaces take the mac address of stack.
Layer 2 interfaces of the switch

Show tech fabric command prints out stack related information as well as stack related registers and counters

Consoles

- On switch stacks all console ports are active, redirection occurs to active switch
- Redirected session is a vty(telnet), console on Active is con (console)
- * indicates current active session in use, internal IPs are not in routing tables


```
FE2050#show users
```

Line	User	Host(s)	Idle	Location
* 0 con 0	cisco	idle	00:00:00	
3 vty 2	cisco	idle	00:01:30	192.168.1.5

```
FE2050#show tcp brief
```

TCB	Local Address	Foreign Address	(state)
7FF9CC3DEE20	192.168.1.6.23	192.168.1.5.43170	ESTAB

- Chassis based switches do not use console redirection, use “active” Supervisor
- Standby console can be enabled
- Functionality limited on standby



```
9407R#sh run | sec redundancy
redundancy
mode sso
main-cpu
standby console enable
```

Unexpected Reloads

- When Switch (active/standby/member) unexpectedly reloads a system_report is generated in local crashinfo:
- Crashinfo file contains core file , traces, crashinfo files
- Crashinfo of member switches reachable via dir crashinfo-x:


```
guest@estg-bru-tftp:/tftpboot/cl$ tar -xzvf system-report_1_20190802-112928-UTC.tar.gz
/bootflash/.prst_sync/reload_info
/tmp/FE2050_1_RP_0-bootuplog-20190802-112928-UTC.log
/flash/core/FE2050_1_RP_0_x86_64_crb_linux_iosd_ngwc-universalk9-ms_9699_20190802-UTC.core.gz
/crashinfo/crashinfo_RP_00_00_20190802-112859-UTC
guest@estg-bru-tftp:/tftpboot/cl$ cat bootflash/.prst_sync/reload_info
ReloadReason=Critical software exception, check crashinfo:crashinfo_RP_00_00_20190802-112859-UTC
RET_2_RCALTS=1564745340 RET_2_RTS=11:29:00 UTC Fri Aug 2 2019
```

Crashinfo File

- Part of Crashinfo file contains:
 - Last interactions on CLI
 - Syslog messages

```
guest@estg-bru-tftp:/tftpboot/cl/crashinfo$ more crashinfo_RP_00_00_20190802-112859-UTC
```

```
CMD: 'sh device-tracking counters ' 11:28:53 UTC Fri Aug 2 2019
CMD: 'sh device-tracking counters vlan 1021' 11:28:57 UTC Fri Aug 2 2019
*Aug  2 11:28:57.445: %HA_EM-6-LOG: catchall: show device-tracking counters vlan 1021
CMD: 'sh device-tracking counters vlan 1022' 11:28:58 UTC Fri Aug 2 2019
*Aug  2 11:28:57.445: %HA_EM-6-LOG: catchall: show device-tracking counters vlan 1021
CMD: 'sh device-tracking counters vlan 1022' 11:28:58 UTC Fri Aug 2 2019
Exception to IOS Thread:
Frame pointer 0x7FD1C728CA28, PC = 0x562DFBC20D36
```



CSCvp72220 crash at sisf_show_counters after entering
show device-tracking counters command

- Memory Utilization
 - Buffer Pool Utilization
 - Uptime and IOS-XE version
- Looking through Crashinfo file can give clues about the state of the system prior to the unexpected reload



Packet Drops

Ethernet controller Statistics

```
9500_1#show controllers ethernet-controller te 1/0/1
```

Transmit	TenGigabitEthernet1/0/1	Receive
38633194582 Total bytes	130046767680	Total bytes
92351261 Unicast frames	100227375	Unicast frames
34062623315 Unicast bytes	116307995084	Unicast bytes
689338 Multicast frames	2043062	Multicast frames
4570570691 Multicast bytes	13738772084	Multicast bytes
9 Broadcast frames	8	Broadcast frames
576 Broadcast bytes	512	Broadcast bytes
0 System FCS error frames	0	IpgViolation frames
0 MacUnderrun frames	0	MacOverrun frames
..		
2241744 65 to 127 byte frames	11667872	65 to 127 byte frames
2838336 128 to 255 byte frames	3177806	128 to 255 byte frames
80722296 256 to 511 byte frames	1593500	256 to 511 byte frames

Ethernet Statistics
also in show tech

Error statistics
from interface

```
Switch#show interfaces gigabitEthernet 1/0/1 counters errors
```

Port	Align-Err	FCS-Err	Xmit-Err	Rcv-Err	UnderSize	OutDiscards
Gi1/0/1	0	0	0	0	0	0

Port	Single-Col	Multi-Col	Late-Col	Excess-Col	Carri-Sen	Runts
Gi1/0/1	0	0	0	0	0	0

Ethernet Interface Utilization

```
Switch#show interfaces | inc line|rate
```

```
Vlan1 is up, line protocol is up , Autostate Enabled
  5 minute input rate 0 bits/sec, 0 packets/sec
  5 minute output rate 0 bits/sec, 0 packets/sec
GigabitEthernet0/0 is administratively down, line protocol is down
  5 minute input rate 0 bits/sec, 0 packets/sec
  5 minute output rate 0 bits/sec, 0 packets/sec
GigabitEthernet1/0/1 is up, line protocol is up (connected)
  5 minute input rate 103000 bits/sec, 174 packets/sec
  5 minute output rate 3879000 bits/sec, 324 packets/sec
GigabitEthernet1/0/2 is down, line protocol is down (notconnect)
  5 minute input rate 0 bits/sec, 0 packets/sec
  5 minute output rate 0 bits/sec, 0 packets/sec
GigabitEthernet1/0/3 is down, line protocol is down (notconnect)
  5 minute input rate 0 bits/sec, 0 packets/sec
  5 minute output rate 0 bits/sec, 0 packets/sec
```

Average rate over interval set
with load interval
Does not always reflect current
load

```
9500 show controllers utilization
```

Port	Receive Utilization	Transmit Utilization
------	---------------------	----------------------

Tel/0/1	3	2
---------	---	---

Tel/0/2	2	3
---------	---	---

.

Tel/0/16	0	0
----------	---	---

Fo1/1/1	0	0
---------	---	---

Fo1/1/2	0	0
---------	---	---

Total Ports : 18

Total Ports Receive Bandwidth Percentage Utilization : 0


Total Ports Transmit Bandwidth Percentage Utilization : 0

Average Switch Percentage Utilization : 0

Bandwidth in %
Current load

Drops due to egress buffering

- WTD drops due to buffering most common cause of drops
- Data center servers typically high bandwidth 10Gb/s+
- Increasing number of buffers could lead to Jitter and Latency
- Protocols like TCP increase window size till congestion experienced



Tail Drops

- Tail drops occur when exceeding buffer thresholds.

```
9300_1#show interfaces gigabitEthernet 1/0/1 | inc output drops
Input queue: 0/2000/0/0 (size/max/drops/flushes); Total output drops: 1277
9300_1#show controllers ethernet-controller gig 1/0/1 | inc Excess Def
1277 Excess Defer frames          0 Collision fragments
SNMP:
SNMPv2-SMI::enterprises.9.2.2.1.1.27.8 = INTEGER: 1277
```

- Buffer allocation per class can be changed inside service-policy
- Global multiplier to increase buffers up to 1200%
qos softmax-queue-multiplier <percentage>
- Increasing buffers increases maximum, buffers allocated based on availability
- 17.1.1 introduces CLI to monitor high watermark utilization on ports

QoS Hardware configuration

```
9300_1#sh plat hard fed switch active qos queue config interface gi 1/0/1
Asic:0 Core:1 DATA Port:0 GPN:1 LinkSpeed:0x1
```

DTS			Hardmax			Softmax			PortSMin		GlblSMin		PortStEnd		
0	1	5	200		3200	5	500	0	0	6	9600				
1	1	4	0	13	4800	5	750	2	300	6	9600				
Priority			Shaped/shared			weight		shaping_step		sharpedWeight					
0		0	Shared			50		0		0					
1		0	Shared			75		0		0					
Port			Port			Port		Port							
Priority			Shaped/shared			weight		shaping_step							
2			Shaped			254		255							
Weight0		Max_Th0	Min_Th0		Weigth1		Max_Th1	Min_Th1		Weight2		Max_Th2	Min_Th2		
0		0	2709				3028				3400		0		
1		0	3825		0		0		4275		0		0		

- Hardmax.
Reserved buffers
- Softmax.
From Global Pool

- Queue mode
- Queue limit:
Step/weight * speed

- Drop thresholds
per queue in buffers

QoS hardware statistics

```
9300_1#sh platform hardware fed switch active qos queue stats interface gigabitEthernet 1/0/1
```

AQM Global counters

GlobalHardLimit: 7976 | GlobalHardBufCount: 0
GlobalSoftLimit: 11872 | GlobalSoftBufCount: 0

set plat hard fed active qos port-monitor interface <if>

High Watermark Soft Buffers: 0 <--- clear on read

Asic:0 Core:1 DATA Port:0 Hardware Enqueue Counters

Q	Buffers (Count)	Enqueue-TH0 (Bytes)	Enqueue-TH1 (Bytes)	Enqueue-TH2 (Bytes)	Qpolicer (Bytes)
0	0	0	385820	46085690	0
1	0	0	0	0	0


Asic:0 Core:1 DATA Port:0 Hardware Drop Counters

Q	Drop-TH0 (Bytes)	Drop-TH1 (Bytes)	Drop-TH2 (Bytes)	SBufDrop (Bytes)	QebDrop (Bytes)	QpolicerDrop (Bytes)
0	0	0	0	0	0	0
1	0	0	0	0	0	0

- At UADP level there are 8 Queues/3 Thresholds
- Enqueue/Drop Counters available per queue/per threshold
- Buffers (count) show currently assigned buffers to Queue (256 bytes)
- High water mark counter (17.x) shows highest watermark since last output

Buffer tuning , lab setup

- 10Gb/s Server, Debian 8.5
- 1Gb/s Client, Debian 8.5
- Interfaces on 9300 default
- Client and server use iperf3
- Both OS & iperf3 using default settings
- Achieved bandwidth , drops, high watermark & max Congestion window noted and



```
michelpe@DUT1:~$ iperf3 -v
iperf 3.1.3
Linux debian1 3.16.0-4-amd64 #1 SMP Debian 3.16.7-ckt25-2 (2016-04-08) x86_64
```


Buffer tuning , example

Cwnd = Congestion window

```
michelpe@DUT1:~$ iperf3 -c 10.200.10.100
```

```
Connecting to host 10.200.10.100, port 5201
```

```
[ 4] local 10.200.10.200 port 37127 connected to 10.200.10.100 port 5201
```

[ID]	Interval		Transfer	Bandwidth	Retr	Cwnd
[4]	0.00-1.00	sec	113 MBytes	951 Mbits/sec	6	167 KBytes
[4]	1.00-2.00	sec	112 MBytes	941 Mbits/sec	0	198 KBytes
[4]	2.00-3.00	sec	112 MBytes	941 Mbits/sec	0	225 KBytes
[4]	3.00-4.00	sec	112 MBytes	941 Mbits/sec	0	247 KBytes
[4]	4.00-5.00	sec	112 MBytes	941 Mbits/sec	1	195 KBytes
[4]	5.00-6.00	sec	112 MBytes	942 Mbits/sec	0	229 KBytes
[4]	6.00-7.00	sec	112 MBytes	941 Mbits/sec	0	250 KBytes
[4]	7.00-8.00	sec	112 MBytes	941 Mbits/sec	0	262 KBytes
[4]	8.00-9.00	sec	112 MBytes	941 Mbits/sec	1	208 KBytes
[4]	9.00-10.00	sec	112 MBytes	941 Mbits/sec	0	233 KBytes

Cwnd increases till drops occur

[ID]	Interval		Transfer	Bandwidth	Retr	
[4]	0.00-10.00	sec	1.10 GBytes	942 Mbits/sec	8	sender
[4]	0.00-10.00	sec	1.10 GBytes	941 Mbits/sec		receiver

Test Results

	No Service Policy					AutoQos Policy				
	Default	200	400	800	1200	Default	200	400	800	1200
Cwnd	273	421	990	1840	2003	117	233	421	762	1200
Drops	18	22	19	50	47	23	20	12	20	38
Speed	942	944	945	945	945	933	938	938	940	941
Buffers	1200	2400	4800	9600	14400	500	1000	2000	4000	6000
Watermark	1197	2394	4795	9499	10171	497	994	1195	3997	5999

- Increasing number of buffers increases max achieved congestion window
- Drops happen in burst, most drops happen at initial ramp up of Cwnd

Packet Parting packet drops

```
9300_1#sh platform hardware fed switch active fwd-asic drops exceptions
```

```
****EXCEPTION STATS ASIC INSTANCE 0 (asic/core 0/0)****
```

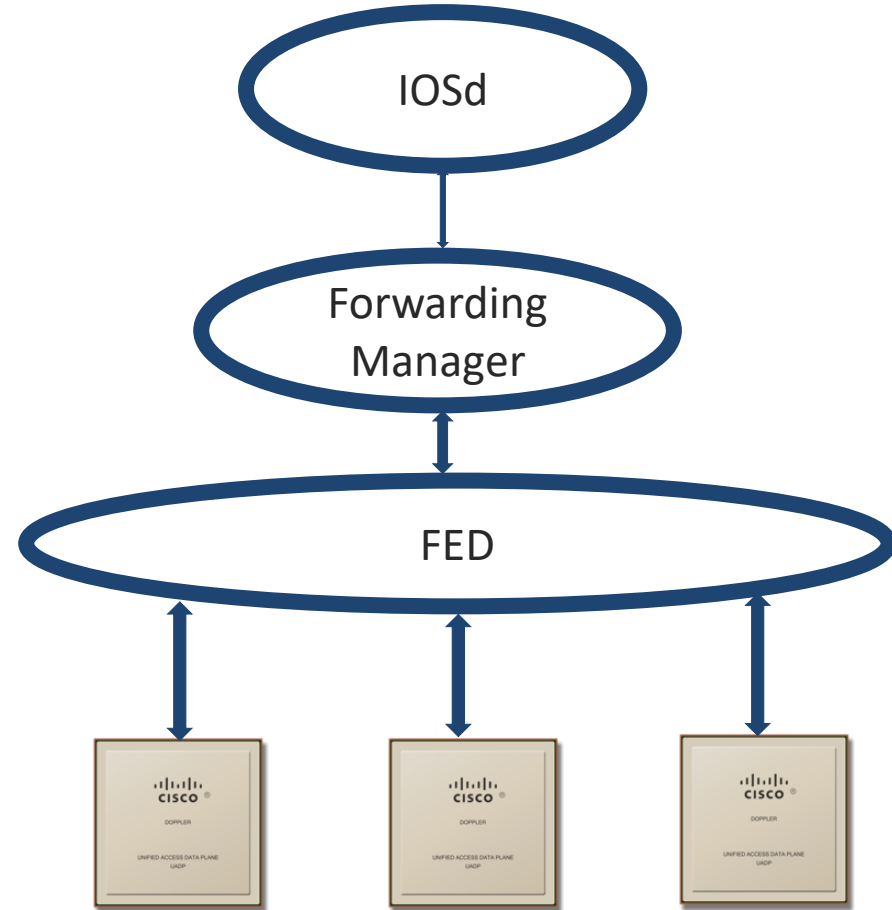
=====				
Asic/core	NAME	prev	current	delta
=====				
0 0	NO_EXCEPTION	35364016	35364108	92
0 0	IPV4_CHECKSUM_ERROR	0	0	0
0 0	ROUTED_AND_IP_OPTIONS_EXCEPTION	2	2	0
0 0	CTS_FILTERED_EXCEPTION	0	0	0
0 0	SIA_TTL_ZERO	0	0	0
0 0	ALLOW_NATIVE_EXCEPTION_COUNT	0	0	0
0 0	ALLOW_DOT1Q_EXCEPTION_COUNT	0	0	0
0 0	ALLOW_PRIORITY_TAGGED_EXCEPTION_COUNT	0	0	0
0 0	ALLOW_UNKNOWN_ETHER_TYPE_EXCEPTION	0	0	0
0 0	IP_SOURCE_GUARD_VIOLATION	0	0	0

- Every packet passing through Port Asic gets parsed,
Layer 2 switched packets also get fully parsed
- Exception drops are counted per Asic , no per port statistics

Forwarding Issues

Troubleshooting Forwarding

- UADP responsible for all hardware forwarding. CPU is not directly involved in forwarding.
- IOS-XE 16+ uses Forwarding manager, Forwarding manager process manages forwarding related information but only on a Platform Independent level. Forwarding manager is on all platforms running IOS-XE16+
- FED (Forwarding Engine Driver) process is Platform Dependant layer, interfaces with Forwarding manager and responsible for all tables on UADP asics



Forwarding troubleshooting, the easy way

- Show forward supported since 2900/3500XL switches , upto 3750 family only software emulation of forwarding results were used
- UADP introduced HW captures of lookup results during various stages of packet forwarding
- CLI: *"Show platform hardware fed switch <ingress switch> forward ... "*
- Supports Input using packet capture file or packet parameters
- Readability of results greatly improved since 16.9 with summary option
- Packets introduced for captures inserted at Ingress and removed at Egress before sending out on the wire.

Running Show platform hardware fed forward

```
9300_1#sh monitor capture file flash:icmp.pcap packet-number 11
Starting the packet display ..... Press Ctrl + Shift + 6 to exit
 11    5.006009 10.100.10.100 b^F^R 10.200.10.200 ICMP 98 Echo (ping) request id=0x262f
```

Verify capture

- Using the packet capture the show forward can be executed
- Always execute it on the switch where the packet ingresses

```
9300_1#sh plat hard fed 1 forward int gi 1/0/1 pcap flash:icmp.pcap num 11 data
Show forward is running in the background. After completion, syslog will be generated.
```

- Once completed a syslog gets generated and results will be available
- Can only run one show forward at a time

```
*Jan 27 10:07:35.009: %SHFWD-6-PACKET_TRACE_DONE: Switch 1 R0/0: fed: Packet Trace
Complete: Execute (show platform hardware fed switch <> forward last summary|detail)
```

Show platform forward results

```
9300_1#sh platform hardware fed switch active forward last summary
```

```
Input Packet Details:
```

```
###[ Ethernet ]###
```

```
dst      = a0:f8:49:10:48:51
src      = 00:0c:29:4d:9e:16
type     = 0x800
```

```
###[ IP ]###
```

```
version  = 4L
ihl       = 5L
tos       = 0x0
len       = 84
id        = 46165
flags     = DF
frag      = 0L
ttl       = 64
proto     = icmp
chksum    = 0x5bfc
src       = 10.100.10.100
dst       = 10.200.10.200
options   = ''
```

First section shows the
exact packet
being send for analysis

Show platform forward results

Next section shows
associated data

Ingress:

Port	: GigabitEthernet1/0/1
Global Port Number	: 1
Local Port Number	: 1
Asic Port Number	: 0
Asic Instance	: 1
Vlan	: 100
Mapped Vlan ID	: 6
STP Instance	: 4
BlockForward	: 0
BlockLearn	: 0
L3 Interface	: 50
IPv4 Routing	: enabled
IPv6 Routing	: enabled
Vrf Id	: 0
Adjacency:	
Station Index	: 185
Destination Index	: 21358
Rewrite Index	: 34
Replication Bit Map	: 0x10 ['coreData']

Ingress port

Vlan and STP info

IPv4/v6 setting.
Vrf-id 0 = GRT

Show platform forward results

Forwarding Decision

```
Adjacency:
  Station Index      : 185
  Destination Index  : 21358
  Rewrite Index      : 34
  Replication Bit Map : 0x10    ['coreData']
Decision:
  Destination Index    : 21358
  Rewrite Index       : 34
  Dest Mod Index     : 0        [IGR_FIXED_DMI_NULL_VALUE]
  CPU Map Index      : 0        [CMI_NULL]
  Forwarding Mode     : 3        [Other or Tunnel]
  Replication Bit Map :         ['coreData']
  Winner             :          L3FWDIPV4 LOOKUP
  Qos Label          : 1
  SGT                 : 0
  DGTID              : 0
```

Destination Index, where to
send

Rewrite Index, how to
rewrite

Forwarding mode, how to
treat the packet

Show platform forward results

```
Egress:
Possible Replication      :
  Port                   : TenGigabitEthernet1/1/7
Output Port Data         :
  Port                   : TenGigabitEthernet1/1/7
    Global Port Number   : 59
    Local Port Number    : 59
    Asic Port Number     : 58
    Asic Instance        : 0
    Unique RI            : 34
    Rewrite Type         : 9      [L3_UNICAST_IPV4]
    Mapped Rewrite Type  : 9      [L3_UNICAST_IPV4]
    Vlan                 : 0
    Mapped Vlan ID       : 0
```

Packet removed before sending out to the wire.

Last section showing Egress Packet

```
Output Packet Details:
  Port                               : TenGigabitEthernet1/1/7
###[ Ethernet ]###
  dst      = 00:50:56:92:48:d8
  src      = a0:f8:49:10:48:66
  type     = 0x800
###[ IP ]###
  version  = 4L
  ihl      = 5L
  tos      = 0x0
  len      = 84
  id       = 46165
  flags    = DF
  frag     = 0L
  ttl      = 63
  proto    = icmp
  chksum   = 0x5cfc
  src      = 10.100.10.100
  dst      = 10.200.10.200
  options  = ''
```

Show platform hardware fed forward detail results

- Flash contains shfwd<>.log with detailed information
- Same information gathered with
show platform hardware fed switch <switch> forward last detail
- Detail info containing raw information regarding lookups

```
9300_1#sh platform hardware fed switch 1 forward last detail
```

```
-----  
Starting IPP capture
```

```
=====
```

```
IppDefaultClientTable[4]
```

```
-----
```

defaultClientLeIndex	1
----------------------	---

```
=====
```

```
IppClientLeAd[1]
```

```
-----
```

LEAD_CLIENT_ANCHORED	0
LEAD_CLIENT_AUTH_BEHAVIOR_TAG	0
LEAD_CLIENT_CLIENT_GROUP	0

Layer 2 Forwarding. Verifying STP state

- Show spanning tree gives IOSd view of Spanning Tree
- Hardware forwarding states can be checked *per switch* on FED layer
- Outputs will show what interface are in forwarding state and if traffic will be tagged or untagged
- Flood list indicates what Ports will receive flooded traffic on this switch

```
9300_1#show platform hardware fed switch 1 vlan 100 egress
VLAN STP State in hardware
vlan id is:: 100
Interfaces in forwarding state: : Te1/1/7(Untagged), Gi1/0/1(Untagged)
9300_1#show platform hardware fed switch 1 vlan 100 ingress
VLAN STP State in hardware
vlan id is:: 100
Interfaces in forwarding state: : Te1/1/7(Untagged), Gi1/0/1(Untagged)
flood list: : Te1/1/7, Gi1/0/1
```

Layer 2 Forwarding, IOSd mac address tables

```
9300_1#sh mac address-table vlan 100
      Mac Address Table
-----
Vlan    Mac Address      Type        Ports
----    -
  100    000c.294d.9e16    DYNAMIC     Gi1/0/1
  100    0050.5692.adb3    DYNAMIC     Gi1/0/1
  100    0050.5692.e9aa    DYNAMIC     Gi1/0/1
  100    a0f8.4910.4851    STATIC      V1100
Total Mac Addresses for this criterion: 4
```

- Show mac address table contains a system wide mac table on IOSd
- Types can be static, dynamic, drop
- Mac Address of SVI interfaces also showing in mac address table

FED MATM Mac Address Table

```
9300_1#sh platform software fed switch 1 matm macTable vlan 100
```

VLAN	MAC	Type	Seq#	EC_Bi	Flags	machandle	siHandle	diHandle	*a_time	*e_time	ports
100	a0f8.4910.4851	0x8002	0	99817	64	0x7f91986dfcd8	0x7f9198dad78	0x0	0	0	Vlan100
100	0050.5692.e9aa	0x1	347	0	0	0x7f9199054668	0x7f9199020798	0x7f91986e4a58	300	234	Gi1/0/1
100	0050.5692.adb3	0x1	352	0	0	0x7f91990144a8	0x7f9199020798	0x7f91986e4a58	300	71	Gi1/0/1
100	000c.294d.9e16	0x1	364	0	0	0x7f919900e9d8	0x7f9199070018	0x7f91986e4a58	300	290	Gi1/0/1

```
*a_time=aging_time(secs) *e_time=total_elapsed_time(secs)
```

```
Type:
```

MAT_DYNAMIC_ADDR	0x1	MAT_STATIC_ADDR	0x2	MAT_CPU_ADDR	0x4	MAT_DISCARD_ADDR	0x8
MAT_ALL_VLANS	0x10	MAT_NO_FORWARD	0x20	MAT_IPMULT_ADDR	0x40	MAT_RESYNC	0x80
MAT_DO_NOT_AGE	0x100	MAT_SECURE_ADDR	0x200	MAT_NO_PORT	0x400	MAT_DROP_ADDR	0x800
MAT_DUP_ADDR	0x1000	MAT_NULL_DESTINATION	0x2000	MAT_DOT1X_ADDR	0x4000	MAT_ROUTER_ADDR	0x8000
MAT_WIRELESS_ADDR	0x10000	MAT_SECURE_CFG_ADDR	0x20000	MAT_OPQ_DATA_PRESENT	0x40000	MAT_WIRED_TUNNEL_ADDR	0x80000
MAT_DLR_ADDR	0x100000	MAT_MRP_ADDR	0x200000	MAT_MSRRP_ADDR	0x400000	MAT_LISP_LOCAL_ADDR	0x800000
MAT_LISP_REMOTE_ADDR	0x1000000	MAT_VPLS_ADDR	0x2000000				

- Every FED has its own Mac address table.
- Type Field indicates the type of mac address using a bitmap
- Sequence number of an entry changing would indicated relearning

Layer 3 Forwarding. Routing protocols

```
9300_1#ping 10.48.91.151
Sending 5, 100-byte ICMP Echos to 10.48.91.151, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms
9300_1#sh ip route 10.48.91.151
Routing entry for 10.48.91.128/25
  Known via "isis", distance 115, metric 40, type level-2
  Redistributing via isis
  Last update from 172.31.250.30 on TenGigabitEthernet1/1/6, 6d22h ago
  Routing Descriptor Blocks:
    * 172.31.250.30, from 172.31.255.21, 6d22h ago, via TenGigabitEthernet1/1/6
      Route metric is 40, traffic share count is 1
9300_1#sh ip arp 172.31.250.30
```

Protocol	Address	Age (min)	Hardware Addr	Type	Interface
Internet	172.31.250.30	56	2c4f.523b.c142	ARPA	TenGigabitEthernet1/1/6

- Check Routing Table for correct next hop (Routing Information Base)
- ARP table display rewrite information for next hop (destination mac)

Cisco Express Forwarding (the FIB)

```
9300_1#sh ip cef 10.48.91.128/25 internal
10.48.91.128/25, epoch 7, RIB[I], refcnt 6, per-destination sharing
sources: RIB
feature space:
  IPRM: 0x00028000
  Broker: linked, distributed at 4th priority
ifnums:
  TenGigabitEthernet1/1/6(85): 172.31.250.30
path list 7F3B3265DE78, 139 locks, per-destination, flags 0x4D [shble, hvsh, rif, hwn]
path 7F3B32181A60, share 1/1, type attached nexthop, for IPv4
  nexthop 172.31.250.30 TenGigabitEthernet1/1/6, IP adj out of TenGigabitEthernet1/1/6,
  addr 172.31.250.30 7F3B33B02738
output chain:
  IP adj out of TenGigabitEthernet1/1/6, addr 172.31.250.30 7F3B33B02738
```

Source of route, RIB

Next hop

```
9300_1#sh adjacency 172.31.250.30 detail
Protocol Interface Address
IP      TenGigabitEthernet1/1/6 172.31.250.30 (89)
0 packets, 0 bytes
epoch 0
sourced in sev-epoch 11
Encap length 14
2C4F523BC142A0F8491048500800
L2 destination address byte offset 0
L2 destination address byte length 6
Link-type after encap: ip
ARP
```

Adjacency -> rewrite info

Platform CEF tables (RP)

```
9300_1#sh platform software ip switch ac R0 cef prefix 10.48.91.128/25
```

Forwarding Table

Prefix/Len	Next Object	Index
------------	-------------	-------

10.48.91.128/25	OBJ_ADJACENCY	0x24
-----------------	---------------	------

```
9300_1#sh platform software adjacency switch active R0 index 0x24
```

Number of adjacency objects: 17

Adjacency id: 0x24 (36)

Interface: **TenGigabitEthernet1/1/6**, IF index: 85, Link Type: MCP_LINK_IP

Encap: **2c:4f:52:3b:c1:42:a0:f8:49:10:48:50:8:0**

Encap Length: 14, Encap Type: MCP_ET_ARPA, MTU: 9100

Flags: no-l3-inject

Incomplete behavior type: None

Fixup: unknown

Fixup_Flags_2: unknown

Nexthop addr: **172.31.250.30**

IP FRR MCP_ADJ_IPFRR_NONE 0

OM handle: 0x348066bc48

Prefix points to
Adjacency 0x24

Similar output should be present on standby RP and the FP processes

FED Routing tables

```
9300_1#sh platform software fed switch 1 ip route 10.48.91.128/25
vrf      dest          htm          flags      SGT      DGID MPLS  Last-modified
---      -
0        10.48.91.128/25 0x7f9199010f78 0x0        0        0          2020/01/14 13:49:42.054
FIB: prefix_hdl:0x5a00026d, mpls_ecr_prefix_hdl:0
===== OCE chain =====
ADJ:objid:36 {link_type:IP ifnum:0x55, si:0x9b00003d, IPv4: 172.31.250.30 }
=====
MPLS info: mpls_ecr_scale_prefix_adj:0, mpls_lspa_hdl:0
9300_1#sh platform software fed switch 1 ip adj | inc dest|--|172.31.250.30
dest      if_name      dst_mac      si_hdl      ri_hdl      pd_flags adj_id  Last-modified
---      -
172.31.250.30 Te1/1/6      2c4f.523b.c142 0x7f9198 0x7f9198 0x0          0x24      14:09:12.058
```

- FED layer has its own copy of the IP routing table and rewrite information
- In a stacked environment every switch has its own FED process.
Important to check Ingress and Egress switch
- FED programs TCAM to facilitate forwarding
- Every VRF has its own unique number, 0 is Global Routing Table

Thank you



Possibilities

#CiscoLive