



Possibilities

#CiscoLive

Cisco Live

EVPN Easy Troubleshoot

Diego Zorrilla, XR Tech Lead
DGTL-TSCSPG-603

CISCO *Live!*

#CiscoLive





Agenda

- EVPN Basic Recap.
- Requeriments for Troubleshooting EVPN
- Case Study 1 - Broken Connectivity
- Case Study 2 - Multicast Duplication

EVPN Basic Recap.

EVPN Basic Recap.

- **EVPN Instance (EVI)**

- Its our VPN

- **Ethernet Segment**


- Null
- Auto or Set

- **Topology**

- Single Home
- Dual Home A/S
- Dual Home A/A

- **BGP**

- AFI = L2VPN (25) EVPN (70)
- Route Type = 1 to 5
- Route Attributes = RD + ESI + MAC + Label + RT



Troubleshooting EVPN

Minimum Requirements for Troubleshooting

• Topology

- Single Home
- Dual Home Active / Standby
- Dual Home Active / Active

• MACs + IPs


- Source & Destination

• Flow Type

- Unicast
- BUM (Multicast / Broadcast / Unknown Unicast)

• EVI & ESI

• DF & Non-DF




Case Study 1

Broken Connectivity between hosts in same DC

Broken Connectivity between hosts in same DC

Problem Description

- Service Provider A had an up and running EVPN network.
- SP had a MW where services were migrated to an existing L2VPN but with new VLAN and everything was left working.
- Next day SP saw that the service that was migrated is not working anymore.



Broken Connectivity between hosts in same DC

Troubleshooting


- Are we learning the MACs in L2VPN BD?
- Are we learning the MACs in the right AC?
- Who is the Designated Forwarder (DF)?
- What PE is receiving the ARP request?

AC and L2VPN config for PE-1 and PE-2

```
!  
interface Bundle-Ether2.1000 l2transport  
encapsulation dot1q 1000  
rewrite ingress tag pop 1 symmetric  
!  
interface Bundle-Ether2.3000 l2transport  
encapsulation dot1q 3000  
rewrite ingress tag pop 1 symmetric  
!  
l2vpn  
bridge group EVPN_BG_1234  
bridge-domain EVPN_BD_1234  
interface Bundle-Ether2.1000  
!  
interface Bundle-Ether2.3000  
!  
evi 1234  
!
```

AC and L2VPN config for PE-3

```
!  
interface Bundle-Ether1.2000 l2transport  
encapsulation dot1q 2000  
rewrite ingress tag pop 1 symmetric  
!  
!  
l2vpn  
bridge group EVPN_BG_1234  
bridge-domain EVPN_BD_1234  
interface Bundle-Ether1.2000  
!!  
evi 1234  
!
```



Case Study 1 : Broken Connectivity between hosts in same DC

Troubleshooting Outputs

CE-2

show arp vrf 2000

Protocol	Address	Age (min)	Hardware Addr	Type	Interface	
Internet	10.20.30.10	6	aaaa.aaaa.1000	ARPA	Port-channel1.2000.	<<<<< CE-1
Internet	10.20.30.30	6	cccc.cccc.3000	ARPA	Port-channel1.2000.	<<<<< CE-3

ping vrf 2000 10.20.30.10

<<< From CE-2 to CE-1

Sending 5, 100-byte ICMP Echos to 10.20.30.10, timeout is 2 seconds:
!!!!

ping vrf 2000 10.20.30.30

<<< From CE-2 to CE-3

Sending 5, 100-byte ICMP Echos to 10.20.30.30, timeout is 2 seconds:
!!!!

XR_PE-3


show l2vpn forwarding bridge-domain EVPN_BG_1234:EVPN_BD_1234 mac-address location 0/0/cpu0

Mac Address	Type	Learned from/Filtered on
aaaa.aaaa.1000	EVPN	BD id: 1
cccc.cccc.3000	EVPN	BD id: 1
b0b0.b0b0.2000	dynamic	BE1.2000

Mac Address	Type	Learned from/Filtered on
aaaa.aaaa.1000	EVPN	BD id: 1
cccc.cccc.3000	EVPN	BD id: 1
b0b0.b0b0.2000	dynamic	BE1.2000

show evpn evi vpn-id 1234 mac

VPN-ID	Encap	MAC address	IP address	NextHop	Label
123	MPLS	aaaa.aaaa.1000	::	192.168.0.1	24030
123	MPLS	aaaa.aaaa.1000	::	192.168.0.2	24030
123	MPLS	b0b0.b0b0.2000	::	Bundle-Ether1.2000	24020
123	MPLS	cccc.cccc.3000	::	192.168.0.1	24030
123	MPLS	cccc.cccc.3000	::	192.168.0.2	24030



Case Study 1 : Broken Connectivity between hosts in same DC

Troubleshooting Outputs

```
CE-1
show arp vrf 1000
Protocol Address      Age (min) Hardware Addr  Type Interface
Internet 10.20.30.10         -        aaaa.aaaa.1000  ARPA  Vlan1000
Internet 10.20.30.20         0        b0b0.b0b0.2000  ARPA  Vlan1000
Internet 10.20.30.30         0        Incomplete      ARPA  <<<<<<<<<< CE-3
```

```
ping vrf 1000 10.20.30.20      <<< From CE-1 to CE-2
Sending 5, 100-byte ICMP Echos to 10.20.30.20, timeout is 2 seconds:
!!!!
```


```
ping vrf 1000 10.20.30.30      <<< Connectivity Failure Between CE-1 and CE-3
Sending 5, 100-byte ICMP Echos to 10.20.30.30, timeout is 2 seconds:
.....
```

```
XR_PE-1
show l2vpn forwarding bridge-domain EVPN_BG_1234:EVPN_BD_1234 mac-address location 0/0/cpu0
```

Mac Address	Type	Learned from/Filtered on
b0b0.b0b0.2000	EVPN	BD id: 2
aaaa.aaaa.1000	dynamic	BE2.1000
cccc.cccc.3000	dynamic	BE2.3000 <<< MAC from CE-3 in PE-1

```
show evpn evi vpn-id 1234 mac
```

VPN-ID	Encap	MAC address	IP address	Nexthop	Label
123	MPLS	aaaa.aaaa.1000	::	Bundle-Ether2.1000	24030
123	MPLS	b0b0.b0b0.2000	::	192.168.0.5	24020
123	MPLS	cccc.cccc.3000	::	Bundle-Ether2.3000	24030 <<< MAC from CE-3 in PE-1



Case Study 1 : Broken Connectivity between hosts in same DC

Troubleshooting Outputs

```
CE-1
show arp vrf 1000
Internet 10.20.30.30      0 Incomplete  ARPA      <<<<< CE-3
```

```
ping vrf 1000 10.20.30.30
Sending 5, 100-byte ICMP Echos to 10.20.30.30, timeout is 2 seconds:
.....
Success rate is 0 percent (0/5)
```

16001 = Transport Label for PE1
24034 = Multicast Label for EVI 1234 in PE1
24039 = Split Horizon Label for ESI BE2 on PE1

XR_PE-2

```
show mpls forwarding prefix 192.168.0.1/32
```

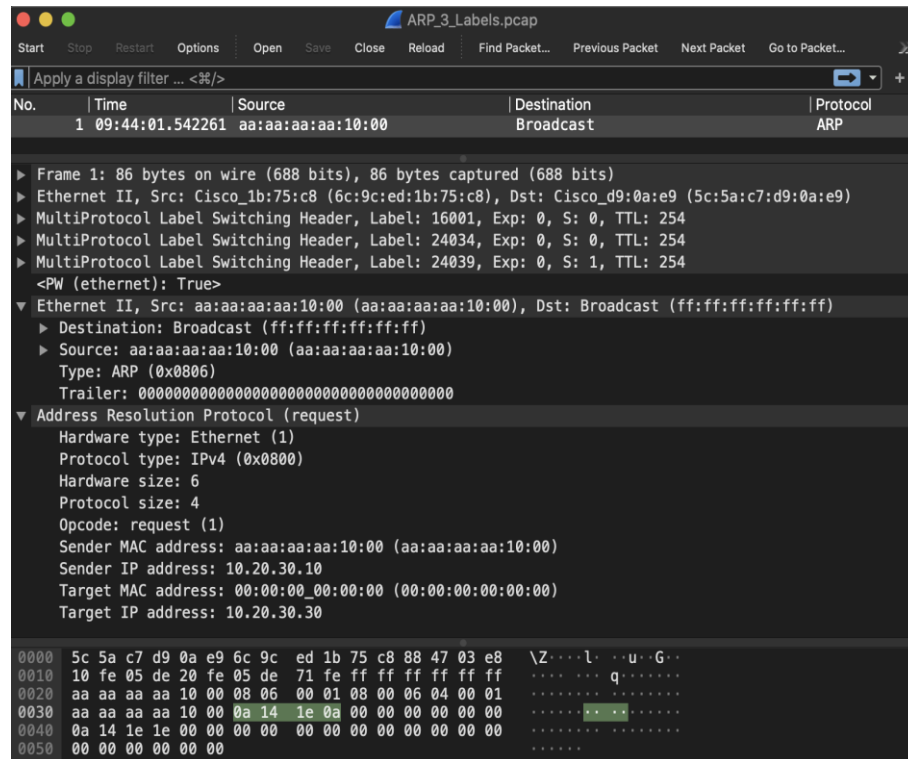
```
16001 16001    SR Pfx (idx 1)    Gi0/0/0/0
```

```
show evpn evi vpn-id 1234 inclusive-multicast detail
```

```
1234  MPLS  0    192.168.0.1
NextHop: 192.168.0.1
Label : 24034
```

```
show evpn ethernet-segment interface bundle-Ether 2 carving detail
```

```
Remote SHG labels : 1
24039 : nexthop 192.168.0.1
```



Case Study 1 : Broken Connectivity between hosts in same DC

Troubleshooting Outputs

XR_PE-1

show evpn ethernet-segment interface bundle-Ether 2 carving detail

Ethernet Segment Id	Interface	Nexthops
0011.2211.2211.2211	BE2	192.168.0.1

192.168.0.2

Main port :
Interface name : Bundle-Ether2

State : Up

Redundancy : Active

ESI type : 0

Value : 11.2211.2211.2211.2211

Topology :
Operational : MH, All-active

Service Carving : Manual

Primary : 1

Pri : 1234

Secondary : 1

Sec : 4000

Service Carving Results:

Forwarders : 2

Elected : 1

EVI E : 1234

Not Elected : 1

EVI NE : 4000

Local SHG label : 24039

Remote SHG labels : 1

24028 : nexthop 192.168.0.2

Access signal mode: Bundle OOS (Default)

XR_PE-2

show evpn ethernet-segment interface bundle-Ether 2 carving detail

Ethernet Segment Id	Interface	Nexthops
0011.2211.2211.2211	BE2	192.168.0.1

192.168.0.2

Main port :
Interface name : Bundle-Ether2

State : Up

Redundancy : Active

ESI type : 0

Value : 11.2211.2211.2211.2211

Topology :
Operational : MH, All-active

Service Carving : Manual

Primary : 1

Pri : 4000

Secondary : 1

Sec : 1234

Service Carving Results:

Forwarders : 2

Elected : 1

EVI E : 4000

Not Elected : 1

EVI NE : 1234

Local SHG label : 24028

Remote SHG labels : 1


24039 : nexthop 192.168.0.1

Access signal mode: Bundle OOS (Default)

Broken Connectivity between hosts in same DC

Root Cause

- ARP Request from VLAN 1000 reach PE-2.
- PE-2 is non-DF for EVI 1234 so traffic between ACs from same BE will not be bridged.
- PE-2 will forward the traffic to PE-1 with Multicast Label and Split Horizon Label.
- When PE-1 receives the packet it will be dropped because of the Split Horizon Label.



Broken Connectivity between hosts in same DC


Solution

1.- Having 2 Bundles between LAN and PEs.

AC-1 for Vlan 1000 in BE-X.

AC-2 for Vlan 2000 in BE-Y.

In this case if BUM traffic from AC-1 in BE-X reaches Non-DF the packet will be forwarded to PE-2 and it will forward to AC-2 since the SHL is for BE-X and not BE-Y.




2.- Using an Active / Standby Solution.

AC for Vlan A and Vlan B in BE1.

ARP request reaches always DF (because of A/S setup).

ARP will be bridged from AC-A to AC-B.




Case Study 2

Duplicate Multicast Traffic being Received

Duplicate Multicast Traffic being Received

Problem Description


- CU X had requirement to run Multicast in existing L2VPN, and because of this, it enhances the network from Single Home to Dual Home.
- During initial test CU sees CC errors on receiver. Deeper analysis shows that CC errors is because of duplicate multicast traffic.



Duplicate Multicast Traffic being Received

Troubleshooting

- Who is the Source?
- What PEs are part of the EVI?
- What are the labels for Multicast and Split Horizon Label for those PEs in the EVI?
- Are we learning the MACs



Case Study 2 : Duplicate Multicast Traffic being Received

Troubleshooting Outputs

```
PE-1 = 192.168.0.1
PE-2 = 192.168.0.2
!
interface Bundle-Ether2.4000 l2transport
 encapsulation dot1q 4000
 rewrite ingress tag pop 1 symmetric
!
l2vpn
 bridge group EVI_4000
 bridge-domain EVI_4000
 interface Bundle-Ether2.4000
!
 evi 4000
!
evpn
 evi 4000
 control-word-disable
 advertise-mac
!
interface Bundle-Ether2
 ethernet-segment
 identifier type 0 11.22.11.22.11.22.11
!
```

```
CE-1
ping 232.11.22.33 source 192.168.40.2 repeat 3333 tim 0
Sending 3333, 100-byte ICMP Echos to 232.11.22.33
.....

XR_PE-2
show interface bundle-ether 2.4000 <<< AC to CE-1
Bundle-Ether2.4000 is up, line protocol is up
    3334 packets input, 393370 bytes

show interface gig 0/0/0/0 <<< CORE
GigabitEthernet0/0/0/1 is up, line protocol is up
    10020 packets output, 1375179 bytes,

show mpls forwarding
```

Local Label	Outgoing Label	Prefix or ID	Outgoing Interface	Next Hop	Bytes Switched
16001	16001	SR Pfx (idx 1)	Gi0/0/0/0	20.20.20.5	383670
16008	16008	SR Pfx (idx 8)	Gi0/0/0/0	20.20.20.5	383670
16009	16009	SR Pfx (idx 9)	Gi0/0/0/0	20.20.20.5	383670

```
XR_PE-2
show evpn evi vpn-id 4000 inclusive-multicast detail
```

VPN-ID	Encap	EtherTag	Originating IP
4000	MPLS	0	192.168.0.1
			Nexthop: 192.168.0.1
			Label : 24008
			Source : Remote
4000	MPLS	0	192.168.0.2
			Nexthop: ::
			Label : 24018
			Source : Local
4000	MPLS	0	192.168.0.8
			Nexthop: 192.168.0.8
			Label : 24122
			Source : Remote
4000	MPLS	0	192.168.0.9
			Nexthop: 192.168.0.9
			Label : 24120
			Source : Remote

Case Study 2 : Duplicate Multicast Traffic being Received

Troubleshooting Outputs


P Router

```

flow monitor-map FMM_EVPN_LABELS
record mpls ipv4-fields
cache entries 1000000
cache timeout active 120
cache timeout rate-limit 2500
!
sampler-map SM_EVPN_1of1
random 1 out-of 1
!
interface GigabitEthernet0/0/0/1
description CONNECT_TO_PE2
cdp
flow mpls monitor FMM_EVPN_LABELS sampler SM_EVPN_1of1 ingress
    
```

show flow monitor FMM_EVPN_LABELS cache format table location 0/0/cpu0

Prefix/Length	Label1-EXP-S	Label2-EXP-S	Label3-EXP-S	InputInterface	OutputInterface	ForwardStatus	ByteCount	PacketCount	
192.168.0.8/32	16008-0-0	24122-0-1	-	Gi0/0/0/1	Gi0/0/0/3	Fwd	406706	3334	<<< From PE-2 to PE-3
192.168.0.9/32	16009-0-0	24120-0-1	-	Gi0/0/0/1	Gi0/0/0/4	Fwd	406706	3334	<<< From PE-2 to PE-4
192.168.0.1/32	16001-0-0	24008-0-0	24039-0-1	Gi0/0/0/1	Gi0/0/0/0	Fwd	406706	3334	<<< From PE-2 to PE-1



Case Study 2 : Duplicate Multicast Traffic being Received

Troubleshooting Outputs

```
XR_PE-1
GigabitEthernet0/0/0/1    <<< CORE Interface
3349 packets input, 454601 bytes, 1 total input drops

RSV_EGR_L2_EVPN_ESI_DROP    3336

Bundle-Ether2.4000    <<< AC to CE-1
0 packets output, 0 bytes


XR_PE-3
GigabitEthernet0/0/0/0    <<< CORE Interface
3394 packets input, 444923 bytes

Bundle-Ether1.4000    <<< AC to CE-2
3337 packets output, 393586 bytes

XR_PE-4
GigabitEthernet0/0/0/0    <<< CORE Interface
3394 packets input, 444923 bytes

Bundle-Ether1.4000    <<< AC to CE-2
3337 packets output, 393586 bytes

CE-2
6797 packets input, 795577 bytes <<< Duplicate Multicast Packets
```



Case Study 2 : Duplicate Multicast Traffic being Received

Troubleshooting Outputs

PE-3 = 192.168.0.8

```
!  
l2vpn  
  bridge group EVI_4000  
  bridge-domain EVI_4000  
  interface Bundle-Ether1.4000  
  !  
  evi 4000  
  !  
evpn  
!  
interface Bundle-Ether1  
  ethernet-segment  
  identifier type 0 88.99.88.99.88.99.88.99.88  
!
```

PE-4 = 192.168.0.9

```
!!  
l2vpn  
  bridge group EVI_4000  
  bridge-domain EVI_4000  
  interface Bundle-Ether1.4000  
  !  
  evi 4000  
  !  
evpn  
!  
interface Bundle-Ether1  
  ethernet-segment  
  identifier type 0 89.89.89.89.89.89.89.89.89  
!
```

PE-3

show evpn ethernet-segment interface bundle-Ether 1 carving detail

Ethernet Segment Id	Interface	Nexthops
0088.9988.9988.9988.9988	BE1	192.168.0.8

Main port :
 Interface name : Bundle-Ether1
 Interface MAC : bc5a.561d.8cdf
ESI type : 0
 Value : 88.9988.9988.9988.9988
ES Import RT : 8899.8899.8899 (from ESI)
Topology :
 Operational : SH
 Configured : All-active (AApF) (default)
Service Carving : Auto-selection
Peering Details :
 192.168.0.8 [MOD:P:00]
Service Carving Results:
 Forwarders : 2
 Elected : 2
 EVI E : 1000, 4000
Local SHG label : 24022

PE-4

show evpn ethernet-segment interface bundle-Ether 1 carving detail


Ethernet Segment Id	Interface	Nexthops
0089.8989.8989.8989.8989	BE1	192.168.0.9

Main port :
 Interface name : Bundle-Ether1
 Interface MAC : bc5a.561d.68de
IfHandle : 0x08004024
ESI type : 0
 Value : 89.8989.8989.8989.8989
ES Import RT : 8989.8989.8989 (from ESI)
Topology :
 Operational : SH
 Configured : All-active (AApF) (default)
Service Carving : Auto-selection
Peering Details :
 192.168.0.9 [MOD:P:00]
Service Carving Results:
 Forwarders : 2
 Elected : 2
 EVI E : 1000, 4000
Local SHG label : 24017

Duplicate Multicast Traffic being Received

Root Cause and Solution

- "Remote" PEs will receive the multicast, and only Designated Forwarder (DF) will flood the Multicast.
- If two PEs working in Dual Home Active / Active do not have same ESI, both will be DF for the ESI and both will forward multicast frames to the LAN, duplicated traffic will happen.



Appendix

Swiss Army Knife

EVPN Basic Configuration for Single Home

BGP

```
router bgp 65500
!  
address-family l2vpn evpn
!  
neighbor-group Route_Reflector
address-family l2vpn evpn
!
```

AC

```
interface Bundle-Ether1.1000 l2transport
encapsulation dot1q 1000
rewrite ingress tag pop 1 symmetric
!
```

EVPN

```
evpn
evi 1000
control-word-disable
advertise-mac
!
```

L2VPN

```
l2vpn
bridge group EVPN_BG
bridge-domain EVPN_BD
interface Bundle-Ether1.1000
!  
evi 1000
!
```

EVPN Basic Configuration for Dual Home A/A (1 of 2)

	PE-1	PE-2
BGP	<pre>router bgp 65500 ! address-family l2vpn evpn !</pre>	<pre>router bgp 65500 ! address-family l2vpn evpn !</pre>
AC	<pre>neighbor-group Route_Reflector address-family l2vpn evpn !</pre>	<pre>neighbor-group Route_Reflector address-family l2vpn evpn !</pre>
EVPN	<pre>interface Bundle-Ether1.1000 l2transport encapsulation dot1q 1000 rewrite ingress tag pop 1 symmetric !</pre>	<pre>interface Bundle-Ether1.1000 l2transport encapsulation dot1q 1000 rewrite ingress tag pop 1 symmetric !</pre>
L2VPN	<pre>evpn evi 1000 control-word-disable advertise-mac ! l2vpn bridge group EVPN_BG bridge-domain EVPN_BD interface Bundle-Ether1.1000 ! evi 1000 !</pre>	<pre>evpn evi 1000 control-word-disable advertise-mac ! l2vpn bridge group EVPN_BG bridge-domain EVPN_BD interface Bundle-Ether1.1000 ! evi 1000 !</pre>

EVPN Basic Configuration for Dual Home A/A (2 of 2)

	PE-1	PE-2	
MLACP	redundancy	redundancy	
	iccp	iccp	
	group 12	group 12	<<< Same Values for P-1 and P-2
	mlacp node 1	mlacp node 2	<<< Dif Values in P-1 and P-2
	mlacp system mac 0111.0111.0111	mlacp system mac 0111.0111.0111	<<< Same Values for P-1 and P-2
	mlacp system priority 1	mlacp system priority 1	
	mode singleton	mode singleton	
	backbone	backbone	
	interface xxxx	interface xxxx	<<< CORE Link(s)
	!	!	
EVPN	interface Bundle-Ether1	interface Bundle-Ether1	
	lacp switchover suppress-flaps 300	lacp switchover suppress-flaps 300	
	mlacp iccp-group 12	mlacp iccp-group 12	
	mlacp port-priority 1	mlacp port-priority 1	
	!	!	
	evpn	evpn	
	interface Bundle-Ether1	interface Bundle-Ether1	
	ethernet-segment	ethernet-segment	
	identifier type 0 12.12.12.12.12.12.12.12	identifier type 0 12.12.12.12.12.12.12.12	<<< Same Values for P-1 and P-2
	!	!	

EVPN Optional Configs.

EVPN

Configuration	Action	Default
evpn		
evi 1000		
advertise-mac	Advertise local mac addr in BGP	Default is NOT to advertise.
unknown-unicast-suppression	Avoid flooding of unknown unicast frames through EVPN.	Default is to flood.
control-word-disable	Add CW on packets.	Default is CW enable.
load-balancing flow-label static	Add Extra label to allow load balancing on core links.	Default is no Flow Label.
!		
interface bundle-ether1		
ethernet-segment		
load-balancing-mode single-active	Only one PE handles the traffic, per EVI.	Default is Active / Active.
service-carving manual		Controls DF and Non-DF.
primary 1-10 secondary 11-20	Manual load balancing per EVI	Default is Even on PE-X and Odd on PE-Y
!		
group 1234		
core interface xxxxxx		
Interface yyyyy <<< Int. to LAN		
Core isolation group 1234		

Swiss Army Knife

EVPN:

```
show evpn ethernet-segment
show evpn ethernet-segment esi xxxxx carving detail
show evpn ethernet-segment virtual vfi xxxxx detail
show evpn ethernet-segment virtual neighbor xxxxx pw-id xxxxx detail
show evpn ethernet-segment detail
show evpn ethernet-segment interface ... detail
show evpn evi inclusive-multicast detail
show evpn evi detail
show evpn evi vpn <evi id> mac <mac> detail
show evpn evi mac private
show evpn evi ead
show evpn internal-label
show evpn summary
```

L2VPN:

```
show l2vpn xconnect
show l2vpn xconnect detail
show l2vpn bridge-domain summary
show l2vpn mac-learning mac all location <loc>
show l2vpn forwarding bridge-domain xxxxx evpn inclusive-multicast detail location <loc>
show l2vpn forwarding bridge-domain <BD>:<BG> mac-address location <loc>
show l2vpn bridge-domain bd-name <bd name> detail
show l2vpn forwarding protection main-interface location <loc>
show l2vpn forwarding bridge-domain mac location <loc>
show l2vpn forwarding bridge-domain xxxxx detail location <loc>
show l2vpn forwarding xc <xc id> detail location <loc>
```

BGP:

```
show bgp l2vpn evpn summary
show bgp l2vpn evpn bridge-domain <bd name>
show bgp l2vpn evpn bridge-domain <bd name> rnh
show bgp l2vpn evpn bridge-domain <bd name> <prefix>
show bgp l2vpn evpn route-type <1|2|3|4|5>
show bgp rt l2vpn evpn
show bgp vrf-db table all
show bgp vrf-db table <table id>
show bgp l2vpn evpn ethernet-ad
show bgp l2vpn evpn route-type mac-advertisement
show bgp l2vpn evpn rd ...
```

L2RIB:

```
show l2route evpn mac all
show l2route evpn mac evi <topo id>
show l2route summar
```

DEBUG:

```
debug l2vpn forwarding platform ...
debug l2vpn forwarding bcdl location <> debug l2vpn forwarding mac locat
```

Thank you



Possibilities

#CiscoLive